

AVR Assembly Programming

Version 1.2

Dazed_N_Confused

2026-05-22

Contents

1	Why Assembly? Toolchain Setup	1
1.1	What Is Assembly Language?	1
1.2	Target: ATtiny3217	1
1.3	The AVR-GAS Toolchain	2
1.4	The Build Pipeline	3
1.5	First Program: Blink	3
1.6	Building and Flashing	6
1.7	A Minimal Makefile	7
1.8	What Just Happened — Mental Model	8
1.9	Summary	9
1.10	Exercises	9
2	AVR Architecture Overview	10
2.1	Harvard Architecture	10
2.2	ATtiny3217 Memory Map	11
2.3	The Register File	12
2.4	The Status Register (SREG)	13
2.5	The Stack and Stack Pointer	15
2.6	The Program Counter (PC)	16
2.7	I/O Register Space	17
2.8	Putting It Together — A Register-State Walk-Through	17
2.9	Key Architectural Constraints Summary	18
2.10	Summary	18
2.11	Exercises	19
3	AVR-AS Syntax and Directives	20
3.1	File Extensions: .S vs .s	20
3.2	Comments	21
3.3	Instruction Syntax	21
3.4	Number Literals	21
3.5	Expressions	21
3.6	Labels	22
3.7	Sections	24
3.8	Data Directives	25
3.9	Constant Directives	26
3.10	The .org Directive	27
3.11	Flash Data on AVRxt (ATtiny3217)	27
3.12	The .include Directive	28
3.13	Macro Directives	28
3.14	Full Section Layout Example	29
3.15	Examining the Output with avr-objdump	30

3.16	Quick Reference	30
3.17	Summary	31
3.18	Exercises	31
4	Instruction Set Basics	33
4.1	Instruction Encoding	33
4.2	Instruction Categories	34
4.3	Data Transfer — Moving Values Around	35
4.4	Arithmetic Instructions	36
4.5	Logic Instructions	39
4.6	Comparison Instructions	40
4.7	Conditional Branches — Full Table	41
4.8	Unsigned vs Signed Comparisons	41
4.9	Flag Effects Summary	42
4.10	Worked Example: Compare, Clamp, Negate	42
4.11	Operand Restrictions Cheat-Sheet	43
4.12	Summary	44
4.13	Exercises	44
5	Load and Store	45
5.1	The Memory-Register Model	45
5.2	Addressing Modes	45
5.3	Direct Addressing: LDS and STS	46
5.4	Pointer Registers: X, Y, Z	47
5.5	Indirect Addressing: LD and ST	47
5.6	Displacement Addressing: LDD and STD	48
5.7	Cycle and Size Comparison	49
5.8	Reading Flash: LPM	49
5.9	Reading Flash: Mapped Access (AVRxt Only)	50
5.10	Pointer Arithmetic	51
5.11	Atomicity on ATtiny3217	52
5.12	Full Addressing Mode Reference	52
5.13	Worked Example: Reverse a String In-Place	53
5.14	Summary	54
5.15	Exercises	55
6	Arithmetic and Logic: 16-bit and Beyond	56
6.1	Multi-Byte Addition and Subtraction	56
6.2	Shift and Rotate Instructions	58
6.3	SWAP — Swap Nibbles	59
6.4	MUL — Unsigned 8×8 Multiply	59
6.5	MULS — Signed 8×8 Multiply	60
6.6	MULSU — Signed × Unsigned Multiply	61
6.7	Multiply Summary	61
6.8	Division — Software Algorithm	61
6.9	Worked Example: 8×8 Unsigned Multiply to 16-bit Result	64
6.10	Signed Arithmetic Pitfalls	64
6.11	Useful Idioms	65
6.12	Summary	66
6.13	Exercises	66
7	Branches and Control Flow	67
7.1	Unconditional Jumps	67
7.2	Conditional Branches — Mechanics	68
7.3	Complete Branch Instruction Table	68

7.4	Skip Instructions	69
7.5	Structured Control Flow Patterns	71
7.6	Worked Example: Sum an Array	73
7.7	Jump Tables (switch/case)	74
7.8	Polling Loops	75
7.9	Branch Optimisation Notes	76
7.10	Summary	77
7.11	Exercises	77
8	Subroutines and the Stack	79
8.1	CALL and RCALL	79
8.2	RET and RETI	80
8.3	PUSH and POP	80
8.4	The GCC AVR ABI	81
8.5	Prologue and Epilogue Patterns	83
8.6	Stack Depth Analysis	84
8.7	Worked Example: Recursive Factorial	85
8.8	Tail Call Optimisation	86
8.9	Complete Function Template	87
8.10	Summary	88
8.11	Exercises	88
9	GPIO	90
9.1	GPIO Fundamentals	90
9.2	Classic DDRx / PORTx / PINx (ATmega-Style)	91
9.3	ATtiny3217 Port Architecture	91
9.4	IN — Read I/O Register	93
9.5	OUT — Write I/O Register	93
9.6	SBI and CBI — Set / Clear Bit in I/O Register	94
9.7	SBIC and SBIS — Skip if Bit Clear / Set	94
9.8	Configuring a GPIO Output	95
9.9	Configuring a GPIO Input with Pull-Up	96
9.10	Debounce	96
9.11	Worked Example: Button-Controlled LED	97
9.12	Summary	99
9.13	Exercises	100
10	Interrupts	101
10.1	Why Interrupts?	101
10.2	The Interrupt Mechanism	101
10.3	SEI and CLI	102
10.4	RETI — Return from Interrupt	103
10.5	The Interrupt Vector Table	104
10.6	ISR Prologue and Epilogue	105
10.7	Configuring Pin Interrupts on ATtiny3217	106
10.8	ISR Latency and Timing	107
10.9	Worked Example: Pin Interrupt Toggles LED	108
10.10	Summary	110
10.11	Exercises	111
11	Timer/Counter	112
11.1	Timer Fundamentals	112
11.2	TCA0 on ATtiny3217	113
11.3	Waveform Generation Modes	115
11.4	Computing Timer Values	116

11.5 TCA0 Interrupt Vectors	116
11.6 The ISR	117
11.7 Main: Timer Initialization Sequence	118
11.8 Worked Example: 1 Hz Blink	119
11.9 Extending: Normal Mode with OVF Interrupt	120
11.10 Extending: Single-Slope PWM	121
11.11 Summary	121
11.12 Exercises	122
12 USART (Serial Communication)	123
12.1 UART Framing	123
12.2 USART0 Registers on ATtiny3217	124
12.3 Baud Rate Calculation	126
12.4 TX/RX Pins on ATtiny3217	126
12.5 Polling Transmit	127
12.6 Polling Receive	128
12.7 ISR-Driven Receive	129
12.8 Initialization Sequence	130
12.9 Worked Example: Echo Terminal	131
12.10 Circular Buffer for High-Throughput RX	132
12.11 Summary	133
12.12 Exercises	134
13 SPI & TWI (I2C) Overview	135
13.1 SPI Fundamentals	135
13.2 SPI0 Registers on the ATtiny3217	136
13.3 SPI Initialisation (Master Mode)	137
13.4 Transferring a Byte	138
13.5 The 25LC010A SPI EEPROM	138
13.6 Example: Write Byte to SPI EEPROM	140
13.7 TWI (I ² C) Fundamentals	141
13.8 TWI0 Master Mode Registers on ATtiny3217	142
13.9 TWI Initialisation	144
13.10 Master Write Transaction (Polling)	144
13.11 Example: Write to PCF8574 I ² C I/O Expander	146
13.12 SPI vs TWI Comparison	146
13.13 Summary	147
13.14 Exercises	147
14 Mixing C and Assembly	149
14.1 Why Mix C and Assembly?	149
14.2 The GCC AVR Calling Convention	150
14.3 Writing an Assembly Subroutine for C	151
14.4 Modern AVR Header Conventions	152
14.5 Example: CRC-8 Subroutine	152
14.6 Calling the Assembly Routine from C	154
14.7 Cycle Count Comparison	155
14.8 Inline Assembly: <code>asm volatile</code>	155
14.9 Inline Assembly Examples	156
14.10 Memory Clobber	157
14.11 Building Mixed Projects	158
14.12 Common Mistakes	159
14.13 Summary	160
14.14 Exercises	160

15 Optimisation Techniques	162
15.1 Cycle Counting with avr-objdump	162
15.2 Avoid SRAM — Keep Variables in Registers	163
15.3 Loop Unrolling	165
15.4 Lookup Tables in Flash (PROGMEM)	167
15.5 Putting It Together: Cycle Comparison	169
15.6 Using avr-objdump Effectively	169
15.7 Common Pitfalls	170
15.8 Summary	171
15.9 Exercises	171
16 Self-Programming Flash: NVMCTRL on tinyAVR 1-Series	173
16.1 tinyAVR 1-Series vs Classic AVR: No SPM, NVMCTRL Instead	173
16.2 NVMCTRL Overview	174
16.3 The Page Buffer and Write Sequence	176
16.4 CCP — Configuration Change Protection	177
16.5 Flash Write Routine in Assembly	178
16.6 Practical Firmware Update: UART-Driven Self-Update	181
16.7 BOOTEND Fuse: Protecting a Bootloader Region	184
16.8 Building and Flashing	185
16.9 Common Pitfalls	187
16.10 Summary	188
16.11 Exercises	189
17 Bit Manipulation & Fixed-Point Math	191
17.1 Shift Operations	191
17.2 Shift-and-Add Multiplication	192
17.3 Fixed-Point Q8.8 Arithmetic	195
17.4 BCD Arithmetic	197
17.5 Binary to BCD Conversion	198
17.6 Complete Example: Sensor Reading to 7-Segment Display	200
17.7 Instruction Reference for This Chapter	202
17.8 Common Pitfalls	203
17.9 Summary	204
17.10 Exercises	205
18 Real-Time Patterns	206
18.1 Real-Time Fundamentals	206
18.2 Cooperative Task Scheduler Design	207
18.3 TCB0 Configuration — 1 ms Tick	208
18.4 ATtiny3217 Interrupt Vector Table	210
18.5 ISR Prologue and Epilogue	211
18.6 ISR Latency Analysis	212
18.7 Complete Two-Task Scheduler	213
18.8 Timing Analysis — Worst-Case Response	217
18.9 Non-Blocking Task Extension	218
18.10 Scaling to More Tasks	219
18.11 Summary	219
19 CORDIC on AVR	221
19.1 Why Use CORDIC on AVR?	221
19.2 The Core Idea	222
19.3 Fixed-Point Choices for AVR	222
19.4 The Angle Table	223
19.5 Rotation Mode for Sine and Cosine	224

19.6 Complete Integer Reference Implementation	225
19.7 How This Maps to AVR Assembly	226
19.8 Vectoring Mode: Magnitude and Angle	228
19.9 Example 1: DDS / NCO Sine Generator	228
19.10Example 2: Tilt Angle from Two Axes	229
19.11Accuracy, Flash, and Cycle Tradeoffs	229
19.12Common Pitfalls	230
19.13Summary	230
19.14Exercises	231
20 ATtiny3217 Register Reference	232
20.1 CPU Registers	232
20.2 CLKCTRL — Clock Controller	233
20.3 GPIO Registers	234
20.4 Curiosity Nano Pin Assignments	237
20.5 TCA0 — Timer/Counter Type A (16-bit)	237
20.6 TCB0 / TCB1 — Timer/Counter Type B (16-bit)	239
20.7 USART0	240
20.8 SPI0	241
20.9 TWI0 — Two-Wire Interface (I2C)	242
20.10ADC0	243
20.11NVMCTRL — Non-Volatile Memory Controller	245
20.12Interrupt Vector Table	246
20.13Fuses	246
21 AVR Instruction Set Summary	248
21.1 Arithmetic and Logic	248
21.2 Shift and Rotate	249
21.3 Bit Operations	249
21.4 Compare and Test	249
21.5 Branch Instructions	250
21.6 Load and Store	251
21.7 I/O Instructions	252
21.8 Miscellaneous	252
21.9 Instruction Encoding Summary	252
21.10GCC ABI Register Conventions	252
22 GAS Directive Reference	254
22.1 Section Directives	254
22.2 Symbol Directives	254
22.3 Data Directives	255
22.4 Alignment Directives	255
22.5 Conditional Assembly	255
22.6 Macro Directives	256
22.7 Include Directives	257
22.8 Listing and Debug Directives	257
22.9 Expression Operators	257
22.10Common Patterns	258
23 Linker Script Walkthrough	259
23.1 Anatomy of a Linker Script	259
23.2 MEMORY Region Flags	260
23.3 Address Spaces on AVR	261
23.4 VMA vs LMA	261
23.5 KEEP()	262

23.6 Bootloader Placement	262
23.7 Parameter Page at Fixed Flash Address	262
23.8 .noinit Section	263
23.9 Useful Linker Symbols	263
23.10 Inspecting Linker Output	264
23.11 Common Linker Errors	264
24 avrdude Flashing Cheat-Sheet	265
24.1 Basic Syntax	265
24.2 Most-Used Commands	266
24.3 Common Programmers	267
24.4 Device Part Numbers	268
24.5 Fuse Byte Reference (ATmega328P)	268
24.6 Generating Intel HEX from ELF	269
24.7 Complete Build + Flash Script	269
24.8 Troubleshooting	270

Chapter 1

Why Assembly? Toolchain Setup

1.1 What Is Assembly Language?

Every program your computer runs is ultimately a sequence of binary numbers: machine code that the CPU decodes and executes. Assembly language is the human-readable form of that machine code. You are no longer asking a compiler to choose the instructions for you; you are choosing them yourself.

Higher-level languages (C, Python, Rust) are *compiled* or *interpreted* — they add layers of abstraction that trade control for convenience. Assembly gives you almost no abstraction: you decide every instruction, every register, and often every cycle count.

1.1.1 When do you actually need assembly?

Situation	Why assembly wins
Tight timing constraints	You control the exact cycle count
Hardware bring-up / no OS	No runtime support — you write everything
Interrupt service routines	Worst-case latency is deterministic
Extremely small flash budget	No compiler overhead or padding
Learning how CPUs work	Forces deep understanding of architecture
Hand-optimising hot paths	Compiler output is a starting point, not the ceiling

In the embedded world, assembly is not just historical trivia. It still shows up in bootloaders, startup code, interrupt paths, cycle-counted loops, and the smallest pieces of firmware where you cannot afford a vague understanding of what the CPU will do next.

1.2 Target: ATtiny3217

This book uses the **ATtiny3217** — an 8-bit AVR from Microchip's tinyAVR 1-series. It is a modern, capable microcontroller with:

Feature	Detail
Architecture	AVRxt (enhanced AVR core)
Flash	32 KB
SRAM	2 KB (at 0x3800–0x3FFF)

Feature	Detail
Clock	Up to 20 MHz internal oscillator
Programming	UPDI (1-wire, replaces ISP)
I/O	New-style PORT peripheral (not classic DDRx/PORTx)
Packages	24-pin VQFN, SOIC, SSOP

The **ATTiny3217 Curiosity Nano** development board is a USB-powered board with an onboard debugger/programmer. It has one yellow LED on **PA3** and exposes all device pins. This book uses this board for all hardware examples.

That choice is deliberate. The chip is modern enough to show the newer AVR peripheral model, but small enough that the architecture is still easy to hold in your head.

Note: All concepts carry over to other tinyAVR 1-series devices (ATTiny3216, ATTiny1617, etc.) and with small adjustments to any AVR.

1.3 The AVR-GAS Toolchain

We use the **GNU Assembler** (`avr-as`), part of the `avr-gcc` toolchain. This is the same assembler `avr-gcc` uses under the hood when it compiles C for AVR. Using it directly keeps the toolchain familiar while removing the compiler as a middleman.

The toolchain has four main tools:

```
avr-as      Assembler: .S source → .o object file
avr-ld      Linker:    .o object(s) → .elf executable
avr-objcopy Converter: .elf → .hex (Intel HEX for flashing)
avr-objdump Inspector: disassemble, inspect sections, count cycles
```

You do not need to memorise all four on day one. The practical workflow is: assemble source into an object file, link it into an executable image, convert that image into something the programmer can flash, then inspect the result when you want to verify what really happened.

1.3.1 Installing the Toolchain

Debian / Ubuntu / Kali:

```
sudo apt install gcc-avr binutils-avr avr-libc avrdude
```

`avr-libc` is needed for the device header files (`<avr/io.h>`) that define register names and addresses.

Arch Linux:

```
sudo pacman -S avr-gcc avr-binutils avr-libc avrdude
```

macOS (Homebrew):

```
brew tap osx-cross/avr
brew install avr-gcc avrdude
```

Verify installation:

```
avr-as --version
avr-ld --version
avr-objcopy --version
```

Expected output (version numbers may differ):

GNU assembler (GNU Binutils) 2.42
 GNU ld (GNU Binutils) 2.42
 GNU objcopy (GNU Binutils) 2.42

1.4 The Build Pipeline

Source goes through four stages before reaching the MCU:

```

blink.S
|
|   avr-as -mmcu=attiny3217 -o blink.o blink.S
|   ▼
blink.o      ← relocatable object (ELF, machine code + symbol table)
|
|   avr-gcc -mmcu=attiny3217 -nostartfiles -o blink.elf blink.o
|   ▼
blink.elf    ← linked executable (ELF, absolute addresses assigned)
|
|   avr-objcopy -O ihex blink.elf blink.hex
|   ▼
blink.hex    ← Intel HEX – what avrdude writes to flash
|
|   pymcuprog write -d attiny3217 -t uart -f blink.hex
|   ▼
ATTiny3217 flash ← running program
  
```

If you are new to embedded builds, this pipeline is worth pausing over. The source file is not flashed directly. It is first turned into machine code, then assigned final addresses, then repackaged into a format the programmer can write into the device's flash memory.

1.4.1 Why two link steps?

We use `avr-gcc` as the linker driver rather than calling `avr-ld` directly because it already knows the right device-specific details: the memory layout, the default linker script, and the expected section placement. We could drive the linker manually, but that adds complexity before it teaches us anything useful.

The `-nostartfiles` flag matters because we are not building a normal C program. We do not want the C runtime to insert its own startup sequence; we want to write the reset path ourselves so we can see exactly how execution begins.

1.5 First Program: Blink

The classic “hello world” of embedded is not printing text. It is making a pin change state in a visible, repeatable way. Here, that means toggling the board LED forever.

Hardware: ATtiny3217 Curiosity Nano, LED0 on PA3.

File: `ch01_intro/src/blink.S`

1.5.1 Understanding the ATtiny3217 PORT Peripheral

The ATtiny3217 uses a **new-style** PORT peripheral, unlike older AVR devices (ATmega328P etc.) which used DDRx, PORTx, PINx registers in low I/O space.

On ATtiny3217, each port has a block of memory-mapped registers. PORTA lives at base address **0x0400**:

Offset	Register	Function
0x00	PORTA_DIR	Direction: 1 = output, 0 = input
0x01	PORTA_DIRSET	Write 1 to set pin as output (non-destructive)
0x02	PORTA_DIRCLR	Write 1 to set pin as input (non-destructive)
0x03	PORTA_DIRTGL	Write 1 to toggle direction
0x04	PORTA_OUT	Output value
0x05	PORTA_OUTSET	Write 1 to drive pin HIGH (non-destructive)
0x06	PORTA_OUTCLR	Write 1 to drive pin LOW (non-destructive)
0x07	PORTA_OUTTGL	Write 1 to toggle pin (atomic!)
0x08	PORTA_IN	Read current pin state

The SET/CLR/TGL variants are key: they let you change one pin without reading the register first, eliminating the classic read-modify-write race condition.

That is one of the nicest features of the newer AVR GPIO design. Instead of reading the whole port, modifying one bit in software, and writing the whole byte back, the peripheral can perform the bit change for you.

Because these registers are above address 0x3F, we cannot use IN/OUT instructions (which only reach 0x00–0x3F). We use STS (store to SRAM address) instead. The header <avr/io.h> defines all register names for us.

So even in this tiny program, the hardware layout is already shaping the instruction choices. We are not using STS by style preference; we are using it because that is the instruction family that can reach these addresses.

1.5.2 The Code

```
#include <avr/io.h>

/* — Vector table ————— */
.section .vectors, "ax", @progbits
.global __vectors

__vectors:
    jmp     main                ; reset vector: first instruction after power-on
```

When the ATtiny3217 resets, execution begins at **byte address 0x0000**. The CPU does not “look for main” by name. It simply fetches whatever instruction is stored at the reset vector and executes it. We place a JMP main there so the first thing the chip does is branch into our real program.

That is an important mental shift if you are coming from C: function names are for the assembler and linker, not for the CPU. The CPU only sees addresses and instructions.

```
.section .text

main:
    /* — Stack pointer init ————— */
    ldi     r16, lo8(RAMEND)      ; RAMEND = 0x3FFF → lo8 = 0xFF
    out     SPL, r16
    ldi     r16, hi8(RAMEND)      ; hi8 = 0x3F
    out     SPH, r16
```

SP (the stack pointer) must be valid before any PUSH, CALL, or RCALL. If it is wrong, the CPU will store return addresses and saved registers in the wrong place, and the program will collapse quickly.

On the ATtiny3217, reset hardware already loads `SP = RAMEND`, so these writes are technically redundant in this exact program. We still do them because they make the startup state explicit, and because this is the pattern you need on many older AVR parts where the hardware does not fully set the stack up for you.

`RAMEND` expands to `0x3FFF` (last byte of ATtiny3217 SRAM). `lo8()` and `hi8()` are assembler expressions that extract the low and high bytes of a 16-bit value.

`SPL` and `SPH` are in low I/O space (addresses `0x3D` and `0x3E`), so we reach them with `OUT`.

```
ldi    r16, (1 << 3)      ; bitmask for PA3 = 0x08
sts    PORTA_DIRSET, r16   ; set PA3 as output
```

`LDI` loads a constant into a register. `STS` stores that register value to a memory-mapped peripheral address. `PORTA_DIRSET` expands to `0x0401`, and writing `0x08` there means “set bit 3 of the direction register.”

In plain language: we are telling the chip that PA3 is now an output pin. We are not yet turning the LED on or off; we are only telling the hardware that this pin should be driven by the CPU rather than treated as an input.

`loop:`

```
ldi    r16, (1 << 3)
sts    PORTA_OUTTGL, r16 ; atomically toggle PA3
```

`PORTA_OUTTGL` flips PA3's output state. If it was high, it becomes low; if it was low, it becomes high. The useful part is that the flip happens inside the hardware peripheral. We do not have to read the old state first, and there is no risk of losing a concurrent change while doing a software read-modify-write.

```
ldi    r19, 25
delay_outer:
ldi    r25, 0
ldi    r24, 0
delay_inner:
sbiw   r24, 1          ; subtract 1 from r25:r24 (2 cycles)
brne   delay_inner     ; branch if not zero      (2 cycles)
dec    r19
brne   delay_outer

rjmp   loop
```

`SBIW` subtracts an immediate from a 16-bit register pair. Here, `r25:r24` acts as the inner delay counter. It counts down from 65535 to 0. Each time it reaches zero, `BRNE` stops looping, `DEC r19` reduces the outer counter by one, and the whole process repeats until the outer loop also finishes.

This is not a precise timer. It is a simple busy-wait delay that burns CPU cycles. That makes it ideal for a first example because you can understand it with nothing more than arithmetic and branches.

After reset, the ATtiny3217 does not normally run at the full oscillator rate. It starts from `OSC20M` with a **divide-by-6 prescaler**. That gives **3.3 MHz** if the `FUSE.OSCCFG` frequency select is 20 MHz, or **2.66 MHz** if it is 16 MHz. Assuming the 20 MHz setting, the loop timing is roughly: $25 \times 65536 \times 4 \text{ cycles} \div 3,330,000 \approx 2 \text{ seconds}$ per toggle (about 1 Hz blink). If you later change the clock prescaler, oscillator source, or fuse setting, the blink rate changes proportionally.

That last point is worth remembering early: software delay loops are only as accurate as your clock configuration. Change the clock, and the visible timing changes too.

Instruction introduced: `LDI Rd, K` — Load immediate value `K` into register `Rd`. `Rd` must be `r16–r31`. 1 cycle. Affects no flags.

Instruction introduced: **OUT** **OUT A, Rr** — Write register Rr to I/O address A (0x00–0x3F). 1 cycle. Affects no flags.

Instruction introduced: **STS** **STS k, Rr** — Store register Rr to SRAM address k (0x0000–0xFFFF). 2 cycles. Affects no flags. 32-bit instruction (two 16-bit words).

Instruction introduced: **SBIW** **SBIW Rd, K** — Subtract immediate K (0–63) from register pair Rd+1:Rd. Valid pairs: r25:r24, r27:r26 (X), r29:r28 (Y), r31:r30 (Z). 2 cycles. Affects C, Z, N, V, S flags.

Instruction introduced: **BRNE** **BRNE k** — Branch if Z flag is clear (last result was not zero). 2 cycles if taken, 1 cycle if not. Range: ± 63 words from current PC.

Instruction introduced: **DEC** **DEC Rd** — Decrement register Rd by 1. 1 cycle. Affects Z, N, V, S flags. Does NOT affect C flag — careful when using in multi-byte arithmetic.

Instruction introduced: **RJMP** **RJMP k** — Relative jump. $PC = PC + k + 1$. Range: ± 2047 words. 2 cycles. Affects no flags.

Instruction introduced: **JMP** **JMP k** — Absolute jump to any address in 64K word program space. 3 cycles. 32-bit instruction. Affects no flags.

1.6 Building and Flashing

1.6.1 Step 1: Assemble

```
cd book/ch01_intro/src
avr-as -mmcu=attiny3217 -o blink.o blink.S
```

-mmcu=attiny3217 tells the tools exactly which device we are targeting. That choice affects instruction availability, predefined symbols, and which register definitions appear through `#include <avr/io.h>`.

No output on success. On error, `avr-as` prints the filename, line number, and problem.

1.6.2 Step 2: Link

```
avr-gcc -mmcu=attiny3217 -nostartfiles -o blink.elf blink.o
```

The linker is the stage where your program stops being a pile of pieces and becomes a real image with fixed addresses. It decides where `.vectors` and `.text` live in flash and resolves any symbol references between sections.

-nostartfiles suppresses the C runtime startup (`crt0.o`). Our `.vectors` section and `main` label take its place.

1.6.3 Step 3: Convert to HEX

```
avr-objcopy -O ihex blink.elf blink.hex
```

The `.elf` is the rich build artifact for tools and humans. The `.hex` is the lean transport format for programming the chip. Same program, different packaging.

1.6.4 Step 4: Inspect (optional but recommended)

```
avr-objdump -d blink.elf
```

Sample output:

```
blink.elf:      file format elf32-avr
```

Disassembly of section .vectors:

```
00000000 <__vectors>:
   0:  0c 94 02 00      jmp      0x4 <__ctors_end>
```

Disassembly of section .text:

```
00000004 <__ctors_end>:
   4:  0f ef           ldi      r16, 0xFF          ; RAMEND low byte
   6:  0d bf           out      0x3d, r16         ; SPL
   8:  0f e3           ldi      r16, 0x3F          ; RAMEND high byte
  a:  0e bf           out      0x3e, r16         ; SPH
  c:  08 e0           ldi      r16, 0x08          ; PIN3 mask
  e:  00 93 01 04     sts      0x0401, r16        ; PORTA_DIRSET
  ...
```

This is one of the healthiest habits you can build early: do not blindly trust source code. Inspect what the toolchain actually produced. `avr-objdump` shows the addresses, the machine-code bytes, and the final instruction sequence that will land in flash.

On this device/linker setup, `avr-objdump` may label the first `.text` address as `__ctors_end`; that is still the same address as our main label in this minimal program.

1.6.5 Step 5: Flash

ATtiny3217 Curiosity Nano (onboard debugger via USB):

```
pymcuprog write -d attiny3217 -t uart -f blink.hex
```

UPDI adapter (e.g. FT232R or CH340 with SerialUPDI):

```
avrdude -c serialupdi -p attiny3217 -P /dev/ttyUSB0 -b 115200 \
-U flash:w:blink.hex:i
```

Successful flash output:

```
avrdude: AVR device initialized and ready to accept instructions
avrdude: writing flash (40 bytes):
avrdude: 40 bytes of flash written
avrdude: verifying flash memory against blink.hex:
avrdude: 40 bytes of flash verified
avrdude done. Thank you.
```

The LED on PA3 should now blink at approximately 1 Hz.

1.7 A Minimal Makefile

Typing the build steps by hand is useful once because it teaches the pipeline. Typing them forever is just friction. Save this as `ch01_intro/src/Makefile`:

```
MCU      = attiny3217
TARGET   = blink
SRCS     = blink.S
```

```

AS      = avr-as
CC      = avr-gcc
OBJCOPY = avr-objcopy
OBJDUMP = avr-objdump
PROG    = serialupdi

ASFLAGS = -mmcu=$(MCU)
LDFLAGS = -mmcu=$(MCU) -nostartfiles

.PHONY: all flash disasm clean

all: $(TARGET).hex

$(TARGET).o: $(TARGET).S
    $(AS) $(ASFLAGS) -o $@ $<

$(TARGET).elf: $(TARGET).o
    $(CC) $(LDFLAGS) -o $@ $<

$(TARGET).hex: $(TARGET).elf
    $(OBJCOPY) -O ihex $< $@

flash: $(TARGET).hex
    pymcuprog write -d $(MCU) -t uart -f $<

disasm: $(TARGET).elf
    $(OBJDUMP) -d $<


clean:
    rm -f $(TARGET).o $(TARGET).elf $(TARGET).hex

```

Now build with `make`, flash with `make flash`, inspect with `make disasm`.

1.8 What Just Happened — Mental Model

Before moving on, fix this picture in your mind. This is the first real mental model of the whole book:



CPU reads instructions from Flash.

CPU reads/writes variables in SRAM.

Registers (r0-r31) are inside the CPU — fastest possible storage.

The program counter starts at 0x0000 after reset, fetches the `JMP` at the reset vector, lands in `main`, runs straight through the setup code, and then spins forever around `loop`. Nothing mystical is happening. The

CPU is just fetching, executing, and updating state one instruction at a time.

1.9 Summary

Concept	Key point
Assembly	One instruction = one CPU operation. No abstraction.
AVR-GAS	GNU Assembler, invoked as <code>avr-as</code> . GAS syntax.
Build pipeline	<code>.S</code> → assemble → <code>.o</code> → link → <code>.elf</code> → convert → <code>.hex</code> → flash
ATtiny3217 PORT	New-style: memory-mapped registers, use STS/LDS. OUTTGL is atomic.
Stack pointer	Must be valid before any subroutine call; on ATtiny3217 reset already loads <code>SP = RAMEND</code> .
Default clock	$\text{OSC20M} \div 6$ after reset: 3.3 MHz or 2.66 MHz, depending on FUSE.OSCCFG.

1.10 Exercises

1. Change the blink rate to 2 Hz (halve the delay loop count). Recalculate the required `r19` value.
 2. Use `PORTA_OUTSET` and `PORTA_OUTCLR` instead of `PORTA_OUTTGL`. Make the LED on for a longer period than it is off (asymmetric blink).
 3. Run `avr-objdump -d blink.elf` and find:
 - The byte address of the `rjmp loop` instruction.
 - The encoding of `sts 0x0407, r16` — how many bytes does it take?
 - The total size of the `.vectors` section.
 4. Add a second LED on PA5. Both LEDs should blink in opposite phases (one on while the other is off).
-

Next: Chapter 2 — AVR Architecture Overview


Chapter 2

AVR Architecture Overview

Before writing a single instruction, you need a mental model of what the AVR CPU is actually doing: where code lives, where data lives, what the registers are for, and what hidden state changes after each instruction. Without that model, assembly feels like memorising magic spells. With it, most instructions become straightforward and predictable.

2.1 Harvard Architecture

The AVR is a **Harvard** architecture. This means it has **two separate memory spaces** — one for instructions (flash) and one for data (SRAM) — connected to the CPU by separate buses.



Why does this matter to the programmer?

Because on AVR, "memory" is not one big uniform place. Where a value lives determines how you access it.

When your program is running, the CPU fetches instructions from **flash** automatically. You do not manually load the next instruction yourself; the program counter and instruction fetch hardware handle that in the

background. That is why your code can live safely in non-volatile flash while your working data changes constantly in SRAM.

Your **variables**, your **stack**, and most of the peripheral register space you work with as data live in **SRAM/data space**. That is the memory area you read and write with ordinary load/store instructions such as LD, LDS, ST, and STS. So if you push a register, call a subroutine, store a loop counter, or write to a peripheral register, you are operating in the data space, not in program flash.

Flash is different. It holds your program, and on many AVR parts it also holds constant tables or strings that you do not want to waste SRAM on. But flash is not ordinary read/write RAM. You cannot treat it like a normal destination for ST and expect the program image to change. Self-programming flash uses the NVM controller and the SPM path, which is a special mechanism covered later in the book.

This also affects **read-only data** such as lookup tables. On classic AVR, reading data back from flash requires LPM, because flash is on the program side of the machine. On AVRxt parts such as the ATtiny3217, flash is also mapped into the data address space, so some reads can be done through that mapping with ordinary data-memory instructions. Chapter 5 shows both styles and when each applies.

The practical rule is simple:

- If the value is temporary, changing, or part of the stack, think **SRAM/data space**.
- If the value is your code or a fixed table that should survive power loss, think **flash/program space**.

This separation is why you cannot simply think “an address is an address.” On AVR, you must use the instruction that matches the address space you are trying to access.

2.2 ATtiny3217 Memory Map

The ATtiny3217 uses the **AVRxt** core (also called avrxmega3). It has an important difference from classic AVR (ATmega328P): **flash is also mapped into the data address space** at offset 0x8000. This means you can read flash with ordinary LD instructions — no need for LPM.

That feature is convenient, but do not let it blur the architectural picture. Flash is still program memory. The chip is simply giving you a second way to read it. Thinking in terms of “program space” versus “data space” is still the right model.

Data Address Space (what LD/ST/LDS/STS see):

0x0000	I/O registers (IN/OUT range)	0x0000-0x003F (64 bytes)
0x003F		
0x0040	Extended I/O / Peripherals (LDS/STS only, not IN/OUT)	0x0040-0x0FFF
0x0FFF		
0x1000	(reserved)	
0x37FF		
0x3800	SRAM – 2 KB	0x3800-0x3FFF
0x3FFF		
0x4000	(reserved)	
0x7FFF		
0x8000	Flash – 32 KB (read-only) mapped into data space	0x8000-0xFFFF
0xFFFF		

Program Address Space (what the PC addresses, separate bus):

0x0000 Flash – 16K words (32 KB)
Byte addr = word addr × 2
0x3FFF
(word addresses; byte addresses go to 0x7FFF)

Key numbers to memorise:

Symbol	Value	Meaning
FLASHEND	0x7FFF	Last byte address in flash
RAMSTART	0x3800	First SRAM byte
RAMEND	0x3FFF	Last SRAM byte (init SP here)
MAPPED_FLASH_START	0x8000	Flash base in data address space

Key point: When you see a linker script or header file refer to `MAPPED_PROGMEM_START`, it means 0x8000 — where flash appears in the data bus. On classic AVR, flash is *not* in the data bus at all.

If you are new to memory maps, read the diagram like this:

- The **low addresses** are the control area: CPU and peripheral registers.
- The **middle SRAM block** is your working memory.
- The **high mapped-flash block** is read-only program memory made visible from the data side on AVRxt devices.

That one picture explains why some operations use IN/OUT, some use LDS/STS, and some use flash-specific mechanisms.


2.3 The Register File

The AVR CPU has **32 general-purpose 8-bit registers**, named `r0` through `r31`. They live inside the CPU — accessing them takes **0 extra cycles**. Compare that to SRAM: a load (LDS) takes 3 cycles; a store (STS) takes 2.

This is one of the main reasons assembly on AVR can be so efficient. If you can keep a value in a register, you avoid repeated trips to SRAM. A large part of “good AVR assembly” is really just good register management.

Register file layout:

r0		General purpose. Some instructions can only use r0 (e.g. MUL stores result in r1:r0).
r1		
r2		
r3		
...		
r15		
r16		← LDI can only load immediate into r16-r31. Most two-operand arithmetic works on any register, but immediate operands restrict you to r16-r31.
r17		
...		
r23		
r24		SBIW/ADIW work on r25:r24, r27:r26, r29:r28, r31:r30.
r25		
r26	XL	7



2.3.1 Register Pairs (16-bit Pointers)

The top six registers form three **pointer pairs** used for indirect memory access:

Pair	Registers	Name	Special use
X	r27:r26	X-pointer	Indirect load/store
Y	r29:r28	Y-pointer	Indirect load/store, stack frame pointer
Z	r31:r30	Z-pointer	Indirect load/store, LPM, indirect jumps/calls

When you see `LD r0, X` in assembly, it means “load the byte at the address stored in the X register pair into r0.” We cover this in depth in Chapter 5.

In practice, think of X, Y, and Z as your built-in address registers. The ordinary registers hold values; the pointer registers hold locations.

2.3.2 LDI Restriction

LDI — load immediate — **only works on r16–r31**. This is one of the most common beginner mistakes.

```
ldi r5, 42      ; ERROR: illegal operand – r5 not allowed with LDI
ldi r16, 42     ; OK
```

If you need a constant in r0–r15, load it into r16–r31 first, then MOV it:

```
ldi r16, 42
mov r5, r16     ; now r5 = 42
```


That may feel arbitrary at first, but it is simply part of the AVR instruction encoding. Some instructions have fewer bits available to describe operands, so they only work on part of the register file. You will see this kind of limit often in small instruction-set architectures.


2.4 The Status Register (SREG)

SREG is an 8-bit register at I/O address `0x3F` (accessible with IN/OUT). Every arithmetic and logic instruction updates some of its bits. Branch instructions read these bits to decide whether to jump.

You can think of SREG as the CPU's short-term memory of “what just happened.” Did the last result become zero? Did it overflow? Was there a carry? Branches do not inspect your registers directly; they inspect these flag bits.

SREG bit layout:





2.4.1 Flag by Flag

C — Carry Set when an addition produces a result > 255 (unsigned overflow) or a subtraction borrows. Also used by shift instructions (the bit shifted out). Used for multi-byte arithmetic (ADC, SBC).

```
ldi r16, 0xFF
ldi r17, 0x01
add r16, r17    ; result = 0x00, C=1 (carry out), Z=1
```

Z — Zero Set when the result of an operation is exactly zero. The most-used flag. BRNE (branch if not equal) branches when $Z=0$; BREQ branches when $Z=1$.

```
ldi r16, 1
dec r16        ; r16 = 0, Z=1
brne somewhere ; NOT taken - Z is set
breq somewhere ; taken   - Z is set
```

N — Negative Set when bit 7 of the result is 1 (result is negative in two's complement).

V — Overflow Set when two's-complement arithmetic overflows. For example, adding two positive numbers and getting a negative result.

```
ldi r16, 0x7F    ; +127 (largest positive signed 8-bit)
ldi r17, 0x01
add r16, r17     ; result = 0x80 = -128, V=1 (signed overflow)
```

S — Sign $S = N \text{ XOR } V$. This is the “true” sign of the result after accounting for overflow. Use BRGE/BRLT (which test S) for signed comparisons, not N alone.

H — Half-carry Carry out of bit 3 into bit 4. It matters for nibble-oriented arithmetic and some compare/add/subtract cases. AVR does **not** have a DAA instruction, so you usually inspect H only in specialised code.

T — Transfer A single-bit scratch pad. BST copies a bit from a register into T; BLD copies T into a register bit. Useful for bit manipulation without clobbering other flags.

I — Global Interrupt Enable When $I=1$, the CPU responds to interrupts. SEI sets it; CLI clears it. Entering an ISR automatically clears I; RETI restores it. Chapter 10 covers this in detail.

2.4.2 Reading and Saving SREG

SREG is a normal I/O register. You can save and restore it:

```
in    r0, SREG        ; save SREG
cli                    ; disable interrupts (clear I bit)
; ... critical section ...
out   SREG, r0         ; restore SREG (re-enables interrupts if I was set)
```

This pattern appears constantly in ISR prologues and critical sections.

If you remember only one thing about SREG, remember this: flags are part of program state. If your code depends on them, preserve them deliberately.

Instructions introduced: IN / OUT (revisited) IN Rd, A — Load I/O register at address A into Rd. 1 cycle. OUT A, Rr — Write Rr to I/O register at address A. 1 cycle. Range: A = 0x00–0x3F only. For higher addresses, use LDS/STS.

Instruction introduced: SEI SEI — Set Global Interrupt Enable (I bit in SREG). 1 cycle.


Instruction introduced: CLI CLI — Clear Global Interrupt Enable. 1 cycle. Next instruction always executes before any interrupt is served (pipeline guarantee).

2.5 The Stack and Stack Pointer

The stack is a **last-in, first-out** (LIFO) region of SRAM. The AVR stack **grows downward** — pushing data decrements the stack pointer first, then writes.

That “grows downward” detail matters. On AVR, the top of SRAM is the natural starting point for the stack, and each push moves toward lower addresses. Beginners often imagine the stack growing upward because that feels more intuitive. AVR does the opposite.

SRAM at 0x3800–0x3FFF:



SP is a 16-bit register implemented as two I/O registers:

- SPL at I/O address 0x3D (low byte of SP)
- SPH at I/O address 0x3E (high byte of SP)

On the ATtiny3217, reset hardware loads SP with RAMEND automatically. You will still see startup code write it explicitly, because that is harmless and keeps examples portable to older AVR parts where software initialisation is required.

So there are really two questions:

- What value should SP start with? RAMEND.
- Who sets it? On ATtiny3217, the hardware does; on some older AVR parts, startup code must do it.

Both bytes are in the IN/OUT address range, so explicit initialisation uses OUT:

```
ldi r16, lo8(RAMEND)    ; 0xFF
out SPL, r16
ldi r16, hi8(RAMEND)    ; 0x3F
out SPH, r16
```

2.5.1 PUSH and POP

```
push r16    ; SP -= 1, then write r16 to SRAM[SP] (2 cycles)
pop  r16    ; read r16 from SRAM[SP], then SP += 1 (2 cycles)
```

Push decrements first, then writes. Pop reads first, then increments.

2.5.2 How CALL Uses the Stack

CALL / RCALL automatically pushes the **return address** (PC+1) onto the stack so RET can find it:

```
rcall my_function    ; push PC+1, jump to my_function
; RET in my_function pops the return address and jumps back here
```

On ATtiny3217 (16K word address space), return addresses are 2 bytes (16-bit). On devices with >128KB flash, they are 3 bytes. Always account for this when calculating stack depth.

This is why even a program with “no local variables” still uses stack space. Every nested call consumes stack bytes for the return address, and any saved registers add to that cost.

2.6 The Program Counter (PC)

The Program Counter is a register inside the CPU that holds the **word address** of the next instruction to execute. It is not directly accessible via IN/OUT — you control it only through jump, branch, and call instructions.

In other words, the PC is where control flow lives. Arithmetic changes data; branches and calls change the PC.

```
After reset:      PC = 0x0000 (first word in flash)
After JMP main:   PC = address of first instruction in main
After RCALL foo:  PC = address of foo; old PC+1 on stack
After RET:        PC = popped from stack
```

The PC is 14 bits wide on ATtiny3217 (addressing 0x0000–0x3FFF in word units, covering the full 32KB flash).

2.6.1 Instruction Fetch Pipeline

AVR uses a **single-stage pipeline**: while the CPU executes instruction N, it fetches instruction N+1 from flash. This is why most instructions execute in 1 cycle — the fetch is free. Instructions that change the PC (branches, jumps, calls, returns) take 2–4 cycles because the pipeline must be flushed and a new fetch begun.

You do not need to become a pipeline expert to use AVR well. The practical lesson is simpler: straight-line code is cheap, but changing direction costs extra cycles.

Cycle:	1	2	3	4	
Fetch:	N	N+1	N+2	N+3	
Execute:		N	N+1	N+2	← 1 cycle/instruction (normal)

With a taken branch at instruction N:

Fetch:	N	N+1	target	target+1	
Execute:		N	---	target	← N+1 fetch wasted; branch costs 2 cycles

This is why BRNE costs 1 cycle when not taken (no pipeline flush) and 2 cycles when taken (flush + re-fetch from target).

2.7 I/O Register Space

The ATtiny3217 peripherals are controlled through memory-mapped I/O registers. They fall into two zones:

Address range	Access method	Examples
0x0000–0x003F	IN / OUT (fast)	VPORTB_OUT, SREG, SPL, SPH, CCP
0x0040–0x0FFF	LDS / STS (2 words)	PORTB_DIR, TCA0_CNT, USART0_RXDATA1

The low 64 addresses (0x00–0x3F) are the “classic” I/O space reachable with 1-word IN/OUT instructions in 1 cycle. The low half of that space (0x00–0x1F) is also reachable by SBI, CBI, SBIC, and SBIS. Everything above 0x3F needs the 2-word LDS/STS instructions, which take 3 and 2 cycles respectively.

This is one of the most important performance distinctions on AVR. Two registers may both control hardware, but if one is in low I/O space and the other is in extended space, the instruction choices and timing differ.

The ATtiny3217’s important low I/O registers:

I/O Addr	Name	Description
0x05	VPORTB_OUT	Virtual PORTB output register (fast GPIO access)
0x04	VPORTB_DIR	Virtual PORTB direction register
0x3F	SREG	Status Register
0x3E	SPH	Stack Pointer High
0x3D	SPL	Stack Pointer Low
0x34	CCP	Configuration Change Protection (unlock protected registers)
0x1F	GPIOR3	General Purpose I/O Register 3 (scratch)
0x1E	GPIOR2	General Purpose I/O Register 2 (scratch)
0x1D	GPIOR1	General Purpose I/O Register 1 (scratch)
0x1C	GPIOR0	General Purpose I/O Register 0 (scratch)

VPORTx_* mirrors let you do single-cycle GPIO with IN/OUT and bit instructions, while the full PORTx_* peripheral block provides the richer DIRSET/OUTCLR/OUTTGL registers in extended I/O space.

GPIOR0–3 are four registers with no peripheral function — pure scratch RAM accessible via IN/OUT in a single cycle. Useful for storing flags or counters that need fast ISR access.

Think of them as “fast little mailboxes” between parts of your program. They are not magic, but they can be handy when one-byte state needs to be very cheap to access.

2.8 Putting It Together — A Register-State Walk-Through

Let’s trace the CPU state through the blink program from Chapter 1, cycle by cycle through the setup section:

The goal here is not to memorise every state transition. It is to get used to thinking in terms of registers, memory side effects, and whether flags changed.

```
; State at reset:
; PC = 0x0000
; SP = 0x3FFF on ATtiny3217 reset hardware
; SREG = 0x00
; Treat general registers as undefined until your code sets them

ldi r16, lo8(RAMEND)    ; r16 = 0xFF. PC → next instr. 1 cycle.
out SPL, r16            ; explicit re-init of SP low byte. 1 cycle.
ldi r16, hi8(RAMEND)    ; r16 = 0x3F. PC → next. 1 cycle.
out SPH, r16            ; explicit re-init of SP high byte. 1 cycle.

ldi r16, (1 << 3)       ; r16 = 0x08. 1 cycle.
```

```
sts PORTA_DIRSET, r16 ; Write 0x08 to 0x0401.
                      ; PA3 now output. 2 cycles.
```

```
; State now:
; r16 = 0x08, SP = 0x3FFF, PORTA_DIR bit 3 = 1, all flags unchanged
```

Flags: LDI, OUT, STS do **not** affect SREG at all. Only arithmetic, logic, and bit operations change flags.

2.9 Key Architectural Constraints Summary

These are the constraints you will bump into constantly. They look dry, but they explain a large share of beginner errors.

Constraint	Rule
LDI range	r16–r31 only
IN/OUT range	I/O addresses 0x00–0x3F only
SBIW/ADIW pairs	r25:r24, r27:r26, r29:r28, r31:r30 only
Pointer pairs	X = r27:r26, Y = r29:r28, Z = r31:r30
MUL result	Always in r1:r0
Stack direction	Grows downward; PUSH decrements first
Return address	2 bytes on stack (ATtiny3217, ≤128KB flash)
Interrupt latency	4 cycles minimum (flush + vector fetch)
Flash read via LD	Use address + 0x8000 (mapped flash, AVRxt)
Flash read via LPM	Use Z = byte address, reads via program bus

2.10 Summary

Harvard = two buses. Flash for code. SRAM for data.

ATtiny3217 memory map:

```
0x0000–0x003F  I/O (IN/OUT)
0x0040–0x0FFF  Peripherals (LDS/STS)
0x3800–0x3FFF  SRAM (2 KB)
0x8000–0xFFFF  Flash mapped in data space (AVRxt feature)
```

32 registers (r0–r31):

```
r0–r15: general, no LDI
r16–r31: general + LDI + most immediates
r26:r27 = X, r28:r29 = Y, r30:r31 = Z (pointer pairs)
```

SREG (0x3F): C Z N V S H T I – set by ALU, read by branches.

SP = SPH:SPL, initialise to RAMEND. Stack grows down.

On ATtiny3217, reset hardware already loads SP = RAMEND.

PC = word address, 14 bits. Pipeline prefetch = 1 cycle/instr normal.

2.11 Exercises

1. What is the data-bus address of the first byte of flash on ATtiny3217? What is the data-bus address of the last byte of flash?
2. After `ldi r16, 0x80` and `ldi r17, 0x80` and `add r16, r17`, which SREG flags are set? Work it out by hand, then verify with:

```
avr-as -mmcu=attiny3217 -o /tmp/t.o - <<'EOF'
.text
ldi r16, 0x80
ldi r17, 0x80
add r16, r17
EOF
avr-objdump -d /tmp/t.o
```

(The disassembly won't show flag values — you must reason through them.)

3. Write a short assembly snippet that:
 - Saves SREG to r0
 - Disables interrupts
 - Does some work (e.g. `ldi r16, 42`)
 - Restores SREG (and thus the interrupt state)
4. Calculate the maximum stack depth for a program that calls a function, which calls another function. How many bytes does the stack use just for return addresses?
5. Why can't you write `ldi r15, 0xFF`? What two-instruction sequence achieves the same result?

Next: Chapter 3 — AVR-AS Syntax and Directives

Chapter 3

AVR-AS Syntax and Directives

Assembly source is more than a list of instructions. Before the CPU can run anything, the assembler has to answer basic questions: where does the code go, what data should be emitted, which names are constants, and which labels refer to real addresses? Those decisions are controlled by **directives** — assembler commands that begin with a dot.

This chapter is about that “language around the instructions.” If Chapter 1 showed how to make the chip do something, Chapter 3 explains how to write assembly source in a way the toolchain can understand and organise correctly.

3.1 File Extensions: **.S** vs **.s**

GAS accepts two extensions:

Extension	Preprocessor	Use
.S	Yes — C preprocessor runs first	Always use this
.s	No — raw GAS input	Rarely; only for generated files

The uppercase **.S** lets you use `#include`, `#define`, `#ifdef`, and all other C preprocessor directives. Since we always want `#include <avr/io.h>` for register names, **always use **.S****.

For a beginner, the practical rule is simple: if you are hand-writing AVR assembly for this book, name the file with an uppercase **.S** and do not think about it further.

The build flow preprocesses **.S** files before assembly. If you want to inspect that preprocessed output yourself, use the compiler driver:

```
# Inspect the preprocessed source that avr-as will assemble:
avr-gcc -x assembler-with-cpp -mmcu=attiny3217 -E blink.S
```

You normally do not need to run this manually — `avr-as -mmcu=attiny3217 file.S` handles **.S** preprocessing for you.

3.2 Comments

GAS accepts three comment styles. In .S files all three work:

```
; This is a comment – semicolon style, common in AVR assembly
```

```
// This is a comment – C++ line style
```

```
/* This is a
   block comment – C style */
```

Semicolon is traditional in AVR assembly and what you'll see most. Use block comments for file headers and section separators.

Style matters more in assembly than in many other languages because the code is so compact. A short comment can carry a lot of explanatory weight.

3.3 Instruction Syntax

AVR GAS uses a **destination-first, source-second** convention:

```
mnemonic destination, source
```

Examples:

```
mov r16, r17      ; r16 ← r17
add r16, r17      ; r16 ← r16 + r17
ldi r16, 42       ; r16 ← 42
sts 0x0421, r16    ; mem[0x0421] ← r16
ld r16, X         ; r16 ← mem[X]
```

This matches GCC output and the Atmel/Microchip instruction set manual.

3.4 Number Literals

GAS accepts numbers in four bases:

```
ldi r16, 42       ; decimal
ldi r16, 0x2A     ; hexadecimal
ldi r16, 0b00101010 ; binary
ldi r16, 052      ; octal (leading zero – easy to mistake for decimal!)
```

All four produce the same value (42). Avoid octal — the leading-zero trap causes real bugs. Hex and binary are clearest for hardware work.

Character literals also work:

```
ldi r16, 'A'      ; r16 = 65
ldi r16, '\n'     ; r16 = 10
```

3.5 Expressions

GAS evaluates constant expressions at assembly time. You can use:

```
ldi r16, (1 << 4)           ; bit shift – evaluates to 16
ldi r16, (1 << 4) | (1 << 2) ; bitwise OR – evaluates to 20
ldi r16, RAMEND & 0xFF      ; mask – evaluates to 0xFF
ldi r16, RAMEND / 256       ; division – evaluates to 0x3F
ldi r16, 10 * 5 - 8         ; arithmetic – evaluates to 42
```

These are evaluated at **assembly time** — no runtime cost. They exist only to make source code readable.

That is an important distinction. These expressions do not generate arithmetic instructions. The assembler does the math once while building the file, and the CPU never sees the expression itself.

3.5.1 Built-in Expression Functions

GAS provides helper functions for splitting 16-bit values into bytes:

```
ldi r16, lo8(RAMEND)        ; low byte of 0x3FFF = 0xFF
ldi r16, hi8(RAMEND)        ; high byte of 0x3FFF = 0x3F
ldi r16, lo8(0x1234)        ; = 0x34
ldi r16, hi8(0x1234)        ; = 0x12
```

Two more functions matter for AVRxt flash access (see Section 3.11):

```
; pm(label) – convert byte address to word address (for IJMP/ICALL targets)
ldi ZL, lo8(pm(my_func))    ; ZL = word address low byte of my_func
ldi ZH, hi8(pm(my_func))    ; ZH = word address high byte of my_func
```

3.6 Labels

A label marks a position in the output. It ends with a colon:

```
main:
    ldi r16, 0xFF

loop:
    dec r16
    brne loop        ; reference label by name – no colon when referencing
```

Labels have two uses:

1. **Jump/branch targets** — as above
2. **Data addresses** — when placed before `.byte`, `.word`, etc.

You can think of a label as a bookmark the assembler turns into an address. Sometimes that address is where execution should jump. Sometimes it is where data begins.

```
table:
    .byte 1, 2, 3, 4    ; "table" now equals the address of byte 1
```

3.6.1 Global vs Local Labels

By default, labels are **local to the object file** — the linker cannot see them unless you explicitly export them:

```
.global main        ; export "main" so the linker can use it as entry point
.global __vectors    ; export vector table symbol

main:                ; now visible to linker
...

```

```
internal_helper:    ; NOT global – invisible outside this file
    ...
```

3.6.2 GAS Local Labels (Numeric)

Naming every loop, branch target, and temporary label creates a pollution of unique names. GAS solves this with **numeric local labels**:

```
1:                ; define label "1"
    dec r16
    brne 1b        ; "1b" = most recent backward occurrence of label "1"

    tst r16
    breq 2f        ; "2f" = next forward occurrence of label "2"
    ldi r16, 0xFF
2:                ; define label "2"
    ...
```

Rules:

- Any number 0–9 (or longer: 10:, 255:) can be a local label
- Nb = nearest backward occurrence of N:
- Nf = nearest forward occurrence of N:
- The same number can be reused anywhere in the file — each definition is distinct. GAS resolves 1b to the *nearest* preceding 1:.

This is invaluable in loops:

```
ldi r18, 10        ; outer loop count
1:                ; outer loop top
    ldi r17, 0
    ldi r16, 0
2:                ; inner loop top
    sbiw r16, 1
    brne 2b        ; back to inner "2:"
    dec r18
    brne 1b        ; back to outer "1:"
```

No naming conflict — 1b and 2b refer to the right labels unambiguously.

This style looks strange at first, but it becomes natural quickly. Named labels are best for important program locations such as `main` or `uart_tx`. Numeric labels are best for short internal loops where a descriptive name would add more clutter than clarity.

3.6.3 File-Scoped Local Labels (`.L` prefix)

Labels starting with `.L` are stripped from the symbol table by the linker and do not appear in debug info. Use them for internal subroutine labels you don't want to export:

```
my_subroutine:
.Lloop:
    dec r16
    brne .Lloop
    ret
```

`.Lloop` will not appear in `avr-nm` output or debugger symbol lists.

3.7 Sections

A **section** is a named region of output. The linker arranges sections into the final memory map. Every assembly statement belongs to the *current* section.

If labels are bookmarks, sections are containers. They tell the linker what kind of thing you are defining: executable code, read-only flash data, SRAM data, zero-initialised storage, and so on.

3.7.1 Core Sections

```
.section .text      ; code – goes to flash
.section .data      ; initialised SRAM data – startup code must copy it in
.section .bss       ; zero-initialised SRAM data – startup code must clear it
.section .rodata    ; read-only constants – stay in flash
```

Shorthand forms (equivalent):

```
.text      ; same as .section .text
.data      ; same as .section .data
.bss       ; same as .section .bss
```

You can switch between sections multiple times in one file:

```
.text
foo:
    ldi r16, 5
    ret

.data
my_var: .byte 0

.text
bar:
    ldi r16, 10
    ret
```

The assembler concatenates all `.text` chunks together, all `.data` chunks together, etc. The order of sections in the output is determined by the linker script, not the source order.

That is why switching sections in the middle of a file is normal. You are not describing the final physical layout directly; you are describing which bytes belong to which category.

3.7.2 Section Flags and Types (Full Form)

The full `.section` directive takes optional flags:

```
.section name, "flags", @type
```

Flag	Meaning
a	Allocatable (takes space in memory)
x	Executable
w	Writable

Type	Meaning
@progbits	Contains data/code

Type	Meaning
@nobits	Takes space but has no data in file (BSS)

Most of the time you use the shorthand and GAS fills in reasonable defaults. You need the full form for special sections like `.vectors`:

```
.section .vectors, "ax", @progbits
```

This tells the linker: place this section at the start of flash (the linker script maps `.vectors` to address 0x0000), it contains code (x), and it is allocatable (a).

You do not need to memorise all the flags immediately. What matters is knowing that the full form exists when you need something more specific than the common defaults.

3.8 Data Directives

These directives emit bytes into the current section.

That phrase “emit bytes” is worth taking literally. Directives such as `.byte` and `.word` do not create variables in the high-level-language sense. They cause exact byte patterns to be written into the object file.

3.8.1 `.byte` — 8-bit values

```
.byte 0x42          ; one byte: 0x42
.byte 1, 2, 3, 4    ; four bytes
.byte 'H', 'i', '!' ; character literals
.byte (1 << 3)       ; expression – evaluates to 8
```

3.8.2 `.word` — 16-bit values (little-endian)

```
.word 0x1234        ; emits 0x34, 0x12 (little-endian)
.word 0, 1, 2       ; six bytes: three 16-bit words
```

Warning: On AVR, `.word` is 16 bits. In non-AVR GAS targets `.word` can be 32 bits. Never assume word size when reading unfamiliar assembly.

3.8.3 `.long` — 32-bit values

```
.long 0x12345678    ; emits 78 56 34 12 (little-endian)
```

3.8.4 `.ascii` and `.asciz` — Strings

```
.ascii "Hello"      ; 5 bytes: H e l l o – NO null terminator
.asciz "Hello"      ; 6 bytes: H e l l o 0x00 – WITH null terminator
.string "Hello"     ; alias for .asciz
```

Almost always use `.asciz` or `.string` so C string functions work correctly.

If you forget the terminating zero, the bytes are still valid assembly data, but any code that expects a C-style string will keep reading past the intended end.

3.8.5 .skip and .space — Reserve Bytes

```
buf:    .skip 16          ; reserve 16 bytes (content undefined in .text/.data)
buf:    .skip 16, 0       ; reserve 16 bytes, filled with 0x00
```

In .bss, content is always zeroed by startup code regardless of fill value.

3.8.6 .balign — Alignment

```
.balign 2          ; align to 2-byte boundary (pad with 0x00 if needed)
.balign 4          ; align to 4-byte boundary
```

Required before .word or .long data in some contexts. Also used to align code for performance (though on AVR this matters less than on ARM).

3.9 Constant Directives

3.9.1 .equ — Define a Constant

```
.equ LED_PIN, 4
.equ LED_MASK, (1 << LED_PIN) ; expressions work
.equ DELAY, 5000
```

.equ defines a **symbol** with a value. It cannot be redefined later. Use it for things that truly never change: pin numbers, masks, magic values.

In practice, .equ is how you turn anonymous numeric constants into named ideas. 0x08 is just a number; LED_MASK tells the reader what the number means.

Using the constant:

```
ldi r16, LED_MASK ; assembler substitutes the value at assembly time
```

3.9.2 .set — Reassignable Constant

```
.set count, 0
.set count, count + 1 ; OK — .set can be redefined
.set count, count + 1 ; count = 2 now
```

.set is useful for building values incrementally or for conditional definitions. Equivalent to #define + #undef patterns but within GAS.

3.9.3 #define — Preprocessor Macro

Because .S files go through cpp, you can also use:

```
#define LED_PIN 4
#define LED_MASK (1 << LED_PIN)
#define SET_BIT(reg, bit) ((reg) |= (1 << (bit))) ; functional macro
```

#define is processed *before* GAS sees the file; .equ is processed *by* GAS. For simple constants, .equ is idiomatic in assembly. For complex macros or conditional compilation, #define/#ifdef are necessary.

If you are unsure which to choose, prefer .equ for plain assembly constants. Reach for #define only when you really need preprocessor behavior.

3.10 The .org Directive

.org sets the **location counter** — the current output offset within the section. It can move the location counter **forward**, creating padding, but it cannot move it backward.

This is a more advanced tool than it first appears. It is useful when a layout must match hardware-defined slots, such as interrupt vectors, but it is also an easy way to create confusing gaps if used casually.

```
.org 0x0000
    jmp main           ; lands exactly at byte address 0

.org 0x0004
    jmp isr_int0       ; next 4-byte JMP slot

.org 0x0008
    jmp isr_pcint0     ; next slot
```

If you use JMP, each vector entry consumes 4 bytes. If you use RJMP, each entry consumes 2 bytes, so the .org spacing changes accordingly. The exact vector table is in the datasheet and in <avr/iotn3217.h> as _VECTOR(n) macros. Chapter 10 builds the full vector table.

So the real lesson is not “always use .org.” The lesson is: use it when the hardware requires exact placement, and understand the instruction size when you do.

Warning: .org without context is relative to the *start of the current section*. In .vectors, .org 0 means “start of the vector section”, which the linker places at flash address 0. In .text, .org 0x100 would create a 0x100-byte gap from the section's start — usually not what you want.

3.11 Flash Data on AVRxt (ATtiny3217)

This is where AVRxt differs from classic AVR in an important way.

This section is easy to skim past, but it matters because it changes how you think about constants in flash. On older AVR parts, flash reads feel like a special case. On AVRxt, they can look much more like ordinary data reads.

Classic AVR (ATmega328P): Flash is NOT in the data address space. To read a byte from flash at runtime, you must use the LPM instruction with the Z register pointing to the flash byte address.

AVRxt (ATtiny3217): Flash IS mapped in the data address space at offset MAPPED_PROGMEM_START (= 0x8000). You can read flash with ordinary LD instructions by adding this offset to the flash address.

```
.section .rodata

table:
    .byte 10, 20, 30, 40, 50

.section .text

read_table_entry:           ; r24 = index, returns value in r25
    ldi ZL, lo8(table + MAPPED_PROGMEM_START)
    ldi ZH, hi8(table + MAPPED_PROGMEM_START)
    add ZL, r24             ; Z now points to table[r24] in data space
    adc ZH, r1              ; propagate carry (r1 = 0 by convention)
    ld r25, Z               ; read from flash via data bus
    ret
```

The assembler expression `table + MAPPED_PROGMEM_START` adds 0x8000 to the flash byte address of `table`. The result fits in 16 bits because the table starts in the first 32KB of flash.

Conceptually, you are taking a flash address and translating it into the corresponding data-space view of the same bytes.

Old LPM method still works on AVRxt and is sometimes needed for compatibility:

```
ldi ZL, lo8(table)           ; Z = flash byte address (no +0x8000)
ldi ZH, hi8(table)
lpm r25, Z                   ; read from flash via program bus
```

Both methods read the same data. On AVRxt, LD with mapped address is simpler and more natural. LPM is covered in Chapter 5.

That means you should understand both styles, even if the mapped-flash method is the one you use most often on the ATtiny3217.

3.12 The .include Directive

```
.include "macros.inc"        ; include a GAS file
```

Inserts the contents of another file at the point of the `.include`. The included file is processed by GAS (not cpp), so it cannot use `#define`.

For files that need preprocessor features, use `#include` instead:

```
#include <avr/io.h>           ; cpp include – register definitions
#include "my_macros.h"        ; cpp include – your own header
#include "data_tables.S"      ; GAS include – another assembly file
```

3.13 Macro Directives

GAS has its own macro system separate from the C preprocessor.

```
.macro delay_ms ms_count
    ldi r19, \ms_count
.Lmacro_outer_\@:
    ldi r18, 0
    ldi r17, 0
.Lmacro_inner_\@:
    sbiw r17, 1
    brne .Lmacro_inner_\@
    dec r19
    brne .Lmacro_outer_\@
.endm
```

Using the macro:

```
delay_ms 100                ; expands inline – \ms_count → 100
delay_ms 500                ; second expansion – \@ ensures unique label names
```

Key syntax:

- `\param` — substitute parameter value
- `\@` — unique invocation counter (prevents label collisions across expansions)

- `.endm` — end macro definition

GAS macros are powerful but can obscure what code is actually emitted. Prefer subroutines (`rcall/ret`) over macros for anything that runs more than once, to save flash space.

This is a good place to be disciplined. Macros are tempting because they feel convenient, but they duplicate code at every use site. In assembly, that cost is often visible.

3.14 Full Section Layout Example

This is the structure of a well-organised AVR assembly file:

```
/*
 * myprogram.S - description
 * Build: avr-as -mmcu=attiny3217 -o myprogram.o myprogram.S
 */

#include <avr/io.h>

/* — Constants — */
.equ    MY_PIN,    5
.equ    MY_MASK,   (1 << MY_PIN)

/* — Read-only data (flash) — */
.section .rodata

my_table:
    .byte 1, 2, 4, 8, 16, 32, 64, 128

my_string:
    .asciz "ATtiny3217"

/* — Uninitialised SRAM — */
.section .bss

rx_buffer: .skip 64      ; 64-byte ring buffer
rx_head:   .skip 1
rx_tail:   .skip 1

/* — Initialised SRAM (requires startup copy code or CRT startup) — */
.section .data

state:     .byte 0x01     ; starts as 1 in SRAM

/* — Interrupt vector table — */
.section .vectors, "ax", @progbits
.global __vectors

__vectors:
    jmp     main          ; 0x0000 reset

/* — Code — */
.section .text
```

```
.global main

main:
    ldi r16, lo8(RAMEND)
    out SPL, r16
    ldi r16, hi8(RAMEND)
    out SPH, r16
    ; ... rest of program
```

3.15 Examining the Output with `avr-objdump`

After assembling, inspect what the directives produced:

```
avr-as -mmcu=attiny3217 -o syntax_demo.o syntax_demo.S
avr-gcc -mmcu=attiny3217 -nostartfiles -o syntax_demo.elf syntax_demo.o
avr-objdump -s -d syntax_demo.elf
```

The `-s` flag dumps the **raw contents** of every section as hex. Example output for `.rodata`:

Contents of section `.rodata`:

```
0074 01000101 01000101 01010100 48656c6c  ....Hell
0084 6f204156 5200                o AVR.
```

Cross-reference that output with the source: `morse_dit = [01, 00]`, `morse_dah = [01, 01, 01, 00]`, `greeting = "Hello AVR\0"`. Each byte appears in order. This is how you confirm that your directives produced exactly the layout you intended.

The `-d` flag shows the **disassembly** of `.text`:

```
00000008 <main>:
    8: 0f ef          ldi     r16, 0xFF
    a: 0d bf          out     0x3d, r16
    ...
   1a: e0 e7          ldi     r30, 0x70    ; ZL = lo8(morse_dit + 0x8000)
   1c: f0 e8          ldi     r31, 0x80    ; ZH = hi8(...)
```

Notice `ZH = 0x80` — the `0x8000` mapped flash offset is baked into the immediate at assembly time.

That is a perfect example of why inspecting output is useful: you can see a source-level idea become a real machine-level value.

3.16 Quick Reference

Directive	Purpose	Example
<code>.equ</code>	Define constant	<code>.equ PIN, 4</code>
<code>.set</code>	Reassignable constant	<code>.set i, 0</code>
<code>.byte</code>	Emit bytes	<code>.byte 1, 2, 3</code>
<code>.word</code>	Emit 16-bit words	<code>.word 0x1234</code>
<code>.long</code>	Emit 32-bit words	<code>.long 0xDEADBEEF</code>
<code>.ascii</code>	String, no null	<code>.ascii "hi"</code>
<code>.asciz</code>	String with null	<code>.asciz "hi"</code>
<code>.skip</code>	Reserve bytes	<code>.skip 16</code>

Directive	Purpose	Example
<code>.balign</code>	Align location counter	<code>.balign 2</code>
<code>.org</code>	Set location counter	<code>.org 0x0000</code>
<code>.global</code>	Export symbol	<code>.global main</code>
<code>.section</code>	Switch/create section	<code>.section .text</code>
<code>.include</code>	Include GAS file	<code>.include "f.S"</code>
<code>.macro/.endm</code>	Define GAS macro	<code>.macro name p1</code>

Expression	Meaning	Result for RAMEND=0x3FFF
<code>lo8(x)</code>	Low byte	<code>0xFF</code>
<code>hi8(x)</code>	High byte	<code>0x3F</code>
<code>pm(label)</code>	Word address	<code>addr / 2</code>
<code>(1 << n)</code>	Bit mask	<code>(1 << 4) = 0x10</code>

3.17 Summary

`.S` extension → `cpp` runs first → `#include` works.

Sections:

```
.text    = flash code
.rodata  = flash constants (read via mapped addr + 0x8000 on AVRxt)
.data    = SRAM, initialised if your startup code/CRT copies it
.bss     = SRAM, zero-init if your startup code/CRT clears it
.vectors = must be first in flash, holds the interrupt vector table
```

Labels:

```
name:    = define label
.lname:  = local, stripped from symbol table
1:       = numeric – reusable, resolved by 1f / 1b
```

Constants: `.equ` (fixed), `.set` (reassignable), `#define` (preprocessor).

AVRxt flash read: address + `MAPPED_PROGMEM_START` (0x8000), then LD.

3.18 Exercises

1. Write a `.rodata` section containing a table of the first 8 powers of 2 (1, 2, 4, 8, ...). Write code that reads `table[r24]` using the mapped flash address and returns the value in `r25`.
2. Add a `.bss` variable `byte_count` and a `.data` variable `max_count` initialised to 10. Write code that increments `byte_count` using LDS/ STS until it equals `max_count`, then halts.
3. Rewrite the Chapter 1 blink delay using a GAS `.macro` named `delay` that takes one parameter (outer loop count). Show that two calls `delay 25` and `delay 50` both assemble without label conflicts.
4. Assemble `syntax_demo.S` and run:


```
avr-objdump -s syntax_demo.elf
```

Verify the bytes in `.rodata` match the `morse_dit`, `morse_dah`, and `greeting` definitions in the source.

5. What happens if you put `.word 0x1234` inside a `.bss` section? Try it and inspect the linker output. Why does the linker warn or error?

Next: Chapter 4 — Instruction Set Basics

Chapter 4

Instruction Set Basics

The AVR instruction set has roughly 135 instructions, but a much smaller core set does most of the real work in ordinary firmware. This chapter focuses on that core: how instructions are encoded, which register restrictions matter, how arithmetic and logic actually affect the flags, and how to read the CPU's behaviour from the instruction sequence.

The goal is not to memorise a catalog. The goal is to make the instruction set feel predictable.

4.1 Instruction Encoding

Every AVR instruction is encoded as one or two 16-bit words (2 or 4 bytes) stored in flash. The CPU fetches one word per clock cycle. Most instructions are a single word and execute in 1–2 cycles.

That encoding detail matters because it affects both code size and speed. On AVR, the shape of an instruction is often the reason for its operand limits. If an instruction only has a few bits available to describe a register or an immediate, the hardware simply cannot support every possible combination.

4.1.1 Single-Word Instructions (16-bit)

The vast majority. The 16 bits encode the opcode, destination register, and source register or small immediate.

LDI Rd, K – Load Immediate

Encoding: 1110 KKKK dddd KKKK

$\begin{array}{c} \text{-----} \quad \text{-----} \\ \text{L} \quad \text{opcode} \quad \text{L} \quad \text{d} = \text{Rd} - 16 \quad (\text{only r16-r31} \rightarrow \text{d} = 0-15) \\ \text{K} = 8\text{-bit immediate split across bits} \end{array}$

ldi r20, 0xAB:

d = 20 - 16 = 4 = 0b0100

K = 0xAB = 0b10101011

→ 1110 1010 0100 1011 = 0xEA4B

ADD Rd, Rr – Add without Carry

Encoding: 0000 11rd dddd rrrr

d = Rd (0-31), r = Rr (0-31)

add r16, r17:

d = 16 = 0b10000, r = 17 = 0b10001
 → 0000 1101 0000 0001 = 0x0D01

4.1.2 Two-Word Instructions (32-bit)

These need a full 16-bit address or program address that doesn't fit in the first word:

Instruction	Why 32-bit
LDS Rd, k	Data address k is 16 bits — needs second word
STS k, Rr	Same
JMP k	Program address k is 22 bits — needs second word
CALL k	Same

STS k, Rr — Store to SRAM

Word 1: 1001 001r rrrr 0000 (r = Rr, 5 bits)

Word 2: kkkk kkkk kkkk kkkk (16-bit address k)

sts PORTB_DIRSET, r16:

PORTB_DIRSET = 0x0421

Word 1: 0x9200 | (16 << 4) = 0x9200

Word 2: 0x0421

→ 4 bytes total in flash

This is why STS costs 2 cycles and takes twice the flash of OUT. Always use OUT when the I/O address is ≤ 0x3F.

This is a recurring AVR theme: the shorter instruction is usually faster, but only if the target address fits the shorter encoding.

4.1.3 Cycle Counts at a Glance

Category	Typical cycles	Notes
Register ops (ADD, AND, MOV...)	1	Pipeline: fetch overlaps
Loads from SRAM (LD, LDS)	2–3	Memory access latency
Stores to SRAM (ST, STS)	2	
IN / OUT	1	Low I/O space only
Taken branch	2	Pipeline flush
Not-taken branch	1	No flush
RCALL / RET	4	Stack push/pop + flush
CALL / JMP (absolute)	4–5	32-bit instr + flush
MUL	2	Result in r1:r0
NOP	1	Explicit no-op

4.2 Instruction Categories

AVR Instruction Categories										
Data Transfer	MOV	MOVW	LDI	LD	ST	LDS	STS	LPM	PUSH	POP

	IN OUT
Arithmetic	ADD ADC SUB SBC SUBI SBCI ADIW SBIW INC DEC NEG MUL MULS MULSU
Logic	AND ANDI OR ORI EOR COM TST CLR SER
Shift / Rotate	LSL LSR ROL ROR ASR SWAP
Bit Operations	SBI CBI SBIC SBIS SBRC SBRS BST BLD SEC CLC SEI CLI BSET BCLR
Branch / Jump	RJMP JMP IJMP RCALL CALL ICALL RET RETI BREQ BRNE BRCS BRCC BRSH BRLO BRMI BRPL BRGE BRLT BRVS BRVC BRIE BRID CPSE
MCU Control	NOP SLEEP WDR BREAK

Chapters 5–8 cover Load/Store, Arithmetic, Branches, and Subroutines in depth. This chapter covers the core register-to-register operations that form the building blocks of everything else.

So think of this chapter as the grammar of AVR instructions. Later chapters combine these pieces into bigger programming patterns.

4.3 Data Transfer — Moving Values Around

4.3.1 MOV — Copy Register

`MOV Rd, Rr` ; `Rd ← Rr`

Copies the 8-bit value from `Rr` to `Rd`. Source unchanged. No flags affected. Works on any register `r0`–`r31`.

`MOV` is boring in the best possible way. It does exactly what it says, and because it does not touch flags, it is safe to use in the middle of code that depends on the current condition state.

```
ldi r16, 42
mov r0, r16      ; r0 = 42, r16 still = 42
mov r5, r0       ; r5 = 42 (MOV works on r0-r15 unlike LDI)
```

1 cycle. Does not affect `SREG`.

4.3.2 MOVW — Copy Register Pair

`MOVW Rd, Rr` ; `Rd+1:Rd ← Rr+1:Rr`

Copies two registers in one instruction. `Rd` and `Rr` must be even-numbered: `r0`, `r2`, `r4`, ... `r30`.

```
; Copy X pointer (r27:r26) into r1:r0
movw r0, r26      ; r1:r0 = r27:r26

; Copy Z into Y for a secondary pointer
movw r28, r30     ; r29:r28 = r31:r30 (Y = Z)
```

1 cycle. Does not affect `SREG`. Saves one instruction vs two `MOV`s.

Whenever you are copying a pointer pair or a 16-bit value, `MOVW` is worth remembering because it is both shorter and clearer than two separate moves.

4.3.3 LDI — Load Immediate

`LDI Rd, K` ; $Rd \leftarrow K$ ($K = 0..255$)

Load a constant into a register. **r16–r31 only**.

```
ldi r16, 0xFF ; OK
ldi r20, (1 << 4) ; OK — expression evaluated at assembly time
ldi r0, 42 ; ERROR: r0 not allowed
```

1 cycle. Does not affect SREG.

4.3.4 CLR and SER — Clear and Set All Bits

These are **pseudo-instructions** (GAS aliases, not separate opcodes):

```
clr r16 ; r16 = 0x00 — assembled as EOR r16, r16
ser r16 ; r16 = 0xFF — assembled as LDI r16, 0xFF (r16–r31 only)
```

CLR works on any register (r0–r31) because it uses EOR. SER only works on r16–r31 because it uses LDI.

This is a good example of how AVR pseudo-instructions work. They look like real instructions in source, but the assembler is translating them into more basic operations for you.

CLR sets Z=1, clears N, V, S. SER affects no flags.

4.4 Arithmetic Instructions

4.4.1 ADD — Add Without Carry

`ADD Rd, Rr` ; $Rd \leftarrow Rd + Rr$

Adds two registers, stores result in Rd. Affects: H, S, V, N, Z, C.

If you are coming from a higher-level language, read that carefully: the result replaces the original destination register. AVR arithmetic is almost always destructive to the destination operand.

```
ldi r16, 100
ldi r17, 50
add r16, r17 ; r16 = 150, C=0, Z=0, N=0
```

Unsigned overflow (result > 255): C=1, result wraps.

```
ldi r16, 200
ldi r17, 100
add r16, r17 ; 200+100=300 → r16=44, C=1
```

4.4.2 ADC — Add with Carry

`ADC Rd, Rr` ; $Rd \leftarrow Rd + Rr + C$

Identical to ADD but adds the carry flag too. Essential for **multi-byte arithmetic** — carry from the low byte automatically propagates into the high byte:

```

; 16-bit add: r17:r16 += r19:r18
add r16, r18          ; low bytes first – may set C
adc r17, r19          ; high bytes – C from previous ADD is included

; 24-bit add: r2:r1:r0 += r5:r4:r3
add r0, r3
adc r1, r4
adc r2, r5            ; each ADC picks up carry from the byte below

```

This low-byte-first pattern is one of the most important habits in AVR assembly. Once it becomes automatic, larger integer arithmetic stops feeling special.

4.4.3 SUB — Subtract Without Carry

```
SUB Rd, Rr          ; Rd ← Rd - Rr
```

Subtracts Rr from Rd. C=1 means borrow (unsigned underflow).

That “C means borrow” convention is easy to trip over because it feels backwards if you learned carry only from addition. On AVR subtraction, carry set means the subtraction had to borrow.

```

ldi r16, 10
ldi r17, 20
sub r16, r17          ; 10-20 = -10 → r16=246 (0xF6), C=1 (borrow)

```

4.4.4 SBC — Subtract with Carry (Borrow)

```
SBC Rd, Rr          ; Rd ← Rd - Rr - C
```

Like ADC for addition, SBC propagates borrow in multi-byte subtraction:

```

; 16-bit subtract: r17:r16 -= r19:r18
sub r16, r18          ; low bytes
sbc r17, r19          ; high bytes – includes borrow from low byte

```

Again, the pattern is the same as multi-byte addition: low byte first, then work upward using the carry/borrow chain.

4.4.5 SUBI — Subtract Immediate

```
SUBI Rd, K          ; Rd ← Rd - K (r16-r31 only)
```

Subtract a constant without needing a second register. **No ADDI exists** — use SUBI with the negated value, or ADIW/SBIW for 16-bit register pairs.

That missing ADDI surprises many beginners. It is not an omission in the book; it is a real feature of the AVR ISA.

```

ldi r16, 50
subi r16, 7           ; r16 = 43
subi r16, -3          ; r16 = 46 (subtract -3 = add 3)

```

4.4.6 SBCI — Subtract Immediate with Carry

```
SBCI Rd, K          ; Rd ← Rd - K - C (r16-r31 only)
```

Pair with SUBI for 16-bit immediate subtraction:

```

; r17:r16 -= 300 (0x012C)
subi r16, lo8(300)    ; r16 -= 0x2C
sbci r17, hi8(300)    ; r17 -= 0x01 - C

```

4.4.7 ADIW — Add Immediate to Word

ADIW Rd, K ; $Rd+1:Rd \leftarrow Rd+1:Rd + K$ ($K = 0..63$)

Adds a small constant to a 16-bit register pair in one instruction. Only works on r25:r24, r27:r26 (X), r29:r28 (Y), r31:r30 (Z).

```
adiw r24, 1 ; r25:r24 += 1
adiw XL, 4 ; X pointer += 4 (advance 4 bytes)
adiw ZL, 16 ; Z pointer += 16
```

2 cycles. Affects C, Z, N, V, S.

ADIW is especially nice for pointer arithmetic, because X, Y, and Z all live in the supported register-pair set.

4.4.8 SBIW — Subtract Immediate from Word

SBIW Rd, K ; $Rd+1:Rd \leftarrow Rd+1:Rd - K$ ($K = 0..63$)

Same register restrictions as ADIW. Used in delay loops (Chapter 1) and pointer adjustment.

```
sbiw r24, 1 ; r25:r24 -= 1 — sets Z when result = 0
sbiw ZL, 8 ; Z -= 8
```

4.4.9 INC and DEC — Increment / Decrement

INC Rd ; $Rd \leftarrow Rd + 1$
 DEC Rd ; $Rd \leftarrow Rd - 1$

Works on r0–r31. **Critical gotcha:** INC and DEC do not affect the carry flag (C). This catches people using DEC in multi-byte loops.

```
; This loop is WRONG for multi-byte work:
ldi r16, 0
ldi r17, 0
dec r16 ; r16 = 255, C unchanged — cannot chain with SBC!
```

```
; Correct 16-bit decrement: use SBIW
sbiw r16, 1 ; r17:r16 -= 1, sets Z correctly
```

Flags affected by INC/DEC: S, V, N, Z — but NOT C.

That is why INC and DEC are great for single-byte counters and a bad fit for arithmetic that needs carry propagation.

4.4.10 NEG — Two's Complement Negate

NEG Rd ; $Rd \leftarrow 0x00 - Rd$

Negates the register (flip all bits, then add 1). NEG 0 → 0 with C=0. All other values: C=1.

```
ldi r16, 10
neg r16 ; r16 = 246 (0xF6), which is -10 in two's complement
neg r16 ; r16 = 10 again
```

Useful for: converting subtraction direction, computing absolute value.

Whenever you see NEG, think “arithmetic sign change.” Whenever you see COM, think “bitwise inversion.” They are related, but not interchangeable.

4.5 Logic Instructions

4.5.1 AND / ANDI — Bitwise AND

```
AND Rd, Rr          ; Rd ← Rd AND Rr          (r0-r31)
ANDI Rd, K           ; Rd ← Rd AND K           (r16-r31 only)
```

Use AND to **clear** (mask off) specific bits:

```
ldi r16, 0xAB        ; 0b10101011
andi r16, 0x0F        ; r16 = 0x0B (0b00001011) – upper nibble cleared
andi r16, 0b1111110  ; clear bit 0
```

Affects: S, V(=0), N, Z. Clears V.

Bit masking is one of the most common operations in embedded work. In practice: AND clears bits, OR sets bits, and EOR toggles bits.

4.5.2 OR / ORI — Bitwise OR

```
OR Rd, Rr           ; Rd ← Rd OR Rr
ORI Rd, K            ; Rd ← Rd OR K            (r16-r31 only)
```

Use OR to **set** specific bits:

```
ldi r16, 0x2A
ori r16, 0x80        ; set bit 7 → r16 = 0xAA
ori r16, (1 << 3)    ; set bit 3
```

4.5.3 EOR — Bitwise XOR (Exclusive OR)

```
EOR Rd, Rr          ; Rd ← Rd XOR Rr
```

Use EOR to **toggle** bits or **invert** a value:

```
ldi r16, 0xFF
ldi r17, 0x0F
eor r16, r17         ; r16 = 0xF0 – lower nibble toggled
eor r16, r16         ; r16 = 0x00 – self-XOR always gives 0 (this is CLR)
```

4.5.4 COM — One's Complement (Bitwise NOT)

```
COM Rd               ; Rd ← 0xFF - Rd (flip every bit)
```

```
ldi r16, 0b10101010
com r16              ; r16 = 0b01010101
```

Always sets C=1. Note: COM ≠ NEG. COM flips bits; NEG negates (COM + 1).

4.5.5 TST — Test Register (Non-Destructive)

```
TST Rd               ; Rd AND Rd → set N, Z flags; Rd unchanged
```

Tests whether a register is zero or negative without modifying it:

```
lds r16, some_var
tst r16
breq var_is_zero     ; Z=1 if r16 was 0
brmi var_is_neg      ; N=1 if bit 7 was set (negative signed)
```

Assembled as `AND Rd, Rd`.

TST is one of those instructions that mostly exists to make your intent more obvious to a human reader.

4.6 Comparison Instructions

These set flags **without storing any result**. Always followed by a conditional branch.

This is the AVR version of "ask a question, then branch on the answer." The question is encoded in the compare; the answer appears in the flags.

4.6.1 CP — Compare

`CP Rd, Rr` ; $Rd - Rr \rightarrow$ flags only, result discarded

Exactly like SUB but throws away the result. Use before BREQ, BRNE, BRLO, BRGE, etc.

That makes CP conceptually simple: it performs a subtraction only for the purpose of updating flags.

```
ldi r16, 42
ldi r17, 100
cp r16, r17 ; flags set as if 42 - 100
brlo r16_smaller ; BRLO branches if C=1 (unsigned: Rd < Rr)
```

4.6.2 CPC — Compare with Carry

`CPC Rd, Rr` ; $Rd - Rr - C \rightarrow$ flags only

For multi-byte comparison (after comparing low bytes with CP):

```
; Compare r17:r16 with r19:r18 (16-bit unsigned)
cp r16, r18 ; compare low bytes
cpc r17, r19 ; compare high bytes including borrow
brlo less_than ; taken if r17:r16 < r19:r18 unsigned
```

4.6.3 CPI — Compare with Immediate

`CPI Rd, K` ; $Rd - K \rightarrow$ flags only (r16-r31 only)

```
ldi r16, 10
cpi r16, 10
breq exactly_ten ; taken
```

4.6.4 CPSE — Compare, Skip if Equal

`CPSE Rd, Rr` ; if $Rd == Rr$: skip next instruction

Skips exactly one instruction when equal. Used for compact conditional code:

```
cpse r16, r17
rjmp not_equal ; skipped if r16 == r17
; falls through here if equal
```


4.7 Conditional Branches — Full Table

Every branch instruction tests one or more SREG flags. The mnemonic tells you what condition causes the branch.

Once you stop reading branch mnemonics as magic words and start reading them as “test this flag pattern,” control flow becomes much easier to reason about.

Mnemonic	Full name	Condition	Flag(s)
BREQ	Branch if Equal	$Z = 1$	Z
BRNE	Branch if Not Equal	$Z = 0$	Z
BRCS	Branch if Carry Set	$C = 1$	C
BRCC	Branch if Carry Cleared	$C = 0$	C
BRSH	Branch if Same or Higher	$C = 0$	C $\leftarrow$ unsigned $\geq$
BRLO	Branch if Lower	$C = 1$	C $\leftarrow$ unsigned $<$
BRMI	Branch if Minus	$N = 1$	N
BRPL	Branch if Plus	$N = 0$	N
BRGE	Branch if $\geq$ (signed)	$S = 0$	$N \oplus V \leftarrow$ signed $\geq$
BRLT	Branch if $<$ (signed)	$S = 1$	$N \oplus V \leftarrow$ signed $<$
BRVS	Branch if Overflow Set	$V = 1$	V
BRVC	Branch if Overflow Clear	$V = 0$	V
BRIE	Branch if Interrupts En.	$I = 1$	I
BRID	Branch if Interrupts Dis.	$I = 0$	I
BRHS	Branch if Half-carry Set	$H = 1$	H
BRHC	Branch if Half-carry Clr	$H = 0$	H
BRTS	Branch if T Set	$T = 1$	T
BRTC	Branch if T Clear	$T = 0$	T

All branch instructions: **± 63 words** range. If the target is farther, use JMP or RJMP with an inverted condition:

```
; Branch to far_label if r16 == r17 – but far_label > 63 words away:
cp    r16, r17
brne skip           ; if NOT equal, skip the jump
jmp  far_label
skip:
```

This branch-range limit is another place where instruction encoding leaks into everyday programming. The branch is short because the displacement field is short.

4.8 Unsigned vs Signed Comparisons

This trips up almost every beginner. Unsigned and signed comparisons use different branch instructions after the same CP.

```
ldi r16, 0xF0      ; = 240 unsigned, or -16 signed
ldi r17, 0x10      ; = 16 unsigned, or 16 signed

cp    r16, r17     ; 0xF0 - 0x10 = 0xE0. C=0 (no borrow), N=1, V=0, S=1
```

```
brlo below_u    ; C=1? NO → not taken.  Correct: 240 > 16 unsigned.
brlt below_s    ; S=1? YES → taken.    Correct: -16 < 16 signed.
```

Rule: **BRLO/BRSH for unsigned. BRLT/BRGE for signed.**

If you take only one rule from this section, take that one. The CPU does not know whether you intended a byte to mean 240 or -16. Your branch choice tells it how to interpret the flags.

4.9 Flag Effects Summary

The table below covers every instruction in this chapter. “—” means the flag is not affected.

Instruction	C	Z	N	V	S	H	
ADD	•	•	•	•	•	•	
ADC	•	•	•	•	•	•	
SUB	•	•	•	•	•	•	
SBC	•	•	•	•	•	•	
SUBI	•	•	•	•	•	•	
SBCI	•	•	•	•	•	•	
ADIW	•	•	•	•	•	—	
SBIW	•	•	•	•	•	—	
INC	—	•	•	•	•	—	← C NOT affected
DEC	—	•	•	•	•	—	← C NOT affected
NEG	•	•	•	•	•	•	
AND / ANDI	0	•	•	0	•	—	← V cleared
OR / ORI	0	•	•	0	•	—	← V cleared
EOR	0	•	•	0	•	—	← V cleared
COM	1	•	•	0	•	—	← C always 1, V cleared
TST	0	•	•	0	•	—	
CP / CPI	•	•	•	•	•	•	
CPC	•	•	•	•	•	•	
MOV/MOVW/LDI	—	—	—	—	—	—	
CLR	0	1	0	0	0	—	

4.10 Worked Example: Compare, Clamp, Negate

Goal: read an 8-bit value, clamp it to 0–100 range, then negate it if a flag register bit is set.

This example matters because it combines several ideas you will use constantly: compare, branch, clamp by assignment, test a control bit, and optionally apply an arithmetic transform.

```
/*
 * clamp_and_negate:
 *   Input:  r24 = raw value (unsigned), r25 = flags (bit 0 = negate)
 *   Output: r24 = clamped value, optionally negated
 *   Clobbers: r16
 */
clamp_and_negate:
    /* Step 1: clamp to 100 max */
    cpi    r24, 100          ; compare r24 with 100
```

```

    brlo    clamped                ; if r24 < 100 (unsigned), skip clamp
    ldi     r24, 100                ; else clamp to 100
clamped:

/* Step 2: clamp to 0 min – already unsigned so r24 >= 0 always.
 * If the raw value were signed we'd add BRGE check here. */

/* Step 3: negate if bit 0 of r25 is set */
sbrs    r25, 0                    ; Skip if Bit in Register Set (bit 0 of r25)
ret                                           ; bit 0 clear – return as-is
neg     r24                        ; negate: r24 = 0 - r24
ret

```

New instructions introduced:

SBRs Rr, b Skip next instruction if bit b of Rr is set. 1 cycle (not skipped), 2 cycles (skipped over 1-word instr), 3 cycles (skipped over 2-word instr). Does not affect SREG.

Trace through the example:

Call: r24=150, r25=0b00000001 (negate flag set)

```

cpi r24, 100    → 150-100=50, C=0, Z=0
brlo clamped    → C=0: NOT taken (150 >= 100, no branch)
ldi r24, 100    → r24 = 100
clamped:
sbrs r25, 0     → bit 0 of r25 is 1: SKIP next instr
                  (ret is skipped)
neg r24         → r24 = 0 - 100 = 156 (0x9C = -100 signed)
ret

```

Walking through examples like this is how you build intuition. Do not just read the final code; read the flags and the decision points as they evolve.

4.11 Operand Restrictions Cheat-Sheet

This is worth turning into instinct:

Instruction	Rd range	Rr range	Immediate range
MOV	r0-r31	r0-r31	–
MOVW	r0,r2..r30	r0,r2..r30	– (even only)
LDI	r16-r31	–	0-255
CLR	r0-r31	–	–
SER	r16-r31	–	–
ADD/ADC	r0-r31	r0-r31	–
SUB/SBC	r0-r31	r0-r31	–
SUBI/SBCI	r16-r31	–	0-255
ADIW/SBIW	r24,X,Y,Z	–	0-63
INC/DEC	r0-r31	–	–
NEG/COM	r0-r31	–	–
AND/OR/EOR	r0-r31	r0-r31	–
ANDI/ORI	r16-r31	–	0-255
CP	r0-r31	r0-r31	–
CPI	r16-r31	–	0-255

CPC	r0-r31	r0-r31	—
CPSE	r0-r31	r0-r31	—
SBRS/SBRC	r0-r31	—	bit 0-7

4.12 Summary

Instructions: 1 or 2 words (16 or 32 bits).

Two-word: LDS, STS, JMP, CALL — need full 16-bit address.

Register rules:

LDI, SUBI, SBCI, ANDI, ORI, SER, CPI → r16-r31 only.

ADIW, SBIW → r24, X(r26), Y(r28), Z(r30) only.

Everything else → r0-r31.

Carry chains: ADD+ADC, SUB+SBC, SUBI+SBCI — always low byte first.

INC/DEC do NOT touch C — never use for multi-byte counters.

NEG ≠ COM: NEG = two's complement, COM = bitwise invert.

Comparisons:

Unsigned: BRLO (<), BRSH (≥), BRCS/BRCC

Signed: BRLT (<), BRGE (≥), BRMI/BRPL

Branch range: ±63 words. Use JMP for farther targets.

4.13 Exercises

1. Add two 24-bit numbers stored in r2:r1:r0 and r5:r4:r3. Store the result in r2:r1:r0. What happens to the carry if the result exceeds 24 bits?
 2. Write a sequence to compute $r16 = r16 * 4$ without using MUL. (Hint: shifting left by 1 is the same as multiplying by 2.) Use only ADD instructions.
 3. Write a loop that counts from 0 to 255 in r16, then wraps and repeats forever. Use INC and BREQ. Why is BREQ the right branch here?
 4. Given $r16 = 0xAB$ and $r17 = 0x3C$, predict the result and all six flag values (C, Z, N, V, S, H) after each operation — before assembling:
 - ADD r16, r17
 - SUB r16, r17 (fresh r16=0xAB each time)
 - AND r16, r17
 - EOR r16, r17
 5. Write clamp_and_negate to also handle the case where r24 starts as 0 and the negate flag is set. What does NEG 0 produce and is that correct for your use case?
 6. Why does SUBI r16, -1 add 1 to r16? What is the bit pattern of -1 as an 8-bit two's complement number, and what does subtracting it do?
-

Next: Chapter 5 — Load and Store

Chapter 5


Load and Store

Every useful program moves data. In AVR assembly, **every data movement is explicit**: you choose the exact instruction that moves each byte, the exact addressing mode, and often the exact pointer update behaviour. This chapter covers the main ways to get data into registers, get it back out to memory, and read constants from flash on the ATtiny3217.

5.1 The Memory-Register Model

The AVR CPU can only operate on data that is in **registers** (r0–r31). To work on a byte in SRAM, you must load it into a register first, operate, then store it back. There is no “operate directly on memory” instruction.

That rule is so central that it is worth repeating in plain language: if the value is not in a register yet, the ALU cannot touch it.



Every instruction in this chapter is either a load (memory → register) or a store (register → memory).

Once you see that pattern, a lot of AVR code becomes easier to read. A routine is often just alternating phases: bring bytes into registers, do work there, write results back out.

5.2 Addressing Modes

The AVR supports five ways to specify a memory address:

Mode	Syntax	Address source	Instruction
Direct	LDS Rd, k	16-bit constant k in instruction	LDS / STS

Mode	Syntax	Address source	Instruction
Indirect	LD Rd, X	Value in X, Y, or Z pointer pair	LD / ST
Post-increment	LD Rd, X+	X (used), then X += 1	LD / ST
Pre-decrement	LD Rd, -X	X -= 1, then X (used)	LD / ST
Indirect + displacement	LDD Rd, Y+q	Y + constant q (0–63)	LDD / STD

These are not five completely different ideas. They are variations on one theme: where does the address come from, and should the pointer change as a side effect?

5.3 Direct Addressing: LDS and STS

```
LDS Rd, k      ; Rd ← SRAM[k]      (k = 16-bit address)
STS k, Rr      ; SRAM[k] ← Rr
```

The address is encoded directly in the instruction as a second 16-bit word. Both are **32-bit (two-word) instructions**.

That makes LDS/STS very convenient to write, but relatively expensive in both flash space and cycles.

```
.section .bss
my_var: .skip 1      ; reserve one byte in SRAM

.section .text
ldi r16, 0xAB
sts my_var, r16      ; write 0xAB to my_var's SRAM address
lds r17, my_var      ; r17 = 0xAB
```

Cycle costs:

Instruction	Cycles
LDS Rd, k	3
STS k, Rr	2

When to use LDS/STS: accessing a single named variable. For loops over arrays, indirect addressing (Section 5.4) is faster and smaller.

So LDS/STS are great for “go touch this exact place once.” They are a poor fit for repeated traversal through memory.

5.3.1 I/O Registers Above 0x3F

STS/LDS are the only way to reach peripheral registers with addresses above 0x3F. This includes all the PORT, timer, and USART registers on ATtiny3217:

```
ldi r16, (1 << 4)
sts PORTB_DIRSET, r16      ; PORTB_DIRSET = 0x0421, above 0x3F → STS
```

For addresses 0x00–0x3F, always prefer IN/OUT (1 cycle, 1 word) over LDS/STS (3/2 cycles, 2 words).

That is one of the easiest AVR optimisations to make because it is mechanical: if the register is in low I/O space, use the I/O instructions.

5.4 Pointer Registers: X, Y, Z

The three pointer register pairs are 16-bit registers built from adjacent general-purpose registers:

```
X = r27 : r26      XH : XL
Y = r29 : r28      YH : YL
Z = r31 : r30      ZH : ZL
```

To load an address into a pointer pair, use two LDI instructions:

```
; Point X at buf (a label in .bss)
ldi XL, lo8(buf)      ; XL = r26 = low byte of buf's address
ldi XH, hi8(buf)      ; XH = r27 = high byte
```

Now X holds the 16-bit address of buf and can be used with LD/ST.

This is the beginning of pointer-based programming in AVR. Instead of naming an address directly in every instruction, you put the address in X, Y, or Z once and then operate through that pointer.

5.5 Indirect Addressing: LD and ST

5.5.1 Basic Indirect — No Change to Pointer

```
LD Rd, X      ; Rd ← SRAM[X]      (X unchanged)
LD Rd, Y      ; Rd ← SRAM[Y]
LD Rd, Z      ; Rd ← SRAM[Z]
ST X, Rr      ; SRAM[X] ← Rr      (X unchanged)
ST Y, Rr
ST Z, Rr
```

2 cycles each. Use when you want to read/write the same address multiple times without moving the pointer.

This is the pointer equivalent of “stay where you are and keep working here.”

```
; Read, modify, write a single SRAM byte
ldi XL, lo8(my_var)
ldi XH, hi8(my_var)
ld r16, X      ; r16 = my_var
ori r16, 0x80  ; set bit 7
st X, r16      ; write back — X still points to my_var
```

5.5.2 Post-Increment — Walk Forward Through Memory

```
LD Rd, X+      ; Rd ← SRAM[X], then X += 1
ST X+, Rr      ; SRAM[X] ← Rr, then X += 1
```

The increment happens **after** the access. X is ready to point to the next byte. Perfect for forward array traversal:

This is one of the most natural addressing modes on AVR. It lets the memory access and the pointer update happen in the same instruction.

```
; Zero-fill buf (8 bytes)
ldi XL, lo8(buf)
ldi XH, hi8(buf)
ldi r16, 0
```

```
ldi r18, 8
1:
    st  X+, r16      ; write 0, advance X
    dec r18
    brne 1b
```

5.5.3 Pre-Decrement — Walk Backward Through Memory

```
LD Rd, -X      ; X -= 1, then Rd ← SRAM[X]
ST -X, Rr      ; X -= 1, then SRAM[X] ← Rr
```

The decrement happens **before** the access. To walk an array backwards, point X one past the last element:

Pre-decrement often feels odd the first time you see it, but it matches how stacks and reverse traversals naturally work.

```
; Sum array backwards into r17
ldi XL, lo8(buf + 8) ; one past end
ldi XH, hi8(buf + 8)
clr r17
ldi r18, 8
1:
    ld  r16, -X      ; X--, then load
    add r17, r16
    dec r18
    brne 1b
```

5.5.4 All Three Pointers, All Three Modes

X, Y, and Z all support LD/ST, LD+/ST+, and LD-/ST-. The only difference is which register pair holds the address:

So the real question is rarely “which instruction family do I need?” More often it is “which pointer register should own this job?”

```
ld r16, X+      ; load via X, post-increment
ld r16, Y+      ; load via Y, post-increment
ld r16, Z+      ; load via Z, post-increment — also used for LPM
```

5.6 Displacement Addressing: LDD and STD

Y and Z support a **constant displacement** added to the pointer without modifying it. X does **not** support displacement.

```
LDD Rd, Y+q      ; Rd ← SRAM[Y + q]    (Y unchanged, q = 0–63)
LDD Rd, Z+q      ; Rd ← SRAM[Z + q]
STD Y+q, Rr      ; SRAM[Y + q] ← Rr
STD Z+q, Rr
```

2 cycles. Displacement q must be a constant 0–63 — it cannot be a register.

Displacement mode is useful because it lets one pointer stay parked at a base address while you reach nearby fields by constant offset.

Primary use: struct member access. Point Y at the base of a struct, then access each member by offset:


```

; Struct layout in SRAM:
;   offset 0: status (1 byte)
;   offset 1: count  (1 byte)
;   offset 2: value  (2 bytes, little-endian)
;   offset 4: flags  (1 byte)

.equ STATUS_OFF, 0
.equ COUNT_OFF, 1
.equ VALUE_OFF, 2
.equ FLAGS_OFF, 4

; Point Y at the struct
ldi YL, lo8(my_struct)
ldi YH, hi8(my_struct)

; Read all fields without moving Y
ldd r16, Y+STATUS_OFF    ; r16 = status
ldd r17, Y+COUNT_OFF    ; r17 = count
ldd r18, Y+VALUE_OFF     ; r18 = value low byte
ldd r19, Y+VALUE_OFF+1   ; r19 = value high byte
ldd r20, Y+FLAGS_OFF     ; r20 = flags

```

Secondary use: stack frame pointer. The compiler uses Y as the frame pointer, accessing local variables via LDD r, Y+offset. Chapter 8 covers this in the context of subroutines.

In other words, LDD/STD are the AVR equivalent of “base pointer + fixed offset” addressing.

5.7 Cycle and Size Comparison

Mode	Instruction	Cycles	Words	Pointer moves?
Direct	LDS / STS	3 / 2	2	No
Indirect	LD / ST	2	1	No
Post-increment	LD+ / ST+	2	1	+1 after
Pre-decrement	LD- / ST-	2	1	-1 before
Displacement	LDD / STD	2	1	No (Y/Z only)
I/O register	IN / OUT	1	1	N/A (addr 0-3F)

Rule of thumb:

- **Named variable, one access:** LDS/STS
- **Named variable, repeated access:** load address into pointer, then LD/ST
- **Array traversal:** LD+/ST+ or LD-/ST-
- **Struct fields:** LDD/STD with Y as base
- **I/O registers $\leq 0x3F$:** IN/OUT

Those rules are worth internalising because they cover most real choices in ordinary firmware.

5.8 Reading Flash: LPM

On all AVR devices, LPM reads one byte from the **program bus** (flash) into a register. Z holds the byte address of the flash location to read.

This is the classic AVR way to read constants stored in flash. If you learned AVR on older parts, LPM is probably the method you saw first.

```
LPM Rd, Z          ; Rd ← Flash[Z]      (Z unchanged)
LPM Rd, Z+         ; Rd ← Flash[Z], Z++ (post-increment)
LPM                ; r0 ← Flash[Z]      (deprecated form)
```

3 cycles. Z must be the byte address (not word address).

That byte-address detail matters because some other AVR concepts, like jump targets, use word-oriented program addressing instead.

```
.section .rodata
lookup:
    .byte 0, 1, 1, 2, 3, 5, 8, 13 ; Fibonacci

.section .text

    ; Load lookup[3] into r16
    ldi ZL, lo8(lookup + 3)
    ldi ZH, hi8(lookup + 3)
    lpm r16, Z                ; r16 = 2 (lookup[3])
```

Copying a flash table to SRAM with LPM:

```
    ldi ZL, lo8(lookup)
    ldi ZH, hi8(lookup)
    ldi XL, lo8(dest_buf)
    ldi XH, hi8(dest_buf)
    ldi r18, 8

copy:
    lpm r16, Z+                ; load from flash, Z++
    st X+, r16                 ; store to SRAM, X++
    dec r18
    brne copy
```

5.9 Reading Flash: Mapped Access (AVRxt Only)

On AVRxt (ATtiny3217), flash is mapped into the **data address space** at offset `MAPPED_PROGMEM_START = 0x8000`. You can read it with ordinary LD instructions — no LPM needed.

This is one of the nicest quality-of-life improvements on the ATtiny3217. It lets flash reads look much more like normal memory reads.

```
LD Rd, Z          ; Rd ← DataMem[Z]
```

If $Z \geq 0x8000$, the data bus reads from flash. If $Z < 0x8000$, it reads SRAM or I/O registers. The CPU doesn't care which — the address selects the bus.

That means one instruction family, LD, can read from very different physical resources depending on the address you place in the pointer.

```
.section .rodata
lookup:
    .byte 0, 1, 1, 2, 3, 5, 8, 13
```

```
.section .text

; Load lookup[3] into r16 via mapped flash
ldi ZL, lo8(lookup + MAPPED_PROGMEM_START + 3)
ldi ZH, hi8(lookup + MAPPED_PROGMEM_START + 3)
ld r16, Z ; r16 = 2
```

Copying with mapped LD — same loop as SRAM copy:

```
ldi ZL, lo8(lookup + MAPPED_PROGMEM_START)
ldi ZH, hi8(lookup + MAPPED_PROGMEM_START)
ldi XL, lo8(dest_buf)
ldi XH, hi8(dest_buf)
ldi r18, 8
```

`copy_mapped:`

```
ld r16, Z+ ; load from mapped flash, Z++
st X+, r16 ; store to SRAM, X++
dec r18
brne copy_mapped
```

5.9.1 LPM vs Mapped LD — Which to Use?

	LPM	Mapped LD
Available on	All AVR	AVRxt only (tinyAVR 1-series, megaAVR 0-series)
Cycles	3	2
Words	1	1
Z value	Flash byte address	Flash byte address + 0x8000
Pointer modes	Z, Z+ only	X, Y, Z — all modes including displacement
Flash > 32KB	Needs RAMPZ	Not applicable (mapped only covers 32KB)

On ATtiny3217 with ≤ 32 KB flash: **prefer mapped LD**. It is faster, uses all pointer modes, and reads flash and SRAM the same way — simpler mental model.

LPM is still worth knowing because it is portable across AVR families, but for this chip, mapped LD is usually the cleaner choice.

5.10 Pointer Arithmetic

5.10.1 Advancing a Pointer by N Bytes

Post-increment advances by 1 per cycle. To skip N bytes, use ADIW:

```
adiw ZL, 4 ; Z += 4 (skip 4 bytes)
adiw XL, 16 ; X += 16
```

ADIW only handles values 0–63. For larger steps, use ADD/ADC on the pair directly:

So pointer arithmetic on AVR comes in two flavors: compact/specialized (ADIW/SBIW) and general/manual (ADD/ADC).

```
; Z += 200
ldi r16, 200
```

```
add ZL, r16
ldi r16, 0
adc ZH, r16      ; propagate carry into high byte
```

Or use r1 (which the ABI keeps at 0) for the carry:

```
ldi r16, 200
add ZL, r16
adc ZH, r1       ; r1 = 0 by convention – just propagates carry
```

5.10.2 Computing an Array Element Address

```
; Address of array[i]: base + i * element_size (element_size = 1 here)
; r24 = index, Z = array base
add ZL, r24      ; Z += index (1-byte elements)
adc ZH, r1       ; carry
ld r16, Z        ; load array[i]
```

For 2-byte elements:

```
; i * 2 = i << 1
lsl r24          ; r24 *= 2 (left shift)
rol r25          ; shift carry into r25 (if index > 127)
add ZL, r24
adc ZH, r25
```

This is the same reasoning as in multi-byte arithmetic: build the low byte result first, then propagate carry into the high byte.

5.11 Atomicity on ATtiny3217

Some AVR families implement XCH, LAC, LAS, and LAT, but the ATtiny3217's AVRxt core does **not**. If you need interrupt-safe bit updates on GPIO, use the dedicated peripheral registers such as PORTx_OUTSET, PORTx_OUTCLR, and PORTx_OUTTGL instead of a read-modify-write sequence.

This matters because “load-modify-store” is not automatically safe when an ISR or another hardware path can touch the same state concurrently.

```
ldi r16, (1 << 3)
sts PORTA_OUTSET, r16 ; set PA3 without reading PORTA_OUT first
sts PORTA_OUTCLR, r16 ; clear PA3 atomically
sts PORTA_OUTTGL, r16 ; toggle PA3 atomically
```

For ordinary SRAM bytes shared with an ISR, there is no single instruction on this MCU that performs an atomic arbitrary read-modify-write. Use a critical section (IN/CLI/OUT SREG) when the update must not be interrupted.

So the practical split is:

- for GPIO, prefer the hardware set/clear/toggle registers;
- for shared SRAM state, protect the update with interrupt control when needed.

5.12 Full Addressing Mode Reference

Instruction	Addr calc	Pointer after	Cycles	Words
-------------	-----------	---------------	--------	-------

LDS	Rd, k	k	—	3	2
STS	k, Rr	k	—	2	2
LD	Rd, X	X	unchanged	2	1
LD	Rd, X+	X	X + 1	2	1
LD	Rd, -X	X - 1	X - 1	2	1
ST	X, Rr	X	unchanged	2	1
ST	X+, Rr	X	X + 1	2	1
ST	-X, Rr	X - 1	X - 1	2	1
LD	Rd, Y	Y	unchanged	2	1
LD	Rd, Y+	Y	Y + 1	2	1
LD	Rd, -Y	Y - 1	Y - 1	2	1
LDD	Rd, Y+q	Y + q	unchanged	2	1
ST	Y, Rr	Y	unchanged	2	1
ST	Y+, Rr	Y	Y + 1	2	1
ST	-Y, Rr	Y - 1	Y - 1	2	1
STD	Y+q, Rr	Y + q	unchanged	2	1
LD	Rd, Z	Z	unchanged	2	1
LD	Rd, Z+	Z	Z + 1	2	1
LD	Rd, -Z	Z - 1	Z - 1	2	1
LDD	Rd, Z+q	Z + q	unchanged	2	1
ST	Z, Rr	Z	unchanged	2	1
ST	Z+, Rr	Z	Z + 1	2	1
ST	-Z, Rr	Z - 1	Z - 1	2	1
STD	Z+q, Rr	Z + q	unchanged	2	1
LPM	Rd, Z	Flash[Z]	unchanged	3	1
LPM	Rd, Z+	Flash[Z]	Z + 1	3	1
IN	Rd, A	I/O[A] (0-3F)	—	1	1
OUT	A, Rr	I/O[A] (0-3F)	—	1	1
PUSH	Rr	SRAM[SP-1]	SP - 1	2	1
POP	Rd	SRAM[SP]	SP + 1	2	1

5.13 Worked Example: Reverse a String In-Place

Reverse an ASCII string stored in SRAM. Classic two-pointer approach.

This example is useful because it combines several addressing modes in one small routine: forward scan with Z+, fixed access through X and Z, then converging pointers for the swap loop.

```

/*
 * str_reverse – reverse a null-terminated string in SRAM in-place.
 * Input:  X = pointer to first byte of string
 * Clobbers: X, Z, r16, r17, r18
 */
str_reverse:
    ; Find the end: walk Z to the null terminator
    movw ZL, XL          ; Z = X (copy pointer)

```

```

1:
    ld  r16, Z+           ; load byte, Z++
    tst r16               ; null?
    brne 1b              ; no – keep walking

    ; Z is now one past the null. Step back twice: past null, to last char.
    sbiw ZL, 2           ; Z now points to last character

    ; Now swap X[front] and Z[back], moving inward
2:
    cp   XL, ZL           ; has front crossed back? (low bytes)
    cpc  XH, ZH           ; (high bytes)
    brsh done            ; BRSH: unsigned >=, front >= back → done

    ld  r16, X            ; r16 = front char
    ld  r17, Z            ; r17 = back char
    st  X+, r17           ; store back char at front, advance front
    st  Z, r16            ; store front char at back
    sbiw ZL, 1           ; retreat back pointer
    rjmp 2b

done:
    ret

```

Trace with "AVR\0" at address 0x3810:

Initial: X = 0x3810 ('A'), Z = 0x3812 ('R')

Iter 1: swap 'A' and 'R' → X=0x3811, Z=0x3811

Iter 2: XL=0x11 >= ZL=0x11 → BRSH taken → done

Result in SRAM: "RVA\0" OK

Notice how little arithmetic the routine needs. Most of the real work is being done by the addressing modes themselves.

5.14 Summary

LDS/STS – direct, 16-bit address in instruction, 3/2 cycles, 2 words.
Use for named variables, single accesses.

LD/ST – indirect via X, Y, or Z. 2 cycles, 1 word.
no mode – pointer unchanged
X+/Y+/Z+ – post-increment (forward traversal)
-X/-Y/-Z – pre-decrement (backward traversal)

LDD/STD – Y or Z plus constant offset 0-63. Struct member access.
X does NOT support displacement.

LPM – flash read via program bus. Z = byte address. 3 cycles.
Mapped LD – AVRxt: LD with Z = flash addr + 0x8000. 2 cycles. Preferred.

ADIW/SBIW – adjust pointer by constant 0-63 in 2 cycles.

ADD+ADC – adjust pointer by any 16-bit value.

ATtiny3217 has no XCH/LAC/LAS/LAT. Use PORTx set/clear/toggle registers for GPIO, and critical sections for shared SRAM updates.

5.15 Exercises

1. Write a subroutine `memset(X, val, n)` that fills `n` bytes starting at address `X` with value `val` (in `r16`). Use `ST X+` in the loop.
 2. Write `memcpy(X, Z, n)` that copies `n` bytes from `Z` to `X`. Both source and destination are SRAM. Then modify it so the source is flash (use mapped LD for ATtiny3217).
 3. Given the struct:


```
offset 0: type    (1 byte)
offset 1: length  (1 byte)
offset 2: data    (4 bytes)
```

Write code using LDD/STD with `Y` to read `type`, increment `length`, and zero out the data field — all without moving `Y`.
 4. What is the maximum array size (in elements, 1 byte each) you can access with LDD `Y+q`? What happens if you need element 64? Write code that handles this using ADIW and basic LD.
 5. Assemble `loadstore.S` and run `avr-objdump -s loadstore.elf`. Find the `.rodata` section. Verify the bytes match 1, 2, 4, 8, 16, 32, 64, 128. Then find the address of `buf_lpm` in the `.bss` section — what address does it get?
 6. Why does the `str_reverse` subroutine use CP+CPC instead of just CP to compare `X` and `Z`? What would go wrong with only CP `XL, ZL`?
-

Next: Chapter 6 — Arithmetic and Logic (16-bit and beyond)

Chapter 6

Arithmetic and Logic: 16-bit and Beyond

Chapter 4 introduced 8-bit arithmetic. Real programs quickly outgrow that. Addresses are 16-bit, timers are often 16-bit, counters spill past 255, and scaled sensor values rarely fit in a single byte for long. This chapter shows how AVR handles arithmetic once one byte is no longer enough.

The key idea is simple: AVR does not magically become a 16-bit machine when you want a 16-bit result. You build larger arithmetic explicitly out of 8-bit pieces, and the carry flag is what ties those pieces together.

6.1 Multi-Byte Addition and Subtraction

The AVR has no 16-bit ADD instruction. You build it from two 8-bit instructions using the carry flag as a bridge between bytes.

That sounds more awkward than it is. Once you trust the carry chain, multi-byte arithmetic becomes a repeatable pattern rather than a special trick.

6.1.1 The Carry Chain Pattern

Always start with the least-significant byte. Never with the most-significant.

```
; 16-bit add: r25:r24 += r23:r22
add r24, r22      ; low bytes – may produce carry
adc r25, r23      ; high bytes – ADC includes carry from ADD

; 16-bit subtract: r25:r24 -= r23:r22
sub r24, r22      ; low bytes – may produce borrow (C=1)
sbc r25, r23      ; high bytes – SBC subtracts borrow
```

The carry flag is the pipeline between bytes. Break the chain and the high byte has the wrong value.

That is why instruction order matters here. Arithmetic is not just about the operands; it is also about preserving the flag state from one byte to the next.

6.1.2 24-bit and 32-bit

The pattern extends to any width:

```
; 32-bit add: r3:r2:r1:r0 += r7:r6:r5:r4
add r0, r4
adc r1, r5
```



```

adc r2, r6
adc r3, r7

; 32-bit subtract: r3:r2:r1:r0 -= r7:r6:r5:r4
sub r0, r4
sbc r1, r5
sbc r2, r6
sbc r3, r7

```

After the chain: if C=1, unsigned overflow occurred (result too large for N bytes). If V=1, signed overflow occurred.

So after a multi-byte operation, the flags still matter. The CPU is telling you whether the final N-byte result wrapped or overflowed in the signed sense.

6.1.3 Adding a Small Constant to a 16-bit Register Pair

ADIW handles constants 0–63 in 2 cycles:

```

adiw r24, 1      ; r25:r24 += 1  - 2 cycles, 1 word
adiw ZL, 16      ; Z += 16

```

For constants 64–255, use SUBI+SBCI with negated values — because there is no ADDI:

```

; r25:r24 += 100
subi r24, lo8(-100)    ; subtract -100 = add 100
sbc r25, hi8(-100)    ; propagate carry

```

For constants > 255, load into a register pair and use ADD+ADC:

```

; r25:r24 += 1000
ldi r22, lo8(1000)
ldi r23, hi8(1000)
add r24, r22
adc r25, r23

```

This gives you three practical tools for “add a constant”:

- ADIW when the constant is small and the register pair is supported
- SUBI/SBCI when the constant fits in one byte but ADIW is not enough
- ADD/ADC when you want the general case

6.1.4 Negating a 16-bit Value

Two's complement negate: invert all bits, add 1.

```

; negate r25:r24
com r24      ; ~low
com r25      ; ~high
adiw r24, 1  ; + 1 (or: ldi r16,1 / add r24,r16 / adc r25,r1)

```

Or equivalently using NEG + SBC:

```

neg r24      ; low = 0 - low (sets C if r24 was non-zero)
neg r25      ; high = 0 - high
sbc r25, 0    ; subtract borrow - corrects high byte

```

The second version is more common in compiler output.

The reason is that it fits naturally into the carry/borrow model the CPU already uses for subtraction.

6.2 Shift and Rotate Instructions

Shifts move bits left or right. The bit shifted out goes into the carry flag. Rotates put the old carry flag back in on the vacated side.

If addition and subtraction are about numeric value, shifts and rotates are about bit position. They are the bridge between arithmetic and bit-level work.

LSL Rd – Logical Shift Left

C ← [7][6][5][4][3][2][1][0] ← 0

Rd *= 2 (unsigned). Old bit 7 → C.

LSR Rd – Logical Shift Right

0 → [7][6][5][4][3][2][1][0] → C

Rd /= 2 (unsigned). Old bit 0 → C. Bit 7 forced to 0.

ASR Rd – Arithmetic Shift Right

b7 → [7][6][5][4][3][2][1][0] → C

Rd /= 2 (signed). Bit 7 preserved (sign extension).

ROL Rd – Rotate Left through Carry

C ← [7][6][5][4][3][2][1][0] ← C(old)

Like LSL but old C enters at bit 0.

ROR Rd – Rotate Right through Carry

C(old) → [7][6][5][4][3][2][1][0] → C

Like LSR but old C enters at bit 7.

All shift/rotate instructions: **1 cycle**. They affect C, Z, N, V, S, H and T are not affected.

That makes them cheap enough to use constantly, especially for scaling by powers of two or for moving carries across byte boundaries.

6.2.1 16-bit Shift Left (Multiply by 2)

```
lsl r24      ; low byte: bit 7 → C
rol r25      ; high byte: C → bit 0 (carry from low byte propagates)
```

6.2.2 16-bit Shift Right, Logical (Unsigned Divide by 2)

```
lsr r25      ; high byte: bit 0 → C (MSB filled with 0)
ror r24      ; low byte: C → bit 7
```

6.2.3 16-bit Shift Right, Arithmetic (Signed Divide by 2)

```
asr r25      ; high byte: sign bit preserved, bit 0 → C
ror r24      ; low byte: C → bit 7
```

6.2.4 Shifting by N Bits

GAS does not have a multi-bit shift instruction (unlike x86's SHL r, cl). Repeat for each bit:

```
; r16 <= 3 (multiply by 8)
lsl r16
lsl r16
lsl r16
```

For large shifts on 8-bit values, SWAP is faster when shifting by 4:

```
; r16 <=> 4
swap r16      ; swap nibbles: bits 7:4 ↔ bits 3:0
andi r16, 0xF0 ; zero lower nibble (was upper before swap)

; r16 >=> 4 (logical)
swap r16
andi r16, 0x0F ; zero upper nibble
```

This is a good reminder that on AVR, repeated simple instructions often replace one “fancier” instruction from a larger ISA.

6.2.5 Extracting a Bit Field

Example: extract bits 5:3 of r16 into bits 2:0 of r17:

```
mov r17, r16
andi r17, 0b00111000 ; mask bits 5:3
lsr r17               ; shift right 3 times
lsr r17
lsr r17               ; r17 = bits 5:3 in positions 2:0
```

6.3 SWAP — Swap Nibbles

SWAP Rd ; bits 7:4 ↔ bits 3:0

1 cycle. Does not affect any flags. Pure bit rearrangement.

```
ldi r16, 0xAB
swap r16 ; r16 = 0xBA
```

Primary uses:

- Fast $\times 16$ or $\div 16$ on values that fit in a nibble
- BCD digit manipulation
- Building two-nibble values from separate sources:

```
; Pack two 4-bit values (r16 low nibble, r17 low nibble) into one byte
andi r16, 0x0F ; keep only low nibble of r16
swap r17       ; move r17's low nibble to high position
andi r17, 0xF0 ; keep only high nibble of r17
or r16, r17    ; r16 = r17_nibble : r16_nibble
```

SWAP is a good example of an instruction that looks niche until you start doing display code, BCD work, or byte packing. Then it becomes surprisingly useful.

6.4 MUL — Unsigned 8×8 Multiply

MUL Rd, Rr ; r1:r0 ← Rd × Rr (unsigned)

2 cycles. Result always in r1:r0 regardless of operands.

The fixed result location is the most important thing to remember about MUL. The product does not stay near the source operands. It always lands in r1:r0.

```
ldi r16, 12
ldi r17, 25
mul r16, r17      ; r1:r0 = 300 = 0x012C
movw r24, r0      ; copy result to a usable register pair
clr r1            ; RESTORE r1 = 0 – this is mandatory (see below)
```

Why must you clear r1?

By the GCC AVR ABI, r1 must always be 0 when calling C functions or when C code calls assembly. MUL sets r1 to the high byte of the product. If you forget CLR r1, the next multiply or any code that assumes r1=0 (like ADC Rr, r1 for carry propagation) will produce wrong results.

Develop the habit: every MUL/MULS/MULSU is immediately followed by CLR r1 unless you are certain no C code is involved and you consume r1 before any carry-chain arithmetic.

Even if you are not mixing with C today, it is still a good discipline. The “r1 should be zero” convention is useful across the whole codebase.

6.4.1 Self-Multiply (Square)

```
; r25:r24 = r16 * r16 (r16 squared)
mul r16, r16
movw r24, r0
clr r1
```

MUL allows both operands to be the same register.

6.4.2 16-bit Result Scaling

A common embedded pattern: scale a 16-bit value by an 8-bit fraction using the high byte of the product as a “multiply then shift right 8” operation:

This is a classic fixed-point trick: keep the multiply, throw away the low byte, and treat the high byte as the scaled result.

```
; Approximate: result = (val * fraction) >> 8
; val in r16, fraction in r17 (0x00=0%, 0xFF≈100%, 0x80=50%)
mul r16, r17
mov r16, r1      ; take only the high byte → effectively >> 8
clr r1
```

6.5 MULS — Signed 8×8 Multiply

MULS Rd, Rr ; r1:r0 ← (signed)Rd × (signed)Rr

Both operands treated as signed (-128..127). Result is signed 16-bit. **r16–r31 only** for both operands.

```
ldi r16, 0xF6      ; -10 in two's complement
ldi r17, 5
muls r16, r17      ; r1:r0 = -50 = 0xFFCE
movw r24, r0
clr r1

ldi r16, 0xF6      ; -10
ldi r17, 0xFB      ; -5
muls r16, r17      ; r1:r0 = 50 = 0x0032
```

```
movw r24, r0
clr r1
```

6.6 MULSU — Signed × Unsigned Multiply

MULSU Rd, Rr ; r1:r0 ← (signed)Rd × (unsigned)Rr

Rd (r16–r23) is signed; Rr (r16–r23) is unsigned. Result is signed 16-bit.

Used in fixed-point arithmetic where one operand is a coefficient (signed scaling factor) and the other is a measurement (unsigned raw value):

That signed/unsigned mix shows up often in calibration and correction code.

```
; Scale a raw ADC reading (0-255, unsigned) by a signed calibration
; coefficient (-128..127), result in r25:r24.
; Example: reading=200 (0xC8), coeff=-3 (0xFD)
ldi r16, 0xFD ; signed coefficient -3
ldi r17, 200 ; unsigned reading
mulsu r16, r17 ; r1:r0 = -600 = 0xFDA8
movw r24, r0
clr r1
```

6.7 Multiply Summary

Instruction	Rd range	Rr range	Signed?	Result
MUL Rd, Rr	r0–r31	r0–r31	both unsigned	r1:r0 unsigned
MULS Rd, Rr	r16–r31	r16–r31	both signed	r1:r0 signed
MULSU Rd, Rr	r16–r23	r16–r23	Rd signed, Rr unsigned	r1:r0 signed

All: 2 cycles. Result always in r1:r0. Always CLR r1 after.

MULS/MULSU affect: C, Z. MUL affects: C, Z.

6.8 Division — Software Algorithm

AVR has no divide instruction. Division is implemented with a shift-subtract loop: the same long division algorithm you learned in school, but in binary.

So division is not something the hardware does for you in one step. You are trading simplicity of concept for cost in cycles.

6.8.1 Unsigned 8-bit Division

dividend / divisor → quotient, remainder

The algorithm processes one bit at a time from MSB to LSB:

For each bit (MSB to LSB):

1. Shift MSB of dividend into remainder LSB (remainder $\ll=$ 1, r[0] = d_msb)
2. If remainder $\geq$ divisor:
 - remainder -= divisor
 - set current quotient bit = 1
- Else:
 - quotient bit = 0

If that seems abstract, think of it this way: each iteration asks, "Can the current remainder afford one more divisor?" If yes, subtract and record a 1 in the quotient. If not, record a 0 and continue.

Assembly implementation (clean restoring version):

```
/*
 * div8u – unsigned 8-bit division
 * Input:  r16 = dividend, r17 = divisor
 * Output: r18 = quotient, r19 = remainder
 * Clobbers: r20
 * Cycles: ~9 × 8 = ~72 worst case
 */
div8u_clean:
    clr r19                ; remainder = 0
    ldi r20, 8             ; 8 bits
1:
    lsl r16                ; MSB of dividend → C
    rol r19                ; remainder = (r19 << 1) | C
    sec                   ; tentatively: quotient bit = 1
    cp r19, r17
    brsh 2f               ; remainder >= divisor: confirmed, subtract
    clc                   ; remainder < divisor: quotient bit = 0
    rjmp 3f
2:
    sub r19, r17          ; remainder -= divisor
3:
    rol r18                ; shift quotient bit (C) into r18
    dec r20
    brne 1b
    ret
```

Trace: 7 / 3:

```
bit 7: remainder=0→0, C(lsl)=0, 0<3 → C=0 → r18=0b00000000
bit 6: remainder=0→0, C=0, 0<3 → C=0 → r18=0b00000000
bit 5: remainder=0→0, C=0, 0<3 → C=0 → r18=0b00000000
bit 4: remainder=0→0, C=0, 0<3 → C=0 → r18=0b00000000
bit 3: remainder=0→0, C=0, 0<3 → C=0 → r18=0b00000000
bit 2: remainder=0→1, C=1, 1<3 → C=0 → r18=0b00000000
bit 1: remainder=1→3, C=1, 3>=3 → sub=0, C=1 → r18=0b00000001
bit 0: remainder=0→1, C=1, 1<3 → C=0 → r18=0b00000010
```

quotient = 0b00000010 = 2 OK

remainder = 1 OK (7 = 2×3 + 1)

Tracing a few divisions by hand is worth the time. After that, the loop stops looking magical and starts looking mechanical.

6.8.2 Unsigned 16-bit Division

Extend the pattern to 16-bit dividend, 8-bit divisor (common for scaled sensor values):

This is the same algorithm again, just stretched over more bits.

```

/*
 * div16u8 – divide 16-bit by 8-bit
 *   Input:  r25:r24 = dividend, r16 = divisor
 *   Output: r25:r24 = quotient, r18 = remainder
 *   Clobbers: r19, r20
 */
div16u8:
    clr  r18                ; remainder = 0
    ldi  r19, 16            ; 16 bits
    clr  r20                ; quotient low
    clr  r21                ; quotient high
1:
    lsl  r24                ; shift MSB of dividend → C
    rol  r25
    rol  r18                ; remainder <= 1, | C
    sec
    cp   r18, r16
    brsh 2f
    clc
    rjmp 3f
2:
    sub  r18, r16
3:
    rol  r20                ; quotient <= 1, insert bit in C
    rol  r21
    dec  r19
    brne 1b
    movw r24, r20           ; quotient → r25:r24
    ret

```

6.8.3 Power-of-2 Division — Just Shift

If the divisor is a compile-time power of 2, skip the loop entirely:

```

; Unsigned: r16 /= 8  (= 2^3)
lsr r16
lsr r16
lsr r16

; Signed: r16 /= 4  (rounds toward -∞, which is floor division)
asr r16
asr r16

; 16-bit unsigned: r25:r24 /= 4
lsr r25
ror  r24
lsr r25
ror  r24

```

Always use shifts for power-of-2 division. The loop costs ~72 cycles; three LSRs cost 3.

That is not a small optimization. It is the difference between “expensive software division” and “almost free.”

6.9 Worked Example: 8×8 Unsigned Multiply to 16-bit Result

Full annotated example with verification by objdump.

```
/*
 * Input:  r24 = a (0-255), r22 = b (0-255)
 * Output: r25:r24 = a * b (0-65025)
 * ABI-compliant: returns 16-bit result in r25:r24.
 */
mul8_to_16:
    mul    r24, r22      ; r1:r0 = a * b
    movw   r24, r0       ; r25:r24 = result
    clr    r1            ; restore ABI r1=0
    ret
```

Disassembly of this function (avr-objdump -d):

```
<mul8_to_16>:
 0: 86 9f      mul    r24, r22
 2: c0 01      movw   r24, r0
 4: 11 24      eor    r1, r1      ; CLR alias
 6: 08 95      ret
```

MUL is a single 16-bit word (2 bytes). The two-cycle cost comes from the hardware multiplier computing the 16-bit product in parallel with the next fetch. MOVW is another single word. The entire routine is 5 instructions, 8 bytes, 6 cycles.

This is a good example of why knowing the ISA matters. A tiny hand-written routine can be both obvious and efficient once you know the dedicated instruction exists.

6.10 Signed Arithmetic Pitfalls

6.10.1 Two Mistakes with Signed Division

Mistake 1: Using LSR (logical) instead of ASR (arithmetic) to divide a signed value.

```
ldi  r16, 0xFE      ; = -2 signed
lsr  r16            ; r16 = 0x7F = 127 ← WRONG (unsigned right shift)
asr  r16            ; r16 = 0xFF = -1 ← CORRECT (-2 / 2 = -1)
```

Mistake 2: Using BRLO/BRSH (unsigned) instead of BRLT/BRGE (signed) for signed comparisons. Covered in Chapter 4 but worth repeating whenever you work with signed arithmetic.

The CPU does not know your intent. A byte is just a bit pattern until your instruction choice tells the machine whether to treat it as signed or unsigned.

6.10.2 Overflow Detection

After a signed operation, V=1 means overflow. The result is wrong as a signed value:


```
ldi r16, 100      ; +100
ldi r17, 100      ; +100
add r16, r17      ; result = 200, but 200 > 127 → V=1 (signed overflow)
brvs overflow_handler
```

After unsigned operations, C=1 means overflow (result > 255 / 65535).

That signed/unsigned split is one of the most common sources of subtle bugs in embedded arithmetic.

6.11 Useful Idioms

6.11.1 Test if a value is zero without clobbering flags from the previous op

```
tst r16           ; Z=1 if r16=0, N=1 if negative – no flags from ADD etc.
```

Strictly speaking, TST does set flags. The real point is that it lets you ask “is this register zero or negative right now?” without changing the register value itself.

6.11.2 Conditional absolute value

```
; |r16|
sbrcc r16, 7      ; skip if bit 7 clear (positive)
neg r16           ; negate if negative
```

6.11.3 Saturating add (clamp at 255 instead of wrapping)

```
add r16, r17      ; add
brcc 1f           ; no carry: result fits in 8 bits
ser r16           ; carry: clamp to 0xFF
1:
```

This is a very common embedded pattern because sensor and UI code often prefers clamping over wraparound.

6.11.4 Saturating subtract (clamp at 0 instead of wrapping)

```
sub r16, r17
brcc 1f           ; no borrow (C=0): result >= 0
clr r16           ; borrow: clamp to 0
1:
```

6.11.5 Sign extension: 8-bit signed → 16-bit signed

```
; r16 (signed) → r17:r16 (signed 16-bit)
clr r17
sbrcc r16, 7      ; if bit 7 clear: r17 = 0 (positive, done)
ser r17           ; bit 7 set: r17 = 0xFF (negative, sign extend)
; r17:r16 is now the correct signed 16-bit representation
```

6.12 Summary

Multi-byte arithmetic:

Always low byte first. Use ADC/SBC to chain carry between bytes.

ADIW/SBIW: $\pm 0-63$ to 16-bit pair, 2 cycles.

SUBI+SBCI: add any constant (negate the immediate).

Shifts:

LSL/LSR – logical (unsigned), 0 inserted.

ASR – arithmetic (signed), sign bit preserved.

ROL/ROR – rotate through carry – chains shifts across bytes.

SWAP – nibble swap, 1 cycle, use for $\times 16$ / $\div 16$.

Multiply:

MUL – unsigned $\times$ unsigned $\rightarrow$ r1:r0. Always CLR r1 after.

MULS – signed $\times$ signed $\rightarrow$ r1:r0. r16-r31 only.

MULSU – signed $\times$ unsigned $\rightarrow$ r1:r0. r16-r23 only.

All: 2 cycles. Result in r1:r0.

Division: no hardware DIV. Use shift-subtract loop (~ 72 cycles for 8-bit).

Power-of-2 divisors: use LSR/ASR instead – $O(\log_2 n)$ cycles.

Signed pitfalls:

Use ASR not LSR for signed right shift.

Use BRLT/BRGE not BRLO/BRSH for signed branch.

V flag = signed overflow, C flag = unsigned overflow.

6.13 Exercises

- Write a 16-bit multiply: $r25:r24 = r25:r24 \times r23:r22$ (both inputs unsigned). You cannot use MUL for the full 16×16 product directly (it only handles 8×8). Decompose into four 8×8 multiplies and accumulate. Hint: $(a_{hi} \times 256 + a_{lo}) \times (b_{hi} \times 256 + b_{lo})$.
- Extend div16u8 to handle divide-by-zero explicitly. Choose a policy (for example: quotient = 0xFFFF, remainder = dividend low byte, or branch to an error handler) and document it.
- Write a function that counts the number of 1-bits in r16 (popcount). Use LSL + ADC r17, r1 pattern: after each LSL, C holds the shifted-out bit; ADC r17, r1 adds C to the accumulator (r1=0).
- Implement saturating 16-bit addition: $r25:r24 + r23:r22$, clamped to 0xFFFF if overflow. The 16-bit carry flag is in C after the second ADC.
- Verify the div8u_clean routine:
 - Trace $100 / 7$ by hand (expected: quotient=14, remainder=2).
 - Assemble and simulate: use `avr-objdump -d` to confirm the instruction sequence, then trace the bit loop manually.
- Without MUL, compute $r16 \times 10$ using only shifts and adds. ($10 = 8 + 2 = (1 \ll 3) + (1 \ll 1)$.)

Next: Chapter 7 — Branches and Control Flow

Chapter 7

Branches and Control Flow

Every non-trivial program has decisions and loops. In assembly, all of them reduce to one fundamental act: change the program counter, or deliberately let it continue to the next instruction. This chapter covers the main ways AVR does that — unconditional jumps, conditional branches, and single-instruction skips — and shows how to map C-style `if`, `while`, `for`, and `switch` into assembly patterns you can read and write fluently.

The key mental shift is that high-level control flow is just structured PC manipulation.

7.1 Unconditional Jumps

7.1.1 RJMP — Relative Jump

`RJMP k` ; PC ← PC + k + 1 (k = signed offset, -2048..+2047 words)

2 cycles. 1 word. Range: **±2047 words** (±4094 bytes) from the instruction.

```
rjmp loop      ; jump to label "loop" — GAS computes the offset
rjmp .         ; infinite loop (jump to current address, offset = -1)
```

RJMP covers most local control-flow transfers — nearly 4K instruction words (just under 8 KB of code bytes). Use it by default.

That “use it by default” rule is practical: RJMP is smaller and faster than JMP, and most control-flow targets are close enough to fit.

7.1.2 JMP — Absolute Jump

`JMP k` ; PC ← k (k = 22-bit program word address)

3 cycles. 2 words (32-bit instruction). Range: entire 64K-word (128KB) flash.

```
jmp main       ; jump anywhere — no range limit
jmp reset_handler
```

Use JMP only when the target is out of RJMP range, or in the vector table where you need a guaranteed 2-word slot. For everything else, RJMP is smaller and faster.

So the usual decision is simple:

- nearby target: RJMP
- far target or fixed-size vector entry: JMP

7.1.3 IJMP — Indirect Jump via Z

IJMP ; PC ← Z (Z holds word address of target)

2 cycles. 1 word. Z contains the **word address** (byte address ÷ 2) of the destination. Use pm(label) to convert a label to its word address:

```
ldi ZL, lo8(pm(target))
ldi ZH, hi8(pm(target))
ijmp ; jump to target
```

Primary use: **jump tables** (switch/case). Covered in Section 7.7.

Indirect jumps are less common in beginner code, but they are the step that takes you from simple branch chains to table-driven dispatch.

7.2 Conditional Branches — Mechanics

All conditional branches work identically:

1. Test one or more bits in SREG.
2. If the condition is true: PC += offset + 1 (branch taken, 2 cycles).
3. If the condition is false: PC += 1 (not taken, 1 cycle).

Range: **±63 words** (±126 bytes). If the target is farther, invert the condition and use JMP/RJMP:

```
; Branch to far_target if r16 == 0 – but far_target > 63 words away:
tst r16
brne skip ; if NOT zero, skip the jump (inverted condition)
jmp far_target ; now unconditionally jump to far target
skip:
```

This pattern — invert condition, skip over jump — is what the compiler emits.

Once you recognize that shape, compiler-generated control flow becomes much easier to read.

7.3 Complete Branch Instruction Table

Mnemonic	Alias for	Condition	Flags tested	Typical use after
BREQ	BRBS 1,k	Z = 1	Z	CP, CPI, TST, SUB
BRNE	BRBC 1,k	Z = 0	Z	CP, CPI, DEC, INC
BRCS	BRBS 0,k	C = 1	C	ADD, SUB, LSL, LSR
BRCC	BRBC 0,k	C = 0	C	ADD, SUB
BRLO	BRCS	C = 1	C	CP, CPI (unsigned <)
BRSH	BRCC	C = 0	C	CP, CPI (unsigned ≥)
BRLT	BRBS 4,k	S = 1 (N ⊕ V)	N, V	CP, CPI (signed <)
BRGE	BRBC 4,k	S = 0 (N ⊕ V)	N, V	CP, CPI (signed ≥)
BRMI	BRBS 2,k	N = 1	N	ADD, SUB (signed neg)
BRPL	BRBC 2,k	N = 0	N	ADD, SUB (signed pos)

BRVS	BRBS 3, k	V = 1	V	ADD, SUB (ovf detect)
BRVC	BRBC 3, k	V = 0	V	ADD, SUB
BRIE	BRBS 7, k	I = 1	I	SEI/CLI context
BRID	BRBC 7, k	I = 0	I	SEI/CLI context
BRHS	BRBS 5, k	H = 1	H	BCD arithmetic
BRHC	BRBC 5, k	H = 0	H	BCD arithmetic
BRTS	BRBS 6, k	T = 1	T	BST/BLD sequences
BRTC	BRBC 6, k	T = 0	T	BST/BLD sequences

BRBS *b*, *k* (Branch if Bit in SREG Set) and BRBC *b*, *k* (Branch if Bit in SREG Clear) are the generic forms. Every named branch is an alias for one of these two — GAS assembles BREQ as BRBS 1, *k* (bit 1 = Z flag).

In practice you almost always write the named forms, because BREQ expresses intent better than “branch if SREG bit 1 is set.”

7.4 Skip Instructions

Skips are single-instruction conditional skips — they bypass exactly the next instruction if a condition is met. No label needed. They are the most compact way to write small conditionals.

If branches are for “go somewhere else,” skips are for “ignore the next thing if this condition holds.”

7.4.1 CPSE — Compare, Skip if Equal

CPSE Rd, Rr ; if Rd == Rr: skip next instruction

1 cycle (no skip), 2 cycles (skip 1-word), 3 cycles (skip 2-word).

That timing detail matters because the skipped instruction still occupies space in the stream even though it does not execute.

```
cpse r16, r17
rjmp not_equal ; skipped if r16 == r17 — falls through to equal path
; equal path here
```

7.4.2 SBRC — Skip if Bit in Register is Clear

SBRC Rr, b ; if bit b of Rr == 0: skip next instruction

Tests bit *b* (0–7) of register Rr (r0–r31). No flags affected.

```
; Process bit 0 of flags register r16
sbrc r16, 0 ; skip if bit 0 clear
rcall handle_flag0 ; only called if bit 0 was set
```

7.4.3 SBRS — Skip if Bit in Register is Set

SBRS Rr, b ; if bit b of Rr == 1: skip next instruction

```
sbrs r16, 7 ; skip if bit 7 set (value is negative signed)
rjmp positive_path ; only taken if bit 7 was clear
; negative path falls through here
```

7.4.4 SBIC — Skip if Bit in I/O Register is Clear

SBIC A, b ; if bit b of I/O[A] == 0: skip (A = 0x00–0x1F only)

Tests a bit in an I/O register in a single cycle without loading into a register first. The address range is limited to 0x00–0x1F (the lower half of the IN/OUT space).

```
; Wait until a flag bit in GPIOR0 is set by an ISR
wait_for_flag:
    sbis GPIOR0, 0 ; skip if bit 0 of GPIOR0 is set
    rjmp wait_for_flag ; not set: loop
; bit 0 now set – proceed
```

7.4.5 SBIS — Skip if Bit in I/O Register is Set

SBIS A, b ; if bit b of I/O[A] == 1: skip (A = 0x00–0x1F only)

```
; Check if a GPIOR flag is clear
sbic GPIOR0, 2 ; skip if bit 2 clear
rcall clear_handler ; only if bit 2 was set
```

7.4.6 SBIC/SBIS Address Limit on ATtiny3217

On ATtiny3217, SBIC/SBIS reach I/O addresses 0x00–0x1F only. The PORT peripheral registers (PORTB_IN, etc.) are at 0x0420+ — far beyond this limit. To test a PORT pin, you must load it first:

```
; Read PA3 state into r16, test bit 3
lds r16, PORTA_IN ; load port input register
sbrc r16, 3 ; skip if PA3 is low
rcall pa3_high_handler
```

The registers that ARE reachable with SBIC/SBIS on ATtiny3217:

I/O Addr	Register	Useful bits
0x1C	GPIOR0	User-defined flags (bit 0–7)
0x1D	GPIOR1	User-defined flags
0x1E	GPIOR2	User-defined flags
0x1F	GPIOR3	User-defined flags

GPIOR0–3 are key scratch flag registers in low I/O space: readable in 1 cycle with SBIC/SBIS, no load required. VPORTx_IN is the other key case: it gives you fast bit tests on real pins without touching the slower PORTx_IN registers in extended I/O space.

7.4.7 Skip vs Branch — When to Use Which

Situation	Use
Skip 1 instruction on a bit condition	SBRC/SBRS/SBIC/SBIS
Skip 1 instruction on equality	CPSE
Branch to a label	BREQ, BRNE, etc.
More than 1 instruction in the conditional body	Branch

Skips are best for guard clauses and flag-polling loops. Branches are best for bodies with multiple instructions.

That is the simplest rule for choosing between them: one-instruction body, consider a skip; larger body, use a branch.

7.5 Structured Control Flow Patterns

7.5.1 if / else

```
// C:
if (a >= b) { result = a; } else { result = b; } // max(a, b)

; r16 = a, r17 = b, result → r18
; For max(a,b), choosing a when equal is fine.
cp   r16, r17           ; a - b
brsh then_branch        ; if a >= b unsigned: take then-branch
; else-branch: b >= a, so max is b
mov  r18, r17
rjmp end_if
then_branch:
mov  r18, r16
end_if:
```

Pattern: Test → branch to else (or over then) → then-body → rjmp end → else-body → end label. Invert the C condition for the branch direction.

That “invert the condition” step is one of the most common assembly habits. You often branch around the code you do not want rather than branch directly into every desired path.

7.5.2 if without else (guard clause)

```
if (r16 == 0) return;

tst  r16
brne continue        ; NOT zero: skip the return
ret                  ; zero: return early
continue:
```

For this pattern, a branch is clearer than forcing CPSE into it:

```
tst  r16
breq early_return
rjmp continue
early_return:
ret
continue:
```

This is a good example of a broader rule: the shortest control-flow sequence is not always the clearest one.

7.5.3 while loop

```
while (condition) { body; }

; C: while (r16 != 0) { r16--; }
while_test:
    tst  r16
    breq while_done    ; condition false: exit
    dec  r16           ; body
```

```
    rjmp while_test
while_done:
```

Pre-test — loop may execute zero times. Costs one extra RJMP per iteration.

That is why a literal C-style while is often not the most efficient assembly shape.

7.5.4 do-while loop (preferred in assembly)

```
do { body; } while (condition);
; C: do { r16--; } while (r16 != 0);
do_body:
    dec r16          ; body
    brne do_body     ; condition: not zero → repeat
```

This is the most efficient loop form. No RJMP back — the branch IS the loop-back. One instruction of overhead per iteration instead of two. The compiler emits this whenever it can prove the loop body runs at least once.

This is one of the most useful structural transformations in assembly: if you can legally turn a loop into do-while form, you usually should.

Always restructure assembly loops as do-while when the body runs at least once:

```
; Count-down loop: r18 iterations, guaranteed >= 1
ldi r18, N
loop:
    ; body
    dec r18
    brne loop          ; one branch instruction, no rjmp
```

7.5.5 for loop

```
for (i = 0; i < N; i++) { body; }
```

Count **down** rather than up — comparison against 0 is free (BRNE after DEC):

```
; Count-down (preferred):
ldi r18, N
1:
    ; body
    dec r18
    brne 1b

; Count-up (necessary when index used inside body):
clr r18
1:
    ; body (r18 = current index)
    inc r18
    cpi r18, N          ; costs one extra CPI per iteration
    brne 1b
```

That is why countdown loops appear everywhere in AVR code. Zero is the cheapest comparison target on the machine.

7.5.6 Nested loops

```

for (i = 8; i != 0; i--)
    for (j = 64; j != 0; j--)
        body;

ldi r19, 8                ; outer counter
outer:
    ldi r18, 64           ; inner counter – reload every outer iteration
    inner:
        ; body
        dec r18
        brne inner
    dec r19
    brne outer

```

Critical: reload the inner counter at the top of the outer loop body, **before** the inner loop starts. A common mistake is loading it once before the outer loop — it reaches zero after the first outer iteration and never loops again.

This bug is common because the loop still looks nested at a glance. The reload placement is what actually makes it nested.

7.6 Worked Example: Sum an Array

Sum 8 unsigned bytes from flash, store 16-bit result in SRAM.

```

/* sum_array – sums ARRAY_LEN bytes from flash
 * Uses mapped flash (AVRxt): Z = flash addr + MAPPED_PROGMEM_START
 * Result in r25:r24
 */

.section .rodata
test_array:
    .byte 10, 20, 30, 40, 50, 60, 70, 80    ; sum = 360 = 0x0168
.equ ARRAY_LEN, 8

.section .text

sum_array:
    ldi ZL, lo8(test_array + MAPPED_PROGMEM_START)
    ldi ZH, hi8(test_array + MAPPED_PROGMEM_START)
    clr r24                ; sum_lo = 0
    clr r25                ; sum_hi = 0
    ldi r18, ARRAY_LEN     ; counter
loop:
    ld r16, Z+             ; byte = flash[Z], Z++
    add r24, r16           ; sum_lo += byte
    adc r25, r1            ; sum_hi += carry (r1 = 0)
    dec r18
    brne loop              ; do-while: no rjmp needed

    ; r25:r24 = 360 = 0x0168
    ret

```

Step-by-step trace:

Iter	r18	Z points to	r16	r25:r24
1	7	[20]	10	0x000A
2	6	[30]	20	0x001E
3	5	[40]	30	0x003C
4	4	[50]	40	0x0064
5	3	[60]	50	0x0096
6	2	[70]	60	0x00D2
7	1	[80]	70	0x0118
8	0	—	80	0x0168 OK

Why `ADC r25, r1` and not `ADC r25, r17`? After `DEC r18`, the N and Z flags reflect r18, but C is unchanged from the `ADD`. Using r1 (`=0`) adds only the carry bit, correctly propagating any overflow from the low byte without needing a dedicated zero register.

This is a good example of how flag discipline and dataflow interact. The right instruction is not just about values in registers; it is about which flag bit you are intentionally consuming next.

7.7 Jump Tables (switch/case)

For multi-way branches, a jump table is faster than a chain of comparisons. Each table entry holds the word address of a handler. `IJMP` jumps to the address in Z.

Conceptually, a jump table turns “which case is this?” into array indexing.

```
switch (cmd) {
    case 0: cmd0(); break;
    case 1: cmd1(); break;
    case 2: cmd2(); break;
    default: cmd_default();
}

; r16 = command index (0, 1, or 2)

; Bounds check first:
cpi r16, 3          ; 3 = number of cases
brsh jtable_default ; >= 3: out of range

; Build Z = pm(jtable) + r16
; Each RJMP in the table is 1 word, so index directly.
ldi ZL, lo8(pm(jtable))
ldi ZH, hi8(pm(jtable))
add ZL, r16          ; offset by case index
adc ZH, r1            ; carry
ijmp                 ; jump to table[r16] - Z holds word address

jtable:
    rjmp cmd0        ; case 0
    rjmp cmd1        ; case 1
    rjmp cmd2        ; case 2
```

```

jtable_default:
    rcall cmd_default
    rjmp jtable_done

cmd0: rcall handler0 ; rjmp table_done implicit via fall-through pattern
      rjmp jtable_done
cmd1: rcall handler1
      rjmp jtable_done
cmd2: rcall handler2
      rjmp jtable_done
jtable_done:

```

How it works:

`pm(jtable)` gives the **word address** of `jtable`. Adding the case index (in words) moves `Z` to the corresponding `RJMP` instruction. `IJMP` then executes that `RJMP`, which jumps to the actual handler. Each entry is 1 word (`RJMP` = 1 word) so the index scales correctly.

That two-step jump looks strange the first time you see it, but it buys you compact entries and simple indexing.

Cycle cost: bounds check (2) + `LDI`×2 (2) + `ADD`+`ADC` (2) + `IJMP` (2) + `RJMP` from table (2) = **10 cycles**. A chain of `CPI`/`BREQ` for 3 cases costs up to $3 \times (1+2) = 9$ cycles in the worst case — comparable. Jump tables win decisively with more cases.

So jump tables are not automatically better for tiny switches. They become attractive once the case count is large enough that linear comparison chains stop being cheap.

7.7.1 Large Jump Tables with JMP Entries

If handlers are far from the table (`RJMP` out of range), use `JMP` (2 words) instead. Then each table entry is 2 words, so multiply the index by 2:

```

; Index × 2 for 2-word entries
lsl r16 ; r16 *= 2
ldi ZL, lo8(pm(jtable))
ldi ZH, hi8(pm(jtable))
add ZL, r16
adc ZH, r1
ijmp

jtable:
    jmp cmd0 ; 2 words each
    jmp cmd1
    jmp cmd2

```

7.8 Polling Loops

A polling loop spins until a condition becomes true — set by hardware or an ISR.

Polling is simple and sometimes appropriate, but it is also the most CPU-hungry waiting strategy because the core keeps executing the test loop continuously.

```

; Wait for USART0 receive complete (bit 7 of USART0_STATUS = RXCIF)
; Cannot use SBIS: address too high. Must use LDS + SBRS.
wait_rxc:

```

```

lds r16, USART0_STATUS    ; load USART status
sbrs r16, 7                ; skip if RXCIF bit set
rjmp wait_rxc              ; not set: loop
lds r16, USART0_RXDATA    ; read the received byte

```

For low I/O registers (GPOR), SBIS eliminates the load:

```

; ISR sets bit 0 of GPOR0 when data ready
wait_ready:
    sbis GPOR0, 0          ; 1 cycle: skip if bit 0 set
    rjmp wait_ready        ; 2 cycles: loop back
; 3 cycles per iteration vs 5 cycles with LDS+SBRs

```

That difference is exactly why low-I/O flag registers are so useful on AVR.

7.9 Branch Optimisation Notes

7.9.1 Do-while saves one instruction per loop

```

; while: 3 instructions per iteration (body + dec + brne + rjmp)
; do-while: 2 instructions per iteration (body + dec + brne)

```

For a loop that runs 1000 times, that's 1000 saved cycles.

Over one or two iterations this does not matter. Over hot loops, it matters a lot.

7.9.2 Count down, not up

DEC r18 + BRNE is **two instructions with no extra comparison**. Counting up needs INC + CPI N + BRNE — three instructions.

This is one of those patterns that looks like a small optimization until you write enough firmware to see it everywhere.

7.9.3 Use SBRC/SBRs for single-bit tests

Instead of:

```

andi r16, (1 << 3)
brne bit3_set

```

Write:

```

sbrs r16, 3
rjmp bit3_clear
; bit3_set path falls through

```

Saves one instruction and doesn't destroy r16.

Preserving the register value is often as important as saving the instruction.

7.9.4 Branch delay slots do not exist on AVR

Unlike ARM or MIPS, the AVR has no branch delay slot. The instruction after a branch is NOT executed if the branch is taken. No surprises.

That makes AVR control flow pleasantly literal: if you branch away, the next sequential instruction does not secretly run first.

7.10 Summary

Unconditional:

- RJMP – ± 2047 words, 2 cycles, 1 word. Default choice.
- JMP – full range, 3 cycles, 2 words. For far targets or vector table.
- IJMP – Z = word address, 2 cycles. For jump tables.

Conditional branches: ± 63 words, 2 cycles taken / 1 not taken.

- Unsigned: BRLO (<), BRSH ($\geq$), BREQ (=), BRNE ($\neq$)
- Signed: BRLT (<), BRGE ($\geq$), BRMI (neg), BRPL (pos)
- Far target: invert condition, skip over JMP.

Skip instructions (1-3 cycles, no label needed):

- CPSE Rd, Rr – skip if Rd == Rr
- SBRC Rr, b – skip if bit b of Rr clear
- SBRS Rr, b – skip if bit b of Rr set
- SBIC A, b – skip if bit b of I/O[A] clear ($A \leq 0x1F$)
- SBIS A, b – skip if bit b of I/O[A] set ($A \leq 0x1F$)

Loop patterns:

- do-while: most efficient – branch is the loop-back (no extra RJMP)
- for: count down with DEC+BRNE – no CPI needed
- nested: reload inner counter at start of each outer iteration

Jump tables: IJMP with Z = pm(table) + index. 0(1) for any case count.

Low-I/O targets such as `VPORTx\_IN` and `GPIOR0-3` are reachable with `SBIC`/`SBIS`. `GPIORx` make especially good fast flag bytes between ISR and main loop.

7.11 Exercises

1. Rewrite the array-sum example to sum **signed** 8-bit values into a signed 16-bit result. What changes? (Hint: ADC r25, r1 is still correct for unsigned carry, but how do you handle sign extension of each negative byte before adding?)
2. Write a function `find_min(X, n)` that finds the minimum byte value in an array of n bytes at SRAM address X. Return the minimum in r24. Use a do-while loop.
3. Extend the jump table example to 8 cases. What is the cycle cost of the dispatch compared to a chain of 8 CPI/BREQ pairs (worst case)?
4. Write a `strlen` function: given X pointing to a null-terminated string in SRAM, return the length (not counting null) in r24. Use a do-while loop with LD X+ and TST.
5. Implement a `memchr` function: scan n bytes at X for value r22, return pointer to first match in X (or 0 if not found). Use CPSE to detect a match.
6. The polling loop `sbis GPIOR0, 0 / rjmp wait` burns CPU cycles continuously. Name one hardware peripheral on ATtiny3217 that would let you sleep the CPU instead of polling, waking only when the condition is true. (Covered in Chapter 10 — but think about it now.)

Next: Chapter 8 — Subroutines and the Stack

Chapter 8

Subroutines and the Stack

Subroutines let you package a piece of code once and reuse it from many places. The stack is what makes that possible: it remembers where execution should resume after the called code finishes. This chapter explains the mechanics of calls and returns, the GCC AVR ABI so your assembly can cooperate with C, frame pointers for local variables, and a recursive example that shows how all of those pieces interact.

This is the chapter where assembly starts to feel less like a straight-line script and more like structured software.

8.1 CALL and RCALL

```
RCALL k      ; push PC+1 → stack, PC ← PC + k + 1  (±2047 words)
CALL  k      ; push PC+1 → stack, PC ← k           (full range)
```

Both push the **return address** (the word address of the instruction after the call) onto the stack, then jump. On ATtiny3217 ($\leq 128\text{KB}$ flash), return addresses are **2 bytes** (16-bit word address).

That return address is the whole reason the stack exists in normal function calls. Without it, the CPU would have no idea where to come back after the callee finishes.

RCALL: 3 cycles, 1 word. Range ± 2047 words — sufficient for calls within the same file. Use this by default.

CALL: 4 cycles, 2 words. Full 64K-word address space. Use when the target is out of RCALL range.

So this chapter follows the same practical rule as jumps:

- nearby function: RCALL
- far function: CALL

```
rcall my_function      ; call nearby – 3 cycles
call  far_function     ; call anywhere – 4 cycles
```

The stack after RCALL:

Before call:	After RCALL my_function: (return addr = word addr of next instr)
SP → [free]	SP → ret_hi ret_lo [free before call]

RCALL pushes the low byte of the return address first (SP--), then the high byte (SP--). On return, POP is reversed.

You usually do not need to care about the byte order while writing ordinary calls, but it matters when you are reasoning about exact stack layout.

8.1.1 ICALL — Indirect Call via Z

ICALL ; push PC+1, PC ← Z (Z = word address of target)

Like IJMP but saves a return address. Used for function pointers:

```
ldi  ZL, lo8(pm(my_func))
ldi  ZH, hi8(pm(my_func))
icall                   ; call my_func via pointer in Z
```

Indirect calls are the function-call version of jump tables and callbacks: the destination is chosen at runtime instead of being fixed in the source.

8.2 RET and RETI

RET ; PC ← stack (pop 2 bytes), SP += 2
RETI ; PC ← stack, SP += 2, SREG I-bit ← 1

RET: 4 cycles. Pops the return address and jumps to it.

RETI: same as RET but also re-enables interrupts (sets I in SREG). Used exclusively at the end of interrupt service routines. Chapter 10 covers ISRs.

Every CALL/RCALL must have a matching RET. Mismatched push/pop of the return address will send execution to a garbage address — one of the hardest bugs to diagnose.

This is why stack bugs are so destructive. A wrong data value is one problem; a wrong return address means execution can wander into nonsense.

8.3 PUSH and POP

PUSH Rr ; SP ← SP - 1, SRAM[SP] ← Rr (2 cycles)
POP Rd ; Rd ← SRAM[SP], SP ← SP + 1 (2 cycles)

PUSH: decrement first, then write. SP ends up pointing to the pushed byte. **POP: read from SP, then increment.** SP ends up pointing to the byte above the popped value.

That order is worth internalising because many stack diagrams only make sense once you remember that AVR updates the pointer before a push and after a pop.

```
push r16       ; save r16
push r17       ; save r17
; ... use r16, r17 freely ...
pop r17        ; restore r17 FIRST (LIFO – last pushed, first popped)
pop r16        ; restore r16
```

Stack grows **downward**. After two pushes at SP=0x3FFF:


```

0x3FFF  [original top – SP was here]
0x3FFE  r16 value  ← after first push
0x3FFD  r17 value  ← SP after second push
0x3FFC  [free]

```

8.4 The GCC AVR ABI

When assembly code must call C functions or be called from C, both sides must agree on which registers hold arguments and which are preserved across calls. This agreement is the **ABI** (Application Binary Interface).

Without the ABI, two correct pieces of code could still fail to work together. The ABI is the contract that keeps separately written code interoperable.

8.4.1 Argument Passing

Arguments are passed in registers, allocated **from r25 downward**, starting with the first (leftmost) argument:

C function: `foo(uint8_t a, uint16_t b, uint8_t c)`

Register:
 r24 r23:r22 r20

Rules:

- 1-byte argument: single register (r24, r22, r20, r18...)
- 2-byte argument: register pair (r25:r24, r23:r22, r21:r20...)
- 4-byte argument: 4 registers (r25:r24:r23:r22...)
- Arguments fill slots downward; odd-size arguments consume a full pair

When registers run out (more than ~8 argument bytes), remaining arguments go on the stack — pushed right-to-left before the call.

The practical implication is that many small functions never touch the stack for arguments at all. Everything lives in registers unless the argument list gets large.

Common signatures:

```

void f(uint8_t a)                    a → r24
void f(uint16_t a)                  a → r25:r24
void f(uint8_t a, uint8_t b)        a → r24, b → r22
void f(uint8_t a, uint16_t b)       a → r24, b → r23:r22
void f(uint16_t a, uint16_t b)      a → r25:r24, b → r23:r22

```

8.4.2 Return Values

Return type	Register
void	—
uint8_t / int8_t	r24
uint16_t / int16_t / pointer	r25:r24
uint32_t / int32_t	r25:r24:r23:r22 (r22=LSB)
uint64_t	r25:r24:r23:r22:r21:r20:r19:r18

8.4.3 Register Classes

This is the most important table in this chapter. Get it wrong and either your callee corrupts the caller's data, or you waste cycles saving registers you didn't need to.

In other words: this table is not bookkeeping trivia. It is the difference between correct reusable code and subtle corruption bugs.

Register	Class	Rule
r0	Scratch	Caller must save. Compiler uses freely.
r1	Zero register	Must ALWAYS be 0x00. If you write it (e.g. after MUL), CLR r1 immediately.
r2-r17	Call-preserved	Callee saves/restores if used.
r18-r27	Call-clobbered	Caller saves if needed across a call.
r28 (YL)	Call-preserved	Callee saves/restores if used.
r29 (YH)	Call-preserved	Used as frame pointer – callee saves.
r30 (ZL)	Call-clobbered	Caller saves if needed.
r31 (ZH)	Call-clobbered	Caller saves if needed.
SREG	–	Not automatically saved. ISRs must save it.
SP	–	Must be consistent on entry and exit.

Call-clobbered (r18–r27, r30–r31): the callee may destroy these. If the caller needs their values after the call, the caller must push them before the call and pop them after.

Call-preserved (r2–r17, r28–r29): the callee must save these before using them and restore them before returning. The caller can rely on them surviving a call.

That gives you two different programming styles:

- temporary scratch work belongs naturally in call-clobbered registers
- long-lived values belong naturally in call-preserved registers

```
; Caller perspective: preserve r20 across a call
push r20
rcall some_function      ; may destroy r18-r27, r30-r31
pop r20                  ; r20 is back to its pre-call value

; Callee perspective: use r14 (call-preserved)
my_func:
    push r14              ; must save
    ; ... use r14 ...
    pop r14               ; must restore
    ret
```

8.4.4 The r1 = 0 Rule

r1 must always be 0x00 except momentarily during a MUL/MULS/MULSU instruction sequence. The compiler generates ADC Rd, r1 to propagate carry without a separate zero register. If r1 is non-zero, those adds produce silently wrong results.

This rule feels strange at first because it treats one general-purpose register as special. But once you accept it, a lot of AVR idioms make more sense.

```
mul r16, r17              ; r1:r0 = product – r1 may be non-zero now
movw r24, r0              ; copy result
clr r1                    ; restore r1 = 0 – mandatory
```

8.5 Prologue and Epilogue Patterns

8.5.1 Minimal (no local variables, few registers)

```
my_func:
    push r14           ; save call-preserved regs used
    push r15
    ; ... body ...
    pop  r15           ; restore in reverse order
    pop  r14
    ret
```

This is the ideal small-function shape: save only what you must, do the work, restore it, and return.

8.5.2 With Stack Frame (local variables)

Local variables can be allocated on the stack using Y as a frame pointer. The Y register pair (r29:r28) is call-preserved, so the callee must save and restore it.

This is the moment where stack use becomes more than “where return addresses go.” The stack also becomes temporary storage for each active call frame.

Stack layout after full prologue (N bytes of locals):

```
Higher address (old SP):  [ caller's data ]
                        [ return address hi ] ← pushed by RCALL
                        [ return address lo ]
                        [ saved r29 (YH)    ] ← PUSH r29
                        [ saved r28 (YL)    ] ← PUSH r28
                        [ local N           ] ← Y+N
                        [ local 2           ] ← Y+2
                        [ local 1           ] ← Y+1 ← SP (new top)
```

Prologue:

```
my_func:
    push r28           ; save Y
    push r29
    in  r28, SPL        ; Y = current SP (points to saved r29)
    in  r29, SPH
    sbiw r28, N         ; allocate N bytes of locals
    out SPH, r29
    out SPL, r28        ; writing SPL completes the update
```

Order matters: Write SPH before SPL. On AVR, writing SPL automatically blocks interrupts for a few instructions (or until the next I/O write), which makes software SP updates safe when done in this order.

That is a hardware rule, not style advice. If you update SP in the wrong order, you risk an inconsistent stack pointer during an interrupt-sensitive window.

Access locals:

```
ldi  r16, 42
std  Y+1, r16          ; local1 = 42
ldd  r17, Y+1          ; r17 = local1
```

Epilogue:

```

    adiw r28, N          ; deallocate locals – Y points to saved r29 again
    out SPH, r29
    out SPL, r28
    pop r29              ; restore Y
    pop r28
    ret

```

8.5.3 When Do You Need a Frame?

Situation	Solution
No locals, ≤ 6 arguments	Use r18–r25 directly, no frame
A few locals that fit in call-preserved regs	Push/pop those regs
Many locals, or recursion	Full frame with Y pointer
Recursion	Always use the stack — registers don't survive recursive calls

So the frame pointer is not mandatory for every function. It is a tool you use when simple register-only management is no longer enough.

8.6 Stack Depth Analysis

Knowing how deep the stack goes prevents stack overflow (which silently corrupts SRAM).

That “silently” is the important word. Small microcontrollers usually do not raise an exception when the stack collides with data. They just start corrupting memory.

For each function, stack usage = (bytes for return address) + (bytes pushed in prologue) + (bytes of local allocation):

ATtiny3217: return address = 2 bytes (16-bit word address)

simple_add: 2 (ret addr) + 0 (no pushes) + 0 (no locals) = 2 bytes

with_frame: 2 (ret) + 2 (push r28/r29) + 4 (locals) = 8 bytes

factorial(n): each call frame = 2 (ret addr) + 1 (push r24 for n) = 3 bytes

depth for factorial(8) = $8 \times 3 = 24$ bytes of stack

Total SRAM = 2 KB (0x3800–0x3FFF). Stack starts at 0x3FFF and grows down. With .bss and .data consuming some SRAM, calculate:

Available for stack $\approx$ RAMEND - end_of_allocated_SRAM + 1 - safety_margin

Inspect with:

```
avr-objdump -t subroutines.elf | grep -E "\.bss|\.data"
```

```
avr-nm --size-sort subroutines.elf
```

For programs with bounded recursion depth (like factorial), calculate the maximum frame count $\times$ bytes per frame. For programs with unbounded recursion on a microcontroller: you have a stack overflow bug by design.

This is one reason recursion is often treated cautiously in embedded work. It is not forbidden, but it needs a real depth bound.

8.7 Worked Example: Recursive Factorial

$\text{factorial}(n) = n \times \text{factorial}(n-1)$, base case $\text{factorial}(1) = \text{factorial}(0) = 1$.

ABI: n passed in $r24$ (`uint8_t`), result returned in $r25:r24$ (`uint16_t`).

For $n \leq 8$: result fits in `uint16_t` ($8! = 40320 = 0x9D80$). For $n \geq 13$: result overflows `uint16_t` ($13! = 6,227,020,800 > 65535$).

That explicit range note is part of writing honest low-level code: if a routine has a numerical limit, document it.

```
/*
 * factorial – recursive uint16_t factorial(uint8_t n)
 * Input:  r24 = n
 * Output: r25:r24 = n!
 */
factorial:
    /* Base case: n <= 1 → return 1 */
    cpi    r24, 1
    brlo   .Lfact_base    ; n < 2: return 1

    /* Save n before it's overwritten by the recursive return value */
    push   r24             ; SP -= 1, SRAM[SP] = n

    /* Compute factorial(n-1) */
    dec    r24             ; r24 = n - 1
    rcall  factorial       ; r25:r24 = factorial(n-1)
                                ; --- recursive call ---
                                ; When we return here, r25:r24 = (n-1)!
                                ; r24 (our saved n) is still on the stack

    /* Restore n */
    pop    r16             ; r16 = original n (8-bit)

    /* Compute n × factorial(n-1) */
    ; r25:r24 = factorial(n-1), r16 = n
    ; Use 8 × 16 → 16-bit multiply:
    mul    r16, r24        ; r1:r0 = n × fact_lo
    movw   r22, r0         ; save partial result
    mul    r16, r25        ; r1:r0 = n × fact_hi (contributes to byte 1)
    add    r23, r0         ; r23 += (n × fact_hi) – upper byte of result
    clr    r1              ; restore r1 = 0 (ABI)
    movw   r24, r22        ; r25:r24 = final result
    ret

.Lfact_base:
    clr    r25
    ldi    r24, 1          ; return 1
    ret
```

8.7.1 Call Stack Trace for factorial(4)

```
factorial(4):
    push r24=4           ; stack: [4]
    call factorial(3):
```

```

push r24=3          ; stack: [4, 3]
call factorial(2):
  push r24=2        ; stack: [4, 3, 2]
  call factorial(1):
    n < 2: return 1    r25:r24 = 0x0001
  pop r16=2          ; stack: [4, 3]
  2 × 1 = 2          r25:r24 = 0x0002
  return
pop r16=3            ; stack: [4]
3 × 2 = 6            r25:r24 = 0x0006
return
pop r16=4            ; stack: []
4 × 6 = 24           r25:r24 = 0x0018
return

```

Result: 0x0018 = 24 OK

Stack depth at peak (inside factorial(1) call from factorial(4)):

- 4 frames × (2 bytes ret addr + 1 byte push r24) = **12 bytes**

That is the real cost of recursion in concrete terms: each level adds another copy of the call frame.

8.7.2 The mul8x16 Helper

The multiply step deserves its own routine for clarity:

```

/*
 * mul8x16 – 8-bit × 16-bit → 16-bit (truncated, lower 16 bits only)
 * Input:  r16 = 8-bit factor, r25:r24 = 16-bit factor
 * Output: r25:r24 = result (lower 16 bits)
 * Clobbers: r0, r1 (cleared), r22, r23
 *
 * Math: (r16) × (r25:r24)
 *       = r16×r24 + (r16×r25) × 256
 *
 * byte 0: (r16×r24).lo
 * byte 1: (r16×r24).hi + (r16×r25).lo
 * byte 2+: not computed (truncated)
 */
mul8x16:
  mul    r16, r24          ; r1:r0 = r16 × r24 (low byte of 16-bit factor)
  movw   r22, r0           ; r23:r22 = partial result
  mul    r16, r25          ; r1:r0 = r16 × r25 (high byte of 16-bit factor)
  add    r23, r0           ; r23 += contribution from high byte × r16
  clr    r1
  movw   r24, r22          ; return in r25:r24
  ret

```

8.8 Tail Call Optimisation

A **tail call** is a call that is the last action before returning. Instead of RCALL + RET, you can RJMP — the callee's RET becomes your return:

```

; Without tail call:
my_func:
    ; ... setup ...
    rcall other_func    ; call
    ret                ; immediately return after

; With tail call (saves one RCALL+RET = 7 cycles, 2 stack bytes):
my_func:
    ; ... setup ...
    rjmp other_func     ; jump; other_func's RET returns to MY caller

```

This works only when:

1. The result of other_func is passed directly to the caller (no post-processing)
2. No call-preserved registers need restoring after the call
3. SP and Y are in the correct state for other_func to use

Tail calls are most valuable in tight loops and deep call chains to avoid stack overflow.

The idea is simple: if you were going to return immediately after a call anyway, do not create an extra return layer you do not need.

8.9 Complete Function Template

Copy-paste starting point for any well-behaved function:

Templates like this matter in assembly because they reduce the chance of forgetting one push, one pop, or one ABI rule.

```

/*
 * my_function — brief description
 * Input:  r24 = first_arg (uint8_t)
 *         r23:r22 = second_arg (uint16_t)
 * Output: r25:r24 = result (uint16_t)
 * Preserved: r2-r17, r28-r29 (Y) per ABI
 * Clobbers: r18-r27, r30-r31 (call-clobbered per ABI)
 */
.global my_function
my_function:
    /* — Prologue ————— */
    push r14                ; save any call-preserved regs you use
    push r15
    ; push r28                ; uncomment if using Y as frame pointer
    ; push r29
    ; in  r28, SPL            ; Y = SP
    ; in  r29, SPH
    ; sbi r28, N              ; allocate N local bytes
    ; out SPH, r29
    ; out SPL, r28

    /* — Body ————— */
    ; ... your code here ...
    ; Arguments: r24 = first_arg, r23:r22 = second_arg
    ; Result:    build in r25:r24 before returning

```

```

/* — Epilogue ————— */
; adiw r28, N           ; if frame was allocated
; out SPH, r29
; out SPL, r28
; pop r29               ; if Y was saved
; pop r28
pop r15                 ; restore in reverse push order
pop r14
ret

```

8.10 Summary

RCALL k – call ± 2047 words, 3 cycles, 1 word. Default.
 CALL k – call full range, 4 cycles, 2 words. For far targets.
 ICALL – call via Z (word address), 3 cycles. For function pointers.
 RET – pop return address, 4 cycles.
 RETI – RET + set I flag. ISR exit only.

Stack: grows downward. PUSH: SP--, write. POP: read, SP++.
 Return address: 2 bytes on ATtiny3217.

GCC AVR ABI:

Arguments: left to right in register pairs descending from r25:r24
 examples: r24, r23:r22, r20, r19:r18 ...
 Return: r24 (8-bit), r25:r24 (16-bit)
 Clobbered: r0, r18-r27, r30-r31
 Preserved: r1(=0), r2-r17, r28-r29

r1 must always be 0. CLR r1 after every MUL/MULS/MULSU.

Frame pointer Y (r29:r28):

Prologue: push r28/r29, Y=SP, SBIW Y N, update SP.
 Access: LDD/STD Y+1 through Y+N.
 Epilogue: ADIW Y N, restore SP, pop r29/r28.

Stack usage per frame = 2 (ret addr) + pushes + SBIW allocation.

Tail call: replace RCALL+RET with RJMP to save frame + cycles.

8.11 Exercises

1. Write a function `uint16_t sum_range(uint8_t start, uint8_t end)` that returns the sum of all integers from start to end inclusive. Use the ABI: start in r24, end in r22. Verify `sum_range(1, 10) = 55`.
2. The factorial routine uses `rcall mul8x16` inside the recursive path. How many bytes of stack does `factorial(5)` consume at peak depth? Count: return addresses, pushed r24 values, and the `mul8x16` frame.
3. Rewrite `factorial` without recursion using a loop. Compare stack usage. Which is preferable for $n \leq 8$ on an embedded system?

4. Write a function `void swap(uint8_t *a, uint8_t *b)` where `a` is passed in `r25:r24` and `b` in `r23:r22` (pointers to SRAM). Swap the bytes at those addresses in-place using three XOR operations (no temporary byte on the stack).
5. Annotate `with_frame` from the source file with the exact stack layout (addresses, values) after every push, the `SBIW`, every `STD`, and the `ADIW`. Assume `SP` starts at `0x3FFF`.
6. Write a calling sequence in `main` that calls `factorial(8)` and then checks if the result equals `40320` (`0x9D80`). Branch to an error handler if it does not. Which branch instruction do you use and why?

Next: Chapter 9 — GPIO

Chapter 9

GPIO

GPIO (General-Purpose Input/Output) is the mechanism that connects the CPU to the physical world: sensors, LEDs, switches, relays, and anything else wired to a pin. This chapter covers the four dedicated GPIO instructions — IN, OUT, SBI, CBI — together with SBIC/SBIS for pin-state testing, and builds up to a fully debounced button-controlled LED.

The ATtiny3217 uses a **dual-register model** (VPORT + PORT) that differs from the classic DDRx/PORTx/PINx model found on ATmega parts. Both models are covered here; the worked example targets ATtiny3217.

That difference matters because a lot of AVR examples online assume the older ATmega naming. If you only memorize the old names, modern tinyAVR code will look stranger than it really is.

9.1 GPIO Fundamentals

Every GPIO pin has three properties the firmware controls:

Direction 0 = input (pin floats / follows external signal)
 1 = output (pin is driven high or low by the CPU)

Output 0 = drive low (when direction = output)
value 1 = drive high (when direction = output)

When direction = input:
0 = no pull-up (pin floats if nothing drives it)
1 = weak pull-up to VCC (prevents floating)

Input Read-only: reflects actual voltage on pin

On all AVR devices, GPIO is controlled through I/O registers. The exact register names and addresses differ between classic ATmega and the newer tinyAVR 1-series, but the concepts are identical.

So before worrying about the register map, keep the simple mental model in mind: configure direction, choose driven value or pull-up behavior, then read the pin state back when needed.

VPORTB_IN	0x06	0x0006	8-bit	R
VPORTB_INTFLAGS	0x07	0x0007	8-bit	R/W(W1C)
VPORTC_DIR	0x08	0x0008	8-bit	R/W
VPORTC_OUT	0x09	0x0009	8-bit	R/W
VPORTC_IN	0x0A	0x000A	8-bit	R
VPORTC_INTFLAGS	0x0B	0x000B	8-bit	R/W(W1C)

W1C = Write 1 to Clear (write 1 to a flag bit to clear it)

Note on tinyAVR 1-series addressing: On classic ATmega, there is a 0x20 offset between I/O addresses and memory addresses (I/O 0x00 = memory 0x0020). On ATtiny3217, this offset does not exist — I/O address N = memory address N. `_SFR_MEM8(0x0000)` and I/O address 0x00 are the same location.

VPORT registers are **aliases** of the full PORT registers. Writing to VPORTA_DIR is identical to writing to PORTA_DIR.

The big idea is that VPORT gives you the fast instruction-set view of a port, while the full PORT block gives you the richer configuration interface.

9.3.2 PORT — Full Port Registers (Configuration and Atomic Operations)

The full PORT peripheral sits in the extended memory range (above 0x3F) and cannot be accessed with IN/OUT/SBI/CBI. Use LDS/STS.

PORTA base: 0x0400

PORTB base: 0x0420

PORTC base: 0x0440

Offset	Register	Description
+0x00	DIR	Direction (0=in, 1=out) – full byte read/write
+0x01	DIRSET	Atomic direction set (write 1 to set bits)
+0x02	DIRCLR	Atomic direction clear (write 1 to clear bits)
+0x03	DIRTGL	Atomic direction toggle
+0x04	OUT	Output value – full byte read/write
+0x05	OUTSET	Atomic output set (write 1 to set bits)
+0x06	OUTCLR	Atomic output clear (write 1 to clear bits)
+0x07	OUTTGL	Atomic output toggle (write 1 to toggle bits)
+0x08	IN	Input (read-only: actual pin voltage)
+0x09	INTFLAGS	Interrupt flags – write 1 to clear a flag
+0x0A	PORTCTRL	Port control
+0x10	PIN0CTRL	Pin 0: ISC[2:0] + PULLUPEN + INVEN
+0x11	PIN1CTRL	Pin 1 configuration
...		
+0x17	PIN7CTRL	Pin 7 configuration

The **SET/CLR/TGL** variants perform atomic read-modify-write in hardware: a single STS `PORTA_OUTSET, r16` can set individual bits without disabling interrupts. Use these for multi-bit changes where atomicity matters (e.g., if an ISR also touches the same port).

This is one of the nicest parts of the newer AVR GPIO design. Instead of forcing software to read a whole port byte, modify one bit, and write the whole byte back, the peripheral can do the bit update internally.

The **PINnCTRL** registers at offsets 0x10–0x17 hold:

- ISC[2:0] (bits 2:0) — interrupt sense (covered in Chapter 10)

- PULLUPEN (bit 3) — enables the internal pull-up resistor
 - INVEN (bit 7) — inverts the pin logic
-

9.4 IN — Read I/O Register

IN Rd, A

Operands: Rd $\in$ {R0..R31}, A $\in$ {0..63} (I/O address)

Operation: Rd $\leftarrow$ I/O(A)

Flags: None

Cycles: 1

Encoding: 1011 0AA dddd AAAA

IN reads one byte from an I/O register into a general-purpose register. On ATtiny3217, VPORT registers sit at I/O addresses 0x00–0x0B and are within this range.

Think of IN as “take a snapshot of this register so the CPU can inspect it.”

```
in  r16, VPORTA_IN      ; r16 = current state of PORTA pins (8 bits)
in  r17, VPORTB_OUT     ; r17 = current output values of PORTB
in  r18, SREG           ; r18 = Status Register (a critical use of IN)
```

One common idiom: read SREG before a critical section, write it back after.

This shows up here because GPIO work and interrupt control often meet in the same small critical sections.

```
in  r24, SREG           ; save interrupt enable state
cli                               ; disable interrupts (Chapter 10)
; ... critical section ...
out SREG, r24           ; restore interrupt enable state exactly
```

9.5 OUT — Write I/O Register

OUT A, Rr

Operands: A $\in$ {0..63}, Rr $\in$ {R0..R31}

Operation: I/O(A) $\leftarrow$ Rr

Flags: None

Cycles: 1

Encoding: 1011 1AAr rrrr AAAA

OUT writes one byte from a register to an I/O register.

Where IN captures a register value, OUT pushes a whole register value back into the I/O space in one step.

```
out VPORTA_OUT, r16      ; drive all 8 PORTA output pins at once
out VPORTB_DIR, r17      ; set direction for all 8 PORTB pins
out SPL, r16             ; set Stack Pointer Low byte
```

Writing all 8 bits at once with OUT is often less convenient than the single-bit SBI/CBI instructions. Use OUT when you want to change multiple pins simultaneously (e.g., output a byte to an 8-bit bus).

So OUT is not the default for every GPIO task. It is the right tool when the whole byte matters.

9.6 SBI and CBI — Set / Clear Bit in I/O Register

SBI A, b

Operands: $A \in \{0..31\}$, $b \in \{0..7\}$ ← I/O addresses 0x00–0x1F only
 Operation: $I/O(A).b \leftarrow 1$
 Flags: None
 Cycles: 2

CBI A, b

Operands: $A \in \{0..31\}$, $b \in \{0..7\}$
 Operation: $I/O(A).b \leftarrow 0$
 Flags: None
 Cycles: 2

SBI and CBI are **atomic** single-bit operations. On ATtiny3217, the hardware internally does read-modify-write in one indivisible step: no interrupt can observe a partial state between the read and the write.

That atomicity is why these instructions are still so valuable even though the newer PORT peripheral has richer register options.

Range restriction: SBI/CBI only reach I/O addresses 0x00–0x1F (the lower 32 I/O registers). The VPORT registers (0x00–0x0B) are all within this range. Full PORT registers (0x0400+) require LDS/STS.

```
sbi VPORTA_DIR, 3      ; set PA3 as output  (VPORTA_DIR = I/O 0x00)
cbi VPORTA_OUT, 3      ; drive PA3 low     (VPORTA_OUT = I/O 0x01)
sbi VPORTA_OUT, 3      ; drive PA3 high
sbi VPORTB_DIR, 3      ; set PB3 as output
cbi VPORTB_DIR, 5      ; set PB5 as input
```

When to use SBI/CBI vs. full PORT OUTSET/OUTCLR:

Situation	Instruction	Why
Toggle one pin, no ISR conflict	SBI/CBI on VPORT	1 cycle faster
Set/clear bits, ISR may also touch port	STS PORTA_OUTSET	Guaranteed atomic
Set multiple bits at once	STS PORTA_OUTSET	Single write, no loop

If you remember only one distinction here, remember this one:

- SBI/CBI are ideal for one bit in low I/O space
- OUTSET/OUTCLR/OUTTGL are ideal for atomic updates in the full port block

9.7 SBIC and SBIS — Skip if Bit Clear / Set

SBIC A, b

Operands: $A \in \{0..31\}$, $b \in \{0..7\}$
 Operation: if $I/O(A).b = 0$, skip next instruction
 Flags: None
 Cycles: 1 (no skip), 2 (skip 1-word instruction), 3 (skip 2-word)

SBIS A, b

Operands: $A \in \{0..31\}$, $b \in \{0..7\}$
 Operation: if $I/O(A).b = 1$, skip next instruction

Flags: None
 Cycles: 1 (no skip), 2 (skip 1-word instruction), 3 (skip 2-word)

These are the fastest way to branch on a single I/O pin state. Instead of IN + SBRC/SBRS, a single instruction reads and conditionally skips in one operation.

That is why SBIC and SBIS feel so natural in polling loops: test plus control flow happens in one instruction.

Wait for a button press (PB2 active-low):

```
.Lwait:
    sbis VPORTB_IN, 2      ; skip if PB2 = 1 (not pressed)
    rjmp .Lpressed        ; PB2 = 0: button is pressed
    rjmp .Lwait
```

Equivalent without SBIS:

```
in    r16, VPORTB_IN      ; 1 cy: read all PORTB pins
sbrs  r16, 2              ; 1 cy: skip if bit 2 = 1
rjmp  .Lpressed           ; 2 cy: branch if bit was 0 (pressed)
rjmp  .Lwait
```

The SBIC/SBIS version saves one register (r16 untouched) and avoids the IN instruction.

Small savings like that matter in tight loops and ISR-adjacent code.

Skip behavior for 2-word instructions: SBIC/SBIS skip exactly one instruction slot. If the instruction immediately after is a 2-word instruction (JMP, CALL, LDS, STS, LDD, STD), the skip costs 3 cycles (the full 2-word instruction is skipped). Skip over a 1-word instruction costs 2 cycles.

So the timing of a skip depends not just on the skip instruction itself, but also on what comes immediately after it.

```
; Good: skip over a 1-word RJMP (common pattern)
sbic VPORTB_IN, 2      ; skip if PB2 = 0 (pressed)
rjmp .Lnot_pressed     ; 1 word – skipped in 2 cycles
```

```
; Also fine: skip over a 2-word STS
sbis VPORTA_IN, 3       ; skip if PA3 = 1
sts  PORTA_OUTTGL, r16  ; 2 words – skipped in 3 cycles
```

9.8 Configuring a GPIO Output

Two equivalent approaches on ATtiny3217:

Method A — SBI on VPORT (single pin, 2 cycles):

```
sbi  VPORTA_DIR, 3      ; PA3 = output
sbi  VPORTA_OUT, 3      ; initial value: high (Curiosity Nano LED off)
```

Method B — STS with PORT DIRSET/OUTCLR (atomic, any number of pins):

```
ldi  r16, (1 << 3)
sts  PORTA_DIRSET, r16   ; PA3 = output, other bits unchanged
sts  PORTA_OUTSET, r16   ; PA3 driven high: Curiosity Nano LED off
```

Both methods produce identical hardware results for a single pin. Method B is preferred when configuring multiple pins in a single write:

That means you can choose the style based on clarity and scaling, not because one is inherently more “correct” for a single pin.

```
ldi r16, (1 << 3) | (1 << 5) ; configure PA3 and PA5
sts PORTA_DIRSET, r16        ; both pins become outputs
sts PORTA_OUTCLR, r16       ; both driven low
```

9.9 Configuring a GPIO Input with Pull-Up

On ATtiny3217, pin direction defaults to input (0) at reset, so only the pull-up needs configuration. Pull-up lives in the PINnCTRL register, which is above the I/O space — LDS/STS required.

PINnCTRL register layout:

```
Bit: 7      6:3      3      2:0
      INVEN  reserved PULLUPEN ISC[2:0]
```

```
PULLUPEN (bit 3) = 0: pull-up off (pin floats if nothing drives it)
                  = 1: pull-up on  (pin pulled to VCC through ~100 kΩ)
ISC[2:0]         = 0x00: interrupt disabled, input buffer enabled
                  (other values in Chapter 10)
PORT_PULLUPEN_bm = 0x08
```

The floating input problem: If a digital input pin is not connected to anything and has no pull-up, its voltage floats between GND and VCC. The CPU reads random logic levels that change with temperature, noise, and nearby switching signals. Always enable the pull-up when using a button wired to GND, or drive the pin from a source you control.

This is one of the first hardware realities software people run into on a microcontroller: an unconnected input is not neutral, it is unreliable.

```
; PB2 as input with pull-up (button wired to GND)
; Direction is already 0 (input) – nothing to do for DIR.
ldi r16, PORT_PULLUPEN_bm      ; = 0x08
sts PORTB_PIN2CTRL, r16        ; PORTB_PIN2CTRL = memory 0x0432

; PB3 as input, NO pull-up (externally driven signal, e.g. UART RX)
ldi r16, 0
sts PORTB_PIN3CTRL, r16        ; ISC = disabled, PULLUPEN = 0
```

Reading a pin:

```
in r16, VPORTB_IN              ; read all 8 PORTB pins
sbrc r16, 2                    ; skip if bit 2 = 0 (button pressed)
; ... bit 2 = 1: button not pressed ...


; Or more directly:
sbis VPORTB_IN, 2              ; skip if PB2 = 1 (not pressed)
rjmp .Lpressed
```

9.10 Debounce

Mechanical switch contacts **bounce**: when you press or release a button, the contacts make and break contact several times in rapid succession before settling. The CPU is fast enough to see each bounce as a separate

button event.

So debouncing is not a cosmetic improvement. It is what turns a noisy physical event into a single logical event.



The standard software fix: after detecting a press, wait a few milliseconds (bounce settles in < 10 ms for most switches), then re-check. Do the same on release.

Simple time-delay debounce:

```
; Wait for press
.Lpoll:
    sbis VPORTB_IN, BTN_BIT    ; skip if not pressed
    rjmp .Lpressed
    rjmp .Lpoll

.Lpressed:
    rcall delay_20ms           ; wait for bounce to settle

    ; Optional: confirm still pressed (guards against glitches)
    sbic VPORTB_IN, BTN_BIT    ; skip if still pressed (PB2 = 0)
    rjmp .Lpoll                ; was a glitch - restart

    ; Confirmed: genuine press
    ; ... handle event ...

    ; Wait for release
.Lwait_release:
    sbis VPORTB_IN, BTN_BIT    ; skip if released (PB2 = 1)
    rjmp .Lwait_release
    rcall delay_20ms           ; debounce the release
    rjmp .Lpoll
```

For typical tactile buttons, 20 ms is a safe debounce window. Some mechanical switches need 50 ms.

That delay is a human-timescale compromise: long enough to outlast bounce, but short enough that the user does not notice it.

9.11 Worked Example: Button-Controlled LED

LED0 on PA3 (Curiosity Nano, active-low). External button on PB2 (active-low, internal pull-up). Each confirmed press toggles the LED state.

This example combines almost every GPIO concept in one small program: output setup, input pull-up, active-low logic, skip-based polling, debounce, and atomic output update.

Hardware:

ATtiny3217

VCC
|
PA3 — Curiosity Nano LED0 (on board, active-low)

PB2 — [10 kΩ internal pull-up] — VCC

PB2 — [button] — GND

Source file: src/gpio.S

```
#include <avr/io.h>

.equ LED_BIT,      3          ; PA3, Curiosity Nano LED0
.equ BTN_BIT,      2          ; PB2
.equ DEBOUNCE_COUNT, 16666    ; ~20 ms at 3.3 MHz (sbiw+brne loop)

.section .vectors, "ax", @progbits
.global __vectors
__vectors:
    rjmp main

.section .text
.global main

main:
    ldi r16, lo8(RAMEND)
    out SPL, r16
    ldi r16, hi8(RAMEND)
    out SPH, r16

    sbi VPORTA_DIR, LED_BIT    ; PA3 = output
    sbi VPORTA_OUT, LED_BIT    ; LED off (active-low)

    ldi r16, PORT_PULLUPEN_bm ; 0x08
    sts PORTB_PIN2CTRL, r16    ; PB2: input + pull-up, no interrupt

.Lloop:
    sbis VPORTB_IN, BTN_BIT    ; skip if PB2 = 1 (not pressed)
    rjmp .Lpressed
    rjmp .Lloop

.Lpressed:
    rcall delay_20ms
    sbic VPORTB_IN, BTN_BIT    ; skip if PB2 is still low (still pressed)
    rjmp .Lloop                ; bounce/glitch: ignore

    ldi r16, (1 << LED_BIT)
    sts PORTA_OUTTGL, r16      ; atomic toggle PA3

.Lwait_release:
    sbis VPORTB_IN, BTN_BIT    ; skip if PB2 = 1 (released)
    rjmp .Lwait_release
    rcall delay_20ms
```

```

    rjmp .Lloop

delay_20ms:
    ldi r26, lo8(DEBOUNCE_COUNT)
    ldi r27, hi8(DEBOUNCE_COUNT)
.Ldly_loop:
    sbiw r26, 1           ; 2 cycles
    brne .Ldly_loop       ; 2 cycles taken
    ret

```

9.11.1 Why PORTA_OUTTGL instead of SBI/CBI?

SBI VPORTA_OUT, 3 sets the bit; CBI VPORTA_OUT, 3 clears it. Neither alone toggles — you'd need to read the current state first. The naive approach:

```

; Non-atomic toggle – dangerous if ISR touches PORTA_OUT
in    r16, VPORTA_OUT
ldi   r17, (1 << LED_BIT)
eor   r16, r17           ; flip LED bit
out   VPORTA_OUT, r16    ; write back

```

Three instructions. An interrupt that fires between IN and OUT and modifies another PA pin will have its change overwritten by the OUT. The PORTA_OUTTGL approach is atomic: the hardware performs the read, XOR, and write in one step.

That is the real lesson: the shorter code is not just shorter, it is safer when something else can touch the same port state.

9.11.2 Build and Inspect

```

# .S (uppercase) is preprocessed before assembly.
# Both avr-as file.S and avr-gcc -c file.S work; avr-gcc is convenient here.
avr-gcc -mmcu=attiny3217 -c -o gpio.o src/gpio.S
avr-gcc -mmcu=attiny3217 -nostartfiles -o gpio.elf gpio.o
avr-objdump -d gpio.elf      # inspect disassembly
avr-objcopy -O ihex gpio.elf gpio.hex  # generate flashable HEX

```

Check the disassembly around .Lpressed to confirm that sts PORTA_OUTTGL encodes as a 2-word sts instruction (opcode 0x92..) with the 16-bit address 0x0407.

9.12 Summary

IN Rd, A	– read I/O register A (0-63) into Rd.	1 cycle.
OUT A, Rr	– write Rr to I/O register A (0-63).	1 cycle.
SBI A, b	– set bit b of I/O register A (0-31).	2 cycles. Atomic.
CBI A, b	– clear bit b of I/O register A (0-31).	2 cycles. Atomic.
SBIC A, b	– skip next if I/O(A).b = 0.	1/2/3 cycles.
SBIS A, b	– skip next if I/O(A).b = 1.	1/2/3 cycles.

SBI/CBI/SBIC/SBIS range: I/O addresses 0x00-0x1F only.

VPORT registers (0x00-0x0B) are within this range: all six instructions apply.

Full PORT registers (0x0400+) need LDS/STS.

ATtiny3217 GPIO quick-reference:

```

Configure output:  sbi VPORTA_DIR, n    or  sts PORTA_DIRSET, r16
Drive high:       sbi VPORTA_OUT, n    or  sts PORTA_OUTSET, r16
Drive low:        cbi VPORTA_OUT, n    or  sts PORTA_OUTCLR, r16
Toggle output:    sts PORTA_OUTTGL, r16 (atomic)
Configure input:  (DIR default = 0 – no action required)
Enable pull-up:   sts PORTB_PIN2CTRL, r16 (PORT_PULLUPEN_bm = 0x08)
Read input:       in r16, VPORTB_IN    or  sbic/sbis VPORTB_IN, n

```

Debounce: after detecting edge, wait ~20 ms, re-check, wait for release,
wait ~20 ms again before returning to poll loop.

9.13 Exercises

1. Reconfigure the example so the LED is on PC0 instead of PA3. Which registers change? Which instructions change? Is there any difference in the generated machine code size?
 2. Write a routine `toggle_n_times` that accepts a count in `r24` and toggles PA3 that many times with a 100 ms delay between each toggle. Use `PORTA_OUTTGL` and a subroutine-based delay.
 3. The debounce loop calls `delay_20ms` twice: once after press, once after release. Under what button-press pattern could the LED be toggled without a real press if you removed the release debounce?
 4. Rewrite the button poll using `SBIC` instead of `SBIS`. Draw the new control flow. Are the two approaches equivalent in terms of cycles per loop iteration?
 5. How many cycles does the `delay_20ms` subroutine take from `RCALL` to the final `RET`? Show the sum: prologue, loop body $\times N$, final loop exit, return.
 6. Add a second button on PA5 (active-low, pull-up). When BTN1 (PB2) is pressed, turn the LED on; when BTN2 (PA5) is pressed, turn it off. Use only `SBIC/SBIS` — no `IN` instruction.
-

Next: Chapter 10 — Interrupts

Chapter 10

Interrupts

Polling burns CPU cycles checking whether something has happened. An **interrupt** lets hardware signal the CPU directly: the CPU pauses whatever it was doing, runs a short handler (ISR), and resumes exactly where it left off. This chapter covers the interrupt mechanism, the ATtiny3217 vector table, the SEI/CLI/RETI instructions, and the discipline required to write a correct ISR.

This is where the program stops being a single straight-line thread of control and starts reacting to the outside world asynchronously.

10.1 Why Interrupts?

The polling loop from Chapter 9:

```
.Lloop:
    sbis VPORTB_IN, BTN_BIT    ; check button every iteration
    rjmp .Lpressed
    rjmp .Lloop
```

This consumes 100% of the CPU checking one pin. In a real application, the CPU may need to: update a display, process serial data, manage a state machine — all while remaining responsive to button presses. Polling every possible event source and handling each one in priority order becomes unwieldy.

Interrupts invert the model: the CPU runs its main work and is **notified** when an event occurs. The ISR handles the event in a bounded time, then execution resumes with main code unmodified.

That last requirement is the hard part. An interrupt is only useful if the main code can continue as though nothing important was damaged by the pause.

10.2 The Interrupt Mechanism

10.2.1 What Happens When an Interrupt Fires

On ATtiny3217, think of interrupt entry as a short fixed hardware sequence: finish the current instruction, clear I, push the PC, then execute the vector entry. With the RJMP-based vector table used in this book, the first ISR instruction starts after roughly **5 cycles**, plus any extra cycles needed to finish a multi-cycle instruction already in flight.

That means interrupt latency is not arbitrary. It is mostly predictable, which is one reason small microcontrollers work well for real-time event handling.

1. CPU finishes the current instruction.
2. I flag in SREG is cleared (prevents re-entry / nested interrupts).
3. PC is pushed onto the stack (2 bytes on ATtiny3217).
4. The interrupt vector entry executes (`RJMP` in this chapter's layout).
5. CPU begins the first instruction of the ISR body.

When the ISR executes RETI:

1. PC is popped from the stack ($SP += 2$).
2. I flag in SREG is set (re-enables interrupts).
3. CPU resumes the instruction after where it was interrupted.

The interrupted code is **unaware** an interrupt occurred — as long as the ISR leaves all registers and SREG exactly as it found them.

That “as long as” is the whole discipline of ISR writing.

10.2.2 The I Flag (Global Interrupt Enable)

SREG bit layout:

```

Bit:  7   6   5   4   3   2   1   0
      I   T   H   S   V   N   Z   C
      ↑
      Global Interrupt Enable
      1 = interrupts enabled
      0 = interrupts disabled (masked)

```

The I flag is a master switch. Even if a peripheral has its interrupt configured and pending, the CPU will not respond until $I = 1$. Setting I with SEI and clearing it with CLI are the only ways to control this bit (other than OUT SREG, Rr and RETI).

So enabling a peripheral interrupt is never the whole story. The peripheral must be configured, and the global interrupt gate must also be open.

10.2.3 Interrupt Priority

On ATtiny3217, lower vector address = higher priority. RESET (0x0000) has the highest priority. If two interrupts become pending simultaneously, the one with the lower vector address fires first.

Only one interrupt runs at a time by default (I is cleared on entry). To enable nested interrupts (an ISR interrupted by a higher-priority ISR), call SEI inside the ISR — but this requires careful stack depth analysis and is rarely needed in practice.

For most firmware, “no nesting unless you have a strong reason” is the right default.

10.3 SEI and CLI

SEI

```

Operation: SREG.I ← 1
Flags:     I = 1
Cycles:    1
Note:      The instruction immediately after SEI executes before any
           pending interrupt fires.

```

CLI

Operation: $SREG.I \leftarrow 0$
 Flags: $I = 0$
 Cycles: 1
 Note: Guaranteed to prevent any interrupt from firing after CLI.

The one-instruction delay after SEI is important:

```
; Guaranteed: the RJMP executes before any pending interrupt fires.
sei
rjmp .Lidle
```

That behavior prevents awkward races where an interrupt could fire immediately between SEI and the next instruction.

10.3.1 Critical Sections

CLI/SEI create a **critical section** — a code region that cannot be interrupted:

```
cli                ; disable interrupts – begin critical section
; ... read or write shared variables that the ISR also touches ...
sei                ; re-enable interrupts – end critical section
```

A cleaner pattern that preserves the interrupt state rather than unconditionally re-enabling:

```
in  r24, SREG      ; save current SREG (including I flag)
cli                ; disable interrupts
; ... critical section ...
out SREG, r24      ; restore SREG exactly (may or may not re-enable)
```

This is safe if the caller had interrupts disabled — SEI would have incorrectly re-enabled them, but restoring SREG preserves their disabled state.

This is the pattern to prefer in reusable code, because it does not assume that interrupts were enabled before the critical section began.

10.4 RETI — Return from Interrupt

RETI

Operation: $PC \leftarrow \text{stack (pop 2 bytes)}, SP \leftarrow SP + 2, SREG.I \leftarrow 1$
 Flags: $I = 1$
 Cycles: 4

RETI differs from RET in exactly one architectural way: it **sets the I flag**, re-enabling interrupts. Never use RET to exit an ISR:

Using RET at end of ISR:	Using RETI at end of ISR:
I flag stays 0 after return.	I flag set to 1 after return.
No more interrupts fire.	Normal interrupt operation resumes.
System appears to "freeze" (all interrupts silently ignored forever).	

The gap between executing the `OUT SREG, r24` (in the epilogue) and RETI is safe: RETI atomically pops PC and sets I. No interrupt can fire in that gap.

That is why restoring SREG before RETI is correct even though it can look, at first glance, like it might reopen interrupts too early.

10.5 The Interrupt Vector Table

10.5.1 Layout on ATtiny3217

The vector table occupies the first 62 bytes of flash (31 entries $\times$ 2 bytes each). Every entry is one RJMP instruction — JMP (4 bytes) would overflow the 2-byte slot.

Byte addr	Vector num	Peripheral	Trigger condition
0x0000	—	RESET	Power-on, external reset, WDT
0x0002	1	CRCSCAN_NMI	CRC scan failure (non-maskable)
0x0004	2	BOD_VLM	Voltage level monitor
0x0006	3	PORTA_PORT	Any PORTA pin with ISC configured
0x0008	4	PORTB_PORT	Any PORTB pin with ISC configured
0x000A	5	PORTC_PORT	Any PORTC pin with ISC configured
0x000C	6	RTC_CNT	RTC counter/compare
0x000E	7	RTC_PIT	RTC periodic interrupt
0x0010	8	TCA0_OVF	Timer/Counter A overflow
0x0012	9	TCA0_HUNF	TCA high-byte underflow (split)
0x0014	10	TCA0_CMP0	TCA compare channel 0
0x0016	11	TCA0_CMP1	TCA compare channel 1
0x0018	12	TCA0_CMP2	TCA compare channel 2
0x001A	13	TCB0_INT	Timer/Counter B0
0x001C	14	TCB1_INT	Timer/Counter B1
0x001E	15	TCD0_OVF	Timer/Counter D overflow
0x0020	16	TCD0_TRIG	TCD trigger
0x0022	17	AC0_AC	Analog Comparator 0
0x0024	18	AC1_AC	Analog Comparator 1
0x0026	19	AC2_AC	Analog Comparator 2
0x0028	20	ADC0_RESRDY	ADC result ready
0x002A	21	ADC0_WCOMP	ADC window compare
0x002C	22	ADC1_RESRDY	ADC1 result ready
0x002E	23	ADC1_WCOMP	ADC1 window compare
0x0030	24	TWI0_TWIS	TWI slave
0x0032	25	TWI0_TWIM	TWI master
0x0034	26	SPI0_INT	SPI
0x0036	27	USART0_RXC	USART RX complete
0x0038	28	USART0_DRE	USART data register empty
0x003A	29	USART0_TXC	USART TX complete
0x003C	30	NVMCTRL_EE	EEPROM ready

10.5.2 Defining the Vector Table in .S

```
.section .vectors, "ax", @progbits
.global __vectors
__vectors:
    rjmp main          /* 0x0000: RESET — must be first */
    rjmp isr_default   /* 0x0002: CRCSCAN_NMI */
```



```

rjmp isr_default      /* 0x0004: BOD_VLM                */
rjmp isr_default      /* 0x0006: PORTA_PORT                */
rjmp portb_isr        /* 0x0008: PORTB_PORT ← our handler */
rjmp isr_default      /* 0x000A: PORTC_PORT                */
/* ... continue for all 30 remaining vectors ... */

```

Each consecutive RJMP is exactly 2 bytes, so the *n*th entry lands at byte address *n*×2 automatically — no `.org` directives needed.

Unhandled vectors should point to a catch-all handler that just RETIs, not to `main` (re-running `main` on a spurious interrupt would re-initialize your whole program). A simple trap for debugging: `rjmp .` (infinite loop).

Which default you choose depends on your goal:

- `reti` if you want unused interrupts to be harmless
- `rjmp .` if you want unexpected interrupts to halt the system for debugging

```

isr_default:
    reti                    /* ignore unexpected interrupts silently */

```

10.6 ISR Prologue and Epilogue

An ISR may fire at any point during the main code: between any two instructions. The main code has no way to know whether an ISR ran during a particular period. Therefore:

Every register written by an ISR must be saved before use and restored before RETI.

SREG must always be saved. Any instruction that modifies flags (arithmetic, logic, shifts, compares, even SBIW) will corrupt the flags that the interrupted code was relying on.

These are the core ISR rules. Violating them may not fail immediately, which is why ISR bugs can be so frustrating to diagnose.

10.6.1 Minimal ISR (uses r16 only)

```

my_isr:
    push    r16              ; save r16 (2 cycles, 1 byte stack)
    in      r16, SREG        ; save SREG into r16
    push    r16              ; save SREG on stack (2 cycles, 1 byte)

    ; ... ISR body, may freely use r16 ...

    pop     r16              ; restore SREG value into r16
    out     SREG, r16        ; restore SREG
    pop     r16              ; restore r16
    reti

```

10.6.2 Why This Order?

Step	Stack contents (grows down)	SP points to
(on ISR entry)	[ret_hi][ret_lo]	top of ret addr
push r16	[ret_hi][ret_lo][r16]	saved r16
in r16, SREG	r16 = SREG value	
push r16	[ret_hi][ret_lo][r16][SREG_val] ← SP	

```

... body ...
pop r16          r16 = SREG_val, stack shrinks
out SREG, r16    SREG restored – I flag = 0 (was 0 when ISR entered)
pop r16          r16 = original r16 value
reti            pop return addr, set I=1, jump back

```

OUT SREG, r16 writes I=0 (the I flag was 0 when IN r16, SREG captured it, because the CPU cleared I on interrupt entry). This is intentional: SREG is restored to its pre-interrupt state, and RETI then sets I=1. The two steps together (restore SREG, then RETI sets I) are why you cannot use RET — RET would leave I=0.

That subtle interaction is one of the main reasons RETI exists as a distinct instruction instead of reusing ordinary RET.

10.6.3 Saving Multiple Registers

```

my_isr:
    push    r16
    push    r17
    push    r18
    in      r16, SREG
    push    r16                ; SREG last – now on top of stack

    ; ... ISR body using r16, r17, r18 ...

    pop     r16                ; SREG value
    out     SREG, r16
    pop     r18                ; restore in reverse order
    pop     r17
    pop     r16
    reti

```

10.6.4 ISR Stack Budget

Every PUSH costs 2 cycles and 1 byte of stack. The two-register minimal ISR (r16 + SREG) uses:

- 2 bytes for the return address (pushed automatically by CPU)
- 2 bytes for the two PUSH instructions

Total: 4 bytes of stack per invocation. If main code uses 200 bytes of stack and ISRs nest, add the ISR frame size to your stack budget.

So ISR design is not just about speed. It is also about bounded stack usage.

10.7 Configuring Pin Interrupts on ATtiny3217

Classic ATmega devices have dedicated INT0/INT1 external interrupt pins. ATtiny3217 takes a different approach: **any pin** can generate an interrupt by configuring its PINnCTRL register.

This is more flexible than the old dedicated-pin model, but it also means you think in terms of per-pin configuration rather than special global interrupt pins.

10.7.1 PINnCTRL Layout

```

Bit:    7      6:4      3      2:0
        INVEN  reserved  PULLUPEN  ISC[2:0]

```

PULLUPEN = 1: enable internal pull-up (bit 3 = 0x08)

ISC[2:0]: Interrupt Sense Configuration

```
0x00 = PORT_ISC_INTDISABLE_gc - interrupt disabled, input buffer on
0x01 = PORT_ISC_BOTHEDGES_gc  - interrupt on both edges (rise + fall)
0x02 = PORT_ISC_RISING_gc     - interrupt on rising edge (low → high)
0x03 = PORT_ISC_FALLING_gc    - interrupt on falling edge (high → low)
0x04 = PORT_ISC_INPUT_DISABLE_gc - input buffer off (ADC pins, saves power)
0x05 = PORT_ISC_LEVEL_gc      - interrupt while pin is low (level triggered)
```

For a button wired to GND (active-low, fires on press = falling edge):

```
ldi r16, (PORT_PULLUPEN_bm | PORT_ISC_FALLING_gc) ; = 0x08 | 0x03 = 0x0B
sts PORTB_PIN2CTRL, r16 ; PB2: pull-up + falling edge interrupt
```

10.7.2 PORT INTFLAGS — Clearing Interrupt Flags

When an interrupt fires, the corresponding bit in PORT_INTFLAGS is set. The flag is **sticky**: it stays set until explicitly cleared. If you return from the ISR without clearing the flag, the interrupt fires again immediately.

That sticky-flag behavior is easy to forget the first time you write an ISR. If the flag is not cleared, the hardware still believes the event is pending.

Clearing: write 1 to the flag bit (write-1-to-clear, or W1C semantics).

Method A — STS with full PORT address (any bit):

```
ldi r16, (1 << BTN_BIT) ; bit 2 = PB2
sts PORTB_INTFLAGS, r16 ; write 1 to bit 2 → clears flag
```

Method B — SBI on VPORTB_INTFLAGS (single bit, faster):

```
sbi VPORTB_INTFLAGS, BTN_BIT ; VPORTB_INTFLAGS = I/O 0x07
; SBI writes 1 to bit 2 → clears flag
```

SBI writing a 1 to a W1C register clears the flag. This is counterintuitive (SBI normally “sets” a bit) but correct: the hardware interprets a write-1 as “clear this flag”.

This is one of those peripheral rules you simply have to memorize: W1C means “write one to clear,” even if the instruction name sounds like the opposite.

The flag for PORTB_PORT covers all PORTB pins. If multiple pins are configured with ISC ≠ 0, read PORTB_INTFLAGS first to determine which pin(s) fired before clearing.

10.8 ISR Latency and Timing

With the RJMP vector style used here, the CPU reaches the first ISR instruction in roughly **5 cycles** after the current instruction completes. If the interrupt arrives during a multi-cycle instruction, that instruction must finish first, so total observed latency is higher.

```
Finish current instruction
Push 2-byte PC to stack
Execute vector-table RJMP
Begin ISR body
```

At 3.3 MHz, one cycle is about 303 ns, so 5 cycles is about 1.5 us before the ISR body starts, not counting any extra wait for a multi-cycle interrupted instruction or synchronizer delay in the peripheral itself.

That is fast enough for many GPIO, timer, and serial-response tasks, but still slow enough that cycle counting matters in tight timing work.

The minimal ISR prologue shown here (PUSH r16, IN r16, SREG, PUSH r16) adds **5 cycles** before the first useful ISR work:

- PUSH = 2 cycles
- IN = 1 cycle
- PUSH = 2 cycles

For time-critical ISRs, minimize the prologue: save only what you use.

In ISR code, every unnecessary instruction is paid on every event.

10.9 Worked Example: Pin Interrupt Toggles LED

Curiosity Nano LED0 on PA3, external button on PB2 (falling edge). The ISR toggles the LED; main loops forever doing nothing.

This is intentionally minimal so the interrupt path is easy to see in isolation.

Source file: src/interrupts.S

```
#include <avr/io.h>

.equ LED_BIT,    3
.equ BTN_BIT,    2
.equ BTN_CTRL,   (PORT_PULLUPEN_bm | PORT_ISC_FALLING_gc) /* 0x0B */

/* — Vector table — */
.section .vectors, "ax", @progbits
.global __vectors
__vectors:
    rjmp main          /* 0x0000: RESET */
    rjmp isr_default   /* 0x0002: CRCSCAN_NMI */
    rjmp isr_default   /* 0x0004: BOD_VLM */
    rjmp isr_default   /* 0x0006: PORTA_PORT */
    rjmp portb_isr     /* 0x0008: PORTB_PORT */
    rjmp isr_default   /* 0x000A: PORTC_PORT */
    /* ... (remaining 25 vectors omitted for brevity, all isr_default) .. */

/* — Catch-all ISR — */
.section .text

isr_default:
    reti

/* — PORTB ISR — */
/*
 * Prologue:    PUSH r16, IN r16, SREG, PUSH r16      (5 cycles, 2 stack)
 * Body:        LDI r16, STS PORTA_OUTTGL, SBI VPORTB_INTFLAGS
 * Epilogue:    POP r16, OUT SREG-r16, POP r16, RETI  (9 cycles, 0 stack)
 */
portb_isr:
    push r16
```

```

    in    r16, SREG
    push r16

    ldi   r16, (1 << LED_BIT)
    sts   PORTA_OUTTGL, r16      /* toggle PA3 / LED0          */
    sbi   VPORTB_INTFLAGS, BTN_BIT /* clear flag (W1C via SBI) */

    pop   r16
    out   SREG, r16
    pop   r16
    reti

/* — main ————— */
.global main

main:
    ldi   r16, lo8(RAMEND)
    out   SPL, r16
    ldi   r16, hi8(RAMEND)
    out   SPH, r16

    sbi   VPORTA_DIR, LED_BIT    /* PA3 = output          */
    sbi   VPORTA_OUT, LED_BIT    /* LED off (active-low)  */

    ldi   r16, BTN_CTRL
    sts   PORTB_PIN2CTRL, r16    /* PB2: pull-up + falling edge interrupt */

    sei                                     /* enable global interrupts */

.Lidle:
    rjmp  .Lidle                 /* CPU idles, ISR handles everything */

```

10.9.1 Execution Flow Walkthrough

Power-on:

RESET vector (0x0000) → RJMP main → main executes
 Configures PA3 output, PB2 pull-up + ISC, SEI
 Enters .Lidle loop

User presses button:

PB2 falls from 1 → 0 (falling edge)
 PORTB_PIN2CTRL ISC detects edge → sets PORTB_INTFLAGS bit 2
 CPU detects pending interrupt (I=1, so enabled)
 Finishes current RJMP .Lidle instruction
 Clears I, pushes PC onto stack
 Fetches PORTB_PORT vector (0x0008) → RJMP portb_isr
 Executes portb_isr:
 push r16 / save SREG
 sts PORTA_OUTTGL → PA3 toggles (LED changes state)
 sbi VPORTB_INTFLAGS, 2 → clears flag
 restore SREG / pop r16
 RETI → pop PC, set I=1, jump back to .Lidle

User presses button again:

Same sequence, LED toggles back

10.9.2 Why Not Debounce in the ISR?

The Chapter 9 polling example used a 20 ms delay loop after each press. That loop busy-waits, which is unacceptable inside an ISR (other time-critical events would be missed, and ISR latency would explode to 20 ms).

Options for interrupt-based debouncing:

1. **Hardware debounce:** RC + Schmitt trigger on the button line.
2. **Timer debounce:** in the ISR, start a timer; the real handler fires when the timer expires (Chapter 11). Until the timer expires, disable the pin interrupt.
3. **Level-based re-arm:** change ISC to ISC_LEVEL or use a state machine in main that re-enables the interrupt after a delay.

For this example, the button is held long enough that multiple edges do not cause visible problems. Production firmware should use method 2 or 3.

That is an important distinction between a teaching example and production code: the example demonstrates the interrupt mechanism cleanly, not the final debounce architecture you would ship.

10.9.3 Build and Inspect

```
avr-gcc -mmcu=attiny3217 -c -o interrupts.o src/interrupts.S
avr-gcc -mmcu=attiny3217 -nostartfiles -o interrupts.elf interrupts.o
avr-objdump -d interrupts.elf
avr-objcopy -O ihex interrupts.elf interrupts.hex
```

Inspect the disassembly to confirm:

- The vector table starts at address 0x0000 with a sequence of rjmps.
- portb_isr is at the address pointed to by the entry at 0x0008.
- The sts PORTA_OUTTGL inside portb_isr is a 4-byte instruction.
- sbi VPORTB_INTFLAGS, 2 is a 2-byte instruction (opcode 0x9ABF or similar — check with avr-objdump -d).

10.10 Summary

SEI — set I flag (enable global interrupts). 1 cycle.
 CLI — clear I flag (disable global interrupts). 1 cycle.
 RETI — return from ISR: pop PC, set I=1. 4 cycles. Never use RET here.

Interrupt response sequence (ATtiny3217, roughly 5 cycles with RJMP vectors):
 finish current instruction → clear I / push PC → execute vector RJMP → ISR

ISR discipline:

1. PUSH all registers used. (2 cy, 1 stack byte each)
2. IN r16, SREG → PUSH r16. (save SREG — always)
3. ISR body.
4. POP r16 → OUT SREG, r16. (restore SREG)
5. POP all registers in reverse order.
6. RETI. (not RET!)

Vector table: .section .vectors — sequential RJMP instructions.

2 bytes each. Entry n at byte address $n \times 2$. RESET at `0x0000`.
 Unhandled vectors $\rightarrow$ `isr_default: reti`.

ATtiny3217 pin interrupt config:

```
sts PORTx_PINnCTRL, r16
    PORT_PULLUPEN_bm      = 0x08 (bit 3: pull-up)
    PORT_ISC_FALLING_gc    = 0x03 (bits 2:0: falling edge)
    PORT_ISC_RISING_gc     = 0x02
    PORT_ISC_BOTHEDGES_gc = 0x01
```

Clearing INTFLAGS (W1C):

```
sts PORTx_INTFLAGS, r16      (STS to full PORT – any bit)
sbi VPORTx_INTFLAGS, n      (faster – SBI writes 1  $\rightarrow$  clears flag)
```

Must be done before RETI or interrupt fires again.

10.11 Exercises

1. The current ISR saves r16 and SREG. If the ISR body also needed to use r17 and r18, where exactly (before/after which existing instruction) should the PUSH r17 and PUSH r18 go? Write the complete modified prologue and epilogue.
 2. What happens if you replace `reti` with `ret` at the end of `portb_isr`? Describe the symptom you would observe and explain why it occurs in terms of the I flag.
 3. The idle loop is `rjmp .Lidle`. If you replaced it with the AVR sleep instruction (opcode 0x9588), the CPU would enter idle sleep mode and halt until the next interrupt. What changes would be required in `main` to enable sleep mode? (Hint: the SLPCTRL peripheral needs to be configured — look at Chapter 1's linker-script walkthrough or the ATtiny3217 datasheet section on SLPCTRL.)
 4. Configure PB3 for a second button that also fires on falling edge. In `portb_isr`, read `PORTB_INTFLAGS` to determine which pin fired, clear only that flag, and call one of two handlers: `led_on` or `led_off`. Write the complete modified ISR.
 5. Using `avr-objdump -d interrupts.elf`, count the exact number of bytes used by the vector table. Does it match $31 \text{ vectors} \times 2 \text{ bytes} = 62 \text{ bytes}$? What byte address does the first non-vector instruction appear at?
 6. The ISR uses `sts PORTA_OUTTGL, r16` (a 4-byte instruction, 2 cycles). Could you replace it with a 2-byte instruction using the VPORT registers? If so, write the alternative. If not, explain why.
-

Next: Chapter 11 — Timer/Counter

Chapter 11

Timer/Counter

Timers are the heartbeat of embedded firmware. They let you generate precise waveforms, schedule periodic work, measure elapsed time, and produce PWM signals — all without busy-waiting and without tying up the CPU. This chapter covers the TCA0 (Timer/Counter Type A) peripheral on the ATtiny3217, its three waveform-generation modes, and the interrupts each mode can generate. The worked example produces a 1 Hz LED blink driven entirely by a CMP0 compare-match interrupt.

This is where you start handing time-sensitive repetition over to hardware instead of burning CPU cycles in delay loops.

11.1 Timer Fundamentals

11.1.1 What a Timer Actually Is

A timer is a hardware counter clocked by a divided version of the CPU clock. Each clock tick, the counter register (CNT) increments by 1. When CNT reaches a programmed limit (TOP), something happens — an interrupt fires, an output pin toggles, CNT resets — depending on the mode. The CPU is not involved; the peripheral manages everything in hardware.

That is the mental model to keep in view: a timer is just a counter plus rules about what to do when the counter reaches certain values.

The key equation:

$$f_{\text{interrupt}} = f_{\text{cpu}} / (\text{prescaler} \times (\text{TOP} + 1))$$

Rearranged for TOP given a desired frequency:

$$\text{TOP} = (f_{\text{cpu}} / (\text{prescaler} \times f_{\text{desired}})) - 1$$

Most timer configuration problems reduce to this equation plus careful attention to which register plays the role of TOP in the chosen mode.

11.1.2 Prescaler

The prescaler divides the CPU clock before it reaches the counter. Without it, a 16-bit counter at 20 MHz would overflow in 3.2 ms — useful for microsecond timing but not for a 1 Hz blink. Available prescaler values on TCA0:

Prescaler	Divides by	Max period @ 3.333 MHz (16-bit counter)
-----------	------------	---

DIV1	1	19.7 ms	(65535 ticks)
DIV2	2	39.3 ms	
DIV4	4	78.6 ms	
DIV8	8	157.3 ms	
DIV16	16	314.6 ms	
DIV64	64	1.26 s	
DIV256	256	5.03 s	
DIV1024	1024	20.1 s	

For a 1 Hz blink, DIV1024 gives 20 seconds of headroom with a 16-bit counter — plenty of room.

So the prescaler is what lets one timer peripheral cover both fast and slow timescales.

11.2 TCA0 on ATtiny3217

ATtiny3217 has one Timer/Counter Type A (TCA0) and two Timer/Counter Type B (TCB0, TCB1). TCA0 is the most versatile:

- 16-bit counter (CNT)
- Three compare channels (CMP0, CMP1, CMP2) each with an output pin option
- Three waveform generation modes
- Single mode (16-bit) or split mode (two independent 8-bit timers)

This chapter uses **single mode** (the default).

That is the easiest mode to reason about because the whole 16-bit counter stays together as one timer.

11.2.1 TCA0 Register Map

Register	Address	Width	Purpose
TCA0_SINGLE_CTRLA	0x0A00	8	Prescaler, enable
TCA0_SINGLE_CTRLB	0x0A01	8	Waveform generation mode, compare enable
TCA0_SINGLE_CTRLC	0x0A02	8	Compare output value override
TCA0_SINGLE_CTRLD	0x0A03	8	Split mode enable (leave 0 for single)
TCA0_SINGLE_CTRLCLR	0x0A04	8	Command: clear CNT, stop, restart
TCA0_SINGLE_CTRLSET	0x0A05	8	Command: set (write-to-set)
TCA0_SINGLE_INTCTRL	0x0A06	8	Interrupt enable bits
TCA0_SINGLE_INTFLAGS	0x0A07	8	Interrupt flags (W1C)
TCA0_SINGLE_DBGCTRL	0x0A08	8	Debug run override
TCA0_SINGLE_EVCTRL	0x0A09	16	Event control
TCA0_SINGLE_CNT	0x0A20	16	Current counter value (L at 0x20, H at 0x21)
TCA0_SINGLE_PER	0x0A26	16	Period (TOP in normal/PWM mode)
TCA0_SINGLE_CMP0	0x0A28	16	Compare 0 (also TOP in FRQ mode)
TCA0_SINGLE_CMP1	0x0A2A	16	Compare 1
TCA0_SINGLE_CMP2	0x0A2C	16	Compare 2
TCA0_SINGLE_PERBUF	0x0A36	16	Period buffer (double-buffered update)
TCA0_SINGLE_CMP0BUF	0x0A38	16	CMP0 buffer
TCA0_SINGLE_CMP1BUF	0x0A3A	16	CMP1 buffer
TCA0_SINGLE_CMP2BUF	0x0A3C	16	CMP2 buffer

All TCA0 registers are outside the I/O space (addresses $\geq 0x0060$), so they require LDS/STS — not IN/OUT. The 16-bit registers use two consecutive bytes: the L suffix register at the lower address must be read or written first (low byte first for 16-bit accesses on AVR).

That means timer code is one of the places where the full memory-mapped peripheral model becomes very visible.

11.2.2 CTRLA — Control A

Bit:	7	6:4	3:1	0
	—	—	—	—
	RUNSTDBY	—	CLKSEL	ENABLE

CLKSEL[3:1]:

TCA_SINGLE_CLKSEL_DIV1_gc	= (0x00 << 1)	prescaler /1
TCA_SINGLE_CLKSEL_DIV2_gc	= (0x01 << 1)	prescaler /2
TCA_SINGLE_CLKSEL_DIV4_gc	= (0x02 << 1)	prescaler /4
TCA_SINGLE_CLKSEL_DIV8_gc	= (0x03 << 1)	prescaler /8
TCA_SINGLE_CLKSEL_DIV16_gc	= (0x04 << 1)	prescaler /16
TCA_SINGLE_CLKSEL_DIV64_gc	= (0x05 << 1)	prescaler /64
TCA_SINGLE_CLKSEL_DIV256_gc	= (0x06 << 1)	prescaler /256
TCA_SINGLE_CLKSEL_DIV1024_gc	= (0x07 << 1)	prescaler /1024

ENABLE (bit 0):

1 = timer running
0 = timer stopped (CNT holds its value)

Write ENABLE last after configuring the mode and TOP — avoids glitches where the timer runs briefly with incorrect settings.

This is a recurring embedded pattern: configure first, enable last.

11.2.3 CTRLB — Control B

Bit:	7	6	5	4	3:0
	CMP2EN	CMP1EN	CMP0EN	ALUPD	WGMODE[3:0]

CMP2EN/CMP1EN/CMP0EN: 1 = compare output drives the associated pin

ALUPD: auto-lock update of double-buffered registers on OVF/TOP

WGMODE[3:0]:

TCA_SINGLE_WGMODE_NORMAL_gc	= 0x00	Normal – count to PER, interrupt on OVF
TCA_SINGLE_WGMODE_FRQ_gc	= 0x02	Frequency – count to CMP0, toggle pin, interrupt on CMP0
TCA_SINGLE_WGMODE_SINGLESLOPE_gc	= 0x03	PWM – count to PER, duty via CMPn
TCA_SINGLE_WGMODE_DSTOP_gc	= 0x05	Dual-slope PWM (top)
TCA_SINGLE_WGMODE_DSBOTH_gc	= 0x06	Dual-slope (top + bottom)
TCA_SINGLE_WGMODE_DSBOTTOM_gc	= 0x07	Dual-slope (bottom)


11.2.4 INTCTRL and INTFLAGS

Bit:	4	3	2	0
	CMP2	CMP1	CMP0	OVF

INTCTRL: writing 1 to a bit enables that interrupt.

INTFLAGS: bit set when interrupt fires. Clear by writing 1 (W1C).

TCA_OVF_bm = 0x01 Overflow / Underflow
TCA_CMP0_bm = 0x10 Compare 0 match



Duty cycle = $\text{CMP0} / \text{PER}$

CNT counts to PER. WO0 is high from 0 to CMP0, then low from CMP0 to PER. Duty cycle is $\text{CMP0} / \text{PER}$. OVF interrupt fires when CNT wraps. This is the standard PWM mode for motor control, LED dimming, etc.

PWM mode is where the timer stops being just a scheduler and starts being a waveform generator.

PWM frequency:

$f_{\text{pwm}} = f_{\text{cpu}} / (\text{prescaler} \times (\text{PER} + 1))$
 Duty = $\text{CMP0} / \text{PER}$ (0.0 = 0%, 1.0 = 100%)

11.4 Computing Timer Values

For a 1 Hz interrupt from a 3.333 MHz CPU using FRQ mode and prescaler 1024:

$\text{TOP} = (f_{\text{cpu}} / (\text{prescaler} \times f_{\text{desired}})) - 1$
 $= (3333333 / (1024 \times 1)) - 1$
 $= 3255.2 - 1$
 $= 3254.2 \rightarrow \text{round to } 3254$

Actual frequency check:

$f = 3333333 / (1024 \times (3254 + 1))$
 $= 3333333 / 3332096$
 $= 1.000371 \text{ Hz (error } < 0.04\%)$

The 16-bit value 3254 in hex is 0x0CB6:

$3254 = 0x0CB6$
 $\text{lo8}(3254) = 0xB6$
 $\text{hi8}(3254) = 0x0C$

Low byte must be written first. The ATtiny3217 uses a 16-bit write latch: writing CMP0L latches the value, and writing CMP0H completes the transfer to the 16-bit register atomically. Reversing the order produces a glitch where the intermediate (wrong) value is briefly seen by the timer.

This is not stylistic ceremony. The write order is part of the hardware interface contract.

```

ldi r16, lo8(3254)
sts TCA0_SINGLE_CMP0L, r16    /* writes 0xB6, latches */
ldi r16, hi8(3254)
sts TCA0_SINGLE_CMP0H, r16    /* writes 0x0C, transfers 0x0CB6 atomically */
  
```

11.5 TCA0 Interrupt Vectors

From the ATtiny3217 vector table (Chapter 10):

Byte addr	Vector	Peripheral	Trigger
0x0010	8	TCA0_OVF	Overflow (normal/PWM modes)
0x0012	9	TCA0_HUNF	Split mode high-byte underflow
0x0014	10	TCA0_CMP0	Compare 0 match
0x0016	11	TCA0_CMP1	Compare 1 match
0x0018	12	TCA0_CMP2	Compare 2 match

In FRQ mode, the **TCA0_CMP0** interrupt (vector 10, address 0x0014) fires on every compare match. TCA0_OVF is not generated in FRQ mode.

So the interrupt source depends on the chosen mode, not just on the peripheral name.

To place `tca0_cmp0_isr` at 0x0014, the vector table needs `rjmp tca0_cmp0_isr` as vector entry 10 (counting RESET as entry 0):

```
__vectors:
    rjmp main          /* 0x0000: RESET */
    rjmp isr_default   /* 0x0002: CRCSCAN */
    rjmp isr_default   /* 0x0004: BOD */
    rjmp isr_default   /* 0x0006: PORTA */
    rjmp isr_default   /* 0x0008: PORTB */
    rjmp isr_default   /* 0x000A: PORTC */
    rjmp isr_default   /* 0x000C: RTC_CNT */
    rjmp isr_default   /* 0x000E: RTC_PIT */
    rjmp isr_default   /* 0x0010: TCA0_OVF */
    rjmp isr_default   /* 0x0012: TCA0_HUNF */
    rjmp tca0_cmp0_isr /* 0x0014: TCA0_CMP0 ← */
```

11.6 The ISR

The TCA0 CMP0 ISR follows the same prologue/epilogue discipline from Chapter 10. It adds one step: **clear the INTFLAGS bit before returning**, or the interrupt fires again immediately.

```
tca0_cmp0_isr:
    push r16
    in r16, SREG
    push r16

    ldi r16, (1 << LED_BIT)
    sts PORTA_OUTTGL, r16 /* toggle LED0 on PA3 (active low) */

    ldi r16, TCA_CMP0_bm /* = 0x10 */
    sts TCA0_SINGLE_INTFLAGS, r16 /* clear CMP0 flag (W1C) */

    pop r16
    out SREG, r16
    pop r16
    reti
```

PORTA_OUTTGL is a toggle register — writing a 1 to bit *n* toggles PORTAn without affecting other bits. It is at address 0x0407 (outside I/O space, requires STS). Writing $(1 \ll \text{LED_BIT}) = 0x08$ toggles only PA3.

TCA0_SINGLE_INTFLAGS uses W1C semantics: writing TCA_CMP0_bm (bit 4 = 1) clears only the CMP0 flag. If you needed to clear all flags simultaneously: `ldi r16, 0xFF; sts TCA0_SINGLE_INTFLAGS, r16.`

As with every interrupt source on AVR, forgetting to clear the flag turns one interrupt into an endless interrupt loop.

11.7 Main: Timer Initialization Sequence

Order matters. The safe sequence is:

1. Set CTRLB (waveform mode) – timer is not running yet
2. Write CMP0L then CMP0H – latch and load TOP value
3. Set INTCTRL – enable the interrupt we want
4. Set CTRLA – prescaler + enable bit (starts the timer)
5. SEI – enable global interrupts

Starting the timer before configuring CMP0 would cause the timer to run briefly with CNT=0 and CMP0=0, firing immediately. Starting before INTCTRL is set means the first match is missed silently (no interrupt enabled yet).

Initialization order matters because hardware begins reacting as soon as the enable bit is set.

main:

```

    ldi    r16, lo8(RAMEND)
    out    SPL, r16
    ldi    r16, hi8(RAMEND)
    out    SPH, r16

    sbi    VPORTA_DIR, LED_BIT      /* PA3 = output (LED0 on Curiosity Nano) */
    sbi    VPORTA_OUT, LED_BIT      /* LED off initially (active low: HIGH = off) */

    /* 1. Waveform mode */
    ldi    r16, TCA_SINGLE_WGMODE_FRQ_gc
    sts    TCA0_SINGLE_CTRLB, r16

    /* 2. TOP value: CMP0 = 3254 */
    ldi    r16, lo8(3254)
    sts    TCA0_SINGLE_CMP0L, r16
    ldi    r16, hi8(3254)
    sts    TCA0_SINGLE_CMP0H, r16

    /* 3. Enable CMP0 interrupt */
    ldi    r16, TCA_CMP0_bm
    sts    TCA0_SINGLE_INTCTRL, r16

    /* 4. Start timer: prescaler 1024, enable */
    ldi    r16, (TCA_SINGLE_CLKSEL_DIV1024_gc | TCA_SINGLE_ENABLE_bm)
    sts    TCA0_SINGLE_CTRLA, r16

    /* 5. Global interrupt enable */
    sei

.Lidle:
    rjmp   .Lidle

```

11.8 Worked Example: 1 Hz Blink

Full source: `src/timer_ctc.S`

The complete program combines the vector table, the CMP0 ISR, and main into one file. No shared state is needed between main and the ISR — main initializes hardware and parks in the idle loop; the ISR does all the work.

That makes this example a good demonstration of timer-driven design: main sets the policy, the timer peripheral generates the schedule, and the ISR performs the periodic action.

11.8.1 Register Usage

r16: scratch — used in prologue/epilogue and ISR body
(saved/restored each invocation)

No other registers needed. The simplest possible ISR.

LED0 on Curiosity Nano is active low (PA3 high = LED off, PA3 low = LED on). PORTA_OUTTGL toggles PA3 on every CMP0 match — alternating on/off at 1 Hz.

11.8.2 Timing Analysis

Event	Cycles
TCA0 CMP0 match (hardware)	0
Interrupt request to CPU	1-2 (sync)
CPU finishes current RJMP	2
Push PC, clear I, fetch vector	3-5
RJMP tca0_cmp0_isr	1
Total to ISR first instruction	~7-9
ISR prologue (PUSH + IN + PUSH)	6 cycles (2+1+2+1)
Toggle (LDI + STS)	3 cycles (1+2)
Clear flag (LDI + STS)	3 cycles (1+2)
ISR epilogue (POP + OUT + POP + RETI)	9 cycles
Total ISR duration	~21 cycles

At 3.333 MHz: 21 cycles $\times$ 300 ns $\approx$ 6.3 μ s ISR runtime per blink

The timer period is 3333333 / (1024 $\times$ 3255) $\approx$ 1.0 s. The ISR runtime ($\sim$ 6.3 μ s) is 0.00063% of the period — nearly zero CPU overhead for the blink.

This is exactly why timers are so valuable. The CPU only pays for the tiny ISR runtime, not for the whole one-second wait.

11.8.3 Build Commands

```
avr-gcc -mmcu=attiny3217 -c -o timer_ctc.o src/timer_ctc.S
avr-gcc -mmcu=attiny3217 -nostartfiles -o timer_ctc.elf timer_ctc.o
avr-objcopy -O ihex timer_ctc.elf timer_ctc.hex
```

```
# Inspect disassembly
avr-objdump -d timer_ctc.elf
```

11.8.4 Verifying with avr-objdump

Expected output (excerpts):

```
00000000 <__vectors>:
  0: XX XX rjmp .+N    ; → main
  4: ...
 14: XX XX rjmp .+N    ; → tca0_cmp0_isr  (TCA0_CMP0 vector at 0x0014)

<tca0_cmp0_isr>:
  push r16
  in  r16, 0x3f      ; SREG = 0x3f in I/O space
  push r16
  ldi r16, 0x08      ; (1 << 3) = PA3
  sts 0x0407, r16    ; PORTA_OUTTGL
  ldi r16, 0x10      ; TCA_CMP0_bm
  sts 0x0a07, r16    ; TCA0_SINGLE_INTFLAGS
  pop r16
  out 0x3f, r16      ; restore SREG
  pop r16
  reti
```

Check that `tca0_cmp0_isr` is at the address that vector `0x0014` points to.

11.9 Extending: Normal Mode with OVF Interrupt

For cases where you want a fixed period and don't need compare-match output, use Normal mode with the OVF interrupt. Change two registers:

```
/* Normal mode – OVF fires when CNT reaches PER */
ldi r16, TCA_SINGLE_WGMODE_NORMAL_gc /* = 0x00 */
sts TCA0_SINGLE_CTRLB, r16

ldi r16, lo8(3254)
sts TCA0_SINGLE_PERL, r16 /* PER (not CMP0) */
ldi r16, hi8(3254)
sts TCA0_SINGLE_PERH, r16

ldi r16, TCA_OVF_bm /* = 0x01: overflow interrupt */
sts TCA0_SINGLE_INTCTRL, r16
```

The ISR changes slightly — clear OVF instead of CMP0:

```
tca0_ovf_isr:
  push r16
  in  r16, SREG
  push r16

  ldi r16, (1 << LED_BIT)
  sts PORTA_OUTTGL, r16

  ldi r16, TCA_OVF_bm /* = 0x01 */
  sts TCA0_SINGLE_INTFLAGS, r16 /* clear OVF flag */

  pop r16
```



```

out    SREG, r16
pop    r16
reti

```

And the vector changes to entry 8 (0x0010 — TCA0_OVF):

```

rjmp   tca0_ovf_isr    /* 0x0010: TCA0_OVF */
rjmp   isr_default     /* 0x0012: TCA0_HUNF */
rjmp   isr_default     /* 0x0014: TCA0_CMP0 */

```

Both modes produce identical 1 Hz behavior. FRQ mode is marginally simpler (no PER register to configure separately).

That kind of tradeoff is common in timer work: several modes can solve the same problem, but one usually fits the mental model more directly.

11.10 Extending: Single-Slope PWM

To dim an LED with 50% duty cycle at ~16.3 kHz using PA0 (WO0 pin):

```

/* PER = 204 → f_pwm = 333333 / (1 × 205) ≈ 16.26 kHz */
/* CMP0 = 102 → duty = 102/204 ≈ 50% */

ldi    r16, TCA_SINGLE_WGMODE_SINGLESLOPE_gc /* CTRLB: SS-PWM + CMP0EN */
ori    r16, TCA_SINGLE_CMP0EN_bm
sts    TCA0_SINGLE_CTRLB, r16

ldi    r16, lo8(204)
sts    TCA0_SINGLE_PERL, r16
ldi    r16, hi8(204)
sts    TCA0_SINGLE_PERH, r16

ldi    r16, lo8(102)
sts    TCA0_SINGLE_CMP0L, r16
ldi    r16, hi8(102)
sts    TCA0_SINGLE_CMP0H, r16

/* Start timer with prescaler 1 (no division) */
ldi    r16, (TCA_SINGLE_CLKSEL_DIV1_gc | TCA_SINGLE_ENABLE_bm)
sts    TCA0_SINGLE_CTRLA, r16

```

No interrupt needed for PWM — the hardware toggles WO0 automatically. The PA0 pin must be set as output (SBI VPORTA_DIR, 0) before enabling the timer output, or WO0 is disconnected from the pin.

This is an important milestone: once PWM is configured, the peripheral can keep producing the waveform even if the CPU goes off and does something unrelated.

To change duty cycle at runtime, write a new value to CMP0BUF (the buffer). The ALUPD bit in CTRLB causes the buffer to transfer to CMP0 on the next OVF, preventing glitches mid-cycle.

11.11 Summary

TCA0 modes:

NORMAL (0x00): CNT → PER → reset; OVF interrupt.
 FRQ (0x02): CNT → CMP0 → reset; CMP0 interrupt; WO0 toggles (if enabled).
 SS-PWM (0x03): CNT → PER → reset; WO0 high from 0→CMP0, low CMP0→PER; OVF int.

Frequency formula:

$$f = f_{\text{cpu}} / (\text{prescaler} \times (\text{TOP} + 1))$$

$$\text{TOP} = (f_{\text{cpu}} / (\text{prescaler} \times f)) - 1$$

16-bit register write order: LOW byte first, then HIGH byte (latch).

Initialization order:

1. CTRLB (mode) 2. CMP0/PER (TOP) 3. INTCTRL 4. CTRLA (start) 5. SEI

Interrupt flag clear: write TCA_CMP0_bm (or TCA_OVF_bm) to TCA0_SINGLE_INTFLAGS before RETI – or interrupt fires again immediately.

ISR discipline: identical to Chapter 10.

PUSH all used regs → IN/PUSH SREG → body → POP/OUT SREG → POP regs → RETI.

ATtiny3217 TCA0 vectors:

0x0010 TCA0_OVF – overflow
 0x0014 TCA0_CMP0 – compare 0 match (FRQ mode trigger)
 0x0016 TCA0_CMP1
 0x0018 TCA0_CMP2

11.12 Exercises

1. The example uses prescaler 1024 and CMP0 = 3254. Calculate the CMP0 value needed for a 2 Hz blink with the same prescaler. What is the exact frequency error as a percentage?
2. Modify the example to blink the LED three times fast (150 ms on, 150 ms off) then pause for 700 ms. First reconfigure the timer for a 150 ms tick, then use a counter variable in SRAM to step through the pattern.
3. Calculate PER and CMP0 values for a 1 kHz PWM signal on WO0 with 25% duty cycle, using prescaler DIV1 at 3.333 MHz. Write the complete initialization sequence in assembly.
4. The ISR clears TCA0_SINGLE_INTFLAGS with an STS instruction. STS is a 4-byte, 2-cycle instruction. TCA0_SINGLE_INTFLAGS is at 0x0A07 — not in the VPORT I/O range (0x00–0x1F), so SBI cannot be used here. Confirm this by checking whether 0x0A07 mapped to any VPORT address. What is the smallest instruction count you can use to clear the flag?
5. Add a second compare channel (CMP1) to the FRQ-mode example. CMP1 should fire at 2 Hz (while CMP0 still fires at 1 Hz). What constraint does FRQ mode place on CMP1's relationship to CMP0? (Hint: re-read the FRQ mode description — CNT resets at CMP0, not PER.)
6. Using `avr-objdump -d timer_ctc.elf`, verify that the `sts TCA0_SINGLE_CMP0L` instruction is encoded as a 4-byte instruction. What are the two 16-bit words of its encoding? (Refer to the AVR instruction set manual for the STS 32-bit format.)

Next: Chapter 12 — USART (Serial Communication)

Chapter 12

USART (Serial Communication)


USART (Universal Synchronous/Asynchronous Receiver/Transmitter) is the simplest way to communicate between a microcontroller and a PC. One wire transmits (TX), one receives (RX), and data flows as a timed sequence of bits. No clock line is needed — both sides agree on a baud rate (bits per second) in advance. This chapter covers the ATtiny3217's USART0 peripheral: register setup, polling transmit, polling receive, ISR-driven receive, and a complete echo terminal example.

This is often the first peripheral that makes a microcontroller feel “interactive” because it gives you a direct text path to a terminal.


12.1 UART Framing

12.1.1 The Bit Stream

USART uses NRZ (Non-Return-to-Zero) encoding: the line is high (idle) when no data is being sent. Each byte is framed with:



Timeline at 9600 baud (1 bit = 104.2 μ s):



One character transmission = 1 start + 8 data + 1 stop = **10 bit periods**. At 9600 baud: $10 \times 104.2 \mu\text{s} = 1.042 \text{ ms}$ per byte $\rightarrow$ maximum ~ 960 bytes/second.

That simple 10-bit framing cost is worth remembering. Even before software overhead, UART throughput is lower than the raw baud number suggests.

12.1.2 8N1 Format

8N1 means:

- 8 data bits (LSB first)
- N — no parity bit
- 1 stop bit

This is the most common UART configuration and the ATtiny3217 USART0 default after reset.

So for ordinary 8-bit serial text, the default framing is already the framing you want.

12.2 USART0 Registers on ATtiny3217

Register	Address	Purpose
USART0_RXDATA	0x0800	Receive data low byte (read = receive byte)
USART0_RXDATAH	0x0801	Receive data high byte (error flags + 9th bit)
USART0_TXDATA	0x0802	Transmit data (write = queue byte to send)
USART0_STATUS	0x0804	Status flags (RXCIF, TXCIF, DREIF, etc.)
USART0_CTRLA	0x0805	Interrupt enables, RS-485 mode
USART0_CTRLB	0x0806	TX/RX enable, start-of-frame detect, CLK2X
USART0_CTRLC	0x0807	Frame format: baud mode, char size, parity, stop bits
USART0_BAUDL	0x0808	Baud rate register low byte
USART0_BAUDH	0x0809	Baud rate register high byte
USART0_DBGCTRL	0x080A	Debug run override
USART0_EVCTRL	0x080B	Event input enable
USART0_TXPLCTRL	0x080C	TX PastLimiter control
USART0_RXPLCTRL	0x080D	RX PastLimiter control

All outside I/O space: require LDS/STS.

That means USART code on ATtiny3217 naturally uses the “full peripheral register” access style rather than the short IN/OUT style.

12.2.1 STATUS — Status Register (0x0804)

Bit:	7	6	5	4	3	2	1	0
	RXCIF	TXCIF	DREIF	FERR	BUFOVF	PERR	—	WFB

- RXCIF — RX Complete Interrupt Flag: 1 when unread byte in RXDATA.
Cleared automatically by reading RXDATA.
- TXCIF — TX Complete Interrupt Flag: 1 when last byte shifted out fully.
Cleared by writing 1 (W1C) or enabling TXCIE.
- DREIF — Data Register Empty Flag: 1 when TXDATA is ready for new byte.
Cleared automatically when TXDATA is written.
On reset, DREIF = 1 (TX is ready immediately).
- FERR — Frame Error: start/stop bit violation in received byte.
- BUFOVF — Buffer Overflow: byte arrived before RXDATA was read.
- PERR — Parity Error.

For polling TX: wait until DREIF = 1, then write TXDATA. For polling RX: wait until RXCIF = 1, then read RXDATA. For ISR-driven RX: enable RXCIE in CTRLA; ISR fires when RXCIF sets.

Those three sentences are the whole USART mental model in miniature: one flag for “ready to transmit,” one flag for “data received,” and one interrupt-enable bit if you want hardware to notify you instead of polling.

12.2.2 CTRLB — Control B (0x0806)

Bit:	7	6	5	4	3	2:1	0
	RXEN	TXEN	SFDEN	ODME	RXMODE	—	CLK2X

RXEN – 1 = receive enabled (USART monitors RXD pin)
 TXEN – 1 = transmit enabled (USART drives TXD pin)
 RXMODE – 0 = normal (16× oversampling), 1 = CLK2X (8×), 2 = LINAUTO, 3 = GENAUTO

Bit constants (from `avr/io.h`):

USART_TXEN_bm = 0x40
 USART_RXEN_bm = 0x80
 USART_RXCIE_bm is in CTRLA (see below)

12.2.3 CTRLA — Control A / Interrupt Enables (0x0805)

Bit:	7	6	5	4	3:2	1	0
	RXCIE	TXCIE	DREIE	RXSIE	LBME	–	RS485

RXCIE – RX Complete Interrupt Enable: fires ISR when RXCIF = 1
 TXCIE – TX Complete Interrupt Enable
 DREIE – Data Register Empty Interrupt Enable

Bit constants:

USART_RXCIE_bm = 0x80
 USART_TXCIE_bm = 0x40
 USART_DREIE_bm = 0x20

12.2.4 CTRLC — Frame Format (0x0807)

Bit:	7:6	5:4	3	2:1	0
	CMODE	–	SBMODE	PMODE	CHSIZE

CMODE (bits 7:6): communication mode

0x00 = USART_CMODE_ASYNCHRONOUS_gc – async (normal UART)
 0x01 = USART_CMODE_SYNCHRONOUS_gc – sync
 0x02 = USART_CMODE_IRCOM_gc – IrDA
 0x03 = USART_CMODE_MSPI_gc – SPI master

SBMODE (bit 3): stop bits

0 = 1 stop bit
 1 = 2 stop bits

PMODE (bits 2:1): parity

0x00 = none
 0x02 = even
 0x03 = odd

CHSIZE:

0x03 = 8-bit (USART_CHSIZE_8BIT_gc) – most common

Default after reset = 0x03 = async, 8-bit, no parity, 1 stop = 8N1.
 No write needed if 8N1 is desired.

12.3 Baud Rate Calculation

12.3.1 The Formula

USART0 on ATtiny3217 uses a 16-bit baud register (BAUDL + BAUDH) with a fractional format. For **asynchronous normal mode** (16× oversampling):

$$\begin{aligned}\text{BAUD_REG} &= (\text{f_cpu} \times 64) / (16 \times \text{baud_rate}) \\ &= (\text{f_cpu} \times 4) / \text{baud_rate}\end{aligned}$$

The $\times 64$ and $\div 16$ come from the internal fractional divider format — the top 12 bits of BAUD_REG are the integer part, the bottom 4 bits are the fractional part (1/16 increments).

You do not need to memorize the internal fractional details to use the peripheral well, but they explain why the formula looks stranger than a simple clock-divider equation.

For 9600 baud at 3.333 MHz:

$$\begin{aligned}\text{BAUD_REG} &= (3333333 \times 64) / (16 \times 9600) \\ &= 213333312 / 153600 \\ &= 1389.0 \quad (\text{exact to 4 significant figures}) \\ &\rightarrow 0x056D\end{aligned}$$

$$\begin{aligned}\text{Actual baud rate} &= (3333333 \times 64) / (16 \times 1389) \\ &= 213333312 / 22224 \\ &= 9600.0 \text{ bps} \quad (< 0.01\% \text{ error} - \text{negligible})\end{aligned}$$

Common baud rates at 3.333 MHz:

Baud	BAUD_REG (decimal)	BAUD_REG (hex)	Error
9600	1389	0x056D	0.00%
19200	694	0x02B6	0.04%
38400	347	0x015B	0.06%
57600	231	0x00E7	0.16%
115200	116	0x0074	0.16%

Writing BAUD (low byte first):

```
ldi  r16, lo8(1389)    /* = 0x6D */
sts  USART0_BAUDL, r16
ldi  r16, hi8(1389)    /* = 0x05 */
sts  USART0_BAUDH, r16
```

As with other 16-bit peripheral values on AVR, the write order is part of the hardware interface, not just style.

12.4 TX/RX Pins on ATtiny3217

USART0 uses **PORTB** by default:

Function	Pin	Direction	Notes
----------	-----	-----------	-------

TXD	PB2	Output	Must set VPORTB_DIR bit 2 = 1
RXD	PB3	Input	Leave as input (default); no pull-up needed

The USART peripheral drives TXD when TXEN=1, but only if the pin is configured as output. If PB2 is left as input, TXD is disconnected from the pin driver and the line floats.

That is an easy detail to miss: enabling the USART transmitter is not the same thing as configuring the port pin direction.

```
sbi   VPORTB_DIR, 2      /* PB2 = output for TXD */
/* PB3 (RXD) = input by default – no configuration needed */
```

Alternative pin mapping (PORTMUX_USARTROUTEA register) allows moving USART0 to other pins, but this chapter uses the default mapping.

12.5 Polling Transmit

The simplest TX approach: spin until DREIF = 1, then write the byte.

Polling transmit is usually acceptable because the program already knows when it wants to send something.

AVR Instruction: LDS Rd, A – load from SRAM/peripheral
 SBRS Rd, b – skip next if bit b of Rd is set

```
/*
 * usart_tx_byte – transmit r16 via polling
 * Clobbers: r17
 */
usart_tx_byte:
    push    r17
.lwait_dreif:
    lds     r17, USART0_STATUS
    sbrc    r17, USART_DREIF_bp    /* bit 5 – skip if DREIF = 1          */
    rjmp     .lwait_dreif          /* DREIF = 0: TX busy, loop        */
    sts     USART0_TXDATA, r16     /* write byte: clears DREIF, starts TX */
    pop     r17
    ret
```

USART_DREIF_bp is the bit position (5) of DREIF in the STATUS register. SBRS r17, 5 skips the RJMP when bit 5 is 1 (DREIF set = TX ready).

So the loop is doing exactly one question over and over: “is the transmit data register empty yet?”

12.5.0.1 Why Not SBIS?

SBIS and SBIC only work with I/O space registers (addresses 0x00–0x1F in I/O space, or 0x20–0x3F in data space). USART0_STATUS is at 0x0804 — far outside I/O space. Use LDS + SBRS for peripheral registers.

12.5.0.2 Cycle Cost of Polling

Each loop iteration: LDS (2 cy) + SBRS (1 cy, not skipping) + RJMP (2 cy) = 5 cycles. At 3.333 MHz, 5 cycles = 1.5 μ s per iteration. At 9600 baud, one bit period is 104 μ s — the polling loop runs ~69 times per bit period. Acceptable for simple programs; for power-sensitive applications, use DREIE interrupt instead.

That is the tradeoff in one paragraph: polling is simple, but it spends CPU time checking a condition that usually is not ready yet.

12.5.1 Transmitting a String from Flash

```
/*
 * usart_tx_string – send bytes from flash
 * ZL:ZH = pointer to first byte, r18 = length
```

```

* Clobbers: r16, r17, r18, Z
*/
usart_tx_string:
    tst    r18
    breq   .Ltx_done
.Ltx_loop:
    lpm    r16, Z+                /* load byte from flash, post-increment Z */
    rcall  usart_tx_byte
    dec    r18
    brne   .Ltx_loop
.Ltx_done:
    ret

```

String data in flash:

```

.Lready_str:
    .byte  'R', 'e', 'a', 'd', 'y', '\r', '\n'
.equ  READY_LEN, 7

```

Calling it:

```

ldi    ZL, lo8(.Lready_str)
ldi    ZH, hi8(.Lready_str)
ldi    r18, READY_LEN
rcall  usart_tx_string

```

The `.byte` directive places raw bytes in the `.text` section at the label address. `LPM Z+` loads each byte from flash and advances `Z`. `READY_LEN` is a compile-time constant counting the bytes (`7 = 5 chars + CR + LF`).

This is a nice place where Chapter 5's flash-reading model becomes immediately useful in application code.

12.6 Polling Receive

Wait until `RXCIF = 1`, then read `RXDATAL`. Reading clears the flag.

Polling receive is conceptually symmetric with polling transmit, but it is much more wasteful in practice because the CPU does not know when the next incoming byte will arrive.

```

/*
* usart_rx_byte - receive one byte into r16 (polling)
* Clobbers: r17
*/
usart_rx_byte:
    push   r17
.Lwait_rxcif:
    lds    r17, USART0_STATUS
    sbrc   r17, USART_RXCIF_bp    /* bit 7 - skip if RXCIF = 1 (byte ready) */
    rjmp   .Lwait_rxcif
    lds    r16, USART0_RXDATAL    /* read byte - clears RXCIF */
    pop    r17
    ret

```

`USART_RXCIF_bp = 7` (bit position). `SBRSC r17, 7` skips the loop branch when the RX byte is ready.

12.6.1 Error Checking

Before reading RXDATAH, the error flags in STATUS should be checked:

```
.Lwait_rxcif:
    lds    r17, USART0_STATUS
    sbrs   r17, USART_RXCIF_bp
    rjmp   .Lwait_rxcif
    lds    r17, USART0_RXDATAH    /* read high byte first for error flags */
    andi   r17, (USART_FERR_bm | USART_BUFOVF_bm | USART_PERR_bm)
    brne   .Lrx_error            /* jump to error handler if any flag set */
    lds    r16, USART0_RXDATAH    /* no error – read data */
```

RXDATAH contains the error flags (FERR, BUFOVF, PERR) and optionally the 9th data bit. For 8N1, only the error bits matter. RXDATAH must be read before RXDATAH — reading RXDATAH consumes the byte and discards the error flags.

That read order is another hardware contract: get the status context first, then consume the byte.

For the simple echo example, error checking is omitted for clarity.

12.7 ISR-Driven Receive

Polling TX is fine for transmit — the CPU controls when it sends. But polling RX is wasteful: the CPU spins doing nothing until a byte arrives, unable to perform other work. The RXCIE interrupt solves this: the ISR fires only when a byte is ready, and main code runs uninterrupted between bytes.

This is one of the clearest examples of when interrupts are worth the added complexity.

12.7.1 USART0_RXC Vector

From the ATtiny3217 vector table (Chapter 10):

Byte	addr	Vector	Peripheral	Trigger
0x0036	27	USART0_RXC		RXCIF = 1 (byte received)
0x0038	28	USART0_DRE		DREIF = 1 (TX data register empty)
0x003A	29	USART0_TXC		TXCIF = 1 (TX shift register empty)

The USART0_RXC vector is at byte address 0x0036. In the vector table, it is the 28th entry (counting RESET as entry 0):

```
__vectors:
    rjmp   main          /* 0x0000: entry 0  RESET    */
    ...           /* entries 1-26 */
    rjmp   usart0_rxc_isr /* 0x0036: entry 27 USART0_RXC */
    rjmp   isr_default   /* 0x0038: entry 28 USART0_DRE */
    rjmp   isr_default   /* 0x003A: entry 29 USART0_TXC */
    rjmp   isr_default   /* 0x003C: entry 30 NVMCTRL_EE */
```

12.7.2 RXC ISR

Reading RXDATAH inside the ISR **automatically clears RXCIF**. No explicit flag clear is needed (unlike TCA0 INTFLAGS which requires W1C).

That makes USART receive interrupts a little cleaner than many timer or GPIO interrupt sources.

```

usart0_rxc_isr:
    push    r16
    push    r17
    in      r16, SREG
    push    r16

    lds     r16, USART0_RXDATA    /* read byte – clears RXCIF automatically */
    rcall   usart_tx_byte         /* echo it back (polling TX) */

    pop     r16
    out     SREG, r16
    pop     r17
    pop     r16
    reti

```

Note r17 is pushed/popped because `usart_tx_byte` clobbers it. The ISR calls a subroutine — this is legal but requires that the callee's clobbers are also saved in the prologue.

This is a good example of Chapter 8's ABI rules showing up in a real peripheral handler.

12.7.2.1 Calling Subroutines from ISRs

Calling `RCALL` inside an ISR pushes another 2 bytes onto the stack. The total stack depth for this ISR:

On ISR entry (CPU pushes return addr):	2 bytes
PUSH r16:	1 byte
PUSH r17:	1 byte
PUSH r16 (SREG):	1 byte
RCALL usart_tx_byte (CPU pushes addr):	2 bytes
PUSH r17 (in usart_tx_byte):	1 byte
Total stack depth:	8 bytes

On ATtiny3217, SRAM is 2 KB. Typical use of a few hundred bytes of stack is safe, but deep call trees inside ISRs require careful accounting.

So ISR design is not just about correctness and latency; it is also about bounded stack growth.

12.7.2.2 Blocking TX Inside an ISR

The ISR calls `usart_tx_byte`, which polls `DREIF`. While polling, the ISR is running and global interrupts are disabled (`I=0` after ISR entry). This is safe here because at 9600 baud the TX wait is at most one byte period (~1.042 ms) and no other time-critical ISR is competing. In higher-throughput designs, use a circular buffer: ISR puts byte in buffer, `DREIE` interrupt sends from buffer — the `RXC` ISR returns immediately.

That distinction is important: a simple echo ISR is fine for a teaching example, but blocking inside an ISR does not scale well.

12.8 Initialization Sequence

```

main:
    /* Stack */
    ldi     r16, lo8(RAMEND)
    out     SPL, r16
    ldi     r16, hi8(RAMEND)

```

```

out    SPH, r16

/* TXD pin = output */
sbi    VPORTB_DIR, TXD_BIT          /* PB2 */

/* Baud rate: 9600 at 3.333 MHz → BAUD_REG = 1389 = 0x056D */
ldi    r16, BAUD_LO                 /* 0x6D */
sts    USART0_BAUDL, r16
ldi    r16, BAUD_HI                 /* 0x05 */
sts    USART0_BAUDH, r16

/* CTRLC defaults = 8N1 async – no write needed */

/* Enable TX + RX */
ldi    r16, (USART_TXEN_bm | USART_RXEN_bm)
sts    USART0_CTRLB, r16            /* TXEN + RXEN */
/* RXCIE lives in CTRLA – separate write */
ldi    r16, USART_RXCIE_bm
sts    USART0_CTRLA, r16

sei

/* Send startup message */
ldi    ZL, lo8(.Lready_str)
ldi    ZH, hi8(.Lready_str)
ldi    r18, READY_LEN
rcall  usart_tx_string

.Lidle:
rjmp   .Lidle

```

Note on CTRLB vs CTRLA: TXEN and RXEN are in CTRLB (0x0806). RXCIE is in CTRLA (0x0805). They are separate registers — two STS writes.

This is exactly the sort of small register-detail bug that can waste a lot of debug time if you assume all enable bits live together.

12.9 Worked Example: Echo Terminal

Full source: `src/usart_echo.S`

The program transmits `Ready\r\n` on startup, then echoes every received byte back to the sender. No main-loop processing — all receive logic is in the ISR.

That makes the data path very easy to reason about: receive a byte, trigger the ISR, send the same byte back.

12.9.1 Data Flow

```

PC terminal → USB-UART adapter → RXD (PB3) → USART0_RXDATA
                                     RXCIF = 1 → USART0_RXC ISR fires
                                               reads RXDATA (clears RXCIF)
                                               calls usart_tx_byte(r16)
                                               polls DREIF → writes TXDATA
TXD (PB2) → USB-UART adapter → PC terminal (sees same char echoed back)

```

12.9.2 Build and Test

```
# Build
avr-gcc -mmcu=attiny3217 -c -o usart_echo.o src/usart_echo.S
avr-gcc -mmcu=attiny3217 -nostartfiles -o usart_echo.elf usart_echo.o
avr-objcopy -O ihex usart_echo.elf usart_echo.hex

# Flash
avrdude -c serialupdi -p attiny3217 -P /dev/ttyUSB0 -U flash:w:usart_echo.hex

# Connect serial terminal at 9600 8N1
# Linux/macOS:
screen /dev/ttyUSB0 9600
# or:
minicom -b 9600 -D /dev/ttyUSB0
```

12.9.3 Disassembly Check

```
avr-objdump -d usart_echo.elf
```

Key things to verify:

- Vector at 0x0036 (rjmp usart0_rxc_isr) — counts 27 entries from RESET.
- sts 0x0808, r16 for BAUDL, sts 0x0809, r16 for BAUDH.
- sts 0x0806, r16 for CTRLB, sts 0x0805, r16 for CTRLA.
- lds r16, 0x0800 for RXDATAL inside the ISR.
- lds r17, 0x0804 for STATUS in the TX polling loop.

12.10 Circular Buffer for High-Throughput RX

The simple ISR copies one byte directly. At 115200 baud with rapid bursts, the CPU may not finish the TX echo before the next byte arrives — BUFOVF sets and bytes are lost.

A circular buffer decouples RX from TX:

SRAM layout:

```
rx_buf:    .space 16      ; 16-byte circular buffer
rx_head:    .space 1      ; write index (ISR writes here)
rx_tail:    .space 1      ; read index (main reads here)
```

The ISR writes to rx_buf[rx_head] and increments rx_head (modulo 16). Main code checks rx_head != rx_tail to know a byte is ready, reads from rx_buf[rx_tail], and increments rx_tail.

This is the standard way to decouple “data arrived now” from “main code is ready to process it now.”

```
/* ISR stores byte in circular buffer */
usart0_rxc_isr:
    push    r16
    push    r17
    push    ZL
    push    ZH
    in      r16, SREG
    push    r16

    lds     r16, USART0_RXDATAL    /* read byte, clear RXCIF */
```

```

lds    r17, rx_head          /* load write index          */
ldi    ZL, lo8(rx_buf)
ldi    ZH, hi8(rx_buf)
add    ZL, r17               /* Z = &rx_buf[rx_head]    */
adc    ZH, __zero_reg__
st     Z, r16                /* store byte              */

inc    r17
andi   r17, 0x0F             /* modulo 16 (power-of-2 wrap) */
sts    rx_head, r17

pop    r16
out    SREG, r16
pop    ZH
pop    ZL
pop    r17
pop    r16
reti

```

Full circular buffer implementation is left as Exercise 5.

12.11 Summary

USART0 pins (default): TXD = PB2 (output), RXD = PB3 (input)

Baud rate register:

$\text{BAUD_REG} = (\text{f_cpu} \times 64) / (16 \times \text{baud})$ [async normal mode]

At 3.333 MHz, 9600 bps $\rightarrow$ $\text{BAUD_REG} = 1389 = 0x056D$

Write: BAUDL first (low byte), then BAUDH.

Registers to configure:

USART0_BAUDL / BAUDH – baud rate (write low byte first)
 USART0_CTRLA – frame format (default = 8N1, no write needed)
 USART0_CTRLB – TXEN_bm (0x40) | RXEN_bm (0x80)
 USART0_CTRLA – RXCIE_bm (0x80) for ISR-driven RX

STATUS flags:

DREIF (bit 5) – TX data register empty: 1 = ready for new byte

RXCIF (bit 7) – RX complete: 1 = byte ready in RXDATA1

Both checked with: LDS Rd, USART0_STATUS then SBRS/SBRC

Polling TX:

wait DREIF = 1 $\rightarrow$ STS USART0_TXDATA1, Rr $\rightarrow$ DREIF auto-clears

Polling RX:

wait RXCIF = 1 $\rightarrow$ LDS Rd, USART0_RXDATA1 $\rightarrow$ RXCIF auto-clears

ISR-driven RX (USART0_RXC, vector 27, address 0x0036):

Enable: USART_RXCIE_bm in USART0_CTRLA

ISR reads RXDATA1 $\rightarrow$ clears RXCIF automatically

No manual flag clear needed (unlike TCA0 INTFLAGS)

ISR rules: same as always.

Save r16+SREG minimum. Save any register clobbered by callees.

Stack depth = 2 (ret addr) + saved regs + callee frames.

12.12 Exercises

1. Calculate BAUD_REG for 115200 baud at 3.333 MHz. What is the percentage error? If the error exceeds 2%, the link is likely to be unreliable — is 115200 feasible at this clock speed?
 2. The polling TX function checks DREIF each iteration. At 9600 baud and 3.333 MHz, how many poll iterations occur while waiting for a byte to finish transmitting? (Use the LDS+SBRS+RJMP cycle count.)
 3. The ISR calls `usart_tx_byte`, which polls DREIF. During this polling, could a second USART0_RXC interrupt fire? Explain why or why not. What would happen to the second received byte?
 4. Modify `usart_echo.S` to echo each byte with its hex representation. Receiving A (0x41) should transmit the string 41\r\n. Write the `byte_to_hex` subroutine that converts r16 (0x00–0xFF) into two ASCII hex digits in r17 (high nibble) and r16 (low nibble).
 5. Implement the 16-byte circular buffer described in Section 12.10. Write:
 - the SRAM variable declarations (`.bss` section)
 - the full ISR using the buffer
 - a `main_loop` that calls a `process_byte` subroutine for each byte in the buffer (use `rx_head != rx_tail` as the ready check)
 6. Using `avr-objdump -d usart_echo.elf`, locate the LDS r17, USART0_STATUS instruction in `usart_tx_byte`. What is its 32-bit encoding? What address (0x0804) appears in which bytes of the instruction?
-

Next: Chapter 13 — SPI & TWI (I2C) Overview

Chapter 13

SPI & TWI (I2C) Overview

Serial communication protocols let microcontrollers talk to sensors, memory chips, displays, and other devices using only a handful of wires. This chapter covers the two most common synchronous serial buses used in embedded systems: SPI (Serial Peripheral Interface) and TWI (Two-Wire Interface, Microchip's name for I²C). Both are present on the ATtiny3217 as hardware peripherals; both are accessed through memory-mapped registers using LDS/STS, exactly as with USART0 in Chapter 12.

This chapter is less about talking to a terminal and more about talking to other chips.

13.1 SPI Fundamentals

SPI is a full-duplex, synchronous serial bus. Four wires connect one master to one or more slave devices:

Master (ATtiny3217)		Slave (25LC010A EEPROM)
MOSI (PA1)	→	SI (Master Out Slave In)
MISO (PA2)	←	SO (Master In Slave Out)
SCK (PA3)	→	SCK (clock, generated by master)
CS (PA4)	→	C $\bar{S}$ (chip select, active-low)

Full duplex: every clock pulse simultaneously shifts one bit out on MOSI and samples one bit in on MISO. An 8-bit transfer is 8 clock pulses. If the slave has nothing meaningful to return (a pure command), the byte received by the master is discarded (or a dummy byte 0xFF is sent to clock the slave's response out).

That full-duplex behavior is one of the first things that makes SPI feel different from UART. You are always sending and receiving at the same time, even if one direction is only carrying dummy data.

Chip select: SPI has no built-in addressing. Instead, the master drives CS low to select a particular slave before the transfer begins, then drives CS high after the transfer ends. Each slave needs its own CS line.

So SPI trades protocol simplicity for pin count: the bus itself is simple, but every slave needs a dedicated select signal.

13.1.1 SPI Clock Polarity and Phase (Mode)

The idle state of SCK and the edge on which data is sampled define four modes:


Mode	CPOL	CPHA	SCK idle	Data captured on
0	0	0	low	rising edge ← 25LC010A compatible

1	0	1	low	falling edge	
2	1	0	high	falling edge	
3	1	1	high	rising edge	← also 25LC010A compatible

The 25LC010A EEPROM used in this chapter works in Mode 0 or Mode 3. We use Mode 0 (CPOL=0, CPHA=0) — SCK idles low, data captured on the rising edge.

In practice, “which SPI mode?” is one of the first things you check in a datasheet when a device does not respond correctly.

13.1.2 Timing Diagram (Mode 0)



Data is set up before the rising edge and sampled on the rising edge.

13.2 SPI0 Registers on the ATtiny3217

The ATtiny3217 SPI0 peripheral uses five memory-mapped registers:

Register	Address	Purpose
SPI0_CTRLA	0x08C0	Master/slave, clock, data order, enable
SPI0_CTRLB	0x08C1	Buffer mode, slave-select, SPI mode
SPI0_INTCTRL	0x08C2	Interrupt enables
SPI0_INTFLAGS	0x08C3	Transfer-complete flag (IF), write-collision (WRCOL)
SPI0_DATA	0x08C4	Transmit / receive data register

13.2.1 SPI0_CTRLA — Control A

Bit:	7	6	5	4	3	2	1	0
	—	DORD	MASTER	CLK2X	—	PRESC[1:0]	ENABLE	

Field	Bits	Description
DORD	6	Data order: 0 = MSB first (normal), 1 = LSB first
MASTER	5	1 = master mode, 0 = slave mode
CLK2X	4	Double the clock speed (halves effective prescaler)
PRESC	2:1	Clock prescaler: 00=/4, 01=/16, 10=/64, 11=/128
ENABLE	0	1 = enable SPI peripheral

At 3.333 MHz with PRESC=DIV4 and CLK2X=0: SCK = 3.333 MHz / 4 $\approx$ **833 kHz**. The 25LC010A maximum SCK is 2 MHz (3.3 V) — well within spec.

13.2.2 SPI0_CTRLB — Control B

Bit: 7 6 5 4 3 2 1 0
 BUFEN BUFWR – – – SSD MODE[1:0]

Field	Bits	Description
BUFEN	7	Enable buffer mode (more complex; not used here)
BUFWR	6	Buffer write mode enable (buffer mode only)
SSD	2	Slave Select Disable — set to 1 when driving CS manually
MODE	1:0	SPI mode: 00=Mode0, 01=Mode1, 10=Mode2, 11=Mode3

SSD must be set when driving CS as a GPIO (as we do in this chapter). If SSD=0 and the hardware SS pin (PA4) is driven low by external logic, the ATtiny3217 will switch itself to slave mode — a common source of confusion.

This is one of those details that feels obscure until it breaks the whole bus.

13.2.3 SPI0_INTFLAGS — Interrupt / Status Flags (non-buffered mode)

Bit: 7 6 5 4 3 2 1 0
 IF – – – – – – WRCOL

Flag	Bit	Set when	Cleared by
IF	7	8-bit transfer complete	Reading SPI0_DATA
WRCOL	0	Write collision (SPI0_DATA written mid-transfer)	Reading SPI0_DATA

IF at bit position SPI_IF_bp (= 7). Poll with SBRS.

13.2.4 SPI0_DATA — Transmit / Receive

Reading SPI0_DATA after IF is set returns the received byte and clears IF. Writing SPI0_DATA (when IF is set or before the first transfer) queues the byte to be shifted out. Both transmit and receive share the same register address — write to transmit, read to receive.

13.3 SPI Initialisation (Master Mode)

Before any transfer, configure the three output pins, leave MISO as input, set SSD, set MODE, then enable as master.

That order is deliberate: get the pins and mode right first, then turn the peripheral on.

```
/* Pin directions: MOSI (PA1), SCK (PA3), CS (PA4) → output */
ldi  r16, (1 << 1) | (1 << 3) | (1 << 4)
out  VPORTA_DIR, r16

/* CS high (deselect) */
sbi  VPORTA_OUT, 4
```

```

/* SPI0_CTRLB: SSD=1 (manual CS), MODE=0 (SPI mode 0) */
ldi  r16, SPI_SSD_bm
sts  SPI0_CTRLB, r16

/* SPI0_CTLRA: master, DIV4 prescaler, MSB first, enable */
ldi  r16, (SPI_MASTER_bm | SPI_ENABLE_bm)
sts  SPI0_CTLRA, r16

SPI_MASTER_bm = (1 << SPI_MASTER_bp) = 0x20. SPI_ENABLE_bm = (1 << SPI_ENABLE_bp) = 0x01. Both
are defined in <avr/io.h> for the ATtiny3217.

```

13.4 Transferring a Byte

The hardware handles the 8-bit shift; the CPU's job is to:

1. Write the byte to SPI0_DATA (starts the transfer).
2. Poll SPI0_INTFLAGS until IF = 1 (transfer done).
3. Read SPI0_DATA (clears IF; returns received byte).

Once you see SPI transfers this way, most device protocols reduce to “send a command byte, maybe send an address byte, maybe send or receive data bytes.”

```

/*
 * spi_transfer – exchange one byte over SPI
 * In:  r16 = byte to send
 * Out: r16 = byte received
 * Clobbers: r17
 */
spi_transfer:
    sts  SPI0_DATA, r16          /* write tx byte → starts 8-bit transfer */
.lwait_if:
    lds  r17, SPI0_INTFLAGS
    sbrs r17, SPI_IF_bp          /* spin until IF (bit 7) = 1 */
    rjmp .lwait_if
    lds  r16, SPI0_DATA          /* read received byte – clears IF */
    ret

```

Cycle count at 3.333 MHz with SCK = f/4:

Phase	Duration
8 SPI bits	8 × 4 CPU cycles = 32 cycles (SCK)
Poll loop	typically 1–3 iterations = 6 cycles
STS+LDS	2 + 2 = 4 cycles
Total	≈ 42 CPU cycles (12.6 μs)

13.5 The 25LC010A SPI EEPROM

The 25LC010A is a 1-Kbit (128 × 8-bit) SPI EEPROM. It communicates using single-byte commands followed by optional address and data bytes.

13.5.1 Key Characteristics

Parameter	Value
Capacity	128 bytes (1 Kbit)
Address width	7 bits (0x00 – 0x7F)
Page size	16 bytes (for page write)
SPI modes	Mode 0 and Mode 3
Max SCK (3.3 V)	2 MHz
Write cycle time	5 ms maximum
Supply voltage	1.8 V – 5.5 V

13.5.2 Command Set

Opcode	Hex	Bytes after opcode	Description
READ	0x03	addr, [data...]	Read one or more bytes
WRITE	0x02	addr, data[...]	Write one byte (or page)
WREN	0x06	–	Set Write Enable Latch (WEL)
WRDI	0x04	–	Clear Write Enable Latch
RDSR	0x05	[dummy]	Read Status Register
WRSR	0x01	data	Write Status Register

Write Enable Latch (WEL): the EEPROM resets WEL after every write. You must send WREN (its own CS-framed transaction) before every WRITE command.

That makes accidental writes less likely, but it also means your software has to treat “enable writing” as part of every write sequence, not as a one-time setup step.

13.5.3 Status Register (returned by RDSR)

Bit:	7	6	5	4	3	2	1	0
	–	–	–	–	BP1	BP0	WEL	WIP

Flag	Bit	Meaning
WIP	0	Write In Progress — 1 while write cycle is running
WEL	1	Write Enable Latch — 1 after WREN, 0 after write completes
BP	3:2	Block-protect bits (control write-protect regions)

Poll WIP = 0 before issuing another WRITE command.

In other words, the WRITE command only starts the EEPROM's internal work. The chip still needs time afterwards to commit the byte into nonvolatile storage.

13.5.4 CS Framing Rules

Each command must be bracketed by CS low/high transitions. Raising CS before a complete transfer aborts the current command. The typical patterns:

This is the other half of SPI correctness besides mode. Many slave devices care not just about the bytes, but about the exact CS framing around them.

Write Enable: CS↓ → WREN(0x06) → CS↑
 Write byte: CS↓ → WRITE(0x02) → ADDR → DATA → CS↑
 Read status (poll WIP): CS↓ → RDSR(0x05) → DUMMY → CS↑

read status byte from DUMMY transfer

13.6 Example: Write Byte to SPI EEPROM

The full annotated source is at `ch13_spi_twi/src/spi_eeprom.S`. Here we walk through the key subroutines.

13.6.1 `eeeprom_wren` — Send Write Enable

```
eeeprom_wren:
    rcall cs_assert          /* CS low */
    ldi    r16, EEPROM_WREN /* 0x06 */
    rcall spi_transfer       /* shift out opcode (received byte discarded) */
    rcall cs_deassert        /* CS high — WEL is now set */
    ret
```

13.6.2 `eeeprom_write_byte` — Write One Byte

```
/*
 * r24 = 8-bit address, r22 = data byte
 */
eeeprom_write_byte:
    rcall eeeprom_wren      /* WEL must be set before each write */

    rcall cs_assert
    ldi    r16, EEPROM_WRITE /* 0x02 */
    rcall spi_transfer       /* opcode */
    mov    r16, r24
    rcall spi_transfer       /* address */
    mov    r16, r22
    rcall spi_transfer       /* data */
    rcall cs_deassert        /* CS↑ latches the write; 5 ms cycle starts */
    ret
```

13.6.3 `eeeprom_wait_wip` — Poll WIP Until Write Done

```
eeeprom_wait_wip:
.Lwip_poll:
    rcall cs_assert
    ldi    r16, EEPROM_RDSR /* 0x05 */
    rcall spi_transfer       /* send RDSR opcode */
    ldi    r16, 0xFF         /* dummy: clocks out the status register */
    rcall spi_transfer       /* r16 = status byte after return */
    rcall cs_deassert
    sbrc   r16, WIP_BIT      /* skip rjmp if WIP = 0 (write done) */
    rjmp   .Lwip_poll
    ret
```

After `spi_transfer` with the dummy byte, `r16` holds the status register. `SBRC r16, 0` skips the `RJMP` when bit 0 (WIP) is clear — loop exits.

This is a common SPI pattern: send a command, then send a dummy byte only to clock the response back from the slave.

13.6.4 main — Putting It Together

```
main:
    /* SP, pins, SPI0 init (see full source) */

    ldi    r24, 0x00          /* EEPROM address          */
    ldi    r22, 0xA5          /* data byte               */
    rcall  eeprom_write_byte /* WREN + WRITE command sequence */
    rcall  eeprom_wait_wip   /* wait up to 5 ms for write cycle */

.Lhalt:
    rjmp   .Lhalt
```

13.6.5 Build and Inspect

```
# Assemble and link
avr-as -mmcu=attiny3217 -o spi_eeprom.o spi_eeprom.S
avr-ld -m avr35 -o spi_eeprom.elf spi_eeprom.o

# Or with GCC driver (handles include paths automatically):
avr-gcc -mmcu=attiny3217 -nostartfiles -o spi_eeprom.elf spi_eeprom.S

# Disassemble – verify instruction encoding
avr-objdump -d spi_eeprom.elf

# Generate flashable HEX
avr-objcopy -O ihex spi_eeprom.elf spi_eeprom.hex
```

Expected objdump output for spi_transfer:

```
00000084 <spi_transfer>:
 84: 00 93 c4 08    sts    0x08c4, r16    ; SPI0_DATA ← tx byte
 88: 10 91 c3 08    lds    r17, 0x08c3    ; read SPI0_INTFLAGS
 8c: 17 ff          sbrs   r17, 7         ; wait for IF (bit 7)
 8e: fc cf          rjmp   0x88           ; loop
 90: 00 91 c4 08    lds    r16, 0x08c4    ; read received byte, clears IF
 94: 08 95          ret
```

(Exact addresses vary; the LDS/STS with 16-bit I/O addresses are 4-byte instructions — note the c4 08 bytes encode address 0x08C4 in little-endian.)

That is another good reminder that on AVR, peripheral access cost depends on where the register lives.


13.7 TWI (I²C) Fundamentals

TWI uses two wires: SDA (data) and SCL (clock). Both lines are open-drain with pull-up resistors — any device can pull a line low, but no device drives it high (the pull-up does that). This allows multi-master arbitration and clock stretching by slaves.

This is a very different bus philosophy from SPI. TWI uses fewer wires and built-in addressing, but the protocol is more stateful and the electrical behavior is less direct.

13.7.1 Bus Topology

VCC




Every device has a 7-bit address (0x00–0x7F). The master sends the address at the start of each transaction; only the matching slave responds.

So TWI solves the “which slave am I talking to?” problem in-band, with an address byte, instead of with separate CS wires.

13.7.2 Transaction Structure

A master write transaction (master sends data to slave) looks like:



- **START (S)**: SDA falls while SCL is high.
- **Address + R/W**: 7-bit address followed by a Read(1) or Write(0) bit.
- **ACK (A)**: slave pulls SDA low during the 9th SCK pulse to acknowledge.
- **STOP (P)**: SDA rises while SCL is high.

13.7.3 I²C Speeds

Mode	Max SCL	Common use
Standard	100 kHz	Most sensors, EEPROMs
Fast	400 kHz	High-speed peripherals
Fast-Plus	1 MHz	ATtiny3217 supports this

13.8 TWI0 Master Mode Registers on ATtiny3217

The ATtiny3217's TWI0 peripheral has separate register sets for master and slave operation. We use master-mode registers only.

Register	Address	Purpose
TWI0_CTRLA	0x08A0	SDA hold time, fast-mode-plus enable
TWI0_MCTRLA	0x08A2	Master enable, interrupts, smart mode
TWI0_MCTRLB	0x08A3	Flush, ACK action, bus command (STOP etc.)
TWI0_MSTATUS	0x08A4	WIF, RIF, CLKHOLD, RXACK, ARBLOST, BUSERR, BUSSTATE
TWI0_MBAUD	0x08A5	SCL baud rate register
TWI0_MADDR	0x08A6	Send START + address (write to initiate)
TWI0_MDATA	0x08A7	Master data (read = receive, write = transmit)

13.8.1 TWI0_MCTRLA — Master Control A

Bit: 7 6 5 4 3 2 1 0
 RIEN WIEN – QCEN TIMEOUT[1:0] SMEN ENABLE

Field	Description
RIEN	Read Interrupt Enable (fires on RIF)
WIEN	Write Interrupt Enable (fires on WIF)
QCEN	Quick Command Enable (ACK sent automatically)
TIMEOUT	Bus timeout (00 = disabled)
SMEN	Smart Mode Enable (ACK sent automatically on read)
ENABLE	Enable TWI master

For polling mode (no interrupts): set only TWI_ENABLE_bm.

13.8.2 TWI0_MSTATUS — Master Status

Bit: 7 6 5 4 3 2 1 0
 RIF WIF CLKHOLD RXACK ARBLOST BUSERR BUSSTATE[1:0]

Flag	Description
RIF	Read Interrupt Flag — byte received and ready in MDATA
WIF	Write Interrupt Flag — address/data write accepted, ACK received
CLKHOLD	Clock hold — SCL is being stretched; write MDATA/MCTRLB to release
RXACK	Received ACK: 0 = ACK (slave acknowledged), 1 = NACK
ARBLOST	Arbitration lost (multi-master: another master won the bus)
BUSERR	Bus error (illegal START/STOP condition detected)
BUSSTATE	00=unknown, 01=idle, 10=owner, 11=busy

WIF is the flag you poll most: it sets after the master successfully sends an address byte or a data byte. Check RXACK immediately after — if RXACK=1 the slave sent NACK (not present or busy).

That RXACK polarity is easy to misread because success is encoded as zero.

13.8.3 TWI0_MBAUD — Baud Rate Register

$$f_{\text{SCL}} = f_{\text{CPU}} / (10 + 2 \times \text{BAUD} + f_{\text{CPU}} \times T_{\text{rise}})$$

Ignoring T_{rise} (assume short wires, small capacitance):

$$\text{BAUD} = (f_{\text{CPU}} / f_{\text{SCL}} - 10) / 2$$

$$\text{At 3.333 MHz for 100 kHz: } \text{BAUD} = (3333333/100000 - 10) / 2 = (33.3 - 10) / 2 \approx 11$$

`.equ TWI_BAUD_100K, 11        /* 3.333 MHz → ~100 kHz SCL        */`

13.8.4 TWI0_MCTRLB — Master Control B (Bus Commands)

Write to TWI0_MCTRLB to issue a bus command after each transaction phase. The MCMD field (bits 1:0) controls the state machine:

MCMD value	Symbolic name	Effect
0b00	TWI_MCMD_NOACT_gc	No action

0b01	TWI_MCMD_REPSTART_gc	Repeated START (re-address without STOP)
0b10	TWI_MCMD_RECVTRANS_gc	ACK received byte + clock next (read op)
0b11	TWI_MCMD_STOP_gc	Issue STOP condition, release bus

To send STOP after a write: `sts TWI0_MCTRLB, TWI_MCMD_STOP_gc`

13.9 TWI Initialisation

```
/* TWI0 pins (default mux): SDA = PB0, SCL = PB1
 * Both open-drain – no DDR configuration needed.
 * The TWI peripheral takes control of the pins when ENABLE = 1.
 * External 4.7 kΩ pull-ups to VCC required. */

/* Force bus state to IDLE before we start */
ldi r16, TWI_BUSSTATE_IDLE_gc /* 0x01 */
sts TWI0_MSTATUS, r16

/* Set baud rate: 100 kHz at 3.333 MHz */
ldi r16, TWI_BAUD_100K /* 11 */
sts TWI0_MBAUD, r16

/* Enable master */
ldi r16, TWI_ENABLE_bm /* 0x01 */
sts TWI0_MCTRLA, r16
```

Writing `TWI_BUSSTATE_IDLE_gc` (= 0x01) to `TWI0_MSTATUS` forces the bus state machine to IDLE. Without this, after reset the state is UNKNOWN and the peripheral will not initiate transactions.

This is one of those TWI-specific startup details that feels odd until you remember the peripheral is tracking bus state, not just shifting bytes.

13.10 Master Write Transaction (Polling)

A complete write transaction — send one data byte to a slave device:

Step	Action
1.	Write (slave_addr << 1 0) to <code>TWI0_MADDR</code> → generates START + address
2.	Wait for WIF set
3.	Check <code>RXACK == 0</code> (slave ACK); if <code>RXACK == 1</code> → NACK, abort
4.	Write data byte to <code>TWI0_MDATA</code>
5.	Wait for WIF set
6.	Check <code>RXACK == 0</code>
7.	Write <code>TWI_MCMD_STOP_gc</code> to <code>TWI0_MCTRLB</code> → generates STOP

That list is longer than the SPI case because TWI has more protocol structure: START, address, ACK/NACK, data, ACK/NACK, STOP.

In assembly (using `r24` = slave address, `r22` = data byte):

```
/*
 * twi_write_byte – write one byte to TWI slave
 * r24 = 7-bit slave address (unshifted, 0x00-0x7F)
```



```

* r22 = data byte
* Returns: r24 = 0 on success, 1 on NACK
* Clobbers: r16, r17
*/
twi_write_byte:
    /* Send START + address + W (R/W bit = 0) */
    lsl    r24                /* shift left: addr occupies bits[7:1], bit0=0 */
    sts    TWI0_MADDR, r24    /* write triggers START + 9-bit address phase */

    /* Wait for WIF */
.Lwif_addr:
    lds    r17, TWI0_MSTATUS
    sbrs   r17, TWI_WIF_bp    /* WIF = bit 6 */
    rjmp   .Lwif_addr

    /* Check RXACK */
    sbrc   r17, TWI_RXACK_bp /* RXACK = bit 4; skip ret if ACK (bit = 0) */
    rjmp   .Ltwi_nack

    /* Send data byte */
    sts    TWI0_MDATA, r22

    /* Wait for WIF again */
.Lwif_data:
    lds    r17, TWI0_MSTATUS
    sbrs   r17, TWI_WIF_bp
    rjmp   .Lwif_data

    /* Check RXACK for data */
    sbrc   r17, TWI_RXACK_bp
    rjmp   .Ltwi_nack

    /* Send STOP */
    ldi    r16, TWI_MCMD_STOP_gc /* 0x03 */
    sts    TWI0_MCTRLB, r16

    ldi    r24, 0             /* return 0 = success */
    ret

.Ltwi_nack:
    ldi    r16, TWI_MCMD_STOP_gc
    sts    TWI0_MCTRLB, r16 /* release bus even on failure */
    ldi    r24, 1             /* return 1 = NACK error */
    ret

```

13.10.1 SBRC vs SBRS for RXACK

RXACK = 0 means ACK received (success). The condition for *success* is bit = 0, so we use SBRC (Skip if Bit in Register is Cleared) to skip the NACK handler: if the bit is cleared (ACK), skip over `rjmp .Ltwi_nack`.

This is the opposite of USART's SBRS on DREIF — read the bit polarity carefully for each flag.

That is a general embedded rule: never assume status bits use the same success polarity across peripherals.

13.11 Example: Write to PCF8574 I²C I/O Expander

The PCF8574 is a popular 8-bit I²C I/O expander. Its 7-bit address is `0b0100xxx` where `xxx` is set by hardware pins A2:A0. With A2:A0 = 0b000, the address is `0x20`.

To write a byte to the PCF8574's output latch, a single data byte is sent after the address — no register address needed.

That simplicity is why expanders like the PCF8574 are common first TWI devices: the bus protocol is still there, but the device-side command format is minimal.

```
START | 0x20 W | ACK | 0b10101010 | ACK | STOP
      (slave)                (slave)
```

This drives outputs P7..P0 = 1,0,1,0,1,0,1,0 (alternating).

```
.equ PCF8574_ADDR, 0x20
```

```
main:
```

```
/* TWI0 init (as shown in 13.9) */
ldi r16, TWI_BUSSTATE_IDLE_gc
sts TWI0_MSTATUS, r16
ldi r16, TWI_BAUD_100K
sts TWI0_MBAUD, r16
ldi r16, TWI_ENABLE_bm
sts TWI0_MCTRLA, r16
```

```
/* Write 0xAA to PCF8574 */
```

```
ldi r24, PCF8574_ADDR /* 0x20 */
ldi r22, 0b10101010 /* output data */
rcall twi_write_byte
```

```
.Lhalt:
```

```
rjmp .Lhalt
```

13.11.1 Build

```
avr-gcc -mmcu=attiny3217 -nostartfiles -o pcf8574.elf pcf8574.S
```

```
avr-objcopy -O ihex pcf8574.elf pcf8574.hex
```

```
avrdude -c serialupdi -p t3217 -P /dev/ttyUSB0 -U flash:w:pcf8574.hex
```

13.12 SPI vs TWI Comparison

Feature	SPI	TWI (I ² C)
Wires	4 (MOSI, MISO, SCK, CS×n)	2 (SDA, SCL) + pull-ups
Multi-slave	One CS pin per slave	Addressed (7-bit)
Duplex	Full duplex	Half duplex
Max speed	833 kHz (f/4), up to 2 MHz	100 kHz / 400 kHz / 1 MHz
Protocol overhead	None (raw shift register)	START/STOP/ACK overhead
Hardware	Shift register + IF flag	State machine (WIF/RIF)
Complexity	Low	Medium
Typical devices	EEPROM, flash, ADC, display	Sensor ICs, I/O expanders

The practical split is usually this: use SPI when you want speed and simple byte streaming, and use TWI when you want several slower peripherals to share the same two wires.

13.13 Summary

SPI0 setup (master, mode 0, f/4, manual CS):

```
SPI0_CTRLB ← SPI_SSD_bm           ; disable hardware SS
SPI0_CTRLA ← SPI_MASTER_bm | SPI_ENABLE_bm
```

SPI transfer:

```
STS SPI0_DATA, r16      ; start transfer
poll SPI_IF_bp in SPI0_INTFLAGS
LDS r16, SPI0_DATA      ; get received byte, clears IF
```

25LC010A write sequence:

1. CS↓, send WREN (0x06), CS↑
2. CS↓, send WRITE (0x02), addr, data, CS↑
3. CS↓, send RDSR (0x05), dummy, CS↑ – check WIP (bit 0)
4. Repeat step 3 until WIP = 0 (max 5 ms)

TWI0 setup (master, 100 kHz):

```
TWI0_MSTATUS ← TWI_BUSSTATE_IDLE_gc ; force idle
TWI0_MBAUD   ← 11                    ; 100 kHz at 3.333 MHz
TWI0_MCTRLA  ← TWI_ENABLE_bm
```

TWI master write:

```
STS TWI0_MADDR, (addr << 1)          ; START + address + W
wait WIF, check RXACK == 0
STS TWI0_MDATA, data_byte
wait WIF, check RXACK == 0
STS TWI0_MCTRLB, TWI_MCMD_STOP_gc    ; STOP
```

RXACK polarity: 0 = ACK (slave present), 1 = NACK (slave absent/busy)

13.14 Exercises

1. The `spi_transfer` function polls `SPI_IF_bp` in `SPI0_INTFLAGS`. At 3.333 MHz with `SCK = f/4`, how many CPU cycles elapse from writing `SPI0_DATA` until `IF` becomes set? How many poll iterations does the CPU execute during that time? (Use the `LDS+SBR5+RJMP` cycle count.)
2. The 25LC010A has a 16-byte page buffer. A page write starts with address 0x00 and fills up to 16 bytes before CS↑. Modify `eeeprom_write_byte` into `eeeprom_write_page` that accepts a pointer (`Z`) and a byte count (`r20`, max 16). Send WREN, then one WRITE transaction with all bytes inside the same CS-framed transfer. How does the CS framing differ from the single-byte version?
3. Change the SPI clock to the maximum the 25LC010A allows (2 MHz). Which combination of PRESC and CLK2X achieves the highest SCK without exceeding 2 MHz at 3.333 MHz CPU clock? Write the new `SPI0_CTRLA` value.
4. In `twi_write_byte`, the `BUSSTATE` field of `TWI0_MSTATUS` is not checked before sending. What value does it hold between transactions (after a STOP)? Under what conditions might starting a transaction

with `BUSSTATE = BUSY` cause a problem? How would you guard against it?

5. Write `twi_read_byte`: read one byte from a TWI slave after sending its address with `R/W = 1`. After receiving the byte, send `NACK + STOP` (to tell the slave you don't want more bytes). Which `MCMD` value performs "receive byte + send NACK"? Which `MSTATUS` flag indicates the byte is ready?
6. Using `avr-objdump -d spi_eeprom.elf`, find the `STS SPI0_DATA, r16` instruction. What is its 32-bit encoding? Identify which bytes encode the destination address (`0x08C4`) and which bytes encode the source register (`r16 = R16 = register 16`). Recall that `STS` with a 16-bit address uses the encoding `1001 001d dddd 0000 | kkkk kkkk kkkk kkkk`.

Next: Chapter 14 — Mixing C and Assembly

Chapter 14

Mixing C and Assembly

Pure assembly programs give you total control but can become unmanageable as projects grow. Pure C programs are easier to maintain but sometimes leave performance or size on the table. The practical approach is to mix them: write the bulk of your firmware in C, then drop into assembly for the few routines where it matters — tight ISRs, cycle-counted delay loops, hardware-specific bit manipulation, or an algorithm where the compiler's output is genuinely suboptimal.

This chapter covers three techniques:

1. **Separate assembly file** — write a `.S` subroutine, call it from C.
2. **Extended inline assembly** — embed hand-written instructions inside a C function with `asm volatile(...)`.
3. **Calling convention detail** — which registers the compiler expects you to preserve, which it permits you to destroy.

14.1 Why Mix C and Assembly?

The practical reasons:

Situation	Tool
Single performance-critical routine	Separate <code>.S</code> file, called from C
Access a register with no C syntax	<code>asm volatile</code> (e.g., <code>SWAP</code> , <code>ROL</code>)
Guarantee instruction count / timing	Separate <code>.S</code> or inline <code>asm</code>
Atomic multi-step hardware sequence	Inline <code>asm</code> (prevents reordering)
Large program logic	C — readability wins

Writing the *entire* program in assembly buys diminishing returns beyond a few hundred lines. Compilers are good. Write C; reach for assembly when you can *measure* the gain.

That is the mindset for this chapter: assembly is a tool for the narrow parts where control matters more than readability, not a badge of honor to use everywhere.

14.2 The GCC AVR Calling Convention

GCC for AVR follows a documented ABI (Application Binary Interface). You must match it exactly when writing assembly routines that C calls or are called by.

14.2.1 Register Classification

Registers	Role	Must callee preserve?
R0	Scratch	No (caller-saved); may be clobbered
R1	Always zero (zero reg)	No – but you MUST restore R1 to 0 if your asm uses R1 as scratch
R2–R17	Call-saved	YES – push before use, pop before ret
R18–R27	Call-clobbered	No (caller-saved) – free to destroy
R28 (YL)	Call-saved	YES – push before use, pop before ret
R29 (YH)	Call-saved	YES – push before use, pop before ret
R30 (ZL)	Call-clobbered	No
R31 (ZH)	Call-clobbered	No

Call-saved means the *callee* (your assembly routine) is responsible for leaving those registers unchanged when it returns. If you need to use R2–R17 or R28–R29, push them at the top of your function and pop them before RET.

Call-clobbered means the *caller* (C code) knows these registers may be destroyed by any function call. The compiler saves them itself if it needs them across a call site.

If you remember only one rule, remember this one: C code is willing to let you trash call-clobbered registers, but it expects call-saved registers to come back exactly as they were.

14.2.2 Argument Passing

Arguments are allocated to registers starting from **R25 going down**, in left-to-right parameter order. Each argument starts in an even-numbered register position, which is why 8-bit arguments appear in R24, R22, R20, and so on.

C type	Width	Registers used
uint8_t / char	1 byte	R24 (first arg), R22 (second), R20 (third) ...
uint16_t / ptr	2 bytes	R25:R24 (first arg), R23:R22 (second) ...
uint32_t	4 bytes	R25:R24:R23:R22 (first arg)

First argument always uses the highest available register pair, so:

```
uint8_t crc8(const uint8_t *data, uint8_t len);
//                ↑                ↑
//                R25:R24          R22
```

- data (pointer, 16-bit) → **R25:R24** (high byte in R25, low byte in R24)
- len (uint8_t, 8-bit) → **R22**

This is the point that usually surprises people the first time they inspect a mixed C/assembly call site: an 8-bit first argument often lands in R24, but a pointer or uint16_t first argument consumes the full R25:R24 pair.

14.2.3 Return Values

C return type	Registers
uint8_t / char	R24
uint16_t / ptr	R25:R24
uint32_t	R25:R24:R23:R22
void	—

Return the value in R24 (or R25:R24 for 16-bit) and the calling C code picks it up automatically.

14.2.4 The Zero Register Rule

GCC's generated code relies on R1 always being zero. It uses R1 to zero-initialize memory and as a zero source for SBIW-equivalents. If your assembly routine uses R1 as a scratch register, **restore it to zero before returning**:

```
eor r1, r1          /* restore zero register before RET          */
ret
```

Failure to do this causes bizarre bugs that are very hard to trace — the next C function that relies on R1 = 0 will see garbage.

Those bugs feel random because the mistake usually happens in one function and the visible failure appears later in unrelated compiler-generated code.

14.3 Writing an Assembly Subroutine for C

The workflow:

1. Write `function_name.S` with `.global function_name`.
2. Declare extern prototype in a C header (or directly in the `.c` file).
3. Compile both files together with `avr-gcc`.

14.3.1 Rules for the .S File

```
#include <avr/io.h>          /* gives register names and bit masks      */

.section .text
.global my_function          /* MUST match C prototype name exactly  */
my_function:
    /* your code */
    ret
```

- **.global** exports the symbol so the linker can resolve it.
- **#include <avr/io.h>** works in `.S` files because GCC preprocesses them (the `.S` extension, uppercase, triggers preprocessing; `.s` does not).
- Use only **R18–R27**, **R30–R31** freely. Push/pop **R2–R17**, **R28–R29** if used.
- **Do not call SEI or CLI** unless the function is explicitly documented as doing so — the caller may be inside a critical section.

In other words, a stand-alone assembly routine should behave like a normal C function: take inputs where the ABI says, return outputs where the ABI says, and avoid surprising the caller with hidden global side effects.

14.4 Modern AVR Header Conventions

Microchip's newer AVR headers expose peripherals in two equivalent styles: flat register names that map cleanly to assembly, and structured names that are idiomatic in C.

```
PORTA.DIRSET = PIN3_bm; /* structured C style */
USART0.CTRLB = USART_TXEN_bm | USART_RXEN_bm;
TCA0.SINGLE.CTRLA = TCA_SINGLE_ENABLE_bm;
```

The same registers appear in assembly as flat symbols from `<avr/io.h>`:

```
sts    PORTA_DIRSET, r16
sts    USART0_CTRLB, r16
sts    TCA0_SINGLE_CTRLA, r16
```

For mixed C/assembly work, the important suffixes are:

Suffix	Meaning
<code>_bm</code>	bit mask: already shifted, e.g. <code>USART_TXEN_bm = 0x40</code>
<code>_bp</code>	bit position: bit number for SBI/CBI/SBRS, e.g. <code>USART_DREIF_bp = 5</code>
<code>_gc</code>	group configuration constant for a bit field, e.g. <code>TCA_SINGLE_WGMODE_FRQ_gc</code>

Use `_bm` when ORing or ANDing whole register values in C or loading a mask in assembly. Use `_bp` only when an instruction wants a literal bit number. Use `_gc` names for multi-bit fields instead of open-coded shifts: they document the selected hardware mode and survive refactors better than raw constants.

This naming scheme looks noisy at first, but it solves a real problem: it tells you whether a symbol is meant to be used as a mask, a bit number, or an already-shifted field value.

Newer AVR C code also uses module types such as `PORT_t`, `USART_t`, or `TCB_t` * to write reusable helpers that can target any peripheral instance:

```
static inline void tcb_enable(volatile TCB_t *tc)
{
    tc->CTRLA |= TCB_ENABLE_bm;
}
```

That style does not change the hardware naming underneath. It is only a C-side view over the same memory-mapped registers your assembly accesses directly.

14.5 Example: CRC-8 Subroutine

CRC-8/SMBUS uses polynomial 0x07. The algorithm XORs each byte into an accumulator, then processes each bit: shift left, and if the MSB was 1, XOR with the polynomial.

C equivalent (for reference):

```
uint8_t crc8_c(const uint8_t *data, uint8_t len) {
    uint8_t crc = 0;
    while (len--) {
        crc ^= *data++;
        for (uint8_t i = 0; i < 8; i++) {
            if (crc & 0x80)
                crc = (crc << 1) ^ 0x07;
            else
                crc = (crc << 1);
        }
    }
    return crc;
}
```



```

        crc <=<= 1;
    }
}
return crc;
}

```

Assembly implementation (crc8.S):

```
#include <avr/io.h>
```

```
.equ CRC8_POLY, 0x07
```

```
.section .text
```

```
.global crc8
```

```
crc8:
```

```

    movw    ZL, r24          /* Z ← data pointer from R25:R24      */
    mov     r20, r22         /* r20 ← len                          */
    ldi     r24, 0x00        /* CRC accumulator = 0                */
    ldi     r25, CRC8_POLY   /* polynomial constant (0x07)         */
                                /* R25 is now free: pointer already in Z */
    tst     r20
    breq    .Lcrc_done      /* len = 0: return 0                  */

```

```
.Lcrc_byte:
```

```

    ld      r18, Z+          /* load byte from SRAM, advance pointer */
    eor     r24, r18         /* crc ^= byte                          */
    ldi     r19, 8           /* 8 bits to process                   */

```

```
.Lcrc_bit:
```

```

    lsl     r24              /* shift MSB into carry                */
    brcc    .Lcrc_no_xor    /* carry = 0: no feedback              */
    eor     r24, r25        /* carry = 1: XOR with polynomial      */

```

```
.Lcrc_no_xor:
```

```

    dec     r19
    brne    .Lcrc_bit

```

```

    dec     r20
    brne    .Lcrc_byte

```

```
.Lcrc_done:
```

```

    ret                                /* return value in R24                */

```

Register use breakdown:

R24 – CRC accumulator, also return register (no conflict)
 R25 – polynomial constant 0x07 (reused after pointer copied to Z)
 R18 – current data byte from memory
 R19 – bit loop counter (0..7)
 R20 – byte loop counter (len down to 0)
 Z (R31:R30) – pointer into data array

All are call-clobbered — no push/pop needed.

Why reuse R25? After `MOVW ZL, R24`, the pointer lives in Z (= R31:R30). R24 and R25 are free; R24 becomes the CRC accumulator (and the return register), R25 becomes the polynomial constant. One fewer `LDI` per bit iteration versus loading 0x07 from a constant — minor but illustrates thinking about register reuse.

That kind of reuse is where hand-written assembly still earns its keep: not in magic instructions, but in being deliberate about where every value lives.

14.6 Calling the Assembly Routine from C

main.c:

```
#include <stdint.h>

/* Declare the assembly function.
 * "extern" is optional in C (functions are extern by default) but explicit
 * here for clarity. The prototype MUST match the assembly ABI exactly. */
uint8_t crc8(const uint8_t *data, uint8_t len);

static uint8_t msg[] = {0x01, 0x02, 0x03, 0x04};
volatile uint8_t result; /* volatile: prevent optimiser eliding the call */

int main(void) {
    result = crc8(msg, sizeof(msg));
    /* Expected: 0x84 (CRC-8/SMBUS of that message) */
    for (;;)
        return 0;
}
```

14.6.1 Build Commands

```
# Compile and assemble in one step, then link
avr-gcc -mmcu=attiny3217 -O2 -o crc_demo.elf main.c crc8.S

# Inspect what the compiler generated for the call site
avr-objdump -d crc_demo.elf | grep -A 20 '<main>'

# Generate flashable HEX
avr-objcopy -O ihex crc_demo.elf crc_demo.hex
```

14.6.2 Inspecting the Call Site

avr-objdump shows one typical instruction sequence around the `crc8()` call when `msg` has a fixed SRAM address:

```
; C: result = crc8(msg, sizeof(msg));

ldi    r24, lo8(msg)          ; low byte of &msg[0]
ldi    r25, hi8(msg)          ; high byte of &msg[0]
ldi    r22, 4                 ; R22 = len = 4 (second arg)
call   <crc8>                 ; branch to our assembly
sts    result, r24             ; store R24 (return value) into 'result'
```

You can see exactly how GCC places arguments into R25:R24 and R22 — matching the calling convention described in Section 14.2.

That is why `avr-objdump` is such a useful reality check in mixed-language work. It lets you confirm the ABI with actual generated code instead of memory.

14.7 Cycle Count Comparison

Compare GCC's output vs the hand-written assembly for the CRC inner loop.

C version inner loop (avr-gcc -O2, ATtiny3217 target):

```
; crc ^= byte (r24 = crc, r20 = *ptr)
eor    r24, r20          ; 1 cycle
; 8-bit loop:
lsl    r24                ; 1 cycle
brcc   .lno_xor           ; 1 cycle (2 if branch taken)
eor    r24, r18           ; 1 cycle (r18 loaded with 0x07 from a saved register)
.lno_xor:
dec    r19                ; 1 cycle
brne   .lcrc_bit          ; 1 (2 if taken)
```

The compiler's output is nearly identical to the hand-written version — demonstrating that -O2 is strong for arithmetic-only loops. The hand-written version wins when the compiler has to spill registers for a larger function or when the polynomial is a compile-time constant the compiler doesn't pre-load into a register.

This is an important lesson: if the disassembly already looks clean, assembly may not buy you much. Inspect first, then decide.

14.8 Inline Assembly: `asm volatile`

For short sequences (1–4 instructions) that must be embedded inside a C function — especially when instruction ordering matters — use `asm volatile`.

14.8.1 Basic Syntax

```
asm volatile (
    "instruction1" "\n\t"
    "instruction2" "\n\t"
    : output_operands
    : input_operands
    : clobbers
);
```

- `"\n\t"` separates instructions (newline for the assembler, tab for alignment in the generated `.s` file). Multiple instructions in one string literal.
- **Output operands:** `"=r" (var)` — write result into any register, store in `var`.
- **Input operands:** `"r" (expr)` — load expression into any register.
- **Clobbers:** tell GCC which registers/memory this `asm` modifies.

Inline assembly is not just “assembly inside C.” It is a contract with the compiler: you must describe what goes in, what comes out, and what gets disturbed, or the optimizer will make assumptions that break your code.

14.8.2 Constraint Letters

Constraint	Meaning
r	Any general-purpose register
d	Upper registers (R16–R31, required for LDI)
I	6-bit I/O address (0–63), for IN/OUT
n	Compile-time constant (any value)
=	Write-only output (register not read first)
+	Read-write operand (register read AND written)
0–9	Same register as operand N (tie operands together)

14.9 Inline Assembly Examples

14.9.1 Example 1 — NOP (Delay One Cycle)

```
asm volatile ("nop");
```

No operands, no clobbers. GCC inserts exactly one NOP instruction. `volatile` prevents the optimiser from removing it (a pure NOP has no visible side-effect).

Without `volatile`, GCC is allowed to conclude that a `nop` changes nothing and delete it.

14.9.2 Example 2 — Read and Disable Interrupts Atomically

```
static inline uint8_t save_and_cli(void) {
    uint8_t sreg;
    asm volatile (
        "in  %0, %1"   "\n\t"
        "cli"          "\n\t"
        : "=r" (sreg)
        : "I"  (_SFR_IO_ADDR(SREG))
        : "memory"
    );
    return sreg;
}
```

`_SFR_IO_ADDR(SREG)` converts the memory-mapped SREG address to its 6-bit I/O address (required by IN). The "I" constraint enforces this. `"=r"` allocates any register for `sreg`; the `%0` placeholder is replaced with that register name by the assembler.

Restore interrupts:

```
static inline void restore_sreg(uint8_t sreg) {
    asm volatile (
        "out %0, %1"
        :
        : "I" (_SFR_IO_ADDR(SREG)), "r" (sreg)
        : "memory"
    );
}
```

Two input operands, no outputs. `%0` = SREG I/O address (constant in the instruction encoding), `%1` = the register holding `sreg`.

The split into `save` and `restore` also keeps the generated code easy to audit: one helper captures the old interrupt state, and the other puts that exact state back.

14.9.3 Example 3 — SWAP Nibbles

AVR has a dedicated SWAP instruction (1 cycle) that exchanges the high and low nibbles of a byte. C has no equivalent; the compiler generates two shifts and an OR. Inline assembly is cleaner:

```
static inline uint8_t swap_nibbles(uint8_t val) {
    asm volatile (
        "swap %0"
        : "+r" (val)    /* "+r": val is read in, written out, same register */
    );
    return val;
}
```

The "+r" constraint means the register is both read (input) and written (output). GCC picks one register, loads val into it, runs SWAP, then reads the result from the same register — stored back into val.

14.9.4 Example 4 — Atomic 16-bit Increment

Incrementing a uint16_t in C compiles to two instructions. An ISR can fire between them, corrupting the value. Use save_and_cli/restore_sreg:

```
volatile uint16_t packet_count;

void count_packet(void) {
    uint8_t sreg = save_and_cli(); /* save I-flag, disable interrupts */
    packet_count++;                /* two-instruction critical section */
    restore_sreg(sreg);           /* restore I-flag (re-enables if set) */
}
```

This pattern is correct even when called from an ISR (interrupts already disabled — CLI is a no-op if already clear; restore_sreg restores the old value, not unconditionally re-enables). The "memory" clobber keeps GCC from moving memory accesses into or out of the critical section.

That last point matters as much as the CLI itself. A critical section is only correct if the compiler is forced to keep the protected loads and stores inside the region you intended.

14.10 Memory Clobber

Sometimes inline assembly reads or writes memory without explicit operands — for example, an instruction that clears a peripheral flag as a side effect. Tell GCC with "memory" in the clobber list:

```
asm volatile (
    "lpm    r0, Z+" "\n\t"
    :
    : "r0", "r30", "r31", "memory"
);
```

"memory" is a special clobber that forces GCC to flush all cached values back to memory before the asm and reload them after. Without it, the compiler might reorder loads/stores across the asm block.

So "memory" is about compiler ordering, not CPU ordering. It does not emit a hardware fence; it tells GCC to stop pretending memory stayed unchanged.

14.11 Building Mixed Projects

14.11.1 Makefile Fragment

```
MCU      = attiny3217
CC       = avr-gcc
CFLAGS   = -mmcu=$(MCU) -O2 -Wall
OBJCOPY  = avr-objcopy

SRCS_C   = main.c
SRCS_S   = crc8.S

OBJS     = $(SRCS_C:.c=.o) $(SRCS_S:.S=.o)

crc_demo.elf: $(OBJS)
    $(CC) $(CFLAGS) -o $@ $^

%.o: %.c
    $(CC) $(CFLAGS) -c -o $@ $<

%.o: %.S
    $(CC) $(CFLAGS) -c -o $@ $<

crc_demo.hex: crc_demo.elf
    $(OBJCOPY) -O ihex $< $@
```

avr-gcc handles both .c and .S files — it invokes avr-as for .S after running the C preprocessor (so #include and .equ-equivalents work).

That is why mixed projects usually build more smoothly through avr-gcc than by invoking avr-as and avr-ld separately for every file.

14.11.2 Verifying the ABI with objdump

```
# Show symbol table – confirm crc8 is exported
avr-nm crc_demo.elf | grep crc8

# Disassemble to verify calling convention at call site
avr-objdump -d crc_demo.elf
```

Look for the sequence before call <crc8>:

```
movw  r24, <ptr_reg>    ; first arg: pointer in R25:R24
ldi   r22, N            ; second arg: len in R22
call  <crc8>
; ... R24 holds return value after return
```

If you see mov r22, r24 or similar before the call, double-check your prototype matches the .S argument layout.

When the prototype is wrong, the code may still compile and link cleanly. The failure shows up only at runtime, which is why ABI mismatches are easy to miss.

14.12 Common Mistakes

14.12.1 Mistake 1 — Wrong Register for LDI

```
ldi    r15, 0x07      /* ERROR: LDI only works for R16-R31 */
```

Fix: use R16–R31. For call-saved registers (R16–R17), push/pop them.

14.12.2 Mistake 2 — Forgetting .global

```
.section .text
my_function:      /* no .global – linker cannot see this symbol */
    ret
```

Fix: add .global my_function before the label.

14.12.3 Mistake 3 — Wrong Argument Register

A common error is assuming the first argument is in R24 when it is actually a pointer (16-bit) and therefore in R25:R24:

```
; uint8_t sum(const uint8_t *arr, uint8_t n);
sum:
    ld    r18, r24      /* WRONG: r24 holds low byte of pointer, not byte */
```

Fix: movw ZL, r24 to get the full pointer into Z, then ld r18, Z+.

14.12.4 Mistake 4 — Corrupting a Call-Saved Register

```
my_func:
    ldi    r16, 0        /* R16 is call-saved – must be preserved!          */
    /* ... */
    ret                /* caller sees r16 changed: undefined behaviour    */
```

Fix: push R16 at the top, pop before RET.

14.12.5 Mistake 5 — Leaving R1 Non-Zero

R1 must always be zero when code returns to the C runtime. GCC relies on R1 as a zero source for generated code. The ATtiny3217 has no MUL instruction, but R1 can still be corrupted by inline assembly or manual shifts that use R1 as a temporary. Restore it before returning:

```
my_func:
    /* ... uses R1 as temp scratch ... */
    eor    r1, r1        /* restore zero register before RET          */
    ret
```

On devices that do have MUL (ATmega series), MUL writes R1:R0 and R1 must be explicitly cleared after use:

```
mul    r24, r25      /* NOT available on ATtiny3217 */
mov    r24, r0
eor    r1, r1        /* restore zero register */
ret
```

14.13 Summary

GCC AVR calling convention:

First arg: R25:R24 (16-bit) or R24 (8-bit)

Second arg: R23:R22 (16-bit) or R22 (8-bit)

Return: R24 (8-bit), R25:R24 (16-bit)

Call-saved (callee preserves): R2-R17, R28-R29

Call-clobbered (free to use): R18-R27, R30-R31

R0: scratch; R1: MUST be zero on return (restore with EOR R1, R1)

Writing ASM for C:

```
.global function_name      ← export the symbol
#include <avr/io.h>         ← register names (works in .S via preprocessor)
Use call-clobbered regs freely
PUSH/POP call-saved regs if used
```

Build:

```
avr-gcc -mmcu=attiny3217 -O2 -o out.elf main.c routine.S
```

Inline assembly template:

```
asm volatile (
    "instruction %0, %1" "\n\t"
    : "=r" (output_var)      /* output operands */
    : "r" (input_expr)       /* input operands */
    : "r18"                  /* clobbers */
);
```

Constraints:

```
"r" = any register
"d" = upper register (R16-R31), needed for LDI
"I" = 6-bit I/O address (for IN/OUT)
"=" = write-only output
"+" = read-write (same reg for in and out)
"memory" = flush/reload all memory across the asm block
```

14.14 Exercises

1. Write an assembly subroutine `add16` with C prototype:

```
uint16_t add16(uint16_t a, uint16_t b);
```

Add `a` and `b` using `ADD + ADC` for the two register pairs. Identify which registers hold `a` and `b` on entry, and which register pair holds the return value.

2. GCC sometimes passes a `uint32_t` argument in `R25:R24:R23:R22`. Write an assembly subroutine `count_bits32` that counts the number of set bits (popcount) in a 32-bit value. Return the count in `R24`. Hint: process each of the four bytes with a loop over bits, accumulating into `R24`.
3. Write an assembly subroutine `mul8` with prototype `uint16_t mul8(uint8_t a, uint8_t b)`. The ATtiny3217 has hardware multiply; implement this version using shift-and-add anyway so the algorithm is clear and portable to AVR parts without `MUL`. Arguments arrive in `R24` (`a`) and `R22` (`b`); return the 16-bit result in `R25:R24`. How many cycles does your implementation take for the worst-case input (255×255)?

4. Rewrite the `swap_nibbles` inline example (Section 14.8) without inline assembly — using only C bit-shifts and masks. Compare the disassembly of both versions with `avr-objdump -d`. How many cycles does each version take at -O2?
5. In the `save_and_cli / restore_sreg` pattern, what does `asm volatile` guarantee, and what extra guarantee does the "memory" clobber add? Explain why both matter when protecting a critical section around `packet_count++`.
6. Verify the CRC-8 implementation: a. Assemble `crc_demo.elf` and run it in `simavr`:

```
simavr -m attiny3217 crc_demo.elf
```

b. Inspect the value of `result` in SRAM after the program halts. Using `avr-objdump --syms crc_demo.elf`, find the address of `result`, then dump that address with `simavr`'s built-in GDB interface. c. Confirm the result is `0x84` (CRC-8/SMBUS of `{0x01, 0x02, 0x03, 0x04}`).

Next: Chapter 15 — Optimisation Techniques

Chapter 15

Optimisation Techniques

Assembly gives you complete control over cycle counts, register allocation, and memory layout. That control is only useful if you know where to look for waste and which transformations actually help. This chapter covers four practical techniques, each with a concrete measurement strategy so you can verify the gain rather than guess at it.

The techniques:

1. **Cycle counting** — read the disassembly, count cycles, build a baseline.
2. **Avoid SRAM** — keep hot variables in registers; SRAM loads/stores cost cycles.
3. **Loop unrolling** — trade code size for fewer branch and counter instructions.
4. **Flash lookup tables** — replace computation with a single LPM instruction.

Each technique applies to different situations. Measure first, optimise second.

That sentence is the whole chapter in one line. Optimisation on AVR works best when it is concrete: count instructions, count cycles, and compare versions.

15.1 Cycle Counting with `avr-objdump`

Every AVR instruction has a documented cycle count (1 or 2 for most, 3–4 for branch-taken, 4–5 for CALL/RET). Count cycles in the inner loop — that is where time is spent.

The important habit is to stop thinking in vague terms like “this feels fast” and start thinking in exact terms like “this loop costs 7 cycles per byte.”

15.1.1 Workflow

1. Compile with debug info so `objdump` can correlate source lines

```
avr-gcc -mmcu=attiny3217 -O2 -g -o bench.elf bench_main.c memcpy_naive.S
```

2. Disassemble – see instructions and addresses

```
avr-objdump -d bench.elf
```

Example output for `memcpy_naive`:

```
00000142 <memcpy_naive>:
142:  fc 01          movw    r30, r24      ; 1 cycle  Z=dst... wait,
                                ; see §15.2 – dst in R25:R24
144:  b7 01          movw    r22, r22      ; ...
```

The raw hex `fc 01` is the 16-bit instruction word. `avr-objdump` decodes it to the mnemonic (`movw`), operands, and address offset. You count cycles by looking each instruction up in the AVR Instruction Set Manual; the table below covers the instructions used in this chapter.

15.1.2 Cycle Reference (common instructions)

Instruction	Cycles	Notes
NOP	1	
MOV Rd, Rr	1	
MOVW Rd, Rr	1	copies two registers at once
LDI Rd, K	1	upper regs only (R16-R31)
ADD / ADC	1	
AND / OR / EOR	1	
INC / DEC	1	
LSR / LSL / ROR	1	
TST Rd	1	same as AND Rd,Rd, sets flags
BREQ / BRNE (not taken)	1	
BREQ / BRNE (taken)	2	
RJMP	2	
LD Rd, Z+	2	post-increment; LD Rd, Z = 1 on XMEGA
ST X+, Rr	2	post-increment
LPM Rd, Z+	3	load from flash, post-increment
PUSH / POP	2	(1 extra for XMEGA series vs classic)
CALL / RCALL	4 / 3	(RCALL saves 1 word vs CALL)
RET	5	(4 on some cores)

15.1.3 Annotating the Inner Loop

Write the cycle count beside each instruction in a comment:

```
.Lnaive_loop:
    ld    r18, Z+      /* 2 cycles: load byte, advance Z      */
    st    X+, r18      /* 2 cycles: store byte, advance X      */
    dec   r20          /* 1 cycle: decrement counter          */
    brne  .Lnaive_loop /* 2 cycles (taken), 1 (fall-through)  */
```

Inner loop body (taken path): $2 + 2 + 1 + 2 = 7$ cycles per byte.

On the last iteration BRNE falls through (1 cycle instead of 2), saving 1 cycle.

For 16 bytes: $7 \times 15 + 6 = 111$ cycles (+ prologue/epilogue ≈ 10 cycles).

This is your **baseline** to beat.

Without a baseline, an “optimisation” is just a guess that happened to produce different code.

15.2 Avoid SRAM — Keep Variables in Registers

SRAM access via LDS/STS costs **2 cycles** for the instruction plus an additional cycle for the SRAM access itself on classic AVR (3 total). LD through a pointer register (X/Y/Z) is 1–2 cycles. Registers are free.

15.2.1 The Problem

A naive loop that spills a loop variable to SRAM on every iteration:

```
/* C that the compiler might spill if register pressure is high */
volatile uint8_t counter = 0;
for (uint8_t i = 0; i < 100; i++) {
    counter++;
}
```

If counter is volatile (as shown), the compiler **must** load and store it on every iteration — it cannot keep it in a register. Unnecessary volatile on internal variables is a performance anti-pattern in AVR firmware.

Beginners often use volatile as if it means “important variable.” It does not. It means “this value may change outside the compiler’s view, so reload and rewrite it exactly as written.”

15.2.2 The Rule

- Mark variables volatile only when they are accessed by ISRs or hardware.
- Internal loop counters, accumulators, and intermediate results: keep in registers, never declare volatile.
- In hand-written assembly, simply *never allocate a hot variable to an SRAM address* — assign it a register and keep it there.

15.2.3 Register Allocation Strategy

Plan register use before writing the first instruction. Write it out:

```
Routine: block_sum – sum N bytes starting at ptr
R24 = accumulator (also return register – convenient)
Z (R31:R30) = ptr (from R25:R24 via MOVW ZL, R24)
R22 = len
R18 = temp byte from load
Clobbered: R18, R22, R24, Z – all call-clobbered, no PUSH needed
```

Example:

```
#include <avr/io.h>

.section .text
.global block_sum
/* uint8_t block_sum(const uint8_t *ptr, uint8_t len);
 * Returns sum of len bytes (wraps on overflow). */
block_sum:
    movw    ZL, r24          /* Z = ptr */
    ldi     r24, 0           /* accumulator = 0 (reuses the arg register) */
    tst     r22
    breq    .Lbs_done
.Lbs_loop:
    ld      r18, Z+          /* load byte */
    add     r24, r18         /* accumulate */
    dec     r22              /* len-- */
    brne    .Lbs_loop
.Lbs_done:
    ret                    /* return value already in R24 */
```

No SRAM touched after entry. Every variable lives in a register for the entire duration of the routine.

That is usually the first big win in small AVR routines: not exotic instructions, just avoiding repeated trips out to SRAM.

15.2.4 When Registers Run Out

The AVR has 32 general-purpose registers but only 12 are freely usable without saving (R18–R27, R30–R31). For routines that genuinely need more:

1. PUSH/POP call-saved registers (R2–R17, R28–R29) — cheap if done once per call, not per iteration.
2. Use the Y register (R29:R28) as a frame pointer to access function-local variables on the stack — same pattern the compiler uses.
3. Decompose the routine into smaller subroutines, each with fewer live variables at once.

In practice, that third option is often the cleanest. If a routine needs so many live values that everything spills, it may be doing too much at once.

15.3 Loop Unrolling

Each loop iteration pays overhead: DEC (1 cycle) + BRNE (2 cycles taken) = **3 overhead cycles**. For a 2-cycle loop body, overhead is 60% of total cost. Unrolling processes multiple elements per iteration, amortising that overhead.

15.3.1 Naive vs Unrolled: memcpy

Naive inner loop (1 byte per iteration):

Cycles per byte:

```
ld    r18, Z+    = 2
st    X+, r18    = 2
dec   r20        = 1
brne                     = 2 (taken)
```

7 cycles/byte

Unrolled x4 inner loop (4 bytes per iteration):

`.Lfast_block:`

```
ld    r18, Z+    /* 2 */
ld    r19, Z+    /* 2 */
ld    r22, Z+    /* 2 */
ld    r23, Z+    /* 2 */
st    X+, r18    /* 2 */
st    X+, r19    /* 2 */
st    X+, r22    /* 2 */
st    X+, r23    /* 2 */
dec   r20        /* 1 */
brne   .Lfast_block /* 2 (taken) */
```

Cycles for 4 bytes: $4 \times 2 + 4 \times 2 + 1 + 2 = 19$. Cycles per byte: $19 / 4 = 4.75$ cycles/byte.

Overhead dropped from $3/7 = 43\%$ to $3/19 = 16\%$.

That is why unrolling helps most when the original loop body is very small. If the useful work is only a couple of instructions, branch overhead dominates.

15.3.2 The Tail Problem

16 bytes unrolled x4 works cleanly: $16 / 4 = 4$ full blocks, no remainder. But len is arbitrary. For len = 13: 3 full blocks + 1 remaining byte.

The solution: compute `tail = len & 3` (remainder), `blocks = len >> 2` (quotient), run the block loop, then clean up the tail with a small loop.

This is the part people forget when they first hear “just unroll the loop.” The fast path gets simpler, but the general-case control flow gets a little more complicated because arbitrary lengths rarely divide evenly.

```
memcpy_fast:
    movw    ZL, r22          /* Z = src          */
    movw    XL, r24          /* X = dst          */
    tst     r20
    breq    .Lfast_done

    mov     r21, r20
    andi    r21, 0x03        /* tail = len & 3   */
    lsr     r20
    lsr     r20              /* blocks = len >> 2 */
    breq    .Lfast_tail      /* skip block loop if len < 4 */

.Lfast_block:
    ld      r18, Z+
    ld      r19, Z+
    ld      r22, Z+
    ld      r23, Z+
    st      X+, r18
    st      X+, r19
    st      X+, r22
    st      X+, r23
    dec     r20
    brne    .Lfast_block

.Lfast_tail:
    tst     r21
    breq    .Lfast_done
    ld      r18, Z+
    st      X+, r18
    dec     r21
    brne    .Lfast_tail

.Lfast_done:
    ret
```

Full source at `src/memcpy_fast.S`.

15.3.3 Measuring the Gain

```
avr-gcc -mmcu=attiny3217 -O2 -o bench.elf bench_main.c \
    memcpy_naive.S memcpy_fast.S lut_sin.S
```

```
avr-objdump -d bench.elf | grep -A 40 '<memcpy_naive>'
avr-objdump -d bench.elf | grep -A 40 '<memcpy_fast>'
```

Count cycles in each inner loop. For `len = 16`:

```
memcpy_naive: 7 cycles/byte × 16 = 112 cycles + overhead ≈ 122 total
memcpy_fast:  4.75 cycles/byte × 16 = 76 cycles + overhead ≈ 88 total
Speedup: ~1.4×
```

For larger buffers (e.g. 256 bytes), the overhead becomes negligible and the speedup approaches $7 / 4.75 \approx 1.47\times$. Unrolling $\times 8$ would approach $7 / 4.375 = 1.6\times$ at the cost of doubling code size in the block loop.

So unrolling is never free. You are buying speed with flash bytes and with a little more control-flow complexity.

15.3.4 When NOT to Unroll

- **Small loop counts** (< 4 iterations): the prologue code for computing blocks and tail may cost more than it saves.
- **Flash-constrained devices**: every unrolled instruction costs 2 bytes of flash. An ATtiny85 has 8 KB; unrolling a $\times 16$ loop wastes 30+ bytes.
- **When the compiler already does it**: GCC at -O3 unrolls automatically. Check the compiler output first.

That last point matters. Hand-unrolling code the compiler already optimized is an easy way to make the source harder to maintain for no measurable gain.

15.4 Lookup Tables in Flash (PROGMEM)

Some computations are expensive but have a small, bounded input domain. If the input fits in 8 bits (256 possible values), precompute the result for every input and store it in flash. A single LPM instruction (3 cycles) replaces an arbitrary number of arithmetic instructions.

This is one of the cleanest optimisation trades in embedded work: spend memory once so you do not have to spend cycles every time.

15.4.1 AVR Flash Access: LPM

LPM loads one byte from program memory (flash) into a register using the Z pointer. Three forms:

```
LPM Rd, Z      ; load flash[Z] into Rd – 3 cycles, Z unchanged
LPM Rd, Z+     ; load flash[Z] into Rd, then Z++ – 3 cycles
LPM            ; load flash[Z] into R0 – 3 cycles (legacy form)
```

Z is a byte address into flash (not a word address). For a table at label mytable:

```
ldi    ZL, lo8(mytable) /* low byte of flash address          */
ldi    ZH, hi8(mytable) /* high byte of flash address         */
add    ZL, r24          /* add index (unsigned byte offset)      */
adc    ZH, r1           /* propagate carry (r1 = 0)              */
lpm    r24, Z           /* load table[index]                     */
```

lo8() and hi8() are assembler built-ins that extract the low and high bytes of a symbol address. Because AVR flash addresses are in a separate address space, you must use these — a plain .word mytable would give the wrong address form.

15.4.2 Placing a Table in Flash

Use the .progmem.data section so the linker places it in flash, not SRAM:

```
.section .progmem.data, "a", @progbits
.global my_table
my_table:
    .byte 0x00, 0x01, 0x04, 0x09 /* etc. */
```

Alternatively in C with PROGMEM:

```
#include <avr/pgmspace.h>
const uint8_t my_table[256] PROGMEM = { 0, 1, 4, 9, /* ... */ };
```

Both land in flash. The assembly form is cleaner for tables built by hand.

The important idea is not the syntax. It is that the table lives in program memory, so it does not consume precious SRAM at startup.

15.4.3 Example: 8-bit Sine Table

Computing $\sin(\text{angle})$ in 8-bit AVR arithmetic requires multiplication, shifts, and polynomial approximation — dozens of cycles. A 256-entry table covers all possible inputs; lookup is 5 instructions + 3-cycle LPM.

Table layout: 256 entries, each byte is $128 + 127 \times \sin(n \times 2\pi / 256)$, unsigned. Angle 0 = 0°, angle 64 = 90°, angle 128 = 180°, angle 192 = 270°.

Sine lookup routine (src/lut_sin.S):

```
#include <avr/io.h>

/* uint8_t sin8(uint8_t angle);
 * Returns 0..255: 128 = 0.0, 255 = +1.0, 0 = -1.0 */

.section .progmem.data,"a",@progbits
.global sin_table
sin_table:
    .byte 128,131,134,137,140,143,146,149 /* 0°-22.5° (first 8 of 256) */
    /* ... full table in src/lut_sin.S ... */

.section .text
.global sin8
sin8:
    ldi    ZL, lo8(sin_table)
    ldi    ZH, hi8(sin_table)
    add    ZL, r24                /* Z = sin_table + angle */
    adc    ZH, r1                 /* carry propagation */
    lpm    r24, Z                 /* r24 = table[angle] */
    ret
```

Cycle count: 1 + 1 + 1 + 1 + 3 + 5 = 12 cycles (including RCALL overhead from the caller). A polynomial sine computation on AVR takes 80–200 cycles.

15.4.4 Flash Address Warning: Devices with > 64 KB Flash

LPM accesses only the low 64 KB of flash. For devices with larger flash (ATmega2560 etc.), use ELPM with the RAMPZ register to address beyond 64 KB. The ATtiny3217 has 32 KB flash — LPM is sufficient.

15.4.5 Table Size and Flash Budget

```
256-entry uint8_t table = 256 bytes of flash
256-entry uint16_t table = 512 bytes of flash
256-entry uint32_t table = 1,024 bytes of flash
```

On a 32 KB device (ATtiny3217), a 256-byte table consumes 0.8% of flash. For a CRC-32 table (256 × 4 bytes = 1 KB), verify you have the budget before committing.

15.5 Putting It Together: Cycle Comparison

Build the benchmark and inspect it:

```
avr-gcc -mmcu=attiny3217 -O2 -g \
    -o bench.elf bench_main.c memcpy_naive.S memcpy_fast.S lut_sin.S

# Symbol sizes – flash consumed by each routine
avr-nm --size-sort --print-size bench.elf | grep -E 'memcpy|sin8|sin_table'
```

```
# Disassembly of all three
avr-objdump -d bench.elf
```

Expected avr-nm output (approximate):

```
00000100 00000010 T sin8          ; 16 bytes of code
00000110 00000100 T block_sum     ; ...
00000124 00000018 T memcpy_naive  ; 24 bytes
0000013c 00000034 T memcpy_fast   ; 52 bytes (unrolled costs code size)
00008200 00000100 T sin_table     ; 256 bytes in flash (PROGMEM section)
```

Summary comparison for len = 16 copy:

Routine	Cycles (16 bytes)	Code size
memcpy_naive	~122	24 bytes
memcpy_fast	~88	52 bytes
Savings	~28%	-28 bytes

And for sin8 vs equivalent polynomial:

Routine	Cycles	Code size
sin8 (LUT)	12	16 bytes code + 256 bytes table
sin_poly (est.)	~150	~80 bytes code

The LUT uses substantially more flash overall but runs about 12.5× faster.

That trade is often excellent on AVR, because flash is usually less scarce than CPU time in code that runs repeatedly.

15.6 Using avr-objdump Effectively

15.6.1 Key Flags

```
# Full disassembly with source interleaving (requires -g at compile time)
avr-objdump -dS bench.elf

# Show only the .text section (skip PROGMEM tables in disassembly)
avr-objdump -j .text -d bench.elf

# Dump PROGMEM section to verify table contents
avr-objdump -j .progmem.data -s bench.elf
```

```
# Symbol sizes – find largest functions / tables
avr-nm --size-sort --print-size bench.elf

# Segment sizes: text (flash), data (SRAM init), bss (SRAM zero)
avr-size bench.elf
```

15.6.2 Reading the Disassembly

```
00000142 <memcpy_naive>:
142:  fc 01          movw    r30, r24
144:  b7 01          movw    r22, r22
146:  2d 23          and     r18, r29
...
↑           ↑           ↑
└──────────┴──────────┘ mnemonic + operands (human readable)
└────────────────────────┘ 16-bit instruction word (little-endian hex)
└────────────────────────┘ byte address in flash (always even)
```

Addresses are byte addresses but instructions are 16-bit words, so consecutive instructions differ by 2 (or 4 for 32-bit instructions: LDS, STS, CALL, JMP).

15.6.3 Computing Total Cycle Count for a Path

Trace the execution path through the disassembly, note each instruction's cycle count from the reference table in Section 15.1, and sum. For loops, multiply by the iteration count. This is the most reliable way to verify your optimisation — far more trustworthy than timing with a stopwatch on real hardware where external effects dominate.

You do not need perfect lab instrumentation to reason about short AVR routines. For deterministic straight-line code and tight loops, the instruction timings already tell you almost everything you need.

15.7 Common Pitfalls

15.7.1 Pitfall 1 — Optimising the Wrong Loop

Profile first. In most firmware, 90% of CPU time is spent in < 10% of the code. Spending a week hand-optimising a function that runs once at startup is wasted. Use Timer0 or a logic analyser to find the hot path.

The point is not to avoid optimisation. It is to aim it where the CPU actually spends its time.

15.7.2 Pitfall 2 — Unrolling Beyond the I-Cache

Classic AVR has no instruction cache — flash access is 1 cycle per word with no caching. Excessive unrolling doesn't hurt due to cache misses, but it does consume flash you may need for other code. Stay proportionate.

15.7.3 Pitfall 3 — LPM Pointer Arithmetic Wrapping

If `sin_table` starts at flash address `0xFF00` and your angle index is `0x80`, then `ZL = 0x00 + 0x80 = 0x80`, carry = 0 — correct. But if `sin_table` is at `0xFF80` and index is `0x80`, `ZL = 0x80 + 0x80 = 0x00`, carry = 1, `ZH` increments — still correct. The ADC `ZH`, `R1` carry propagation handles this. What it does **not** handle: if `ZH` overflows past `0xFF`, you wrap into flash page 0. Ensure your table fits below the 64 KB boundary.

That warning does not matter on the ATtiny3217, but it matters if you carry the same technique to larger AVR parts later.

15.7.4 Pitfall 4 — Counting Cycles for a Branch-Taken vs Not-Taken

Conditional branches cost **1 cycle when not taken, 2 cycles when taken**. For a loop that runs N times, the branch is taken $N-1$ times and falls through once:

$$\text{Total branch cycles} = (N-1) \times 2 + 1 \times 1 = 2N - 1$$

Forgetting this causes off-by-one errors in cycle budgets.

15.8 Summary

Technique	When to use	Typical gain
Cycle counting	Always – build a baseline	(measurement tool)
Avoid SRAM	Hot variables in tight loops	2-3 cycles/access
Loop unrolling	High-iteration loops, loop body < 4 instructions	30-50% vs naive
LUT in flash	Fixed-domain functions with expensive computation	10-100× vs compute

Workflow:

1. `avr-gcc -O2` (let compiler do its job first)
2. `avr-objdump -dS` → identify inner loops
3. Count cycles per iteration by hand
4. Apply one technique; re-count
5. `avr-size` → verify flash budget not exceeded

Avoid premature optimisation. Measure, change one thing, measure again. AVR assembly optimisation is a precise discipline — cycles are countable, flash is measurable, and guesswork is unnecessary.

When you keep the process that disciplined, optimisation stops being mystical. It becomes a sequence of small engineering decisions with visible results.

15.9 Exercises

1. Count the exact cycle cost of the `memcpy_naive` inner loop for `len = 8`. Include the prologue (`MOVW, TST, BREQ`) and epilogue (`RET`). Compare to `memcpy_fast` for the same length.
2. Modify `memcpy_fast` to unroll $\times 8$ instead of $\times 4$. What is the new cycle count per byte in the inner loop? What is the new code size? At what value of `len` does the $\times 8$ version break even against $\times 4$?
3. Write a lookup table routine `gamma8(uint8_t x)` that returns the gamma-2.2 corrected value of x (compute the 256-entry table offline with $\text{round}(255 \times (n/255)^{(1/2.2)})$ for $n = 0..255$, embed as `.byte` directives). How many flash bytes does it consume? How does cycle count compare to a floating-point equivalent?
4. The `sin8` routine uses `ADC ZH, r1` to propagate carry. Why is `r1` always safe to use as a zero source here? What would break if you used `ADC ZH, r18` instead (assume `r18` holds an arbitrary value)?

5. Use `avr-objdump -j .progmem.data -s bench.elf` to dump the `sin_table` section. Verify the first and last bytes match the expected values (128 and 124). Explain why the last entry is 124 rather than 128.
6. A firmware routine is spending 70% of CPU time in a 6-instruction loop that runs 200 times. The loop body is:

```
ld    r18, Z+      ; 2
add   r24, r18     ; 1
adc   r25, r1      ; 1
dec   r20          ; 1
brne  .Lloop       ; 2 (taken)
```

Calculate the total cycle count. Now unroll $\times 2$: write the unrolled loop body, calculate the new cycle count for 200 iterations, and compute the speedup.

Next: Chapter 16 — Bootloaders

Chapter 16

Self-Programming Flash: NVMCTRL on tinyAVR 1-Series

Classic AVR devices (ATmega328P and friends) use a dedicated boot section and the SPM instruction to write flash. The **ATtiny3217**, like other modern tinyAVR 1-series parts, takes a different approach: the SPM instruction does not exist, there is no boot section, and there is no BOOTRST fuse. Instead, a memory-mapped peripheral called **NVMCTRL** (Non-Volatile Memory Controller) handles all flash writes from anywhere in the application.

This chapter covers everything needed to write assembly code that programs its own flash on the ATtiny3217 Curiosity Nano board:

1. Why AVRXMEGA3 replaced SPM with NVMCTRL.
2. How NVMCTRL registers and commands work.
3. The page buffer and the exact write sequence.
4. Configuration Change Protection (CCP) — the timing-critical unlock.
5. A complete 64-byte page-write routine in assembly.
6. A UART-driven self-update loop that receives Intel HEX and writes it to flash live.
7. The BOOTEND fuse: carving out a protected region for the updater itself.
8. Building, uploading, and common pitfalls.

This is one of the more hardware-specific chapters in the book. The goal is to turn “the chip can rewrite its own flash” from a mysterious feature into a sequence you can reason about step by step.

16.1 tinyAVR 1-Series vs Classic AVR: No SPM, NVMCTRL Instead

16.1.1 Architecture Family: AVRXMEGA3

The ATtiny3217 belongs to the **AVRXMEGA3** sub-architecture. The three key differences that affect self-programming are:

Feature	ATmega328P (classic)	ATtiny3217 (AVRXMEGA3)
Self-program instr.	SPM (dedicated opcode)	none – use ST to flash
Boot section	yes (BOOTSZ fuses)	none
BOOTRST fuse	yes	none
Flash write ctrl.	SPMCSR register	NVMCTRL peripheral
Flash address space	separate word-address	byte-address, unified

Page size	128 bytes (64 words)	64 bytes (32 words)
Write from app?	boot section only	anywhere (subject to BOOTEND fuse)
UPDI programming	no	yes (PA0, replaces ISP)

The AVRXMEGA3 memory map is **unified**: flash, SRAM, and peripherals all share the same byte-addressed data space. Flash byte addresses start at `0x8000` when accessed as data, but the CPU's program counter addresses flash from `0x0000`. You will see both conventions in this chapter.

16.1.2 What Replaces SPM?

On tinyAVR 1-series, you fill the flash **page buffer** by writing normally to flash addresses using ST (Store Indirect). There is no special opcode. The NVMCTRL peripheral watches these writes and accumulates them in a 64-byte buffer. When you issue a command to NVMCTRL_CTRLA, it commits the buffer to the actual flash array.

Classic AVR SPM flow:	tinyAVR NVMCTRL flow:
load R1:R0 ← data word	(nothing special)
write SPMCSR ← mode	wait FBUSY = 0
SPM ← fills buffer	ST Z+, rn ← fills buffer (just write!)
(repeat for each word)	(repeat for each word)
write SPMCSR ← PGMCR	CCP unlock (0x9D → 0x0034)
SPM ← erases page	STS NVMCTRL_CTRLA ← 0x03
write SPMCSR ← PGWRT	wait FBUSY = 0
SPM ← writes page	

The NVMCTRL approach is simpler to understand: write your data to the target flash addresses, then send the “erase+write” command. The hardware does the rest.

The important mental shift is that self-programming no longer looks like a special instruction stream. It looks like ordinary memory-mapped peripheral control wrapped around a careful write sequence.

16.1.3 Programming Interface: UPDI

The ATtiny3217 has no ISP pins. All external programming is done over **UPDI** (Unified Program and Debug Interface), a single-wire protocol on **PA0**. The Curiosity Nano board's on-board debugger bridges USB to UPDI automatically. Use `pymcuprog` for the on-board debugger:

```
pymcuprog write -d attiny3217 -t uart -f firmware.hex
```

Once the self-update firmware is running, the device reprograms itself — no external programmer needed for firmware updates.

16.2 NVMCTRL Overview

16.2.1 Flash Parameters

ATtiny3217 flash layout

Total flash:	32 KB (32768 bytes)
Page size:	64 bytes
Number of pages:	512
Flash data space:	0x8000-0xFFFF (byte addresses when read as data)
Flash program:	0x0000-0x7FFF (word addresses for PC / LPM)

SRAM: 0x3800-0x3FFF (2 KB)

The 64-byte page size is critical: every NVMCTRL write operation works on exactly one 64-byte page. Partial writes within a page are allowed (the page buffer is pre-loaded from flash before programming), but the hardware always erases and rewrites the full 64 bytes.

That last point matters because it explains why page alignment and buffering show up everywhere in self-programming code. Flash is not rewritten one byte at a time, even if your firmware update logic receives bytes one at a time.

16.2.2 NVMCTRL Register Map

All NVMCTRL registers are outside the I/O space and require LDS/STS:

Register	Address	Width	Purpose
NVMCTRL_CTRLA	0x1000	1 byte	Command register – write command here
NVMCTRL_CTRLB	0x1001	1 byte	Config: BOOTLOCK, APCWP write-protect
NVMCTRL_STATUS	0x1002	1 byte	Status flags: FBUSY, EEBUSY, WRERROR
NVMCTRL_INTCTRL	0x1003	1 byte	Interrupt enable: EEREADY (bit 0)
NVMCTRL_INTFLAGS	0x1004	1 byte	Interrupt flags (write 1 to clear)
NVMCTRL_DATA	0x1006	2 bytes	Data register (used for EEPROM word)
NVMCTRL_ADDR	0x1008	4 bytes	Address register (used for EEPROM)

For flash programming, you only need NVMCTRL_CTRLA and NVMCTRL_STATUS.

16.2.3 NVMCTRL_STATUS — Status Register (0x1002)

Bit:	7	6	5	4	3	2	1	0
	–	–	–	–	–	WRERROR	EEBUSY	FBUSY

FBUSY – Flash Busy: 1 while a flash erase/write operation is in progress.

Poll this bit; do not issue a new command while FBUSY = 1.

EEBUSY – EEPROM Busy: 1 while an EEPROM operation is in progress.

WRERROR – Write Error: set if a write was attempted to a protected region.

Write 1 to clear (W1C).

16.2.4 NVMCTRL Commands (written to NVMCTRL_CTRLA)

Value	Name	Description
0x00	NOCMD	No operation (also clears WRERROR)
0x01	PAGEWRITE	Write page buffer to flash (no erase first)
0x02	PAGEERASE	Erase one flash page (Z must point into page)
0x03	PAGEERASEWRITE	Erase + write in one atomic step – use this
0x04	PAGEBUFERASE	Clear the page buffer (all bytes → 0xFF)
0x05	CHIPERASE	Erase entire flash (except BOOTEND region)
0x06	EEERASE	Erase EEPROM
0x07	FUSEWRITE	Write fuse bytes (restricted, rarely needed)

PAGEERASEWRITE (0x03) is the command you will use for firmware updates. It atomically erases the target page and writes the page buffer in one operation — safer than doing them separately, because a power failure between separate erase and write steps cannot leave a page half-erased.

In practice, that combined command is what makes the sequence robust enough to use in the field. Fewer separate steps means fewer partially completed states.

16.3 The Page Buffer and Write Sequence

16.3.1 How the Page Buffer Works

The NVMCTRL page buffer is a 64-byte staging area internal to the NVM controller. It is **not** directly addressable — you populate it by writing to the corresponding flash addresses in the data address space ($0x8000 + \text{flash_byte_address}$).

When the Z pointer points into flash address space and you execute `ST Z+, Rn`, the byte goes into the page buffer at the offset corresponding to Z's address. The hardware handles this automatically whenever the NVM controller recognises the write target as flash.

Page buffer write – step by step:

Flash page at byte address `0x0000` appears in data space at `0x8000`.

```
Z = 0x8000      ; points to byte 0 of flash page 0
st Z+, r16      ; → page buffer[0] = r16
st Z+, r17      ; → page buffer[1] = r17
...             ; fill all 64 bytes
st Z+, r30      ; → page buffer[63] = r30
                ; Z = 0x8040 (past end of page)
```

The buffer retains its contents until you issue a command or an explicit `PAGEBUFFERASE`. If you do not fill all 64 bytes, the remaining bytes are pre-loaded from the existing flash content during the `PAGEERASEWRITE` command — so partial page updates work correctly.

That behavior is what makes incremental updates possible. Your code does not have to rebuild the untouched bytes of the page by hand before every write.

16.3.2 Complete Flash Page Write Sequence

Step 1: Wait for any previous NVM operation to finish.
Poll `NVMCTRL_STATUS.FBUSY` until it reads 0.

Step 2: Issue `NOCMD` to clear any stale error state.
`STS NVMCTRL_CTRLA, 0x00`

Step 3: Fill the page buffer.
Set `Z = 0x8000 + (target_page × 64)` ; data-space address
`ST Z+, r0 ; ST Z+, r1 ; ...` ; 64 sequential writes

Step 4: Issue `PAGEERASEWRITE` (CCP-protected).
`LDI r16, 0x9D` ; CCP_SPM_gc key
`STS 0x0034, r16` ; write to CCP register – starts 4-cycle window
`LDI r16, 0x03` ; `PAGEERASEWRITE` command
`STS NVMCTRL_CTRLA, r16` ; MUST be within 4 cycles of CCP write

Step 5: Wait for `FBUSY` to clear again.
Poll `NVMCTRL_STATUS.FBUSY` until 0.

16.3.3 The Z Pointer Convention

On ATtiny3217 the Z pointer (R31:R30) is a 16-bit register. Flash page 0 lives at data address `0x8000`, which is beyond the 16-bit range... but only if you use Z as a pure 16-bit address. The AVRXMEGA3 architecture provides the **RAMPZ** register (at I/O address `0x3B`) to extend Z to 24 bits for `ELPM` and `LPM` instructions.

However, for ST writes to flash in the page buffer fill loop, the hardware uses only the lower 16 bits of Z relative to the page boundary — the NVMCTRL page buffer offset is determined by `Z & 0x3F` (the 6 low bits).

In practice, the simplest approach is to set Z to the flash **byte address** (not the data-space address) and let the hardware resolve which page buffer slot to populate:

```
/* To write page at flash byte address 0x0200 (page 8): */
ldi ZL, lo8(0x0200)
ldi ZH, hi8(0x0200)
/* Now ST Z+, Rn writes to page buffer slots starting at offset 0 */
```

The NVMCTRL recognises the write targets and routes them to the page buffer automatically. The examples in Sections 16.5 and 16.6 use this convention.

That convention keeps the assembly easier to follow because the same flash byte address can be used consistently through buffering, writing, and later jumps.

16.4 CCP — Configuration Change Protection

16.4.1 Why CCP Exists

Some NVMCTRL commands (and other sensitive operations like changing clock settings) can corrupt firmware if executed accidentally — for example, if a runaway pointer writes to NVMCTRL_CTRLA, or if a flash write executes during a brownout. **CCP** (Configuration Change Protection) prevents accidental execution by requiring a timed unlock sequence before each protected write.

So CCP is not there to make your life difficult. It is there because an accidental flash-control write is serious enough that the hardware demands a deliberate two-step action.

16.4.2 The CCP Register (0x0034)

Address: 0x0034 (in I/O space: use OUT or STS)

CCP keys:

```
0x9D - CCP_SPM_gc    Unlocks NVM Self-Programming commands
0xD8 - CCP_IOREG_gc  Unlocks protected I/O register writes (e.g. CLKCTRL)
```

After writing a key, there is a 4-instruction-cycle window in which the protected write must occur. The window closes automatically after 4 cycles, regardless of whether the write happened. Any interrupt that fires during the window resets the protection.

16.4.3 The 4-Cycle Rule

The CCP window opens the moment you write the key and closes 4 clock cycles later. You must write the protected register (in this case NVMCTRL_CTRLA) within those 4 cycles. STS takes 2 cycles; two back-to-back STS instructions consume exactly 4 cycles — the window is just large enough:

Cycle:	1	2	3	4	
	STS CCP		STS CTRL		
	writes key		writes cmd		← must complete within 4 cycles
			window closes		

Assembly idiom:

```
ldi    r16, 0x03          /* PAGEERASEWRITE command      */
ldi    r17, 0x9D          /* CCP_SPM_gc key          */
sts    0x0034, r17        /* open the CCP window (cycle 1-2) */
sts    NVMCTRL_CTRLA, r16 /* protected write (cycle 3-4)    */
```

Critical: do not put any instruction between STS CCP and STS NVMCTRL_CTRLA — not even a NOP. Interrupts must be disabled globally (CLI) before this sequence to prevent an ISR from consuming cycles inside the window.

This is one of the rare cases where instruction spacing is not just an optimization detail. It is part of the functional correctness of the code.

16.4.4 C Equivalent

In C, do not hand-write `CPU.CCP = key; reg = value;` unless you are also guaranteeing the instruction spacing yourself. The recommended `avr-libc` helpers emit the protected sequence in inline assembly so the compiler cannot insert extra instructions into the 4-cycle window. For I/O-protected registers use `ccp_write_io()` or `_PROTECTED_WRITE(...)`; for NVM/SPM-protected writes use `_PROTECTED_WRITE_SPM(...)`.

The `avr-libc` macro `_PROTECTED_WRITE(reg, val)` generates this exact sequence:

```
/* C equivalent of the CCP unlock + write: */
_PROTECTED_WRITE(NVMCTRL_CTRLA, NVMCTRL_CMD_PAGEERASEWRITE_gc);

/* Expands roughly to (inline asm): */
__asm__ volatile(
    "out %[ccp], %[key]" "\n\t"
    "sts %[reg], %[val]"
    : : [ccp] "I" (_SFR_IO_ADDR(CCP)),
      [key] "r" ((uint8_t)CCP_SPM_gc),
      [reg] "n" (NVMCTRL_CTRLA_address),
      [val] "r" ((uint8_t)NVMCTRL_CMD_PAGEERASEWRITE_gc)
    );
```

In pure assembly you write it directly, as shown in the code listings below.

That is also why the `avr-libc` helpers are worth using from C: they protect the timing-sensitive part from compiler reordering and from innocent-looking source changes.

16.5 Flash Write Routine in Assembly

The routine `nvm_write_page` writes one 64-byte page. It expects:

- **Z** (R31:R30): flash byte address of the target page (must be 64-byte aligned).
- **X** (R27:R26): data address of the 64-byte source buffer in SRAM.

It clobbers R16, R17, R18 (scratch) and leaves Z past the end of the written page. The source buffer pointer X is incremented and left past the last byte.

Leaving both pointers advanced is a small design choice, but it makes the routine easier to chain when writing consecutive pages from a stream.

16.5.1 Register Allocation

Z (R31:R30) flash destination – page base address, then advanced
 X (R27:R26) SRAM source pointer – 64-byte buffer of new content
 R16 scratch: NVMCTRL command, status poll byte
 R17 loop counter (64 iterations)
 R18 data byte being copied into page buffer

16.5.2 Full Listing: `src/nvmctrl_write.S`

See `src/nvmctrl_write.S` for the complete annotated source. Key routines:

```
/*
 * nvm_wait – spin until NVMCTRL_STATUS.FBUSY = 0
 *
 * Clobbers: R16
 * Cycles per loop: LDS (2) + SBRC (1 or 2) + RJMP (2) = 5 cycles per poll
 */
nvm_wait:
    lds    r16, NVMCTRL_STATUS          /* 2 cy: load status          */
    sbrc   r16, NVMCTRL_FBUSY_bp        /* 1 cy: skip next if FBUSY=0 */
    rjmp   nvm_wait                    /* 2 cy: still busy, loop     */
    ret                                     /* 4 cy                      */
```

SBRC Rd, b — Skip if Bit in Register is Cleared. Takes 1 cycle when not skipping, 2 when skipping (pipeline flush). NVMCTRL_FBUSY_bp = 0 (bit position 0 of NVMCTRL_STATUS).

```
/*
 * nvm_write_page – write 64-byte SRAM buffer to one flash page
 *
 * Entry: Z = flash byte address (64-byte aligned)
 *        X = SRAM pointer to 64-byte source buffer
 * Exit: Z advanced by 64, X advanced by 64
 * Clobbers: R16, R17, R18
 */
nvm_write_page:
    /* Step 1: wait for any previous NVM operation */
    rcall nvm_wait                /* 4 + 5n cy (n polls)      */

    /* Step 2: clear previous error (NOCMD) */
    ldi   r16, NVMCTRL_CMD_NOCMD_gc /* r16 = 0x00              */
    sts   NVMCTRL_CTRLA, r16        /* 2 cy                     */

    /* Step 3: fill page buffer from SRAM source */
    ldi   r17, 64                  /* loop counter = page size */
.Lfill_loop:
    ld     r18, X+                  /* 2 cy: load byte from SRAM */
    st     Z+, r18                  /* 2 cy: write byte to flash page buf */
    dec    r17                      /* 1 cy: decrement counter    */
    brne   .Lfill_loop             /* 2 cy taken / 1 cy fall-through */
    /* Total fill: 63 × 7 + 6 = 447 cy */

    /* Step 4: rewind Z to page start for PAGEERASEWRITE */
    sbiwr ZL, 64                   /* 2 cy: Z -= 64            */

    /* Step 5: CCP-protected PAGEERASEWRITE command */
```

```

cli                /* 1 cy: disable interrupts          */
ldi   r16, NVMCTRL_CMD_PAGEERASEWRITE_gc /* r16 = 0x03 */
ldi   r17, CCP_SPM_gc /* r17 = 0x9D */
sts   CPU_CCP, r17 /* 2 cy: unlock CCP window (cycles 1-2) */
sts   NVMCTRL_CTRLA, r16 /* 2 cy: issue command (cycles 3-4) */
sei                /* 1 cy: re-enable interrupts        */

/* Step 6: wait for completion */
rcall nvm_wait

/* Z points to page start; advance past the page for caller convenience */
adiw  ZL, 64 /* 2 cy */
ret

```

16.5.3 Instruction Notes

Two instructions appear for the first time in this chapter:

SBIW Rd, K – Subtract Immediate from Word
 Operation: $Rd+1:Rd \leftarrow Rd+1:Rd - K$ ($K = 0..63$)
 Encoding: 1001 0111 KKdd KKKK (2 bytes)
 Registers: $d \in \{24, 26, 28, 30\}$ only ($R25:R24, X, Y, Z$)
 Flags: C, Z, N, V, S
 Cycles: 2

ADIW Rd, K – Add Immediate to Word
 Operation: $Rd+1:Rd \leftarrow Rd+1:Rd + K$ ($K = 0..63$)
 Encoding: 1001 0110 KKdd KKKK (2 bytes)
 Registers: $d \in \{24, 26, 28, 30\}$ only
 Flags: C, Z, N, V, S
 Cycles: 2

SBIW ZL, 64 rewinds Z to the start of the page after the fill loop advanced it by 64. **ADIW ZL, 64** advances Z to the start of the next page at the end of the routine. Both execute in 2 cycles and operate on the 16-bit Z pair atomically — no carry propagation needed.

That is exactly the sort of detail assembly lets you express cleanly: a page is 64 bytes, so subtract 64 once as a word operation instead of manually fixing up two separate bytes.

16.5.4 Side-by-Side C Equivalent

```

/* C equivalent of nvm_write_page() */
void nvm_write_page(uint16_t flash_addr, const uint8_t *src) {
    uint8_t volatile *dst = (uint8_t volatile *) (0x8000 + flash_addr);

    /* Wait for NVM ready */
    while (NVMCTRL.STATUS & NVMCTRL_FBUSY_bm) {}

    /* Clear error */
    NVMCTRL_CTRLA = NVMCTRL_CMD_NOCMD_gc;

    /* Fill page buffer */
    for (uint8_t i = 0; i < 64; i++)
        dst[i] = src[i]; /* plain write → page buffer */
}

```

```

/* Issue erase+write (CCP-protected) */
_PROTECTED_WRITE(NVMCTRL.CTRLA, NVMCTRL_CMD_PAGEERASEWRITE_gc);

/* Wait for completion */
while (NVMCTRL.STATUS & NVMCTRL_FBUSY_bm) {}
}

```

16.6 Practical Firmware Update: UART-Driven Self-Update

With `nvm_write_page` in hand, we can build a self-update loop. The device receives new firmware as Intel HEX records over `USART0`, accumulates complete 64-byte pages in SRAM, then writes each page to flash.

At that point the problem stops being “how do I write flash?” and becomes a system problem: receive bytes reliably, group them into pages, commit them, and avoid damaging the updater itself.

16.6.1 Intel HEX Recap

```

:LLAAAATT[DD...]CC
| | | | |
| | | | | checksum (2 hex ASCII chars)
| | | | | data bytes (LL × 2 hex chars)
| | | | | record type: 00=data, 01=EOF
| | | | | address (4 hex ASCII chars, big-endian)
| | | | | byte count LL (2 hex ASCII chars)
↑ start code ':'

```

Checksum: two's complement of $(LL + \text{addr_hi} + \text{addr_lo} + \text{type} + \text{sum}(\text{data}))$
 The sum of all bytes in the record including checksum = $0x00 \pmod{256}$.

16.6.2 Architecture of `selfupdate.S`

```

selfupdate_main:
|
|— init: stack, USART0 (PB2 TX, PB3 RX, 9600 baud at 3.333 MHz)
|— LED0 (PA3) on → indicates update mode active
|— send banner "SU\r\n"
|
|— record_loop:
|   |— wait for ':' (discards noise before start code)
|   |— recv_hex_byte × 4: byte_count, addr_hi, addr_lo, type
|   |— if type = 0x01 (EOF): send "OK\r\n", wait FBUSY, IJMP to 0x0000
|   |— if type = 0x00 (data):
|       |— set Z = record address (flash byte address)
|       |— receive LL data bytes → SRAM page buffer
|       |— if buffer ≥ 64 bytes: call nvm_write_page
|       |— verify checksum → send '.' (ACK) or '!' (NACK)
|   — repeat

```

16.6.3 Register Allocation for `selfupdate.S`

```

Z (R31:R30)  flash write address (maintained across records)
X (R27:R26)  SRAM page buffer pointer (buf_base .. buf_base+63)
R16          scratch / UART byte / NVMCTRL command

```

R17 checksum accumulator (zeroed at each record start)
 R18 byte count from record header
 R19 record type (00 or 01)
 R24 address low byte from record header
 R25 address high byte from record header
 R20, R21 scratch in helpers
 R15 saved high nibble in recv_hex_byte

16.6.4 USART0 Initialisation at 3.333 MHz

The ATtiny3217's default clock is 3.333 MHz (20 MHz internal oscillator with the default /6 prescaler in CLKCTRL). The USART0 fractional baud register for 9600 baud:

$$\begin{aligned}
 \text{BAUD_REG} &= (\text{f_cpu} \times 64) / (16 \times \text{baud}) \\
 &= (3333333 \times 64) / (16 \times 9600) \\
 &= 213333312 / 153600 \\
 &= 1389 \quad (0x056D)
 \end{aligned}$$

```

.equ BAUD_REG, 1389                                /* 9600 baud @ 3.333 MHz, error < 0.01% */

usart_init:
    sbi   VPORTB_DIR, 2                            /* PB2 = TXD output */
    ldi   r16, lo8(BAUD_REG)
    sts   USART0_BAUDL, r16
    ldi   r16, hi8(BAUD_REG)
    sts   USART0_BAUDH, r16
    ldi   r16, USART_TXEN_bm | USART_RXEN_bm
    sts   USART0_CTRLB, r16                        /* enable TX and RX */
    ret
    /* CTRLC defaults to 8N1 – no write needed */
  
```

Note: VPORTB_DIR is at I/O address 0x04 (within sbi/cbi range) — we use sbi rather than sts to set the single bit without disturbing other PORTB direction bits.

16.6.5 Full Listing: src/selfupdate.S

The complete source is at src/selfupdate.S. The core hex-receive loop and page-write call are shown below for reference:

```

.Ldata_loop:
    rcall recv_hex_byte                            /* receive next data byte from HEX record */
    st    X+, r16                                  /* store in SRAM page buffer */
    add   r17, r16                                  /* accumulate checksum */
    dec   r18                                       /* decrement byte count */

    /* check if page buffer is full (64 bytes written) */
    mov   r20, XL
    andi  r20, 0x3F                                /* XL & 63 – non-zero means not at boundary */
    brne  .Lcheck_done                             /* not at 64-byte boundary: continue */

    /* page buffer is full: save X, rewind Z 64 bytes, write page, restore */
    movw  r20, XL                                  /* save X */
    sbiw  ZL, 64                                    /* Z back to page start */
    rcall nvm_write_page                           /* writes page, advances Z by 64 */
    movw  XL, r20                                  /* restore X */
  
```

```
.Lcheck_done:
    tst    r18
    brne   .Ldata_loop           /* more bytes in this record      */
```

16.6.6 Jumping to the Updated Application

After the EOF record is received and the last page is written:

```
.Leof:
    rcall  recv_hex_byte          /* consume EOF checksum byte (not verified) */
    rcall  nvm_wait               /* ensure final write is complete           */
    /* turn off LED0 (PA3) to signal update complete */
    sbi    VPORTA_OUT, 3          /* PA3 high = LED off (active-low)          */
    /* send "OK\r\n" */
    ldi    r16, 'O'
    rcall  usart_tx
    ldi    r16, 'K'
    rcall  usart_tx
    ldi    r16, '\r'
    rcall  usart_tx
    ldi    r16, '\n'
    rcall  usart_tx
    /* jump to new application */
    clr    ZL
    clr    ZH
    ijmp                                /* PC ← Z = 0x0000 (new firmware entry)    */
```

IJMP (Indirect Jump): sets the program counter to Z (2 cycles). With Z = 0, execution continues from the reset vector of the freshly written firmware.

Conceptually, that final jump is just a handoff. The updater finishes its work, then transfers control to the program it just installed.

16.6.7 Uploading Firmware

Use the existing `src/send_hex.py` helper, or any serial terminal. The selfupdate firmware sends `SU\r\n` instead of `BL\r\n`:

```
# Flash the self-update firmware via the Curiosity Nano debugger (one-time):
pymcuprog write -d attiny3217 -t uart -f selfupdate.hex
```

```
# Reset the board – LED0 turns on (active-low, so PA3 = 0 = LED on)
# Then immediately send application firmware:
python3 src/send_hex.py /dev/ttyUSB0 app.hex
```

Expected output:

```
Waiting for bootloader on /dev/ttyUSB0...
Bootloader ready
  record 1/18 OK
  record 2/18 OK
  ...
  record 18/18 OK
Upload complete. Application starting.
```

LED0 turns off when the update completes, and the new firmware starts.

16.7 BOOTEND Fuse: Protecting a Bootloader Region

16.7.1 The Problem

The self-update firmware in Section 16.6 writes to any flash address, including address 0x0000 — which is where the self-update firmware itself lives. If the update process is interrupted after erasing page 0 but before writing the new content, the device is bricked (no valid firmware at the reset vector). It also means buggy application firmware could accidentally overwrite the updater.

16.7.2 BOOTEND: Defining a Protected Region

The ATtiny3217 provides the **BOOTEND** fuse (in the FUSES.BOOTEND byte). It divides flash into a protected **boot section** and an unprotected **application section**:

FUSES.BOOTEND value:

- 0x00 — no boot section (entire flash is application section)
- 0x01 — 256 bytes protected (pages 0-3)
- 0x02 — 512 bytes protected (pages 0-7)
- 0x04 — 1024 bytes protected (pages 0-15)
- ...
- 0xFF — entire flash protected

Boot section: 0x0000 .. (BOOTEND × 256 - 1)


Application section: (BOOTEND × 256) .. 0x7FFF

With BOOTEND = 0x04 (1 KB boot section), the self-update firmware occupies pages 0-15 (addresses 0x0000-0x03FF). The application firmware lives at 0x0400 onward. The NVMCTRL enforces this boundary: writing to the boot section from application code triggers WRERROR.

This is the safety boundary that turns a self-programming demo into something you can trust more. Without it, the updater can erase the very code that is needed to recover from a failed update.

16.7.3 Flash Layout With BOOTEND = 0x04

ATtiny3217 flash (32 KB) with BOOTEND = 0x04



16.7.4 Application Entry Point

When BOOTEND = 0x04, the application reset vector is at 0x0400. The self-update firmware must jump there (not to 0x0000) after a successful update:


```

/* With BOOTEND = 0x04: application starts at 0x0400 */
.equ APP_START, 0x0400

.Leof:
    rcall recv_hex_byte
    rcall nvm_wait
    ldi    ZL, lo8(APP_START)
    ldi    ZH, hi8(APP_START)
    ijmp                      /* jump to application section start */

```

The application must be built with its `.text` section starting at `0x0400`:

```

avr-gcc -mmcu=attiny3217 -nostartfiles \
        -Wl,--section-start=.text=0x0400 \
        -o app.elf app.S

```

16.7.5 BOOTEND Lock (APCWP)

NVMCTRL_CTRLB.APCWP (Application Code Write Protect, bit 0) is a run-time lock that prevents the **application** section from being written even by the boot code. Setting this bit after the update completes provides defence-in-depth:

```

/* After successful update, lock application section against further writes */
ldi    r16, 0x01          /* APCWP bit */
sts    NVMCTRL_CTRLB, r16

```

This is cleared on every reset, so it does not permanently lock the device.

Think of APCWP as a temporary guardrail for the current boot session, not as a replacement for fuse-based protection.

16.8 Building and Flashing

16.8.1 Building nvmctrl_write.S (standalone test)

```

# Assemble the NVM write routine alone (produces an ELF for inspection)
avr-gcc -mmcu=attiny3217 -nostartfiles -nodefaultlibs \
        -o nvmctrl_write.elf src/nvmctrl_write.S

# Check sizes
avr-size nvmctrl_write.elf

# Disassemble to verify instruction encoding
avr-objdump -d nvmctrl_write.elf

```

16.8.2 Building selfupdate.S

```

# Assemble the self-update firmware
# Without BOOTEND protection: starts at 0x0000
avr-gcc -mmcu=attiny3217 -nostartfiles -nodefaultlibs \
        -o selfupdate.elf src/selfupdate.S

# With BOOTEND=0x04: starts at 0x0000, app section starts at 0x0400
# The selfupdate firmware itself is the same; linker placement is 0x0000

```

```
avr-gcc -mmcu=attiny3217 -nostartfiles -nodefaultlibs \
    -o selfupdate.elf src/selfupdate.S

# Size check – must fit within chosen boot section (1 KB = BOOTEND 0x04)
avr-size selfupdate.elf
```

Expected size output:

text	data	bss	dec	hex	filename
486	0	4	490	1ea	selfupdate.elf

486 bytes fits comfortably within the 1 KB (1024 byte) boot section.

16.8.3 Converting to HEX

```
avr-objcopy -O ihex selfupdate.elf selfupdate.hex
avr-objcopy -O ihex app.elf app.hex
```

16.8.4 Flashing with pymcuprog

```
# Program selfupdate firmware via UPDI (requires Curiosity Nano USB connection)
pymcuprog write -d attiny3217 -t uart -f selfupdate.hex
```

```
# Set BOOTEND fuse to 0x04 (1 KB boot section)
# WARNING: verify the exact pymcuprog fuse syntax for your installed version.
# BOOTEND is fuse byte 8; write value 0x04.
```

```
# Read back fuse to verify
pymcuprog read -d attiny3217 -t uart
```

Warning: BOOTEND is in fuse byte 8 on ATtiny3217. An incorrect value can make the device appear to run no code. Always read back fuses after writing. UPDI can always recover the device regardless of fuse settings (unlike ISP on ATmega parts where the BOOTRST fuse can complicate recovery).

That recovery path is important operationally: experimenting with self-update code is much less risky when you know UPDI can still get you back in.

16.8.5 Verifying the Disassembly

After building, verify the CCP sequence is encoded correctly:

```
avr-objdump -d selfupdate.elf | grep -A5 "CPU_CCP\|0034"
```

Look for two back-to-back sts instructions with addresses 0x0034 (CCP) and 0x1000 (NVMCTRL_CTRLA). No instruction should appear between them in the listing.

16.8.6 Building an Application for the Protected Layout

```
# Application must be linked at 0x0400 when BOOTEND = 0x04
avr-gcc -mmcu=attiny3217 -nostartfiles -nodefaultlibs \
    -Wl,--section-start=.text=0x0400 \
    -o app.elf app.S
avr-objcopy -O ihex app.elf app.hex
```

```
# Send via self-update firmware:
python3 src/send_hex.py /dev/ttyUSB0 app.hex
```

16.9 Common Pitfalls

16.9.1 Pitfall 1 — CCP Window Violated

Symptom: NVMCTRL_CTRLA write appears to succeed (no error) but no flash write occurs; WRERROR is set in STATUS.

Cause: An instruction was inserted between STS CPU_CCP, r16 and STS NVMCTRL_CTRLA, r16, or an interrupt fired inside the window and consumed cycles.

Fix: Use CLI before the CCP sequence and SEI immediately after the protected write. Ensure no instruction sits between the two STS calls.

```
/* WRONG – ldi between STS instructions wastes a cycle inside the window */
sts  CPU_CCP, r16
ldi  r16, 0x03          /* ← this is inside the window, shrinks margin */
sts  NVMCTRL_CTRLA, r16

/* CORRECT – load command value BEFORE writing CCP key */
ldi  r16, 0x03          /* load command first (outside window) */
...
ldi  r17, CCP_SPM_gc
sts  CPU_CCP, r17        /* open window */
sts  NVMCTRL_CTRLA, r16 /* issue command (within window) */
```

Note the subtlety: if r16 holds the CCP key when you write CPU_CCP, you must reload r16 with the command before the protected write, or use two registers. The corrected pattern above uses two registers to keep the command value ready before opening the CCP window.

That is a good example of why “almost correct” is not good enough here. A sequence can look visually right and still fail because one register was reused at the wrong moment.

16.9.2 Pitfall 2 — Page Buffer Not Aligned

Symptom: The first few bytes of a written page are corrupted or contain old flash content.

Cause: The flash address passed to `nvm_write_page` is not 64-byte aligned. The page buffer fill loop always fills from byte 0 to byte 63 of the chosen page. If Z points into the middle of a page, the first ST Z+, Rn writes to a mid-page buffer slot, and the slots before it are pre-loaded from existing flash content (which may be stale).

Fix: Always pass a 64-byte aligned address:

```
/* Align Z to 64-byte page boundary: clear bits 5:0 */
andi ZL, 0xC0          /* mask: 1100 0000 – clear bits 5:0 */
```

Page alignment bugs are especially deceptive because most of the write may look correct. Only the offsets near the page boundary reveal the mistake.

16.9.3 Pitfall 3 — Writing to Flash During Interrupt

Symptom: Application ISR fires while `nvm_write_page` is between the page buffer fill and the PAGEERASEWRITE command. The ISR reads stale flash content via LPM from the page being written.

Cause: The page buffer is not yet committed. If an ISR attempts to read instruction or data from flash in the middle of a PAGEERASEWRITE, it reads the erased (0xFF) state during the erase phase.

Fix: Disable global interrupts before the entire write sequence and re-enable after:

```
cli
rcall nvm_write_page
sei
```

The `nvm_write_page` routine already disables interrupts around the CCP write, but a complete CLI/SEI around the entire page write is safer if you have ISRs active.

16.9.4 Pitfall 4 — Forgetting to Wait for FBUSY Before Each Command

Symptom: Random corruption of flash pages, or `WRERROR` flag set frequently.

Cause: A second `PAGEERASEWRITE` command was issued before the first completed. `NVMCTRL_STATUS.FBUSY` must be 0 before issuing any command.

Fix: The `nvm_wait` routine at the start of `nvm_write_page` ensures this. If you issue `NVMCTRL` commands from multiple places in your code, always call `nvm_wait` first.

16.9.5 Pitfall 5 — Writing the Boot Section From Application Code

Symptom: Write to pages 0–(BOOTEND-based boundary) triggers `WRERROR` but no actual write.

Cause: With `BOOTEND` set, `NVMCTRL` enforces the boundary. Application code cannot write the boot section via `NVMCTRL`; only `UPDI` or code running in the boot section itself can.

Fix: This is intentional behaviour — the boot section protection is working correctly. If you need to update the self-update firmware, use `UPDI` with `pymcuprog`.

In other words, `WRERROR` here is a feature, not a malfunction.

16.9.6 Pitfall 6 — UPDI Port Floating (PA0)

Symptom: `avrdude` cannot connect; “timeout communicating with programmer.”

Cause: On the Curiosity Nano, PA0 is used for `UPDI`. If you configure PA0 as a GPIO output in your firmware, `UPDI` communication breaks. The on-board on-board debugger pulls PA0 to its own `UPDI` driver.

Fix: Never configure PA0 as GPIO output. Leave PA0 at its power-on default (`UPDI` function). If your application needs an extra GPIO and PA0 was used, remap to a different pin.

16.9.7 Pitfall 7 — Baud Rate at Non-Default Clock

Symptom: Garbage characters received; UART communication fails after changing `CLKCTRL` prescaler.

Cause: The `BAUD_REG` constant was calculated for 3.333 MHz (default clock). If your firmware changes `CLKCTRL_MCLKCTRLB` to select a different frequency before calling `usart_init`, the baud rate will be wrong.

Fix: Recalculate `BAUD_REG` for the actual `f_cpu`, or initialise `USART` before changing the clock. The self-update firmware deliberately runs at the default clock to avoid this dependency.

16.10 Summary

ATtiny3217 self-programming (AVRXMEGA3):

Flash geometry:

32 KB total, 64-byte pages, 512 pages
Data-space flash base address: 0x8000

Page buffer fill: ST Z+, Rn (plain store instruction)

NVMCTRL registers:

NVMCTRL_CTRLA 0x1000 command register
 NVMCTRL_STATUS 0x1002 FBUSY (bit 0), WRERROR (bit 2)

Write sequence for one 64-byte page:

1. Poll NVMCTRL_STATUS until FBUSY = 0
2. STS NVMCTRL_CTRLA, 0x00 (NOCMD – clear errors)
3. 64 × ST Z+, Rn (fill page buffer)
4. SBIW ZL, 64 (rewind Z to page base)
5. CLI
 STS CPU_CCP, 0x9D (unlock CCP)
 STS NVMCTRL_CTRLA, 0x03 (PAGEERASEWRITE – within 4 cycles)
 SEI
6. Poll NVMCTRL_STATUS until FBUSY = 0

CCP (Configuration Change Protection):

Register: 0x0034 (CPU_CCP)
 Key for NVMCTRL: 0x9D (CCP_SPM_gc)
 Window: 4 clock cycles
 Rule: STS CPU_CCP must be immediately followed by STS NVMCTRL_CTRLA

BOOTEND fuse (fuse byte 8):

0x00 = no protection
 0x04 = 1 KB boot section (selfupdate.S fits here, ~486 bytes)
 App section starts at BOOTEND × 256

Key instructions introduced:

SBIW Rd, K – subtract immediate from 16-bit register pair (2 cy)
 ADIW Rd, K – add immediate to 16-bit register pair (2 cy)
 IJMP – indirect jump through Z (2 cy)

Build:

```
avr-gcc -mmcu=attiny3217 -nostartfiles -nodefaultlibs \  

  -o selfupdate.elf src/selfupdate.S  

avr-objcopy -O ihex selfupdate.elf selfupdate.hex  

pymcuprog write -d attiny3217 -t uart -f selfupdate.hex
```

16.11 Exercises

1. The `nvm_wait` polling loop costs 5 cycles per iteration at 3.333 MHz. A `PAGEERASEWRITE` operation on ATtiny3217 takes approximately 2.5 ms. How many loop iterations does `nvm_wait` execute while waiting? At 3.333 MHz, what is the duration of each iteration in microseconds? Is a timer-based timeout (to detect a stuck NVM controller) practical within the same loop?
2. The CCP window is exactly 4 clock cycles. The sequence `STS CPU_CCP, r16 + STS NVMCTRL_CTRLA, r16` uses $2 + 2 = 4$ cycles. If the CPU were running at 20 MHz (no prescaler), would the same two back-to-back STS instructions still fit within the window? What if you ran at 32 MHz with an external crystal — would you need to change the assembly?
3. The `nvm_write_page` routine fills the page buffer with 64 iterations of `LD r18, X+ + ST Z+, r18`. Each pair takes $2 + 2 = 4$ cycles. Calculate the total cycle count for the fill loop plus the surrounding

overhead (wait, NOCMD write, SBIW, CCP sequence, final wait entry). What is the minimum time in milliseconds required to write 16 pages (1 KB)?

4. The self-update firmware in `selfupdate.S` does not verify the EOF record's checksum. Add checksum verification: if the EOF checksum is incorrect, send '!', turn LED0 on (it was turned off on successful EOF), and loop back to `record_loop` without jumping to the application. What register holds the checksum accumulator at the point of the EOF handler, and what condition do you test?
5. The current `selfupdate.S` jumps to address `0x0000` on EOF (the reset vector). If `BOOTEND = 0x04`, the application section begins at `0x0400`. Modify the EOF handler to: (a) check that at least one page in the application section was written (track the highest written address in a register), (b) jump to `0x0400` instead of `0x0000`. Show the assembly changes.
6. Implement a **write-verify** step: after writing each page with `nvm_write_page`, read back the 64 bytes using `LPM Z+, r16` and compare against the SRAM buffer. If any byte mismatches, retransmit '!' (NACK) and re-request the record. Where in the `selfupdate.S` main loop would you add this? What additional register or memory is needed?
7. The `CHIPERASE` command (`0x05`) erases all flash not protected by `BOOTEND` in a single operation — much faster than erasing page by page. Modify `selfupdate.S` to issue `CHIPERASE` when the first data record is received (before starting the page-fill loop). What must you do with the byte you received (the first record's data) after `CHIPERASE` completes? Sketch the change to the record handling code.
8. Using `avr-objdump -d selfupdate.elf`, locate the two `STS` instructions that form the CCP sequence. Record their addresses. What is the byte offset between them? Confirm that no other instruction appears at an address between them. Now deliberately insert a `NOP` between them in the source and rebuild — does `avr-objdump` show the `NOP` between the two `STS` instructions? Describe the hardware consequence of this change.

Next: Chapter 17 — Bit Manipulation and Fixed-Point Arithmetic

Chapter 17

Bit Manipulation & Fixed-Point Math

AVR has no hardware divider and no floating-point unit. The ATtiny3217 does include the AVR hardware multiplier, but shift-and-add multiplication remains useful for understanding the machine model and for AVR parts that omit MUL. When computation must be fast and deterministic — sensor scaling, display formatting, signal processing — the answers live in bit manipulation and fixed-point arithmetic.

This chapter covers:

1. **Shift-and-add multiply** — multiplication using only shifts and adds; useful as a portable fallback and as a teaching algorithm.
2. **Fixed-point Q8.8 arithmetic** — representing and operating on fractional numbers in 16-bit register pairs.
3. **BCD arithmetic** — binary-coded decimal for human-readable display.
4. **Practical example** — convert a 16-bit binary sensor reading to BCD for a four-digit 7-segment display, using all three techniques.

The common thread is that all of these techniques trade mathematical elegance for predictability. On a small MCU, that is usually a good trade.

17.1 Shift Operations

AVR provides five shift/rotate instructions. All take a single register operand and execute in **1 cycle**.

Instruction	Operation	Flags
LSL Rd	$C \leftarrow Rd7$; $Rd \leftarrow Rd \ll 1$; $Rd0 = 0$	C, Z, N, V, H
LSR Rd	$Rd7 = 0$; $Rd \leftarrow Rd \gg 1$; $C \leftarrow Rd0$	C, Z, N, V
ASR Rd	$Rd7$ unchanged; $Rd \leftarrow Rd \gg 1$; $C \leftarrow Rd0$	C, Z, N, V
ROL Rd	$C \leftarrow Rd7$; $Rd \leftarrow Rd \ll 1$; $Rd0 = \text{old } C$	C, Z, N, V, H
ROR Rd	$Rd7 = \text{old } C$; $Rd \leftarrow Rd \gg 1$; $C \leftarrow Rd0$	C, Z, N, V

LSL/LSR: logical shift, zero fill. LSL Rd multiplies by 2; LSR Rd divides by 2 (unsigned).

ASR: arithmetic shift right — sign-extends during right shift, preserving the sign of a signed value. ASR Rd divides a signed byte by 2 (rounds toward negative infinity).

ROL/ROR: rotate through carry. The carry flag acts as a 9th bit. This is how multi-byte shifts are chained:

That is the key idea behind most “wide” arithmetic on AVR. You do not get a 16-bit or 32-bit shift instruction; you build one byte at a time and let carry connect the pieces.

```

/* Logical left shift a 16-bit value in R25:R24 by 1 bit */
lsl  r24      /* shift low byte left; bit 7 → C      */
rol  r25      /* shift high byte left; C → bit 0      */

```

Multi-byte right shift (logical):

```

/* Logical right shift R25:R24 by 1 bit */
lsr  r25      /* shift high byte right; bit 0 → C      */
ror  r24      /* shift low byte right; C → bit 7      */

```

17.1.1 Register Pair Diagram — 16-bit LSL

Before: R25[b7..b0] R24[b7..b0] └ (bit 0 = 0 after LSL)

After:

```

R24: b6 b5 b4 b3 b2 b1 b0 0  ← LSL r24 (b7 of R24 → C)
R25: b6 b5 b4 b3 b2 b1 b0 C  ← ROL r25 (C → bit 0 of R25)
C   = old R25[b7]

```

17.2 Shift-and-Add Multiplication

The MUL instruction ($8 \times 8 \rightarrow 16$ unsigned) exists on ATtiny3217 and many other AVR devices. Some smaller or older AVR parts omit hardware multiply, so the shift-and-add algorithm is still worth knowing.

The shift-and-add algorithm works on any AVR: for each bit in the multiplier, if the bit is set, add the multiplicand (shifted appropriately) to the accumulator.

17.2.1 Algorithm — Unsigned $8 \times 8 \rightarrow 16$

```

product = 0
for bit = 0 to 7:
    if multiplier[bit] == 1:
        product += multiplicand << bit
    (equivalently: shift multiplicand left each iteration,
     add if LSB of multiplier is 1, then shift multiplier right)

```

17.2.2 Implementation

Instruction	Operands	Cycles	Flags
LSL Rd	Rd	1	C,Z,N,V,H
LSR Rd	Rd	1	C,Z,N,V
ROL Rd	Rd	1	C,Z,N,V,H
ROR Rd	Rd	1	C,Z,N,V
ADD Rd,Rr	Rd, Rr	1	C,Z,N,V,H
ADC Rd,Rr	Rd, Rr	1	C,Z,N,V,H
BRCC k	k (−64..63)	1/2	—
DEC Rd	Rd	1	Z,N,V

```

/*
 * mul8u — unsigned 8×8 → 16 multiply (shift-and-add)
 *
 * Entry: R16 = multiplicand

```



```

*      R17 = multiplier
* Exit:  R25:R24 = R16 × R17 (unsigned 16-bit product)
* Clobbers: R16, R17, R18, R20
* Cycles: worst-case ~56
*/
mul8u:
    clr    r24          /* product = 0 */
    clr    r25
    clr    r20          /* multiplicand high byte = 0 */
    ldi    r18, 8       /* 8 bits to process */
.Lmul8u_loop:
    lsr    r17          /* multiplier >>= 1; old bit 0 → C */
    brcc   .Lmul8u_skip
    add    r24, r16      /* product_lo += multiplicand_lo */
    adc    r25, r20      /* product_hi += multiplicand_hi + carry */
.Lmul8u_skip:
    lsl    r16          /* multiplicand_lo <<= 1; old bit 7 → C */
    rol    r20          /* multiplicand_hi <<= 1; C → bit 0 */
    dec    r18
    brne   .Lmul8u_loop
    ret

```

The key detail is that the multiplicand is tracked as a 16-bit value (R20:R16), so left shifts preserve the carry-out from the low byte.

That is what makes the routine correct for products larger than 255. If you shifted only the low byte, the upper half of the multiplicand would vanish.

Trace — 3×5 (i.e. $0b00000011 \times 0b00000101 = 15$):

r16=3, r17=5, r20=0, r24=0, r25=0

Iter 1: lsr r17 → r17=2, C=1 (bit 0 of 5)
 add r24, r16 → r24=3; adc r25, r20 → r25=0
 lsl r16 → r16=6, C=0; rol r20 → r20=0

Iter 2: lsr r17 → r17=1, C=0 (bit 0 of 2)
 skip
 lsl r16 → r16=12; rol r20 → r20=0

Iter 3: lsr r17 → r17=0, C=1 (bit 0 of 1)
 add r24, r16 → r24=3+12=15; adc r25, r20 → r25=0
 lsl r16 → r16=24; rol r20 → r20=0

Iters 4-8: r17=0, C=0 each time → all skip
 Result: r25:r24 = $0x000F$ = 15 OK

17.2.3 C Equivalent

```

uint16_t mul8u(uint8_t a, uint8_t b) {
    uint16_t product = 0;
    uint16_t multiplicand = a;
    for (int i = 0; i < 8; i++) {
        if (b & 1) product += multiplicand;
        multiplicand <<= 1;
        b >>= 1;
    }
}

```

```

    }
    return product;
}

```

17.2.4 Signed Multiply — mul8s (8×8 → 16 signed)

Signed multiply extends naturally: run the unsigned loop for 7 bits, then treat the MSB of the multiplier as negative (two's complement weight -128):

```

/*
 * mul8s - signed 8×8 → 16 multiply
 *
 * Entry:  R16 = multiplicand (signed)
 *         R17 = multiplier   (signed)
 * Exit:   R25:R24 = R16 × R17 (signed 16-bit product)
 * Clobbers: R16, R17, R18, R20
 */
mul8s:
    clr    r24
    clr    r25
    /* sign-extend R16 into R20 */
    mov    r20, r16
    lsl    r20          /* old bit 7 → C                */
    sbc    r20, r20      /* r20 = 0xFF if negative, 0x00 if positive        */
    ldi    r18, 7        /* process 7 unsigned bits                        */
.Lmul8s_loop:
    lsr    r17
    brcc   .Lmul8s_skip
    add    r24, r16
    adc    r25, r20
.Lmul8s_skip:
    lsl    r16
    rol    r20
    dec    r18
    brne   .Lmul8s_loop
    /* bit 7 of multiplier: subtract instead of add (two's complement) */
    lsr    r17          /* bit 7 → C                */
    brcc   .Lmul8s_done
    sub    r24, r16      /* product -= (multiplicand << 7)                */
    sbc    r25, r20
.Lmul8s_done:
    ret

```

SBC Rd, Rr subtracts Rr and the carry flag from Rd. Here, after `lsr r17`, C holds the sign bit of the original multiplier. If set, we subtract the current (shifted) multiplicand — the two's complement correction for the negative weight of bit 7.

This is one of those places where two's complement stops being abstract. The top bit is not just “another positive weight”; it contributes a negative term, so the final step has to subtract instead of add.


17.3 Fixed-Point Q8.8 Arithmetic

Fixed-point represents a real number as an integer scaled by a power of two. **Q8.8** uses a 16-bit register pair where the high byte is the integer part and the low byte is the fractional part:

The advantage is that the CPU still performs ordinary integer arithmetic. The “fractional” meaning comes from how you interpret the bits, not from special hardware support.

Q8.8 value = register pair / 256

Bit layout (R25:R24 example):



Range: -128.0 to +127.99609375 (signed)
0.0 to 255.99609375 (unsigned)

Resolution: $1/256 \approx 0.0039$

17.3.1 Converting to Q8.8

```

/* Integer 42 → Q8.8: shift left 8 bits (multiply by 256) */
ldi    r24, 0           /* fractional part = 0 */
ldi    r25, 42          /* integer part = 42 */
/* R25:R24 = 0x2A00 = 10752 raw = 42.0 in Q8.8 */

/* Fractional constant 0.5 → Q8.8 = 128 = 0x0080 */
ldi    r24, 128
clr    r25
/* R25:R24 = 0x0080 = 0.5 in Q8.8 */

/* Convert from Q8.8 back to integer (truncate) */
mov    r16, r25          /* integer part = high byte */

```

17.3.2 O8.8 Addition

Identical to 16-bit integer addition — no adjustment needed:

That simplicity is one of the main reasons fixed-point is attractive. Once two values share the same scaling, addition and subtraction are just normal integer operations.

```
/* R25:R24 += R23:R22 (Q8.8 + Q8.8) */
add    r24, r22
adc     r25, r23
```

17.3.3 Q8.8 Multiplication ($\mathbb{Q}8.8 \times \mathbb{Q}8.8 \rightarrow \mathbb{Q}8.8$)

Two Q8.8 values multiplied give a Q16.16 result in a 32-bit space. To get a Q8.8 result, extract bits [23:8] of the 32-bit product — effectively shift right 8 bits:

```
A = a15..a0 (Q8.8)    B = b15..b0 (Q8.8)
A × B = 32-bit Q16.16
Q8.8 result = (A × B) >> 8 = bits [23:8] of the 32-bit product
```

In practice, build the 32-bit product as four 8×8 partial products:

A = AH:AL (high byte : low byte)

B = BH:BL

$$A \times B = AH \times BH \text{ (weight } 2^{16}) + AH \times BL \text{ (weight } 2^8) + AL \times BH \text{ (weight } 2^8) + AL \times BL \text{ (weight } 2^0)$$

Q8.8 result = shift right 8 → take bits [23:8]:

bits [15:8] of (AH×BL + AL×BH) + AH×BH (lower byte) + carry from AL×BL

Since ATtiny3217 has no MUL, use mul8u from Section 17.2. The same partial product decomposition applies:

The bookkeeping looks tedious, but the idea is straightforward: break the 16-bit multiply into byte-sized pieces, then keep only the middle 16 bits that correspond to the Q8.8 result.

```
/*
 * mulq88 – unsigned Q8.8 × Q8.8 → Q8.8 using mul8u
 *
 * Entry:  R25:R24 = A (Q8.8),  R23:R22 = B (Q8.8)
 * Exit:   R25:R24 = A × B (Q8.8, truncated)
 */
mulq88:
    push    r14
    push    r15

    mov     r14, r24    /* AL                      */
    mov     r15, r25    /* AH                      */

    /* AL × BL */
    mov     r16, r14
    mov     r17, r22
    rcall   mul8u
    mov     r16, r25    /* high byte of AL×BL contributes at weight 2^8 */
    clr     r17

    /* AL × BH */
    mov     r16, r14
    mov     r17, r23
    rcall   mul8u
    add     r16, r24
    adc     r17, r25

    /* AH × BL */
    mov     r16, r15
    mov     r17, r22
    rcall   mul8u
    add     r16, r24
    adc     r17, r25

    /* AH × BH */
    mov     r16, r15
    mov     r17, r23
    rcall   mul8u
    add     r17, r24    /* low byte of AH×BH contributes at weight 2^16 */

    movw    r24, r16
```

```

pop    r15
pop    r14
ret

```

Full listing in `src/fixedpoint.S`.

17.4 BCD Arithmetic

Binary-Coded Decimal stores each decimal digit as a 4-bit nibble. Two formats:

Unpacked BCD: one digit per byte (low nibble used, high nibble = 0)

Decimal 47 → byte 0x04, byte 0x07

Packed BCD: two digits per byte

Decimal 47 → single byte 0x47

17.4.1 Decimal Adjust on AVR

Unlike some 8-bit CPUs, AVR does **not** provide a general DAA instruction for packed-BCD correction. After a binary ADD/ADC, packed BCD must be corrected manually using the half-carry and carry flags.

17.4.2 H Flag and BCD Correction

The H (half-carry) flag is set when a carry propagates from bit 3 to bit 4 during addition — i.e., when the low nibble overflows. This is exactly the condition that requires BCD correction:

After binary ADD of two packed BCD bytes:

If low nibble > 9 OR H flag set: add 0x06 to correct low nibble

If high nibble > 9 OR C flag set: add 0x60 to correct high nibble

```

/*
 * bcd_add – add two single packed BCD bytes
 *
 * Entry:  R16 = BCD byte A (e.g. 0x47 for decimal 47)
 *         R17 = BCD byte B
 * Exit:   R16 = BCD sum (low byte), C = 1 if result > 99
 * Clobbers: R18
 */
bcd_add:
    add    r16, r17      /* binary add */
    ldi    r18, 0
    /* check low nibble: H flag set or low nibble > 9 */
    brhs   .lbcd_fix_low /* H set → low nibble overflowed */
    mov    r18, r16
    andi   r18, 0x0F
    cpi    r18, 10
    brlo   .lbcd_check_high
.lbcd_fix_low:
    subi   r16, -6        /* add 6 to low nibble (subi x, -N = addi x, N) */
.lbcd_check_high:
    brcs   .lbcd_fix_high /* carry from low correction means high overflowed */
    mov    r18, r16
    andi   r18, 0xF0
    cpi    r18, 0xA0

```

```

    brlo  .lbcd_done
.lbcd_fix_high:
    subi  r16, -0x60    /* add 0x60 to high nibble          */
.lbcd_done:
    ret

```

BRHS (Branch if Half-carry Set) and BRHC (Branch if Half-carry Clear) branch directly on the H flag.

That matters because packed BCD correction is really nibble arithmetic. The H flag tells you whether the low decimal digit overflowed even when the full byte did not.

Instruction	Branch condition	Cycles
BRHS k	H = 1	1/2
BRHC k	H = 0	1/2

17.5 Binary to BCD Conversion

Converting a multi-digit binary number to BCD for display is a recurring task. The **double-dabble** (shift-and-add-3) algorithm converts an N-bit binary number to packed BCD in N iterations, requiring no division:

Algorithm:

Place binary value in a shift register alongside N/4 BCD digits (all zero).

Repeat N times:

For each BCD digit: if digit > 4, add 3 to that digit.

Shift the entire combined register left by 1 bit.

After N iterations, BCD digits contain the decimal representation.

The “add 3” step ensures that a BCD digit that is about to overflow (≥ 5) will produce the correct carry into the next BCD digit after the shift (since shifting left doubles, and $5 \times 2 = 10$ wraps in BCD; pre-adding 3 makes $(5+3) \times 2 = 16$, which carries correctly as 6 in the upper digit and 0 in the lower).

If double-dabble feels magical at first, focus on that one fact: digits 5..9 need help before the shift so that a binary left shift turns into a valid decimal carry.

17.5.1 16-bit Binary to 5-Digit BCD

A 16-bit value (0–65535) needs 5 BCD digits (0–9 each), packed into 3 bytes (two packed + one partial):

Packed BCD storage: [tens-thousands][thousands|hundreds][tens|ones]

Example: 65535 $\rightarrow$ 0x06, 0x55, 0x35

```

/*
 * bin16_to_bcd – convert 16-bit unsigned binary to 5 packed BCD digits
 *
 * Entry:  R25:R24 = 16-bit value (0-65535)
 * Exit:   R21:R20:R19 = packed BCD (5 digits)
 *         R21[3:0] = ten-thousands digit
 *         R20[7:4] = thousands digit
 *         R20[3:0] = hundreds digit
 *         R19[7:4] = tens digit
 *         R19[3:0] = ones digit
 * Clobbers: R16, R17, R22
 */
bin16_to_bcd:

```

```

    clr    r19          /* BCD digits [tens|ones] = 0 */
    clr    r20          /* BCD digits [thousands|hundreds] = 0 */
    clr    r21          /* BCD digit [ten-thousands] = 0 */
    ldi    r16, 16      /* loop counter: 16 bits to process */

.Lbcd_loop:
    /* For each BCD nibble: if > 4, add 3 */
    rcall  bcd_adjust   /* adjust R21:R20:R19 */

    /* Shift binary value and BCD digits left by 1 */
    lsl    r24          /* binary low byte: bit 7 → C */
    rol    r25          /* binary high byte: C → bit 0; bit 7 → C */
    rol    r19          /* BCD byte 0: C in, bit 7 → C */
    rol    r20          /* BCD byte 1: C in, bit 7 → C */
    rol    r21          /* BCD byte 2: C in (no further shift needed) */

    dec    r16
    brne   .Lbcd_loop
    ret

/*
 * bcd_adjust – add 3 to any BCD nibble > 4 in R21:R20:R19
 */
bcd_adjust:
    /* Check and adjust each nibble */
    /* R19 low nibble (ones) */
    mov    r22, r19
    andi   r22, 0x0F
    cpi    r22, 5
    brlo   .Lbcd_adj_r19h
    subi   r19, -3      /* add 3 to low nibble */

.Lbcd_adj_r19h:
    /* R19 high nibble (tens) */
    mov    r22, r19
    andi   r22, 0xF0
    cpi    r22, 0x50
    brlo   .Lbcd_adj_r20l
    subi   r19, -0x30

.Lbcd_adj_r20l:
    /* R20 low nibble (hundreds) */
    mov    r22, r20
    andi   r22, 0x0F
    cpi    r22, 5
    brlo   .Lbcd_adj_r20h
    subi   r20, -3

.Lbcd_adj_r20h:
    /* R20 high nibble (thousands) */
    mov    r22, r20
    andi   r22, 0xF0
    cpi    r22, 0x50
    brlo   .Lbcd_adj_r21
    subi   r20, -0x30

.Lbcd_adj_r21:
    /* R21 low nibble (ten-thousands) – max digit is 6, never > 4 until late */

```

```

mov    r22, r21
andi   r22, 0x0F
cpi    r22, 5
brlo   .Lbcd_adj_done
subi   r21, -3
.Lbcd_adj_done:
ret

```

17.5.2 Build and Test

```

avr-gcc -mmcu=attiny3217 -nostartfiles \
        -o bin2bcd.elf bin16_to_bcd.S
avr-objdump -d bin2bcd.elf

```

Simulate with simavr:

```
simavr -m attiny3217 bin2bcd.elf
```

17.6 Complete Example: Sensor Reading to 7-Segment Display

A 10-bit ADC reading (0–1023) from a temperature sensor is scaled to a °C value (0–99.9°C, one decimal place), then displayed on a four-digit 7-segment display connected over SPI.

17.6.1 Scaling: ADC → Temperature

Assume linear sensor: 0 → 0.0°C, 1023 → 99.9°C.

```
temp_Q8 = (adc_reading × 25344) >> 16
```

Where 25344 = round(99.9 / 1023 × 65536) — maps the full ADC range to 99.9 scaled as a 16-bit integer. The result is in units of 0.1°C.

Or, equivalently with Q8.8:

```
scale_factor = 99.9 × 256 / 1023 ≈ 24.99... ≈ Q8.8 value 0x18FF
temp_Q88 = adc_reading_Q88 × scale_factor_Q88 (Q8.8 multiply)
```

We use the simpler fixed shift approach here to avoid an extra Q8.8 multiply.

The full example is in `src/display.S`. Below are the key routines.

This example is intentionally end-to-end. It shows why bit tricks and numeric formats matter: eventually the value has to become something a human can read.

17.6.2 Register Map for display.S

R25:R24	ADC reading (10-bit, 0–1023)
R23:R22	scaled temperature × 10 (e.g. 365 = 36.5°C)
R21:R20:R19	BCD digits (five digits packed)
R18	digit index (0–3)
R17	segment data byte to SPI
R16	scratch
Z (R31:R30)	pointer to 7-segment LUT in flash

17.6.3 7-Segment LUT

```
.section .rodata
seg7_lut:
    /* segments: gfedcba (bit 6..0), active high */
    .byte 0x3F /* 0: abcdef = 0b0111111 */
    .byte 0x06 /* 1: bc      = 0b0000110 */
    .byte 0x5B /* 2: abdeg  = 0b1011011 */
    .byte 0x4F /* 3: abcdg  = 0b1001111 */
    .byte 0x66 /* 4: bcfg   = 0b1100110 */
    .byte 0x6D /* 5: acdfg  = 0b1101101 */
    .byte 0x7D /* 6: acdefg = 0b1111101 */
    .byte 0x07 /* 7: abc     = 0b0000111 */
    .byte 0x7F /* 8: all    = 0b1111111 */
    .byte 0x6F /* 9: abcdfg = 0b1101111 */
```

17.6.4 Scale ADC to Temperature

The implementation in `src/display.S` uses a 32-bit intermediate for:

```
temp10 = (adc × 999 + 512) >> 10
```

That keeps the result accurate across the full 0-1023 ADC range and avoids the overflow problems that appear in a 16-bit-only sketch.

That is a practical fixed-point lesson by itself: the final answer may fit in 16 bits, but an intermediate step often needs more headroom.

17.6.5 Extract BCD Digits and Drive SPI

```
/*
 * display_temp – convert R23:R22 (temp×10) to BCD and send to SPI display
 *
 * Digit layout: [D3][D2][D1].[D0] (D1D0 = ones and tenths, decimal point
 *                  after digit D1)
 * Entry:  R23:R22 = temperature × 10 (0-999)
 */
display_temp:
    movw r24, r22
    rcall bin16_to_bcd /* R21:R20:R19 = BCD digits */

    /* temp×10 = abc in decimal corresponds to display "ab.c".
       hundreds digit = R20[3:0]
       tens digit     = R19[7:4]
       ones digit     = R19[3:0]
       For example, 365 (meaning 36.5°C) displays as " 36.5". */

    /* Send digit 3 (leading digit; blank if zero) */
    mov r16, r20
    andi r16, 0x0F
    breq .ldisp_blank_d3
    rcall send_digit
    rjmp .ldisp_d2
.ldisp_blank_d3:
    ldi r16, 0x00 /* blank segment */
    rcall spi_send
```

```

.Ldisp_d2:
    /* tens digit */
    mov    r16, r19
    swap   r16
    andi   r16, 0x0F
    rcall  send_digit    /* send plain digit */

    /* ones digit with decimal point */
    mov    r16, r19
    andi   r16, 0x0F
    rcall  send_digit_dp /* send digit + decimal point bit */

    /* tenths digit comes from the ones digit of temp*10.
       In the full source, preserve that nibble before reusing it. */
    ret

/*
 * send_digit – look up 7-seg pattern for digit in R16, send via SPI
 */
send_digit:
    ldi    ZL, lo8(seg7_lut)
    ldi    ZH, hi8(seg7_lut)
    add    ZL, r16    /* index into LUT */
    adc    ZH, r1
    lpm    r17, Z      /* r17 = segment pattern */
    rjmp   spi_send

/*
 * spi_send – send R17 via SPI (polling)
 */
spi_send:
    sts    SPDR, r17    /* load SPI data register */

.Lspi_wait:
    lds    r16, SPSR
    sbrc   r16, SPIF     /* wait for SPIF (transfer complete) bit */
    rjmp   .Lspi_wait
    ret

```

17.6.6 Build

```

# display.S in this directory targets ATmega328P SPI hardware
avr-gcc -mmcu=atmega328p -nostartfiles \
    -o display.elf display.S fixedpoint.S bin16_to_bcd.S
avr-objcopy -O ihex display.elf display.hex
avr-size display.elf

```

17.7 Instruction Reference for This Chapter

Instruction	Syntax	Cycles	Flags	Notes
LSL	lsl Rd	1	C,Z,N,V,H	$Rd \leftarrow Rd \ll 1$; $C \leftarrow Rd7$
LSR	lsr Rd	1	C,Z,N,V	$Rd \leftarrow Rd \gg 1$; $C \leftarrow Rd0$
ASR	asr Rd	1	C,Z,N,V	Signed $\gg 1$, sign extends
ROL	rol Rd	1	C,Z,N,V,H	Rotate left through C
ROR	ror Rd	1	C,Z,N,V	Rotate right through C
SWAP	swap Rd	1	—	Exchange high/low nibbles
MUL	mul Rd,Rr	2	Z,C	$R1:R0 \leftarrow Rd \times Rr$ unsigned
MULS	muls Rd,Rr	2	Z,C	Signed $\times$ signed
MULSU	mulsu Rd,Rr	2	Z,C	Signed $\times$ unsigned
BRHS	brhs k	1/2	—	Branch if H set
BRHC	brhc k	1/2	—	Branch if H clear
ADIW	adiw Rd,K	2	Z,C,N,V,S	Add immediate (0–63) to register pair

SWAP Rd is extremely useful for packed BCD: it moves the high nibble to the low position for digit extraction without shifting:

That is much cleaner than doing four single-bit shifts just to expose one decimal digit.

```
/* Extract high nibble of R16 as a value 0-15 */
swap r16          /* high nibble → low nibble          */
andi r16, 0x0F    /* mask off high bits          */
```

17.8 Common Pitfalls

17.8.1 Pitfall 1 — Using R1 Without Zeroing After MUL

MUL writes R1:R0. The GCC ABI and avr-libc assume R1 is always zero. After any MUL sequence, restore:

```
mul r16, r17
mov r24, r0      /* use result */
clr r1           /* MUST restore before calling any C code or library */
```

17.8.2 Pitfall 2 — ASR vs LSR on Signed Values

LSR fills the MSB with 0, which is wrong for negative values:

```
-8 in 8-bit two's complement: 0b11111000 = 0xF8
LSR → 0b01111100 = 0x7C = +124  WRONG
ASR → 0b11111100 = 0xFC = -4   CORRECT (-8 / 2)
```

Always use ASR for signed right shifts.

If you pick the wrong shift instruction on signed data, the result is not “slightly off.” It is usually catastrophically wrong.

17.8.3 Pitfall 3 — Q8.8 Overflow

Q8.8 integer range is ± 127 (signed) or 0–255 (unsigned). Multiplication of two large Q8.8 values easily overflows:

$200.0 \times 200.0 = 40000.0$ – exceeds unsigned Q8.8 max (255.99)

Check that the mathematical range fits before choosing Q8.8. For larger ranges, use Q4.12 (4-bit integer, 12-bit fraction) or Q12.4.

Choosing a fixed-point format is mostly a range-versus-resolution trade. More fraction bits give you finer steps; more integer bits give you more headroom.

17.8.4 Pitfall 4 — BCD Correction Order

The BCD correction must check low nibble **before** high nibble. Correcting the high nibble first can leave the low nibble invalid after the correction carry propagates:

```
/* WRONG: high nibble corrected first */
subi r16, -0x60 /* might corrupt low nibble */
subi r16, -0x06

/* CORRECT: low nibble first */
subi r16, -0x06
subi r16, -0x60
```

17.8.5 Pitfall 5 — Shift Count for Multi-Byte Values

AVR shifts operate on single registers only. There is no `lsl r25:r24` opcode. Multi-byte shifts require one instruction per byte:

```
/* 24-bit left shift (R22:R21:R20) */
lsl r20
rol r21
rol r22
```

Forget the middle `rol` and bits disappear silently.

That is a recurring AVR pattern: multi-byte arithmetic usually fails by losing a carry in the middle, not by producing an obvious syntax error.

17.9 Summary

Shift recap:

- LSL / LSR – logical shift (unsigned), zero fill
- ASR – arithmetic right shift (signed), sign extend
- ROL / ROR – rotate through C; chain with LSL/LSR for multi-byte shifts
- SWAP – swap nibbles in 1 cycle; essential for packed BCD

Shift-and-add multiply (no hardware MUL):

- 8 iterations: `lsr multiplier`; `brcs skip`; add accumulator; shift 16-bit multiplicand
- ~56 cycles for 8×8 vs 2 cycles for MUL – use MUL when available

Fixed-point Q8.8:

- `value = raw_int / 256`
- add: same as 16-bit integer add
- multiply: 4 partial products, extract bits [23:8] of 32-bit result

BCD conversion (double-dabble):

- 16 iterations of: add-3-if-nibble>4, then shift-left
- No division, no LUT, predictable cycle count at the cost of a relatively long loop

Binary-to-BCD use cases:

ADC reading → decimal display

Timer counter → MM:SS display

Any sensor value needing human-readable output

Build:

```
avr-gcc -mmcu=attiny3217 -nostartfiles -o ch17.elf *.S
avr-objdump -d ch17.elf    # verify cycle counts
avr-size ch17.elf         # check flash budget
```

17.10 Exercises

1. The `mul8u` routine takes ~56 cycles worst-case. For a specific use case — multiplying an 8-bit value by the constant 17 — write a dedicated routine using only `LSL/ADD/MOV` that computes the result in fewer than 6 cycles. (Hint: $17 = 16 + 1$.)
 2. Q4.12 format stores 4 integer bits and 12 fractional bits in 16 bits. Write the Q4.12 multiply: two Q4.12 values → Q4.12 result. Where do you extract bits from the 32-bit intermediate product? What is the new maximum integer value?
 3. The `bin16_to_bcd` routine calls `bcd_adjust` 16 times — once per shift. Each `bcd_adjust` call checks all 5 nibbles. How many of those 5-nibble checks do useful work on iteration 1 (when all BCD bytes are 0)? Optimise `bcd_adjust` to skip nibbles that cannot yet be > 4 based on loop iteration count. How many instruction words does this save?
 4. Implement `bcd16_add` — add two 4-digit packed BCD values (two bytes each, representing 0000–9999) with carry out. Use the H flag and manual correction. Test with $9999 + 0001$ (result: 0000, carry = 1).
 5. The `seg7_lut` uses active-high segment encoding. Many 7-segment driver ICs use active-low (a 0 bit turns the segment on). Rewrite `send_digit` to invert the LUT byte before sending. Use a single `COM` Rd instruction (one's complement, 1 cycle) — do not rewrite the LUT itself.
 6. Profile `display_temp` using `avr-objdump -d`. Count the cycles for the path where no digit is blank. Then count cycles for the path where the hundreds digit is zero (blanked). Which path is longer and by how much?
 7. The double-dabble loop shifts the binary input value leftward until it is consumed. After 16 iterations, $R25:R24 = 0$. Verify this is always true regardless of the input value, and explain why it does not cause a problem that $R25:R24$ is destroyed during the conversion.
-

Next: Chapter 18 — Real-Time Patterns

Chapter 18

Real-Time Patterns

Real-time embedded systems must respond to events within bounded, predictable time. An AVR running bare-metal assembly has no OS scheduler, no kernel, and no hidden overhead — every cycle is accounted for. That is both its strength and its responsibility.

This chapter covers:

1. **Cooperative task scheduling** — a minimal round-robin scheduler in assembly; no dynamic allocation, no context switch overhead beyond register saves.
2. **Deterministic ISR latency analysis** — counting cycles from interrupt trigger to first ISR instruction; identifying and eliminating latency sources.
3. **Example: two-task scheduler with TCB0** — a TCB0 periodic interrupt fires every 1 ms; cooperative task switching drives an LED blinker and a USART heartbeat sender on the ATtiny3217 Curiosity Nano.

The main goal is not to build a tiny RTOS. It is to show how timing, interrupts, and task structure fit together when you want behavior that stays bounded and explainable.

18.1 Real-Time Fundamentals

A real-time system is not necessarily fast — it is **predictable**. Two key properties:

Deadline: a task must complete before a specific time. Missing a deadline is a failure, regardless of whether the computation eventually finishes.

Determinism: given the same inputs, execution time is bounded and known in advance. The bound must hold under all interrupt patterns and all data values.

18.1.1 Two Scheduling Models

Preemptive: a higher-priority task can interrupt a lower-priority task at any instruction boundary. The interrupted task's full register state must be saved and restored. This is what FreeRTOS and similar RTOSes implement.

Cooperative: tasks run to completion (or to an explicit yield point) before the scheduler runs the next task. No mid-task preemption. Context switch is cheaper; latency is bounded only by the longest task segment between yields.

That last clause is the defining tradeoff. Cooperative systems are simple because nothing interrupts a task halfway through, but that also means one task can delay everything behind it if it runs too long.

This chapter implements cooperative scheduling. It is appropriate when:

- All tasks have known, short run times.
- The system has only a few tasks (two to eight is typical on an 8-bit AVR).
- Worst-case response time = sum of all task segments + scheduler overhead, and that total meets all deadlines.

If you can write that sum down and it fits comfortably inside your timing budget, cooperative scheduling is often the cleanest answer on an AVR.

18.2 Cooperative Task Scheduler Design

18.2.1 Task Representation

Each task is a subroutine that runs, does some work, and returns. The scheduler calls each task in rotation. A task that needs to wait (e.g., for a UART byte) must either poll with a non-blocking check and return immediately, or use a flag set by an ISR.

Scheduler main loop:

```
loop forever:
    call task_0
    call task_1
    call task_2
    ...
```

This is the simplest possible cooperative scheduler. Its properties:

- **No stack switch:** each task runs on the same stack, called directly.
- **No saved context:** tasks return normally; all register state is the task's own responsibility (save what you use).
- **Tick-driven variant:** tasks are only called when a timer tick has elapsed, providing a fixed time base.

Notice what is missing: there is no saved "task state" in CPU registers and no separate per-task stack. These tasks are really just well-behaved subroutines that run on a schedule.

18.2.2 Tick-Driven Scheduler

A TCB0 periodic interrupt fires every 1 ms (the *tick*). Each tick sets a flag. The main loop spins until the flag is set, clears it, then runs all tasks:

```
volatile uint8_t tick_flag;

ISR(TCB0_INT) {
    tick_flag = 1;
}

main() {
    while (1) {
        while (!tick_flag);    /* wait for tick */
        tick_flag = 0;
        task_led();
        task_uart();
    }
}
```

In assembly, `tick_flag` is a byte in SRAM (`.bss`). The ISR sets it; the main loop polls and clears it.

That one-byte handoff is a classic embedded pattern: ISR records that an event happened, main-loop code does the slower work later.

18.2.3 Task Counter Pattern

Each task maintains its own counter and acts only when its counter reaches its period:

```
task_led:
    led_counter++
    if led_counter < LED_PERIOD: return
    led_counter = 0
    toggle LED
    return

task_uart:
    uart_counter++
    if uart_counter < UART_PERIOD: return
    uart_counter = 0
    send heartbeat byte
    return
```

With a 1 ms tick:

- `LED_PERIOD = 500` → LED toggles every 500 ms (1 Hz blink)
- `UART_PERIOD = 1000` → UART byte sent every 1 s

This pattern scales well because each task keeps its own notion of time without needing a complex global scheduler.

18.3 TCB0 Configuration — 1 ms Tick

On the ATtiny3217 Curiosity Nano the CPU runs at **3.333 MHz** (internal 20 MHz oscillator divided by 6, set by `CLKCTRL_MCLKCTRLB`).

TCB0 in **Periodic Interrupt mode** (`CNTMODE = 0`) counts up from 0, reloads from `CCMP` when it matches, and fires an interrupt on each match. With no prescaler (`CLK_PER`):

```
CCMP = (F_CPU / F_tick) - 1
      = (3 333 333 / 1000) - 1
      = 3333 - 1
      = 3332 (0x0D04)
```

18.3.1 TCB0 Register Map

Address	Name	Description
0x0A40	TCB0_CTRLA	Control A – clock select, enable
0x0A41	TCB0_CTRLB	Control B – mode select
0x0A44	TCB0_INTCTRL	Interrupt enable (bit 0: CAPT)
0x0A45	TCB0_INTFLAGS	Interrupt flags (bit 0: CAPT; write 1 to clear)
0x0A4A	TCB0_CNTL	Counter low byte
0x0A4B	TCB0_CNTH	Counter high byte
0x0A4C	TCB0_CCMP	Compare/capture match low byte

0x0A4D TCB0_CCMPH Compare/capture match high byte

TCB0_CTRLA bit layout:

Bit	Name	Description
0	ENABLE	1 = timer runs
3:1	CLKSEL	000 = CLK_PER (no prescaler) 001 = CLK_PER/2 110 = TCA0 CLK (divide-down from TCA0 prescaler output)

TCB0_CTRLB bit layout:

Bit	Name	Description
2:0	CNTMODE	000 = Periodic Interrupt (INT mode) 001 = Timeout Check 010 = Input Capture on Event 011 = Input Capture Frequency Measurement 100 = Input Capture Pulse-Width Measurement 101 = Input Capture Frequency and Pulse-Width Measurement 110 = Single Shot 111 = 8-bit PWM

18.3.2 Initialization Code

```

/* tcb0_init - configure TCB0 for 1 ms periodic interrupt
 *
 * F_CPU = 3 333 333 Hz, CCMP = 3332 = 0x0D04
 * Clobbers: R16, R17
 */
tcb0_init:
    /* Set CCMP = 3332 (low byte first, high byte second - TCB locks on CCMPH write) */
    ldi    r16, lo8(3332)          /* 0x04 */
    sts    TCB0_CCMPH, r16
    ldi    r16, hi8(3332)          /* 0x0D */
    sts    TCB0_CCMPH, r16

    /* Mode: periodic interrupt (CNTMODE = 0) */
    ldi    r16, 0
    sts    TCB0_CTRLB, r16

    /* Enable CAPT interrupt */
    ldi    r16, TCB_CAPT_bm        /* bit 0 */
    sts    TCB0_INTCTRL, r16

    /* Enable timer, clock = CLK_PER (no prescaler) */
    ldi    r16, TCB_ENABLE_bm      /* bit 0 */
    sts    TCB0_CTRLA, r16

    ret

```

Instruction	Cycles
LDS Rd, k	3
STS k, Rr	3

LDS/STS access SRAM-mapped peripherals (addresses > 0x001F). Each costs 3 cycles on avrxmega3.

That is an easy place to lose cycles in timer setup and ISR code on newer AVR parts: many peripheral registers are no longer reachable with the short I/O instructions.

18.4 ATtiny3217 Interrupt Vector Table

Vectors are 4 bytes each (one JMP instruction). The table starts at address 0x0000 in flash.

Byte	addr	Vector #	Source	Description
0x0000		0	RESET	Power-on, external, WDT, BOD reset
0x0004		1	CRCSCAN_NMI	CRC scan non-maskable interrupt
0x0008		2	BOD_VLM	Brown-out detector voltage level monitor
0x000C		3	PORTA_PORT	Port A pin-change
0x0010		4	PORTB_PORT	Port B pin-change
0x0014		5	PORTC_PORT	Port C pin-change
0x0018		6	RTC_CNT	RTC counter overflow/compare
0x001C		7	RTC_PIT	RTC periodic interrupt timer
0x0020		8	TCA0_LUNF	TCA0 low byte underflow (split mode)
0x0024		9	TCA0_OVF/HUNF	TCA0 overflow / high byte underflow
0x0028		10	TCA0_CMP0	TCA0 compare match 0
0x002C		11	TCA0_CMP1	TCA0 compare match 1
0x0030		12	TCA0_CMP2	TCA0 compare match 2
0x0034		13	TCB0_INT	TCB0 capture/compare (our tick ISR)
0x0038		14	TCB1_INT	TCB1 capture/compare
0x003C		15	TWI0_TWIS	TWI0 slave
0x0040		16	TWI0_TWIM	TWI0 master
0x0044		17	SPI0_INT	SPI0
0x0048		18	USART0_RXC	USART0 receive complete
0x004C		19	USART0_DRE	USART0 data register empty
0x0050		20	USART0_TXC	USART0 transmit complete
0x0054		21	PORTMUX	Port multiplexer
0x0058		22	ADC0_RESRDY	ADC0 result ready
0x005C		23	ADC0_WCOMP	ADC0 window comparator match
0x0060		24	AC0_AC	Analog comparator 0
0x0064		25	DAC0_DACREADY	DAC0 ready

In assembly, the vector table uses jmp for each slot:

```
.section .vectors, "ax", @progbits
    jmp reset_handler      /* 0x0000 RESET */
    jmp default_isr        /* 0x0004 CRCSCAN_NMI */
    jmp default_isr        /* 0x0008 BOD_VLM */
    jmp default_isr        /* 0x000C PORTA_PORT */
    jmp default_isr        /* 0x0010 PORTB_PORT */
    jmp default_isr        /* 0x0014 PORTC_PORT */
    jmp default_isr        /* 0x0018 RTC_CNT */
    jmp default_isr        /* 0x001C RTC_PIT */
    jmp default_isr        /* 0x0020 TCA0_LUNF */
    jmp default_isr        /* 0x0024 TCA0_OVF */
    jmp default_isr        /* 0x0028 TCA0_CMP0 */
    jmp default_isr        /* 0x002C TCA0_CMP1 */
    jmp default_isr        /* 0x0030 TCA0_CMP2 */
```

```

jmp    tcb0_isr          /* 0x0034 TCB0_INT ← our tick          */
.fill 12 * 4, 1, 0xFF    /* remaining vectors – unused          */

```

.fill count, size, value pads with the given byte value. Here it is only a compact placeholder for omitted vectors in the listing. For production code, prefer populating every vector explicitly (for example, with jmp default_isr) rather than relying on erased-flash filler bytes.

The explicit-table style is a little longer, but it makes interrupt ownership obvious when you inspect the binary later.

18.5 ISR Prologue and Epilogue

Every ISR on AVR must:

1. **Save SREG** — instructions inside the ISR affect flags; they must be restored before RETI.
2. **Save any registers it uses** — if the ISR uses R16–R29 or any other register that the interrupted code relies on.
3. **Acknowledge the interrupt source correctly** — for TCB0, write 1 to TCB0_INTFLAGS bit 0 (TCB_CAPT_bm). Other peripherals may clear their flags by different means, such as a data-register read.
4. **Restore saved registers and SREG.**
5. **RETI** — return and re-enable the global interrupt flag (I bit in SREG).

That save/restore discipline is the price you pay for asynchronous execution. If the ISR leaves flags or registers disturbed, the interrupted code resumes in an impossible state.

18.5.1 Minimal Tick ISR

The tick ISR only sets a 1-byte flag — tick_flag in SRAM:

```

/*
 * tcb0_isr – TCB0 periodic interrupt (1 ms tick)
 *
 * Sets tick_flag = 1.
 * Saves/restores: SREG, R16.
 * Latency from interrupt trigger to first useful instruction: see §18.6.
 */
tcb0_isr:
    push    r16          /* 2 cycles: save R16 to stack          */
    in      r16, _SFR_IO_ADDR(SREG) /* 1 cycle: save SREG                  */
    push    r16          /* 2 cycles                             */

    /* clear interrupt flag (write 1 to TCB_CAPT_bm = bit 0) */
    ldi     r16, TCB_CAPT_bm
    sts     TCB0_INTFLAGS, r16 /* 3 cycles                             */

    /* set tick flag */
    ldi     r16, 1
    sts     tick_flag, r16 /* 3 cycles                             */

    pop     r16          /* 2 cycles: restore SREG value          */
    out     _SFR_IO_ADDR(SREG), r16 /* 1 cycle: restore SREG                */
    pop     r16          /* 2 cycles: restore R16                 */
    reti    /* 4 cycles: return + set I bit          */

```

Total ISR body: $2+1+2+1+3+1+3+2+1+2+4 = 22$ **cycles** from first instruction to end of RETI.

RETI costs 4 cycles (same as RET on avrxmega3) and sets the I bit atomically on return.

That automatic re-enable is why you almost always want RETI only at the true end of the handler, after every register and flag has been restored.

18.6 ISR Latency Analysis

ISR latency = cycles from interrupt trigger condition to the first instruction of the ISR executing.

18.6.1 Sources of Latency

On ATtiny3217 (avrxmega3), the CPU-side interrupt response after the request has been recognized is:

Source	Cycles
Peripheral/request synchronization	0–2
Finish current instruction	1–4
Push PC to stack	2
Jump from vector to ISR (jmp)	3
Total (best case)	6
Total (worst case with 2 sync cycles and a 4-cycle instruction)	11

Microchip documents a minimum CPU interrupt response of 6 cycles on devices with more than 8 KB of flash: 1 cycle to finish a single-cycle instruction, 2 cycles to push the PC, and 3 cycles for the vector jmp. Peripheral-level synchronization can add a small device-specific delay before the CPU begins that sequence. In practice:

Interrupt asserted during instruction N

- request is synchronized into the CPU domain (0–2 cycles)
- instruction N completes (1–4 cycles total, depending on instruction length)
- CPU pushes PC (2 cycles)
- CPU executes JMP at vector (3 cycles)
- first ISR instruction executes (cycle 6–11 from assertion)

This is the right way to think about interrupt latency on AVR: not as one magic number, but as a short chain of small, countable delays.

18.6.2 Push/Pop Overhead

The prologue push r16 / in r16, SREG / push r16 costs **5 cycles** before any useful work. For a time-critical ISR (e.g., one that must sample a pin within N cycles), count:

```
latency_to_useful_work = ISR_entry_latency + prologue_cycles
                        = 6 (best) + 5 = 11 cycles minimum
                        = 11 (worst) + 5 = 16 cycles worst case
```

At 3.333 MHz, 16 cycles = **4.8 μs** worst-case latency to the first useful instruction.

18.6.3 Eliminating Sources of Latency

Technique	Cycles saved
Keep ISRs short — move work to main loop	Reduces re-entry latency
Use VPORT for output (1-cycle sbi/cbi)	0 — saves body cycles, not entry
Avoid long instructions (LPM, CALL) near critical ISR entry points	0–4
Reduce number of registers pushed in prologue	2 per register pair
Use R0 and R1 for ISR scratch (no push needed for GCC-ABI single-purpose ISR)	4

For the tick ISR, latency is unimportant — the tick fires every 1 ms and a few microseconds of jitter is negligible. Latency matters for ISRs that must capture hardware events (e.g., TCB0 in input-capture mode, or a pin-change ISR measuring pulse width).

So always ask “latency relative to what?” before optimizing an ISR. A scheduler tick and an edge-timestamp ISR care about very different things.

18.6.4 Measuring Latency

Toggle a debug pin at the start and end of the ISR, then use an oscilloscope or logic analyser to measure ISR entry and execution width relative to another observable system event:

```
tcb0_isr:
    sbi    VPORTB_OUT, 3      /* debug pin PB3 high – mark ISR entry      */
    push  r16
    ...
    cbi    VPORTB_OUT, 3      /* debug pin PB3 low  – mark ISR exit      */
    reti
```

SBI VPORTB_OUT, 3 executes in **1 cycle** and does not affect SREG — safe before saving SREG, as long as the debug pin is not used by the application.

18.7 Complete Two-Task Scheduler

18.7.1 Hardware Setup — Curiosity Nano

LED0 = PA3, active low (HIGH = off, LOW = on)
 USART0 TX = PB2 (to on-board debugger virtual COM port)
 TCB0 tick = every 1 ms (CCMP = 3332, CLK_PER, F_CPU = 3.333 MHz)

18.7.2 Register Allocation

R2 – tick_flag mirror (ISR writes SRAM; main reads via LDS)
 R24 – scratch / parameter
 R25 – scratch
 R26:R27 (X) – not used (reserved for future tasks)

SRAM variables (.bss):
 tick_flag 1 byte – set by ISR, cleared by main loop
 led_cnt 2 bytes – 16-bit counter for LED task
 uart_cnt 2 bytes – 16-bit counter for UART task

Using SRAM variables (accessed via LDS/STS) ensures the ISR and main loop share data correctly — SRAM is in the data address space accessible to both.

It also keeps the design honest: anything shared between asynchronous code paths should live in memory, not only in a caller's temporary register plan.

18.7.3 Source Code

```

/* scheduler.S – two-task cooperative scheduler, ATtiny3217 Curiosity Nano
 *
 * Tasks:
 *   task_led – toggle LED0 (PA3) every 500 ms
 *   task_uart – send 'H' over USART0 every 1000 ms
 *
 * Tick: TCB0 periodic interrupt, 1 ms
 * F_CPU: 3 333 333 Hz (20 MHz / 6, default Curiosity Nano clock)
 */

#include <avr/io.h>

#define F_CPU      333333UL
#define LED_PIN    3          /* PA3, Curiosity Nano LED0 */
#define LED_PERIOD 500        /* ticks (= 500 ms) */
#define UART_PERIOD 1000      /* ticks (= 1 s) */

/* USART0 baud: 9600 at 3.333 MHz
 * BAUD register = 64 * F_CPU / (16 * baud) = 64 * 333333 / (16 * 9600)
 *                = 21333333 / 153600 ≈ 1389 = 0x056D
 */
#define BAUD_9600  1389

/* — Vector table — */
.section .vectors, "ax", @progbits
    jmp reset_handler      /* 0x0000 RESET */
    jmp default_isr        /* 0x0004 CRCSCAN_NMI */
    jmp default_isr        /* 0x0008 BOD_VLM */
    jmp default_isr        /* 0x000C PORTA_PORT */
    jmp default_isr        /* 0x0010 PORTB_PORT */
    jmp default_isr        /* 0x0014 PORTC_PORT */
    jmp default_isr        /* 0x0018 RTC_CNT */
    jmp default_isr        /* 0x001C RTC_PIT */
    jmp default_isr        /* 0x0020 TCA0_LUNF */
    jmp default_isr        /* 0x0024 TCA0_OVF */
    jmp default_isr        /* 0x0028 TCA0_CMP0 */
    jmp default_isr        /* 0x002C TCA0_CMP1 */
    jmp default_isr        /* 0x0030 TCA0_CMP2 */
    jmp tcb0_isr           /* 0x0034 TCB0_INT */
    .fill (26 - 14) * 4, 1, 0xFF /* remaining 12 vectors unused */

/* — SRAM variables — */
.section .bss
tick_flag: .byte 0          /* 1 = tick has fired; 0 = not yet */
led_cnt:   .word 0          /* LED task counter (16-bit) */
uart_cnt:  .word 0          /* UART task counter (16-bit)

/* — TCB0 ISR — */
.section .text

```

```

tcb0_isr:
    push    r16
    in      r16, _SFR_IO_ADDR(SREG)
    push    r16

    ldi     r16, TCB_CAPT_bm
    sts     TCB0_INTFLAGS, r16    /* clear interrupt flag          */

    ldi     r16, 1
    sts     tick_flag, r16        /* signal tick to main loop    */

    pop     r16
    out     _SFR_IO_ADDR(SREG), r16
    pop     r16
    reti

default_isr:
    reti                          /* ignore unexpected interrupts */

/* — Initialization ————— */
reset_handler:
    /* Stack pointer: top of SRAM = 0x3FFF */
    ldi     r16, hi8(RAMEND)
    out     _SFR_IO_ADDR(SPH), r16
    ldi     r16, lo8(RAMEND)
    out     _SFR_IO_ADDR(SPL), r16

    /* LED0: PA3 output, initial state HIGH (LED off – active low) */
    sbi     VPORTA_DIR, LED_PIN
    sbi     VPORTA_OUT, LED_PIN

    /* USART0: set TX pin PB2 as output */
    sbi     VPORTB_DIR, 2

    rcall   tcb0_init
    rcall   usart0_init

    sei                          /* enable global interrupts    */

    rjmp    main_loop

/* — TCB0 setup ————— */
tcb0_init:
    ldi     r16, lo8(3332)
    sts     TCB0_CCMPH, r16
    ldi     r16, hi8(3332)
    sts     TCB0_CCMPL, r16
    ldi     r16, 0
    sts     TCB0_CTRLB, r16      /* CNTMODE = 0: periodic interrupt */
    ldi     r16, TCB_CAPT_bm
    sts     TCB0_INTCTRL, r16    /* enable CAPT interrupt        */
    ldi     r16, TCB_ENABLE_bm
    sts     TCB0_CTRLA, r16      /* CLK_PER, enable              */

```

```

    ret

/* — USART0 setup ————— */
usart0_init:
    /* BAUD = 1389 = 0x056D */
    ldi    r16, lo8(BAUD_9600)
    sts    USART0_BAUDL, r16
    ldi    r16, hi8(BAUD_9600)
    sts    USART0_BAUDH, r16
    /* CTRLB: TX enable (TXEN bit 6) */
    ldi    r16, USART_TXEN_bm
    sts    USART0_CTRLB, r16
    /* CTRLC: async USART, 8N1 (default reset value = 0x03 = 8-bit, 1 stop) */
    /* No explicit write needed – reset value is correct */
    ret

/* — Main loop ————— */
main_loop:
    /* wait for tick */
.Lwait_tick:
    lds    r16, tick_flag
    tst    r16
    breq   .Lwait_tick

    /* clear flag */
    clr    r16
    sts    tick_flag, r16

    /* run tasks */
    rcall  task_led
    rcall  task_uart

    rjmp   main_loop

/* — task_led ————— */
task_led:
    /* increment 16-bit counter */
    lds    r24, led_cnt
    lds    r25, led_cnt+1
    adiw   r24, 1
    sts    led_cnt, r24
    sts    led_cnt+1, r25

    /* compare to LED_PERIOD (500) */
    cpi    r24, lo8(LED_PERIOD)
    ldi    r16, hi8(LED_PERIOD)
    cpc    r25, r16
    brlo   .lled_done          /* counter < period: nothing to do */

    /* reset counter */
    clr    r24
    clr    r25
    sts    led_cnt, r24
    sts    led_cnt+1, r25

```



```

    /* toggle PA3 via PORTA_OUTTGL (1-byte write toggles the pin) */
    ldi    r16, (1 << LED_PIN)
    sts    PORTA_OUTTGL, r16

.Lled_done:
    ret

/* — task_uart ————— */
task_uart:
    /* increment 16-bit counter */
    lds    r24, uart_cnt
    lds    r25, uart_cnt+1
    adiw   r24, 1
    sts    uart_cnt, r24
    sts    uart_cnt+1, r25

    /* compare to UART_PERIOD (1000) */
    cpi    r24, lo8(UART_PERIOD)
    ldi    r16, hi8(UART_PERIOD)
    cpc    r25, r16
    brlo   .Luart_done

    /* reset counter */
    clr    r24
    clr    r25
    sts    uart_cnt, r24
    sts    uart_cnt+1, r25

    /* send 'H' (0x48): wait for DREIF (Data Register Empty Flag) */
.Luart_wait_dreif:
    lds    r16, USART0_STATUS
    sbrs   r16, USART_DREIF_bp    /* skip next if DREIF=1 (ready) */
    rjmp   .Luart_wait_dreif

    ldi    r16, 'H'
    sts    USART0_TXDATA1, r16

.Luart_done:
    ret

```

18.7.4 Build

```

MCU=attiny3217
avr-gcc -mmcu=${MCU} -nostartfiles -o scheduler.elf scheduler.S
avr-size scheduler.elf
avr-objcopy -O ihex scheduler.elf scheduler.hex
avrdude -p t3217 -c serialupdi -P /dev/ttyUSB0 -U flash:w:scheduler.hex:i

```

18.8 Timing Analysis — Worst-Case Response

With two tasks per tick, the cooperative overhead before the main loop re-checks `tick_flag` is:

```
scheduler_overhead = wait_tick_poll (1 iteration) + clear_flag
                    = (3 + 1 + 1) + (1 + 3) = ~9 cycles
```

```
task_led_max = lds×4 + adiw + sts×4 + cpi + ldi + cpc + brlo + ... toggle path
              ≈ 4×3 + 1 + 4×3 + 1 + 1 + 1 + 2 + clr×2 + sts×4 + ldi + sts
              ≈ 12 + 1 + 12 + 5 + 8 + 4×3 + 1 + 3 = ~54 cycles (toggle path)
              ≈ 28 cycles (no-toggle path)
```

```
task_uart_max (no-send path) ≈ 28 cycles (same structure as task_led)
task_uart_send_fast_path ≈ 28 + 3 + 1 + 3 = ~35 cycles
```

```
total_cycles_per_tick (when USART is immediately ready) ≈ 9 + 54 + 35 = ~98 cycles maximum
time_at_3.333MHz = 98 / 3 333 333 ≈ 29 μs
```

The tick period is 1 ms = 3333 cycles. On the fast path, the scheduler consumes ≈ 98 of those cycles per tick — less than 3% CPU utilisation.

Worst-case tick latency (time between ISR firing and main loop responding):

```
worst_case_tick_latency_fast_path ≈ task_led_max + task_uart_send_fast_path
                                   ≈ 54 + 35 = ~89 cycles ≈ 27 μs
```

This occurs when both tasks do their full work in the tick that just preceded the one being measured. The scheduler handles one tick at a time; if both tasks fire simultaneously on the same tick, `tick_flag` is set for the next tick while the current work finishes.

Important caveat: the scheduler's UART task is still **blocking**. If the code reaches `.Luart_wait_dreif` just after the USART has become busy, the loop can stall for almost a full character time. At 9600 baud with 8N1, one frame is about $10/9600\text{ s} \approx 1.04\text{ ms}$, which is longer than the 1 ms tick. So the fast-path cycle count above is useful for instruction budgeting, but it is **not** a strict worst-case bound for the blocking UART design.

That is the most important design lesson in the chapter: a system is not truly bounded just because its arithmetic paths are short. One blocking peripheral wait can dominate the whole schedule.

18.9 Non-Blocking Task Extension

The fix is to decouple task execution from peripheral transmit timing. The non-blocking variant in `src/scheduler_nb.S` uses:

- a small SRAM ring buffer,
- a short `uart_enqueue` routine used by the task, and
- the `USART0_DRE` interrupt to feed bytes into `TXDATA` whenever the UART is ready for another byte.

Conceptually:

```
task_uart:
    once per second:
        enqueue "Hbeat\r\n"
        enable DRE interrupt

ISR(USART0_DRE):
    if buffer empty:
        disable DRE interrupt
        return
    write next queued byte to TXDATA
    advance tail index
```

This keeps `task_uart` short and bounded: it copies bytes into the queue and returns. The USART hardware timing is handled asynchronously by the DRE ISR. In the example implementation, bytes are dropped if the ring buffer is full; that policy keeps the task non-blocking and deterministic.

This is a typical real-time trade: it is often better to lose optional output than to let one slow device destroy the timing guarantees of the whole system.

Full implementation in `src/scheduler_nb.S`.

18.10 Scaling to More Tasks

The round-robin scheduler scales by adding more `rcall task_N` instructions in the main loop. Each task must be **short** — its worst-case execution time must fit within the tick budget:

```
sum(task_worst_case_cycles) < tick_period_cycles
sum ≤ 3333 cycles (1 ms at 3.333 MHz)
```

For a task that exceeds this (e.g., a DFT computation), split it across multiple ticks using a state machine:

```
dft_task:
    /* process one input sample per tick – spread N-point DFT over N ticks */
    switch (dft_state):
        SAMPLE_0: compute partial sum for k=0; dft_state = SAMPLE_1; return
        SAMPLE_1: compute partial sum for k=1; dft_state = SAMPLE_2; return
        ...
```

Breaking work into per-tick slices is one of the most useful habits in small embedded systems. It keeps long algorithms compatible with a simple scheduler.

18.10.1 When to Move to a RTOS

A cooperative scheduler works until either:

1. Any task's worst-case time can exceed the tick period, or
2. Tasks have different priorities that require preemption.

For those cases, FreeRTOS or RIOT OS port for AVR provides preemptive scheduling with proper stack-switching. The assembly patterns in this chapter (ISR prologue/epilogue, timer setup, ISR-to-task signalling) are the same building blocks used by any RTOS port.

So even if you eventually move to a preemptive kernel, the reasoning here still matters. The kernel is built out of the same timer and interrupt mechanics.

18.11 Summary

Topic	Key Point
Cooperative scheduling	Tasks run to completion; scheduler calls them in rotation
Tick source	TCB0 periodic interrupt, CCMP=3332 for 1 ms at 3.333 MHz
ISR entry latency	6–11 cycles (hardware) + 5 cycles (prologue)
SREG save	Mandatory in any ISR that uses arithmetic or branch instructions
RETI	Restores I bit; costs 4 cycles
VPORT	1-cycle SBI/CBI for fast output; does not affect SREG
Tick budget	3333 cycles/ms; cooperative core uses < 100 cycles on the fast path

Topic	Key Point
Scaling	Add <code>rcall task_N</code> ; split long tasks into per-tick state machines

Next: Chapter 19 — CORDIC on AVR

Chapter 19

CORDIC on AVR

CORDIC (COordinate Rotation DIgital Computer) is one of the classic tricks for small processors. It computes trigonometric, inverse-trigonometric, and vector rotation results using only:

- shifts,
- adds/subtracts,
- comparisons, and
- a small table of constants.

That makes it a natural fit for AVR. Even on AVR parts that have MUL, there is still no floating-point unit and no hardware divider. When you need `sin()`, `cos()`, `atan2()`, or vector magnitude with predictable execution time, CORDIC is often a better fit than floating-point math.

This chapter covers:

1. **Why CORDIC is useful on AVR** — where it beats floating point, and where a lookup table is still better.
2. **Rotation mode** — computing sine and cosine from an angle.
3. **Vectoring mode** — reducing a vector to magnitude and angle.
4. **Fixed-point formats** — a practical AVR-friendly choice: Q1.15 for vector components and BAM16 for angles.
5. **Examples** — a complete integer reference implementation and AVR assembly stage examples.

The important theme is the same one that runs through the rest of the book: replace expensive, hard-to-bound math with small deterministic steps.

19.1 Why Use CORDIC on AVR?

CORDIC is attractive on AVR for three practical reasons:

1. **No floating-point hardware.** A call to `sinf()` or `atan2f()` on AVR pulls in a large software library and can cost hundreds or thousands of cycles.
2. **Fixed iteration count.** A 12-stage CORDIC always does 12 stages. That is useful when you care about real-time bounds.
3. **One core, many functions.** The same shift-add structure can be used for sine/cosine, vector rotation, `atan2`, and magnitude approximation.

Typical AVR uses:

- waveform generation in a DDS/NCO,

- phase tracking,
- I/Q demodulation,
- motor-control angle estimation,
- tilt/heading math from sensors, and
- display-oriented trig without linking `libm`.

19.1.1 When Not to Use It

CORDIC is not magic. Sometimes a lookup table is better:

- If you only need `sin()` at 8-bit resolution, a 256-byte flash table is usually faster.
- If you need one-off trig rarely, a library call may be simpler.
- If your MCU has enough flash and your timing budget is loose, a polynomial or LUT approach may be easier to maintain.

So the real question is not "is CORDIC clever?" It is "does repeated deterministic shift-add math fit this problem better than tables or software floating point?"

19.2 The Core Idea

CORDIC works by applying a sequence of tiny rotations whose tangents are powers of two:

`atan(20), atan(2-1), atan(2-2), atan(2-3), ...`

Because `tan(theta_i) = 2-i`, each stage can be written without a multiply:

```
x' = x plus_or_minus (y >> i)
y' = y minus_or_plus (x >> i)
z' = z minus_or_plus atan(2-i)
```

The shifts replace a true multiply by `tan(theta_i)`.

There are two standard modes:

19.2.1 Rotation Mode

Start with a unit vector on the x-axis and rotate it toward the requested angle. After the final stage:

- `x` approximates `cos(theta)`
- `y` approximates `sin(theta)`

19.2.2 Vectoring Mode

Start with an input vector (`x`, `y`) and rotate it toward the x-axis. After the final stage:

- `z` approximates the vector angle (`atan2`)
- `x` approximates the vector magnitude, scaled by the CORDIC gain
- `y` approaches zero

This is why one algorithm can serve both "give me sine/cosine" and "tell me the angle/magnitude of this vector."

19.3 Fixed-Point Choices for AVR

CORDIC is almost always done in fixed-point on AVR.

19.3.1 Vector Format: Q1.15

Use signed 16-bit values for x and y:

Q1.15: signed fraction in the range about -1.0 to +0.99997

```
0x4000 = +0.5
0x7FFF = +0.99997
0x0000 = 0.0
0xC000 = -0.5
0x8000 = -1.0
```

This is a good fit for sine and cosine because those values naturally live in the interval $[-1, +1]$.

19.3.2 Angle Format: BAM16

For angles, a very AVR-friendly choice is **BAM16** (Binary Angular Measure):

```
Full turn = 65536 units
90°       = 16384 = 0x4000
180°      = 32768 = 0x8000
360°      = wraps to 0x0000
```

Why BAM16 works well:

- addition and wraparound are natural in 16-bit unsigned arithmetic,
- phase accumulators use the same representation,
- quadrants are easy to identify from the top two bits.

19.3.3 CORDIC Gain

The repeated pseudo-rotations scale the vector by a constant gain:

```
1 / K ≈ 1.646760258
K      ≈ 0.607252935
```

For a 12-stage sine/cosine CORDIC, preload:

```
K_Q15 = round(0.607252935 × 32768) = 19897 = 0x4DB9
```

That means rotation mode typically starts with:

```
x = 0x4DB9 ; +0.60725 in Q1.15
y = 0x0000
```

and the final result lands near properly scaled sine/cosine outputs.

19.4 The Angle Table

For 12 stages in BAM16:

i	atan(2 ⁻ⁱ)	BAM16 value
0	45.000°	0x2000
1	26.565°	0x12E4
2	14.036°	0x09FB
3	7.125°	0x0511
4	3.576°	0x028B

i	$\text{atan}(2^{-i})$	BAM16 value
5	1.790°	0x0146
6	0.895°	0x00A3
7	0.448°	0x0051
8	0.224°	0x0029
9	0.112°	0x0014
10	0.056°	0x000A
11	0.028°	0x0005

This table is tiny. Even stored as 16-bit words in flash, 12 stages only cost 24 bytes.

19.4.1 Convergence Range

Basic rotation-mode CORDIC converges only for angles near $\pm 99.88^\circ$, not over the full 360° range. On AVR, the usual fix is **quadrant reduction**:

1. Determine the quadrant from bits [15:14] of the BAM16 angle.
2. Reflect the angle into the first quadrant [0° , 90°].
3. Run the CORDIC core there.
4. Apply the final sign correction to sin and cos.

That is why BAM16 is convenient: the top two bits already tell you the quadrant.

19.5 Rotation Mode for Sine and Cosine

Pseudocode using Q1.15 vectors and BAM16 angles:

```
x = K_Q15;
y = 0;
z = reduced_angle;      // now in first quadrant, 0..90 degrees

for (i = 0; i < 12; i++) {
    if (z >= 0) {
        x_new = x - (y >> i);
        y_new = y + (x >> i);
        z_new = z - atan_table[i];
    } else {
        x_new = x + (y >> i);
        y_new = y - (x >> i);
        z_new = z + atan_table[i];
    }
    x = x_new;
    y = y_new;
    z = z_new;
}
```

After the loop:

- x is the first-quadrant cosine
- y is the first-quadrant sine

Then apply quadrant signs:

Quadrant	Reduced angle	cos sign	sin sign
0	angle	+	+
1	180°-angle	-	+
2	angle-180°	-	-
3	360°-angle	+	-

19.5.1 Example Values

With 12 stages and $K_{Q15} = 0x4DB9$:

- 30° (0x1555) gives about $\cos \approx 0.8660$, $\sin \approx 0.5001$
- 45° (0x2000) gives about $\cos \approx 0.7070$, $\sin \approx 0.7072$
- 300° (0xD555) gives about $\cos \approx 0.5001$, $\sin \approx -0.8660$

That is easily good enough for control loops, DDS, UI graphics, and many sensor applications on AVR.

19.6 Complete Integer Reference Implementation

The full reference implementation is in `ch19_cordic/src/cordic_q15.c`. It is written as plain integer C so you can:

- compile it directly with `avr-gcc`,
- inspect the generated AVR code with `avr-objdump`,
- use it as a correctness reference before hand-optimizing assembly.

19.6.1 Core Routine

```
#include <stdint.h>

static const uint16_t cordic_atan_bam16[12] = {
    0x2000, 0x12E4, 0x09FB, 0x0511,
    0x028B, 0x0146, 0x00A3, 0x0051,
    0x0029, 0x0014, 0x000A, 0x0005
};

#define CORDIC_K_Q15 ((int16_t)0x4DB9)

void cordic_sincos_bam16(uint16_t angle, int16_t *cos_q15, int16_t *sin_q15)
{
    uint8_t quadrant = angle >> 14;
    int16_t z;
    int8_t cos_sign = 1;
    int8_t sin_sign = 1;

    switch (quadrant) {
    case 0:
        z = (int16_t)angle;
        break;
    case 1:
        z = (int16_t)(0x8000u - angle);
        cos_sign = -1;
        break;
    
```

```

    case 2:
        z = (int16_t)(angle - 0x8000u);
        cos_sign = -1;
        sin_sign = -1;
        break;
    default:
        z = (int16_t)(0u - angle);
        sin_sign = -1;
        break;
}

int16_t x = CORDIC_K_Q15;
int16_t y = 0;

for (uint8_t i = 0; i < 12; i++) {
    int16_t x_shift = (int16_t)(x >> i);
    int16_t y_shift = (int16_t)(y >> i);
    int16_t x_new;
    int16_t y_new;

    if (z >= 0) {
        x_new = x - y_shift;
        y_new = y + x_shift;
        z -= (int16_t)cordic_atan_bam16[i];
    } else {
        x_new = x + y_shift;
        y_new = y - x_shift;
        z += (int16_t)cordic_atan_bam16[i];
    }

    x = x_new;
    y = y_new;
}

*cos_q15 = (cos_sign > 0) ? x : (int16_t)-x;
*sin_q15 = (sin_sign > 0) ? y : (int16_t)-y;
}

```

19.6.2 Build and Inspect

```

avr-gcc -mmcu=attiny3217 -O2 -c ch19_cordic/src/cordic_q15.c
avr-objdump -d cordic_q15.o

```

This is the right first step on AVR: get the integer algorithm correct, inspect the generated code, then decide whether a hand-written assembly version is worth the maintenance cost.

19.7 How This Maps to AVR Assembly

At the algorithm level, CORDIC is simple. At the assembly level, the main cost is the **variable arithmetic shift** by *i*.

For a general loop, *x >> i* and *y >> i* mean:

1. copy the 16-bit value to a temporary pair,
2. shift it right i times using ASR on the high byte and ROR on the low byte,
3. use the shifted temporary in the add/subtract step.

That is why many hand-tuned AVR CORDIC routines are either:

- partially unrolled, or
- fully unrolled.

When i is a compile-time constant, the shifts become fixed instruction sequences and the loop-control overhead disappears.

19.7.1 Example: One Unrolled Stage ($i = 3$)

Suppose:

- R25:R24 = x
- R23:R22 = y
- R21:R20 = z

One positive-rotation stage for $i = 3$ looks like:

```
/* temp_y = y >> 3 (signed shift) */
movw  r19, r22          /* r19:r18 = y copy          */
asr   r19
ror   r18
asr   r19
ror   r18
asr   r19
ror   r18

/* x = x - temp_y */
sub   r24, r18
sbc   r25, r19

/* temp_x = old x >> 3
 * In a real implementation, preserve old x before modifying it. */
```

The mechanics are exactly the same as any other multi-byte signed shift on AVR: ASR the high byte, ROR the low byte, repeat as many times as needed.

19.7.2 Practical Register Plan

For a hand-written AVR version, a reasonable plan is:

```
R25:R24  x
R23:R22  y
R21:R20  z
R19:R18  temp shifted copy
R17      stage counter i
R16      scratch / table byte
Z        pointer into atan table
```

This is workable, but you can see why register pressure gets real quickly. CORDIC is a good example of a routine where AVR assembly quality depends as much on register planning as on the arithmetic itself.

19.8 Vectoring Mode: Magnitude and Angle

Vectoring mode starts with an input vector (x , y) and rotates it toward the x-axis:

```
for (i = 0; i < 12; i++) {
    if (y >= 0) {
        x_new = x + (y >> i);
        y_new = y - (x >> i);
        z_new = z + atan_table[i];
    } else {
        x_new = x - (y >> i);
        y_new = y + (x >> i);
        z_new = z - atan_table[i];
    }
}
```

After the last stage:

- z approximates $\text{atan2}(y, x)$
- x approximates $K * \text{magnitude}$
- y is near zero

The magnitude is scaled by K , so you either:

- compensate afterwards, or
- calibrate the rest of the application around the scaled value.

19.8.1 Where AVR Uses This

Vectoring mode is especially useful for:

- phase estimation from I/Q samples,
- converting Cartesian sensor axes to angle,
- simple magnitude estimation without a square root.

This is where CORDIC often beats ad-hoc math on AVR: one fixed-iteration loop replaces both $\text{atan2}()$ and $\text{sqrt}(x*x + y*y)$.

19.9 Example 1: DDS / NCO Sine Generator

A common AVR use is a numerically controlled oscillator:

1. keep a 16-bit phase accumulator,
2. add a phase step each sample period,
3. feed the phase into CORDIC rotation mode,
4. send the sine result to a PWM or DAC.

```
phase += phase_step;
cordic_sincos_bam16(phase, &cos_q15, &sin_q15);
pwm_value = (uint8_t)((sin_q15 >> 8) + 128);
```

Why CORDIC can make sense here:

- no large sine table in flash,
- easy frequency control through `phase_step`,
- constant run time every sample.

If you need the very highest sample rate, a LUT is still usually faster. But if flash is tight and you want both sine and cosine from the same phase input, CORDIC is an attractive compromise.

19.10 Example 2: Tilt Angle from Two Axes

Suppose an AVR reads two already-filtered accelerometer channels:

```
x_axis = forward/back tilt component
y_axis = vertical tilt component
```

You want an angle estimate:

```
tilt = atan2(x_axis, y_axis)
```

On a desktop you would call `atan2f()`. On AVR, vectoring CORDIC gives you a fixed-time integer path instead:

1. load (x, y) into Q1.15,
2. run vectoring mode,
3. interpret z in BAM16,
4. convert BAM16 to degrees only if the UI needs it.

That avoids floating-point entirely and keeps the control-loop representation in a format that is cheap to update.

19.11 Accuracy, Flash, and Cycle Tradeoffs

19.11.1 Iteration Count

More stages mean more accuracy and more cycles.

Typical choices on AVR:

- 8 stages: quick, coarse
- 12 stages: good practical middle ground
- 16 stages: near the limit of 16-bit usefulness

With Q1.15 outputs, 12 stages is usually enough because the result resolution is already limited by the 16-bit format.

19.11.2 LUT vs CORDIC

Rough comparison:

Method	Flash	Cycles	Strength
256-byte sine LUT	moderate	very low	fastest for sine only
CORDIC (12 stages)	small table + code	moderate, fixed	sin/cos/atan2/magnitude from one core
libm software float	large	high, variable-looking at source level	easiest source code

The main advantage of CORDIC is not “fewest cycles.” It is that one deterministic integer core can cover several mathematically different jobs.

19.12 Common Pitfalls

19.12.1 Pitfall 1 — Forgetting Quadrant Reduction

Basic rotation mode does not converge over the full 360° range.

If you feed 210° directly into the core without reducing it first, the result will be wrong even if the stage math itself is perfect.

19.12.2 Pitfall 2 — Using Logical Right Shift on Signed Values

The shifts inside CORDIC are signed for x and y.

On AVR assembly:

```
asr    r25
ror    r24
```

is correct for a signed 16-bit right shift.

Using:

```
lsr    r25
ror    r24
```

would break negative vector components.

19.12.3 Pitfall 3 — Ignoring the Gain

CORDIC results are scaled. If you forget to preload with K in rotation mode, your sine/cosine amplitude will be too large. If you forget the gain in vectoring mode, your magnitude estimate will be off by a constant factor.

19.12.4 Pitfall 4 — Overflow in Narrow Formats

Q1.15 is ideal for normalized vectors, but not for arbitrary raw values.

If your input vector can exceed the $[-1, +1]$ range, normalize it or move to a wider fixed-point format before running CORDIC.

19.12.5 Pitfall 5 — Hand-Writing the Loop Before Proving Correctness

CORDIC is simple enough to look obvious and subtle enough to get wrong.

On AVR, prove the math first with a small integer C reference, then compare the assembly against it stage by stage. That is faster than debugging a broken fully-unrolled assembly version from scratch.

19.13 Summary

CORDIC on AVR:

- uses shifts + adds/subtracts + compare + small atan table
- no floating point required
- fixed iteration count → deterministic timing

Useful formats:

- x,y = Q1.15 signed fractions

```
angle    = BAM16 (65536 units per turn)
K_Q15    = 0x4DB9 for 12-stage rotation mode
```

Rotation mode:

```
input    = angle
output   = cos(theta), sin(theta)
```

Vectoring mode:

```
input    = x,y vector
output   = angle and scaled magnitude
```

Main AVR tradeoff:

```
slower than a small LUT for sine only
more flexible than a LUT when you need sin/cos/atan2/magnitude
```

19.14 Exercises

1. Implement an 8-stage CORDIC sine/cosine routine and compare its maximum error against the 12-stage version. How many cycles do you save?
 2. Store the BAM16 atan table in flash and read it with LPM through Z. Compare the flash cost and cycle cost against hard-coding the constants in a fully unrolled routine.
 3. Write an assembly helper that performs a signed 16-bit arithmetic right shift by a variable count in R17. Use ASR/ROR in a loop. How many cycles does it take for counts 0, 1, 4, and 8?
 4. Modify the C reference implementation so that it returns only sine, not cosine. Can the compiler eliminate any useful work at -O2, or does the shared CORDIC state keep both values effectively “free” once the loop runs?
 5. Implement vectoring mode for inputs already normalized to Q1.15 and test it on (0.6, 0.8). What BAM16 angle do you get, and how close is the scaled magnitude to K?
 6. Compare three ways to generate a PWM sine on AVR: a. 256-entry flash LUT b. 12-stage CORDIC c. `sinf()` from software floating point Measure flash usage with `avr-size` and inspect the call paths with `avr-objdump -d`.
-

Next: Appendices

Chapter 20

ATtiny3217 Register Reference

Quick lookup for registers used throughout this book. The ATtiny3217 is an AVRXMEGA3 device. All peripheral registers are memory-mapped and require LDS/STS unless noted otherwise.

VPORT exception: Virtual PORT registers at 0x0000–0x001F are within the I/O space and can be accessed with IN/OUT, SBI/CBI, SBIS/SBIC in a single cycle. All other peripheral registers require LDS/STS.


CPU registers: SREG (0x003F), SPH (0x003E), SPL (0x003D) are accessible with IN/OUT using their I/O addresses directly (no +0x20 offset needed on AVRXMEGA3 — these addresses are already the memory-mapped addresses, and the I/O address equals the memory address for registers at 0x0000–0x003F).

20.1 CPU Registers

Register	Mem Addr	Access	Description
SREG	0x003F	IN/OUT	Status Register (I T H S V N Z C)
SPH	0x003E	IN/OUT	Stack Pointer High
SPL	0x003D	IN/OUT	Stack Pointer Low
CCP	0x0034	LDS/STS	Configuration Change Protection

SREG, SPH, and SPL sit within the range 0x0000–0x003F and are accessible via IN/OUT using the address directly. The CCP register requires LDS/STS.

20.1.1 SREG Bit Layout



20.1.2 Configuration Change Protection (CCP)

Certain registers are protected against accidental writes. Before writing to a protected register, write the appropriate key to CCP. The protected register must then be written within 4 CPU cycles.

Key	Value	Protects
IOREG	0xD8	I/O registers (e.g. CLKCTRL, BOD, OSC settings)
SPM	0x9D	NVM (flash/EEPROM) write commands

Example — disable clock prescaler to run at 20 MHz:

```
ldi    r16, 0xD8          ; CCP key for I/O register changes
sts    0x0034, r16        ; write to CCP
sts    0x0061, r1         ; write 0 to CLKCTRL_MCLKCTRLB (clear PEN)
                          ; must complete within 4 cycles of CCP write
```

20.2 CLKCTRL — Clock Controller

Base address: 0x0060. All registers require LDS/STS and CCP protection where noted.

Register	Mem Addr	CCP?	Description
CLKCTRL_MCLKCTRLA	0x0060	Yes	Main clock source select
CLKCTRL_MCLKCTRLB	0x0061	Yes	Main clock prescaler
CLKCTRL_MCLKLOCK	0x0062	Yes	Lock register (prevents further changes)
CLKCTRL_MCLKSTATUS	0x0063	No	Status (read-only)

20.2.1 MCLKCTRLA — Main Clock Source Select

Bit: 7 6:2 1:0
 CLKOUT — CLKSEL

CLKSEL	Source
0x0	20 MHz internal oscillator ← factory default
0x1	32 KHz internal ULP oscillator
0x2	32.768 KHz external crystal (XOSC32K)
0x3	External clock on CLKI pin

Setting CLKOUT (bit 7) routes the main clock to the CLKOUT pin.

20.2.2 MCLKCTRLB — Main Clock Prescaler

Bit: 7:5 4:1 0
 — PDIV PEN

PEN Prescaler enable (1 = prescaler active, 0 = disabled)

PDIV	Division factor (when PEN=1)
0x0	/2
0x1	/4
0x2	/8
0x3	/16
0x4	/32
0x5	/64

```

0x6  /6  ← factory default (20 MHz / 6 = 3.333 MHz)
0x7  /10
0x8  /12
0x9  /24
0xA  /48

```

Factory default: PDIV=0x6 (divide by 6), PEN=1 → **3.333 MHz** system clock.

To run at full **20 MHz**: write 0x00 to MCLKCTRLB (clears PEN) — requires CCP key 0xD8 first. To run at **10 MHz**: write PEN=1, PDIV=0x1 (divide by 2).

20.2.3 MCLKSTATUS — Status (read-only)

```

Bit:  7      6      5      4      3      2      1      0
      EXTS   XOSC32KS SOSC   OSC32KS OSC20MS -      -      SOSC

```

Bit name	Description
OSC20MS	1 = 20 MHz oscillator is stable and running
OSC32KS	1 = 32 KHz oscillator is stable
SOSC	1 = System oscillator is stable
EXTS	1 = External clock source is active

20.3 GPIO Registers

The AVRXMEGA3 GPIO model is completely different from classic AVR. Each port has a full set of direction/output/input/control registers at a fixed base address. Atomic set/clear/toggle registers eliminate the need for read-modify-write sequences.

20.3.1 Port Base Addresses

Port	Base Addr
PORTA	0x0400
PORTB	0x0420
PORTC	0x0440

20.3.2 Per-Port Register Map (offset from base)

Offset	Register	Description
+0x00	PORTx_DIR	Data direction register (1=output, 0=input)
+0x01	PORTx_DIRSET	Write 1 to set direction bits (atomic)
+0x02	PORTx_DIRCLR	Write 1 to clear direction bits (atomic)
+0x03	PORTx_DIRTGL	Write 1 to toggle direction bits (atomic)
+0x04	PORTx_OUT	Output value register
+0x05	PORTx_OUTSET	Write 1 to set output bits (atomic)
+0x06	PORTx_OUTCLR	Write 1 to clear output bits (atomic)
+0x07	PORTx_OUTTGL	Write 1 to toggle output bits (atomic)
+0x08	PORTx_IN	Input value (reads current pin state)
+0x09	PORTx_INTFLAGS	Pin-change interrupt flags (W1C)
+0x0A	PORTx_PORTCTRL	Port control register
+0x10	PORTx_PIN0CTRL	Pin 0 config (pull-up, invert, ISC)

+0x11	PORTx_PIN1CTRL	Pin 1 config
+0x12	PORTx_PIN2CTRL	Pin 2 config
+0x13	PORTx_PIN3CTRL	Pin 3 config
+0x14	PORTx_PIN4CTRL	Pin 4 config
+0x15	PORTx_PIN5CTRL	Pin 5 config
+0x16	PORTx_PIN6CTRL	Pin 6 config
+0x17	PORTx_PIN7CTRL	Pin 7 config

20.3.3 PORTA Absolute Addresses

Register	Mem Addr	Description
PORTA_DIR	0x0400	Direction
PORTA_DIRSET	0x0401	Set direction (atomic)
PORTA_DIRCLR	0x0402	Clear direction (atomic)
PORTA_DIRTGL	0x0403	Toggle direction (atomic)
PORTA_OUT	0x0404	Output value
PORTA_OUTSET	0x0405	Set output (atomic)
PORTA_OUTCLR	0x0406	Clear output (atomic)
PORTA_OUTTGL	0x0407	Toggle output (atomic)
PORTA_IN	0x0408	Input value
PORTA_INTFLAGS	0x0409	Interrupt flags
PORTA_PORTCTRL	0x040A	Port control
PORTA_PIN0CTRL	0x0410	PA0 config (UPDI pin – do not reconfigure)
PORTA_PIN1CTRL	0x0411	PA1 config (SPI MOSI default)
PORTA_PIN2CTRL	0x0412	PA2 config (SPI MISO default)
PORTA_PIN3CTRL	0x0413	PA3 config (LED0 – Curiosity Nano)
PORTA_PIN4CTRL	0x0414	PA4 config (SPI SS default)
PORTA_PIN5CTRL	0x0415	PA5 config
PORTA_PIN6CTRL	0x0416	PA6 config
PORTA_PIN7CTRL	0x0417	PA7 config

20.3.4 PORTB Absolute Addresses

Register	Mem Addr	Description
PORTB_DIR	0x0420	Direction
PORTB_DIRSET	0x0421	Set direction (atomic)
PORTB_DIRCLR	0x0422	Clear direction (atomic)
PORTB_DIRTGL	0x0423	Toggle direction (atomic)
PORTB_OUT	0x0424	Output value
PORTB_OUTSET	0x0425	Set output (atomic)
PORTB_OUTCLR	0x0426	Clear output (atomic)
PORTB_OUTTGL	0x0427	Toggle output (atomic)
PORTB_IN	0x0428	Input value
PORTB_INTFLAGS	0x0429	Interrupt flags
PORTB_PORTCTRL	0x042A	Port control
PORTB_PIN0CTRL	0x0430	PB0 config (TWI SDA default)
PORTB_PIN1CTRL	0x0431	PB1 config (TWI SCL default)
PORTB_PIN2CTRL	0x0432	PB2 config (USART0 TX – Curiosity Nano)
PORTB_PIN3CTRL	0x0433	PB3 config (USART0 RX – Curiosity Nano)
PORTB_PIN4CTRL	0x0434	PB4 config
PORTB_PIN5CTRL	0x0435	PB5 config
PORTB_PIN6CTRL	0x0436	PB6 config


PORTB_PIN7CTRL 0x0437 PB7 config

20.3.5 PORTC Absolute Addresses

Register	Mem Addr	Description
PORTC_DIR	0x0440	Direction
PORTC_DIRSET	0x0441	Set direction (atomic)
PORTC_DIRCLR	0x0442	Clear direction (atomic)
PORTC_DIRTGL	0x0443	Toggle direction (atomic)
PORTC_OUT	0x0444	Output value
PORTC_OUTSET	0x0445	Set output (atomic)
PORTC_OUTCLR	0x0446	Clear output (atomic)
PORTC_OUTTGL	0x0447	Toggle output (atomic)
PORTC_IN	0x0448	Input value
PORTC_INTFLAGS	0x0449	Interrupt flags
PORTC_PORTCTRL	0x044A	Port control
PORTC_PIN0CTRL	0x0450	PC0 config
PORTC_PIN1CTRL	0x0451	PC1 config
PORTC_PIN2CTRL	0x0452	PC2 config
PORTC_PIN3CTRL	0x0453	PC3 config
PORTC_PIN4CTRL	0x0454	PC4 config
PORTC_PIN5CTRL	0x0455	PC5 config
PORTC_PIN6CTRL	0x0456	PC6 config (not present on 20-pin package)
PORTC_PIN7CTRL	0x0457	PC7 config (not present on 20-pin package)

20.3.6 PINnCTRL Bit Layout

Each PORTx_PINnCTRL register configures one pin independently.



ISC[2:0] Input Sense Configuration

0x0	INTDISABLE	No interrupt; digital input buffer enabled
0x1	BOTHEDGES	Interrupt on both rising and falling edges
0x2	RISING	Interrupt on rising edge only
0x3	FALLING	Interrupt on falling edge only
0x4	INPUT_DISABLE	Digital input buffer disabled (reduces power)
0x5	LEVEL	Interrupt on low level

20.3.7 Virtual PORT Registers (VPORT) — Fast I/O

VPORTs mirror the key PORT registers into I/O space (0x0000–0x001F). They can be accessed with IN/OUT (1 cycle), SBI/CBI, and SBIS/SBIC. Writing to VPORT_IN toggles the corresponding output bit (same as OUTTGL).

Register	Addr	Mirrors
VPORTA_DIR	0x0000	PORTA_DIR (1=output)
VPORTA_OUT	0x0001	PORTA_OUT

VPORTA_IN	0x0002	PORTA_IN	(write 1 to toggle output)
VPORTA_INTFLAGS	0x0003	PORTA_INTFLAGS	
VPORTB_DIR	0x0004	PORTB_DIR	
VPORTB_OUT	0x0005	PORTB_OUT	
VPORTB_IN	0x0006	PORTB_IN	
VPORTB_INTFLAGS	0x0007	PORTB_INTFLAGS	
VPORTC_DIR	0x0008	PORTC_DIR	
VPORTC_OUT	0x0009	PORTC_OUT	
VPORTC_IN	0x000A	PORTC_IN	
VPORTC_INTFLAGS	0x000B	PORTC_INTFLAGS	

Use VPORT registers with IN/OUT for single-cycle access. Use the full PORT registers with LDS/STS when you need atomic set/clear/toggle or pin config.

20.4 Curiosity Nano Pin Assignments

Signal	Port Pin	Mem Addr (PINnCTRL)	Notes
LED0	PA3	0x0413	Active low – drive LOW to light
USART0 TX	PB2	0x0432	To on-board debugger virtual COM port
USART0 RX	PB3	0x0433	From on-board debugger virtual COM port
UPDI	PA0	0x0410	Programming/debug – do not reconfigure

20.4.1 LED and Button Quick Reference

LED0 on (light): STS PORTA_OUTCLR, (1<<3) ; clear PA3 (active low)
 LED0 off (dark): STS PORTA_OUTSET, (1<<3) ; set PA3
 LED0 toggle: STS PORTA_OUTTGL, (1<<3) ; toggle PA3

The ATtiny3217 Curiosity Nano does not provide a dedicated user pushbutton. For button examples, wire an external active-low button to a free GPIO and enable that pin's internal pull-up.

20.5 TCA0 — Timer/Counter Type A (16-bit)

Base address: 0x0A00 (single-slope mode, the default). TCA0 is a general-purpose 16-bit timer supporting normal, CTC, and PWM modes.

Register	Mem Addr	Description
TCA0_SINGLE_CTRLA	0x0A00	Prescaler select (CLKSEL), enable (ENABLE)
TCA0_SINGLE_CTRLB	0x0A01	Waveform mode (WGMODE), compare enables
TCA0_SINGLE_CTRLC	0x0A02	Force compare outputs
TCA0_SINGLE_CTRLD	0x0A03	Split mode enable
TCA0_SINGLE_CTRLCLR	0x0A04	Control E clear
TCA0_SINGLE_CTRLSET	0x0A05	Control E set
TCA0_SINGLE_INTCTRL	0x0A06	Interrupt enable (OVF, CMP0, CMP1, CMP2)
TCA0_SINGLE_INTFLAGS	0x0A07	Interrupt flags (W1C)

TCA0_SINGLE_EVCTRL	0x0A09	Event control
TCA0_SINGLE_DBGCTRL	0x0A0E	Debug control
TCA0_SINGLE_TEMP	0x0A0F	Temp register for 16-bit read/write
TCA0_SINGLE_CNT	0x0A20	Counter value (16-bit, write low byte first)
TCA0_SINGLE_PER	0x0A26	Period / TOP value (16-bit)
TCA0_SINGLE_CMP0	0x0A28	Compare/capture 0 (16-bit)
TCA0_SINGLE_CMP1	0x0A2A	Compare/capture 1 (16-bit)
TCA0_SINGLE_CMP2	0x0A2C	Compare/capture 2 (16-bit)
TCA0_SINGLE_PERBUF	0x0A36	Period buffer (double-buffered)
TCA0_SINGLE_CMP0BUF	0x0A38	Compare 0 buffer
TCA0_SINGLE_CMP1BUF	0x0A3A	Compare 1 buffer
TCA0_SINGLE_CMP2BUF	0x0A3C	Compare 2 buffer

20.5.1 CTRLA — Prescaler and Enable

Bit: 7:4 3:1 0
 — CLKSEL ENABLE

CLKSEL Prescaler

0x0	DIV1	(no prescaling)
0x1	DIV2	
0x2	DIV4	
0x3	DIV8	
0x4	DIV16	
0x5	DIV64	
0x6	DIV256	
0x7	DIV1024	

Example — 1 Hz overflow at 3.333 MHz with DIV64 (prescaled clock = 52,083 Hz, PER = 52082):

```
ldi r16, 0x0B           ; CLKSEL=DIV64 (bits 3:1 = 101), ENABLE=1
sts 0x0A00, r16         ; TCA0_SINGLE_CTRLA
ldi r16, lo8(52082)
sts 0x0A26, r16         ; TCA0_SINGLE_PER low byte
ldi r16, hi8(52082)
sts 0x0A27, r16         ; TCA0_SINGLE_PER high byte
ldi r16, 0x01           ; OVF interrupt enable
sts 0x0A06, r16         ; TCA0_SINGLE_INTCTRL
```

20.5.2 CTRLB — Waveform Mode

WGMODE[2:0] (bits 2:0) Mode

0x0	Normal (count up, overflow at PER)
0x1	Frequency (toggle W0, reset at CMP0)
0x2	Single-slope PWM
0x5	Dual-slope PWM (overflow at BOTTOM and TOP)
0x6	Dual-slope PWM (overflow at TOP only)
0x7	Dual-slope PWM (overflow at BOTTOM only)

20.6 TCB0 / TCB1 — Timer/Counter Type B (16-bit)

Type B timers are 16-bit, single-channel timers suited for periodic interrupts, input capture, and PWM. TCB0 and TCB1 share the same register layout.

TCB0 base: 0x0A40

TCB1 base: 0x0A50

Offset	Register	Description
+0x00	TCBn_CTRLA	Clock select, enable
+0x01	TCBn_CTRLB	Compare/capture mode, output enable
+0x04	TCBn_EVCTRL	Event control
+0x05	TCBn_INTCTRL	CAPT interrupt enable (bit 0)
+0x06	TCBn_INTFLAGS	CAPT interrupt flag (bit 0, W1C)
+0x07	TCBn_STATUS	RUN flag (bit 0, read-only)
+0x08	TCBn_DBGCTRL	Debug control
+0x09	TCBn_TEMP	Temp register for 16-bit read/write
+0x0A	TCBn_CNTL	Counter low byte
+0x0B	TCBn_CNTH	Counter high byte
+0x0C	TCBn_CCMP_L	Compare/capture low byte
+0x0D	TCBn_CCMP_H	Compare/capture high byte

20.6.1 CTRLA — Clock Select and Enable

Bit: 7:3 2:1 0
 — CLKSEL ENABLE

CLKSEL Clock source

0x0	CLKPER	(same as system clock)
0x1	CLKPER/2	
0x2	TCA0	(TCB clocked from TCA0 overflow)

20.6.2 CTRLB — Mode (CMODE)

CMODE[2:0] (bits 2:0) Mode

0x0	Periodic Interrupt ← default, most common
0x1	Timeout Check
0x2	Input Capture on Event
0x3	Input Capture Frequency Measurement
0x4	Input Capture Pulse-Width Measurement
0x5	Input Capture Freq+PW Measurement
0x6	Single Shot
0x7	8-bit PWM

In Periodic Interrupt mode (CMODE=0), the counter counts up to CCMP, fires the CAPT interrupt, then resets to zero. CCMP sets the period (TOP).

Example — 1 kHz periodic interrupt at 3.333 MHz:

```
ldi r16, lo8(3332)      ; CCMP = (f_cpu / f_irq) - 1 = 3332
sts 0x0A4C, r16         ; TCB0_CCMP_L
ldi r16, hi8(3332)
sts 0x0A4D, r16         ; TCB0_CCMP_H
ldi r16, 0x01           ; CAPT interrupt enable
```

```

sts 0x0A45, r16      ; TCB0_INTCTRL
ldi r16, 0x01        ; CLKSEL=CLKPER, ENABLE=1
sts 0x0A40, r16      ; TCB0_CTRLA

```

20.7 USART0

Base address: 0x0800. On the Curiosity Nano board, USART0 is connected to the on-board debugger (PB2=TX, PB3=RX), providing a virtual COM port over USB.

Register	Mem Addr	Description
USART0_RXDATAH	0x0800	Receive data (low byte, 8-bit data here)
USART0_RXDATAH	0x0801	Receive data high + error flags
USART0_TXDATAH	0x0802	Transmit data (write byte here to send)
USART0_TXDATAH	0x0803	Transmit data high (9-bit mode only)
USART0_STATUS	0x0804	Status flags
USART0_CTRLA	0x0805	Interrupt enables
USART0_CTRLB	0x0806	Receiver/transmitter enable
USART0_CTRLC	0x0807	Frame format (default = 8N1)
USART0_BAUDL	0x0808	Baud rate register low byte
USART0_BAUDH	0x0809	Baud rate register high byte
USART0_CTRLD	0x080A	Auto-baud, IRCOM
USART0_DBGCTRL	0x080B	Debug control
USART0_EVCTRL	0x080C	Event control
USART0_TXPLCTRL	0x080D	TX pin LINAUTO
USART0_RXPLCTRL	0x080E	RX pin LINAUTO

20.7.1 STATUS Bits

Bit:	7	6	5	4	3	2	1:0
	RXCIF	TXCIF	DREIF	FERR	BUFOVF	PERR	—

RXCIF Receive Complete Interrupt Flag (1 = unread data in RXDATA)
 TXCIF Transmit Complete Interrupt Flag (1 = TX shift register empty)
 DREIF Data Register Empty Flag (1 = OK to write next byte to TXDATA)
 FERR Frame Error
 BUFOVF Buffer Overflow
 PERR Parity Error

20.7.2 CTRLA — Interrupt Enables

Bit:	7	6	5	4:2	1	0
	RXCIE	TXCIE	DREIE	—	LBME	ABEIE

RXCIE Receive Complete IE
 TXCIE Transmit Complete IE
 DREIE Data Register Empty IE (triggers when ready for next transmit byte)

20.7.3 CTRLB — Enable Bits

Bit:	7	6	5:4	3	2	1:0
	RXEN	TXEN	MPCM	RXMODE	ODME	—

RXEN 1 = enable receiver
 TXEN 1 = enable transmitter
 RXMODE 0x0=NORMAL, 0x1=CLK2X, 0x2=GENAUTO, 0x3=LINAUTO

20.7.4 CTRLC — Frame Format

Default value: 0x03 = 8N1 (8 data bits, no parity, 1 stop bit)

Bit:	7:6	5:4	3	2:0
	CMODE	PMODE	SBMODE	CHSIZE

CMODE 0x0=ASYNCHRONOUS, 0x1=SYNCHRONOUS, 0x2=IRCOM, 0x3=MSPI
 PMODE 0x0=DISABLED, 0x2=EVEN, 0x3=ODD
 SBMODE 0=1 stop bit, 1=2 stop bits
 CHSIZE 0x0=5-bit, 0x1=6-bit, 0x2=7-bit, 0x3=8-bit, 0x4-0x6=reserved, 0x7=9-bit

20.7.5 Baud Rate Formula and Common Values

Normal asynchronous mode:

$$\begin{aligned}
 \text{BAUD_REG} &= (\text{f_cpu} \times 64) / (16 \times \text{baud_rate}) \\
 &= (\text{f_cpu} \times 4) / \text{baud_rate}
 \end{aligned}$$

Baud Rate	f_cpu = 3.333 MHz BAUD_REG (hex)	f_cpu = 20 MHz BAUD_REG (hex)
9600	0x056D (1389)	0x2098 (8344)
19200	0x02B7 (694)	0x104C (4172)
38400	0x015B (347)	0x0826 (2086)
57600	0x00E2 (226)	0x0571 (1393)
115200	0x0074 (116)	0x02B9 (697)

20.8 SPI0

Base address: 0x08C0. Default pin assignments: MOSI=PA1, MISO=PA2, SCK=PA3, SS=PA4.

Register	Mem Addr	Description
SPI0_CTRLA	0x08C0	Control A: DORD, MASTER, CLK2X, PRESC[1:0], ENABLE
SPI0_CTRLB	0x08C1	Control B: BUFEN, SSD, MODE[1:0]
SPI0_INTCTRL	0x08C2	Interrupt enables: IE, SSIE, TXCIE, RXCIE
SPI0_INTFLAGS	0x08C3	Interrupt flags: IF(7), WRCOL(0)
SPI0_DATA	0x08C4	TX/RX data register

20.8.1 CTRLA Bits

Bit:	7	6	5	4:3	2	1	0
	DORD	MASTER	CLK2X	PRESC	—	—	ENABLE

DORD 0=MSB first, 1=LSB first
 MASTER 0=slave, 1=master
 CLK2X 1=double clock speed

PRESC 0x0=DIV4, 0x1=DIV16, 0x2=DIV64, 0x3=DIV128
 ENABLE 1=enable SPI

20.8.2 SPI Clock Rates

PRESC	CLK2X=0	CLK2X=1
0x0	f_cpu / 4	f_cpu / 2
0x1	f_cpu / 16	f_cpu / 8
0x2	f_cpu / 64	f_cpu / 32
0x3	f_cpu / 128	f_cpu / 64

20.9 TWI0 — Two-Wire Interface (I2C)

Base address: 0x08A0. Default pins: SDA=PB0, SCL=PB1. Requires external 4.7 kΩ pull-up resistors on SDA and SCL lines.

Register	Mem Addr	Description
TWI0_CTRLA	0x08A0	SDA hold time, fast-mode plus enable
TWI0_DUALCTRL	0x08A1	Dual-mode control
TWI0_MCTRLA	0x08A2	Master: RIEN, WIEN, QCEN, TIMEOUT, SMEN, ENABLE
TWI0_MCTRLB	0x08A3	Master: FLUSH, ACKACT, MCMD[1:0]
TWI0_MSTATUS	0x08A4	Master status flags
TWI0_MBAUD	0x08A5	Master baud rate
TWI0_MADDR	0x08A6	Master address (write triggers START condition)
TWI0_MDATA	0x08A7	Master data register (read/write)
TWI0_SCTRLA	0x08A8	Slave: DIEN, APIEN, PIEN, PMEN, SMEN, ENABLE
TWI0_SCTRLB	0x08A9	Slave: ACKACT, SCMD[1:0]
TWI0_SSTATUS	0x08AA	Slave status flags
TWI0_SADDR	0x08AB	Slave address
TWI0_SDATA	0x08AC	Slave data
TWI0_SADDRMASK	0x08AD	Slave address mask

20.9.1 MSTATUS Bits

Bit: 7 6 5 4 3 2 1:0
 RIF WIF CLKHOLD RXACK ARBLOST BUSERR BUSSTATE

RIF Read Interrupt Flag
 WIF Write Interrupt Flag
 CLKHOLD Clock hold (SCL stretched by master)
 RXACK Received ACK (0=ACK received, 1=NACK)
 ARBLOST Arbitration lost
 BUSERR Bus error
 BUSSTATE 0x0=UNKNOWN, 0x1=IDLE, 0x2=OWNER, 0x3=BUSY

20.9.2 MCTRLB — Master Command

MCMD[1:0]	Command
0x0	NOACT No action

0x1	REPSTART	Issue repeated START
0x2	RECVTRANS	Continue (ACK) / send byte
0x3	STOP	Issue STOP condition

20.9.3 Baud Rate Formula

$BAUD = (f_{cpu} / f_{scl} - 10) / 2$ (for f_{cpu} in Hz, f_{scl} in Hz)

100 kHz at 3.333 MHz: $BAUD = (3333333 / 100000 - 10) / 2 \approx 12$
 400 kHz at 3.333 MHz: $BAUD = (3333333 / 400000 - 10) / 2 \approx 2$ (use fast-mode+)
 100 kHz at 20 MHz: $BAUD = (20000000 / 100000 - 10) / 2 = 95$
 400 kHz at 20 MHz: $BAUD = (20000000 / 400000 - 10) / 2 = 20$

20.10 ADC0

Base address: 0x0600. The ATtiny3217 ADC is 10-bit (or 8-bit selectable).

Register	Mem Addr	Description
ADC0_CTRLA	0x0600	Enable, resolution, free-run, left-adjust
ADC0_CTRLB	0x0601	Sample accumulation (SAMPNUM)
ADC0_CTRLC	0x0602	Prescaler, reference select, sample capacitor
ADC0_CTRLD	0x0603	Init delay, auto-sample delay
ADC0_CTRLF	0x0604	Window comparator mode
ADC0_SAMPCTRL	0x0605	Sample length extension
ADC0_MUXPOS	0x0606	Input mux (selects channel)
ADC0_COMMAND	0x0608	Write bit 0 (STCONV) to start conversion
ADC0_EVCTRL	0x0609	Event trigger enable
ADC0_INTCTRL	0x060A	RESRDY(0) and WCMP(1) interrupt enables
ADC0_INTFLAGS	0x060B	Interrupt flags (W1C)
ADC0_RESL	0x060C	Result low byte
ADC0_RES	0x060D	Result high byte
ADC0_WINLT	0x060E	Window comparator low threshold
ADC0_WINHT	0x0610	Window comparator high threshold
ADC0_CALIB	0x0612	Calibration

20.10.1 CTRLA Bits

Bit:	7	6:5	4	3	2	1	0
	—	—	LEFTADJ	RESSEL	FREERUN	RUNSTDBY	ENABLE

LEFTADJ 1=result left-adjusted in 16-bit register (8-bit in RESH)
 RESSEL 0=10-bit result, 1=8-bit result
 FREERUN 1=continuous conversions
 ENABLE 1=enable ADC

20.10.2 CTRLC — Prescaler and Reference

Bit:	7:6	5:4	3	2:0
	—	REFSEL	SAMPCAP	PRESC

REFSEL 0x0=INTREF (internal), 0x1=VDDIO2/10, 0x2=reserved, 0x3=VREFA (external)

SAMPCAP 1=reduced sample capacitance (for high-impedance sources)

PRESC ADC clock prescaler (see table below)

PRESC	Division
-------	----------

0x0	DIV2
-----	------

0x1	DIV4
-----	------

0x2	DIV8
-----	------

0x3	DIV16
-----	-------

0x4	DIV32
-----	-------

0x5	DIV64
-----	-------

0x6	DIV128
-----	--------

0x7	DIV256
-----	--------

ADC requires 50–1500 kHz clock. At 3.333 MHz use DIV8 (417 kHz) or DIV16 (208 kHz). At 20 MHz use DIV32 (625 kHz) or DIV64 (312 kHz).

20.10.3 MUXPOS — Input Mux

MUXPOS	Channel
--------	---------

0x00	AIN0 (PA0 – also UPDI, avoid)
------	-------------------------------

0x01	AIN1 (PA1)
------	------------

0x02	AIN2 (PA2)
------	------------

0x03	AIN3 (PA3)
------	------------

0x04	AIN4 (PA4)
------	------------

0x05	AIN5 (PA5)
------	------------

0x06	AIN6 (PA6)
------	------------

0x07	AIN7 (PA7)
------	------------

0x08	AIN8 (PB5)
------	------------

0x09	AIN9 (PB4)
------	------------

0x0A	AIN10 (PB1)
------	-------------

0x0B	AIN11 (PB0)
------	-------------

0x1B	DACREF0
------	---------

0x1C	TEMPSENSE (internal temperature sensor)
------	---

0x1D	GND
------	-----

0x1E	VDDIO2/10
------	-----------

0x1F	INTREF (internal reference voltage)
------	-------------------------------------

20.10.4 Internal Reference Voltage

The internal reference is configured via VREF peripheral (base 0x00A0):

Register	Mem Addr	Description
VREF_CTRLA	0x00A0	ADC0 and DAC0 reference select
VREF_CTRLB	0x00A1	ADC0 reference force enable

VREF_CTRLA bits 2:0 (ADC0REFSEL):

0x0	0.55V
-----	-------

0x1	1.1V
-----	------

0x2	2.5V
-----	------

0x3	4.3V
-----	------

0x4	1.5V
-----	------

Default internal reference: 1.1V.

20.11 NVMCTRL — Non-Volatile Memory Controller

Base address: 0x1000. Writing flash or EEPROM requires writing the CCP_SPM key (0x9D) to CCP before the command, then completing the write within 4 cycles.

Register	Mem Addr	Description
NVMCTRL_CTRLA	0x1000	Command register (write CCP key first)
NVMCTRL_CTRLB	0x1001	Boot lock, application code write protect
NVMCTRL_STATUS	0x1002	Status flags (read-only)
NVMCTRL_INTCTRL	0x1003	EEREADY interrupt enable
NVMCTRL_INTFLAGS	0x1004	EEREADY interrupt flag (W1C)
NVMCTRL_DATA	0x1006	Data register (16-bit)
NVMCTRL_ADDR	0x1008	Address register (16-bit)

20.11.1 CTRLA — NVM Commands

CMD[2:0] (bits 2:0)	Command	
0x00	NOCMD	No operation
0x01	PAGEWRITE	Write page buffer to flash
0x02	PAGEERASE	Erase one page of flash
0x03	PAGEERASEWRITE	Erase and write in one operation
0x04	PAGEBUFCLR	Clear page write buffer
0x05	CHIPERASE	Erase entire chip (flash + EEPROM)
0x06	EEERASE	Erase EEPROM
0x07	FUSEWRITE	Write fuse bytes

20.11.2 STATUS Bits

Bit:	2	1	0
	WRERROR	EEBUSY	FBUSY

FBUSY 1 = flash write/erase in progress
 EEBUSY 1 = EEPROM write/erase in progress
 WRERROR 1 = write error occurred

20.11.3 Memory Layout

Region	Start	End	Size	Notes
Flash	0x0000	0x7FFF	32 KB	Program memory (word-addressed by PC)
SRAM	0x3800	0x3FFF	2 KB	Data memory (avr-ld origin 0x803800)
EEPROM	0x1400	0x14FF	256 B	(memory-mapped for read; use NVMCTRL to write)
Fuses	0x1280	0x128A		Written via UPDI/avrdude

Flash page size: 64 bytes. EEPROM page size: 32 bytes.

20.12 Interrupt Vector Table

All vectors are 4-byte (2-word) entries. Use JMP (4-byte instruction) for targets anywhere in the 32 KB flash, or RJMP + NOP (2+2 bytes) to keep the same slot size.

Vector	Byte Addr	Name	Description
0	0x0000	RESET	Power-on / external / WDT reset
1	0x0004	CRCSCAN_NMI	CRC scan non-maskable interrupt
2	0x0008	BOD_VLM	Brown-out detection voltage level monitor
3	0x000C	PORTA_PORT	Port A pin-change interrupt
4	0x0010	PORTB_PORT	Port B pin-change interrupt
5	0x0014	PORTC_PORT	Port C pin-change interrupt
6	0x0018	RTC_CNT	RTC overflow / compare match
7	0x001C	RTC_PIT	RTC periodic interrupt timer
8	0x0020	TCA0_OVF / TCA0_LUNF	TCA0 overflow (single) / low underflow (split)
9	0x0024	TCA0_HUNF	TCA0 high byte underflow (split mode)
10	0x0028	TCA0_CMP0 / TCA0_LCMP0	TCA0 compare 0 match
11	0x002C	TCA0_CMP1 / TCA0_LCMP1	TCA0 compare 1 match
12	0x0030	TCA0_CMP2 / TCA0_LCMP2	TCA0 compare 2 match
13	0x0034	TCB0_INT	TCB0 capture interrupt
14	0x0038	TCB1_INT	TCB1 capture interrupt
15	0x003C	TWI0_TWIS	TWI0 slave interrupt
16	0x0040	TWI0_TWIM	TWI0 master interrupt
17	0x0044	SPI0_INT	SPI0 transfer complete
18	0x0048	USART0_RXC	USART0 receive complete
19	0x004C	USART0_DRE	USART0 data register empty
20	0x0050	USART0_TXC	USART0 transmit complete
21	0x0054	AC0_AC	Analog comparator interrupt
22	0x0058	ADC0_RESRDY	ADC0 result ready
23	0x005C	ADC0_WCMP	ADC0 window comparator match
24	0x0060	CCL_CCL	Configurable custom logic interrupt
25	0x0064	NVMCTRL_EE	EEPROM ready

Minimal startup with vector table in assembly:

```
.section .vectors,"ax",@progbits
jmp _start          ; 0x0000 RESET – must be first
jmp _unhandled      ; 0x0004 CRCSCAN_NMI
; ... one jmp per vector ...

.section .text
_unhandled:
    rjmp _unhandled    ; spin on unhandled interrupts

_start:
    ; ... init stack, SRAM, then jump to main
```

20.13 Fuses

Fuses are one-time-programmable configuration bytes written via UPDI using avrdude. Factory defaults are shown. Unlike ATmega, there is no separate CKDIV, BOOTRST, or BOOTSZ fuse.

Fuse	Mem Addr	Default	Description
------	----------	---------	-------------

FUSE0	0x1280	0x00	APPEND[7:0]	– application code end page
FUSE1	0x1281	0x00	BOOTEND[7:0]	– boot section end page (0=none)
FUSE2	0x1282	0x02	OSCCFG	– oscillator config
FUSE4	0x1284	0xF6	SYSCFG0	– CRC, reset pin, UPDI pin config
FUSE5	0x1285	0x07	SYSCFG1	– startup time SUT[2:0]
FUSE6	0x1286	0x00	(same as FUSE0 on this device)	
FUSE7	0x1287	0x00	(same as FUSE1)	
FUSE8	0x1288	0xAA	LOCKBIT	

20.13.1 FUSE2 — OSCCFG

Bit: 1 0
 FREQSEL FREQSEL

FREQSEL 0=16 MHz internal oscillator base
 1=20 MHz internal oscillator base ← factory default (0x02)

The clock prescaler in CLKCTRL_MCLKCTRLB further divides this frequency. Factory default: 20 MHz / 6 = 3.333 MHz.

20.13.2 FUSE4 — SYSCFG0

Bit: 7:6 5:4 3:2 1:0
 CRCSRC RSTPINCFG UPDIPINCFG –

RSTPINCFG 0x0=GPIO, 0x1=UPDI, 0x2=RESET (0xF6 default = RESET=GPIO, UPDI active)
 CRCSRC 0x3=no CRC check (default)

20.13.3 Writing Fuses with avrdude

```
avrdude -c serialupdi -p t3217 -P /dev/ttyUSB0 -b 115200 \  

-U fuse2:w:0x02:m
```

```
avrdude -c serialupdi -p t3217 -P /dev/ttyUSB0 -b 115200 \  

-U fuse0:w:0x00:m -U fuse1:w:0x00:m -U fuse2:w:0x02:m \  

-U fuse4:w:0xF6:m -U fuse5:w:0x07:m
```

Warning: Setting RSTPINCFG incorrectly can disable the UPDI programming interface, effectively bricking the device. Never modify FUSE4 unless certain.

Chapter 21

AVR Instruction Set Summary

Complete instruction set for AVR (classic core, enhanced core with MUL). Columns: **Opcode** — mnemonic and operands; **Operation** — what it does; **Cycles** — clock cycles (branch cycles shown as taken/not-taken); **Flags** — SREG bits modified (I T H S V N Z C).

Rd = destination register (R0–R31 unless noted). Rr = source register. K = 8-bit constant (0–255). k = relative branch offset (–64..63 for 7-bit, –2048..2047 for 12-bit). b = bit index (0–7). A = I/O address (0–63). q = displacement (0–63) for Y/Z indirect with displacement.

21.1 Arithmetic and Logic

Opcode	Operation	Cycles	Flags
ADD Rd, Rr	$Rd \leftarrow Rd + Rr$	1	H S V N Z C
ADC Rd, Rr	$Rd \leftarrow Rd + Rr + C$	1	H S V N Z C
ADIW Rd, K	$Rd+1:Rd \leftarrow Rd+1:Rd + K$ (0–63)	2	S V N Z C
SUB Rd, Rr	$Rd \leftarrow Rd - Rr$	1	H S V N Z C
SUBI Rd, K	$Rd \leftarrow Rd - K$	1	H S V N Z C
SBC Rd, Rr	$Rd \leftarrow Rd - Rr - C$	1	H S V N Z C
SBCI Rd, K	$Rd \leftarrow Rd - K - C$	1	H S V N Z C
SBIW Rd, K	$Rd+1:Rd \leftarrow Rd+1:Rd - K$ (0–63)	2	S V N Z C
AND Rd, Rr	$Rd \leftarrow Rd \text{ AND } Rr$	1	S V N Z
ANDI Rd, K	$Rd \leftarrow Rd \text{ AND } K$ (R16–R31)	1	S V N Z
OR Rd, Rr	$Rd \leftarrow Rd \text{ OR } Rr$	1	S V N Z
ORI Rd, K	$Rd \leftarrow Rd \text{ OR } K$ (R16–R31)	1	S V N Z
EOR Rd, Rr	$Rd \leftarrow Rd \text{ XOR } Rr$	1	S V N Z
COM Rd	$Rd \leftarrow 0xFF - Rd$ (bitwise NOT)	1	S V N Z C
NEG Rd	$Rd \leftarrow 0x00 - Rd$ (two's comp)	1	H S V N Z C
INC Rd	$Rd \leftarrow Rd + 1$	1	S V N Z
DEC Rd	$Rd \leftarrow Rd - 1$	1	S V N Z
MUL Rd, Rr	$R1:R0 \leftarrow Rd \times Rr$ (unsigned)	2	Z C
MULS Rd, Rr	$R1:R0 \leftarrow Rd \times Rr$ (signed)	2	Z C
	Rd, Rr: R16–R31 only		
MULSU Rd, Rr	$R1:R0 \leftarrow Rd \times Rr$ (sgn $\times$ uns)	2	Z C
	Rd, Rr: R16–R23 only		
FMUL Rd, Rr	$R1:R0 \leftarrow (Rd \times Rr) \ll 1$ (uns)	2	Z C
FMULS Rd, Rr	$R1:R0 \leftarrow (Rd \times Rr) \ll 1$ (sgn)	2	Z C

FMULSU Rd, Rr $R1:R0 \leftarrow (Rd \times Rr) \ll 1 \text{ (s}\times\text{u)} \ 2$ Z C

21.2 Shift and Rotate

Opcode	Operation	Cycles	Flags
LSL Rd	$Rd \leftarrow Rd \ll 1$; $Rd0=0$; $C \leftarrow Rd7$	1	H S V N Z C
LSR Rd	$Rd \leftarrow Rd \gg 1$; $Rd7=0$; $C \leftarrow Rd0$	1	S V N Z C
ASR Rd	$Rd \leftarrow Rd \gg 1$; $Rd7$ unchanged	1	S V N Z C
ROL Rd	$Rd \leftarrow Rd \ll 1$; $Rd0=C$; $C \leftarrow Rd7$	1	H S V N Z C
ROR Rd	$Rd \leftarrow Rd \gg 1$; $Rd7=C$; $C \leftarrow Rd0$	1	S V N Z C
SWAP Rd	$Rd[7:4] \leftrightarrow Rd[3:0]$	1	—

21.3 Bit Operations

Opcode	Operation	Cycles	Flags
SBI A, b	$I/O[A][b] \leftarrow 1$ (A: 0–31)	2	—
CBI A, b	$I/O[A][b] \leftarrow 0$ (A: 0–31)	2	—
BST Rd, b	$T \leftarrow Rd[b]$	1	T
BLD Rd, b	$Rd[b] \leftarrow T$	1	—
SEC	$C \leftarrow 1$	1	C
CLC	$C \leftarrow 0$	1	C
SEN	$N \leftarrow 1$	1	N
CLN	$N \leftarrow 0$	1	N
SEZ	$Z \leftarrow 1$	1	Z
CLZ	$Z \leftarrow 0$	1	Z
SEI	$I \leftarrow 1$ (global interrupt on)	1	I
CLI	$I \leftarrow 0$ (global interrupt off)	1	I
SES	$S \leftarrow 1$	1	S
CLS	$S \leftarrow 0$	1	S
SEV	$V \leftarrow 1$	1	V
CLV	$V \leftarrow 0$	1	V
SET	$T \leftarrow 1$	1	T
CLT	$T \leftarrow 0$	1	T
SEH	$H \leftarrow 1$	1	H
CLH	$H \leftarrow 0$	1	H
BSET b	$SREG[b] \leftarrow 1$	1	(b)
BCLR b	$SREG[b] \leftarrow 0$	1	(b)

21.4 Compare and Test

Opcode	Operation	Cycles	Flags
CP Rd, Rr	$Rd - Rr$ (flags only)	1	H S V N Z C
CPC Rd, Rr	$Rd - Rr - C$ (flags only)	1	H S V N Z C
CPI Rd, K	$Rd - K$ (flags only, R16–R31)	1	H S V N Z C

TST Rd Rd AND Rd (flags only) 1 S V N Z

21.5 Branch Instructions

Opcode	Condition	Cycles	Notes
RJMP k	—	2	PC ← PC + k + 1; k: -2048..2047
JMP k	—	3	PC ← k; 32-bit instruction
IJMP	—	2	PC ← Z (16-bit)
EIJMP	—	2	PC ← EIND:Z (>128KB flash)
RCALL k	—	3	Push PC, PC ← PC + k + 1
CALL k	—	4	Push PC, PC ← k; 32-bit
ICALL	—	3	Push PC, PC ← Z
EICALL	—	3	Push PC, PC ← EIND:Z
RET	—	4	Pop PC
RETI	—	4	Pop PC, I ← 1
BRBS b, k	SREG[b] = 1	1/2	Branch if bit set; 1=not taken, 2=taken
BRBC b, k	SREG[b] = 0	1/2	Branch if bit clear
/* Named branch mnemonics (aliases for BRBS/BRBC) */			
BREQ k	Z = 1	1/2	Branch if equal
BRNE k	Z = 0	1/2	Branch if not equal
BRCS k	C = 1	1/2	Branch if carry set
BRCC k	C = 0	1/2	Branch if carry clear
BRSH k	C = 0	1/2	Branch if same or higher (unsigned ≥)
BRLO k	C = 1	1/2	Branch if lower (unsigned <)
BRMI k	N = 1	1/2	Branch if minus
BRPL k	N = 0	1/2	Branch if plus
BRGE k	S = 0	1/2	Branch if greater or equal (signed ≥)
BRLT k	S = 1	1/2	Branch if less than (signed <)
BRHS k	H = 1	1/2	Branch if half-carry set
BRHC k	H = 0	1/2	Branch if half-carry clear
BRTS k	T = 1	1/2	Branch if T set
BRTC k	T = 0	1/2	Branch if T clear
BRVS k	V = 1	1/2	Branch if overflow set
BRVC k	V = 0	1/2	Branch if overflow clear
BRIE k	I = 1	1/2	Branch if interrupt enabled
BRID k	I = 0	1/2	Branch if interrupt disabled

21.5.1 Skip Instructions

Opcode	Condition	Cycles	Notes
CPSE Rd, Rr	Rd = Rr → skip next	1/2/3	skip 1 or 2-word instruction
SBRC Rd, b	Rd[b] = 0 → skip	1/2/3	
SBRS Rd, b	Rd[b] = 1 → skip	1/2/3	
SBIC A, b	I/O[A][b] = 0 → skip	2/3/4	A: 0-31 only
SBIS A, b	I/O[A][b] = 1 → skip	2/3/4	A: 0-31 only

Skip cycle counts: 1 = no skip; 2 = skip 1-word instruction; 3 = skip 2-word instruction (LDS, STS, CALL,

JMP). SBIC/SBIS add 1 due to I/O read.

21.6 Load and Store

Opcode	Operation	Cycles	Notes
MOV Rd, Rr	$Rd \leftarrow Rr$	1	
MOVW Rd, Rr	$Rd+1:Rd \leftarrow Rr+1:Rr$	1	Rd, Rr: even registers
LDI Rd, K	$Rd \leftarrow K$ (R16–R31 only)	1	
LDS Rd, k	$Rd \leftarrow \text{SRAM}[k]$	2	k: 16-bit address
STS k, Rr	$\text{SRAM}[k] \leftarrow Rr$	2	k: 16-bit address
/* X register (R27:R26) indirect */			
LD Rd, X	$Rd \leftarrow \text{SRAM}[X]$	2	
LD Rd, X+	$Rd \leftarrow \text{SRAM}[X]; X \leftarrow X+1$	2	
LD Rd, -X	$X \leftarrow X-1; Rd \leftarrow \text{SRAM}[X]$	2	
ST X, Rr	$\text{SRAM}[X] \leftarrow Rr$	2	
ST X+, Rr	$\text{SRAM}[X] \leftarrow Rr; X \leftarrow X+1$	2	
ST -X, Rr	$X \leftarrow X-1; \text{SRAM}[X] \leftarrow Rr$	2	
/* Y register (R29:R28) indirect */			
LD Rd, Y	$Rd \leftarrow \text{SRAM}[Y]$	2	
LD Rd, Y+	$Rd \leftarrow \text{SRAM}[Y]; Y \leftarrow Y+1$	2	
LD Rd, -Y	$Y \leftarrow Y-1; Rd \leftarrow \text{SRAM}[Y]$	2	
LDD Rd, Y+q	$Rd \leftarrow \text{SRAM}[Y+q]$ q: 0–63	2	
ST Y, Rr	$\text{SRAM}[Y] \leftarrow Rr$	2	
ST Y+, Rr	$\text{SRAM}[Y] \leftarrow Rr; Y \leftarrow Y+1$	2	
ST -Y, Rr	$Y \leftarrow Y-1; \text{SRAM}[Y] \leftarrow Rr$	2	
STD Y+q, Rr	$\text{SRAM}[Y+q] \leftarrow Rr$ q: 0–63	2	
/* Z register (R31:R30) indirect */			
LD Rd, Z	$Rd \leftarrow \text{SRAM}[Z]$	2	
LD Rd, Z+	$Rd \leftarrow \text{SRAM}[Z]; Z \leftarrow Z+1$	2	
LD Rd, -Z	$Z \leftarrow Z-1; Rd \leftarrow \text{SRAM}[Z]$	2	
LDD Rd, Z+q	$Rd \leftarrow \text{SRAM}[Z+q]$ q: 0–63	2	
ST Z, Rr	$\text{SRAM}[Z] \leftarrow Rr$	2	
ST Z+, Rr	$\text{SRAM}[Z] \leftarrow Rr; Z \leftarrow Z+1$	2	
ST -Z, Rr	$Z \leftarrow Z-1; \text{SRAM}[Z] \leftarrow Rr$	2	
STD Z+q, Rr	$\text{SRAM}[Z+q] \leftarrow Rr$ q: 0–63	2	

21.6.1 Load from Program Memory (Flash)

Opcode	Operation	Cycles	Notes
LPM	$R0 \leftarrow \text{FLASH}[Z]$	3	Z = byte address
LPM Rd, Z	$Rd \leftarrow \text{FLASH}[Z]$	3	
LPM Rd, Z+	$Rd \leftarrow \text{FLASH}[Z]; Z \leftarrow Z+1$	3	
ELPM	$R0 \leftarrow \text{FLASH}[\text{RAMPZ}:Z]$	3	>64KB flash devices
ELPM Rd, Z	$Rd \leftarrow \text{FLASH}[\text{RAMPZ}:Z]$	3	
ELPM Rd, Z+	$Rd \leftarrow \text{FLASH}[\text{RAMPZ}:Z]; Z++$	3	
SPM	$\text{FLASH}[Z] \leftarrow R1:R0$	varies	Boot section only

21.7 I/O Instructions

Opcode	Operation	Cycles	Notes
IN Rd, A	$Rd \leftarrow I/O[A]$	1	A: 0-63
OUT A, Rr	$I/O[A] \leftarrow Rr$	1	A: 0-63
PUSH Rr	$SRAM[SP] \leftarrow Rr; SP \leftarrow SP-1$	2	
POP Rd	$SP \leftarrow SP+1; Rd \leftarrow SRAM[SP]$	2	

21.8 Miscellaneous

Opcode	Operation	Cycles	Notes
NOP	No operation	1	
SLEEP	Enter sleep mode	1	MCU-defined behaviour
WDR	Reset watchdog timer	1	
BREAK	Stop for on-chip debug system	1	

21.9 Instruction Encoding Summary

Most AVR instructions are 16-bit (one word). 32-bit instructions:

32-bit instructions (2 words):

- LDS Rd, k – word 1: opcode+Rd; word 2: 16-bit address k
- STS k, Rr – word 1: opcode+Rr; word 2: 16-bit address k
- CALL k – word 1: 0x940E+kH; word 2: kL (22-bit address)
- JMP k – word 1: 0x940C+kH; word 2: kL (22-bit address)

All other instructions are 16-bit. The assembler handles encoding automatically.

21.10 GCC ABI Register Conventions

Registers	Role	Save by
R0	Temporary, MUL result	Caller (assume destroyed)
R1	Always zero (GCC)	Callee must restore if modified
R2-R17	Call-saved	Callee (push/pop if used)
R18-R27	Call-clobbered	Caller (assume destroyed)
R28-R29 (Y)	Frame pointer	Callee
R30-R31 (Z)	Call-clobbered	Caller

Return values:

- 8-bit: R24
- 16-bit: R25:R24
- 32-bit: R25:R24:R23:R22 (little-endian: R22=lowest byte)

Arguments (first to last):

arg1: R25:R24 (or R24 for 8-bit)

arg2: R23:R22

arg3: R21:R20

arg4: R19:R18

More: pushed on stack (rightmost first)

After MUL/MULS/MULSU: always restore R1 = 0 before returning or calling any C function. R0 is implicitly scratch.

Chapter 22

GAS Directive Reference

GNU Assembler (`avr-as`) directives used in AVR assembly. Directives begin with `.` and are not instructions — they control how the assembler lays out output.

22.1 Section Directives

Directive	Effect
<code>.section name [, "flags" [, @type]]</code>	Switch to named section. Common flags: a = allocatable w = writable x = executable Types: <code>@progbits</code> (data), <code>@nobits</code> (BSS, zero-init)
<code>.text</code>	Switch to <code>.text</code> (code, flash)
<code>.data</code>	Switch to <code>.data</code> (initialised SRAM variables)
<code>.bss</code>	Switch to <code>.bss</code> (zero-initialised SRAM)
<code>.rodata</code>	Read-only data (constants in flash on AVR)

Common section usage on AVR:

```
.section .text          /* executable code (flash) */
.section .data          /* initialised variables (SRAM; copied from flash at startup) */
.section .bss           /* zero-initialised variables (SRAM) */
.section .rodata        /* constants (flash; access via LPM) */
.section .vectors, "ax", @progbits /* interrupt vector table */
.section .noinit, "aw", @nobits    /* SRAM variables NOT initialised at reset */
```

22.2 Symbol Directives

Directive	Effect
<code>.global symbol</code>	Export symbol (visible to linker / other files)
<code>.local symbol</code>	Mark symbol local (file scope only)

<code>.weak</code>	<code>symbol</code>	Weak symbol: overridden if another file defines it
<code>.type</code>	<code>symbol, @function</code>	Mark symbol as function (for debugging/ELF info)
<code>.type</code>	<code>symbol, @object</code>	Mark symbol as data object
<code>.size</code>	<code>symbol, size</code>	Set symbol size (bytes)
<code>.set</code>	<code>symbol, expr</code>	Define symbol = expr (reassignable)
<code>.equ</code>	<code>symbol, expr</code>	Define symbol = expr (not reassignable; synonym: <code>.equiv</code>)

Example:

```
.global uart_init
.type  uart_init, @function
uart_init:
    /* ... */
    ret
.size  uart_init, . - uart_init /* size = current address minus symbol start */
```

22.3 Data Directives

Directive	Size per item	Example
<code>.byte</code> <code>expr[,...]</code>	1 byte	<code>.byte 0x3F, 42, 'A'</code>
<code>.2byte</code> <code>expr[,...]</code>	2 bytes LE	<code>.2byte 0x1234</code> <i>/* = 0x34, 0x12 in memory */</i>
<code>.word</code> <code>expr[,...]</code>	2 bytes LE	<code>.word 1000</code> <i>/* AVR word = 2 bytes */</i>
<code>.long</code> <code>expr[,...]</code>	4 bytes LE	<code>.long 0xDEADBEEF</code>
<code>.quad</code> <code>expr[,...]</code>	8 bytes LE	<code>.quad 0</code>
<code>.ascii</code> <code>"str"</code>	n bytes	<code>.ascii "hello"</code> <i>/* no null terminator */</i>
<code>.asciz</code> <code>"str"</code>	n+1 bytes	<code>.asciz "hello"</code> <i>/* null terminated */</i>
<code>.string</code> <code>"str"</code>	n+1 bytes	<code>.string "hello"</code> <i>/* alias for .asciz */</i>
<code>.space</code> <code>n [,fill]</code>	n bytes	<code>.space 16, 0xFF</code> <i>/* fill with 0xFF */</i>
<code>.zero</code> <code>n</code>	n zero bytes	<code>.zero 64</code> <i>/* alias for .space n, 0 */</i>
<code>.skip</code> <code>n [,fill]</code>	n bytes	<code>.skip 4</code> <i>/* alias for .space */</i>
<code>.fill</code> <code>reps,size,val</code>		<code>.fill 8, 1, 0xAA</code> <i>/* 8 bytes of 0xAA */</i>

22.4 Alignment Directives

Directive	Effect
<code>.align</code> <code>n</code>	Align to 2 ⁿ byte boundary (with NOP padding in <code>.text</code>)
<code>.balign</code> <code>n [,fill]</code>	Align to n byte boundary
<code>.p2align</code> <code>n</code>	Align to 2 ⁿ boundary (same as <code>.align</code> on most targets)
<code>.org</code> <code>expr</code>	Set location counter to <code>expr</code> (absolute address)

On AVR, `.align 1` aligns to 2-byte boundary (one AVR word). `.align 2` aligns to 4 bytes. Useful before jump tables and 16-bit data.

22.5 Conditional Assembly

Directive	Effect
-----------	--------

```

.ifdef sym      Include if sym is defined
.ifndef sym     Include if sym is NOT defined
.if expr       Include if expr != 0
.else          Else branch
.elif expr      Else-if
.endif         End conditional block
.error "msg"    Abort assembly with message
.warning "msg"  Emit warning, continue

```

Example — select multiply implementation at assemble time:

```

#ifdef NOMUL
    rcall mul8u_shiftadd
#else
    mul    r16, r17
    movw   r24, r0
    clr    r1
#endif

```

When using `avr-gcc` as the front-end, `#ifdef/#define` (C preprocessor) work in `.S` files. For plain `avr-as`, use `.ifdef/.equ`.

22.6 Macro Directives

Directive	Effect
<code>.macro name [params]</code>	Begin macro definition
<code>.endm</code>	End macro definition
<code>.exitm</code>	Exit macro early
<code>.rept n</code>	Repeat enclosed block n times
<code>.endr</code>	End <code>.rept</code> block
<code>.irp sym, values</code>	Iterate: sym takes each value in turn
<code>.endr</code>	
<code>.irpc sym, chars</code>	Iterate over characters of a string
<code>.endr</code>	

Example — macro for a 16-bit load:

```

.macro ldw reg_h, reg_l, value
    ldi  \reg_l, lo8(\value)
    ldi  \reg_h, hi8(\value)
.endm

/* Usage */
ldw    r25, r24, 1999    /* expands to: ldi r24, lo8(1999); ldi r25, hi8(1999) */

```

Example — `.rept` for vector table padding:

```

__vectors:
    rjmp reset_handler
    .rept 25
    rjmp dummy_isr        /* fill remaining 25 vectors */
    .endr

```


22.7 Include Directives

Directive	Effect
-----------	--------

<code>.include "file"</code>	Textually include file at this point
<code>#include <file></code>	C preprocessor include (only when assembled via <code>avr-gcc</code>)

For AVR with `avr-gcc -x assembler-with-cpp` (the default for `.S` files):

```
#include <avr/io.h>          /* defines I/O register names, bit names, RAMEND, etc. */
<avr/io.h> selects the correct device header based on -mmcu= flag.
```

22.8 Listing and Debug Directives

Directive	Effect
-----------	--------

<code>.file "name"</code>	Set source file name (for debug info)
<code>.line n</code>	Set source line number
<code>.loc file line [col]</code>	DWARF location (used by <code>avr-gcc</code> ; rarely hand-written)
<code>.stabs "str",n,n,n,n</code>	STABS debug info (legacy)
<code>.ident "string"</code>	Embed comment string in <code>.comment</code> ELF section

22.9 Expression Operators

GAS expressions can use these operators in directives and immediate operands:

Operator	Example	Meaning
<code>+</code>	<code>.byte 1+2</code>	Addition
<code>-</code>	<code>.byte 5-3</code>	Subtraction
<code>*</code>	<code>.word 16*1000</code>	Multiplication
<code>/</code>	<code>.byte 255/2</code>	Integer division
<code>%</code>	<code>.byte 17%10</code>	Modulo
<code><<</code>	<code>.word 1<<8</code>	Left shift
<code>>></code>	<code>.word 256>>1</code>	Right shift
<code>&</code>	<code>.byte 0xFF & val</code>	Bitwise AND
<code> </code>	<code>.byte flags mask</code>	Bitwise OR
<code>^</code>	<code>.byte a ^ b</code>	Bitwise XOR
<code>~</code>	<code>.byte ~mask</code>	Bitwise NOT
<code>lo8(x)</code>	<code>ldi r16, lo8(val)</code>	Low byte of 16-bit value
<code>hi8(x)</code>	<code>ldi r16, hi8(val)</code>	High byte of 16-bit value
<code>hlo8(x)</code>		Byte 2 (bits 23:16) of 32-bit value
<code>hhi8(x)</code>		Byte 3 (bits 31:24) of 32-bit value
<code>pm(x)</code>	<code>.word pm(label)</code>	Program memory word address (byte <code>addr >> 1</code>)
<code>.</code>	<code>(dot)</code>	Current location counter (address)

22.10 Common Patterns

22.10.1 Define a constant

```
.equ F_CPU, 16000000
.equ BAUD, 9600
.equ UBRR_V, (F_CPU / (16 * BAUD)) - 1 /* = 103 */
```

22.10.2 Reserve SRAM variable

```
.section .bss
my_var: .byte 0 /* 1-byte variable */
my_word: .2byte 0 /* 2-byte variable, aligned */
my_buf: .space 64 /* 64-byte buffer */
```

22.10.3 Initialised variable (copied to SRAM at startup by avr-libc __init)

```
.section .data
init_val: .byte 42
```

22.10.4 Constant in flash (accessed via LPM)

```
.section .rodata
my_table:
    .byte 0, 1, 1, 2, 3, 5, 8, 13 /* Fibonacci */
```

22.10.5 Local labels

```
func:
    ldi    r16, 10
1:      /* local label — referenced as 1f (forward) or 1b (back) */
    dec    r16
    brne   1b /* branch back to label 1 */
    ret
```

Local labels (numeric) can be reused across the file; 1f means “next 1: going forward”, 1b means “previous 1: going backward”. They do not pollute the symbol table.

Chapter 23

Linker Script Walkthrough

The linker script tells `avr-ld` where to place each section in the output ELF file and ultimately in the device's address space. For most projects `avr-gcc` selects a built-in script automatically via `-mmcu=`. Understanding linker scripts is necessary when:

- Writing a bootloader (section placement at a specific flash address).
 - Placing variables in a specific SRAM region.
 - Using `.noinit` for variables that survive a software reset.
 - Adding custom sections (e.g., a parameter page at a known flash address).
-

23.1 Anatomy of a Linker Script

A linker script has three main constructs:

```
MEMORY { }           – declares memory regions (name, origin, length)
SECTIONS { }          – maps input sections to output sections and memory regions
PROVIDE() / ENTRY() – define special symbols and the entry point
```

23.1.1 Minimal AVR Linker Script

```
/* minimal_avr.ld – bare-bones script for ATmega328P */
```

```
MEMORY
{
    flash  (rx)  : ORIGIN = 0x000000, LENGTH = 32K
    sram    (rw!x): ORIGIN = 0x800100, LENGTH = 2K
}

SECTIONS
{
    /* Interrupt vector table – must be at flash origin */
    .vectors :
    {
        KEEP(*(.vectors))
    } > flash

    /* Executable code */
```

```

.text :
{
    *(.text)
    *(.text.*)
} > flash

/* Read-only data (constants accessed via LPM) */
.rodata :
{
    *(.rodata)
    *(.rodata.*)
} > flash

/* Initialised data: stored in flash, copied to SRAM at startup */
.data :
{
    _data_start = .;
    *(.data)
    *(.data.*)
    _data_end = .;
} > sram AT > flash          /* VMA in sram, LMA in flash */

/* Zero-initialised data: only reservation in SRAM */
.bss :
{
    _bss_start = .;
    *(.bss)
    *(.bss.*)
    *(COMMON)
    _bss_end = .;
} > sram

/* Variables preserved across software reset (not zeroed at startup) */
.noinit (NOLOAD) :
{
    *(.noinit)
} > sram

/* Stack: grows downward from RAMEND */
_stack_top = ORIGIN(sram) + LENGTH(sram) - 1;
}

```

23.2 MEMORY Region Flags

```

r  - readable
w  - writable
x  - executable
!  - negate following flags

```

Common combinations:

```

(rx)  - read-execute (flash)
(rw!x) - read-write, not executable (SRAM)

```

(rx!w) – read-execute, not writable (flash, explicit)

23.3 Address Spaces on AVR

AVR uses a **Harvard architecture** with separate address spaces:

Flash (program memory):

avr-ld uses byte addresses with origin 0x000000
 ATmega328P: 0x000000–0x007FFF (32 KB)

SRAM (data memory):

avr-ld uses byte addresses with origin 0x800000 + hardware_offset
 ATmega328P hardware SRAM starts at 0x0100
 avr-ld SRAM origin: 0x800100

Registers: 0x0000–0x001F (R0–R31) – not in linker script
 I/O: 0x0020–0x005F – not in linker script
 Ext I/O: 0x0060–0x00FF – not in linker script
 SRAM: 0x0100–0x08FF (2 KB on ATmega328P)

The 0x800000 offset is an avr-ld convention to distinguish data addresses from code addresses in the ELF file. The hardware only sees the lower 16 bits.

23.4 VMA vs LMA

VMA (Virtual Memory Address): where the section lives at runtime (SRAM for .data).

LMA (Load Memory Address): where the section is stored in the ELF output (flash for .data, since flash is what gets programmed).

```
.data :
{
  _data_start = .;          /* VMA: SRAM address where data lives at runtime */
  *(.data)
  _data_end = .;
} > sram AT > flash        /* LMA: in flash, immediately after .rodata */
```

The startup code (__init in avr-libc) copies .data from its LMA (flash) to its VMA (SRAM) before main() runs:

```
/* avr-libc __init pseudo-code */
ldi ZL, lo8(_data_lma_start) /* source: flash LMA */
ldi ZH, hi8(_data_lma_start)
ldi YL, lo8(_data_start)    /* dest: SRAM VMA */
ldi YH, hi8(_data_start)
1: lpm r0, Z+
st Y+, r0
cpi ZL, lo8(_data_lma_end)
brne 1b
```

If you use -nostartfiles, **you must perform this copy yourself** if any .data variables are used. The simplest approach: avoid initialised variables entirely, or use .equ constants in flash and load them explicitly.

23.5 KEEP()

The garbage-collector (`--gc-sections`) removes unreferenced sections. `KEEP()` prevents this:

```
KEEP(*(.vectors))    /* never discard the vector table */
```

When building with `avr-gcc -Wl,--gc-sections` (common for size optimisation), sections not reachable from the entry point are removed. `KEEP()` forces retention regardless of reachability.

23.6 Bootloader Placement

Place the bootloader at the boot section start address by overriding `.text` origin:

```
avr-gcc -Wl,--section-start=.text=0x7C00 ...
```

Or in a dedicated linker script:

```
/* bootloader.ld */
MEMORY
{
    boot_flash (rx) : ORIGIN = 0x007C00, LENGTH = 1K
    sram        (rw!x): ORIGIN = 0x800100, LENGTH = 2K
}

SECTIONS
{
    .text : { *(.vectors) *(.text) *(.text.*) } > boot_flash
    .bss  : { *(.bss) } > sram
}

avr-gcc -mmcu=atmega328p -nostartfiles \
    -T bootloader.ld \
    -o bootloader.elf bootloader.S
```

23.7 Parameter Page at Fixed Flash Address

Store calibration data at a known, fixed flash address so the application can find it regardless of code size changes:

```
SECTIONS
{
    /* ... normal sections ... */

    .params :
    {
        *(.params)
    } > flash
    . = 0x7000;    /* force next section to start at 0x7000 */
    .calibration :
    {
```

```

    KEEP*(.calibration))
} > flash
}

```

In source:

```

.section .calibration
cal_offset: .word 0x0032    /* factory calibration: ADC offset */
cal_gain:   .word 0x0100    /* Q8.8 gain factor */

```

Access at runtime:

```

ldi    ZL, lo8(cal_offset)
ldi    ZH, hi8(cal_offset)
lpm    r24, Z+
lpm    r25, Z      /* r25:r24 = cal_offset value */

```

23.8 .noinit Section

Variables in `.noinit` are **not zeroed** at startup — they retain whatever value was in SRAM before reset. Useful for reset reason tracking:

```

.section .noinit,"aw",@nobits
reset_reason: .byte 0    /* written before software reset; survives it */

```

Application writes a code before resetting:

```

ldi    r16, 0xAA
sts    reset_reason, r16
/* ... trigger watchdog reset ... */

```

After reset, check `reset_reason` before the BSS clear to determine why the device restarted. BSS clear (zeroing `.bss`) happens in `__init` — `.noinit` is placed in a separate section that `__init` never touches.

23.9 Useful Linker Symbols

These symbols are defined by the default `avr-gcc` linker script and available in assembly via `lo8()`/`hi8()`:

Symbol	Value
<code>__data_start</code>	VMA start of <code>.data</code> in SRAM
<code>__data_end</code>	VMA end of <code>.data</code> in SRAM
<code>__data_load_start</code>	LMA of <code>.data</code> in flash (copy source for <code>__init</code>)
<code>__bss_start</code>	Start of <code>.bss</code> in SRAM
<code>__bss_end</code>	End of <code>.bss</code> in SRAM
<code>__heap_start</code>	Start of heap (after <code>.bss</code>)
<code>__stack</code>	Initial stack pointer value (= <code>RAMEND</code>)
<code>__TEXT_REGION_ORIGIN__</code>	Start of <code>.text</code> (usually 0)
<code>__DATA_REGION_ORIGIN__</code>	Start of SRAM (usually <code>0x800100</code>)

Access in assembly:

```

ldi    r26, lo8(__bss_start)
ldi    r27, hi8(__bss_start)
ldi    r24, lo8(__bss_end - __bss_start) /* BSS length */

```

23.10 Inspecting Linker Output

```
# Show section layout (addresses, sizes)
avr-objdump -h firmware.elf

# Show all symbols with addresses
avr-nm --numeric-sort firmware.elf

# Show memory map (verbose linker output)
avr-gcc -Wl,-Map=firmware.map ... && cat firmware.map

# Show disassembly with source interleaved
avr-objdump -dS firmware.elf

# Verify flash usage
avr-size --mcu=atmega328p --format=avr firmware.elf
```

Sample avr-size output with AVR format:

AVR Memory Usage

Device: atmega328p

Program: 486 bytes (1.5% Full)
(.text + .data + .bootloader)

Data: 3 bytes (0.1% Full)
(.data + .bss + .noinit)

23.11 Common Linker Errors

Error	Cause / Fix
undefined reference to `symbol'	Symbol not defined in any linked file; add the .S or .o file that defines it.
relocation truncated to fit: R_AVR_13_PCREL against `symbol'	Branch offset exceeds range (e.g. RJMP only reaches $\pm 2\text{KB}$); use JMP or reorganise code.
section `.text' will not fit in region `flash'	Total code exceeds flash; reduce size or use a larger device.
cannot find entry symbol reset_handler	Define reset_handler or use -e reset_handler linker flag.
multiple definition of `__vector_N'	Two .S files both define the same ISR vector entry; remove the duplicate.

Chapter 24

avrdude Flashing Cheat-Sheet

avrdude is the standard tool for programming AVR microcontrollers. It communicates with the device via a programmer (USB, serial, SPI ISP) and can read/write flash, EEPROM, and fuse bytes.

24.1 Basic Syntax

```
avrdude -p PARTNO -c PROGRAMMER [-P PORT] [-b BAUD] -U MEMTYPE:OP:FILE[:FORMAT]
```

- p PARTNO Device (e.g. m328p, t3217, m2560). See Section E.4 for list.
- c PROGRAMMER Programmer type. See Section E.3.
- P PORT Port (e.g. /dev/ttyUSB0, /dev/ttyACM0, usb).
- b BAUD Baud rate (for serial programmers; default depends on programmer).
- U Memory operation (can be repeated for multiple operations).
- v Verbose output. -vv for more detail.
- n Dry run: do not write anything.
- e Erase flash before programming.
- F Override signature check (use with caution).

24.1.1 -U Flag Format

MEMTYPE:OP:FILE[:FORMAT]

MEMTYPE:

- flash – program memory
- eeprom – EEPROM
- lfuse – low fuse byte
- hfuse – high fuse byte
- efuse – extended fuse byte
- lock – lock bits
- signature – device signature (read only)
- calibration – oscillator calibration (read only)

OP:

- r – read from device, write to FILE
- w – write FILE to device
- v – verify device matches FILE

FORMAT (optional, auto-detected if omitted):

- i – Intel HEX
- s – Motorola S-record
- r – raw binary
- e – ELF
- h – hex string
- b – byte value (for single-byte ops like fuses)
- m – immediate value (e.g. 0xDE)

24.2 Most-Used Commands

24.2.1 Flash a HEX file

Arduino Uno (ATmega328P, USB-serial bootloader on /dev/ttyUSB0)

```
avrdude -p m328p -c arduino -P /dev/ttyUSB0 -b 115200 \
-U flash:w:firmware.hex:i
```

USBasp programmer (no port needed)

```
avrdude -p m328p -c usbasp \
-U flash:w:firmware.hex:i
```

AVRISP mkII (USB)

```
avrdude -p m328p -c avrispmkII \
-U flash:w:firmware.hex:i
```

24.2.2 Read flash to file

```
avrdude -p m328p -c usbasp \
-U flash:r:backup.hex:i
```

24.2.3 Verify only (no write)

```
avrdude -p m328p -c usbasp \
-U flash:v:firmware.hex:i
```

24.2.4 Erase then write

```
avrdude -p m328p -c usbasp -e \
-U flash:w:firmware.hex:i
```

24.2.5 Write EEPROM

```
avrdude -p m328p -c usbasp \
-U eeprom:w:data.hex:i
```

24.2.6 Read fuse bytes

```
avrdude -p m328p -c usbasp \
-U lfuse:r:-:h -U hfuse:r:-:h -U efuse:r:-:h
```

Output: e.g. 0xff 0xd9 0x07 (lfuse hfuse efuse)

'-' = stdout, 'h' = hex string format

24.2.7 Write fuse bytes

WARNING: incorrect fuse values can brick the device.
 # Always calculate fuses with an online calculator and verify before writing.

```
# Set hfuse = 0xDE (BOOTRST=0, BOOTSZ=10 for 1KB bootloader section)
avrdude -p m328p -c usbasp \
        -U hfuse:w:0xDE:m

# Restore ATmega328P factory fuses (lfuse=0xFF, hfuse=0xD9, efuse=0x07)
avrdude -p m328p -c usbasp \
        -U lfuse:w:0xFF:m \
        -U hfuse:w:0xD9:m \
        -U efuse:w:0x07:m
```

Warning: Fuses control clock source, brown-out detection, bootloader region, and JTAG. Setting CKSEL bits for an external crystal when no crystal is connected will make the device unresponsive — recovery requires an ISP programmer that provides a clock signal (High-Voltage Serial Programming).

24.2.8 Multiple operations in one command

```
avrdude -p m328p -c usbasp \
        -U flash:w:firmware.hex:i \
        -U eeprom:w:eeprom_data.hex:i
```

24.3 Common Programmers

Programmer ID	Description
arduino	Arduino bootloader via serial (auto-reset on DTR)
arduino-ft232r	Arduino via FTDI FT232R cable
avrisp	Atmel AVRISP v1 (serial)
avrisp2	Atmel AVRISP mkII (USB HID); also: avrispmkII
usbasp	USBasp (cheap USB ISP clone, widely available)
usbtiny	USBTiny / Adafruit USBtinyISP
dragon_isp	AVR Dragon in ISP mode
dragon_jtag	AVR Dragon in JTAG mode
jtag2isp	JTAGICE mkII in ISP mode
jtag2	JTAGICE mkII
atmelice_isp	Atmel-ICE in ISP mode
atmelice_pdi	Atmel-ICE in PDI mode (XMEGA)
atmelice_updi	Atmel-ICE in UPDI mode (tinyAVR/megaAVR 0/1/2-series)
pymcuprog	Microchip MPLAB Snap / Curiosity Nano via pymcuprog bridge
serialupdi	UPDI via USB-serial adapter (cheap, 1-wire)

24.3.1 UPDI Devices (ATtiny3217 and tinyAVR 1-series)

The ATtiny3217 uses UPDI (Unified Program and Debug Interface), a single-wire protocol. Use serialupdi with a USB-serial adapter connected to the UPDI pin via a 4.7kΩ resistor:

```
# ATtiny3217 via serialupdi on /dev/ttyUSB0
avrdude -p t3217 -c serialupdi -P /dev/ttyUSB0 \
        -U flash:w:firmware.hex:i
```

```
# Read fuses
```

```
avrdude -p t3217 -c serialupdi -P /dev/ttyUSB0 \
        -U fuse0:r:-:h -U fuse1:r:-:h -U fuse2:r:-:h \
        -U fuse4:r:-:h -U fuse5:r:-:h -U fuse6:r:-:h -U fuse7:r:-:h
```

UPDI wiring (1-wire):

```
USB-serial TX ┌───┐
                │   │ 4.7kΩ — UPDI pin (ATtiny3217 pin 16)
USB-serial RX └───┘
USB-serial GND — GND
```

Alternatively, use the Atmel-ICE or MPLAB PICkit 4 for UPDI.

24.4 Device Part Numbers

Part Number	Device
m328p	ATmega328P (Arduino Uno, Nano)
m328pb	ATmega328PB
m2560	ATmega2560 (Arduino Mega)
m32u4	ATmega32U4 (Arduino Leonardo, Pro Micro)
m168p	ATmega168P
m88p	ATmega88P
m48p	ATmega48P
t3217	ATtiny3217 (book target for tinyAVR 1-series)
t1616	ATtiny1616
t816	ATtiny816
t416	ATtiny416
t85	ATtiny85
t84	ATtiny84
t45	ATtiny45
t13	ATtiny13
x128a4u	ATxmega128A4U

Full list: `avrdude -p ?`

24.5 Fuse Byte Reference (ATmega328P)

24.5.1 Low Fuse (LFUSE)

Bit	Name	Description
7	CKDIV8	Divide clock by 8 (0 = enabled; factory default)
6	CKOUT	Clock output on PB0 (0 = enabled)
5:4	SUT1:0	Start-up time (10 = default 14CK + 65ms)
3:0	CKSEL3:0	Clock source (0010 = internal 8MHz RC; 1111 = external crystal)

Factory default (internal RC, /8): LFUSE = 0x62

No divide (internal 8MHz full speed): LFUSE = 0xE2

External 16MHz crystal (Arduino): LFUSE = 0xFF

24.5.2 High Fuse (HFUSE)

Bit	Name	Description
7	RSTDISBL	Disable reset pin (0 = PC6 becomes I/O; DANGEROUS)
6	DWEN	debugWIRE enable (0 = enabled; disables ISP when set)
5	SPIEN	SPI programming enable (0 = enabled; keep this 0)
4	WDTON	Watchdog always on (0 = always on)
3	EESAVE	Preserve EEPROM during chip erase (0 = preserved)
2:1	BOOTSZ1:0	Boot section size (11=256B, 10=512B, 01=1KB, 00=2KB)
0	BOOTRST	Boot reset vector (0 = reset jumps to boot section)

Factory default: HFUSE = 0xD9

(SPI enabled, WDT off, EEPROM erased, 2KB boot, reset to 0x0000)

Bootloader (1KB boot, reset to boot section): HFUSE = 0xDE

24.5.3 Extended Fuse (EFUSE)

Bit	Name	Description
2:0	BODLEVEL2:0	Brown-out detection level
	111	= BOD disabled (factory)
	110	= 1.8V
	101	= 2.7V
	100	= 4.3V

Factory default: EFUSE = 0x07 (BOD disabled; only bits 2:0 used)

24.5.4 Online Fuse Calculator

The Engbedded AVR Fuse Calculator provides a GUI for calculating fuse values. Search “AVR fuse calculator” — cross-check with the device datasheet before programming.

24.6 Generating Intel HEX from ELF

```
avr-objcopy -O ihex firmware.elf firmware.hex
```

```
# EEPROM section separately
avr-objcopy -O ihex -j .eeprom \
  --set-section-flags=.eeprom=alloc,load \
  --change-section-lma .eeprom=0 \
  firmware.elf eeprom.hex
```

24.7 Complete Build + Flash Script

```
#!/bin/bash
# build_flash.sh - assemble, link, and flash ATmega328P
set -e

MCU=atmega328p
```

```

PORT=/dev/ttyUSB0
BAUD=115200
SRC=firmware.S

avr-gcc -mmcu=${MCU} -nostartfiles -o firmware.elf ${SRC}
avr-size firmware.elf
avr-objcopy -O ihex firmware.elf firmware.hex
avrdude -p ${MCU} -c arduino -P ${PORT} -b ${BAUD} \
        -U flash:w:firmware.hex:i
echo "Done."

# For ATtiny3217 via serialupdi
MCU=attiny3217
PORT=/dev/ttyUSB0

avr-gcc -mmcu=${MCU} -nostartfiles -o firmware.elf firmware.S
avr-objcopy -O ihex firmware.elf firmware.hex
avrdude -p ${MCU} -c serialupdi -P ${PORT} \
        -U flash:w:firmware.hex:i

```

24.8 Troubleshooting

Problem	Likely cause / fix
avrdude: stk500_recv(): timeout	Wrong port, wrong baud, board not reset. Try <code>-P /dev/ttyACM0</code> , check USB cable.
avrdude: Expected signature for ATmega328P is 1E 95 0F	Wrong <code>-p</code> part number; use <code>avrdude -p ?</code> to list all; read signature with <code>-U signature:r:-:h</code> .
avrdude: verification error	Flash write failed; bad connection, insufficient power, or write-protected region.
avrdude: error: usbasp, Error: Could not find USB device	USBasp firmware too old for target; update USBasp firmware, or add <code>-B 10</code> flag (slow SCK).
avrdude: initialization failed rc=-1; check connections and try again, or use <code>-F</code>	Target not responding; check power, RESET pin, JTAG enable fuse (JTAGEN=0 disables ISP on some devices).
Can't open device <code>/dev/ttyUSB0</code> Permission denied	Add user to dialout group: <code>sudo usermod -a -G dialout \$USER</code> Then log out and back in.

24.8.1 USBasp Slow Clock

If ISP communication fails with a freshly-soldered board or a device that has been running at low speed (CKDIV8 fuse set):

```

# Slow down SCK to 125 kHz (-B = bitclock period in µs; -B 8 = 125kHz SCK)
avrdude -p m328p -c usbasp -B 8 \
        -U lfuse:w:0xFF:m # then reprogram fuses for 16MHz crystal

```

24.8.2 Arduino Auto-Reset Circuit

Arduino boards use a 100 nF capacitor on DTR/RST. avrdude -c arduino toggles DTR to reset the MCU before uploading. If the reset does not trigger:

```
# Manual: press reset button 2 seconds before avrdude starts, then release
avrdude -p m328p -c arduino -P /dev/ttyUSB0 -b 115200 \
        -U flash:w:firmware.hex:i
```

Or disable the auto-reset capacitor (cut the RESET EN trace on the Arduino PCB) when running a custom bootloader that does not expect the timing.