

AVR Assembly Programming

Dazed_N_Confused; ChatGPT for datasheet/app-note extraction and typing help

Contents

1	Why Assembly? Toolchain Setup	1
1.1	What Is Assembly Language?	1
1.1.1	When do you actually need assembly?	1
1.2	Target: ATtiny3217	1
1.3	The AVR-GAS Toolchain	2
1.3.1	Installing the Toolchain	2
1.4	The Build Pipeline	3
1.4.1	Why two link steps?	3
1.5	First Program: Blink	3
1.5.1	Understanding the ATtiny3217 PORT Peripheral	3
1.5.2	The Code	4
1.6	Building and Flashing	6
1.6.1	Step 1: Assemble	6
1.6.2	Step 2: Link	6
1.6.3	Step 3: Convert to HEX	6
1.6.4	Step 4: Inspect (optional but recommended)	6
1.6.5	Step 5: Flash	7
1.7	A Minimal Makefile	7
1.8	What Just Happened — Mental Model	8
1.9	Summary	9
1.10	Exercises	9
2	AVR Architecture Overview	10
2.1	Harvard Architecture	10
2.2	ATtiny3217 Memory Map	11
2.3	The Register File	13
2.3.1	Register Pairs (16-bit Pointers)	13
2.3.2	LDI Restriction	14
2.4	The Status Register (SREG)	14
2.4.1	Flag by Flag	14
2.4.2	Reading and Saving SREG	15
2.5	The Stack and Stack Pointer	15
2.5.1	PUSH and POP	16
2.5.2	How CALL Uses the Stack	16
2.6	The Program Counter (PC)	17
2.6.1	Instruction Fetch Pipeline	17
2.7	I/O Register Space	17
2.8	Putting It Together — A Register-State Walk-Through	19
2.9	Key Architectural Constraints Summary	19
2.10	Summary	20
2.11	Exercises	20
3	AVR-AS Syntax and Directives	22
3.1	File Extensions: .S vs .s	22
3.2	Comments	23
3.3	Instruction Syntax	23
3.4	Number Literals	23

3.5	Expressions	24
3.5.1	Built-in Expression Functions	24
3.6	Labels	25
3.6.1	Global vs Local Labels	25
3.6.2	GAS Local Labels (Numeric)	25
3.6.3	File-Scoped Local Labels (.L prefix)	26
3.7	Sections	26
3.7.1	Core Sections	26
3.7.2	Section Flags and Types (Full Form)	27
3.8	Data Directives	28
3.8.1	.byte — 8-bit values	28
3.8.2	.word — 16-bit values (little-endian)	28
3.8.3	.long — 32-bit values	28
3.8.4	.ascii and .asciz — Strings	28
3.8.5	.skip and .space — Reserve Bytes	29
3.8.6	.balign — Alignment	29
3.9	Constant Directives	29
3.9.1	.equ — Define a Constant	29
3.9.2	.set — Reassignable Constant	30
3.9.3	#define — Preprocessor Macro	30
3.10	The .org Directive	30
3.11	Flash Data on AVRxt (ATtiny3217)	31
3.12	The .include Directive	32
3.13	Macro Directives	32
3.14	Full Section Layout Example	33
3.15	Examining the Output with avr-objdump	34
3.16	Quick Reference	34
3.17	Summary	35
3.18	Exercises	35
Chapter 3A: Debugging AVR Assembly with GDB		37
3.19	Why This Chapter Belongs Here	37
3.20	The Files GDB Needs	37
3.21	Keep Symbols That Help Debugging	38
3.22	Starting GDB	39
3.23	Debugging Targets Used Here	39
3.23.1	1. Static ELF Inspection	40
3.23.2	2. ATtiny3217 Curiosity Nano Hardware	40
3.23.3	3. PyAvrOCD Debugging Guide	40
3.24	Core GDB Commands for Assembly	46
3.25	Instruction-Level Stepping	47
3.26	Reading Disassembly	47
3.27	Inspecting Registers and Flags	48
3.28	Inspecting SRAM	48
3.29	Inspecting Flash Data	49
3.30	Breakpoints on Microcontrollers	49
3.31	Watchpoints and Data Breaks	50
3.32	Debugging the Stack	50
3.33	Debugging Interrupts	51
3.34	Debugging Peripheral Code	52
3.35	A Minimal Assembly Debug Example	52
3.36	A Useful .gdbinit for AVR Assembly	53
3.37	Failure Patterns and What to Inspect	54
3.38	Professional Debug Habits	55
3.39	Source Notes	55
4 Instruction Set Basics		56
4.1	Instruction Encoding	56

4.1.1	Single-Word Instructions (16-bit)	56
4.1.2	Two-Word Instructions (32-bit)	57
4.1.3	Cycle Counts at a Glance	57
4.2	Instruction Categories	58
4.3	Data Transfer — Moving Values Around	58
4.3.1	MOV — Copy Register	58
4.3.2	MOVW — Copy Register Pair	59
4.3.3	LDI — Load Immediate	59
4.3.4	CLR and SER — Clear and Set All Bits	59
4.4	Arithmetic Instructions	59
4.4.1	ADD — Add Without Carry	59
4.4.2	ADC — Add with Carry	60
4.4.3	SUB — Subtract Without Carry	60
4.4.4	SBC — Subtract with Carry (Borrow)	61
4.4.5	SUBI — Subtract Immediate	61
4.4.6	SBCI — Subtract Immediate with Carry	61
4.4.7	ADIW — Add Immediate to Word	61
4.4.8	SBIW — Subtract Immediate from Word	62
4.4.9	INC and DEC — Increment / Decrement	62
4.4.10	NEG — Two's Complement Negate	62
4.5	Logic Instructions	63
4.5.1	AND / ANDI — Bitwise AND	63
4.5.2	OR / ORI — Bitwise OR	63
4.5.3	EOR — Bitwise XOR (Exclusive OR)	63
4.5.4	COM — One's Complement (Bitwise NOT)	63
4.5.5	TST — Test Register (Non-Destructive)	63
4.6	Comparison Instructions	64
4.6.1	CP — Compare	64
4.6.2	CPC — Compare with Carry	64
4.6.3	CPI — Compare with Immediate	64
4.6.4	CPSE — Compare, Skip if Equal	64
4.7	Conditional Branches — Full Table	65
4.8	Unsigned vs Signed Comparisons	66
4.9	Flag Effects Summary	66
4.10	Worked Example: Compare, Clamp, Negate	67
4.11	Operand Restrictions Cheat-Sheet	68
4.12	Summary	69
4.13	Exercises	69
5	Load and Store	70
5.1	The Memory-Register Model	70
5.1.1	ATtiny3217 Data Space Is Not Just SRAM	70
5.2	Addressing Modes	71
5.3	Direct Addressing: LDS and STS	71
5.3.1	I/O Registers Above 0x3F	72
5.4	Pointer Registers: X, Y, Z	72
5.5	Indirect Addressing: LD and ST	73
5.5.1	Basic Indirect — No Change to Pointer	73
5.5.2	Post-Increment — Walk Forward Through Memory	73
5.5.3	Pre-Decrement — Walk Backward Through Memory	74
5.5.4	All Three Pointers, All Three Modes	74
5.6	Displacement Addressing: LDD and STD	74
5.7	Cycle and Size Comparison	75
5.8	Reading Flash: LPM	76
5.9	Reading Flash: Mapped Access (AVRxt Only)	76
5.9.1	LPM vs Mapped LD — Which to Use?	77
5.10	Pointer Arithmetic	78
5.10.1	Advancing a Pointer by N Bytes	78

5.10.2	Computing an Array Element Address	78
5.11	Atomicity on ATtiny3217	78
5.12	Full Addressing Mode Reference	79
5.13	Worked Example: Reverse a String In-Place	80
5.14	Summary	80
5.15	Exercises	81
6	Arithmetic and Logic: 16-bit and Beyond	82
6.1	Multi-Byte Addition and Subtraction	82
6.1.1	The Carry Chain Pattern	82
6.1.2	24-bit and 32-bit	83
6.1.3	Adding a Small Constant to a 16-bit Register Pair	83
6.1.4	Negating a 16-bit Value	84
6.2	Shift and Rotate Instructions	85
6.2.1	16-bit Shift Left (Multiply by 2)	85
6.2.2	16-bit Shift Right, Logical (Unsigned Divide by 2)	85
6.2.3	16-bit Shift Right, Arithmetic (Signed Divide by 2)	85
6.2.4	Shifting by N Bits	86
6.2.5	Extracting a Bit Field	86
6.3	SWAP — Swap Nibbles	86
6.4	MUL — Unsigned 8×8 Multiply	86
6.4.1	Self-Multiply (Square)	87
6.4.2	8-bit Fraction Scaling	87
6.5	MULS — Signed 8×8 Multiply	87
6.6	MULSU — Signed × Unsigned Multiply	88
6.7	Fractional Multiply: FMUL, FMULS, FMULSU	88
6.8	Multiply Summary	89
6.9	Division — Software Algorithm	89
6.9.1	Unsigned 8-bit Division	89
6.9.2	Unsigned 16-bit Division	90
6.9.3	Power-of-2 Division — Just Shift	91
6.10	Worked Example: 8×8 Unsigned Multiply to 16-bit Result	91
6.11	Signed Arithmetic Pitfalls	92
6.11.1	Two Mistakes with Signed Division	92
6.11.2	Overflow Detection	92
6.12	Useful Idioms	92
6.12.1	Test if a value is zero without changing the register	92
6.12.2	Conditional absolute value	93
6.12.3	Saturating add (clamp at 255 instead of wrapping)	93
6.12.4	Saturating subtract (clamp at 0 instead of wrapping)	93
6.12.5	Sign extension: 8-bit signed → 16-bit signed	93
6.13	Summary	93
6.14	Exercises	94
7	Branches and Control Flow	95
7.1	Unconditional Jumps	95
7.1.1	RJMP — Relative Jump	95
7.1.2	JMP — Absolute Jump	95
7.1.3	IJMP — Indirect Jump via Z	96
7.2	Conditional Branches — Mechanics	96
7.2.1	Exact Branch Range	96
7.2.2	Signed vs Unsigned Conditions	96
7.3	Complete Branch Instruction Table	97
7.4	Skip Instructions	98
7.4.1	CPSE — Compare, Skip if Equal	98
7.4.2	SBRC — Skip if Bit in Register is Clear	98
7.4.3	SBRSC — Skip if Bit in Register is Set	98
7.4.4	SBIC — Skip if Bit in I/O Register is Clear	98

7.4.5	SBIS — Skip if Bit in I/O Register is Set	98
7.4.6	SBIC/SBIS Address Limit on ATtiny3217	99
7.4.7	Skip Over 2-Word Instructions Carefully	99
7.4.8	Skip vs Branch — When to Use Which	100
7.5	Structured Control Flow Patterns	100
7.5.1	Two-way decision	100
7.5.2	Multi-byte comparisons	100
7.5.3	Guard branch	101
7.5.4	Top-tested loop	101
7.5.5	Bottom-tested loop (preferred in assembly)	101
7.5.6	Counted loop	101
7.5.7	Nested loops	102
7.6	Worked Example: Sum an Array	102
7.7	Jump Tables	103
7.7.1	Address Tables with LPM	104
7.7.2	Large Jump Tables with JMP Entries	105
7.8	Polling Loops	105
7.9	Branch Optimisation Notes	106
7.9.1	Bottom-tested saves one instruction per loop	106
7.9.2	Count down, not up	106
7.9.3	Use SBRC/SBRS for single-bit tests	106
7.9.4	Branch delay slots do not exist on AVR	106
7.10	Summary	106
7.11	Exercises	107
8	Subroutines and the Stack	108
8.1	CALL and RCALL	108
8.1.1	Return Address Layout	108
8.1.2	ICALL	109
8.2	RET and RETI	109
8.3	PUSH and POP	109
8.4	Stack Pointer Facts on ATtiny3217	110
8.4.1	Updating SP Safely	110
8.5	Register Ownership	111
8.5.1	The r1 = 0 Rule	111
8.6	Prologue and Epilogue Patterns	112
8.6.1	Minimal Subroutine	112
8.6.2	Saving Preserved Registers	112
8.6.3	Stack Frame with Y	112
8.6.4	When a Frame Is Worth It	113
8.7	Stack Depth Analysis	113
8.8	Worked Example: Recursive Factorial	114
8.8.1	Stack Trace for factorial(4)	114
8.8.2	8 x 16 Multiply Helper	115
8.9	Tail Jumps	115
8.10	Routine Template	116
8.11	Summary	116
8.12	Exercises	117
9	GPIO	118
9.1	GPIO Fundamentals	118
9.2	Classic DDRx / PORTx / PINx (ATmega-Style)	119
9.3	ATtiny3217 Port Architecture	119
9.3.1	VPORT — Virtual Port (Fast Path)	119
9.3.2	PORT — Full Port Registers (Configuration and Atomic Operations)	120
9.4	IN — Read I/O Register	121
9.5	OUT — Write I/O Register	121
9.6	SBI and CBI — Set / Clear Bit in I/O Register	122

9.7	SBIC and SBIS — Skip if Bit Clear / Set	123
9.8	Configuring a GPIO Output	124
9.9	Configuring a GPIO Input with Pull-Up	124
9.10	Debounce	125
9.11	Worked Example: Button-Controlled LED	126
9.11.1	Why PORTA_OUTTGL instead of SBI/CBI?	127
9.11.2	Build and Inspect	127
9.12	Summary	128
9.13	Exercises	128
10	Interrupts	129
10.1	Why Interrupts?	129
10.2	The Interrupt Mechanism	129
10.2.1	What Happens When an Interrupt Fires	129
10.2.2	The I Flag (Global Interrupt Enable)	130
10.2.3	Interrupt Priority	130
10.3	SEI and CLI	131
10.3.1	Critical Sections	131
10.4	RETI — Return from Interrupt	131
10.5	The Interrupt Vector Table	132
10.5.1	Layout on ATtiny3217	132
10.5.2	Defining the Vector Table in .S	133
10.6	ISR Prologue and Epilogue	133
10.6.1	Minimal ISR (uses r16 only)	133
10.6.2	Why This Order?	134
10.6.3	Saving Multiple Registers	134
10.6.4	ISR Stack Budget	134
10.7	Configuring Pin Interrupts on ATtiny3217	134
10.7.1	PINnCTRL Layout	135
10.7.2	PORT INTFLAGS — Clearing Interrupt Flags	135
10.8	ISR Latency and Timing	136
10.9	Worked Example: Pin Interrupt Toggles LED	136
10.9.1	Execution Flow Walkthrough	137
10.9.2	Why Not Debounce in the ISR?	138
10.9.3	Build and Inspect	138
10.10	Summary	138
10.11	Exercises	139
11	Power and Clock Options	140
11.1	The Default Clock	140
11.2	Clock Sources	140
11.3	Main Prescaler	141
11.4	Configuration Change Protection	141
11.5	Waiting for Clock Switches	142
11.6	Sleep Modes	142
11.6.1	Idle	143
11.6.2	Standby	143
11.6.3	Power-down	143
11.6.4	RTC/PIT clock-source changes	143
11.7	Brown-Out and VLM	143
11.8	Worked Example: Change the Prescaler	143
11.9	Worked Example: Power-Down Wake from PIT	144
11.10	Practical Rules	144
11.11	Exercises	145
12	Timer/Counter	146
12.1	Timer Fundamentals	146
12.1.1	What a Timer Actually Is	146
12.1.2	Prescaler	146

12.2	TCA0 on ATtiny3217	147
12.2.1	TCA0 Register Map	147
12.2.2	CTRLA — Control A	148
12.2.3	CTRLB — Control B	148
12.2.4	INTCTRL and INTFLAGS	148
12.3	Waveform Generation Modes	149
12.3.1	Normal Mode (WGMODE = 0x00)	149
12.3.2	FRQ Mode — Frequency Generation (WGMODE = 0x01)	149
12.3.3	Single-Slope PWM Mode (WGMODE = 0x03)	149
12.4	Computing Timer Values	150
12.5	TCA0 Interrupt Vectors	150
12.6	The ISR	151
12.7	Main: Timer Initialization Sequence	152
12.8	Worked Example: 1 Hz Timer Tick	153
12.8.1	Register Usage	153
12.8.2	Timing Analysis	153
12.8.3	Build Commands	153
12.8.4	Verifying with avr-objdump	153
12.9	Extending: Normal Mode with OVF Interrupt	154
12.10	Extending: Single-Slope PWM	155
12.11	Summary	155
12.12	Exercises	156
13	USART (Serial Communication)	157
13.1	UART Framing	157
13.1.1	The Bit Stream	157
13.1.2	8N1 Format	157
13.2	USART0 Registers on ATtiny3217	158
13.2.1	STATUS — Status Register (0x0804)	158
13.2.2	CTRLB — Control B (0x0806)	158
13.2.3	CTRLA — Control A / Interrupt Enables (0x0805)	159
13.2.4	CTRLC — Frame Format (0x0807)	159
13.3	Baud Rate Calculation	160
13.3.1	The Formula	160
13.4	TX/RX Pins on ATtiny3217	160
13.5	Polling Transmit	161
13.5.1	Transmitting a String from Flash	161
13.6	Polling Receive	162
13.6.1	Error Checking	162
13.7	ISR-Driven Receive	163
13.7.1	USART0_RXC Vector	163
13.7.2	RXC ISR	163
13.8	Initialization Sequence	164
13.9	Worked Example: Echo Terminal	165
13.9.1	Data Flow	165
13.9.2	Build and Test	166
13.9.3	Disassembly Check	166
13.10	Circular Buffer for High-Throughput RX	166
13.11	Summary	167
13.12	Exercises	168
14	SPI & TWI (I2C) Overview	169
14.1	SPI Fundamentals	169
14.1.1	SPI Clock Polarity and Phase (Mode)	169
14.1.2	Timing Diagram (Mode 0)	170
14.2	SPI0 Registers on the ATtiny3217	170
14.2.1	SPI0_CTRLA — Control A	170
14.2.2	SPI0_CTRLB — Control B	171

14.2.3	SPI0_INTFLAGS — Interrupt / Status Flags (non-buffered mode)	171
14.2.4	SPI0_DATA — Transmit / Receive	171
14.3	SPI Initialisation (Master Mode)	171
14.4	Transferring a Byte	172
14.5	The 25LC010A SPI EEPROM	173
14.5.1	Key Characteristics	173
14.5.2	Command Set	173
14.5.3	Status Register (returned by RDSR)	173
14.5.4	CS Framing Rules	173
14.6	Example: Write Byte to SPI EEPROM	174
14.6.1	eeeprom_wren — Send Write Enable	174
14.6.2	eeeprom_write_byte — Write One Byte	174
14.6.3	eeeprom_wait_wip — Poll WIP Until Write Done	174
14.6.4	main — Putting It Together	175
14.6.5	Build and Inspect	175
14.7	TWI (I ² C) Fundamentals	175
14.7.1	Bus Topology	176
14.7.2	Transaction Structure	176
14.7.3	I ² C Speeds	176
14.8	TWI0 Master Mode Registers on ATtiny3217	176
14.8.1	TWI0_MCTRLA — Master Control A	177
14.8.2	TWI0_MSTATUS — Master Status	177
14.8.3	TWI0_MBAUD — Baud Rate Register	177
14.8.4	TWI0_MCTRLB — Master Control B (Bus Commands)	178
14.9	TWI Initialisation	178
14.10	Master Write Transaction (Polling)	178
14.10.1	SBRC vs SBRS for RXACK	180
14.11	Example: Write to PCF8574 I ² C I/O Expander	180
14.11.1	Build	180
14.12	SPI vs TWI Comparison	180
14.13	Summary	181
14.14	Exercises	181
15	CRC and Data Integrity in Assembly	183
15.1	CRC Basics	183
15.2	Software CRC-8 Register Plan	184
15.3	Flash Test Data	184
15.4	CRC-8 Implementation	184
15.5	Standalone Demo	185
15.6	Table-Driven CRC	186
15.7	ATtiny3217 CRCSCAN Peripheral	186
15.8	Starting a CRCSCAN in Assembly	187
15.9	NMI on CRC Failure	188
15.10	Startup CRC Through Fuses	188
15.11	Summary	188
15.12	Exercises	189
16	Optimisation Techniques	190
16.1	Build and Inspect the Chapter Examples	190
16.2	Cycle Reference for This Chapter	191
16.3	Baseline: One-Byte Copy Loop	191
16.4	Keep Hot Values in Registers	192
16.5	Loop Unrolling	193
16.6	Tail Handling	194
16.7	Flash Lookup Tables	195
16.7.1	Why ADC ZH, r1 Works	196
16.7.2	Flash Address Limits	196
16.7.3	Verifying Table Placement	196

16.8	Size Results from the Current Sources	196
16.9	Reading avr-objdump Correctly	197
16.10	Common Errors in Hand Optimisation	197
16.10.1	Optimising Cold Code	197
16.10.2	Ignoring Setup Cost	197
16.10.3	Forgetting the Final Branch	198
16.10.4	Using a Zero Register Without Clearing It	198
16.10.5	Putting Hot State in SRAM	198
16.10.6	Unrolling Until the Source Becomes Fragile	198
16.11	Optimisation Checklist	198
16.12	Source Notes	198
16.13	Exercises	199
17	Self-Programming Flash: NVMCTRL on tinyAVR 1-Series	200
17.1	tinyAVR 1-Series vs Classic AVR: No SPM, NVMCTRL Instead	200
17.1.1	Architecture Family: AVRXMEGA3	200
17.2	Microchip App Notes Used	201
17.2.1	What Replaces SPM?	201
17.2.2	Programming Interface: UPDI	202
17.3	NVMCTRL Overview	202
17.3.1	Flash Parameters	202
17.3.2	NVMCTRL Register Map	202
17.3.3	NVMCTRL_STATUS — Status Register (0x1002)	203
17.3.4	NVMCTRL Commands (written to NVMCTRL_CTRLA)	203
17.4	The Page Buffer and Write Sequence	203
17.4.1	How the Page Buffer Works	203
17.4.2	Complete Flash Page Write Sequence	204
17.4.3	The Z Pointer Convention	204
17.5	CCP — Configuration Change Protection	204
17.5.1	Why CCP Exists	204
17.5.2	The CCP Register (0x0034)	205
17.5.3	The Four-Instruction Rule	205
17.6	Flash Write Routine in Assembly	205
17.6.1	Register Allocation	205
17.6.2	Full Listing: src/nvmctrl_write.S	206
17.6.3	Instruction Notes	207
17.7	Practical Firmware Update: UART-Driven Self-Update	207
17.7.1	Intel HEX Recap	207
17.7.2	Architecture of selfupdate.S	207
17.7.3	Register Allocation for selfupdate.S	208
17.7.4	USART0 Initialisation at 3.333 MHz	208
17.7.5	Full Listing: src/selfupdate.S	208
17.7.6	Jumping to the Updated Application	209
17.7.7	Uploading Firmware	209
17.8	BOOTEND Fuse: Protecting a Bootloader Region	210
17.8.1	The Problem	210
17.8.2	BOOTEND: Defining a Protected Region	210
17.8.3	Flash Layout With BOOTEND = 0x04	210
17.8.4	Application Entry Point	211
17.8.5	Run-Time Locks: BOOTLOCK and APCWP	211
17.9	Building and Flashing	212
17.9.1	Building nvmctrl_write.S (standalone test)	212
17.9.2	Building selfupdate.S	212
17.9.3	Converting to HEX	212
17.9.4	Flashing with pymcuprog	212
17.9.5	Verifying the Disassembly	213
17.9.6	Building an Application for the Protected Layout	213
17.10	Common Pitfalls	213

17.10.1	Pitfall 1 — CCP Window Violated	213
17.10.2	Pitfall 2 — Page Buffer Not Aligned	214
17.10.3	Pitfall 3 — Writing to Flash During Interrupt	214
17.10.4	Pitfall 4 — Forgetting to Wait for FBUSY Before Each Command	214
17.10.5	Pitfall 5 — Writing the Boot Section From Application Code	214
17.10.6	Pitfall 6 — UPDI Port Floating (PA0)	214
17.10.7	Pitfall 7 — Baud Rate at Non-Default Clock	215
17.11	Summary	215
17.12	Exercises	216
18	Bit Manipulation & Fixed-Point Math	217
18.1	Shift Operations	217
18.1.1	Register Pair Diagram — 16-bit LSL	218
18.2	Shift-and-Add Multiplication	218
18.2.1	Algorithm — Unsigned $8 \times 8 \rightarrow 16$	218
18.2.2	Implementation	218
18.2.3	Signed Multiply — <code>mul8s</code> ($8 \times 8 \rightarrow 16$ signed)	219
18.3	Fixed-Point Q8.8 Arithmetic	220
18.3.1	Converting to Q8.8	221
18.3.2	Q8.8 Addition	221
18.3.3	Q8.8 Multiplication ($Q8.8 \times Q8.8 \rightarrow Q8.8$)	221
18.4	BCD Arithmetic	222
18.4.1	Decimal Adjust on AVR	223
18.4.2	H Flag and BCD Correction	223
18.5	Binary to BCD Conversion	224
18.5.1	16-bit Binary to 5-Digit BCD	224
18.5.2	Build and Test	226
18.6	Complete Example: Sensor Reading to 7-Segment Display	226
18.6.1	Scaling: ADC $\rightarrow$ Temperature	226
18.6.2	Register Map for <code>display.S</code>	226
18.6.3	7-Segment LUT	227
18.6.4	Scale ADC to Temperature	227
18.6.5	Extract BCD Digits and Drive SPI	227
18.6.6	Build	228
18.7	Instruction Reference for This Chapter	228
18.8	Common Pitfalls	229
18.8.1	Pitfall 1 — Using R1 Without Zeroing After MUL	229
18.8.2	Pitfall 2 — ASR vs LSR on Signed Values	229
18.8.3	Pitfall 3 — Q8.8 Overflow	230
18.8.4	Pitfall 4 — BCD Correction Order	230
18.8.5	Pitfall 5 — Shift Count for Multi-Byte Values	230
18.9	Summary	230
18.10	Exercises	231
19	Real-Time Patterns	232
19.1	Real-Time Fundamentals	232
19.1.1	Two Scheduling Models	232
19.2	Cooperative Task Scheduler Design	233
19.2.1	Task Representation	233
19.2.2	Tick-Driven Scheduler	233
19.2.3	Task Counter Pattern	234
19.3	TCB0 Configuration — 1 ms Tick	234
19.3.1	TCB0 Register Map	234
19.3.2	Initialization Code	235
19.4	ATtiny3217 Interrupt Vector Table	236
19.5	ISR Prologue and Epilogue	237
19.5.1	Minimal Tick ISR	237
19.6	ISR Latency Analysis	238

19.6.1	Sources of Latency	238
19.6.2	Push/Pop Overhead	239
19.6.3	Eliminating Sources of Latency	239
19.6.4	Measuring Latency	239
19.7	Complete Two-Task Scheduler	239
19.7.1	Hardware Setup — Curiosity Nano	239
19.7.2	Register Allocation	240
19.7.3	Source Code	240
19.7.4	Build	244
19.8	Timing Analysis — Worst-Case Response	244
19.9	Non-Blocking Task Extension	245
19.10	Scaling to More Tasks	245
19.10.1	When to Move to a Preemptive Kernel	246
19.11	Summary	246
20	CORDIC on AVR	247
20.1	Why Use CORDIC on AVR?	247
20.1.1	When Not to Use It	248
20.2	The Core Idea	248
20.2.1	Rotation Mode	248
20.2.2	Vectoring Mode	248
20.3	Fixed-Point Choices for AVR	248
20.3.1	Vector Format: Q1.15	249
20.3.2	Angle Format: BAM16	249
20.3.3	CORDIC Gain	249
20.4	The Angle Table	249
20.4.1	Convergence Range	250
20.5	Rotation Mode for Sine and Cosine	250
20.5.1	Example Values	251
20.6	AVR Assembly Structure	251
20.6.1	Table Layout	251
20.6.2	Core Loop Skeleton	251
20.7	How This Maps to AVR Assembly	252
20.7.1	Example: One Unrolled Stage ($i = 3$)	252
20.7.2	Practical Register Plan	253
20.8	Vectoring Mode: Magnitude and Angle	253
20.8.1	Where AVR Uses This	254
20.9	Example 1: DDS / NCO Sine Generator	254
20.10	Example 2: Tilt Angle from Two Axes	254
20.11	Accuracy, Flash, and Cycle Tradeoffs	255
20.11.1	Iteration Count	255
20.11.2	LUT vs CORDIC	255
20.12	Common Pitfalls	255
20.12.1	Pitfall 1 — Forgetting Quadrant Reduction	255
20.12.2	Pitfall 2 — Using Logical Right Shift on Signed Values	255
20.12.3	Pitfall 3 — Ignoring the Gain	256
20.12.4	Pitfall 4 — Overflow in Narrow Formats	256
20.12.5	Pitfall 5 — Hand-Writing the Loop Before Proving Correctness	256
20.13	Summary	256
20.14	Exercises	257
21	ADC and Analog Input	258
21.1	The ADC Mental Model	258
21.2	ADC0 Register Path	258
21.3	Choosing an Analog Pin	259
21.4	Choosing the Reference	259
21.5	Choosing the ADC Clock	260
21.6	Sampling Time	260

21.7	Polling Conversion Routine	261
21.8	Complete Example: Read AIN4 and Drive an LED	261
21.9	Polling vs Interrupts	263
21.10	Accumulation and Simple Filtering	264
21.11	Free-Running Mode	264
21.12	Hardware Sample Accumulation	265
21.13	Window Comparator	265
21.14	Event-Triggered Conversion	266
21.15	Common Pitfalls	266
21.16	Summary	267
21.17	Exercises	267
22	ATtiny3217 Register Reference	268
22.1	CPU Registers	268
22.1.1	SREG Bit Layout	268
22.1.2	Configuration Change Protection (CCP)	269
22.2	CLKCTRL — Clock Controller	269
22.2.1	MCLKCTRLA — Main Clock Source Select	269
22.2.2	MCLKCTRLB — Main Clock Prescaler	269
22.2.3	MCLKSTATUS — Status (read-only)	270
22.3	SLPCTRL — Sleep Controller	270
22.3.1	CTRLA — Sleep Control	270
22.4	BOD — Brown-Out Detector / Voltage Level Monitor	270
22.4.1	BOD_CTRLA	271
22.4.2	VLM Registers	271
22.5	GPIO Registers	271
22.5.1	Port Base Addresses	271
22.5.2	Per-Port Register Map (offset from base)	271
22.5.3	PORTA Absolute Addresses	272
22.5.4	PORTB Absolute Addresses	272
22.5.5	PORTC Absolute Addresses	273
22.5.6	PINnCTRL Bit Layout	273
22.5.7	Virtual PORT Registers (VPORT) — Fast I/O	273
22.6	Curiosity Nano Pin Assignments	274
22.6.1	LED and Button Quick Reference	274
22.7	TCA0 — Timer/Counter Type A (16-bit)	274
22.7.1	CTRLA — Prescaler and Enable	275
22.7.2	CTRLB — Waveform Mode	275
22.8	TCB0 / TCB1 — Timer/Counter Type B (16-bit)	275
22.8.1	CTRLA — Clock Select and Enable	276
22.8.2	CTRLB — Mode (CMODE)	276
22.9	USART0	276
22.9.1	STATUS Bits	277
22.9.2	CTRLA — Interrupt Enables	277
22.9.3	CTRLB — Enable Bits	277
22.9.4	CTRLC — Frame Format	277
22.9.5	Baud Rate Formula and Common Values	278
22.10	SPI0	278
22.10.1	CTRLA Bits	278
22.10.2	SPI Clock Rates	278
22.11	TWI0 — Two-Wire Interface (I2C)	279
22.11.1	MSTATUS Bits	279
22.11.2	MCTRLB — Master Command	279
22.11.3	Baud Rate Formula	279
22.12	ADC0	279
22.12.1	CTRLA Bits	280
22.12.2	CTRLC — Prescaler and Reference	280
22.12.3	MUXPOS — Input Mux	280

22.12.4	Internal Reference Voltage	281
22.13	NVMCTRL — Non-Volatile Memory Controller	281
22.13.1	CTRLA — NVM Commands	281
22.13.2	STATUS Bits	282
22.13.3	Memory Layout	282
22.14	Interrupt Vector Table	282
22.15	Fuses	283
22.15.1	FUSE2 — OSCCFG	283
22.15.2	FUSE4 — SYSCFG0	283
22.15.3	Writing Fuses with avrdude	283
23	AVR Instruction Set Summary	285
23.1	How to Read the AVR Instruction Set Manual	285
23.1.1	What One Instruction Page Contains	285
23.1.2	Start with Syntax, but Trust the Operand Limits	286
23.1.3	Read “Words” Before “Cycles”	286
23.1.4	Relative Addresses Are Not Absolute Addresses	286
23.1.5	Read the SREG Section as “What Can I Test Next?”	287
23.1.6	Do Not Panic About the Boolean Formulas	287
23.1.7	Cycles Are Family-Specific	288
23.1.8	Skip Instructions Need Extra Attention	288
23.1.9	Learn the Address Spaces Separately	288
23.1.10	Aliases and Pseudo-Instructions Can Hide the Real Opcode	289
23.1.11	A Good Manual-Reading Workflow	289
23.1.12	What This Means for the Rest of the Book	289
23.2	Arithmetic and Logic	289
23.3	Shift and Rotate	290
23.4	Bit Operations	290
23.5	Compare and Test	291
23.6	Branch Instructions	291
23.6.1	Skip Instructions	292
23.7	Load and Store	292
23.7.1	Load from Program Memory (Flash)	293
23.8	I/O Instructions	293
23.9	Miscellaneous	293
23.10	Instruction Encoding Summary	293
23.11	GCC ABI Register Conventions	293
24	GAS Directive Reference	295
24.1	Section Directives	295
24.2	Symbol Directives	295
24.3	Data Directives	296
24.4	Alignment Directives	296
24.5	Conditional Assembly	296
24.6	Macro Directives	297
24.7	Include Directives	297
24.8	Listing and Debug Directives	298
24.9	Expression Operators	298
24.10	Common Patterns	298
24.10.1	Define a constant	298
24.10.2	Reserve SRAM variable	298
24.10.3	Initialised variable (copied to SRAM at startup by avr-libc __init)	299
24.10.4	Constant in flash (accessed via LPM)	299
24.10.5	Local labels	299
25	Linker Script Walkthrough	300
25.1	Anatomy of a Linker Script	300
25.1.1	Minimal AVR Linker Script	300
25.2	MEMORY Region Flags	301

25.3	Address Spaces on AVR	302
25.4	VMA vs LMA	302
25.5	KEEP()	302
25.6	Bootloader Placement	303
25.7	Parameter Page at Fixed Flash Address	303
25.8	.noinit Section	304
25.9	Useful Linker Symbols	304
25.10	Inspecting Linker Output	304
25.11	Common Linker Errors	305
26	avrdude Flashing Cheat-Sheet	306
26.1	Basic Syntax	306
26.1.1	-U Flag Format	306
26.2	Most-Used Commands	307
26.2.1	Flash a HEX file	307
26.2.2	Read flash to file	307
26.2.3	Verify only (no write)	307
26.2.4	Erase then write	307
26.2.5	Write EEPROM	307
26.2.6	Read fuse bytes	307
26.2.7	Write fuse bytes	308
26.2.8	Multiple operations in one command	308
26.3	Common Programmers	308
26.3.1	UPDI Devices (ATtiny3217 and tinyAVR 1-series)	308
26.4	Device Part Numbers	309
26.5	Fuse Byte Reference (ATmega328P)	309
26.5.1	Low Fuse (LFUSE)	309
26.5.2	High Fuse (HFUSE)	309
26.5.3	Extended Fuse (EFUSE)	310
26.5.4	Online Fuse Calculator	310
26.6	Generating Intel HEX from ELF	310
26.7	Complete Build + Flash Script	310
26.8	Troubleshooting	311
26.8.1	USBasp Slow Clock	311
26.8.2	Arduino Auto-Reset Circuit	311

Chapter 1

Why Assembly? Toolchain Setup

1.1 What Is Assembly Language?

Every program your computer runs is ultimately a sequence of binary numbers: machine code that the CPU decodes and executes. Assembly language is the human-readable form of that machine code. You are no longer asking a compiler to choose the instructions for you; you are choosing them yourself.

Higher-level languages (C, Python, Rust) are *compiled* or *interpreted* — they add layers of abstraction that trade control for convenience. Assembly gives you almost no abstraction: you decide every instruction, every register, and often every cycle count.

1.1.1 When do you actually need assembly?

Situation	Why assembly wins
Tight timing constraints	You control the exact cycle count
Hardware bring-up / no OS	No runtime support — you write everything
Interrupt service routines	Worst-case latency is deterministic
Extremely small flash budget	No compiler overhead or padding
Learning how CPUs work	Forces deep understanding of architecture
Hand-optimising hot paths	Compiler output is a starting point, not the ceiling

In the embedded world, assembly is not just historical trivia. It still shows up in bootloaders, startup code, interrupt paths, cycle-counted loops, and the smallest pieces of firmware where you cannot afford a vague understanding of what the CPU will do next.

1.2 Target: ATtiny3217

This book uses the **ATtiny3217** — an 8-bit AVR from Microchip’s tinyAVR 1-series. It is a modern, capable microcontroller with:

Feature	Detail
Architecture	AVRxt (enhanced AVR core)
Flash	32 KB
SRAM	2 KB (at 0x3800–0x3FFF)
Clock	Up to 20 MHz internal oscillator
Programming	UPDI (1-wire, replaces ISP)

Feature	Detail
I/O	New-style PORT peripheral (not classic DDRx/PORTx)
Packages	24-pin VQFN, SOIC, SSOP

The **ATTiny3217 Curiosity Nano** development board is a USB-powered board with an onboard debugger/programmer. It has one yellow LED on **PA3** and exposes all device pins. This book uses this board for all hardware examples.

That choice is deliberate. The chip is modern enough to show the newer AVR peripheral model, but small enough that the architecture is still easy to hold in your head.

Note: All concepts carry over to other tinyAVR 1-series devices (ATTiny3216, ATTiny1617, etc.) and with small adjustments to any AVR.

1.3 The AVR-GAS Toolchain

We use the **GNU Assembler** (`avr-as`), part of the `avr-gcc` toolchain. This is the same assembler `avr-gcc` uses under the hood when it compiles C for AVR. Using it directly keeps the toolchain familiar while removing the compiler as a middleman.

The toolchain has four main tools:

```
avr-as      Assembler: .S source → .o object file
avr-ld      Linker:    .o object(s) → .elf executable
avr-objcopy Converter: .elf → .hex (Intel HEX for flashing)
avr-objdump Inspector: disassemble, inspect sections, count cycles
```

You do not need to memorise all four on day one. The practical workflow is: assemble source into an object file, link it into an executable image, convert that image into something the programmer can flash, then inspect the result when you want to verify what really happened.

1.3.1 Installing the Toolchain

Debian / Ubuntu / Kali:

```
sudo apt install gcc-avr binutils-avr avr-libc avrdude
```

`avr-libc` is needed for the device header files (`<avr/io.h>`) that define register names and addresses.

Arch Linux:

```
sudo pacman -S avr-gcc avr-binutils avr-libc avrdude
```

macOS (Homebrew):

```
brew tap osx-cross/avr
brew install avr-gcc avrdude
```

Verify installation:

```
avr-as --version
avr-ld --version
avr-objcopy --version
```

Expected output (version numbers may differ):

```
GNU assembler (GNU Binutils) 2.42
GNU ld (GNU Binutils) 2.42
```

GNU objcopy (GNU Binutils) 2.42

1.4 The Build Pipeline

Source goes through four stages before reaching the MCU:

```

blink.S
|
|   avr-as -mmcu=attiny3217 -o blink.o blink.S
|
▼
blink.o      ← relocatable object (ELF, machine code + symbol table)
|
|   avr-gcc -mmcu=attiny3217 -nostartfiles -o blink.elf blink.o
|
▼
blink.elf    ← linked executable (ELF, absolute addresses assigned)
|
|   avr-objcopy -O ihex blink.elf blink.hex
|
▼
blink.hex    ← Intel HEX – what avrdude writes to flash
|
|   pymcuprog write -d attiny3217 -t uart -f blink.hex
|
▼
ATTiny3217 flash ← running program

```

If you are new to embedded builds, this pipeline is worth pausing over. The source file is not flashed directly. It is first turned into machine code, then assigned final addresses, then repackaged into a format the programmer can write into the device’s flash memory.

1.4.1 Why two link steps?

We use `avr-gcc` as the linker driver rather than calling `avr-ld` directly because it already knows the right device-specific details: the memory layout, the default linker script, and the expected section placement. We could drive the linker manually, but that adds complexity before it teaches us anything useful.

The `-nostartfiles` flag matters because we are not building a normal C program. We do not want the C runtime to insert its own startup sequence; we want to write the reset path ourselves so we can see exactly how execution begins.

1.5 First Program: Blink

The classic “hello world” of embedded is not printing text. It is making a pin change state in a visible, repeatable way. Here, that means toggling the board LED forever.

Hardware: ATTiny3217 Curiosity Nano, LED0 on PA3.

File: `ch01_intro/src/blink.S`

1.5.1 Understanding the ATTiny3217 PORT Peripheral

The ATTiny3217 uses a **new-style** PORT peripheral, unlike older AVR devices (ATmega328P etc.) which used `DDRx`, `PORTx`, `PINx` registers in low I/O space.

On ATTiny3217, each port has a block of memory-mapped registers. PORTA lives at base address **0x0400**:

Offset	Register	Function
0x00	PORTA_DIR	Direction: 1 = output, 0 = input
0x01	PORTA_DIRSET	Write 1 to set pin as output (non-destructive)
0x02	PORTA_DIRCLR	Write 1 to set pin as input (non-destructive)
0x03	PORTA_DIRTGL	Write 1 to toggle direction
0x04	PORTA_OUT	Output value
0x05	PORTA_OUTSET	Write 1 to drive pin HIGH (non-destructive)
0x06	PORTA_OUTCLR	Write 1 to drive pin LOW (non-destructive)
0x07	PORTA_OUTTGL	Write 1 to toggle pin (atomic!)
0x08	PORTA_IN	Read current pin state

The SET/CLR/TGL variants are key: they let you change one pin without reading the register first, eliminating the classic read-modify-write race condition.

That is one of the nicest features of the newer AVR GPIO design. Instead of reading the whole port, modifying one bit in software, and writing the whole byte back, the peripheral can perform the bit change for you.

Because these registers are above address 0x3F, we cannot use IN/OUT instructions (which only reach 0x00-0x3F). We use STS (store to SRAM address) instead. The header `<avr/io.h>` defines all register names for us.

So even in this tiny program, the hardware layout is already shaping the instruction choices. We are not using STS by style preference; we are using it because that is the instruction family that can reach these addresses.

1.5.2 The Code

```
#include <avr/io.h>

/* — Vector table ————— */
.section .vectors, "ax", @progbits
.global __vectors

__vectors:
    jmp     main           ; reset vector: first instruction after power-on
```

When the ATtiny3217 resets, execution begins at **byte address 0x0000**. The CPU does not “look for main” by name. It simply fetches whatever instruction is stored at the reset vector and executes it. We place a JMP main there so the first thing the chip does is branch into our real program.

That is an important mental shift if you are coming from C: function names are for the assembler and linker, not for the CPU. The CPU only sees addresses and instructions.

```
.section .text

main:
    /* — Stack pointer init ————— */
    ldi     r16, lo8(RAMEND)    ; RAMEND = 0x3FFF → lo8 = 0xFF
    out     SPL, r16
    ldi     r16, hi8(RAMEND)    ; hi8 = 0x3F
    out     SPH, r16
```

SP (the stack pointer) must be valid before any PUSH, CALL, or RCALL. If it is wrong, the CPU will store return addresses and saved registers in the wrong place, and the program will collapse quickly.

On the ATtiny3217, reset hardware already loads SP = RAMEND, so these writes are technically redundant in this exact program. We still do them because they make the startup

state explicit, and because this is the pattern you need on many older AVR parts where the hardware does not fully set the stack up for you.

RAMEND expands to 0x3FFF (last byte of ATtiny3217 SRAM). `lo8()` and `hi8()` are assembler expressions that extract the low and high bytes of a 16-bit value.

SPL and SPH are in low I/O space (addresses 0x3D and 0x3E), so we reach them with OUT.

```
ldi    r16, (1 << 3)      ; bitmask for PA3 = 0x08
sts    PORTA_DIRSET, r16   ; set PA3 as output
```

LDI loads a constant into a register. STS stores that register value to a memory-mapped peripheral address. PORTA_DIRSET expands to 0x0401, and writing 0x08 there means “set bit 3 of the direction register.”

In plain language: we are telling the chip that PA3 is now an output pin. We are not yet turning the LED on or off; we are only telling the hardware that this pin should be driven by the CPU rather than treated as an input.

loop:

```
ldi    r16, (1 << 3)
sts    PORTA_OUTTGL, r16   ; atomically toggle PA3
```

PORTA_OUTTGL flips PA3’s output state. If it was high, it becomes low; if it was low, it becomes high. The useful part is that the flip happens inside the hardware peripheral. We do not have to read the old state first, and there is no risk of losing a concurrent change while doing a software read-modify-write.

```
ldi    r19, 25
delay_outer:
ldi    r25, 0
ldi    r24, 0
delay_inner:
sbw    r24, 1              ; subtract 1 from r25:r24 (2 cycles)
brne   delay_inner        ; branch if not zero      (2 cycles)
dec    r19
brne   delay_outer
rjmp   loop
```

SBIW subtracts an immediate from a 16-bit register pair. Here, r25:r24 acts as the inner delay counter. It counts down from 65535 to 0. Each time it reaches zero, BRNE stops looping, DEC r19 reduces the outer counter by one, and the whole process repeats until the outer loop also finishes.

This is not a precise timer. It is a simple busy-wait delay that burns CPU cycles. That makes it ideal for a first example because you can understand it with nothing more than arithmetic and branches.

After reset, the ATtiny3217 does not normally run at the full oscillator rate. It starts from OSC20M with a **divide-by-6 prescaler**. That gives **3.3 MHz** if the FUSE.OSCCFG frequency select is 20 MHz, or **2.66 MHz** if it is 16 MHz. Assuming the 20 MHz setting, the loop timing is roughly: $25 \times 65536 \times 4 \text{ cycles} \div 3,330,000 \approx 2 \text{ seconds}$ per toggle (about 1 Hz blink). If you later change the clock prescaler, oscillator source, or fuse setting, the blink rate changes proportionally.

That last point is worth remembering early: software delay loops are only as accurate as your clock configuration. Change the clock, and the visible timing changes too.

Instruction introduced: LDI LDI Rd, K — Load immediate value K into register Rd. Rd must be r16–r31. 1 cycle. Affects no flags.

Instruction introduced: OUT OUT A, Rr — Write register Rr to I/O address A (0x00–0x3F). 1 cycle. Affects no flags.

Instruction introduced: STS STS *k*, *Rr* — Store register *Rr* to SRAM address *k* (0x0000–0xFFFF). 2 cycles. Affects no flags. 32-bit instruction (two 16-bit words).

Instruction introduced: SBIW SBIW *Rd*, *K* — Subtract immediate *K* (0–63) from register pair *Rd+1:Rd*. Valid pairs: *r25:r24*, *r27:r26* (X), *r29:r28* (Y), *r31:r30* (Z). 2 cycles. Affects C, Z, N, V, S flags.

Instruction introduced: BRNE BRNE *k* — Branch if Z flag is clear (last result was not zero). 2 cycles if taken, 1 cycle if not. Range: ± 63 words from current PC.

Instruction introduced: DEC DEC *Rd* — Decrement register *Rd* by 1. 1 cycle. Affects Z, N, V, S flags. Does NOT affect C flag — careful when using in multi-byte arithmetic.

Instruction introduced: RJMP RJMP *k* — Relative jump. $PC = PC + k + 1$. Range: ± 2047 words. 2 cycles. Affects no flags.

Instruction introduced: JMP JMP *k* — Absolute jump to any address in 64K word program space. 3 cycles. 32-bit instruction. Affects no flags.

1.6 Building and Flashing

1.6.1 Step 1: Assemble

```
cd book/ch01_intro/src
avr-as -mmcu=attiny3217 -o blink.o blink.S
```

`-mmcu=attiny3217` tells the tools exactly which device we are targeting. That choice affects instruction availability, predefined symbols, and which register definitions appear through `#include <avr/io.h>`.

No output on success. On error, `avr-as` prints the filename, line number, and problem.

1.6.2 Step 2: Link

```
avr-gcc -mmcu=attiny3217 -nostartfiles -o blink.elf blink.o
```

The linker is the stage where your program stops being a pile of pieces and becomes a real image with fixed addresses. It decides where `.vectors` and `.text` live in flash and resolves any symbol references between sections.

`-nostartfiles` suppresses the C runtime startup (`crt0.o`). Our `.vectors` section and `main` label take its place.

1.6.3 Step 3: Convert to HEX

```
avr-objcopy -O ihex blink.elf blink.hex
```

The `.elf` is the rich build artifact for tools and humans. The `.hex` is the lean transport format for programming the chip. Same program, different packaging.

1.6.4 Step 4: Inspect (optional but recommended)

```
avr-objdump -d blink.elf
```

Sample output:

```
blink.elf:      file format elf32-avr
```

Disassembly of section .vectors:

```
00000000 <__vectors>:
   0:  0c 94 02 00      jmp      0x4 <__ctors_end>
```

Disassembly of section .text:

```
00000004 <__ctors_end>:
   4:  0f ef            ldi      r16, 0xFF          ; RAMEND low byte
   6:  0d bf            out      0x3d, r16          ; SPL
   8:  0f e3            ldi      r16, 0x3F          ; RAMEND high byte
  a:  0e bf            out      0x3e, r16          ; SPH
  c:  08 e0            ldi      r16, 0x08          ; PIN3 mask
  e:  00 93 01 04      sts      0x0401, r16        ; PORTA_DIRSET
  ...
```

This is one of the healthiest habits you can build early: do not blindly trust source code. Inspect what the toolchain actually produced. `avr-objdump` shows the addresses, the machine-code bytes, and the final instruction sequence that will land in flash.

On this device/linker setup, `avr-objdump` may label the first `.text` address as `__ctors_end`; that is still the same address as our main label in this minimal program.

1.6.5 Step 5: Flash

ATTiny3217 Curiosity Nano (onboard debugger via USB):

```
pymcuprog write -d attiny3217 -t uart -f blink.hex
```

UPDI adapter (e.g. FT232R or CH340 with SerialUPDI):

```
avrdude -c serialupdi -p attiny3217 -P /dev/ttyUSB0 -b 115200 \
-U flash:w:blink.hex:i
```

Successful flash output:

```
avrdude: AVR device initialized and ready to accept instructions
avrdude: writing flash (40 bytes):
avrdude: 40 bytes of flash written
avrdude: verifying flash memory against blink.hex:
avrdude: 40 bytes of flash verified
avrdude done. Thank you.
```

The LED on PA3 should now blink at approximately 1 Hz.

1.7 A Minimal Makefile

Typing the build steps by hand is useful once because it teaches the pipeline. Typing them forever is just friction. Save this as `ch01_intro/src/Makefile`:

```
MCU      = attiny3217
TARGET   = blink
SRCS     = blink.S

AS       = avr-as
CC       = avr-gcc
OBJCOPY  = avr-objcopy
```

```

OBJDUMP = avr-objdump
PROG     = serialupdi

ASFLAGS = -mmcu=$(MCU)
LDFLAGS = -mmcu=$(MCU) -nostartfiles

.PHONY: all flash disasm clean

all: $(TARGET).hex

$(TARGET).o: $(TARGET).S
    $(AS) $(ASFLAGS) -o $@ $<

$(TARGET).elf: $(TARGET).o
    $(CC) $(LDFLAGS) -o $@ $<

$(TARGET).hex: $(TARGET).elf
    $(OBJCOPY) -O ihex $< $@

flash: $(TARGET).hex
    pymcuprog write -d $(MCU) -t uart -f $<

disasm: $(TARGET).elf
    $(OBJDUMP) -d $<


clean:
    rm -f $(TARGET).o $(TARGET).elf $(TARGET).hex

```

Now build with `make`, flash with `make flash`, inspect with `make disasm`.

1.8 What Just Happened — Mental Model

Before moving on, fix this picture in your mind. This is the first real mental model of the whole book:



CPU reads instructions from Flash.

CPU reads/writes variables in SRAM.

Registers (r0–r31) are inside the CPU – fastest possible storage.

The program counter starts at 0x0000 after reset, fetches the `JMP` at the reset vector, lands in `main`, runs straight through the setup code, and then spins forever around `loop`. Nothing mystical is happening. The CPU is just fetching, executing, and updating state one instruction at a time.

1.9 Summary

Concept	Key point
Assembly	One instruction = one CPU operation. No abstraction.
AVR-GAS	GNU Assembler, invoked as <code>avr-as</code> . GAS syntax.
Build pipeline	<code>.S</code> → assemble → <code>.o</code> → link → <code>.elf</code> → convert → <code>.hex</code> → flash
ATtiny3217 PORT	New-style: memory-mapped registers, use STS/LDS. OUTTGL is atomic.
Stack pointer	Must be valid before any subroutine call; on ATtiny3217 reset already loads <code>SP = RAMEND</code> .
Default clock	$\text{OSC20M} \div 6$ after reset: 3.3 MHz or 2.66 MHz, depending on <code>FUSE.OSCCFG</code> .

1.10 Exercises

1. Change the blink rate to 2 Hz (halve the delay loop count). Recalculate the required `r19` value.
2. Use `PORTA_OUTSET` and `PORTA_OUTCLR` instead of `PORTA_OUTTGL`. Make the LED on for a longer period than it is off (asymmetric blink).
3. Run `avr-objdump -d blink.elf` and find:
 - The byte address of the `rjmp loop` instruction.
 - The encoding of `sts 0x0407, r16` — how many bytes does it take?
 - The total size of the `.vectors` section.
4. Add a second LED on PA5. Both LEDs should blink in opposite phases (one on while the other is off).

Next: Chapter 2 — AVR Architecture Overview

Chapter 2


AVR Architecture Overview

Before writing a single instruction, you need a mental model of what the AVR CPU is actually doing: where code lives, where data lives, what the registers are for, and what hidden state changes after each instruction. Without that model, assembly feels like memorising magic spells. With it, most instructions become straightforward and predictable.

The target device in this book, the ATtiny3217, is a useful concrete example: it is an 8-bit AVR running at up to 20 MHz, with 32 KB of flash, 2 KB of SRAM, 256 bytes of EEPROM, a hardware multiplier, and a single-pin UPDI programming and debugging interface. Those numbers matter because assembly programming is always resource-aware. A “small” stack frame or lookup table is only small relative to the actual memory on the chip.

2.1 Harvard Architecture

The AVR is a **Harvard** architecture. This means it has **two separate memory spaces** — one for instructions (flash) and one for data (SRAM) — connected to the CPU by separate buses.



Why does this matter to the programmer?

Because on AVR, “memory” is not one big uniform place. Where a value lives determines how you access it.

When your program is running, the CPU fetches instructions from **flash** automatically. You do not manually load the next instruction yourself; the program counter and instruction fetch hardware handle that in the background. That is why your code can live safely in non-volatile flash while your working data changes constantly in SRAM.

Your **variables**, your **stack**, and most of the peripheral register space you work with as data live in **SRAM/data space**. That is the memory area you read and write with ordinary load/store instructions such as LD, LDS, ST, and STS. So if you push a register, call a sub-routine, store a loop counter, or write to a peripheral register, you are operating in the data space, not in program flash.

Flash is different. It holds your program, and on many AVR parts it also holds constant tables or strings that you do not want to waste SRAM on. But flash is not ordinary read/write RAM. You cannot treat it like a normal destination for ST and expect the program image to change. On the ATtiny3217, writes to flash and EEPROM go through the Nonvolatile Memory Controller (NVMCTRL), using protected command sequences covered later in the book. Older classic AVR devices expose flash self-programming through SPM; do not assume that model applies unchanged to tinyAVR 1-series devices.

This also affects **read-only data** such as lookup tables. On classic AVR, reading data back from flash requires LPM, because flash is on the program side of the machine. On AVRxt parts such as the ATtiny3217, flash is also mapped into the data address space, so some reads can be done through that mapping with ordinary data-memory instructions. Chapter 5 shows both styles and when each applies.

The practical rule is simple:

- If the value is temporary, changing, or part of the stack, think **SRAM/data space**.
- If the value is your code or a fixed table that should survive power loss, think **flash/program space**.

This separation is why you cannot simply think “an address is an address.” On AVR, you must use the instruction that matches the address space you are trying to access.

2.2 ATtiny3217 Memory Map

The ATtiny3217 uses the **AVRxt** core (the GNU toolchain names this architecture `avrxtmega3`). It has an important difference from classic AVR (ATmega328P): **flash is also mapped into the data address space** at offset 0x8000. This means you can read flash with ordinary LD instructions when you use the mapped address — no need for LPM for that access path.

That feature is convenient, but do not let it blur the architectural picture. Flash is still program memory. The chip is simply giving you a second way to read it. Thinking in terms of “program space” versus “data space” is still the right model.

Data Address Space (what LD/ST/LDS/STS see):

0x0000	I/O registers (IN/OUT range)	0x0000–0x003F (64 bytes)
0x003F		
0x0040	Extended I/O / Peripherals (LDS/STS only, not IN/OUT)	0x0040–0x0FFF
0x0FFF		
0x1000	NVM I/O registers and data	0x1000–0x13FF
0x13FF		
0x1400	EEPROM – 256 bytes	0x1400–0x14FF
0x14FF		
0x1500	(reserved)	
0x37FF		

0x3800	SRAM – 2 KB	0x3800–0x3FFF
0x3FFF		
0x4000	(reserved)	
0x7FFF		
0x8000	Flash – 32 KB (read-only) mapped into data space	0x8000–0xFFFF
0xFFFF		

Program Address Space (what the PC addresses, separate bus):

0x0000	Flash – 16K words (32 KB) Byte addr = word addr × 2	
0x3FFF		

(word addresses; byte addresses go to 0x7FFF)

Key numbers to memorise:

Symbol	Value	Meaning
E2START / EEPROM_START	0x1400	First EEPROM byte in data space
E2END / EEPROM_END	0x14FF	Last EEPROM byte
FLASHEND	0x7FFF	Last byte address in flash
RAMSTART	0x3800	First SRAM byte
RAMEND	0x3FFF	Last SRAM byte (init SP here)
MAPPED_PROGMEM_START	0x8000	Flash base in data address space

Key point: When you see a linker script or header file refer to `MAPPED_PROGMEM_START`, it means 0x8000 — where flash appears in the data bus. On classic AVR, flash is *not* in the data bus at all.

If you are new to memory maps, read the diagram like this:

- The **low addresses** are the control area: CPU and peripheral registers.
- The **NVM/EEPROM region** is non-volatile data and device metadata, not general scratch RAM.
- The **middle SRAM block** is your working memory.
- The **high mapped-flash block** is read-only program memory made visible from the data side on AVRxt devices.

That one picture explains why some operations use IN/OUT, some use LDS/STS, and some use flash-specific mechanisms.

There are three common ways an address can appear in real code:

What you mean	Address base	Typical instruction
Program flash as code	0x0000 word address	PC fetch, RJMP, CALL
Program flash as bytes via program bus	0x0000 byte address	LPM with Z
Program flash mapped as data	0x8000 byte address	LD, LDS

This is the source of many off-by-two and off-by-0x8000 mistakes. If a label is used as a branch target, the CPU treats it as a program address. If the same label is used as table data, you must be clear about whether the code is using LPM addressing or mapped-data addressing.

2.3 The Register File

The AVR CPU has **32 general-purpose 8-bit registers**, named r0 through r31. They are directly connected to the ALU. In one CPU clock, the core can fetch register operands, execute an ALU operation, and write the result back to the register file. Compare that to SRAM: a load (LDS) takes 3 cycles; a store (STS) takes 2.

This is one of the main reasons assembly on AVR can be so efficient. If you can keep a value in a register, you avoid repeated trips to SRAM. A large part of “good AVR assembly” is really just good register management.

The ATtiny3217 instruction set has 135 instructions and includes a hardware multiplier. Multiplication is still slower than ordinary register ALU work, but it is a fixed two-cycle hardware operation rather than a long software shift-and-add routine. Chapter 6 uses that distinction when comparing arithmetic strategies.

Register file layout:

r0		General purpose. Some instructions can only use r0
r1		(e.g. MUL stores result in r1:r0).
r2		
r3		r0 is often used as a temporary by the compiler.
...		By convention, keep r1 = 0x00 at all times when
r15		calling C code (GCC ABI requirement).
r16		← LDI can only load immediate into r16–r31.
r17		Most two-operand arithmetic works on any register,
...		but immediate operands restrict you to r16–r31.
r23		
r24		SBIW/ADIW work on r25:r24, r27:r26, r29:r28, r31:r30.
r25		
r26	XL	] X pointer register (r27:r26) – indirect addressing
r27	XH	
r28	YL	] Y pointer register (r29:r28) – indirect + stack frame
r29	YH	
r30	ZL	] Z pointer register (r31:r30) – indirect + flash read
r31	ZH	

2.3.1 Register Pairs (16-bit Pointers)

The top six registers form three **pointer pairs** used for indirect memory access:

Pair	Registers	Name	Special use
X	r27:r26	X-pointer	Indirect load/store
Y	r29:r28	Y-pointer	Indirect load/store, stack frame pointer
Z	r31:r30	Z-pointer	Indirect load/store, LPM, indirect jumps/calls

When you see LD r0, X in assembly, it means “load the byte at the address stored in the X register pair into r0.” We cover this in depth in Chapter 5.

In practice, think of X, Y, and Z as your built-in address registers. The ordinary registers hold values; the pointer registers hold locations.

2.3.2 LDI Restriction

LDI — load immediate — **only works on r16–r31**. This is one of the most common beginner mistakes.

```
ldi r5, 42      ; ERROR: illegal operand – r5 not allowed with LDI
ldi r16, 42     ; OK
```

If you need a constant in r0–r15, load it into r16–r31 first, then MOV it:

```
ldi r16, 42
mov r5, r16     ; now r5 = 42
```


That may feel arbitrary at first, but it is simply part of the AVR instruction encoding. Some instructions have fewer bits available to describe operands, so they only work on part of the register file. You will see this kind of limit often in small instruction-set architectures.

2.4 The Status Register (SREG)

SREG is an 8-bit register at I/O address **0x3F** (accessible with IN/OUT). Every arithmetic and logic instruction updates some of its bits. Branch instructions read these bits to decide whether to jump.

You can think of SREG as the CPU’s short-term memory of “what just happened.” Did the last result become zero? Did it overflow? Was there a carry? Branches do not inspect your registers directly; they inspect these flag bits.

SREG bit layout:



2.4.1 Flag by Flag

C — Carry Set when an addition produces a result > 255 (unsigned overflow) or a subtraction borrows. Also used by shift instructions (the bit shifted out). Used for multi-byte arithmetic (ADC, SBC).

```
ldi r16, 0xFF
ldi r17, 0x01
add r16, r17    ; result = 0x00, C=1 (carry out), Z=1
```

Z — Zero Set when the result of an operation is exactly zero. The most-used flag. BRNE (branch if not equal) branches when Z=0; BREQ branches when Z=1.

```
ldi r16, 1
dec r16         ; r16 = 0, Z=1
```

```
brne somewhere ; NOT taken – Z is set
breq somewhere ; taken – Z is set
```

N — Negative Set when bit 7 of the result is 1 (result is negative in two’s complement).

V — Overflow Set when two’s-complement arithmetic overflows. For example, adding two positive numbers and getting a negative result.

```
ldi r16, 0x7F ; +127 (largest positive signed 8-bit)
ldi r17, 0x01
add r16, r17 ; result = 0x80 = -128, V=1 (signed overflow)
```

S — Sign $S = N \text{ XOR } V$. This is the “true” sign of the result after accounting for overflow. Use BRGE/BRLT (which test S) for signed comparisons, not N alone.

H — Half-carry Carry out of bit 3 into bit 4. It matters for nibble-oriented arithmetic and some compare/add/subtract cases. AVR does **not** have a DAA instruction, so you usually inspect H only in specialised code.

T — Transfer A single-bit scratch pad. BST copies a bit from a register into T; BLD copies T into a register bit. Useful for bit manipulation without clobbering other flags.

I — Global Interrupt Enable When I=1, the CPU responds to interrupts. SEI sets it; CLI clears it. Entering an ISR automatically clears I; RETI restores it. Chapter 10 covers this in detail.

2.4.2 Reading and Saving SREG

SREG is a normal I/O register. You can save and restore it:

```
in r0, SREG ; save SREG
cli ; disable interrupts (clear I bit)
; ... critical section ...
out SREG, r0 ; restore SREG (re-enables interrupts if I was set)
```

This pattern appears constantly in ISR prologues and critical sections.

If you remember only one thing about SREG, remember this: flags are part of program state. If your code depends on them, preserve them deliberately. The CPU does not automatically save or restore SREG when entering or leaving an ISR; that is your software’s responsibility.

Instructions introduced: IN / OUT (revisited) IN Rd, A — Load I/O register at address A into Rd. 1 cycle. OUT A, Rr — Write Rr to I/O register at address A. 1 cycle. Range: A = 0x00-0x3F only. For higher addresses, use LDS/STS.

Instruction introduced: SEI SEI — Set Global Interrupt Enable (I bit in SREG). 1 cycle.


Instruction introduced: CLI CLI — Clear Global Interrupt Enable. 1 cycle. Next instruction always executes before any interrupt is served (pipeline guarantee).

2.5 The Stack and Stack Pointer

The stack is a **last-in, first-out** (LIFO) region of SRAM. The AVR stack **grows downward** — pushing data decrements the stack pointer first, then writes.

That “grows downward” detail matters. On AVR, the top of SRAM is the natural starting point for the stack, and each push moves toward lower addresses. Beginners often imagine the stack growing upward because that feels more intuitive. AVR does the opposite.

SRAM at 0x3800–0x3FFF:



SP is a 16-bit register implemented as two I/O registers: - SPL at I/O address 0x3D (low byte of SP) - SPH at I/O address 0x3E (high byte of SP)

On the ATtiny3217, reset hardware loads SP with RAMEND automatically. You will still see startup code write it explicitly, because that is harmless and keeps examples portable to older AVR parts where software initialisation is required.

So there are really two questions:

- What value should SP start with? RAMEND.
- Who sets it? On ATtiny3217, the hardware does; on some older AVR parts, startup code must do it.

Both bytes are in the IN/OUT address range, so explicit initialisation uses OUT:

```
ldi r16, lo8(RAMEND)    ; 0xFF
out SPL, r16
ldi r16, hi8(RAMEND)    ; 0x3F
out SPH, r16
```

There is one subtle hardware assist here: when software writes SPL, the CPU automatically disables interrupts for up to four instructions, or until the next I/O-memory write, whichever comes first. That prevents an interrupt from observing a half-updated 16-bit stack pointer between the low-byte and high-byte writes. It is still good practice to initialise the stack before enabling interrupts or making subroutine calls.

2.5.1 PUSH and POP

```
push r16    ; SP -= 1, then write r16 to SRAM[SP] (2 cycles)
pop  r16    ; read r16 from SRAM[SP], then SP += 1 (2 cycles)
```

Push decrements first, then writes. Pop reads first, then increments.

2.5.2 How CALL Uses the Stack

CALL / RCALL automatically pushes the **return address** onto the stack so RET can find it:

```
rcall my_function    ; push PC+1, jump to my_function
; RET in my_function pops the return address and jumps back here
```

On ATtiny3217 (16K word address space), return addresses are 2 bytes (16-bit). The return address stored on the stack is a **word pointer**, not a byte pointer. The least-significant byte is pushed first, at the higher stack address, then the most-significant byte is pushed at the next lower address. On devices with >128KB flash, return addresses are 3 bytes. Always account for this when calculating stack depth.

This is why even a program with “no local variables” still uses stack space. Every nested call consumes stack bytes for the return address, and any saved registers add to that cost.

2.6 The Program Counter (PC)

The Program Counter is a register inside the CPU that holds the **word address** of the next instruction to execute. It is not directly accessible via IN/OUT — you control it only through jump, branch, and call instructions.

In other words, the PC is where control flow lives. Arithmetic changes data; branches and calls change the PC.

```
After reset:      PC = 0x0000    (first word in flash)
After JMP main:   PC = address of first instruction in main
After RCALL foo:  PC = address of foo; old PC+1 on stack
After RET:        PC = popped from stack
```

The PC is 14 bits wide on ATtiny3217 (addressing 0x0000–0x3FFF in word units, covering the full 32KB flash).

Most AVR instructions are one 16-bit word. Some instructions, such as JMP, CALL, LDS, and STS, are two-word instructions. This matters for three reasons:

- The PC counts instruction words, not bytes.
- Skip instructions must skip either one word or two words depending on the next instruction.
- Disassembly addresses may be shown as byte addresses by tools, while the CPU return address is stored as a word pointer.

2.6.1 Instruction Fetch Pipeline

AVR uses a **single-stage pipeline**: while the CPU executes instruction N, it fetches instruction N+1 from flash. This is why most instructions execute in 1 cycle — the fetch is free. Instructions that change the PC (branches, jumps, calls, returns) take 2–4 cycles because the pipeline must be flushed and a new fetch begun.

At the device level, this is the basis for the usual AVR rule of thumb: straight-line register code can approach 1 MIPS per MHz of CLK_CPU. That is not a promise that every instruction is single-cycle; it is a reminder that branches, memory accesses, calls, returns, and peripheral waits are where the extra cycles usually appear.

You do not need to become a pipeline expert to use AVR well. The practical lesson is simpler: straight-line code is cheap, but changing direction costs extra cycles.

```
Cycle:   1       2       3       4
Fetch:   N       N+1     N+2     N+3
Execute:      N       N+1     N+2   ← 1 cycle/instruction (normal)
```

With a taken branch at instruction N:

```
Fetch:   N       N+1     target  target+1
Execute:      N       ---      target ← N+1 fetch wasted; branch costs 2 cycles
```

This is why BRNE costs 1 cycle when not taken (no pipeline flush) and 2 cycles when taken (flush + re-fetch from target).

2.7 I/O Register Space

The ATtiny3217 peripherals are controlled through memory-mapped I/O registers. They fall into two zones:

Address range	Access method	Examples
0x0000–0x003F	IN / OUT (fast)	VPORTB_OUT, SREG, SPL, SPH, CCP
0x0040–0x0FFF	LDS / STS (2 words)	PORTB_DIR, TCA0_CNT, USART0_RXDATA1
0x1000–0x13FF	LDS / STS	NVMCTRL, SIGROW, FUSES, USERROW

The low 64 addresses (0x00–0x3F) are the “classic” I/O space reachable with 1-word IN/OUT instructions in 1 cycle. The low half of that space (0x00–0x1F) is also reachable by SBI, CBI, SBIC, and SBIS. Everything above 0x3F needs the 2-word LDS/STS instructions, which take 3 and 2 cycles respectively.

This is one of the most important performance distinctions on AVR. Two registers may both control hardware, but if one is in low I/O space and the other is in extended space, the instruction choices and timing differ.

The ATtiny3217’s important low I/O registers:

I/O Addr	Name	Description
0x05	VPORTB_OUT	Virtual PORTB output register (fast GPIO access)
0x04	VPORTB_DIR	Virtual PORTB direction register
0x3F	SREG	Status Register
0x3E	SPH	Stack Pointer High
0x3D	SPL	Stack Pointer Low
0x34	CCP	Configuration Change Protection (unlock protected registers)
0x1F	GPOR3	General Purpose I/O Register 3 (scratch)
0x1E	GPOR2	General Purpose I/O Register 2 (scratch)
0x1D	GPOR1	General Purpose I/O Register 1 (scratch)
0x1C	GPOR0	General Purpose I/O Register 0 (scratch)

VPORTx_* mirrors let you do single-cycle GPIO with IN/OUT and bit instructions, while the full PORTx_* peripheral block provides the richer DIRSET/OUTCLR/OUTTGL registers in extended I/O space.

GPOR0–3 are four registers with no peripheral function — pure scratch RAM accessible via IN/OUT in a single cycle. Useful for storing flags or counters that need fast ISR access.

Think of them as “fast little mailboxes” between parts of your program. They are not magic, but they can be handy when one-byte state needs to be very cheap to access.

The extended I/O region is where most modern tinyAVR peripherals live. A few important base addresses are:

Base	Peripheral	Why you care
0x0400	PORTA	Full GPIO configuration and atomic set/clear/toggle
0x0600	ADC0	ADC and Peripheral Touch Controller path
0x0800	USART0	Serial port used in later examples
0x0810	TWI0	I2C-compatible bus controller
0x0820	SPI0	SPI controller
0x0A00	TCA0	16-bit timer/counter
0x1000	NVMCTRL	Flash/EEPROM controller
0x1100	SIGROW	Factory calibration and signature data
0x1280	FUSES	Device configuration fuses

This is why examples often prefer header symbols such as USART0_STATUS or PORTA_DIRSET instead of raw numbers. The symbol tells you which peripheral block you are touching and lets the device header carry the exact address.

2.8 Putting It Together — A Register-State Walk-Through

Let's trace the CPU state through the blink program from Chapter 1, cycle by cycle through the setup section:

The goal here is not to memorise every state transition. It is to get used to thinking in terms of registers, memory side effects, and whether flags changed.

```
; State at reset:
; PC = 0x0000
; SP = 0x3FFF on ATtiny3217 reset hardware
; SREG = 0x00
; Treat general registers as undefined until your code sets them

ldi r16, lo8(RAMEND)    ; r16 = 0xFF. PC → next instr. 1 cycle.
out SPL, r16           ; explicit re-init of SP low byte. 1 cycle.
ldi r16, hi8(RAMEND)    ; r16 = 0x3F. PC → next. 1 cycle.
out SPH, r16           ; explicit re-init of SP high byte. 1 cycle.

ldi r16, (1 << 3)      ; r16 = 0x08. 1 cycle.
sts PORTA_DIRSET, r16   ; Write 0x08 to 0x0401.
                        ; PA3 now output. 2 cycles.

; State now:
; r16 = 0x08, SP = 0x3FFF, PORTA_DIR bit 3 = 1, all flags unchanged
```

Flags: LDI, OUT, STS do **not** affect SREG at all. Only arithmetic, logic, and bit operations change flags.

2.9 Key Architectural Constraints Summary

These are the constraints you will bump into constantly. They look dry, but they explain a large share of beginner errors.

Constraint	Rule
LDI range	r16-r31 only
IN/OUT range	I/O addresses 0x00-0x3F only
SBIW/ADIW pairs	r25:r24, r27:r26, r29:r28, r31:r30 only
Pointer pairs	X = r27:r26, Y = r29:r28, Z = r31:r30
MUL result	Always in r1:r0
Stack direction	Grows downward; PUSH decrements first
Return address	2 bytes on stack (ATtiny3217, ≤128KB flash)
Interrupt response	6 cycles minimum on ATtiny3217 flash sizes > 8 KB
Flash read via LD	Use address + 0x8000 (mapped flash, AVRxt)
Flash read via LPM	Use Z = byte address, reads via program bus
EEPROM address	0x1400-0x14FF in data space, controlled by NVMCTRL

Constraint	Rule
Flash/EEPROM writes	Use NVMCTRL command sequences, not ordinary SRAM stores

2.10 Summary

Harvard = two buses. Flash for code. SRAM for data.

ATtiny3217 memory map:

```

0x0000–0x003F  I/O (IN/OUT)
0x0040–0x0FFF  Peripherals (LDS/STS)
0x1000–0x13FF  NVMCTRL / SIGROW / FUSES / USERROW region
0x1400–0x14FF  EEPROM (256 bytes)
0x3800–0x3FFF  SRAM (2 KB)
0x8000–0xFFFF  Flash mapped in data space (AVRxt feature)

```

32 registers (r0–r31):

```

r0–r15: general, no LDI
r16–r31: general + LDI + most immediates
r26:r27 = X, r28:r29 = Y, r30:r31 = Z (pointer pairs)

```

SREG (0x3F): C Z N V S H T I – set by ALU, read by branches.

SP = SPH:SPL, initialise to RAMEND. Stack grows down.

On ATtiny3217, reset hardware already loads SP = RAMEND.

PC = word address, 14 bits. Pipeline prefetch = 1 cycle/instr normal.

2.11 Exercises

1. What is the data-bus address of the first byte of flash on ATtiny3217? What is the data-bus address of the last byte of flash?
2. After `ldi r16, 0x80` and `ldi r17, 0x80` and `add r16, r17`, which SREG flags are set? Work it out by hand, then verify with:

```

avr-as -mmcu=attiny3217 -o /tmp/t.o - <<'EOF'
.text
ldi r16, 0x80
ldi r17, 0x80
add r16, r17
EOF
avr-objdump -d /tmp/t.o

```

(The disassembly won't show flag values — you must reason through them.)

3. Write a short assembly snippet that:
 - Saves SREG to r0
 - Disables interrupts
 - Does some work (e.g. `ldi r16, 42`)
 - Restores SREG (and thus the interrupt state)
4. Calculate the maximum stack depth for a program that calls a function, which calls another function. How many bytes does the stack use just for return addresses?
5. Why can't you write `ldi r15, 0xFF`? What two-instruction sequence achieves the same result?

Next: Chapter 3 — AVR-AS Syntax and Directives

Chapter 3

AVR-AS Syntax and Directives

Assembly source is more than a list of instructions. Before the CPU can run anything, the assembler has to answer basic questions: where does the code go, what data should be emitted, which names are constants, and which labels refer to real addresses? Those decisions are controlled by **directives** — assembler commands that begin with a dot.

This chapter is about that “language around the instructions.” If Chapter 1 showed how to make the chip do something, Chapter 3 explains how to write assembly source in a way the toolchain can understand and organise correctly.

3.1 File Extensions: .S vs .s

GAS accepts two extensions:

Extension	Preprocessor	Use
.S	Yes — C preprocessor runs first	Always use this
.s	No — raw GAS input	Rarely; only for generated files

The uppercase .S lets you use `#include`, `#define`, `#ifdef`, and all other C preprocessor directives. Since we always want `#include <avr/io.h>` for register names, **always use .S**.

For a beginner, the practical rule is simple: if you are hand-writing AVR assembly for this book, name the file with an uppercase .S and do not think about it further.

The build flow preprocesses .S files before assembly. If you want to inspect that preprocessed output yourself, use the compiler driver:

```
# Inspect the preprocessed source that avr-as will assemble:
avr-gcc -x assembler-with-cpp -mmcu=attiny3217 -E blink.S
```

You normally do not need to run this manually if your build uses `avr-gcc` as the driver, because it sees .S and runs the preprocessor before invoking the assembler. `avr-as` itself does **not** run `cpp`; if you call `avr-as` directly, feed it preprocessed assembly or avoid `#include/#define`.

For hand-written examples in this book, prefer this pattern:

```
avr-gcc -x assembler-with-cpp -mmcu=attiny3217 -c blink.S -o blink.o
```

That keeps device-header symbols such as `PORTA_DIRSET`, `RAMEND`, `SREG`, `SPL`, `SPH`, and `MAPPED_PROGMEM_START` available in the source.

3.2 Comments

GAS accepts three comment styles. In .S files all three work:

```
; This is a comment – semicolon style, common in AVR assembly

// This is a comment – C++ line style

/* This is a
   block comment – C style */
```

Semicolon is traditional in AVR assembly and what you’ll see most. Use block comments for file headers and section separators.

Style matters more in assembly than in many other languages because the code is so compact. A short comment can carry a lot of explanatory weight.

3.3 Instruction Syntax

AVR GAS uses a **destination-first, source-second** convention:

```
mnemonic destination, source
```

Examples:

```
mov r16, r17      ; r16 ← r17
add r16, r17      ; r16 ← r16 + r17
ldi r16, 42       ; r16 ← 42
sts 0x0421, r16    ; mem[0x0421] ← r16
ld r16, X         ; r16 ← mem[X]
```

This matches GCC output and the Atmel/Microchip instruction set manual.

Instruction syntax is not just cosmetic; it reflects fields in the opcode. On AVR, many instructions cannot encode every register or address. The assembler will reject operands that do not fit the instruction format:

```
ldi r15, 0x55      ; ERROR: LDI can only encode r16-r31
ldi r16, 0x55      ; OK

out 0x3F, r16      ; OK: OUT reaches I/O address 0x00-0x3F
out 0x40, r16      ; ERROR: use STS for extended I/O
sts 0x0040, r16    ; OK: direct store to data space

sbi 0x1C, 0        ; OK: SBI/CBI/SBIC/SBIS reach 0x00-0x1F
sbi 0x20, 0        ; ERROR: above bit-addressable low I/O range
```

Microchip’s ATtiny3217 peripheral map puts the four general-purpose I/O registers at 0x1C-0x1F, which makes them directly usable with SBI, CBI, SBIS, and SBIC. Most larger peripherals, such as PORTA at 0x0400 and USART0 at 0x0800, are outside the IN/OUT range and require LDS/STS or pointer-based LD/ST.

3.4 Number Literals

GAS accepts numbers in four bases:

```
ldi r16, 42           ; decimal
ldi r16, 0x2A         ; hexadecimal
ldi r16, 0b00101010   ; binary
ldi r16, 052          ; octal (leading zero – easy to mistake for decimal!)
```

All four produce the same value (42). Avoid octal — the leading-zero trap causes real bugs. Hex and binary are clearest for hardware work.

Character literals also work:

```
ldi r16, 'A'          ; r16 = 65
ldi r16, '\n'          ; r16 = 10
```

3.5 Expressions

GAS evaluates constant expressions at assembly time. You can use:

```
ldi r16, (1 << 4)      ; bit shift – evaluates to 16
ldi r16, (1 << 4) | (1 << 2) ; bitwise OR – evaluates to 20
ldi r16, RAMEND & 0xFF ; mask – evaluates to 0xFF
ldi r16, RAMEND / 256   ; division – evaluates to 0x3F
ldi r16, 10 * 5 - 8     ; arithmetic – evaluates to 42
```

These are evaluated at **assembly time** — no runtime cost. They exist only to make source code readable.

That is an important distinction. These expressions do not generate arithmetic instructions. The assembler does the math once while building the file, and the CPU never sees the expression itself.

3.5.1 Built-in Expression Functions

GAS provides helper functions for splitting 16-bit values into bytes:

```
ldi r16, lo8(RAMEND)    ; low byte of 0x3FFF = 0xFF
ldi r16, hi8(RAMEND)    ; high byte of 0x3FFF = 0x3F
ldi r16, lo8(0x1234)    ; = 0x34
ldi r16, hi8(0x1234)    ; = 0x12
```

Two more functions matter for AVRxt flash access (see Section 3.11):

```
; pm(label) – convert byte address to word address (for IJMP/ICALL targets)
ldi ZL, lo8(pm(my_func)) ; ZL = word address low byte of my_func
ldi ZH, hi8(pm(my_func)) ; ZH = word address high byte of my_func
```

For table addresses, be explicit about which address space you mean:

```
ldi ZL, lo8(table)      ; flash byte address for LPM
ldi ZH, hi8(table)

ldi ZL, lo8(table + MAPPED_PROGMEM_START) ; data-space address for LD
ldi ZH, hi8(table + MAPPED_PROGMEM_START)
```

The ATtiny3217 datasheet describes both views: LPM sees flash starting at 0x0000, while ordinary data-space loads see mapped flash starting at 0x8000. Those two forms are not interchangeable.

3.6 Labels

A label marks a position in the output. It ends with a colon:

```
main:
    ldi r16, 0xFF

loop:
    dec r16
    brne loop           ; reference label by name – no colon when referencing
```

Labels have two uses: 1. **Jump/branch targets** — as above 2. **Data addresses** — when placed before `.byte`, `.word`, etc.

You can think of a label as a bookmark the assembler turns into an address. Sometimes that address is where execution should jump. Sometimes it is where data begins.

```
table:
    .byte 1, 2, 3, 4      ; "table" now equals the address of byte 1
```

3.6.1 Global vs Local Labels

By default, labels are **local to the object file** — the linker cannot see them unless you explicitly export them:

```
.global main           ; export "main" so the linker can use it as entry point
.global __vectors      ; export vector table symbol

main:                  ; now visible to linker
    ...

internal_helper:      ; NOT global – invisible outside this file
    ...
```

For functions, you can also tell the linker/debugger what kind of symbol you are defining:

```
.global main
.type   main, @function
main:
    ; ...
    ret
.size   main, .-main
```

`.type` and `.size` do not change the emitted instructions. They improve symbol meta-data, which makes `avr-objdump`, `avr-nm`, and debuggers report function boundaries more cleanly.

3.6.2 GAS Local Labels (Numeric)

Naming every loop, branch target, and temporary label creates a pollution of unique names. GAS solves this with **numeric local labels**:

```
1:                ; define label "1"
    dec r16
    brne 1b        ; "1b" = most recent backward occurrence of label "1"

    tst r16
    breq 2f        ; "2f" = next forward occurrence of label "2"
    ldi r16, 0xFF
2:                ; define label "2"
    ...
```


Rules: - Any number 0–9 (or longer: 10:, 255:) can be a local label - Nb = nearest backward occurrence of N: - Nf = nearest forward occurrence of N: - The same number can be reused anywhere in the file — each definition is distinct. GAS resolves 1b to the *nearest* preceding 1:.

This is invaluable in loops:

```
ldi r18, 10      ; outer loop count
1:              ; outer loop top
    ldi r17, 0
    ldi r16, 0
2:              ; inner loop top
    sbiw r16, 1
    brne 2b      ; back to inner "2:"
    dec r18
    brne 1b      ; back to outer "1:"
```

No naming conflict — 1b and 2b refer to the right labels unambiguously.

This style looks strange at first, but it becomes natural quickly. Named labels are best for important program locations such as `main` or `uart_tx`. Numeric labels are best for short internal loops where a descriptive name would add more clutter than clarity.

3.6.3 File-Scoped Local Labels (`.L` prefix)

Labels starting with `.L` are stripped from the symbol table by the linker and do not appear in debug info. Use them for internal subroutine labels you don't want to export:

```
my_subroutine:
.Lloop:
    dec r16
    brne .Lloop
    ret
```

`.Lloop` will not appear in `avr-nm` output or debugger symbol lists.

3.7 Sections

A **section** is a named region of output. The linker arranges sections into the final memory map. Every assembly statement belongs to the *current* section.

If labels are bookmarks, sections are containers. They tell the linker what kind of thing you are defining: executable code, read-only flash data, SRAM data, zero-initialised storage, and so on.

3.7.1 Core Sections

```
.section .text      ; code – goes to flash
.section .data      ; initialised SRAM data – startup code must copy it in
.section .bss       ; zero-initialised SRAM data – startup code must clear it
.section .rodata    ; read-only constants – stay in flash
```

For the ATtiny3217, the linker maps those abstract sections onto a concrete device memory map:

Section	Runtime location	ATtiny3217 address range
.text, .rodata, .vectors	Flash/program memory	Code space 0x0000-0x7FFF bytes
.data	SRAM at runtime, initial image in flash	SRAM 0x3800-0x3FFF
.bss	SRAM, no file contents	SRAM 0x3800-0x3FFF
EEPROM data	EEPROM, NVMCTRL-controlled	Data space 0x1400-0x14FF
Mapped flash reads	Flash through data bus	Data space 0x8000-0xFFFF

The assembler does not know your intent beyond the section and bytes you emit. The linker script is what makes .text land in flash and .bss reserve SRAM. If you bypass the normal linker script or use custom sections, inspect the linker map and disassembly.

Shorthand forms (equivalent):

```
.text      ; same as .section .text
.data      ; same as .section .data
.bss       ; same as .section .bss
```

You can switch between sections multiple times in one file:

```
.text
foo:
    ldi r16, 5
    ret

.data
my_var: .byte 0

.text
bar:
    ldi r16, 10
    ret
```

The assembler concatenates all .text chunks together, all .data chunks together, etc. The order of sections in the output is determined by the linker script, not the source order.

That is why switching sections in the middle of a file is normal. You are not describing the final physical layout directly; you are describing which bytes belong to which category.

3.7.2 Section Flags and Types (Full Form)

The full .section directive takes optional flags:

.section name, "flags", @type

Flag	Meaning
a	Allocatable (takes space in memory)
x	Executable
w	Writable

Type	Meaning
@progbits	Contains data/code

Type	Meaning
@nobits	Takes space but has no data in file (BSS)

Most of the time you use the shorthand and GAS fills in reasonable defaults. You need the full form for special sections like `.vectors`:

```
.section .vectors, "ax", @progbits
```

This tells the linker: place this section at the start of flash (the linker script maps `.vectors` to address 0x0000), it contains code (x), and it is allocatable (a).

The vector table is executable code on this AVR target, not a table of raw ISR addresses. Each vector slot contains an instruction, commonly `jmp` handler on ATtiny3217-sized flash parts. That is why `.vectors` uses "ax" rather than plain read-only data flags.

You do not need to memorise all the flags immediately. What matters is knowing that the full form exists when you need something more specific than the common defaults.

3.8 Data Directives

These directives emit bytes into the current section.

That phrase “emit bytes” is worth taking literally. Directives such as `.byte` and `.word` do not create variables in the high-level-language sense. They cause exact byte patterns to be written into the object file.

3.8.1 `.byte` — 8-bit values

```
.byte 0x42           ; one byte: 0x42
.byte 1, 2, 3, 4     ; four bytes
.byte 'H', 'i', '!'  ; character literals
.byte (1 << 3)       ; expression – evaluates to 8
```

3.8.2 `.word` — 16-bit values (little-endian)

```
.word 0x1234         ; emits 0x34, 0x12 (little-endian)
.word 0, 1, 2        ; six bytes: three 16-bit words
```

Warning: On AVR, `.word` is 16 bits. In non-AVR GAS targets `.word` can be 32 bits. Never assume word size when reading unfamiliar assembly.

In flash tables used by `LPM`, remember that bytes still come out in little endian order. If you write `.word 0x1234`, the first byte at the lower address is 0x34 and the second is 0x12. If `Z` points at the label and you execute `lpm r16, Z+`, `r16` receives 0x34 first.

3.8.3 `.long` — 32-bit values

```
.long 0x12345678     ; emits 78 56 34 12 (little-endian)
```

3.8.4 `.ascii` and `.asciz` — Strings

```
.ascii "Hello"       ; 5 bytes: H e l l o – NO null terminator
.asciz "Hello"       ; 6 bytes: H e l l o 0x00 – WITH null terminator
.string "Hello"      ; alias for .asciz
```

Almost always use `.asciz` or `.string` so C string functions work correctly.

If you forget the terminating zero, the bytes are still valid assembly data, but any code that expects a C-style string will keep reading past the intended end.

3.8.5 `.skip` and `.space` — Reserve Bytes

```
buf:    .skip 16           ; reserve 16 bytes (content undefined in .text/.data)
buf:    .skip 16, 0        ; reserve 16 bytes, filled with 0x00
```

In `.bss`, content is always zeroed by startup code regardless of fill value.

Do not use `.skip` in `.text` as a cheap delay or timing pad. It emits bytes into flash; those bytes will be interpreted as instructions if execution falls through them. Use real instructions such as `nop`, or branch around data.

3.8.6 `.balign` — Alignment

```
.balign 2           ; align to 2-byte boundary (pad with 0x00 if needed)
.balign 4           ; align to 4-byte boundary
```

Required before `.word` or `.long` data in some contexts. Also used to align code for performance (though on AVR this matters less than on ARM).

AVR instructions are 16-bit words, and some are 32-bit two-word instructions. Keeping code and 16-bit flash data on even byte boundaries makes disassembly and label arithmetic easier to reason about. It also avoids accidental odd-byte entry points when mixing `.byte` tables and executable code in the same section.

3.9 Constant Directives

3.9.1 `.equ` — Define a Constant

```
.equ LED_PIN, 4
.equ LED_MASK, (1 << LED_PIN) ; expressions work
.equ DELAY, 5000
```

`.equ` defines a **symbol** with a value. It cannot be redefined later. Use it for things that truly never change: pin numbers, masks, magic values.

In practice, `.equ` is how you turn anonymous numeric constants into named ideas. `0x08` is just a number; `LED_MASK` tells the reader what the number means.

Using the constant:

```
ldi r16, LED_MASK ; assembler substitutes the value at assembly time
```

Prefer device-header names over hand-written `.equ` addresses whenever the name exists:

```
#include <avr/io.h>

ldi r16, (1 << 3)
sts PORTA_DIRSET, r16 ; better than .equ PORTA_DIRSET, 0x0401
```

The Microchip device header follows the datasheet address map. Using the header keeps source code tied to the selected `-mmcu=attiny3217` target instead of a copied constant that may be wrong on a different AVR.

3.9.2 .set — Reassignable Constant

```
.set count, 0
.set count, count + 1 ; OK — .set can be redefined
.set count, count + 1 ; count = 2 now
```

.set is useful for building values incrementally or for conditional definitions. Equivalent to #define + #undef patterns but within GAS.

3.9.3 #define — Preprocessor Macro

Because .S files go through cpp, you can also use:

```
#define LED_PIN 4
#define LED_MASK (1 << LED_PIN)
#define SET_BIT(reg, bit) ((reg) |= (1 << (bit))) ; functional macro
```

#define is processed *before* GAS sees the file; .equ is processed by GAS. For simple constants, .equ is idiomatic in assembly. For complex macros or conditional compilation, #define/#ifdef are necessary.

If you are unsure which to choose, prefer .equ for plain assembly constants. Reach for #define only when you really need preprocessor behavior.

3.10 The .org Directive

.org sets the **location counter** — the current output offset within the section. It can move the location counter **forward**, creating padding, but it cannot move it backward.

This is a more advanced tool than it first appears. It is useful when a layout must match hardware-defined slots, such as interrupt vectors, but it is also an easy way to create confusing gaps if used casually.

```
.org 0x0000
    jmp main ; lands exactly at byte address 0

.org 0x0004
    jmp isr_int0 ; next 4-byte JMP slot

.org 0x0008
    jmp isr_pcint0 ; next slot
```

If you use JMP, each vector entry consumes 4 bytes. If you use RJMP, each entry consumes 2 bytes, so the .org spacing changes accordingly. The exact vector table is in the datasheet and in <avr/iotn3217.h> as _VECTOR(n) macros. Chapter 10 builds the full vector table.

The ATtiny3217 has 32 KB of flash, so the normal interrupt vector entries use JMP-sized slots in the examples. Microchip also documents a Compact Vector Table mode through CPUINT.CTRLA.CVT; when enabled, the interrupt table groups normal interrupts into three entries: NMI at vector address 1, priority level 1 at vector address 2, and all priority level 0 interrupts at vector address 3. This book uses the ordinary vector table unless a later chapter explicitly says otherwise.

So the real lesson is not “always use .org.” The lesson is: use it when the hardware requires exact placement, and understand the instruction size when you do.

Warning: .org without context is relative to the *start of the current section*. In .vectors, .org 0 means “start of the vector section”, which the linker places at flash address 0. In .text, .org 0x100 would create a 0x100-byte gap from the section’s start — usually not what you want.

3.11 Flash Data on AVRxt (ATtiny3217)

This is where AVRxt differs from classic AVR in an important way.

This section is easy to skim past, but it matters because it changes how you think about constants in flash. On older AVR parts, flash reads feel like a special case. On AVRxt, they can look much more like ordinary data reads.

Classic AVR (ATmega328P): Flash is NOT in the data address space. To read a byte from flash at runtime, you must use the LPM instruction with the Z register pointing to the flash byte address.

AVRxt (ATtiny3217): Flash IS mapped in the data address space at offset `MAPPED_PROGMEM_START` (= 0x8000). You can read flash with ordinary LD instructions by adding this offset to the flash address.

The concrete ATtiny3217 memory map behind that statement is:

0x0000-0x003F	I/O registers
0x0040-0x0FFF	Extended I/O / peripheral registers
0x1000-0x13FF	NVM I/O registers and data
0x1400-0x14FF	EEPROM, 256 bytes
0x3800-0x3FFF	SRAM, 2 KB
0x8000-0xFFFF	32 KB flash mapped into data space

Only the last range is the mapped view of program flash. EEPROM is a separate non-volatile memory region controlled by `NVMCTRL`; do not put ordinary `.rodata` there by accident.

```
.section .rodata

table:
    .byte 10, 20, 30, 40, 50

.section .text

read_table_entry:
    ; r24 = index, returns value in r25
    ldi ZL, lo8(table + MAPPED_PROGMEM_START)
    ldi ZH, hi8(table + MAPPED_PROGMEM_START)
    add ZL, r24
    ; Z now points to table[r24] in data space
    adc ZH, r1
    ; propagate carry (r1 = 0 by convention)
    ld r25, Z
    ; read from flash via data bus
    ret
```

The assembler expression `table + MAPPED_PROGMEM_START` adds 0x8000 to the flash byte address of `table`. The result fits in 16 bits because the table starts in the first 32KB of flash.

Data-space LD access to mapped flash is convenient, but it is not SRAM. The AVR instruction set manual notes that data-memory cycle counts for LD assume internal RAM and that NVM accesses can add implementation-dependent time. For constant tables this usually does not change the code structure, but it matters when doing tight cycle accounting.

Conceptually, you are taking a flash address and translating it into the corresponding data-space view of the same bytes.

Old LPM method still works on AVRxt and is sometimes needed for compatibility:

```
ldi ZL, lo8(table)
ldi ZH, hi8(table)
lpm r25, Z
```

; Z = flash byte address (no +0x8000)
; read from flash via program bus

Both methods read the same data. On AVRxt, LD with mapped address is simpler and more natural. LPM is covered in Chapter 5.

That means you should understand both styles, even if the mapped-flash method is the one you use most often on the ATtiny3217.

3.12 The .include Directive

```
.include "macros.inc"           ; include a GAS file
```

Inserts the contents of another file at the point of the .include. The included file is processed by GAS (not cpp), so it cannot use #define.

For files that need preprocessor features, use #include instead:

```
#include <avr/io.h>             ; cpp include – register definitions
#include "my_macros.h"          ; cpp include – your own header
#include "data_tables.S"        ; GAS include – another assembly file
```

Use #include <avr/io.h> exactly once near the top of a .S file. That header selects the device-specific header from -mmcu=attiny3217; including <avr/iotn3217.h> directly works, but it makes the source less portable across build-system target changes.

3.13 Macro Directives

GAS has its own macro system separate from the C preprocessor.

```
.macro delay_ms ms_count
    ldi r19, \ms_count
.Lmacro_outer_\@:
    ldi r18, 0
    ldi r17, 0
.Lmacro_inner_\@:
    sbiw r17, 1
    brne .Lmacro_inner_\@
    dec r19
    brne .Lmacro_outer_\@
.endm
```

Using the macro:

```
delay_ms 100                    ; expands inline – \ms_count → 100
delay_ms 500                    ; second expansion – \@ ensures unique label names
```

Key syntax: - \param — substitute parameter value - \@ — unique invocation counter (prevents label collisions across expansions) - .endm — end macro definition

GAS macros are powerful but can obscure what code is actually emitted. Prefer subroutines (rcall/ret) over macros for anything that runs more than once, to save flash space.

This is a good place to be disciplined. Macros are tempting because they feel convenient, but they duplicate code at every use site. In assembly, that cost is often visible.

3.14 Full Section Layout Example

This is the structure of a well-organised AVR assembly file:

```

/*
 * myprogram.S – description
 * Build: avr-gcc -x assembler-with-cpp -mmcu=attiny3217 -c myprogram.S
 */

#include <avr/io.h>

/* — Constants — */
.equ    MY_PIN,    5
.equ    MY_MASK,   (1 << MY_PIN)

/* — Read-only data (flash) — */
.section .rodata

my_table:
    .byte 1, 2, 4, 8, 16, 32, 64, 128

my_string:
    .asciz "ATtiny3217"

/* — Uninitialised SRAM — */
.section .bss

rx_buffer:    .skip 64      ; 64-byte ring buffer
rx_head:      .skip 1
rx_tail:      .skip 1

/* — Initialised SRAM (requires startup copy code or CRT startup) — */
.section .data

state:        .byte 0x01    ; starts as 1 in SRAM

/* — Interrupt vector table — */
.section .vectors, "ax", @progbits
.global __vectors

__vectors:
    jmp      main           ; 0x0000 reset

/* — Code — */
.section .text
.global main

main:
    ldi r16, lo8(RAMEND)
    out SPL, r16
    ldi r16, hi8(RAMEND)
    out SPH, r16
    ; ... rest of program

```

That example deliberately uses `.data` and `.bss` to show the sections, but a bare -nostartfiles assembly program must provide the startup behavior itself. If you define `.data`, somebody must copy its initial bytes from flash to SRAM. If you define `.bss`, somebody must clear it. The `avr-libc` C runtime normally does that before `main`; a minimal hand-written reset handler does not unless you write the code.

3.15 Examining the Output with `avr-objdump`

After assembling, inspect what the directives produced:

```
avr-gcc -x assembler-with-cpp -mmcu=attiny3217 -c syntax_demo.S -o syntax_demo.o
avr-gcc -mmcu=attiny3217 -nostartfiles -o syntax_demo.elf syntax_demo.o
avr-objdump -s -d syntax_demo.elf
```

The `-s` flag dumps the **raw contents** of every section as hex. Example output for `.rodata`:

Contents of section `.rodata`:

```
0074 01000101 01000101 01010100 48656c6c .....Hell
0084 6f204156 5200                                o AVR.
```

Cross-reference that output with the source: `morse_dit = [01, 00]`, `morse_dah = [01, 01, 01, 00]`, `greeting = "Hello AVR\0"`. Each byte appears in order. This is how you confirm that your directives produced exactly the layout you intended.

The `-d` flag shows the **disassembly** of `.text`:

```
00000008 <main>:
      8: 0f ef          ldi     r16, 0xFF
      a: 0d bf          out     0x3d, r16
      ...
     1a: e0 e7          ldi     r30, 0x70    ; ZL = lo8(morse_dit + 0x8000)
     1c: f0 e8          ldi     r31, 0x80    ; ZH = hi8(...)
```

Notice `ZH = 0x80` — the `0x8000` mapped flash offset is baked into the immediate at assembly time.

That is a perfect example of why inspecting output is useful: you can see a source-level idea become a real machine-level value.

Two other inspection commands are useful when debugging directives:

```
avr-objdump -h syntax_demo.elf    # section sizes, VMAs, file offsets
avr-nm -n syntax_demo.elf         # symbols sorted by address
```

Use `-h` to confirm that `.bss` takes SRAM space without occupying file bytes. Use `avr-nm` to confirm which labels survived linking; `.L` local labels should not clutter the public symbol list.

3.16 Quick Reference

Directive	Purpose	Example
<code>.equ</code>	Define constant	<code>.equ PIN, 4</code>
<code>.set</code>	Reassignable constant	<code>.set i, 0</code>
<code>.byte</code>	Emit bytes	<code>.byte 1, 2, 3</code>
<code>.word</code>	Emit 16-bit words	<code>.word 0x1234</code>
<code>.long</code>	Emit 32-bit words	<code>.long 0xDEADBEEF</code>
<code>.ascii</code>	String, no null	<code>.ascii "hi"</code>
<code>.asciz</code>	String with null	<code>.asciz "hi"</code>
<code>.skip</code>	Reserve bytes	<code>.skip 16</code>
<code>.balign</code>	Align location counter	<code>.balign 2</code>
<code>.org</code>	Set location counter	<code>.org 0x0000</code>
<code>.global</code>	Export symbol	<code>.global main</code>

Directive	Purpose	Example
<code>.section</code>	Switch/create section	<code>.section .text</code>
<code>.include</code>	Include GAS file	<code>.include "f.S"</code>
<code>.macro/.endm</code>	Define GAS macro	<code>.macro name p1</code>
<code>.type</code>	Mark symbol kind	<code>.type main, @function</code>
<code>.size</code>	Mark symbol size	<code>.size main, .-main</code>

Expression	Meaning	Result for RAMEND=0x3FFF
<code>lo8(x)</code>	Low byte	0xFF
<code>hi8(x)</code>	High byte	0x3F
<code>pm(label)</code>	Word address	<code>addr / 2</code>
<code>(1 << n)</code>	Bit mask	<code>(1 << 4) = 0x10</code>

3.17 Summary

`.S` extension → `cpp` runs first → `#include` works.

Sections:

```
.text    = flash code
.rodata  = flash constants (read via mapped addr + 0x8000 on AVRxt)
.data    = SRAM, initialised if your startup code/CRT copies it
.bss     = SRAM, zero-init if your startup code/CRT clears it
.vectors = must be first in flash, holds the interrupt vector table
```

Labels:

```
name:    = define label
.Lname:  = local, stripped from symbol table
1:       = numeric – reusable, resolved by 1f / 1b
```

Constants: `.equ` (fixed), `.set` (reassignable), `#define` (preprocessor).

AVRxt flash read: address + `MAPPED_PROGMEM_START` (0x8000), then LD.
 ATtiny3217 data map: EEPROM is 0x1400-0x14FF; SRAM is 0x3800-0x3FFF.
 Low I/O: IN/OUT reach 0x00-0x3F; SBI/CBI/SBIS/SBIC reach 0x00-0x1F.

3.18 Exercises

1. Write a `.rodata` section containing a table of the first 8 powers of 2 (1, 2, 4, 8, ...). Write code that reads `table[r24]` using the mapped flash address and returns the value in `r25`.
2. Add a `.bss` variable `byte_count` and a `.data` variable `max_count` initialised to 10. Write code that increments `byte_count` using LDS/ STS until it equals `max_count`, then halts.
3. Rewrite the Chapter 1 blink delay using a GAS `.macro` named `delay` that takes one parameter (outer loop count). Show that two calls `delay 25` and `delay 50` both assemble without label conflicts.
4. Assemble `syntax_demo.S` and run:

```
avr-objdump -s syntax_demo.elf
```

Verify the bytes in `.rodata` match the `morse_dit`, `morse_dah`, and `greeting` definitions in the source.

5. What happens if you put `.word 0x1234` inside a `.bss` section? Try it and inspect the linker output. Why does the linker warn or error?

Next: Chapter 4 — Instruction Set Basics

Chapter 3A: Debugging AVR Assembly with GDB

Assembly debugging is different from source-level debugging in a high-level language. There is no hidden runtime to explain what happened. The evidence is the program counter, the registers, the status flags, SRAM, I/O registers, the stack pointer, and the instruction stream in flash.

This chapter shows how to debug the assembly style used in this book: GNU AVR assembly syntax in .S files, preprocessed through `avr-gcc` and assembled by `avr-as`. The target board is the ATtiny3217 Curiosity Nano, using its on-board debugger over UPDI. The core habit is simple: build an ELF with symbols, keep labels at meaningful locations, stop at instruction boundaries, inspect the CPU state, and prove what changed after each instruction.

3.19 Why This Chapter Belongs Here

By this point you have seen the source format, sections, labels, symbols, and the build flow that produces an ELF file. That is the right moment to learn debugging. If you wait until after timers, USART, SPI, interrupts, and bootloaders, you will have too many moving parts. If you learn GDB now, every later chapter becomes easier to verify.

The goal is not to memorize every GDB command. The goal is to build a repeatable inspection loop:

1. Build an ELF with useful debug information.
2. Load the ELF into GDB.
3. Stop at a named label.
4. Inspect registers, flags, memory, and disassembly.
5. Step one instruction.
6. Inspect again.
7. Explain exactly what changed.

That loop catches most assembly mistakes faster than guessing.

3.20 The Files GDB Needs

GDB works from an ELF file, not from the final Intel HEX file alone.

The HEX file contains bytes to program into flash. It is good for programming the device, but it does not carry the full symbolic view you want while debugging. The ELF file can contain:

- section addresses

- symbol names such as `main`, `sum_loop`, and `break_here`
- line information connecting source lines to machine instructions
- relocation and address information useful to `objdump`, `avr-nm`, and GDB

For this book's assembly examples, keep the ELF file and debug that:

```
avr-gcc -mmcu=attiny3217 -x assembler-with-cpp -g3 -Wa,--gdwarf-2 \
  -c -o gdb_demo.o src/gdb_demo.S

avr-gcc -mmcu=attiny3217 -nostartfiles -g3 \
  -o gdb_demo.elf gdb_demo.o
```

The important pieces are:

Option	Purpose
<code>-mmcu=attiny3217</code>	Selects the device memory layout and instruction variant.
<code>-x assembler-with-cpp</code>	Treats the file as preprocessed assembly.
<code>-g3</code>	Keeps rich debug information in the object and ELF files.
<code>-Wa,--gdwarf-2</code>	Asks the assembler to emit DWARF line information.
<code>-nostartfiles</code>	Keeps the example's reset/vector code under your control.

Do not strip the ELF you plan to debug. Stripping removes the symbolic information that makes GDB useful.

Use `avr-objdump` beside GDB. It gives a static view of exactly what was linked:

```
avr-objdump -h -t -d -S gdb_demo.elf
```

That command shows section placement, symbols, raw instructions, and source interleaving when line information is present.

3.21 Keep Symbols That Help Debugging

GDB is much easier to use when the assembly file exposes useful names.

Prefer this:

```
.global main
.type    main, @function

main:
    ; setup

sum_loop:
    ; loop body

break_here:
    nop
```

Avoid debugging a file where every important place is an anonymous local label:

```
1:
    ; body
    brne 1b
```

Numeric local labels are fine for tight, obvious control flow, but they are poor debugger landmarks. A professional assembly file has labels at state changes, loop entries, peripheral writes, interrupt entries, error traps, and deliberate debug stops.

For functions and major routines, add `.type name, @function`. For data you intend to inspect, give it a symbol and place it in a named section:

```
.section .bss
.global sum_result
sum_result:
.zero 1
```

Then GDB can inspect it by name:

```
(gdb) x/1xb &sum_result
```

3.22 Starting GDB

Use the AVR-targeted GDB when your toolchain provides it:

```
avr-gdb gdb_demo.elf
```

This chapter assumes `avr-gdb`. It matches the rest of the AVR GNU toolchain and loads the AVR register set and disassembler directly.

Inside GDB, confirm the file:

```
(gdb) info files
(gdb) info functions
(gdb) info variables
```

GDB startup scripts can change behavior. The local GDB tutorial included with this book notes the normal startup-script order: system init file, user init file, then the current directory's `.gdbinit`. If a debug session behaves strangely, start once with startup scripts disabled:

```
avr-gdb --nx gdb_demo.elf
```

Useful startup settings for assembly work:

```
set pagination off
set confirm off
set radix 16
set disassemble-next-line on
display/i $pc
```

The most important line is `display/i $pc`. GDB's own manual recommends it when stepping by machine instruction because it keeps the next instruction visible after every stop.

3.23 Debugging Targets Used Here

This book assumes one assembly syntax and one hardware debug board:

- source files use GNU AVR assembly accepted by `avr-as`
- the build driver is `avr-gcc`
- the debug image is the linked ELF file
- the hardware target is the ATtiny3217 Curiosity Nano and its on-board debugger

3.23.1 1. Static ELF Inspection

This does not run the program. It only inspects the linked ELF:

```
avr-gdb gdb_demo.elf
```

Useful commands:

```
(gdb) disassemble /r main
(gdb) info address sum_loop
(gdb) info files
(gdb) x/12i main
```

This mode is still valuable. Many assembly bugs are visible before execution: wrong section, unexpected instruction encoding, missing symbol, incorrect vector placement, a branch target in the wrong place, or data linked into flash when it was meant for SRAM.

3.23.2 2. ATtiny3217 Curiosity Nano Hardware

The ATtiny3217 uses UPDI, the Unified Program and Debug Interface. Microchip describes UPDI as the single-pin programming and debugging interface for this device family. The ATtiny3217 product documentation lists the device as an 8-bit AVR running up to 20 MHz, with up to 32 KB flash, 2 KB SRAM, 256 bytes of EEPROM, and a single-pin UPDI programming/debug interface.

The ATtiny3217 Curiosity Nano hardware guide documents the board's on-board debugger as the programming and debugging path for the ATtiny3217 using UPDI. That same on-board debugger also provides a virtual serial port over UART and debug GPIO. Those are the observation paths assumed in the rest of this book.

Debugger-to-target connections on the ATtiny3217 Curiosity Nano:

ATtiny3217 pin	Debugger function	Book use
PA0	DBG0, UPDI	program/debug lifeline
PB2	CDC RX, ATtiny3217 USART0 TX	serial output to host
PB3	CDC TX, ATtiny3217 USART0 RX	serial input from host
PA3	debug GPIO1	LED0 and timing/debug marks
PB7	debug GPIO0	optional timing/debug marks

Microchip documents that these debugger connections are tri-stated when the debugger is not actively using the interface. The board also exposes the signals on the edge connector, but the default assumption here is the unmodified Curiosity Nano board with the on-board debugger still connected.

Do not cut the debugger straps for the examples in this book. Microchip notes that disconnecting those straps disables programming, debugging, virtual serial port, and data streaming for the ATtiny3217 mounted on the board.

For this book, the practical rule is simple: treat PA0/UPDI as a debug lifeline. Do not use it as an ordinary GPIO pin in examples unless you also explain how the board can be recovered.

3.23.3 3. PyAvrOCD Debugging Guide

PyAvrOCD is a GDB server for 8-bit AVR targets. GDB does not talk directly to the ATtiny3217 Curiosity Nano. GDB talks to a server over a local TCP port, and the server talks to the board's debug probe over USB. On this board the probe is the on-board nEDBG, and the MCU debug interface is UPDI.

The debug path is:

```

avr-gdb or a GDB GUI
    |
    | GDB remote serial protocol on localhost:2000
    v
pyavrocd
    |
    | USB to the Curiosity Nano nEDBG probe
    v
ATtiny3217 over UPDI on PA0

```

This means there are three separate pieces to debug when a session fails:

- the ELF file and symbols that GDB sees
- the PyAvrOCD server process and its command-line options
- the physical/debug connection from nEDBG to the ATtiny3217 UPDI pin

3.23.3.1 Install and Sanity Check

The PyAvrOCD documentation describes several install paths: release archives, Python package installation with `pipx` or `pip`, and running from the GitHub repository. Any of those are fine as long as `pyavrocd` and `avr-gdb` are available from the shell you use for the book examples.

Check the tools first:

```

pyavrocd --version
pyavrocd --help
avr-gdb --version

```

On Linux, PyAvrOCD may need `udev` rules before a normal user can access Microchip debug probes. If the board is plugged in and PyAvrOCD reports that no compatible tool was discovered, fix host USB permissions before changing code, fuses, or wiring.

3.23.3.2 Curiosity Nano Setup

The ATtiny3217 Curiosity Nano already has the necessary debug probe on the board. For the unmodified board:

- connect the board by USB
- leave the debugger straps intact
- leave PA0/UPDI available for programming and debugging
- do not add capacitive loads, strong resistive loads, or active circuitry to the UPDI line

External UPDI wiring is not needed for this board. PyAvrOCD can also work with external probes on other boards, but the Curiosity Nano case is simpler: USB to the board is the debug connection.

3.23.3.3 Build an ELF for Debugging

Debug the ELF file, not only the HEX file. The ELF carries the sections, symbols, line information, and addresses GDB needs.

Build the companion example from `book/ch03a_gdb_debugging`:

```

avr-gcc -mmcu=attiny3217 -x assembler-with-cpp -g3 -Wa,--gdwarf-2 \
    -c -o gdb_demo.o src/gdb_demo.S

avr-gcc -mmcu=attiny3217 -nostartfiles -g3 \
    -o gdb_demo.elf gdb_demo.o

```

For C or mixed C/assembly projects, PyAvrOCD's documentation recommends debug-friendly compiler settings: use debug information, prefer `-Og` while debugging, avoid

-flto where possible, and do not use -mrelax because it can distort line information. For pure assembly, the equivalent discipline is: keep named labels, keep line information, and do not let the linker rewrite jump layout while you are matching source, symbols, and addresses.

3.23.3.4 Start the PyAvrOCD Server

Run PyAvrOCD from the directory that contains gdb_demo.elf:

```
pyavrocd -d attiny3217 -i updi -t nedbg -p 2000 -e gdb_demo.elf
```

PyAvrOCD should connect to the probe, identify the target, start the server, and listen for a GDB connection on port 2000.

The options used here are deliberately explicit:

Option	Long form	Meaning
-d attiny3217	--device attiny3217	Selects the target MCU. This is the required option.
-i updi	--interface updi	Selects UPDI. This is the ATtiny3217 debug interface.
-t nedbg	--tool nedbg	Selects the Curiosity Nano on-board debugger.
-p 2000	--port 2000	Selects the TCP port where GDB connects. Port 2000 is the default.
-e gdb_demo.elf	--elf-file gdb_demo.elf	Lets PyAvrOCD inspect the ELF and reject known-bad debug setups such as -mrelax.

Only -d is mandatory in the general PyAvrOCD command line. This book keeps the other options visible because visible choices are easier to audit when several probes, ports, or target boards are attached.

Useful variations:

```
pyavrocd -d attiny3217 -i updi -t nedbg -p 2001 -e gdb_demo.elf
pyavrocd -d attiny3217 -i updi -t nedbg -e gdb_demo.elf --verbose debug
pyavrocd -d attiny3217 -i updi -t nedbg -e gdb_demo.elf --reboot-debugger
```

Use a different port if another server still owns port 2000. Use verbose logging when the server starts but the behavior is unclear. Rebooting the debugger can help when the host-side probe state is stale.

3.23.3.5 Connect GDB

Open a second terminal in the same directory:

```
avr-gdb --nx gdb_demo.elf
```

Then connect to PyAvrOCD:

```
set pagination off
set radix 16
set disassemble-next-line on
display/i $pc

target remote :2000
```

At this point GDB is connected to the server, but the target flash may still contain an old image. Load the ELF through PyAvrOCD:

```
load
```

The load command matters. It sends the ELF contents through PyAvrOCD so the target flash matches the symbols, addresses, and disassembly GDB is using. A common embedded-debugging mistake is to rebuild the ELF but forget to program the target. After that mistake, breakpoints and symbols describe one program while the chip runs another one.

Now break at a named label and run:

```
break main
break sum_loop
break break_here
continue
```

Then debug at instruction level:

```
si
info registers
x/8i $pc
x/1xb &sum_result
```

The same commands work from a GUI front end, because the GUI is still driving GDB underneath.

3.23.3.6 Optional: Start a GDB Front End

PyAvrOCD can start another program after the server is up:

```
pyavrocd -d attiny3217 -i updi -t nedbg -p 2000 -e gdb_demo.elf --start gde
```

--start gde launches Gede if it is installed and in PATH. Gede is a GDB front end; PyAvrOCD is still the GDB server. In Gede, configure the program as gdb_demo.elf, use avr-gdb as the debugger, connect to :2000, and run the same startup commands:

```
target remote :2000
load
break main
continue
```

Do not add monitor debugwire enable for the ATtiny3217 Curiosity Nano. That command is for debugWIRE targets. The ATtiny3217 uses UPDI.

3.23.3.7 Monitor Commands

GDB commands control GDB. monitor commands are passed through GDB to PyAvrOCD. They control server behavior.

Useful commands for this book:

Command	Use
monitor help	Print PyAvrOCD monitor help.
monitor info	Show target and debugger state.
monitor version	Show the PyAvrOCD version.
monitor reset	Reset the MCU.
monitor breakpoints all	Permit both software and hardware breakpoints.
monitor breakpoints hardware	Use only hardware breakpoints.

Command	Use
<code>monitor breakpoints software</code>	Use only software breakpoints.
<code>monitor load readbeforewrite</code>	Compare flash pages and skip unchanged writes when loading.
<code>monitor load writeonly</code>	Write flash pages without comparing first.
<code>monitor load noinitialload</code>	Skip the first load when the exact image is already present.
<code>monitor verify enable</code>	Verify flash after loading pages.
<code>monitor verify disable</code>	Disable flash verification.
<code>monitor atexit stay</code>	Leave the target in debug mode when the server exits.
<code>monitor atexit leave</code>	Leave debug mode on exit where the interface supports it.
<code>monitor singlestep safe</code>	Protect single stepping against interrupt surprises.
<code>monitor singlestep interruptible</code>	Allow interrupts to affect single stepping.

Most monitor settings can also be supplied as PyAvrOCD command-line options. For example:

```
pyavrocd -d attiny3217 -i updi -t nedbg -e gdb_demo.elf --verify enable
```

For UPDI targets, PyAvrOCD documents `readbeforewrite` as the default load mode and notes that timers are frozen when the CPU is stopped. Keep both facts in mind when timing-sensitive code behaves differently under debug.

3.23.3.8 Persistent Attach

The normal teaching workflow is a fresh session:

1. rebuild the ELF
2. restart PyAvrOCD
3. connect GDB
4. run load
5. set breakpoints by symbol

For bugs that appear only after the program has been running for a while, you may want to leave the target in debug mode and attach later. Before leaving the first GDB session:

```
monitor atexit stay
```

Later, restart PyAvrOCD with attach mode:

```
pyavrocd -d attiny3217 -i updi -t nedbg -p 2000 --attach
```

Then reconnect:

```
target remote :2000
```

Use attach mode deliberately. It is useful when you need it, but it is a poor default for examples because it preserves target state that a beginner may not remember creating.

3.23.3.9 Troubleshooting PyAvrOCD Sessions

If the server does not start:

- Confirm the board is attached by USB and powered.
- Confirm no other program is using the debug probe.
- On Linux, check udev rules and user permissions.
- If several probes are attached, specify `--tool nedbg` and, if needed, `--usbsn`.
- Use `--verbose debug` and read the first critical error.
- Try `--reboot-debugger` if the probe appears stuck.

If GDB cannot connect:

- Confirm PyAvrOCD is still running.
- Confirm the port number in `target remote :PORT` matches `--port`.
- Use a different port if the old one is still busy.
- Start GDB with `--nx` once to rule out startup-script effects.

If breakpoints do not hit:

- Run `load` after every rebuild.
- Use `info address label` and `disassemble /r label` to confirm the symbol.
- Break at named labels, not guessed numeric addresses.
- Avoid `-mrelax` while debugging.
- If a reset happens while running, set breakpoints again.
- If needed, try `monitor breakpoints hardware`.

If stepping or timing looks wrong:

- Remember that stopped debugging changes timing.
- Do not single-step timing-critical peripheral sequences as proof of real-time behavior.
- On UPDI targets, timers are frozen while the CPU is stopped.
- Use GPIO marks, USART output, a logic analyzer, or an oscilloscope for timing evidence.

If the image looks wrong:

- Rebuild the ELF.
- Restart PyAvrOCD.
- Reconnect GDB.
- Run `load`.
- Use `compare-sections` if your GDB/target combination supports it.
- Confirm with `avr-objdump -d -S gdb_demo.elf`.

3.23.3.10 A Repeatable Debug Script

You can put the GDB side of the session in a small command file, for example `gdb_demo.gdb`:

```
set pagination off
set radix 16
set disassemble-next-line on
display/i $pc

target remote :2000
load
break main
break sum_loop
break break_here
continue
```

Then start the server:

```
pyavrocd -d attiny3217 -i updi -t nedbg -p 2000 -e gdb_demo.elf
```

And start GDB:

```
avr-gdb --nx -x gdb_demo.gdb gdb_demo.elf
```

This makes the debug setup repeatable. If the script stops working, inspect the server output and the first failing GDB command before changing the assembly.

3.24 Core GDB Commands for Assembly

This table is the daily working set.

Command	Use
help command	Show built-in help for one command.
apropos word	Search command help by topic.
file firmware.elf	Load symbol and section information.
info files	Show loaded sections and entry information.
info functions	List known function symbols.
info variables	List known data symbols.
break label	Stop when execution reaches a label.
hbreak label	Request a hardware breakpoint when supported.
delete	Delete breakpoints.
continue	Resume execution until the next stop.
stepi or si	Execute one machine instruction, entering calls.
nexti or ni	Execute one machine instruction, stepping over calls when possible.
finish	Run until the current routine returns.
disassemble /r label	Show instructions and raw opcode bytes.
x/8i \$pc	Examine eight instructions at the program counter.
info registers	Show register state.
p/x \$pc	Print the program counter in hexadecimal.
x/16xb address	Examine 16 bytes of memory in hexadecimal.
set \$reg = value	Change a register, using the target's register names.
set {unsigned char}addr = value	Patch one byte in writable memory.

Register names are target-specific. The GDB manual explicitly points out that `info registers` shows the canonical names for the selected machine. On AVR, check `info registers` before assuming exactly how your GDB build names `r0`, `r1`, `SREG`, `SP`, or `PC`.

3.25 Instruction-Level Stepping

Use instruction stepping when debugging assembly:

```
(gdb) break main
(gdb) continue
(gdb) display/i $pc
(gdb) info registers
(gdb) stepi
(gdb) info registers
```

The difference between `stepi` and `nexti` matters:

Command	Behavior
<code>stepi / si</code>	Execute one instruction. If the instruction calls a subroutine, stop inside it.
<code>nexti / ni</code>	Execute one instruction, but try to step over subroutine calls.

For branch debugging, use `stepi` until you trust the flags. A conditional branch is only as correct as the flags set by the previous instruction.

Example:

```
    dec    r18
    brne   sum_loop
```

After `dec r18`, inspect SREG or the named flags your GDB exposes. The branch does not test `r18` directly; it tests the Z flag produced by `DEC`.

For call/return debugging, inspect the stack pointer before and after `CALL`, `RCALL`, `RET`, and `RETI`. On AVR, the return address is stored on the stack. The return-address size depends on the device's program counter width, so avoid hard-coding a stack-byte assumption across all AVR families. On the ATtiny3217, the flash size is small enough that ordinary examples use a two-byte return address.

3.26 Reading Disassembly

GDB has two common disassembly views:

```
(gdb) disassemble main
(gdb) disassemble /r main
```

The `/r` form shows raw instruction bytes as well as mnemonics. That is useful when you are checking instruction encoding, word order, or whether the linker relaxed a branch or call.

To look around the current program counter:

```
(gdb) x/10i $pc
```

To look at a label:

```
(gdb) x/12i sum_loop
```

Remember that AVR flash is word-oriented at the instruction-set level but ELF tools and GDB often present addresses in bytes. The instruction set manual gives instruction sizes in words; the debugger may display byte addresses. When an address appears to move by 2 for a single instruction, that usually means one 16-bit instruction word.

3.27 Inspecting Registers and Flags

The fastest way to debug AVR assembly is to predict register changes before stepping:

```
(gdb) info registers
(gdb) stepi
(gdb) info registers
```

For a small sequence:

```
ldi    r19, 0x11
ldi    r20, 0x22
add    r19, r20
```

You should predict:

- r19 becomes 0x33
- r20 stays 0x22
- SREG changes according to the ADD result
- PC advances by one instruction after each single-word instruction

For flag-heavy code, write down which instruction last set the flags. Common mistakes include:

- expecting LDI to set flags
- forgetting that INC, DEC, ADD, SUB, CP, CPI, LSL, and shifts affect flags
- placing a harmless-looking instruction between a compare and branch, then accidentally overwriting the flags
- treating carry as a generic error bit instead of the exact carry/borrow state produced by the last arithmetic instruction

When in doubt, single-step and read SREG.

3.28 Inspecting SRAM

Use named symbols for SRAM locations:

```
.section .bss
.global sum_result
sum_result:
    .zero 1
```

Then inspect the byte:

```
(gdb) x/1xb &sum_result
```

Other useful memory formats:

```
(gdb) x/16xb &sum_result    # 16 bytes, hex, byte width
(gdb) x/8xh  &sum_result    # 8 halfwords
(gdb) x/4xw  &sum_result    # 4 words in GDB's display format
```

For I/O registers, use the named address from the device header or a symbol you define yourself. Be aware that reading some peripheral registers has side effects. Status registers often clear flags by writing a one, and some data register reads advance FIFOs or clear interrupt states. A debugger memory window is not always passive.

This is especially important when debugging peripherals in later chapters:

- Read the datasheet before repeatedly inspecting a status or data register.
 - Know whether a flag is cleared by writing one.
 - Know whether a register is synchronized to a peripheral clock domain.
 - Know whether a register is protected by a configuration-change mechanism.
-

3.29 Inspecting Flash Data

AVR has separate program and data address spaces. Instructions live in flash. Ordinary SRAM variables live in data memory. Constants may live in flash and be read by LPM.

GDB can disassemble flash reliably because instructions are the normal execution stream:

```
(gdb) x/8i main
```

Raw flash data can be more target-dependent because different tools expose AVR program memory through different address mappings. If GDB's memory view is not obvious, confirm flash contents with `avr-objdump`:

```
avr-objdump -s -j .text gdb_demo.elf
```

For assembly examples, make flash data easy to find:

```
.section .progmem.data, "a", @progbits
.global table_bytes
table_bytes:
    .byte 0x11, 0x22, 0x33, 0x44
```

Then use both tools:

```
avr-nm -n gdb_demo.elf
avr-objdump -s -j .text gdb_demo.elf
```

The default AVR linker script places program-memory input sections into the final flash image, commonly the `.text` output section. That is why the linked ELF is inspected with `-j .text` even though the source used `.progmem.data`. GDB shows what the current target exposes. `objdump` shows what the ELF contains.

3.30 Breakpoints on Microcontrollers

Breakpoints are easy on a desktop process and less abstract on a microcontroller.

A software breakpoint usually means patching code memory with a trap instruction. On a flash-based MCU, that may require special handling because flash is not ordinary writable SRAM. A hardware breakpoint uses on-chip debug resources. Those resources are limited.

Practical rules:

- Prefer breakpoints at labels, not raw numeric addresses.
- Use `hbreak` when the target requires hardware breakpoints.
- Do not assume unlimited breakpoints on a small MCU.
- Keep a deliberate `break_here:` label near important end states.
- For unexpected interrupts, route unused vectors to a trap label such as `unexpected_interrupt: rjmp unexpected_interrupt`.
- Remove stale breakpoints when changing code layout.

On the Curiosity Nano, treat breakpoints as a scarce debug resource and keep your assembly easy to inspect statically with `avr-objdump` even when the hardware session is not attached.

3.31 Watchpoints and Data Breaks

Watchpoints stop when memory changes:

```
(gdb) watch sum_result
```

On embedded targets, watchpoint support depends on the CPU and debug hardware. If a watchpoint is implemented in hardware, there may only be a small number available. If it is implemented in software, it may slow execution heavily.

When watchpoints are limited, instrument the assembly instead:

```
store_result:
    sts    sum_result, r19
break_after_store:
    nop
```

Then break at `break_after_store`.

The label costs nothing in the final instruction stream unless you add an instruction for it. Labels are free debugging handles.

3.32 Debugging the Stack

Stack bugs are common in assembly because the assembler will not protect you from an unbalanced `PUSH`, `POP`, `CALL`, `RET`, interrupt entry, or manual stack write.

Before stepping through subroutines, initialize the stack pointer:

```
ldi    r16, lo8(RAMEND)
out     SPL, r16
ldi    r16, hi8(RAMEND)
out     SPH, r16
```

Then in GDB:

```
(gdb) info registers
(gdb) p/x $sp
(gdb) x/16xb $sp
```

Register naming can vary, so use `info registers` to confirm whether your GDB build calls the stack pointer `$sp`, shows `SPL/SPH`, or exposes both forms.

Checklist for stack debugging:

- Is `SP` initialized before the first `CALL`, `PUSH`, or interrupt?
- Does every `PUSH` have a matching `POP` on every exit path?
- Is `SREG` saved before enabling nested behavior or modifying flags in an ISR?
- Does the routine return with `RET` while an ISR returns with `RETI`?
- Did the code accidentally store ordinary data into the current stack area?

If a return jumps to nonsense, inspect the bytes at the stack pointer before the return instruction. The return address is the evidence.

3.33 Debugging Interrupts

Interrupt debugging requires more discipline than straight-line code.

Give every vector a visible target:

```
.section .vectors, "ax", @progbits
__vectors:
    jmp     main
    jmp     unexpected_interrupt
    jmp     tca0_ovf_isr
```

Give every ISR a visible entry and exit:

```
.global tca0_ovf_isr
.type   tca0_ovf_isr, @function

tca0_ovf_isr:
    push    r16
    in      r16, SREG
    push    r16

    ; ISR body

tca0_ovf_isr_exit:
    pop     r16
    out     SREG, r16
    pop     r16
    reti
```

Then you can use:

```
(gdb) break tca0_ovf_isr
(gdb) break tca0_ovf_isr_exit
(gdb) continue
```

At ISR entry, inspect:

- PC
- SP
- saved return address bytes
- SREG
- peripheral interrupt flag register
- interrupt enable register

At ISR exit, inspect:

- whether the peripheral flag was cleared correctly
- whether saved registers were restored
- whether SREG was restored
- whether the stack pointer returned to its entry value
- whether the final instruction is RETI, not RET

When an interrupt never fires, debug in this order:

1. Is the peripheral clocked?
 2. Is the peripheral configured?
 3. Is the peripheral interrupt flag set?
 4. Is the peripheral interrupt enable bit set?
 5. Is the global interrupt enable bit set in SREG?
 6. Is the vector table at the address the CPU is using?
 7. Does the vector point to the intended ISR label?
-

3.34 Debugging Peripheral Code

For GPIO, timers, USART, SPI, TWI, and NVM code, the debugger can mislead you if you inspect only the instruction stream. Peripheral state often changes outside the CPU pipeline.

Use a three-part method:

1. Step the instruction that writes the control register.
2. Inspect the written register if it is safe to read.
3. Confirm the external effect when possible.

For example, after writing `PORTA_DIRSET`, read back the relevant direction state. After writing `PORTA_OUTTGL`, inspect the output latch and check the pin or LED. After enabling a timer, inspect the timer's control register, then let time pass and inspect the count or interrupt flag.

For communication peripherals, combine GDB with a second observation path:

- virtual COM port for USART
- logic analyzer for SPI/TWI
- oscilloscope for timing-sensitive pins
- Microchip Data Visualizer when using supported board debug features

GDB tells you what the CPU did. External tools tell you what the circuit saw.

3.35 A Minimal Assembly Debug Example

The companion file `src/gdb_demo.S` is intentionally small. It:

- installs a reset vector
- initializes the stack pointer
- loads four flash bytes with LPM
- accumulates them into an 8-bit sum
- stores the result in SRAM
- stops at a named debug label

The flash table is deliberately placed after the halt loop in `.text`, not immediately after the reset vector. This example still installs only the reset vector because interrupts remain disabled, but the table is no longer sitting where a reader would expect interrupt-vector code.

Build it from this chapter directory:

```
avr-gcc -mmcu=attiny3217 -x assembler-with-cpp -g3 -Wa,--gdwarf-2 \
  -c -o gdb_demo.o src/gdb_demo.S

avr-gcc -mmcu=attiny3217 -nostartfiles -g3 \
  -o gdb_demo.elf gdb_demo.o
```

Inspect it statically:

```
avr-nm -n gdb_demo.elf
avr-objdump -d -S gdb_demo.elf
```

For static inspection, start GDB:

```
avr-gdb --nx gdb_demo.elf
```

Then run these commands:

```

set pagination off
set radix 16
set disassemble-next-line on
display/i $pc

break main
break sum_loop
break break_here

info files
disassemble /r main

```

For hardware confirmation, start PyAvrOCD in one terminal:

```
pyavrocd -d attiny3217 -i updi -t nedbg -p 2000 -e gdb_demo.elf
```

Then connect from GDB in a second terminal and load the ELF into target flash:

```

target remote :2000
load
break main
break sum_loop
break break_here
continue

```

The load command programs the target through the on-board debugger. If you also need a standalone HEX file for a separate programming tool, generate it explicitly:

```
avr-objcopy -O ihex gdb_demo.elf gdb_demo.hex
```

At `sum_loop`, step one instruction at a time:

```

info registers
si
info registers
x/8i $pc

```

At `break_here`, inspect the result:

```
x/1xb &sum_result
```

The expected result is `0xaa` because:

```
0x11 + 0x22 + 0x33 + 0x44 = 0xaa
```

If the result differs, debug from evidence:

- Was Z initialized to the flash table address?
- Did LPM r20, Z+ load the expected byte?
- Did ADD r19, r20 update the accumulator?
- Did DEC r18 set Z only on the final iteration?
- Did BRNE sum_loop branch exactly three times?
- Did STS sum_result, r19 write SRAM?

Do not guess. Step and inspect.

3.36 A Useful .gdbinit for AVR Assembly

For a project directory, this is a good starting point:

```

set pagination off
set confirm off
set radix 16
set disassemble-next-line on
display/i $pc

define regs
    info registers
end

define around
    x/8i $pc
end

define hook-stop
    x/i $pc
end

```

The hook-stop command runs every time execution stops. Keep it short. If it prints too much, stepping becomes noisy.

Start with `--nx` when you want to rule out startup-script effects:

```
avr-gdb --nx gdb_demo.elf
```

Start normally when you want the project helpers:

```
avr-gdb gdb_demo.elf
```

3.37 Failure Patterns and What to Inspect

Symptom	Likely area	First inspection
Program starts at the wrong code	vector table or entry address	<code>info files, x/4i 0</code>
Breakpoint never hits	wrong symbol, optimized/changed layout, wrong target image	<code>info break, info address label, compare-sections</code> if supported
Branch goes the wrong way	flags	<code>info registers</code> before and after the flag-setting instruction
Return jumps to nonsense	stack corruption	<code>p/x \$sp, x/16xb \$sp, inspect push/pop balance</code>
SRAM variable unchanged	store not reached or wrong address	break before STS, inspect pointer/address
LPM reads unexpected value	flash table address or address-space mapping	<code>avr-nm, avr-objdump -s, inspect Z</code>
ISR never runs	interrupt enable chain	peripheral flag, peripheral enable bit, global I bit, vector
Curiosity Nano debugger cannot connect	UPDI wiring/fuse/power	target VCC, GND, PA0/UPDI path, fuse state
Debug works once then stops	UPDI pin reused or fuse changed	check PA0 configuration and fuse writes

The point of the table is not to replace thinking. It gives you the first piece of evidence to collect.

3.38 Professional Debug Habits

Write assembly so it can be debugged:

- Use stable labels at important states.
- Keep one instruction per source line.
- Use named SRAM symbols for values you will inspect.
- Keep reset and interrupt vectors readable.
- Save and restore registers in a visible order.
- Keep deliberate trap labels for impossible paths.
- Build with debug information during development.
- Keep the ELF, map file, and objdump listing for every important test image.
- Record the debugger commands that prove a bug is fixed.

When a bug is hard, make the code more observable. Add a label. Store an intermediate value. Toggle a debug pin. Split a dense block into smaller named blocks. Assembly is not self-explaining unless you make the important states visible.

3.39 Source Notes

Facts in this chapter are based on:

- Microchip ATtiny3217 product documentation for device memory sizes and UPDI.
- Microchip ATtiny3217 Curiosity Nano Hardware User Guide sections describing the on-board debugger, UPDI programming/debugging, virtual serial port, and debug GPIO, including PA0/UPDI, PB2/PB3 CDC UART, PA3 debug GPIO1, and PB7 debug GPIO0 connections.
- Microchip ATtiny3217 Curiosity Nano Hardware User Guide sections describing debugger strap disconnection and the warning that cutting the straps disables programming, debugging, virtual serial port, and data streaming for the on-board ATtiny3217.
- PyAvrOCD documentation for its role as an AVR GDB server, support for Curiosity Nano on-board nEDBG probes, command-line options such as `--device`, `--interface`, `--tool`, and `--port`, and the `avr-gdb` remote debugging flow through `target remote`.
- The local GDB 17.2 manual for `stepi`, `nexti`, `x/i`, `display/i $pc`, `info registers`, and `hook-stop`.
- The local Red Hat GDB tutorial PDF for practical startup-script behavior, `--nx`, `help`, `apropos`, and basic debug-session setup.

Chapter 4

Instruction Set Basics

The AVR instruction set has roughly 135 instructions, but a much smaller core set does most of the real work in ordinary firmware. This chapter focuses on that core: how instructions are encoded, which register restrictions matter, how arithmetic and logic actually affect the flags, and how to read the CPU's behaviour from the instruction sequence.

The goal is not to memorise a catalog. The goal is to make the instruction set feel predictable.

This chapter uses the **AVRxt** timing column from Microchip's AVR Instruction Set Manual, because the ATtiny3217 is an AVRxt-family device. Older AVR, AVRe, and AVRxm parts are machine-code compatible for most instructions, but some cycle counts differ. When cycle timing matters, check the column for the actual CPU family instead of copying a number from a different AVR generation.

4.1 Instruction Encoding

Every AVR instruction is encoded as one or two 16-bit words (2 or 4 bytes) stored in flash. The CPU fetches one word per clock cycle. Most instructions are a single word and execute in 1-2 cycles.

That encoding detail matters because it affects both code size and speed. On AVR, the shape of an instruction is often the reason for its operand limits. If an instruction only has a few bits available to describe a register or an immediate, the hardware simply cannot support every possible combination.

The instruction set manual reports instruction sizes in **words**, not bytes. One word is 16 bits. So "Words: 1" means 2 bytes in flash, and "Words: 2" means 4 bytes in flash. That same word-oriented view shows up in branch offsets and return addresses.

4.1.1 Single-Word Instructions (16-bit)

The vast majority. The 16 bits encode the opcode, destination register, and source register or small immediate.

LDI Rd, K – Load Immediate

Encoding: 1110 KKKK dddd KKKK

└── opcode ─┘ └── d = Rd - 16 (only r16-r31 → d = 0-15)
K = 8-bit immediate split across bits

ldi r20, 0xAB:

d = 20 - 16 = 4 = 0b0100

```

K = 0xAB = 0b10101011
→ 1110 1010 0100 1011 = 0xEA4B
ADD Rd, Rr – Add without Carry
Encoding: 0000 11rd dddd rrrr
          d = Rd (0–31), r = Rr (0–31)

```

```

add r16, r17:
d = 16 = 0b10000, r = 17 = 0b10001
→ 0000 1101 0000 0001 = 0x0D01

```

4.1.2 Two-Word Instructions (32-bit)

These need a full 16-bit address or program address that doesn't fit in the first word:

Instruction	Why 32-bit
LDS Rd, k	Data address k is 16 bits — needs second word
STS k, Rr	Same
JMP k	Program address k is 22 bits — needs second word
CALL k	Same

```

STS k, Rr – Store to SRAM
Word 1: 1001 001r rrrr 0000 (r = Rr, 5 bits)
Word 2: kkkk kkkk kkkk kkkk (16-bit address k)

```

```

sts PORTB_DIRSET, r16:
PORTB_DIRSET = 0x0421
Word 1: 0x9200 | (16 << 4) = 0x9300
Word 2: 0x0421
→ 4 bytes total in flash

```

This is why STS costs 2 cycles and takes twice the flash of OUT. Always use OUT when the I/O address is $\leq 0x3F$.

This is a recurring AVR theme: the shorter instruction is usually faster, but only if the target address fits the shorter encoding.

The same size difference appears in disassembly. OUT 0x3d, r16 is a single 16-bit word. STS 0x0401, r16 is two words: one word for the opcode/register field and one word for the 16-bit data-space address. The programmer sees one source line either way, but flash sees different byte counts.

4.1.3 Cycle Counts at a Glance

Category	Typical cycles	Notes
Register ops (ADD, AND, MOV...)	1	Pipeline: fetch overlaps
Loads from SRAM (LD, LDS)	2–3	Memory access latency
Stores to SRAM (ST, STS)	2	
IN / OUT	1	Low I/O space only
Taken branch	2	Pipeline flush
Not-taken branch	1	No flush
RCALL	2	16-bit PC devices such as ATtiny3217
RET	4	Stack pop + control-flow change
CALL	3	32-bit instruction, 16-bit PC
JMP	3	32-bit instruction, no stack use
MUL	2	Result in r1:r0
NOP	1	Explicit no-op

The table assumes ordinary internal SRAM/data accesses. Microchip notes that data-memory cycle counts for LD, LDS, ST, and STS assume internal RAM; accesses to non-volatile memory through the data space can take extra time depending on the NVMCTRL implementation. That matters later when reading mapped flash on AVRxt with ordinary LD.

4.2 Instruction Categories

AVR Instruction Categories	
Data Transfer	MOV MOVW LDI LD ST LDS STS LPM PUSH POP IN OUT
Arithmetic	ADD ADC SUB SBC SUBI SBCI ADIW SBIW INC DEC NEG MUL MULS MULSU
Logic	AND ANDI OR ORI EOR COM TST CLR SER
Shift / Rotate	LSL LSR ROL ROR ASR SWAP
Bit Operations	SBI CBI SBIC SBIS SBRC SBRS BST BLD SEC CLC SEI CLI BSET BCLR
Branch / Jump	RJMP JMP IJMP RCALL CALL ICALL RET RETI BREQ BRNE BRCS BRCC BRSH BRLO BRMI BRPL BRGE BRLT BRVS BRVC BRIE BRID CPSE
MCU Control	NOP SLEEP WDR BREAK

Chapters 5–8 cover Load/Store, Arithmetic, Branches, and Subroutines in depth. This chapter covers the core register-to-register operations that form the building blocks of everything else.

So think of this chapter as the grammar of AVR instructions. Later chapters combine these pieces into bigger programming patterns.

4.3 Data Transfer — Moving Values Around

4.3.1 MOV — Copy Register

`MOV Rd, Rr` ; $Rd \leftarrow Rr$

Copies the 8-bit value from Rr to Rd. Source unchanged. No flags affected. Works on any register r0–r31.

MOV is boring in the best possible way. It does exactly what it says, and because it does not touch flags, it is safe to use in the middle of code that depends on the current condition state.

```
ldi r16, 42
mov r0, r16      ; r0 = 42, r16 still = 42
mov r5, r0       ; r5 = 42 (MOV works on r0–r15 unlike LDI)
```

1 cycle. Does not affect SREG.

4.3.2 MOVW — Copy Register Pair

MOVW Rd, Rr ; Rd+1:Rd ← Rr+1:Rr

Copies two registers in one instruction. Rd and Rr must be even-numbered: r0, r2, r4, ... r30.

```
; Copy X pointer (r27:r26) into r1:r0
movw r0, r26 ; r1:r0 = r27:r26

; Copy Z into Y for a secondary pointer
movw r28, r30 ; r29:r28 = r31:r30 (Y = Z)
```

1 cycle. Does not affect SREG. Saves one instruction vs two MOVs.

Whenever you are copying a pointer pair or a 16-bit value, MOVW is worth remembering because it is both shorter and clearer than two separate moves.

Because MOVW uses even register numbers, `movw r24, r22` copies r23:r22 into r25:r24. Think “low register named, high register implied.” This is the same low-byte-first convention used by pointer pairs X (r27:r26), Y (r29:r28), and Z (r31:r30).

4.3.3 LDI — Load Immediate

LDI Rd, K ; Rd ← K (K = 0..255)

Load a constant into a register. **r16-r31 only.**

```
ldi r16, 0xFF ; OK
ldi r20, (1 << 4) ; OK — expression evaluated at assembly time
ldi r0, 42 ; ERROR: r0 not allowed
```

1 cycle. Does not affect SREG.

LDI has an 8-bit immediate (K = 0..255) and a 4-bit encoded destination field. That 4-bit field is why only the upper half of the register file is legal: the encoded value is Rd - 16.

4.3.4 CLR and SER — Clear and Set All Bits

These are **pseudo-instructions** (GAS aliases, not separate opcodes):

```
clr r16 ; r16 = 0x00 — assembled as EOR r16, r16
ser r16 ; r16 = 0xFF — assembled as LDI r16, 0xFF (r16-r31 only)
```

CLR works on any register (r0-r31) because it uses EOR. SER only works on r16-r31 because it uses LDI.

This is a good example of how AVR pseudo-instructions work. They look like real instructions in source, but the assembler is translating them into more basic operations for you.

CLR sets Z=1, clears N, V, S. SER affects no flags.

Pseudo-instructions are useful, but do not forget their real flag behavior. CLR is not flag-neutral because it is really EOR Rd,Rd. SER is flag-neutral because it is really LDI Rd,0xFF.

4.4 Arithmetic Instructions

4.4.1 ADD — Add Without Carry

ADD Rd, Rr ; Rd ← Rd + Rr

Adds two registers, stores result in Rd. Affects: H, S, V, N, Z, C.

If you are coming from a higher-level language, read that carefully: the result replaces the original destination register. AVR arithmetic is almost always destructive to the destination operand.

```
ldi r16, 100
ldi r17, 50
add r16, r17          ; r16 = 150, C=0, Z=0, N=0
```

Unsigned overflow (result > 255): C=1, result wraps.

```
ldi r16, 200
ldi r17, 100
add r16, r17          ; 200+100=300 → r16=44, C=1
```

For signed arithmetic, the carry flag is not the overflow test. Signed overflow is reported in V, and signed branches use S (N XOR V). For example, $0x7F + 0x01 = 0x80$ sets V because +127 became -128 in 8-bit two's complement, even though that operation is also just an ordinary 8-bit wraparound.

4.4.2 ADC — Add with Carry

ADC Rd, Rr ; $Rd \leftarrow Rd + Rr + C$

Identical to ADD but adds the carry flag too. Essential for **multi-byte arithmetic** — carry from the low byte automatically propagates into the high byte:

```
; 16-bit add: r17:r16 += r19:r18
add r16, r18          ; low bytes first – may set C
adc r17, r19          ; high bytes – C from previous ADD is included

; 24-bit add: r2:r1:r0 += r5:r4:r3
add r0, r3
adc r1, r4
adc r2, r5            ; each ADC picks up carry from the byte below
```

This low-byte-first pattern is one of the most important habits in AVR assembly. Once it becomes automatic, larger integer arithmetic stops feeling special.

Microchip's manual defines ADC as $Rd \leftarrow Rd + Rr + C$; the carry input is the current C bit before the instruction executes, and the output C bit is then ready for the next byte. That is why no separate "carry variable" exists in multi-byte arithmetic.

4.4.3 SUB — Subtract Without Carry

SUB Rd, Rr ; $Rd \leftarrow Rd - Rr$

Subtracts Rr from Rd. C=1 means borrow (unsigned underflow).

That "C means borrow" convention is easy to trip over because it feels backwards if you learned carry only from addition. On AVR subtraction, carry set means the subtraction had to borrow.

```
ldi r16, 10
ldi r17, 20
sub r16, r17          ; 10-20 = -10 → r16=246 (0xF6), C=1 (borrow)
```

The same instruction can be read three ways:

Interpretation	Example result for 10 - 20
Raw 8-bit bits	0xF6
Unsigned value	246, with C=1 showing borrow

Interpretation	Example result for 10 - 20
Signed value	-10, with N=1 and S depending on V

The CPU only produces bits and flags. Your later branch instruction decides which interpretation matters.

4.4.4 SBC — Subtract with Carry (Borrow)

SBC Rd, Rr ; $Rd \leftarrow Rd - Rr - C$

Like ADC for addition, SBC propagates borrow in multi-byte subtraction:

```
; 16-bit subtract: r17:r16 -= r19:r18
sub r16, r18 ; low bytes
sbc r17, r19 ; high bytes — includes borrow from low byte
```

Again, the pattern is the same as multi-byte addition: low byte first, then work upward using the carry/borrow chain.

4.4.5 SUBI — Subtract Immediate

SUBI Rd, K ; $Rd \leftarrow Rd - K$ (r16–r31 only)

Subtract a constant without needing a second register. **No ADDI exists** — use SUBI with the negated value, or ADIW/SBIW for 16-bit register pairs.

That missing ADDI surprises many beginners. It is not an omission in the book; it is a real feature of the AVR ISA.

```
ldi r16, 50
subi r16, 7 ; r16 = 43
subi r16, -3 ; r16 = 46 (subtract -3 = add 3)
```

SUBI r16, -3 works because the assembler encodes -3 as the 8-bit constant 0xFD. Subtracting 0xFD from an 8-bit register is the same low-byte result as adding 3. This idiom is common, but use it deliberately; it can be less readable than ADIW when you really mean “add to a 16-bit pointer.”

4.4.6 SBCI — Subtract Immediate with Carry

SBCI Rd, K ; $Rd \leftarrow Rd - K - C$ (r16–r31 only)

Pair with SUBI for 16-bit immediate subtraction:

```
; r17:r16 -= 300 (0x012C)
subi r16, lo8(300) ; r16 -= 0x2C
sbci r17, hi8(300) ; r17 -= 0x01 plus any borrow from low byte
```

4.4.7 ADIW — Add Immediate to Word

ADIW Rd, K ; $Rd+1:Rd \leftarrow Rd+1:Rd + K$ ($K = 0..63$)

Adds a small constant to a 16-bit register pair in one instruction. Only works on r25:r24, r27:r26 (X), r29:r28 (Y), r31:r30 (Z).

```
adiw r24, 1 ; r25:r24 += 1
adiw XL, 4 ; X pointer += 4 (advance 4 bytes)
adiw ZL, 16 ; Z pointer += 16
```

2 cycles. Affects C, Z, N, V, S.

ADIW is especially nice for pointer arithmetic, because X, Y, and Z all live in the supported register-pair set.

The immediate is only 6 bits, so K is limited to 0..63. Larger pointer steps need repeated ADIW, register arithmetic, or explicit low/high-byte addition.

4.4.8 SBIW — Subtract Immediate from Word

SBIW Rd, K ; Rd+1:Rd ← Rd+1:Rd - K (K = 0..63)

Same register restrictions as ADIW. Used in delay loops (Chapter 1) and pointer adjustment.

```
sbiw r24, 1 ; r25:r24 -= 1 - sets Z when result = 0
sbiw ZL, 8 ; Z -= 8
```

4.4.9 INC and DEC — Increment / Decrement

INC Rd ; Rd ← Rd + 1

DEC Rd ; Rd ← Rd - 1

Works on r0-r31. **Critical gotcha: INC and DEC do not affect the carry flag (C).** This catches people trying to use DEC as part of a multi-byte borrow chain.

```
; This loop is WRONG for multi-byte work:
ldi r16, 0
ldi r17, 0
dec r16 ; r16 = 255, C unchanged - cannot chain with SBC!

; Correct 16-bit decrement on a supported pair: use SBIW
sbiw r24, 1 ; r25:r24 -= 1, sets Z correctly
```

Flags affected by INC/DEC: S, V, N, Z — but NOT C.

That is why INC and DEC are great for single-byte counters and a bad fit for arithmetic that needs carry propagation.

Microchip calls out another subtlety: for unsigned values after INC or DEC, only BREQ and BRNE are generally reliable tests. Use them for “did the counter wrap/reach zero?” tests. Do not use BRL0/BRSH after INC or DEC and expect a meaningful unsigned less-than/greater-than result, because C was not updated.

4.4.10 NEG — Two’s Complement Negate

NEG Rd ; Rd ← 0x00 - Rd

Negates the register (flip all bits, then add 1). NEG 0 → 0 with C=0. All other values: C=1.

```
ldi r16, 10
neg r16 ; r16 = 246 (0xF6), which is -10 in two's complement
neg r16 ; r16 = 10 again
```

Useful for: converting subtraction direction, computing absolute value.

Whenever you see NEG, think “arithmetic sign change.” Whenever you see COM, think “bit-wise inversion.” They are related, but not interchangeable.

There is one unavoidable two’s-complement corner case: NEG 0x80 leaves the bit pattern 0x80. In signed 8-bit arithmetic, -128 has no positive counterpart representable in 8 bits. The V flag reports that overflow case.

4.5 Logic Instructions

4.5.1 AND / ANDI — Bitwise AND

```
AND  Rd, Rr          ; Rd ← Rd AND Rr          (r0–r31)
ANDI Rd, K            ; Rd ← Rd AND K          (r16–r31 only)
```

Use AND to **clear** (mask off) specific bits:

```
ldi r16, 0xAB        ; 0b10101011
andi r16, 0x0F        ; r16 = 0x0B (0b00001011) – upper nibble cleared
andi r16, 0b11111110 ; clear bit 0
```

Affects: S, V(=0), N, Z. Clears V.

Bit masking is one of the most common operations in embedded work. In practice: AND clears bits, OR sets bits, and EOR toggles bits.

ANDI is restricted to r16–r31, but AND works on all 32 registers. If your value is in r0–r15, either use a mask register with AND, or move the value to an upper register before using ANDI.

4.5.2 OR / ORI — Bitwise OR

```
OR   Rd, Rr          ; Rd ← Rd OR Rr
ORI  Rd, K            ; Rd ← Rd OR K          (r16–r31 only)
```

Use OR to **set** specific bits:

```
ldi r16, 0x2A
ori  r16, 0x80        ; set bit 7 → r16 = 0xAA
ori  r16, (1 << 3)    ; set bit 3
```

ORI has the same upper-register restriction as LDI and ANDI. The opcode has room for an 8-bit immediate and only the reduced destination encoding.

4.5.3 EOR — Bitwise XOR (Exclusive OR)

```
EOR Rd, Rr          ; Rd ← Rd XOR Rr
```

Use EOR to **toggle** bits or **invert** a value:

```
ldi r16, 0xFF
ldi r17, 0x0F
eor r16, r17        ; r16 = 0xF0 – lower nibble toggled
eor r16, r16        ; r16 = 0x00 – self-XOR always gives 0 (this is CLR)
```

4.5.4 COM — One's Complement (Bitwise NOT)

```
COM Rd          ; Rd ← 0xFF - Rd (flip every bit)
```

```
ldi r16, 0b10101010
com r16          ; r16 = 0b01010101
```

Always sets C=1. Note: COM ≠ NEG. COM flips bits; NEG negates (COM + 1).

Because COM always sets C, avoid placing it between the low and high bytes of an ADD/ADC or SUB/SBC carry chain unless changing C is exactly what you intend.

4.5.5 TST — Test Register (Non-Destructive)

```
TST Rd          ; Rd AND Rd → set N, Z flags; Rd unchanged
```

Tests whether a register is zero or negative without modifying it:

```
lds r16, some_var
tst r16
breq var_is_zero      ; Z=1 if r16 was 0
brmi var_is_neg       ; N=1 if bit 7 was set (negative signed)
```

Assembled as AND Rd, Rd.

TST is one of those instructions that mostly exists to make your intent more obvious to a human reader.

Since TST is really AND Rd,Rd, it clears V and sets S from N XOR V. After `tst r16`, BRMI tests bit 7 directly because V is zero.

4.6 Comparison Instructions

These set flags **without storing any result**. Always followed by a conditional branch.

This is the AVR version of “ask a question, then branch on the answer.” The question is encoded in the compare; the answer appears in the flags.

4.6.1 CP — Compare

CP Rd, Rr ; Rd - Rr → flags only, result discarded

Exactly like SUB but throws away the result. Use before BREQ, BRNE, BRLO, BRGE, etc.

That makes CP conceptually simple: it performs a subtraction only for the purpose of updating flags.

```
ldi r16, 42
ldi r17, 100
cp r16, r17      ; flags set as if 42 - 100
brlo r16_smaller ; BRLO branches if C=1 (unsigned: Rd < Rr)
```

4.6.2 CPC — Compare with Carry

CPC Rd, Rr ; Rd - Rr - C → flags only

For multi-byte comparison (after comparing low bytes with CP):

```
; Compare r17:r16 with r19:r18 (16-bit unsigned)
cp r16, r18      ; compare low bytes
cpc r17, r19      ; compare high bytes including borrow
brlo less_than    ; taken if r17:r16 < r19:r18 unsigned
```

4.6.3 CPI — Compare with Immediate

CPI Rd, K ; Rd - K → flags only (r16–r31 only)

```
ldi r16, 10
cpi r16, 10
breq exactly_ten ; taken
```

4.6.4 CPSE — Compare, Skip if Equal

CPSE Rd, Rr ; if Rd == Rr: skip next instruction

Skips exactly one instruction when equal. Used for compact conditional code:

```
cpse r16, r17
rjmp not_equal      ; skipped if r16 == r17
; falls through here if equal
```

CPSE does **not** update SREG. It simply compares and adjusts the program counter if the registers are equal. It also has variable timing: 1 cycle if no skip happens, 2 cycles if it skips a one-word instruction, and 3 cycles if it skips a two-word instruction.

That makes CPSE compact, but not always ideal for cycle-stable code:

```
cpse r16, r17
sts PORTA_OUTTGL, r18    ; two-word instruction: skip path is longer
```

If the next instruction might be two words (LDS, STS, JMP, CALL), count the skip path carefully.

4.7 Conditional Branches — Full Table

Every branch instruction tests one or more SREG flags. The mnemonic tells you what condition causes the branch.

Once you stop reading branch mnemonics as magic words and start reading them as “test this flag pattern,” control flow becomes much easier to reason about.

Mnemonic	Full name	Condition	Flag(s)
BREQ	Branch if Equal	Z = 1	Z
BRNE	Branch if Not Equal	Z = 0	Z
BRCS	Branch if Carry Set	C = 1	C
BRCC	Branch if Carry Cleared	C = 0	C
BRSH	Branch if Same or Higher	C = 0	C ← unsigned ≥
BRLO	Branch if Lower	C = 1	C ← unsigned <
BRMI	Branch if Minus	N = 1	N
BRPL	Branch if Plus	N = 0	N
BRGE	Branch if ≥ (signed)	S = 0	N⊕V ← signed ≥
BRLT	Branch if < (signed)	S = 1	N⊕V ← signed <
BRVS	Branch if Overflow Set	V = 1	V
BRVC	Branch if Overflow Clear	V = 0	V
BRIE	Branch if Interrupts En.	I = 1	I
BRID	Branch if Interrupts Dis.	I = 0	I
BRHS	Branch if Half-carry Set	H = 1	H
BRHC	Branch if Half-carry Clr	H = 0	H
BRTS	Branch if T Set	T = 1	T
BRTC	Branch if T Clear	T = 0	T

All branch instructions: **-64 to +63 words** offset relative to the next instruction. In address terms, the destination must be within PC - 63 through PC + 64 after the branch instruction’s normal PC + 1 step. If the target is farther, use JMP or RJMP with an inverted condition:


```

; Branch to far_label if r16 == r17 – but far_label > 63 words away:
cp   r16, r17
brne skip           ; if NOT equal, skip the jump
jmp  far_label
skip:

```

This branch-range limit is another place where instruction encoding leaks into everyday programming. The branch is short because the displacement field is short.

Conditional branches do not modify SREG. They only read the flags left by an earlier instruction. Any instruction inserted between the compare and the branch must be checked for flag effects:

```

cp   r16, r17
mov  r18, r16      ; OK: MOV does not affect flags
breq equal

cp   r16, r17
andi r18, 0x0F     ; BAD if you meant to branch on CP: ANDI changes Z/N/V/S
breq equal

```

4.8 Unsigned vs Signed Comparisons

This trips up almost every beginner. Unsigned and signed comparisons use different branch instructions after the same CP.

```

ldi r16, 0xF0      ; = 240 unsigned, or -16 signed
ldi r17, 0x10      ; = 16 unsigned, or 16 signed

cp   r16, r17      ; 0xF0 - 0x10 = 0xE0. C=0 (no borrow), N=1, V=0, S=1

brlo below_u       ; C=1? NO → not taken. Correct: 240 > 16 unsigned.
brlt below_s       ; S=1? YES → taken. Correct: -16 < 16 signed.

```

Rule: **BRLO/BRSH for unsigned. BRLT/BRGE for signed.**

If you take only one rule from this section, take that one. The CPU does not know whether you intended a byte to mean 240 or -16. Your branch choice tells it how to interpret the flags.

Microchip’s branch summary says the common tests ($R_d < R_r$, $R_d \geq R_r$, and so on) are meaningful immediately after compare/subtract-style instructions such as CP, CPI, SUB, and SUBI. If the preceding instruction was a logical operation, BRL0 still means “C=1”, but it no longer means “less than” unless that flag pattern was intentionally prepared.

4.9 Flag Effects Summary

The table below covers every instruction in this chapter. “—” means the flag is not affected.

Instruction	C	Z	N	V	S	H
ADD	●	●	●	●	●	●
ADC	●	●	●	●	●	●
SUB	●	●	●	●	●	●
SBC	●	●	●	●	●	●
SUBI	●	●	●	●	●	●

SBCI	●	●	●	●	●	●	
ADIW	●	●	●	●	●	—	
SBIW	●	●	●	●	●	—	
INC	—	●	●	●	●	—	← C NOT affected
DEC	—	●	●	●	●	—	← C NOT affected
NEG	●	●	●	●	●	●	
AND / ANDI	—	●	●	0	●	—	← V cleared, C unchanged
OR / ORI	—	●	●	0	●	—	← V cleared, C unchanged
EOR	—	●	●	0	●	—	← V cleared, C unchanged
COM	1	●	●	0	●	—	← C always 1, V cleared
TST	—	●	●	0	●	—	← same flags as AND Rd,Rd
CP / CPI	●	●	●	●	●	●	
CPC	●	●	●	●	●	●	
MOV/MOVW/LDI	—	—	—	—	—	—	
CLR	—	1	0	0	0	—	← same flags as EOR Rd,Rd

The full SREG also contains I, T, H, S, V, N, Z, and C. This table focuses on the arithmetic/logic flags used in this chapter. SEI, CLI, BSET, BCLR, BST, and BLD manipulate the other status bits and are handled more deeply in the interrupt and bit-manipulation chapters.

4.10 Worked Example: Compare, Clamp, Negate

Goal: read an 8-bit value, clamp it to 0-100 range, then negate it if a flag register bit is set.

This example matters because it combines several ideas you will use constantly: compare, branch, clamp by assignment, test a control bit, and optionally apply an arithmetic transform.

```

/*
 * clamp_and_negate:
 *   Input:  r24 = raw value (unsigned), r25 = flags (bit 0 = negate)
 *   Output: r24 = clamped value, optionally negated
 *   Clobbers: r16
 */
clamp_and_negate:
    /* Step 1: clamp to 100 max */
    cpi    r24, 100          ; compare r24 with 100
    brlo   clamped           ; if r24 < 100 (unsigned), skip clamp
    ldi    r24, 100          ; else clamp to 100
clamped:

    /* Step 2: clamp to 0 min – already unsigned so r24 >= 0 always.
     * If the raw value were signed we'd add BRGE check here. */

    /* Step 3: negate if bit 0 of r25 is set */
    sbrs   r25, 0            ; Skip if Bit in Register Set (bit 0 of r25)
    ret                                ; bit 0 clear – return as-is
    neg    r24                ; negate: r24 = 0 - r24
    ret

```

New instructions introduced:

SBRS Rr, b Skip next instruction if bit b of Rr is set. 1 cycle (not skipped), 2 cycles (skipped over 1-word instr), 3 cycles (skipped over 2-word instr). Does not affect SREG.

The skip family all follows the same timing idea on AVRxt:

Instruction	Test location	Skips when
SBRC Rr,b	Register bit	bit is clear
SBRs Rr,b	Register bit	bit is set
SBIC A,b	Low I/O bit 0x00-0x1F	bit is clear
SBIS A,b	Low I/O bit 0x00-0x1F	bit is set
CPSE Rd,Rr	Register compare	registers equal

All skip exactly one following instruction, not a block. If you need to skip multiple operations, skip over a branch:

```
sbrs r25, 0
rjmp no_negate
neg r24
no_negate:
```

Trace through the example:

Call: r24=150, r25=0b00000001 (negate flag set)

```
cpi r24, 100    → 150-100=50, C=0, Z=0
brlo clamped    → C=0: NOT taken (150 >= 100, no branch)
ldi r24, 100    → r24 = 100
clamped:
sbrs r25, 0     → bit 0 of r25 is 1: SKIP next instr
                  (ret is skipped)
neg r24         → r24 = 0 - 100 = 156 (0x9C = -100 signed)
ret
```

Walking through examples like this is how you build intuition. Do not just read the final code; read the flags and the decision points as they evolve.

4.11 Operand Restrictions Cheat-Sheet

This is worth turning into instinct:

Instruction	Rd range	Rr range	Immediate range
MOV	r0-r31	r0-r31	—
MOVW	r0,r2..r30	r0,r2..r30	— (even only)
LDI	r16-r31	—	0-255
CLR	r0-r31	—	—
SER	r16-r31	—	—
ADD/ADC	r0-r31	r0-r31	—
SUB/SBC	r0-r31	r0-r31	—
SUBI/SBCI	r16-r31	—	0-255
ADIW/SBIW	r24,X,Y,Z	—	0-63
INC/DEC	r0-r31	—	—
NEG/COM	r0-r31	—	—
AND/OR/EOR	r0-r31	r0-r31	—
ANDI/ORI	r16-r31	—	0-255
CP	r0-r31	r0-r31	—
CPI	r16-r31	—	0-255
CPC	r0-r31	r0-r31	—
CPSE	r0-r31	r0-r31	—
SBRs/SBRC	r0-r31	—	bit 0-7

4.12 Summary

Instructions: 1 or 2 words (16 or 32 bits).

Two-word: LDS/STS carry a 16-bit data address; JMP/CALL carry a larger program address field.

AVRxt timing in this chapter: LDS=3 cycles, STS=2 cycles, LD/ST=2 cycles for internal SRAM; mapped NVM can add cycles.

Register rules:

LDI, SUBI, SBCI, ANDI, ORI, SER, CPI → r16–r31 only.

ADIW, SBIW → r24, X(r26), Y(r28), Z(r30) only.

Everything else → r0–r31.

Carry chains: ADD+ADC, SUB+SBC, SUBI+SBCI – always low byte first.

INC/DEC do NOT touch C – do not use them to propagate multi-byte carry/borrow.

NEG ≠ COM: NEG = two's complement, COM = bitwise invert.

Comparisons:

Unsigned: BRL0 (<), BRSH (≥), BRCS/BRCC

Signed: BRLT (<), BRGE (≥)

Sign bit: BRMI (N=1), BRPL (N=0)

Branch range: -64..+63 word offset. Use JMP/RJMP patterns for farther targets.

Skip instructions skip one instruction only and cost 1/2/3 cycles depending on whether no skip, one-word skip, or two-word skip occurs.

4.13 Exercises

1. Add two 24-bit numbers stored in r2:r1:r0 and r5:r4:r3. Store the result in r2:r1:r0. What happens to the carry if the result exceeds 24 bits?
2. Write a sequence to compute $r16 = r16 * 4$ without using MUL. (Hint: shifting left by 1 is the same as multiplying by 2.) Use only ADD instructions.
3. Write a loop that counts from 0 to 255 in r16, then wraps and repeats forever. Use INC and BREQ. Why is BREQ the right branch here?
4. Given $r16 = 0xAB$ and $r17 = 0x3C$, predict the result and all six flag states (C, Z, N, V, S, H) after each operation — before assembling. For instructions that do not affect C or H, write “unchanged” for that flag:
 - ADD r16, r17
 - SUB r16, r17 (fresh $r16=0xAB$ each time)
 - AND r16, r17
 - EOR r16, r17
5. Write `clamp_and_negate` to also handle the case where r24 starts as 0 and the negate flag is set. What does NEG 0 produce and is that correct for your use case?
6. Why does SUBI r16, -1 add 1 to r16? What is the bit pattern of -1 as an 8-bit two's complement number, and what does subtracting it do?

Next: Chapter 5 — Load and Store

Chapter 5


Load and Store

Every useful program moves data. In AVR assembly, **every data movement is explicit**: you choose the exact instruction that moves each byte, the exact addressing mode, and often the exact pointer update behaviour. This chapter covers the main ways to get data into registers, get it back out to memory, and read constants from flash on the ATtiny3217.

5.1 The Memory-Register Model

The AVR CPU can only operate on data that is in **registers** (r0-r31). To work on a byte in SRAM, you must load it into a register first, operate, then store it back. There is no “operate directly on memory” instruction.

That rule is so central that it is worth repeating in plain language: if the value is not in a register yet, the ALU cannot touch it.



Every instruction in this chapter is either a load (memory → register) or a store (register → memory).

Once you see that pattern, a lot of AVR code becomes easier to read. A routine is often just alternating phases: bring bytes into registers, do work there, write results back out.

5.1.1 ATtiny3217 Data Space Is Not Just SRAM

The load/store instructions use the AVR **data address space**, not just the 2 KB internal SRAM block. On ATtiny3217, Microchip’s memory map puts several different resources into that 16-bit address space:

Data address range	Resource
0x0000-0x003F	64 I/O registers

Data address range	Resource
0x0040-0x0FFF	Extended I/O/peripheral range; Microchip documents 960 implemented Ext. I/O registers plus reserved space
0x1000-0x13FF	NVM I/O registers and data
0x1400-0x14FF	EEPROM, 256 bytes
0x3800-0x3FFF	Internal SRAM, 2 KB
0x8000-0xFFFF	Flash mapped into data space, 32 KB

That is why the same LD instruction can read SRAM, an I/O register, EEPROM space, or mapped flash. The address decides which hardware block answers.

For normal variables, the assembler and linker place `.bss` and `.data` in SRAM at 0x3800-0x3FFF. Do not treat EEPROM or flash as scratch SRAM: they are nonvolatile memory regions controlled by NVMCTRL and have different write rules and timing.

5.2 Addressing Modes

The AVR supports five ways to specify a memory address:

Mode	Syntax	Address source	Instruction
Direct	LDS Rd, k	16-bit constant k in instruction	LDS / STS
Indirect	LD Rd, X	Value in X, Y, or Z pointer pair	LD / ST
Post-increment	LD Rd, X+	X (used), then X += 1	LD / ST
Pre-decrement	LD Rd, -X	X -= 1, then X (used)	LD / ST
Indirect + displacement	LDD Rd, Y+q	Y + constant q (0-63)	LDD / STD

These are not five completely different ideas. They are variations on one theme: where does the address come from, and should the pointer change as a side effect?

5.3 Direct Addressing: LDS and STS

```
LDS Rd, k          ; Rd ← SRAM[k]      (k = 16-bit address)
STS k, Rr          ; SRAM[k] ← Rr
```

The address is encoded directly in the instruction as a second 16-bit word. Both are **32-bit (two-word) instructions**.

That makes LDS/STS very convenient to write, but relatively expensive in both flash space and cycles.

```
.section .bss
my_var: .skip 1           ; reserve one byte in SRAM

.section .text
ldi r16, 0xAB
sts my_var, r16          ; write 0xAB to my_var's SRAM address
lds r17, my_var          ; r17 = 0xAB
```

Cycle costs:

Instruction	Cycles
LDS Rd, k	3
STS k, Rr	2

These AVRxt cycle counts assume internal SRAM or ordinary data-memory access. Microchip's instruction manual notes that data-memory accesses to NVM need at least one extra cycle, with the exact timing depending on the device's NVMCTRL implementation. Treat EEPROM and mapped flash as slower than SRAM unless the device data sheet gives a tighter bound for the specific access.

When to use LDS/STS: accessing a single named variable. For loops over arrays, indirect addressing (Section 5.4) is faster and smaller.

So LDS/STS are great for “go touch this exact place once.” They are a poor fit for repeated traversal through memory.

5.3.1 I/O Registers Above 0x3F

IN/OUT cannot reach peripheral registers with addresses above 0x3F. Use data-space instructions (LD/LDS/LDD and ST/STS/STD) for those addresses. For a fixed named peripheral register, LDS/STS are the usual direct form:

```
ldi r16, (1 << 4)
sts PORTB_DIRSET, r16      ; PORTB_DIRSET = 0x0421, above 0x3F → STS
```

For addresses 0x00–0x3F, always prefer IN/OUT (1 cycle, 1 word) over LDS/STS (3/2 cycles, 2 words).

That is one of the easiest AVR optimisations to make because it is mechanical: if the register is in low I/O space, use the I/O instructions.

There is one more I/O-address boundary to remember:

- IN and OUT can address I/O locations 0x00–0x3F.
- SBI, CBI, SBIC, and SBIS can address only 0x00–0x1F.

ATtiny3217 exposes virtual port registers (VPORTx) in the low I/O range so simple GPIO operations can use IN/OUT. The full PORTx peripheral registers live in extended I/O space, so examples such as PORTB_DIRSET use STS.

5.4 Pointer Registers: X, Y, Z

The three pointer register pairs are 16-bit registers built from adjacent general-purpose registers:

X = r27 : r26	XH : XL
Y = r29 : r28	YH : YL
Z = r31 : r30	ZH : ZL

To load an address into a pointer pair, use two LDI instructions:

```
; Point X at buf (a label in .bss)
ldi XL, lo8(buf)      ; XL = r26 = low byte of buf's address
ldi XH, hi8(buf)      ; XH = r27 = high byte
```

Now X holds the 16-bit address of buf and can be used with LD/ST.

This is the beginning of pointer-based programming in AVR. Instead of naming an address directly in every instruction, you put the address in X, Y, or Z once and then operate through that pointer.

Practical convention:

- Use X for simple source or destination pointers.
- Use Y when you want a stable base pointer for LDD/STD, especially for structs or stack-frame-style code.
- Use Z for flash table reads, because LPM requires Z and mapped-flash examples commonly use Z too.

The ATtiny3217 data space fits in 64 KB, so X, Y, and Z are enough. Larger AVR devices can add RAMPX, RAMPY, RAMPZ, or RAMPD to extend addressing beyond 64 KB, but that extra machinery is not part of normal ATtiny3217 code.

5.5 Indirect Addressing: LD and ST

5.5.1 Basic Indirect — No Change to Pointer

```
LD  Rd, X      ; Rd ← SRAM[X]      (X unchanged)
LD  Rd, Y      ; Rd ← SRAM[Y]
LD  Rd, Z      ; Rd ← SRAM[Z]
ST  X, Rr      ; SRAM[X] ← Rr      (X unchanged)
ST  Y, Rr
ST  Z, Rr
```

2 cycles each. Use when you want to read/write the same address multiple times without moving the pointer.

This is the pointer equivalent of “stay where you are and keep working here.”

```
; Read, modify, write a single SRAM byte
ldi XL, lo8(my_var)
ldi XH, hi8(my_var)
ld  r16, X      ; r16 = my_var
ori r16, 0x80   ; set bit 7
st  X, r16      ; write back – X still points to my_var
```

5.5.2 Post-Increment — Walk Forward Through Memory

```
LD  Rd, X+      ; Rd ← SRAM[X], then X += 1
ST  X+, Rr      ; SRAM[X] ← Rr, then X += 1
```

The increment happens **after** the access. The full 16-bit pointer is incremented, so 0x38FF correctly becomes 0x3900. X is ready to point to the next byte. Perfect for forward array traversal:

This is one of the most natural addressing modes on AVR. It lets the memory access and the pointer update happen in the same instruction.

```
; Zero-fill buf (8 bytes)
ldi XL, lo8(buf)
```



```
ldi XH, hi8(buf)
ldi r16, 0
ldi r18, 8
1:
    st X+, r16      ; write 0, advance X
    dec r18
    brne 1b
```

5.5.3 Pre-Decrement — Walk Backward Through Memory

```
LD Rd, -X          ; X -= 1, then Rd ← SRAM[X]
ST -X, Rr          ; X -= 1, then SRAM[X] ← Rr
```

The decrement happens **before** the access. The full 16-bit pointer is decremented before the memory cycle. To walk an array backwards, point X one past the last element:

Pre-decrement often feels odd the first time you see it, but it matches how stacks and reverse traversals naturally work.

```
; Sum array backwards into r17
ldi XL, lo8(buf + 8) ; one past end
ldi XH, hi8(buf + 8)
clr r17
ldi r18, 8
1:
    ld r16, -X      ; X--, then load
    add r17, r16
    dec r18
    brne 1b
```

5.5.4 All Three Pointers, All Three Modes

X, Y, and Z all support LD/ST, LD+/ST+, and LD-/ST-. The only difference is which register pair holds the address:

So the real question is rarely “which instruction family do I need?” More often it is “which pointer register should own this job?”

```
ld r16, X+      ; load via X, post-increment
ld r16, Y+      ; load via Y, post-increment
ld r16, Z+      ; load via Z, post-increment — also used for LPM
```

Avoid using a pointer register as both the pointer and the source or destination register in auto-increment/decrement forms. The AVR instruction manual marks Z auto-update combinations such as LD r30, Z+, LD r31, Z+, LD r30, -Z, LD r31, -Z, ST Z+, r30, ST Z+, r31, ST -Z, r30, and ST -Z, r31 as undefined because r30 and r31 are part of Z itself. The same rule applies to LPM r30, Z+ and LPM r31, Z+.

5.6 Displacement Addressing: LDD and STD

Y and Z support a **constant displacement** added to the pointer without modifying it. X does **not** support displacement.

```
LDD Rd, Y+q      ; Rd ← SRAM[Y + q]   (Y unchanged, q = 0–63)
LDD Rd, Z+q      ; Rd ← SRAM[Z + q]
STD Y+q, Rr      ; SRAM[Y + q] ← Rr
STD Z+q, Rr
```

2 cycles. Displacement *q* must be a constant 0–63 — it cannot be a register.

Displacement mode is useful because it lets one pointer stay parked at a base address while you reach nearby fields by constant offset.

The 0–63 displacement limit is encoded directly in the 16-bit instruction. If a field is at offset 64 or later, adjust the base pointer first with ADIW or use normal pointer arithmetic, then access through *Y* or *Z*.

Primary use: struct member access. Point *Y* at the base of a struct, then access each member by offset:

```
; Struct layout in SRAM:
; offset 0: status (1 byte)
; offset 1: count (1 byte)
; offset 2: value (2 bytes, little-endian)
; offset 4: flags (1 byte)

.equ STATUS_OFF, 0
.equ COUNT_OFF, 1
.equ VALUE_OFF, 2
.equ FLAGS_OFF, 4

; Point Y at the struct
ldi YL, lo8(my_struct)
ldi YH, hi8(my_struct)

; Read all fields without moving Y
ldd r16, Y+STATUS_OFF      ; r16 = status
ldd r17, Y+COUNT_OFF     ; r17 = count
ldd r18, Y+VALUE_OFF      ; r18 = value low byte
ldd r19, Y+VALUE_OFF+1    ; r19 = value high byte
ldd r20, Y+FLAGS_OFF      ; r20 = flags
```

Secondary use: stack frame pointer. The compiler uses *Y* as the frame pointer, accessing local variables via LDD *r*, *Y*+offset. Chapter 8 covers this in the context of subroutines.

In other words, LDD/STD are the AVR equivalent of “base pointer + fixed offset” addressing.

5.7 Cycle and Size Comparison

Mode	Instruction	Cycles	Words	Pointer moves?
Direct	LDS / STS	3 / 2	2	No
Indirect	LD / ST	2	1	No
Post-increment	LD+ / ST+	2	1	+1 after
Pre-decrement	LD- / ST-	2	1	-1 before
Displacement	LDD / STD	2	1	No (<i>Y/Z</i> only)
I/O register	IN / OUT	1	1	N/A (addr 0–3F)

Cycle counts in this table are the AVRxt base costs for internal SRAM and low I/O style accesses. Accesses that hit NVM regions, including EEPROM and mapped flash, can take additional cycles.

Rule of thumb: - **Named variable, one access:** LDS/STS - **Named variable, repeated access:** load address into pointer, then LD/ST - **Array traversal:** LD+/ST+ or LD-/ST- - **Struct fields:** LDD/STD with *Y* as base - **I/O registers ≤ 0x3F:** IN/OUT

Those rules are worth internalising because they cover most real choices in ordinary firmware.

5.8 Reading Flash: LPM

On all AVR devices, LPM reads one byte from the **program bus** (flash) into a register. Z holds the byte address of the flash location to read.

This is the classic AVR way to read constants stored in flash. If you learned AVR on older parts, LPM is probably the method you saw first.

```
LPM Rd, Z          ; Rd ← Flash[Z]          (Z unchanged)
LPM Rd, Z+         ; Rd ← Flash[Z], Z++ (post-increment)
LPM                ; r0 ← Flash[Z]          (deprecated form)
```

3 cycles. Z must be the byte address (not word address).

That byte-address detail matters because some other AVR concepts, like jump targets, use word-oriented program addressing instead.

On ATtiny3217, flash is 32 KB, so ordinary LPM is enough for the whole program memory. ELPM and RAMPZ matter on larger AVR devices with more than 64 KB of program memory, but they do not buy anything on this MCU.

For hand-written flash tables intended for LPM, put the bytes in program-memory-addressed storage such as `.progmem.data`, or in `.text` after executable code. Do not blindly use a `.rodata` symbol as an LPM address with modern `avr-ld` scripts: `.rodata` may have a mapped data-space VMA (`0x8000` plus the flash load address). Check `avr-nm` or `avr-objdump -h` when in doubt.

```
.section .progmem.data, "a", @progbits
lookup:
    .byte 0, 1, 1, 2, 3, 5, 8, 13    ; Fibonacci

.section .text

    ; Load lookup[3] into r16
    ldi ZL, lo8(lookup + 3)
    ldi ZH, hi8(lookup + 3)
    lpm r16, Z                      ; r16 = 2 (lookup[3])
```

Copying a flash table to SRAM with LPM:

```
ldi ZL, lo8(lookup)
ldi ZH, hi8(lookup)
ldi XL, lo8(dest_buf)
ldi XH, hi8(dest_buf)
ldi r18, 8

copy:
    lpm r16, Z+                    ; load from flash, Z++
    st X+, r16                     ; store to SRAM, X++
    dec r18
    brne copy
```

5.9 Reading Flash: Mapped Access (AVRxt Only)

On AVRxt (ATtiny3217), flash is mapped into the **data address space** at offset `MAPPED_PROGMEM_START = 0x8000`. You can read it with ordinary LD instructions — no LPM needed.

This is one of the nicest quality-of-life improvements on the ATtiny3217. It lets flash reads look much more like normal memory reads.

```
LD Rd, Z          ; Rd ← DataMem[Z]
```

If $Z \geq 0x8000$, the data bus reads from flash. If $Z < 0x8000$, it reads SRAM or I/O registers. The CPU doesn't care which — the address selects the bus.

That means one instruction family, LD, can read from very different physical resources depending on the address you place in the pointer.

The ATtiny3217 data sheet states that the whole 32 KB flash is mapped from 0x8000 and is accessible with normal LD/ST instructions as well as LPM. In ordinary application code, use this mapped view mainly for reads. Flash writes are NVMCTRL-controlled operations, not arbitrary byte stores like SRAM writes.

```
.section .progmem.data, "a", @progbits
lookup:
    .byte 0, 1, 1, 2, 3, 5, 8, 13

.section .text

    ; Load lookup[3] into r16 via mapped flash
    ldi ZL, lo8(lookup + MAPPED_PROGMEM_START + 3)
    ldi ZH, hi8(lookup + MAPPED_PROGMEM_START + 3)
    ld  r16, Z          ; r16 = 2
```

Copying with mapped LD — same loop as SRAM copy:

```
ldi ZL, lo8(lookup + MAPPED_PROGMEM_START)
ldi ZH, hi8(lookup + MAPPED_PROGMEM_START)
ldi XL, lo8(dest_buf)
ldi XH, hi8(dest_buf)
ldi r18, 8

copy_mapped:
    ld  r16, Z+          ; load from mapped flash, Z++
    st  X+, r16          ; store to SRAM, X++
    dec r18
    brne copy_mapped
```

5.9.1 LPM vs Mapped LD — Which to Use?

	LPM	Mapped LD
Available on	All AVR	AVRxt only (tinyAVR 1-series, megaAVR 0-series)
Cycles	3	2 base, plus NVM wait time if any
Words	1	1
Z value	Flash byte address	Flash byte address + 0x8000
Pointer modes	Z, Z+ only	X, Y, Z — all modes including displacement
ATtiny3217 coverage	Whole 32 KB flash	Whole 32 KB flash at 0x8000-0xFFFF

On ATtiny3217 with 32 KB flash: **prefer mapped LD** when AVRxt-specific code is acceptable. It uses all pointer modes and reads flash and SRAM with the same instruction family, which is usually the simpler mental model. It is also usually faster than LPM, but remember the NVM wait-state caveat above.

LPM is still worth knowing because it is portable across AVR families, but for this chip, mapped LD is usually the cleaner choice.

5.10 Pointer Arithmetic

5.10.1 Advancing a Pointer by N Bytes

Post-increment advances by one byte per memory access. To skip N bytes, use ADIW:

```
adiw ZL, 4      ; Z += 4 (skip 4 bytes)
adiw XL, 16     ; X += 16
```

ADIW only handles values 0–63. For larger steps, use ADD/ADC on the pair directly:

So pointer arithmetic on AVR comes in two flavors: compact/specialized (ADIW/SBIW) and general/manual (ADD/ADC).

```
; Z += 200
ldi r16, 200
add ZL, r16
ldi r16, 0
adc ZH, r16      ; propagate carry into high byte
```

Or use r1 (which the ABI keeps at 0) for the carry:

```
ldi r16, 200
add ZL, r16
adc ZH, r1      ; r1 = 0 by convention – just propagates carry
```

5.10.2 Computing an Array Element Address

```
; Address of array[i]: base + i * element_size (element_size = 1 here)
; r24 = index, Z = array base
add ZL, r24      ; Z += index (1-byte elements)
adc ZH, r1       ; carry
ld r16, Z        ; load array[i]
```

For 2-byte elements:

```
; i * 2 = i << 1
lsl r24          ; r24 *= 2 (left shift)
rol r25          ; shift carry into r25 (if index > 127)
add ZL, r24
adc ZH, r25
```

This is the same reasoning as in multi-byte arithmetic: build the low byte result first, then propagate carry into the high byte.

5.11 Atomicity on ATtiny3217

Some AVR families implement XCH, LAC, LAS, and LAT, but the ATtiny3217's AVRxt core does **not**. If you need interrupt-safe bit updates on GPIO, use the dedicated peripheral registers such as PORTx_OUTSET, PORTx_OUTCLR, and PORTx_OUTTGL instead of a read-modify-write sequence.

This matters because “load-modify-store” is not automatically safe when an ISR or another hardware path can touch the same state concurrently.

```
ldi r16, (1 << 3)
sts PORTA_OUTSET, r16      ; set PA3 without reading PORTA_OUT first
sts PORTA_OUTCLR, r16     ; clear PA3 atomically
sts PORTA_OUTTGL, r16     ; toggle PA3 atomically
```

For ordinary SRAM bytes shared with an ISR, there is no single instruction on this MCU that performs an atomic arbitrary read-modify-write. Use a critical section (IN/CLI/OUT SREG) when the update must not be interrupted.

So the practical split is:

- for GPIO, prefer the hardware set/clear/toggle registers;
- for shared SRAM state, protect the update with interrupt control when needed.

5.12 Full Addressing Mode Reference

Instruction		Addr calc	Pointer after	Cycles	Words
LDS	Rd, k	k	—	3	2
STS	k, Rr	k	—	2	2
LD	Rd, X	X	unchanged	2	1
LD	Rd, X+	X	X + 1	2	1
LD	Rd, -X	X - 1	X - 1	2	1
ST	X, Rr	X	unchanged	2	1
ST	X+, Rr	X	X + 1	2	1
ST	-X, Rr	X - 1	X - 1	2	1
LD	Rd, Y	Y	unchanged	2	1
LD	Rd, Y+	Y	Y + 1	2	1
LD	Rd, -Y	Y - 1	Y - 1	2	1
LDD	Rd, Y+q	Y + q	unchanged	2	1
ST	Y, Rr	Y	unchanged	2	1
ST	Y+, Rr	Y	Y + 1	2	1
ST	-Y, Rr	Y - 1	Y - 1	2	1
STD	Y+q, Rr	Y + q	unchanged	2	1
LD	Rd, Z	Z	unchanged	2	1
LD	Rd, Z+	Z	Z + 1	2	1
LD	Rd, -Z	Z - 1	Z - 1	2	1
LDD	Rd, Z+q	Z + q	unchanged	2	1
ST	Z, Rr	Z	unchanged	2	1
ST	Z+, Rr	Z	Z + 1	2	1
ST	-Z, Rr	Z - 1	Z - 1	2	1
STD	Z+q, Rr	Z + q	unchanged	2	1
LPM	Rd, Z	Flash[Z]	unchanged	3	1
LPM	Rd, Z+	Flash[Z]	Z + 1	3	1
IN	Rd, A	I/O[A] (0–3F)	—	1	1
OUT	A, Rr	I/O[A] (0–3F)	—	1	1
PUSH	Rr	SRAM[SP]	SP - 1	2	1
POP	Rd	SRAM[SP+1]	SP + 1	2	1

5.13 Worked Example: Reverse a String In-Place

Reverse an ASCII string stored in SRAM. Classic two-pointer approach.

This example is useful because it combines several addressing modes in one small routine: forward scan with Z+, fixed access through X and Z, then converging pointers for the swap loop.

```

/*
 * str_reverse – reverse a null-terminated string in SRAM in-place.
 * Input: X = pointer to first byte of string
 * Clobbers: X, Z, r16, r17, r18
 */
str_reverse:
    ; Find the end: walk Z to the null terminator
    movw ZL, XL                ; Z = X (copy pointer)
1:
    ld    r16, Z+              ; load byte, Z++
    tst   r16                  ; null?
    brne 1b                    ; no – keep walking

    ; Z is now one past the null. Step back twice: past null, to last char.
    sbiw  ZL, 2                ; Z now points to last character

    ; Now swap X[front] and Z[back], moving inward
2:
    cp    XL, ZL                ; has front crossed back? (low bytes)
    cpc   XH, ZH                ; (high bytes)
    brsh done                  ; BRSH: unsigned >=, front >= back → done

    ld    r16, X                ; r16 = front char
    ld    r17, Z                ; r17 = back char
    st    X+, r17               ; store back char at front, advance front
    st    Z, r16                ; store front char at back
    sbiw  ZL, 1                ; retreat back pointer
    rjmp  2b

done:
    ret

```

Trace with “AVR\0” at address 0x3810:

Initial: X = 0x3810 ('A'), Z = 0x3812 ('R')

Iter 1: swap 'A' and 'R' → X=0x3811, Z=0x3811

Iter 2: XL=0x11 >= ZL=0x11 → BRSH taken → done

Result in SRAM: “RVA\0” OK

Notice how little arithmetic the routine needs. Most of the real work is being done by the addressing modes themselves.

5.14 Summary

LDS/STS – direct, 16-bit address in instruction, 3/2 cycles, 2 words.
Use for named variables, single accesses.

LD/ST – indirect via X, Y, or Z. 2 cycles, 1 word.

no mode – pointer unchanged
 X+/Y+/Z+ – post-increment (forward traversal)
 -X/-Y/-Z – pre-decrement (backward traversal)

LDD/STD – Y or Z plus constant offset 0–63. Struct member access.
 X does NOT support displacement.

LPM – flash read via program bus. Z = byte address. 3 cycles.
 Mapped LD – AVRxt: LD with Z = flash addr + 0x8000. 2-cycle base
 cost plus possible NVM wait time. Preferred on ATtiny3217.

ADIW/SBIW – adjust pointer by constant 0–63 in 2 cycles.
 ADD+ADC – adjust pointer by any 16-bit value.

ATtiny3217 data space: I/O 0x0000-0x003F, extended I/O 0x0040-0x0FFF,
 EEPROM 0x1400-0x14FF, SRAM 0x3800-0x3FFF, mapped flash 0x8000-0xFFFF.

ATtiny3217 has no XCH/LAC/LAS/LAT. Use PORTx set/clear/toggle registers for
 GPIO, and critical sections for shared SRAM updates.

5.15 Exercises

1. Write a subroutine `memset(X, val, n)` that fills `n` bytes starting at address `X` with value `val` (in `r16`). Use `ST X+` in the loop.
2. Write `memcpy(X, Z, n)` that copies `n` bytes from `Z` to `X`. Both source and destination are SRAM. Then modify it so the source is flash (use mapped LD for ATtiny3217).
3. Given the struct:
 offset 0: type (1 byte)
 offset 1: length (1 byte)
 offset 2: data (4 bytes)
 Write code using LDD/STD with `Y` to read type, increment length, and zero out the data field — all without moving `Y`.
4. What is the maximum array size (in elements, 1 byte each) you can access with LDD `Y+q`? What happens if you need element 64? Write code that handles this using ADIW and basic LD.
5. Assemble `loadstore.S`, run `avr-nm -n loadstore.elf`, then run `avr-objdump -s -j .text loadstore.elf`. Find `powers_of_two` and verify the bytes match 1, 2, 4, 8, 16, 32, 64, 128. Then find the address of `buf_lpm` in the `.bss` section – what address does it get?
6. Why does the `str_reverse` subroutine use `CP+CPC` instead of just `CP` to compare `X` and `Z`? What would go wrong with only `CP XL, ZL`?

Next: Chapter 6 — Arithmetic and Logic (16-bit and beyond)

Chapter 6

Arithmetic and Logic: 16-bit and Beyond

Chapter 4 introduced 8-bit arithmetic. Real programs quickly outgrow that. Addresses are 16-bit, timers are often 16-bit, counters spill past 255, and scaled sensor values rarely fit in a single byte for long. This chapter shows how AVR handles arithmetic once one byte is no longer enough.

The key idea is simple: AVR does not magically become a 16-bit machine when you want a 16-bit result. You build larger arithmetic explicitly out of 8-bit pieces, and the carry flag is what ties those pieces together.

6.1 Multi-Byte Addition and Subtraction

The AVR has no 16-bit ADD instruction. You build it from two 8-bit instructions using the carry flag as a bridge between bytes.

That sounds more awkward than it is. Once you trust the carry chain, multi-byte arithmetic becomes a repeatable pattern rather than a special trick.

Microchip's instruction summary is useful here because it separates the two classes of arithmetic flags:

Flag	Meaning in this chapter
C	Carry for unsigned addition; borrow for unsigned subtraction
Z	Result is zero
N	Result bit 7 is set for 8-bit ops, or bit 15 for ADIW/SBIW
V	Two's-complement signed overflow
S	$N \text{ xor } V$, used by signed branches such as BRLT/BRGE
H	Half-carry/half-borrow between bits 3 and 4, mainly for BCD-style code

For multi-byte arithmetic, the important flags are C for unsigned carry/borrow and V/S/N for signed interpretation of the final high byte.

6.1.1 The Carry Chain Pattern

Always start with the least-significant byte. Never with the most-significant.

```
; 16-bit add: r25:r24 += r23:r22
add r24, r22      ; low bytes – may produce carry
adc r25, r23      ; high bytes – ADC includes carry from ADD
```

```
; 16-bit subtract: r25:r24 -= r23:r22
sub r24, r22      ; low bytes – may produce borrow (C=1)
sbc r25, r23      ; high bytes – SBC subtracts borrow
```

The carry flag is the pipeline between bytes. Break the chain and the high byte has the wrong value.

That is why instruction order matters here. Arithmetic is not just about the operands; it is also about preserving the flag state from one byte to the next.

6.1.2 24-bit and 32-bit

The pattern extends to any width:

```
; 32-bit add: r3:r2:r1:r0 += r7:r6:r5:r4
add r0, r4
adc r1, r5
adc r2, r6
adc r3, r7
```

```
; 32-bit subtract: r3:r2:r1:r0 -= r7:r6:r5:r4
sub r0, r4
sbc r1, r5
sbc r2, r6
sbc r3, r7
```

After an addition chain: if C=1, unsigned overflow occurred (result too large for N bytes). After a subtraction chain, C=1 means unsigned borrow/underflow. If V=1, signed overflow occurred.

So after a multi-byte operation, the flags still matter. The CPU is telling you whether the final N-byte result wrapped or overflowed in the signed sense.

One subtle detail from the AVR instruction manual: SBC preserves the previous Z value when its own byte result is zero. That is deliberate. It means a multi-byte subtract or compare leaves Z=1 only if **all** bytes in the chain were zero/equal.

6.1.3 Adding a Small Constant to a 16-bit Register Pair

ADIW handles constants 0–63 in 2 cycles, but only on the upper even register pairs:

Low register operand	Full pair
r24	r25:r24
r26 / XL	XH:XL
r28 / YL	YH:YL
r30 / ZL	ZH:ZL

```
adiw r24, 1      ; r25:r24 += 1 – 2 cycles, 1 word
adiw ZL, 16      ; Z += 16
```

SBIW has the same register-pair and immediate limits, but subtracts:

```
sbiw r24, 1      ; r25:r24 -= 1
sbiw ZL, 16      ; Z -= 16
```

ADIW/SBIW update Z, C, N, V, and S. They do not update H. That differs from byte ADD/ADC/SUB/SBC, which do update H.

For constants that fit the target width, use SUBI+SBCI with negated values — because there is no ADDI:

```
; r25:r24 += 100
subi r24, lo8(-100)      ; subtract -100 = add 100
sbc i r25, hi8(-100)     ; propagate carry
```

SUBI and SBCI only work on registers r16-r31. If your value is in a lower register pair, move it or use a register-register add/subtract sequence.

For a 16-bit constant, use the same pattern:

```
; r25:r24 += 1000
subi r24, lo8(-1000)
sbc i r25, hi8(-1000)
```

If the destination pair is not in r16-r31, load the constant into a register pair and use ADD+ADC:

```
; r25:r24 += 1000
ldi r22, lo8(1000)
ldi r23, hi8(1000)
add r24, r22
adc r25, r23
```

This gives you three practical tools for “add a constant”:

- ADIW when the constant is small and the register pair is supported
- SUBI/SBCI when the destination pair is in r16-r31
- ADD/ADC when you want the general case

6.1.4 Negating a 16-bit Value

Two’s complement negate: invert all bits, add 1.

```
; negate r25:r24
com r24      ; ~low
com r25      ; ~high
adiw r24, 1   ; + 1 (or: ldi r16,1 / add r24,r16 / adc r25,r1)
```

Or equivalently using NEG + SBC:

```
neg r25      ; high = 0 - high
neg r24      ; low = 0 - low, C=1 if low was non-zero
sbc r25, r1   ; subtract the borrow from high (r1 must be zero)
```

That order looks backwards, but it is intentional: NEG r24 must run before the final SBC so the carry flag tells whether the low-byte negate borrowed from the high byte.

The reason is that it fits naturally into the carry/borrow model the CPU already uses for subtraction.

A correct register-register form that avoids ADIW is:

```
com r24
com r25
ldi r16, 1
add r24, r16
adc r25, r1      ; r1 must be zero
```

For signed 8-bit values, NEG has one important edge case: 0x80 (-128) cannot be represented as +128 in signed 8-bit two’s complement, so the result stays 0x80 and V is set.

6.2 Shift and Rotate Instructions

Shifts move bits left or right. The bit shifted out goes into the carry flag. Rotates put the old carry flag back in on the vacated side.

If addition and subtraction are about numeric value, shifts and rotates are about bit position. They are the bridge between arithmetic and bit-level work.

LSL Rd – Logical Shift Left

$C \leftarrow [7][6][5][4][3][2][1][0] \leftarrow 0$
 $Rd *= 2$ (unsigned). Old bit 7 $\rightarrow C$.

LSR Rd – Logical Shift Right

$0 \rightarrow [7][6][5][4][3][2][1][0] \rightarrow C$
 $Rd /= 2$ (unsigned). Old bit 0 $\rightarrow C$. Bit 7 forced to 0.

ASR Rd – Arithmetic Shift Right

$b7 \rightarrow [7][6][5][4][3][2][1][0] \rightarrow C$
 $Rd /= 2$ (signed). Bit 7 preserved (sign extension).

ROL Rd – Rotate Left through Carry

$C \leftarrow [7][6][5][4][3][2][1][0] \leftarrow C(\text{old})$
 Like LSL but old C enters at bit 0.

ROR Rd – Rotate Right through Carry

$C(\text{old}) \rightarrow [7][6][5][4][3][2][1][0] \rightarrow C$
 Like LSR but old C enters at bit 7.

All shift/rotate instructions: **1 cycle**. The instruction summary lists C, Z, N, and V as directly affected; S follows from $N \text{ xor } V$. LSL and ROL also affect H because they are aliases for ADD Rd,Rd and ADC Rd,Rd. LSR, ASR, and ROR do not affect H. T and I are not affected. SWAP is separate: it affects no flags.

That makes them cheap enough to use constantly, especially for scaling by powers of two or for moving carries across byte boundaries.

6.2.1 16-bit Shift Left (Multiply by 2)

```
lsl  r24      ; low byte: bit 7 → C
rol  r25      ; high byte: C → bit 0 (carry from low byte propagates)
```

6.2.2 16-bit Shift Right, Logical (Unsigned Divide by 2)

```
lsr  r25      ; high byte: bit 0 → C (MSB filled with 0)
ror  r24      ; low byte: C → bit 7
```

6.2.3 16-bit Shift Right, Arithmetic (Signed Divide by 2)

```
asr  r25      ; high byte: sign bit preserved, bit 0 → C
ror  r24      ; low byte: C → bit 7
```

6.2.4 Shifting by N Bits

AVR does not have a multi-bit shift instruction (unlike x86's `SHL r, cl`). Repeat for each bit:

```
; r16 <=> 3 (multiply by 8)
lsl r16
lsl r16
lsl r16
```

For large shifts on 8-bit values, `SWAP` is faster when shifting by 4:

```
; r16 <=> 4
swap r16      ; swap nibbles: bits 7:4 ↔ bits 3:0
andi r16, 0xF0 ; zero lower nibble (was upper before swap)
```

```
; r16 >=> 4 (logical)
swap r16
andi r16, 0x0F ; zero upper nibble
```

This is a good reminder that on AVR, repeated simple instructions often replace one “fancier” instruction from a larger ISA.

6.2.5 Extracting a Bit Field

Example: extract bits 5:3 of r16 into bits 2:0 of r17:

```
mov  r17, r16
andi r17, 0b00111000 ; mask bits 5:3
lsr  r17              ; shift right 3 times
lsr  r17
lsr  r17              ; r17 = bits 5:3 in positions 2:0
```

6.3 SWAP — Swap Nibbles

`SWAP Rd` ; bits 7:4 ↔ bits 3:0

1 cycle. Does not affect any flags. Pure bit rearrangement.

```
ldi r16, 0xAB
swap r16      ; r16 = 0xBA
```

Primary uses: - Fast $\times 16$ or $\div 16$ on values that fit in a nibble - BCD digit manipulation - Building two-nibble values from separate sources:

```
; Pack two 4-bit values (r16 low nibble, r17 low nibble) into one byte
andi r16, 0x0F      ; keep only low nibble of r16
swap r17             ; move r17's low nibble to high position
andi r17, 0xF0      ; keep only high nibble of r17
or   r16, r17        ; r16 = r17_nibble : r16_nibble
```

`SWAP` is a good example of an instruction that looks niche until you start doing display code, BCD work, or byte packing. Then it becomes surprisingly useful.

6.4 MUL — Unsigned 8×8 Multiply

`MUL Rd, Rr` ; $r1:r0 \leftarrow Rd \times Rr$ (unsigned)

2 cycles. Result always in **r1:r0** regardless of operands.

ATtiny3217 has the hardware multiplier block, so this is a real two-cycle instruction, not a library loop. The multiplier takes two 8-bit inputs and produces a 16-bit result; the integer multiply instructions differ only in how they interpret the inputs.

The fixed result location is the most important thing to remember about MUL. The product does not stay near the source operands. It always lands in r1:r0.

```
ldi r16, 12
ldi r17, 25
mul r16, r17      ; r1:r0 = 300 = 0x012C
movw r24, r0      ; copy result to a usable register pair
clr r1           ; RESTORE r1 = 0 – this is mandatory (see below)
```

Why must you clear r1?

By the GCC AVR ABI, r1 must always be 0 when calling C functions or when C code calls assembly. MUL sets r1 to the high byte of the product. If you forget CLR r1, later code that assumes r1=0 (like ADC Rr, r1 for carry propagation) will produce wrong results.

Develop the habit: every MUL/MULS/MULSU is immediately followed by CLR r1 unless you are certain no C code is involved and you consume r1 before any carry-chain arithmetic.

Even if you are not mixing with C today, it is still a good discipline. The “r1 should be zero” convention is useful across the whole codebase.

6.4.1 Self-Multiply (Square)

```
; r25:r24 = r16 * r16 (r16 squared)
mul r16, r16
movw r24, r0
clr r1
```

MUL allows both operands to be the same register.

6.4.2 8-bit Fraction Scaling

A common embedded pattern: scale an 8-bit value by an 8-bit fraction using the high byte of the product as a “multiply then shift right 8” operation:

This is a classic fixed-point trick: keep the multiply, throw away the low byte, and treat the high byte as the scaled result.

```
; Approximate: result = (val * fraction) >> 8
; val in r16, fraction in r17 (0x00=0%, 0xFF≈100%, 0x80=50%)
mul r16, r17
mov r16, r1      ; take only the high byte → effectively >> 8
clr r1
```

6.5 MULS — Signed 8×8 Multiply

MULS Rd, Rr ; r1:r0 ← (signed)Rd × (signed)Rr

Both operands treated as signed (-128..127). Result is signed 16-bit. **r16-r31 only** for both operands.

```
ldi r16, 0xF6      ; -10 in two's complement
ldi r17, 5
muls r16, r17      ; r1:r0 = -50 = 0xFFCE
```

```

movw r24, r0
clr r1

ldi r16, 0xF6      ; -10
ldi r17, 0xFB      ; -5
muls r16, r17      ; r1:r0 = 50 = 0x0032
movw r24, r0
clr r1

```

6.6 MULSU — Signed × Unsigned Multiply

MULSU Rd, Rr ; r1:r0 ← (signed)Rd × (unsigned)Rr

Rd (r16–r23) is signed; Rr (r16–r23) is unsigned. Result is signed 16-bit.

Used in fixed-point arithmetic where one operand is a coefficient (signed scaling factor) and the other is a measurement (unsigned raw value):

That signed/unsigned mix shows up often in calibration and correction code.

```

; Scale a raw ADC reading (0–255, unsigned) by a signed calibration
; coefficient (-128..127), result in r25:r24.
; Example: reading=200 (0xC8), coeff=-3 (0xFD)
ldi r16, 0xFD      ; signed coefficient -3
ldi r17, 200       ; unsigned reading
mulsu r16, r17     ; r1:r0 = -600 = 0xFDA8
movw r24, r0
clr r1

```

6.7 Fractional Multiply: FMUL, FMULS, FMULSU

The ATtiny3217 multiplier also supports fractional multiply instructions:

```

FMUL   Rd, Rr      ; unsigned fractional
FMULS  Rd, Rr      ; signed fractional
FMULSU Rd, Rr      ; signed × unsigned fractional

```

These still compute an 8×8 product into r1:r0 in 2 cycles, but they shift the 16-bit product left by one bit before storing it. Microchip describes the typical input format as 1.7 fixed point, where the binary point is between bits 6 and 7. The left shift puts the high byte back into the same fractional scale that DSP-style code usually wants.

Operand limits:

Instruction	Rd range	Rr range
FMUL	r16–r23	r16–r23
FMULS	r16–r23	r16–r23
FMULSU	r16–r23	r16–r23

Like the integer multiply family, fractional multiply writes r1:r0 and affects only Z and C. Clear r1 after consuming the product if the surrounding code follows the GCC AVR ABI.

Example: signed Q1.7-style scale. 0x40 represents +0.5, so multiplying 0.5 by 0.5 gives 0.25, represented by 0x20 in the high byte:

```
ldi r16, 0x40 ; +0.5 in signed 1.7 fixed point
ldi r17, 0x40 ; +0.5
fmuls r16, r17 ; r1:r0 = (0x40 * 0x40) << 1 = 0x2000
mov r18, r1 ; high byte = 0x20, which represents +0.25
clr r1
```

Fractional multiply is powerful, but be precise about your fixed-point format. For signed fractional edge cases such as $0x80 * 0x80$, Microchip notes that the shifted result can overflow two's-complement interpretation and software must handle that if it matters.

6.8 Multiply Summary

Instruction	Rd range	Rr range	Signed?	Result
MUL Rd, Rr	r0–r31	r0–r31	both unsigned	r1:r0 unsigned
MULS Rd, Rr	r16–r31	r16–r31	both signed	r1:r0 signed
MULSU Rd, Rr	r16–r23	r16–r23	Rd signed, Rr unsigned	r1:r0 signed
FMUL*	r16–r23	r16–r23	fractional	(product << 1) in r1:r0

All: 2 cycles. Result always in r1:r0. Always CLR r1 after.
All multiply instructions affect only C and Z.

6.9 Division — Software Algorithm

AVR has no divide instruction. Division is implemented with a shift-subtract loop: the same long division algorithm you learned in school, but in binary.

So division is not something the hardware does for you in one step. You are trading simplicity of concept for cost in cycles.

6.9.1 Unsigned 8-bit Division

dividend / divisor → quotient, remainder

The algorithm processes one bit at a time from MSB to LSB:

For each bit (MSB to LSB):

1. Shift MSB of dividend into remainder LSB (remainder <= 1, r[0] = d_msb)
2. If remainder >= divisor:
 - remainder -= divisor
 - set current quotient bit = 1
- Else:
 - quotient bit = 0

If that seems abstract, think of it this way: each iteration asks, “Can the current remainder afford one more divisor?” If yes, subtract and record a 1 in the quotient. If not, record a 0 and continue.

Assembly implementation (clean restoring version):

```
/*
 * div8u – unsigned 8-bit division
 * Input: r16 = dividend, r17 = divisor
 * Output: r18 = quotient, r19 = remainder
 * Clobbers: r20
```



```

*   Cycles: roughly 9-10 cycles per bit, plus call/return if used as a subroutine
*/
div8u_clean:
    clr r18                ; quotient = 0
    clr r19                ; remainder = 0
    ldi r20, 8             ; 8 bits
1:
    lsl r16                ; MSB of dividend → C
    rol r19                ; remainder = (r19 << 1) | C
    lsl r18                ; quotient <=< 1, make room for next bit
    cp r19, r17
    brlo 2f               ; remainder < divisor: quotient bit stays 0
    sub r19, r17           ; remainder -= divisor
    ori r18, 1             ; quotient bit = 1
2:
    dec r20
    brne 1b
    ret

```

Trace: 7 / 3:

```

bit 7: remainder=0→0, 0<3 → quotient bit 0 → r18=0b00000000
bit 6: remainder=0→0, 0<3 → quotient bit 0 → r18=0b00000000
bit 5: remainder=0→0, 0<3 → quotient bit 0 → r18=0b00000000
bit 4: remainder=0→0, 0<3 → quotient bit 0 → r18=0b00000000
bit 3: remainder=0→0, 0<3 → quotient bit 0 → r18=0b00000000
bit 2: remainder=0→1, 1<3 → quotient bit 0 → r18=0b00000000
bit 1: remainder=1→3, 3>=3 → sub→0, quotient bit 1 → r18=0b00000001
bit 0: remainder=0→1, 1<3 → quotient bit 0 → r18=0b00000010

```

```

quotient = 0b00000010 = 2 OK
remainder = 1 OK (7 = 2×3 + 1)

```

Tracing a few divisions by hand is worth the time. After that, the loop stops looking magical and starts looking mechanical.

6.9.2 Unsigned 16-bit Division

Extend the pattern to 16-bit dividend, 8-bit divisor (common for scaled sensor values):

This is the same algorithm again, just stretched over more bits.

```

/*
* div16u8 – divide 16-bit by 8-bit
*   Input: r25:r24 = dividend, r16 = divisor
*   Output: r25:r24 = quotient, r18 = remainder
*   Clobbers: r19, r20, r21, r22, r23
*/
div16u8:
    clr r18                ; remainder low = 0
    clr r22                ; remainder high = 0
    clr r23                ; zero register for 16-bit subtract
    ldi r19, 16            ; 16 bits
    clr r20                ; quotient low
    clr r21                ; quotient high
1:
    lsl r24                ; shift MSB of dividend → C
    rol r25
    rol r18                ; remainder <=< 1, | C
    rol r22                ; keep the 9th remainder bit

```

```

    lsl  r20                ; quotient <= 1
    rol  r21
    tst  r22                ; if high byte is nonzero, remainder >= divisor
    brne 3f
    cp   r18, r16
    brlo 2f
3:
    sub  r18, r16
    sbc  r22, r23
    ori  r20, 1            ; quotient bit = 1
2:
    dec  r19
    brne 1b
    movw r24, r20          ; quotient → r25:r24
    ret

```

Both routines assume the divisor is nonzero. AVR has no hardware divide-by-zero trap; if zero is possible, test it before entering the loop and choose an explicit policy such as saturating the quotient, returning zero, or branching to an error handler.

6.9.3 Power-of-2 Division — Just Shift

If the divisor is a compile-time power of 2, skip the loop entirely:

```

; Unsigned: r16 /= 8 (= 2^3)
lsr r16
lsr r16
lsr r16

; Signed: r16 /= 4 (rounds toward -∞, which is floor division)
asr r16
asr r16

; 16-bit unsigned: r25:r24 /= 4
lsr r25
ror r24
lsr r25
ror r24

```

Always use shifts for power-of-2 division. The loop costs ~72 cycles; three LSRs cost 3.

That is not a small optimization. It is the difference between “expensive software division” and “almost free.”

6.10 Worked Example: 8×8 Unsigned Multiply to 16-bit Result

Full annotated example with verification by objdump.

```

/*
 * Input:  r24 = a (0–255), r22 = b (0–255)
 * Output: r25:r24 = a * b (0–65025)
 * ABI-compliant: returns 16-bit result in r25:r24.
 */
mul8_to_16:
    mul  r24, r22          ; r1:r0 = a * b
    movw r24, r0           ; r25:r24 = result

```

```
clr    r1                ; restore ABI r1=0
ret
```

Disassembly of this function (avr-objdump -d):

```
<mul8_to_16>:
 0: 86 9f          mul     r24, r22
 2: c0 01          movw   r24, r0
 4: 11 24          eor     r1, r1      ; CLR alias
 6: 08 95          ret
```

MUL is a single 16-bit word (2 bytes). The two-cycle cost comes from the hardware multiplier computing the 16-bit product in parallel with the next fetch. MOVW is another single word. The entire routine is 4 instructions, 8 bytes, 8 cycles including RET.

This is a good example of why knowing the ISA matters. A tiny hand-written routine can be both obvious and efficient once you know the dedicated instruction exists.

6.11 Signed Arithmetic Pitfalls

6.11.1 Two Mistakes with Signed Division

Mistake 1: Using LSR (logical) instead of ASR (arithmetic) to divide a signed value.

```
ldi r16, 0xFE          ; = -2 signed
lsr r16                ; r16 = 0x7F = 127 ← WRONG (unsigned right shift)
asr r16                ; r16 = 0xFF = -1 ← CORRECT (-2 / 2 = -1)
```

Mistake 2: Using BRLO/BRSH (unsigned) instead of BRLT/BRGE (signed) for signed comparisons. Covered in Chapter 4 but worth repeating whenever you work with signed arithmetic.

The CPU does not know your intent. A byte is just a bit pattern until your instruction choice tells the machine whether to treat it as signed or unsigned.

6.11.2 Overflow Detection

After a signed operation, V=1 means overflow. The result is wrong as a signed value:

```
ldi r16, 100           ; +100
ldi r17, 100           ; +100
add r16, r17           ; result = 200, but 200 > 127 → V=1 (signed overflow)
brvs overflow_handler
```

After unsigned addition, C=1 means overflow (result > 255 / 65535). After unsigned subtraction, C=1 means borrow/underflow.

That signed/unsigned split is one of the most common sources of subtle bugs in embedded arithmetic.

6.12 Useful Idioms

6.12.1 Test if a value is zero without changing the register

```
tst r16                ; Z=1 if r16=0, N=1 if negative – r16 unchanged
```

TST sets flags. The point is that it lets you ask “is this register zero or negative right now?” without changing the register value itself.

6.12.2 Conditional absolute value

```
; |r16|
sbrc r16, 7      ; skip if bit 7 clear (positive)
neg r16          ; negate if negative
```

6.12.3 Saturating add (clamp at 255 instead of wrapping)

```
add r16, r17      ; add
brcc 1f           ; no carry: result fits in 8 bits
ser r16           ; carry: clamp to 0xFF
1:
```

This is a very common embedded pattern because sensor and UI code often prefers clamping over wraparound.

6.12.4 Saturating subtract (clamp at 0 instead of wrapping)

```
sub r16, r17
brcc 1f           ; no borrow (C=0): result >= 0
clr r16           ; borrow: clamp to 0
1:
```

6.12.5 Sign extension: 8-bit signed → 16-bit signed

```
; r16 (signed) → r17:r16 (signed 16-bit)
clr r17
sbrc r16, 7      ; if bit 7 clear: r17 = 0 (positive, done)
ser r17          ; bit 7 set: r17 = 0xFF (negative, sign extend)
; r17:r16 is now the correct signed 16-bit representation
```

6.13 Summary

Multi-byte arithmetic:

Always low byte first. Use ADC/SBC to chain carry between bytes.

ADIW/SBIW: $\pm 0-63$ to 16-bit pair, 2 cycles.

SUBI+SBCI: add constants to r16-r31 pairs (negate the immediate).

Shifts:

LSL/LSR – logical (unsigned), 0 inserted.

ASR – arithmetic (signed), sign bit preserved.

ROL/ROR – rotate through carry – chains shifts across bytes.

SWAP – nibble swap, 1 cycle, use for $\times 16$ / $\div 16$.

Multiply:

MUL – unsigned $\times$ unsigned $\rightarrow$ r1:r0. Always CLR r1 after.

MULS – signed $\times$ signed $\rightarrow$ r1:r0. r16-r31 only.

MULSU – signed $\times$ unsigned $\rightarrow$ r1:r0. r16-r23 only.

FMUL* – fractional multiply, r16-r23 only, product shifted left once.

All: 2 cycles. Result in r1:r0.

Division: no hardware DIV. Use shift-subtract loop (~72 cycles for 8-bit).
 Power-of-2 divisors: use LSR/ASR instead – $O(\log_2 n)$ cycles.

Signed pitfalls:

Use ASR not LSR for signed right shift.

Use BRLT/BRGE not BRLO/BRSH for signed branch.

V flag = signed overflow, C flag = unsigned carry/borrow.

6.14 Exercises

1. Write a 16-bit multiply: $r25:r24 = r25:r24 \times r23:r22$ (both inputs unsigned). You cannot use MUL for the full 16×16 product directly (it only handles 8×8). Decompose into four 8×8 multiplies and accumulate. Hint: $(a_hi \times 256 + a_lo) \times (b_hi \times 256 + b_lo)$.
 2. Extend `div16u8` to handle divide-by-zero explicitly. Choose a policy (for example: quotient = 0xFFFF, remainder = dividend low byte, or branch to an error handler) and document it.
 3. Write a function that counts the number of 1-bits in `r16` (popcount). Use LSL + ADC `r17, r1` pattern: after each LSL, C holds the shifted-out bit; ADC `r17, r1` adds C to the accumulator (`r1=0`).
 4. Implement saturating 16-bit addition: $r25:r24 + r23:r22$, clamped to 0xFFFF if overflow. The 16-bit carry flag is in C after the second ADC.
 5. Verify the `div8u_clean` routine:
 - Trace $100 / 7$ by hand (expected: quotient=14, remainder=2).
 - Assemble and simulate: use `avr-objdump -d` to confirm the instruction sequence, then trace the bit loop manually.
 6. Without MUL, compute $r16 \times 10$ using only shifts and adds. ($10 = 8 + 2 = (1 \ll 3) + (1 \ll 1)$.)
-

Next: Chapter 7 — Branches and Control Flow

Chapter 7

Branches and Control Flow

Every non-trivial program has decisions and loops. In assembly, all of them reduce to one fundamental act: change the program counter, or deliberately let it continue to the next instruction. This chapter covers the main ways AVR does that — unconditional jumps, conditional branches, and single-instruction skips — and shows the assembly patterns behind decisions, loops, and multi-way dispatch.

The key mental shift is that high-level control flow is just structured PC manipulation.

7.1 Unconditional Jumps

7.1.1 RJMP — Relative Jump

RJMP k ; PC $\leftarrow$ PC + k + 1 (k = signed offset, -2048..+2047 words)

2 cycles. 1 word. Range: **± 2047 words** (± 4094 bytes) from the instruction.

```
rjmp loop      ; jump to label "loop" – GAS computes the offset
rjmp .         ; infinite loop (jump to current address, offset = -1)
```

RJMP covers most local control-flow transfers — nearly 4K instruction words (just under 8 KB of code bytes). Use it by default.

That “use it by default” rule is practical: RJMP is smaller and faster than JMP, and most control-flow targets are close enough to fit.

7.1.2 JMP — Absolute Jump

JMP k ; PC $\leftarrow$ k (k = 22-bit program word address)

3 cycles. 2 words (32-bit instruction). Range: entire 64K-word (128KB) flash.

```
jmp main        ; jump anywhere – no range limit
jmp reset_handler
```

Use JMP only when the target is out of RJMP range, or in the vector table where you need a guaranteed 2-word slot. For everything else, RJMP is smaller and faster.

So the usual decision is simple:

- nearby target: RJMP
- far target or fixed-size vector entry: JMP

7.1.3 IJMP — Indirect Jump via Z

IJMP ; PC $\leftarrow$ Z (Z holds word address of target)

2 cycles. 1 word. Z contains the **word address** (byte address $\div$ 2) of the destination. Use pm(label) to convert a label to its word address:

```
ldi ZL, lo8(pm(target))
ldi ZH, hi8(pm(target))
ijmp          ; jump to target
```

Primary use: **jump tables**. Covered in Section 7.7.

Indirect jumps are less common in beginner code, but they are the step that takes you from simple branch chains to table-driven dispatch.

7.2 Conditional Branches — Mechanics

All conditional branches work identically:

1. Test one or more bits in SREG.
2. If the condition is true: PC += offset + 1 (branch taken, 2 cycles).
3. If the condition is false: PC += 1 (not taken, 1 cycle).

Range: **± 63 words** (± 126 bytes). If the target is farther, invert the condition and use JMP/RJMP:

```
; Branch to far_target if r16 == 0 – but far_target > 63 words away:
tst  r16
brne skip          ; if NOT zero, skip the jump (inverted condition)
jmp  far_target     ; now unconditionally jump to far target
skip:
```

This pattern — invert condition, skip over jump — is the standard way to handle far conditional transfers.

7.2.1 Exact Branch Range

The branch operand is a 7-bit signed word offset: $-64 \dots +63$. Because the CPU adds the offset to PC + 1, the reachable destination range is:

PC - 63 through PC + 64 ; destination instruction word addresses

That is why datasheets often describe conditional branches as reaching PC - 63 to PC + 64, while programmers usually summarize the range as “about ± 64 words.” Both descriptions refer to the same encoding.

All AVR branch and jump offsets are in **instruction words**, not bytes. This matters when you compare a disassembly address, a linker error, and the instruction manual:

- RJMP offset: signed word offset, $-2048 \dots +2047$
- conditional branch offset: signed word offset, $-64 \dots +63$
- flash byte addresses: twice the word address on classic AVR notation

GAS hides most of this with labels, but the distinction matters for jump tables, manual address arithmetic, and datasheet tables that list program addresses in words.

7.2.2 Signed vs Unsigned Conditions

The CPU does not know whether a byte is signed or unsigned. You choose the branch mnemonic that interprets the flags the way your data should be read.

After CP Rd, Rr or CPI Rd, K, the comparison is internally Rd - Rr or Rd - K.

Relation	Unsigned branch	Signed branch
<	BRL0 / BRCS	BRLT
>=	BRSH / BRCC	BRGE
==	BREQ	BREQ
!=	BRNE	BRNE

Example: 0x80 is 128 unsigned but -128 signed.

```
ldi r16, 0x80
cpi r16, 1
brlo unsigned_less    ; false: 128 is not < 1
brlt signed_less      ; true: -128 is < 1
```

This is the most common comparison bug in hand-written assembly: using BRL0/BRSH for signed values or BRLT/BRGE for unsigned values.

7.3 Complete Branch Instruction Table

Mnemonic	Alias for	Condition	Flags tested	Typical use after
BREQ	BRBS 1,k	Z = 1	Z	CP, CPI, TST, SUB
BRNE	BRBC 1,k	Z = 0	Z	CP, CPI, DEC, INC
BRCS	BRBS 0,k	C = 1	C	ADD, SUB, LSL, LSR
BRCC	BRBC 0,k	C = 0	C	ADD, SUB
BRLO	BRCS	C = 1	C	CP, CPI (unsigned <)
BRSH	BRCC	C = 0	C	CP, CPI (unsigned ≥)
BRLT	BRBS 4,k	S = 1 (N≠V)	N, V	CP, CPI (signed <)
BRGE	BRBC 4,k	S = 0 (N≠V)	N, V	CP, CPI (signed ≥)
BRMI	BRBS 2,k	N = 1	N	ADD, SUB (signed neg)
BRPL	BRBC 2,k	N = 0	N	ADD, SUB (signed pos)
BRVS	BRBS 3,k	V = 1	V	ADD, SUB (ovf detect)
BRVC	BRBC 3,k	V = 0	V	ADD, SUB
BRIE	BRBS 7,k	I = 1	I	SEI/CLI context
BRID	BRBC 7,k	I = 0	I	SEI/CLI context
BRHS	BRBS 5,k	H = 1	H	BCD arithmetic
BRHC	BRBC 5,k	H = 0	H	BCD arithmetic
BRTS	BRBS 6,k	T = 1	T	BST/BLD sequences
BRTC	BRBC 6,k	T = 0	T	BST/BLD sequences

BRBS b, k (Branch if Bit in SREG Set) and BRBC b, k (Branch if Bit in SREG Clear) are the generic forms. Every named branch is an alias for one of these two — GAS assembles BREQ as BRBS 1, k (bit 1 = Z flag).

In practice you almost always write the named forms, because BREQ expresses intent better than “branch if SREG bit 1 is set.”

7.4 Skip Instructions

Skips are single-instruction conditional skips — they bypass exactly the next instruction if a condition is met. No label needed. They are the most compact way to write small conditionals.

If branches are for “go somewhere else,” skips are for “ignore the next thing if this condition holds.”

7.4.1 CPSE — Compare, Skip if Equal

CPSE Rd, Rr ; if Rd == Rr: skip next instruction

1 cycle (no skip), 2 cycles (skip 1-word), 3 cycles (skip 2-word).

That timing detail matters because the skipped instruction still occupies space in the stream even though it does not execute.

```
cpse r16, r17
rjmp not_equal ; skipped if r16 == r17 – falls through to equal path
; equal path here
```

CPSE does **not** set SREG. It compares only for the purpose of deciding whether to skip the next instruction. If later code needs Z, C, N, V, or S from the comparison, use CP/CPI instead.

7.4.2 SBRC — Skip if Bit in Register is Clear

SBRC Rr, b ; if bit b of Rr == 0: skip next instruction

Tests bit b (0–7) of register Rr (r0–r31). No flags affected.

```
; Process bit 0 of flags register r16
sbrc r16, 0 ; skip if bit 0 clear
rcall handle_flag0 ; only called if bit 0 was set
```

7.4.3 SBRS — Skip if Bit in Register is Set

SBRS Rr, b ; if bit b of Rr == 1: skip next instruction

```
sbrs r16, 7 ; skip if bit 7 set (value is negative signed)
rjmp positive_path ; only taken if bit 7 was clear
; negative path falls through here
```

7.4.4 SBIC — Skip if Bit in I/O Register is Clear

SBIC A, b ; if bit b of I/O[A] == 0: skip (A = 0x00–0x1F only)

Tests a bit in an I/O register **in a single cycle without loading into a register first**. The address range is limited to 0x00–0x1F (the lower half of the IN/OUT space).

```
; Wait until a flag bit in GPIOR0 is set by an ISR
wait_for_flag:
    sbis GPIOR0, 0 ; skip if bit 0 of GPIOR0 is set
    rjmp wait_for_flag ; not set: loop
; bit 0 now set – proceed
```

7.4.5 SBIS — Skip if Bit in I/O Register is Set

SBIS A, b ; if bit b of I/O[A] == 1: skip (A = 0x00–0x1F only)

```
; Check if a GPIOR flag is clear
sbic GPIOR0, 2      ; skip if bit 2 clear
rcall clear_handler ; only if bit 2 was set
```

7.4.6 SBIC/SBIS Address Limit on ATtiny3217

On ATtiny3217, SBIC/SBIS reach I/O addresses **0x00-0x1F** only. The regular PORT peripheral registers (PORTA_IN, PORTB_IN, etc.) are at 0x0400+ — far beyond this limit. To test a regular PORT pin register, you must load it first:

```
; Read PA3 state into r16, test bit 3
lds r16, PORTA_IN    ; load port input register
sbrc r16, 3          ; skip if PA3 is low
rcall pa3_high_handler
```

The registers that ARE reachable with SBIC/SBIS on ATtiny3217:

I/O Addr	Register	Useful bits
0x02	VPORTA_IN	Fast pin reads for Port A
0x06	VPORTB_IN	Fast pin reads for Port B
0x0A	VPORTC_IN	Fast pin reads for Port C
0x1C	GPIOR0	User-defined flags (bit 0-7)
0x1D	GPIOR1	User-defined flags
0x1E	GPIOR2	User-defined flags
0x1F	GPIOR3	User-defined flags

GPIOR0–3 are key scratch flag registers in low I/O space: readable in 1 cycle with SBIC/SBIS, no load required. VPORTx_IN is the other major use: it gives you fast bit tests on real pins without touching the slower PORTx_IN registers in extended I/O space.

```
; Wait while PA3 is low using the virtual port input register
wait_pa3_high:
    sbis VPORTA_IN, 3      ; VPORTA_IN is I/O address 0x02
    rjmp wait_pa3_high
```

The regular PORTx registers are still the full peripheral interface. The virtual ports map only the common DIR, OUT, IN, and INTFLAGS registers into low I/O space so bit instructions can reach them.

7.4.7 Skip Over 2-Word Instructions Carefully

Skip instructions can skip either a 1-word instruction or a 2-word instruction. On AVRxt, the timing is:

Case	Cycles
Condition false, no skip	1
Skip a 1-word instruction	2
Skip a 2-word instruction	3

This is legal:

```
sbrc r16, 0
jmp far_handler      ; 2-word instruction, skipped as a whole if bit 0 clear
```

But it is easy to misread during review because the skipped instruction takes two words in flash. For small conditionals, prefer skipping a 1-word RJMP or RCALL when the target range allows it.

7.4.8 Skip vs Branch — When to Use Which

Situation	Use
Skip 1 instruction on a bit condition	SBRC/SBRS/SBIC/SBIS
Skip 1 instruction on equality	CPSE
Branch to a label	BREQ, BRNE, etc.
More than 1 instruction in the conditional body	Branch

Skips are best for guard clauses and flag-polling loops. Branches are best for bodies with multiple instructions.

That is the simplest rule for choosing between them: one-instruction body, consider a skip; larger body, use a branch.

7.5 Structured Control Flow Patterns

7.5.1 Two-way decision

```
; r16 = a, r17 = b, result → r18
cp   r16, r17           ; a - b
brsh use_a              ; unsigned a >= b
mov  r18, r17
rjmp done
use_a:
mov  r18, r16
done:
```

Pattern: compare/test -> branch to one path -> execute the other path -> rjmp over the first path -> join at a common label.

That “invert the condition” step is one of the most common assembly habits. You often branch around the code you do not want rather than branch directly into every desired path.

7.5.2 Multi-byte comparisons

For 16-bit values, compare the low byte first with CP, then the high byte with CPC so the carry/borrow flows into the high-byte comparison.

```
; Compare unsigned a = r17:r16 with b = r19:r18
cp   r16, r18           ; low byte
cpc  r17, r19           ; high byte with carry from low-byte compare
brlo a_below_b          ; unsigned a < b
brsh a_same_or_above_b  ; unsigned a >= b
```

For equality, the same CP/CPC sequence leaves Z set only if both bytes are equal:

```
cp   r16, r18
cpc  r17, r19
breq equal16
```

For signed 16-bit compares, use the same CP/CPC sequence but finish with BRLT or BRGE. The signed condition comes from the final high-byte comparison, which is the byte containing the sign bit.

7.5.3 Guard branch

```
tst r16
brne continue      ; NOT zero: skip the return
ret                ; zero: return early
continue:
```

For this pattern, a branch is clearer than forcing CPSE into it:

```
tst r16
breq early_return
rjmp continue
early_return:
ret
continue:
```

This is a good example of a broader rule: the shortest control-flow sequence is not always the clearest one.

7.5.4 Top-tested loop

```
top_test:
    tst r16
    breq top_done      ; condition false: exit
    dec r16            ; body
    rjmp top_test
top_done:
```

Pre-test — loop may execute zero times. Costs one extra RJMP per iteration.

That is why a top-tested loop is often not the most efficient assembly shape.

7.5.5 Bottom-tested loop (preferred in assembly)

```
do_body:
    dec r16            ; body
    brne do_body       ; condition: not zero → repeat
```

This is the most efficient loop form. No RJMP back — the branch IS the loop-back. One instruction of overhead per iteration instead of two.

This is one of the most useful structural transformations in assembly: if the loop body must run at least once, put the test at the bottom.

Use this shape when the body runs at least once:

```
; Count-down loop: r18 iterations, guaranteed >= 1
ldi r18, N
loop:
    ; body
    dec r18
    brne loop          ; one branch instruction, no rjmp
```

7.5.6 Counted loop

Count **down** rather than up — comparison against 0 is free (BRNE after DEC):

```
; Count-down (preferred):
ldi r18, N
1:
```

```

; body
dec r18
brne lb

; Count-up (necessary when index used inside body):
clr r18
l:
; body (r18 = current index)
inc r18
cpi r18, N          ; costs one extra CPI per iteration
brne lb

```

That is why countdown loops appear everywhere in AVR code. Zero is the cheapest comparison target on the machine.

7.5.7 Nested loops

```

ldi r19, 8          ; outer counter
outer:
  ldi r18, 64        ; inner counter – reload every outer iteration
  inner:
    ; body
    dec r18
    brne inner
  dec r19
  brne outer

```

Critical: reload the inner counter at the top of the outer loop body, **before** the inner loop starts. A common mistake is loading it once before the outer loop — it reaches zero after the first outer iteration and never loops again.

This bug is common because the loop still looks nested at a glance. The reload placement is what actually makes it nested.

7.6 Worked Example: Sum an Array

Sum 8 unsigned bytes from flash, store 16-bit result in SRAM.

```

/* sum_array – sums ARRAY_LEN bytes from flash
 * Uses LPM: Z is the flash byte address of test_array.
 * Result in r25:r24
 */

.section .progmem.data
test_array:
  .byte 10, 20, 30, 40, 50, 60, 70, 80    ; sum = 360 = 0x0168
.equ ARRAY_LEN, 8

.section .text

sum_array:
  ldi ZL, lo8(test_array)
  ldi ZH, hi8(test_array)
  clr r24          ; sum_lo = 0
  clr r25          ; sum_hi = 0
  ldi r18, ARRAY_LEN ; counter

```

```

loop:
    lpm r16, Z+           ; byte = flash[Z], Z++
    add r24, r16          ; sum_lo += byte
    adc r25, r1           ; sum_hi += carry (r1 = 0)
    dec r18
    brne loop             ; bottom-tested loop: no rjmp needed

    ; r25:r24 = 360 = 0x0168
    ret

```

LPM reads through the program-memory bus. On AVRxt devices such as the ATtiny3217, flash is also mapped into the data address space starting at 0x8000, but LPM is the portable and explicit way to say “read from flash” in a small example like this.

Step-by-step trace:

Iter	r18	Z points to	r16	r25:r24
1	7	[20]	10	0x000A
2	6	[30]	20	0x001E
3	5	[40]	30	0x003C
4	4	[50]	40	0x0064
5	3	[60]	50	0x0096
6	2	[70]	60	0x00D2
7	1	[80]	70	0x0118
8	0	—	80	0x0168 OK

Why ADC r25, r1 and not ADC r25, r17? ADD r24, r16 sets C if the low-byte sum overflows, and ADC r25, r1 immediately consumes that carry. Using r1 (=0) adds only the carry bit, correctly propagating any overflow from the low byte without needing a dedicated zero register. The later DEC r18 updates N, Z, and V for the loop test, but it does not affect C.

This is a good example of how flag discipline and dataflow interact. The right instruction is not just about values in registers; it is about which flag bit you are intentionally consuming next.

7.7 Jump Tables

For multi-way branches, a jump table is faster than a chain of comparisons. In the compact form below, each table entry is an executable RJMP instruction. IJMP jumps into the selected table entry, and that RJMP then jumps to the handler.

Conceptually, a jump table turns a small command number into an indexed branch.

```

; r16 = command index (0, 1, or 2)

; Bounds check first:
cpi r16, 3           ; 3 = number of table entries
brsh jtable_default  ; >= 3: out of range

; Build Z = pm(jtable) + r16
; Each RJMP in the table is 1 word, so index directly.
ldi ZL, lo8(pm(jtable))
ldi ZH, hi8(pm(jtable))
add ZL, r16           ; offset by selector index
adc ZH, r1            ; carry
ijmp                 ; jump to selected RJMP entry in the table

```

```

jtable:
    rjmp cmd0          ; selector 0
    rjmp cmd1          ; selector 1
    rjmp cmd2          ; selector 2

jtable_default:
    rcall cmd_default
    rjmp jtable_done

cmd0: rcall handler0 ; rjmp table_done implicit via fall-through pattern
      rjmp jtable_done
cmd1: rcall handler1
      rjmp jtable_done
cmd2: rcall handler2
      rjmp jtable_done
jtable_done:

```

How it works:

`pm(jtable)` gives the **word address** of `jtable`. Adding the selector index (in words) moves `Z` to the corresponding `RJMP` instruction. `IJMP` loads the program counter with that word address, so execution continues at the selected table entry. The table entry then executes and jumps to the actual handler. Each entry is 1 word (`RJMP` = 1 word), so the index scales correctly.

That two-step jump looks strange the first time you see it, but it buys you compact entries and simple indexing.

Cycle cost: bounds check (2) + `LDI`×2 (2) + `ADD`+`ADC` (2) + `IJMP` (2) + `RJMP` from table (2) = **10 cycles**. A chain of `CPI`/`BREQ` for 3 entries costs up to $3 \times (1+2) = 9$ cycles on the slowest path — comparable. Jump tables win decisively with more entries.

So jump tables are not automatically better for tiny dispatches. They become attractive once the table is large enough that linear comparison chains stop being cheap.

7.7.1 Address Tables with LPM

Another jump-table style stores raw handler addresses as data and loads the selected entry with `LPM`. That is useful when you want the table to be data rather than executable branch instructions, but it costs more instructions because each address is two bytes and must be loaded into `Z` before `IJMP`.

```

; r16 = selector index, handlers are within 64K words
lsl  r16                ; byte offset = index * 2
ldi  ZL, lo8(addr_table) ; LPM uses a flash byte address
ldi  ZH, hi8(addr_table)
add  ZL, r16
adc  ZH, r16

lpm  r30, Z+            ; low byte of handler word address
lpm  r31, Z             ; high byte of handler word address
ijmp

addr_table:
    .word pm(cmd0)
    .word pm(cmd1)
    .word pm(cmd2)

```

Use the executable-`RJMP` table when handlers are within `RJMP` range and code size matters. Use an address table when you need data-like table handling or when the table is

generated by other tools.

7.7.2 Large Jump Tables with JMP Entries

If handlers are far from the table (RJMP out of range), use JMP (2 words) instead. Then each table entry is 2 words, so multiply the index by 2:

```
; Index × 2 for 2-word entries
lsl r16 ; r16 *= 2
ldi ZL, lo8(pm(jtable))
ldi ZH, hi8(pm(jtable))
add ZL, r16
adc ZH, r1
ijmp

jtable:
    jmp cmd0 ; 2 words each
    jmp cmd1
    jmp cmd2
```

On devices with more than 64K words of program memory, indirect jumps may also need EIJMP and the EIND register. ATtiny3217 has 32 KB of flash, so plain IJMP is enough for all program-memory destinations on this device.

7.8 Polling Loops

A polling loop spins until a condition becomes true — set by hardware or an ISR.

Polling is simple and sometimes appropriate, but it is also the most CPU-hungry waiting strategy because the core keeps executing the test loop continuously.

```
; Wait for USART0 receive complete (bit 7 of USART0_STATUS = RXCIF)
; Cannot use SBIS: address too high. Must use LDS + SBRS.
wait_rxc:
    lds r16, USART0_STATUS ; load USART status
    sbrc r16, 7 ; skip if RXCIF bit set
    rjmp wait_rxc ; not set: loop
    lds r16, USART0_RXDATA1 ; read the received byte
```

For low I/O registers (GPOR), SBIS eliminates the load:

```
; ISR sets bit 0 of GPOR0 when data ready
wait_ready:
    sbis GPOR0, 0 ; 1 cycle: skip if bit 0 set
    rjmp wait_ready ; 2 cycles: loop back
; 3 cycles per iteration vs 5 cycles with LDS+SBRS
```

That difference is exactly why low-I/O flag registers are so useful on AVR.

For real pins on ATtiny3217, prefer VPORtx_IN when the loop only needs a single pin state:

```
; Wait for PB2 high through the virtual port
wait_pb2:
    sbis VPORtb_IN, 2 ; VPORtb_IN is I/O address 0x06
    rjmp wait_pb2
```

Use regular PORTx.IN with LDS when you need the full peripheral register view, or when the code is shared with a device that lacks the same virtual-port layout.

7.9 Branch Optimisation Notes

7.9.1 Bottom-tested saves one instruction per loop

```
; top-tested: 3 instructions per iteration (body + dec + brne + rjmp)
; bottom-tested: 2 instructions per iteration (body + dec + brne)
```

For a loop that runs 1000 times, that's 1000 saved cycles.

Over one or two iterations this does not matter. Over hot loops, it matters a lot.

7.9.2 Count down, not up

DEC r18 + BRNE is **two instructions with no extra comparison**. Counting up needs INC + CPI N + BRNE — three instructions.

This is one of those patterns that looks like a small optimization until you write enough firmware to see it everywhere.

7.9.3 Use SBRC/SBRS for single-bit tests

Instead of:

```
andi r16, (1 << 3)
brne bit3_set
```

Write:

```
sbrs r16, 3
rjmp bit3_clear
; bit3_set path falls through
```

Saves one instruction and doesn't destroy r16.

Preserving the register value is often as important as saving the instruction.

7.9.4 Branch delay slots do not exist on AVR

Unlike ARM or MIPS, the AVR has no branch delay slot. The instruction after a branch is NOT executed if the branch is taken. No surprises.

That makes AVR control flow pleasantly literal: if you branch away, the next sequential instruction does not secretly run first.

7.10 Summary

Unconditional:

- RJMP — ± 2047 words, 2 cycles, 1 word. Default choice.
- JMP — full range, 3 cycles, 2 words. For far targets or vector table.
- IJMP — Z = word address, 2 cycles. For jump tables.

Conditional branches: ± 63 words, 2 cycles taken / 1 not taken.

- Unsigned: BRLO (<), BRSH ($\geq$), BREQ (=), BRNE ($\neq$)
- Signed: BRLT (<), BRGE ($\geq$), BRMI (neg), BRPL (pos)
- Far target: invert condition, skip over JMP.

Skip instructions (1–3 cycles, no label needed):

- CPSE Rd, Rr — skip if Rd == Rr

SBRC	Rr, b	– skip if bit b of Rr clear
SBRs	Rr, b	– skip if bit b of Rr set
SBIC	A, b	– skip if bit b of I/O[A] clear (A ≤ 0x1F)
SBIS	A, b	– skip if bit b of I/O[A] set (A ≤ 0x1F)

Loop patterns:

bottom-tested: most efficient – branch is the loop-back (no extra RJMP)
 counted: count down with DEC+BRNE – no CPI needed
 nested: reload inner counter at start of each outer iteration

Jump tables: IJMP with Z = pm(table) + index. 0(1) dispatch for table entries.

Low-I/O targets such as `VPORTx\_IN` and `GPIOR0–3` are reachable with `SBIC`/`SBIS`. `GPIORx` make especially good fast flag bytes between ISR and main loop.

7.11 Exercises

1. Rewrite the array-sum example to sum **signed** 8-bit values into a signed 16-bit result. What changes? (Hint: ADC r25, r1 is still correct for unsigned carry, but how do you handle sign extension of each negative byte before adding?)
2. Write a function `find_min(X, n)` that finds the minimum byte value in an array of n bytes at SRAM address X. Return the minimum in r24. Use a bottom-tested loop.
3. Extend the jump table example to 8 entries. What is the cycle cost of the dispatch compared to a chain of 8 CPI/BREQ pairs (slowest path)?
4. Write a zero-terminated byte-string length routine: given X pointing to the first byte in SRAM, return the length (not counting the zero terminator) in r24. Use a bottom-tested loop with LD X+ and TST.
5. Implement a byte-search routine: scan n bytes at X for value r22, return pointer to first match in X (or 0 if not found). Use CPSE to detect a match.
6. The polling loop `sbis GPIOR0, 0 / rjmp wait` burns CPU cycles continuously. Name one hardware peripheral on ATtiny3217 that would let you sleep the CPU instead of polling, waking only when the condition is true. (Covered in Chapter 10 — but think about it now.)

Next: Chapter 8 — Subroutines and the Stack

Chapter 8

Subroutines and the Stack

Subroutines let you package a sequence of instructions once and reuse it from many places. The stack is what makes that possible: it records the return address, holds temporary saved registers, and gives each active call its own scratch space when registers are not enough.

This chapter stays at the assembly level. The goal is to understand exactly what RCALL, CALL, RET, PUSH, POP, and stack frames do on ATtiny3217.

8.1 CALL and RCALL

RCALL k ; push return address, PC <- PC + k + 1 (-2048..+2047 words)
CALL k ; push return address, PC <- k (absolute target)

Both instructions push the return address, then branch to the target. The return address is the program-memory word address of the instruction after the call.

On ATtiny3217, the Program Counter is 14 bits wide for 32 KB of flash, so the stored return address is **2 bytes**. Larger AVR devices with wider program counters can use larger return addresses, but this chapter assumes the ATtiny3217.

Instruction	Words	Cycles	Range
RCALL k	1	2	relative, -2048..+2047 words
CALL k	2	3	absolute program address

Use RCALL for nearby subroutines. Use CALL only when the target is outside the relative range or when a fixed 2-word instruction is required.

```
rcall delay_tick ; compact local call
call reset_handler ; absolute call
```

The return address is pushed as a **word pointer**, not as a flash byte address. For example, if the next instruction is at byte address 0x0006, the stored return word address is 0x0003.

8.1.1 Return Address Layout

The ATtiny3217 data sheet specifies that the least significant byte of the return address is pushed first and ends up at the higher SRAM address. The stack grows downward, so after a call the stack looks like this:

Before RCALL, SP = 0x3FFF

```
0x3FFF [free]
0x3FFE [free]
```

After RCALL target, return word address = 0x1234

```
0x3FFF 0x34      ; return low byte, higher address
0x3FFE 0x12      ; return high byte
0x3FFD [free]    ; SP now points here
```

You normally let RET handle this, but the byte order matters when debugging a corrupted stack in memory.

8.1.2 ICALL

ICALL ; push return address, PC <- Z

ICALL calls the word address held in Z (r31:r30). It is the subroutine version of IJMP.

```
ldi ZL, lo8(pm(handler))
ldi ZH, hi8(pm(handler))
icall
```

Use it for dispatch through a table or for a selected handler address. On ATtiny3217, plain ICALL is enough for all program-memory destinations. Devices with more than 64K words of program memory may need EICALL and the EIND register.

8.2 RET and RETI

RET ; pop return address into PC
RETI ; pop return address into PC; complete interrupt return

RET returns from a normal subroutine. RETI returns from an interrupt handler and tells the interrupt controller that interrupt handling is complete. On older AVR cores, RETI also re-enables global interrupts by setting SREG.I; on AVRxt devices such as ATtiny3217, Microchip documents that RETI does not modify SREG.I.

Instruction	Cycles on AVRxt with 16-bit PC	Use
RET	4	normal subroutine return
RETI	4	interrupt return

Do not use RETI for ordinary subroutines. Do not use RET at the end of an interrupt handler; the interrupt controller needs the RETI instruction to finish the interrupt return correctly.

Every call path must return with the stack pointer at the same position it had on subroutine entry. If a subroutine pushes three bytes and pops only two, RET will fetch one byte of ordinary data as part of the return address.

8.3 PUSH and POP

PUSH Rr ; SRAM[SP] <- Rr, then SP decreases by 1
POP Rd ; SP increases by 1, then Rd <- SRAM[SP]

On AVRxt, PUSH is 1 cycle and POP is 2 cycles. Both are 1-word instructions. The stack grows from higher SRAM addresses toward lower SRAM addresses.

```
push r16
push r17
; use r16 and r17
pop r17      ; last pushed, first popped
pop r16
```

Example starting with SP = 0x3FFF:

```
push r16:
  SP = 0x3FFE
  SRAM[0x3FFF] = r16
```

```
push r17:
  SP = 0x3FFD
  SRAM[0x3FFE] = r17
  SRAM[0x3FFF] = r16
```

```
pop r17:
  r17 = SRAM[0x3FFE]
  SP = 0x3FFE
```

The practical rule is simple: pop in the exact reverse order of pushes.

8.4 Stack Pointer Facts on ATtiny3217

The stack pointer is the 16-bit CPU.SP register, exposed as SPL and SPH in I/O space. The ATtiny3217 SRAM range is:

```
SRAM start: 0x3800
SRAM end:   0x3FFF
Size:       2 KB
```

After reset, the stack pointer is automatically set to the highest internal SRAM address. In standalone examples this book still initializes it explicitly:

```
ldi r16, lo8(RAMEND)
ldi r17, hi8(RAMEND)
out SPH, r17
out SPL, r16
```

If you move the stack yourself, set it before any subroutine call executes and before interrupts are enabled. A valid custom stack pointer must point inside SRAM and leave enough room above the lowest possible stack address.

8.4.1 Updating SP Safely

When writing a 16-bit stack pointer value manually, write SPH before SPL:

```
out SPH, r29
out SPL, r28
```

Microchip documents a hardware protection detail: writing SPL automatically disables interrupts for up to four instructions, or until the next I/O-memory write, whichever comes first. Writing high first and low second keeps the observable update coherent during that protected window.

8.5 Register Ownership

A reusable subroutine needs a register contract. This book uses the following assembly convention unless a routine documents otherwise:

Register	Meaning in this book's examples
r0	Scratch, often used by MUL. Not preserved.
r1	Zero register. Keep as 0 except during MUL sequences.
r2-r17	Preserved by the subroutine if used.
r18-r27	Scratch for the subroutine.
r28:r29 (Y)	Preserved if used; preferred frame pointer.
r30:r31 (Z)	Scratch pointer; used by LPM, ICALL, IJMP, tables.
SREG	Not preserved by ordinary routines unless documented.
SP	Must be balanced on every return path.

That gives two useful habits:

- short-lived temporary values go in scratch registers
- values that must survive a nested call should be pushed or kept in preserved registers saved by the callee

Example: a subroutine that uses r14 and r15 must save them:

```
worker:
    push r14
    push r15

    ; body may use r14/r15

    pop  r15
    pop  r14
    ret
```

Example: a caller that needs r20 after a nested call must save it itself:

```
push r20
rcall helper      ; helper may use scratch registers
pop  r20
```

8.5.1 The r1 = 0 Rule

AVR multiply instructions store their result in r1:r0. That means r1 may become non-zero after MUL, MULS, or MULSU.

Many AVR idioms use r1 as a zero register:

```
add r24, r16
adc r25, r1      ; add carry only
```

If r1 is not zero, this sequence is wrong. Clear r1 after every multiply sequence unless the routine immediately consumes and overwrites it before any code can observe it.

```
mul  r16, r17      ; r1:r0 = product
movw r24, r0
clr  r1            ; restore zero register
```

8.6 Prologue and Epilogue Patterns

A prologue is the entry sequence that saves state and allocates stack space. An epilogue restores that state and returns.

8.6.1 Minimal Subroutine

```
small_worker:
    ; use only scratch registers
    ret
```

If a routine only uses scratch registers and has no nested state to preserve, it needs no pushes.

8.6.2 Saving Preserved Registers

```
accumulate:
    push r12
    push r13

    ; body uses r12:r13

    pop r13
    pop r12
    ret
```

The order is part of correctness. Push low/high or high/low as you prefer, but the pops must be the exact reverse.

8.6.3 Stack Frame with Y

Use a stack frame when a routine needs per-call storage that cannot live in registers. Y (r29:r28) is the natural frame pointer because LDD and STD support displacement addressing from Y.

This prologue allocates four local bytes:

```
with_frame:
    push r28
    push r29
    in  r28, SPL      ; Y = SP, currently below saved r29
    in  r29, SPH
    sbiw r28, 4       ; reserve 4 bytes below saved Y
    out SPH, r29
    out SPL, r28
```

After this prologue, Y equals the new SP. The local bytes are addressed as Y+1 through Y+4:

Higher addresses

```
Y+8    return low byte
Y+7    return high byte
Y+6    saved r28
Y+5    saved r29
Y+4    local byte 4
Y+3    local byte 3
Y+2    local byte 2
Y+1    local byte 1
```

Y+0 current SP value, below the allocated local area

Lower addresses

The saved Y bytes are above the local area. After `sbiw`, the active local bytes are normally accessed through Y+1..Y+N. Keep the layout written down when using stack frames. Off-by-one errors here corrupt saved registers or return addresses.

Access example:

```
ldi r16, 10
std Y+1, r16
ldi r16, 20
std Y+2, r16
ldd r17, Y+1
```

Epilogue:

```
adw r28, 4           ; Y/SP points below saved r29; POP pre-increments
out SPH, r29
out SPL, r28
pop r29
pop r28
ret
```

8.6.4 When a Frame Is Worth It

Situation	Pattern
no saved state	no prologue
one or two preserved registers	push/pop only
several per-call temporary bytes	Y frame
recursion	stack storage for every value needed after the nested call
interrupt handler	save SREG and every register the handler modifies

8.7 Stack Depth Analysis

Stack overflow on a small AVR usually has no exception or fault signal. The stack simply overwrites SRAM data. You need to estimate the worst case.

For one active call:

stack bytes = return address bytes + pushed registers + local allocation

For ATtiny3217:

```
return address = 2 bytes
SRAM           = 0x3800..0x3FFF
```

Examples:

```
simple_add:
    2 return bytes + 0 pushes + 0 locals = 2 bytes
```

```
with_frame:
    2 return bytes + 2 saved Y bytes + 4 locals = 8 bytes
```

factorial recursive frame:

2 return bytes + 1 saved n byte = 3 bytes

Peak stack use is the deepest active chain, not the total number of calls over time. If a main routine calls alpha, alpha calls beta, and beta calls gamma, add the frame sizes for all three at the deepest point.

Inspect SRAM allocation with:

```
avr-nm -n subroutines.elf
avr-size -A subroutines.elf
```

Then reserve a safety margin for interrupts. If an interrupt can occur at the deepest subroutine point, add the interrupt handler's pushes and return address to the worst-case stack budget.

8.8 Worked Example: Recursive Factorial

Factorial is a compact example because every recursive level needs to remember one byte: the current n.

```
0! = 1
1! = 1
n! = n * (n - 1)!
```

This routine uses:

```
Input:  r24 = n
Output: r25:r24 = n!
Limit:  n <= 8 for a 16-bit result
```

8! = 40320 = 0x9D80, which fits in 16 bits. 9! = 362880, which does not.

```
factorial:
    cpi    r24, 2
    brlo   .Lbase           ; 0! and 1! return 1

    push   r24               ; save n for after the recursive call
    dec    r24
    rcall  factorial         ; r25:r24 = (n - 1)!

    pop     r16              ; r16 = original n
    rcall  mul8x16           ; r25:r24 = r16 * r25:r24, low 16 bits
    ret

.Lbase:
    clr    r25
    ldi    r24, 1
    ret
```

8.8.1 Stack Trace for factorial(4)

This trace shows only the saved n bytes. Each RCALL also pushes a 2-byte return address.

```
factorial(4):
  push 4
  factorial(3):
    push 3
    factorial(2):
      push 2
      factorial(1):
```

```

    return 1
    pop 2, compute 2 * 1 = 2
    pop 3, compute 3 * 2 = 6
    pop 4, compute 4 * 6 = 24

```

Peak active calls for factorial(4):

```

factorial(4), factorial(3), factorial(2), factorial(1)
4 return addresses = 4 * 2 = 8 bytes
saved n values     = 3 * 1 = 3 bytes
mul8x16 call       = 2 bytes while multiply helper is active

```

At the deepest recursive point before unwinding, the saved-n cost is 3 bytes. During each multiply helper call, one additional return address is active.

8.8.2 8 x 16 Multiply Helper

```

/*
 * mul8x16 - 8-bit x 16-bit -> low 16 bits
 * Input:   r16 = 8-bit factor
 *          r25:r24 = 16-bit factor
 * Output:  r25:r24 = low 16 bits of product
 * Clobbers r0, r1, r22, r23
 */
mul8x16:
    mul    r16, r24          ; low factor byte
    movw   r22, r0           ; r23:r22 = low partial
    mul    r16, r25          ; high factor byte
    add    r23, r0           ; add shifted contribution
    clr    r1
    movw   r24, r22
    ret

```

The helper intentionally truncates the product to 16 bits. That is fine for the documented factorial range, but it would be wrong for a routine promising a full 24-bit result.

8.9 Tail Jumps

A tail jump is a jump to another subroutine as the final action. It avoids creating another return address on the stack.

```

; extra return layer:
wrapper:
    rcall target
    ret

; tail jump:
wrapper:
    rjmp target          ; target's RET returns to wrapper's caller

```

Savings:

```

RCALL + RET replaced by RJMP = 4 cycles saved on ATtiny3217
temporary return address avoided = 2 stack bytes on ATtiny3217

```

That cycle saving is RCALL (2) + RET (4), replaced by RJMP (2).

Use a tail jump only when the current routine has already restored any saved registers and returned SP to its entry value. If there is cleanup left to do, use a normal call and

return.

8.10 Routine Template

Use this as a checklist, not as mandatory boilerplate:

```
.global routine_name
routine_name:
    ; Save preserved registers used by this routine.
    ; push r14
    ; push r15

    ; Optional Y frame.
    ; push r28
    ; push r29
    ; in    r28, SPL
    ; in    r29, SPH
    ; sbiw  r28, N
    ; out   SPH, r29
    ; out   SPL, r28

    ; Body.
    ; Put result in the documented output registers.

    ; Optional Y frame cleanup.
    ; adiw  r28, N
    ; out   SPH, r29
    ; out   SPL, r28
    ; pop   r29
    ; pop   r28

    ; Restore preserved registers in reverse order.
    ; pop   r15
    ; pop   r14
ret
```

For every routine, document:

- input registers
- output registers
- clobbered registers
- preserved registers used internally
- maximum stack use
- numerical limits, if any

8.11 Summary

RCALL k - call relative target, 2 cycles, 1 word, +/-2048-word offset.
 CALL k - call absolute target, 3 cycles, 2 words.
 ICALL - call target in Z, 2 cycles.
 RET - pop return address, 4 cycles.
 RETI - interrupt return; on AVRxt, does not modify SREG.I.

ATtiny3217 return address:

2 bytes, stored as a program-memory word pointer.

Low byte is pushed first and resides at the higher SRAM address.

Stack:

SRAM range: 0x3800..0x3FFF.
 Grows downward.
 PUSH stores at SP, then decreases SP by 1.
 POP increases SP by 1, then loads from SP.
 CALL/RCALL/ICALL decrease SP by 2.
 RET/RETI increase SP by 2.

SP updates:

CPU.SP is SPL/SPH.
 Write SPH before SPL for manual 16-bit updates.
 Writing SPL masks interrupts for up to four instructions or until the next I/O-memory write.

Register discipline:

r1 is the zero register; clear it after multiply instructions.
 Push/pop preserved registers in strict reverse order.
 SP must be balanced on every return path.

Stack usage:

return address bytes + pushed registers + local allocation.
 Add interrupt-handler stack use to the worst-case budget when interrupts can fire at the deepest point.

8.12 Exercises

1. Write `sum_range`: input `r24 = start`, `r22 = end`; output `r25:r24 = sum` of all integers from start through end inclusive. Verify 1 through 10 gives 55.
2. For `with_frame`, draw the stack after `push r28`, after `push r29`, after `sbiw r28, 4`, and after each `std Y+q` store. Assume SP starts at 0x3FFF.
3. Calculate peak stack use for `factorial(8)`, including recursive return addresses, saved `n` bytes, and the active `mul8x16` call during unwinding.
4. Rewrite `factorial` as a counted loop. Compare stack use against the recursive version.
5. Write a byte-swap routine: input `X` points to byte `A` and `Z` points to byte `B`. Swap the two SRAM bytes using one scratch register pair and no stack local storage.
6. Add a stack guard byte at the low end of an allowed stack region. Write a small check routine that branches to `stack_error` if the guard byte has changed.

Next: Chapter 9 - GPIO

Chapter 9

GPIO

GPIO (General-Purpose Input/Output) is the mechanism that connects the CPU to the physical world: sensors, LEDs, switches, relays, and anything else wired to a pin. This chapter covers the four dedicated GPIO instructions — IN, OUT, SBI, CBI — together with SBIC/SBIS for pin-state testing, and builds up to a fully debounced button-controlled LED.

The ATtiny3217 uses a **dual-register model** (VPORT + PORT) that differs from the classic DDRx/PORTx/PINx model found on ATmega parts. Both models are covered here; the worked example targets ATtiny3217.

That difference matters because a lot of AVR examples online assume the older ATmega naming. If you only memorize the old names, modern tinyAVR code will look stranger than it really is.

9.1 GPIO Fundamentals

Every GPIO pin has three properties the firmware controls:

Direction 0 = input (pin floats / follows external signal)
 1 = output (pin is driven high or low by the CPU)

Output value 0 = drive low (when direction = output)
 1 = drive high (when direction = output)

When direction = input:
0 = no pull-up (pin floats if nothing drives it)
1 = weak pull-up to VCC (prevents floating)

Input Readback: reflects actual voltage on pin

On all AVR devices, GPIO is controlled through I/O registers. The exact register names and addresses differ between classic ATmega and the newer tinyAVR 1-series, but the concepts are identical.

So before worrying about the register map, keep the simple mental model in mind: configure direction, choose driven value or pull-up behavior, then read the pin state back when needed.

9.2 Classic DDRx / PORTx / PINx (ATmega-Style)

Most AVR examples online target **ATmega328P** (used on the Arduino Uno). That device uses this three-register model per port:

DDRx	– Data Direction Register.	Bit n: 0 = input, 1 = output.
PORTx	– Output Data / Pull-Up.	Bit n: driven value (output) or pull-up enable (input).
PINx	– Input Data.	Reading bit n gives actual pin state; on many classic AVRs, writing 1 toggles the corresponding PORTx output latch.

On ATmega328P these are in the low I/O space (addresses 0x00–0x1F), so IN, OUT, SBI, CBI, SBIC, and SBIS all work directly:

```
/* ATmega328P example – NOT ATtiny3217 */

/* Configure PB5 (Arduino LED) as output */
sbi  DDRB, 5           ; DDRB bit 5 = 1 (output)

/* Drive LED on */
sbi  PORTB, 5          ; PORTB bit 5 = 1 (high)

/* Drive LED off */
cbi  PORTB, 5          ; PORTB bit 5 = 0 (low)

/* Read button on PD2 (input, pull-up enabled) */
cbi  DDRD, 2           ; DDRD bit 2 = 0 (input)
sbi  PORTD, 2          ; enable pull-up (input + PORTx bit = 1)
in   r16, PIND         ; r16 = all 8 bits of PIND
sbrc r16, 2            ; skip if bit 2 = 0 (button pressed)
rjmp not_pressed
```

ATtiny3217 does not have DDRx/PORTx/PINx registers. It uses the PORT peripheral described in the next section. If you see online examples that reference DDRB or PORTC, they are targeting classic ATmega devices — not the ATtiny3217 used in this book.

That is one of the easiest ways to get misled by otherwise-correct AVR code: it may be valid for the wrong AVR family.

9.3 ATtiny3217 Port Architecture

The ATtiny3217 exposes each physical port through **two register sets**:

9.3.1 VPORT — Virtual Port (Fast Path)

VPORT registers are in the low I/O address space (0x00–0x0B). Because they are below address 0x20, all six GPIO instructions work on them: IN, OUT, SBI, CBI, SBIC, SBIS.

Register	I/O addr	Memory addr	Width	Access
VPORTA_DIR	0x00	0x0000	8-bit	R/W
VPORTA_OUT	0x01	0x0001	8-bit	R/W
VPORTA_IN	0x02	0x0002	8-bit	R/W*
VPORTA_INTFLAGS	0x03	0x0003	8-bit	R/W(W1C)
VPORTB_DIR	0x04	0x0004	8-bit	R/W
VPORTB_OUT	0x05	0x0005	8-bit	R/W
VPORTB_IN	0x06	0x0006	8-bit	R/W*

VPORTB_INTFLAGS	0x07	0x0007	8-bit	R/W(W1C)
VPORTC_DIR	0x08	0x0008	8-bit	R/W
VPORTC_OUT	0x09	0x0009	8-bit	R/W
VPORTC_IN	0x0A	0x000A	8-bit	R/W*
VPORTC_INTFLAGS	0x0B	0x000B	8-bit	R/W(W1C)

W1C = Write 1 to Clear (write 1 to a flag bit to clear it)

R/W* = read pin state; writing 1 toggles the matching OUT bit

Note on tinyAVR 1-series addressing: On classic ATmega, there is a 0x20 offset between I/O addresses and memory addresses (I/O 0x00 = memory 0x0020). On ATtiny3217, this offset does not exist — I/O address N = memory address N. `_SFR_MEM8(0x0000)` and I/O address 0x00 are the same location.

VPORT registers are **aliases** of the full PORT registers. Writing to VPORTA_DIR is identical to writing to PORTA_DIR.

Microchip documents the VPORT mapping as four regular PORT registers:

Regular PORT register Virtual PORT register

PORTx.DIR	VPORTx.DIR
PORTx.OUT	VPORTx.OUT
PORTx.IN	VPORTx.IN
PORTx.INTFLAGS	VPORTx.INTFLAGS

Only those high-traffic registers are mirrored in VPORT. Pin configuration still lives in the full PORTx.PINnCTRL registers, so pull-ups, input inversion, and interrupt-sense setup require LDS/STS.

The big idea is that VPORT gives you the fast instruction-set view of a port, while the full PORT block gives you the richer configuration interface.

9.3.2 PORT — Full Port Registers (Configuration and Atomic Operations)

The full PORT peripheral sits in the extended memory range (above 0x3F) and cannot be accessed with IN/OUT/SBI/CBI. Use LDS/STS.

PORTA base: 0x0400

PORTB base: 0x0420

PORTC base: 0x0440

Offset	Register	Description
+0x00	DIR	Direction (0=in, 1=out) — full byte read/write
+0x01	DIRSET	Atomic direction set (write 1 to set bits)
+0x02	DIRCLR	Atomic direction clear (write 1 to clear bits)
+0x03	DIRTGL	Atomic direction toggle
+0x04	OUT	Output value — full byte read/write
+0x05	OUTSET	Atomic output set (write 1 to set bits)
+0x06	OUTCLR	Atomic output clear (write 1 to clear bits)
+0x07	OUTTGL	Atomic output toggle (write 1 to toggle bits)
+0x08	IN	Input value; write 1 to toggle matching OUT bit
+0x09	INTFLAGS	Interrupt flags — write 1 to clear a flag
+0x0A	Reserved	Reserved on ATtiny3217
+0x0B	Reserved	
+0x0C	Reserved	
+0x0D	Reserved	
+0x0E	Reserved	
+0x0F	Reserved	

+0x10	PIN0CTRL	Pin 0: ISC[2:0] + PULLUPEN + INVEN
+0x11	PIN1CTRL	Pin 1 configuration
...		
+0x17	PIN7CTRL	Pin 7 configuration

The **SET/CLR/TGL** variants perform atomic read-modify-write in hardware: a single STS `PORTA_OUTSET`, `r16` can set individual bits without disabling interrupts. Use these for multi-bit changes where atomicity matters (e.g., if an ISR also touches the same port).

This is one of the nicest parts of the newer AVR GPIO design. Instead of forcing software to read a whole port byte, modify one bit, and write the whole byte back, the peripheral can do the bit update internally.

The **PINnCTRL** registers at offsets 0x10-0x17 hold: - ISC[2:0] (bits 2:0) — interrupt sense (covered in Chapter 10) - PULLUPEN (bit 3) — enables the internal pull-up resistor - INVEN (bit 7) — inverts the pin logic

After reset, Microchip specifies all PORT outputs as tri-stated and the digital input buffers enabled. The PINnCTRL reset value is 0x00, which means no inversion, pull-up disabled, and ISC=0x0 (INTDISABLE: interrupt disabled, input buffer enabled).

Writing 1 to OUTTGL toggles the corresponding OUT bit. Microchip also documents that writing 1 to bit `n` of `PORTx.IN` or `VPORTx.IN` toggles the matching OUT bit, but this book uses OUTTGL because it makes the intent obvious and keeps input reads separate from output writes.

9.4 IN — Read I/O Register

IN Rd, A
 Operands: Rd ∈ {R0..R31}, A ∈ {0..63} (I/O address)
 Operation: Rd ← I/O(A)
 Flags: None
 Cycles: 1
 Encoding: 1011 0AA dddd AAAA

IN reads one byte from an I/O register into a general-purpose register. On ATtiny3217, VPORT registers sit at I/O addresses 0x00-0x0B and are within this range.

Think of IN as “take a snapshot of this register so the CPU can inspect it.”

```
in  r16, VPORTA_IN      ; r16 = current state of PORTA pins (8 bits)
in  r17, VPORTB_OUT     ; r17 = current output values of PORTB
in  r18, SREG           ; r18 = Status Register (a critical use of IN)
```

One common idiom: read SREG before a critical section, write it back after.

This shows up here because GPIO work and interrupt control often meet in the same small critical sections.

```
in  r24, SREG           ; save interrupt enable state
cli                               ; disable interrupts (Chapter 10)
; ... critical section ...
out SREG, r24           ; restore interrupt enable state exactly
```

9.5 OUT — Write I/O Register

OUT A, Rr

Operands: $A \in \{0..63\}$, $Rr \in \{R0..R31\}$
 Operation: $I/O(A) \leftarrow Rr$
 Flags: None
 Cycles: 1
 Encoding: 1011 1AAr rrrr AAAA

OUT writes one byte from a register to an I/O register.

Where IN captures a register value, OUT pushes a whole register value back into the I/O space in one step.

```
out  VPORTA_OUT, r16      ; drive all 8 PORTA output pins at once
out  VPORTB_DIR, r17      ; set direction for all 8 PORTB pins
out  SPL, r16             ; set Stack Pointer Low byte
```

Writing all 8 bits at once with OUT is often less convenient than the single-bit SBI/CBI instructions. Use OUT when you want to change multiple pins simultaneously (e.g., output a byte to an 8-bit bus).

So OUT is not the default for every GPIO task. It is the right tool when the whole byte matters.

9.6 SBI and CBI — Set / Clear Bit in I/O Register

SBI A, b
 Operands: $A \in \{0..31\}$, $b \in \{0..7\}$ ← I/O addresses 0x00–0x1F only
 Operation: $I/O(A).b \leftarrow 1$
 Flags: None
 Cycles: 1 on AVRxt (ATtiny3217), 2 on older AVRe devices

CBI A, b
 Operands: $A \in \{0..31\}$, $b \in \{0..7\}$
 Operation: $I/O(A).b \leftarrow 0$
 Flags: None
 Cycles: 1 on AVRxt (ATtiny3217), 2 on older AVRe devices

SBI and CBI are **atomic** single-bit operations. On ATtiny3217, the hardware internally does read-modify-write in one indivisible step: no interrupt can observe a partial state between the read and the write.

That atomicity is why these instructions are still so valuable even though the newer PORT peripheral has richer register options.

Range restriction: SBI/CBI only reach I/O addresses 0x00–0x1F (the lower 32 I/O registers). The VPORT registers (0x00–0x0B) are all within this range. Full PORT registers (0x0400+) require LDS/STS.

```
sbi  VPORTA_DIR, 3        ; set PA3 as output  (VPORTA_DIR = I/O 0x00)
cbi  VPORTA_OUT, 3        ; drive PA3 low      (VPORTA_OUT = I/O 0x01)
sbi  VPORTA_OUT, 3        ; drive PA3 high
sbi  VPORTB_DIR, 3        ; set PB3 as output
cbi  VPORTB_DIR, 5        ; set PB5 as input
```

When to use SBI/CBI vs. full PORT OUTSET/OUTCLR:

Situation	Instruction	Why
Set or clear one pin in VPORT	SBI/CBI	1-word instruction; 1 cycle on ATtiny3217
Set/clear bits, ISR may also touch port	STS PORTA_OUTSET	Guaranteed atomic
Set multiple bits at once	STS PORTA_OUTSET	Single write, no loop

If you remember only one distinction here, remember this one:

- SBI/CBI are ideal for one bit in low I/O space
- OUTSET/OUTCLR/OUTTGL are ideal for atomic updates in the full port block

9.7 SBIC and SBIS — Skip if Bit Clear / Set

SBIC A, b

Operands: $A \in \{0..31\}$, $b \in \{0..7\}$

Operation: if $I/O(A).b = 0$, skip next instruction

Flags: None

Cycles: 1 (no skip), 2 (skip 1-word instruction), 3 (skip 2-word)

SBIS A, b

Operands: $A \in \{0..31\}$, $b \in \{0..7\}$

Operation: if $I/O(A).b = 1$, skip next instruction

Flags: None

Cycles: 1 (no skip), 2 (skip 1-word instruction), 3 (skip 2-word)

These are the fastest way to branch on a single I/O pin state. Instead of `IN + SBRC/SBRS`, a single instruction reads and conditionally skips in one operation.

That is why SBIC and SBIS feel so natural in polling loops: test plus control flow happens in one instruction.

Wait for a button press (PB2 active-low):

```
.Lwait:
    sbis VPORTB_IN, 2      ; skip if PB2 = 1 (not pressed)
    rjmp .Lpressed        ; PB2 = 0: button is pressed
    rjmp .Lwait
```

Equivalent without SBIS:

```
in    r16, VPORTB_IN      ; 1 cy: read all PORTB pins
sbrs r16, 2               ; 1 cy: skip if bit 2 = 1
rjmp .Lpressed            ; 2 cy: branch if bit was 0 (pressed)
rjmp .Lwait
```

The SBIC/SBIS version saves one register (r16 untouched) and avoids the `IN` instruction.

Small savings like that matter in tight loops and ISR-adjacent code.

Skip behavior for 2-word instructions: SBIC/SBIS skip exactly one instruction slot. If the instruction immediately after is a 2-word instruction (`JMP`, `CALL`, `LDS`, `STS`), the skip costs 3 cycles (the full 2-word instruction is skipped). Skip over a 1-word instruction costs 2 cycles.

So the timing of a skip depends not just on the skip instruction itself, but also on what comes immediately after it.

```
; Good: skip over a 1-word RJMP (common pattern)
sbic VPORTB_IN, 2           ; skip if PB2 = 0 (pressed)
rjmp .Lnot_pressed          ; 1 word – skipped in 2 cycles

; Also fine: skip over a 2-word STS
sbis VPORTA_IN, 3           ; skip if PA3 = 1
sts PORTA_OUTTGL, r16       ; 2 words – skipped in 3 cycles
```

9.8 Configuring a GPIO Output

Two equivalent approaches on ATtiny3217:

Method A — SBI on VPORT (single pin, 1 cycle on ATtiny3217):

```
sbi VPORTA_DIR, 3           ; PA3 = output
sbi VPORTA_OUT, 3           ; initial value: high (Curiosity Nano LED off)
```

Method B — STS with PORT DIRSET/OUTCLR (atomic, any number of pins):

```
ldi r16, (1 << 3)
sts PORTA_DIRSET, r16       ; PA3 = output, other bits unchanged
sts PORTA_OUTSET, r16       ; PA3 driven high: Curiosity Nano LED off
```

Both methods produce identical hardware results for a single pin. Method B is preferred when configuring multiple pins in a single write:

That means you can choose the style based on clarity and scaling, not because one is inherently more “correct” for a single pin.

```
ldi r16, (1 << 3) | (1 << 5) ; configure PA3 and PA5
sts PORTA_DIRSET, r16         ; both pins become outputs
sts PORTA_OUTCLR, r16         ; both driven low
```

9.9 Configuring a GPIO Input with Pull-Up

On ATtiny3217, pin direction defaults to input (0) at reset, so only the pull-up needs configuration. Pull-up lives in the PINnCTRL register, which is above the I/O space — LDS/STS required.

PINnCTRL register layout:

```
Bit: 7      6:4      3      2:0
      INVEN  reserved  PULLUPEN  ISC[2:0]
```

```
PULLUPEN (bit 3) = 0: pull-up off (pin floats if nothing drives it)
                  = 1: pull-up on (pin pulled toward VCC)
ISC[2:0]          = 0x00: interrupt disabled, input buffer enabled
                  = 0x04: interrupt and digital input buffer disabled
                  (other values in Chapter 10)
PORT_PULLUPEN_bm = 0x08
```

The floating input problem: If a digital input pin is not connected to anything and has no pull-up, its voltage floats between GND and VCC. The CPU reads random logic levels that change with temperature, noise, and nearby switching signals. Always enable the

pull-up when using a button wired to GND, add an external pull resistor, or drive the pin from a source you control.

An unconnected digital input is not neutral; it is unreliable.

```
; PB2 as input with pull-up (button wired to GND)
; Direction is already 0 (input) – nothing to do for DIR.
ldi r16, PORT_PULLUPEN_bm      ; = 0x08
sts PORTB_PIN2CTRL, r16        ; PORTB_PIN2CTRL = memory 0x0432

; PB3 as input, NO pull-up (externally driven signal, e.g. UART RX)
ldi r16, 0
sts PORTB_PIN3CTRL, r16        ; ISC = disabled, PULLUPEN = 0
```

Reading a pin:


```
in r16, VPORTB_IN              ; read all 8 PORTB pins
sbrcl r16, 2                   ; skip if bit 2 = 0 (button pressed)
; ... bit 2 = 1: button not pressed ...

; Or more directly:
sbis VPORTB_IN, 2              ; skip if PB2 = 1 (not pressed)
rjmp .Lpressed
```

9.10 Debounce

Mechanical switch contacts **bounce**: when you press or release a button, the contacts make and break contact several times in rapid succession before settling. The CPU is fast enough to see each bounce as a separate button event.

So debouncing is not a cosmetic improvement. It is what turns a noisy physical event into a single logical event.



The standard software fix: after detecting a press, wait a few milliseconds (bounce settles in < 10 ms for most switches), then re-check. Do the same on release.

Simple time-delay debounce:

```
; Wait for press
.Lpoll:
  sbis VPORTB_IN, BTN_BIT      ; skip if not pressed
  rjmp .Lpressed
  rjmp .Lpoll

.Lpressed:
  rcall delay_20ms             ; wait for bounce to settle

  ; Optional: confirm still pressed (guards against glitches)
  sbic VPORTB_IN, BTN_BIT      ; skip if still pressed (PB2 = 0)
```

```

    rjmp .Lpoll                ; was a glitch – restart

    ; Confirmed: genuine press
    ; ... handle event ...

    ; Wait for release
.Lwait_release:
    sbis VPORTB_IN, BTN_BIT    ; skip if released (PB2 = 1)
    rjmp .Lwait_release
    rcall delay_20ms           ; debounce the release
    rjmp .Lpoll

```

For typical tactile buttons, 20 ms is a safe debounce window. Some mechanical switches need 50 ms.

That delay is a human-timescale compromise: long enough to outlast bounce, but short enough that the user does not notice it.

9.11 Worked Example: Button-Controlled LED

LED0 on PA3 (Curiosity Nano, active-low). External button on PB2 (active-low, internal pull-up). Each confirmed press toggles the LED state.

This example combines almost every GPIO concept in one small program: output setup, input pull-up, active-low logic, skip-based polling, debounce, and atomic output update.

Hardware:

ATtiny3217

```

          VCC
          |
PA3 — Curiosity Nano LED0 (on board, active-low)

PB2 —[internal pull-up]—— VCC
PB2 —[button]—— GND

```

Source file: src/gpio.S

```

#include <avr/io.h>

.equ LED_BIT,      3          ; PA3, Curiosity Nano LED0
.equ BTN_BIT,      2          ; PB2
.equ DEBOUNCE_COUNT, 16666    ; ~20 ms at 3.3 MHz (sbiw+brne loop)

.section .vectors, "ax", @progbits
.global __vectors
__vectors:
    rjmp main

.section .text
.global main

main:
    ldi r16, hi8(RAMEND)
    out SPH, r16
    ldi r16, lo8(RAMEND)
    out SPL, r16

```

```

sbi  VPORTA_DIR, LED_BIT    ; PA3 = output
sbi  VPORTA_OUT, LED_BIT    ; LED off (active-low)

ldi  r16, PORT_PULLUPEN_bm ; 0x08
sts  PORTB_PIN2CTRL, r16    ; PB2: input + pull-up, no interrupt

.Lloop:
sbis VPORTB_IN, BTN_BIT     ; skip if PB2 = 1 (not pressed)
rjmp .Lpressed
rjmp .Lloop

.Lpressed:
rcall delay_20ms
sbic VPORTB_IN, BTN_BIT     ; skip if PB2 is still low (still pressed)
rjmp .Lloop                 ; bounce/glitch: ignore

ldi  r16, (1 << LED_BIT)
sts  PORTA_OUTTGL, r16      ; atomic toggle PA3

.Lwait_release:
sbis VPORTB_IN, BTN_BIT     ; skip if PB2 = 1 (released)
rjmp .Lwait_release
rcall delay_20ms
rjmp .Lloop

delay_20ms:
ldi  r26, lo8(DEBOUNCE_COUNT)
ldi  r27, hi8(DEBOUNCE_COUNT)

.Ldly_loop:
sbiw r26, 1                 ; 2 cycles
brne .Ldly_loop             ; 2 cycles taken
ret

```

9.11.1 Why PORTA_OUTTGL instead of SBI/CBI?

SBI VPORTA_OUT, 3 sets the bit; CBI VPORTA_OUT, 3 clears it. Neither alone toggles — you'd need to read the current state first. The naive approach:

```

; Non-atomic toggle – dangerous if ISR touches PORTA_OUT
in   r16, VPORTA_OUT
ldi  r17, (1 << LED_BIT)
eor  r16, r17                ; flip LED bit
out  VPORTA_OUT, r16         ; write back

```

Three instructions. An interrupt that fires between IN and OUT and modifies another PA pin will have its change overwritten by the OUT. The PORTA_OUTTGL approach is atomic: the hardware performs the read, XOR, and write in one step.

That is the real lesson: the shorter code is not just shorter, it is safer when something else can touch the same port state.

9.11.2 Build and Inspect

```

avr-gcc -mmcu=attiny3217 -x assembler-with-cpp -c -o gpio.o src/gpio.S
avr-gcc -mmcu=attiny3217 -nostartfiles -o gpio.elf gpio.o
avr-objdump -d gpio.elf      # inspect disassembly
avr-objcopy -O ihex gpio.elf gpio.hex  # generate flashable HEX

```

Check the disassembly around `.Lpressed` to confirm that `sts PORTA_OUTTGL` encodes as a 2-word `sts` instruction (opcode `0x92..`) with the 16-bit address `0x0407`.

9.12 Summary

<code>IN Rd, A</code>	– read I/O register A (0–63) into Rd.	1 cycle.
<code>OUT A, Rr</code>	– write Rr to I/O register A (0–63).	1 cycle.
<code>SBI A, b</code>	– set bit b of I/O register A (0–31).	1 cycle on AVRxt. Atomic.
<code>CBI A, b</code>	– clear bit b of I/O register A (0–31).	1 cycle on AVRxt. Atomic.
<code>SBIC A, b</code>	– skip next if I/O(A).b = 0.	1/2/3 cycles.
<code>SBIS A, b</code>	– skip next if I/O(A).b = 1.	1/2/3 cycles.

SBI/CBI/SBIC/SBIS range: I/O addresses 0x00–0x1F only.

VPORT registers (0x00–0x0B) are within this range: all six instructions apply.

Full PORT registers (0x0400+) need LDS/STS.

ATTiny3217 GPIO quick-reference:

Configure output:	<code>sbi VPORTA_DIR, n</code>	or	<code>sts PORTA_DIRSET, r16</code>
Drive high:	<code>sbi VPORTA_OUT, n</code>	or	<code>sts PORTA_OUTSET, r16</code>
Drive low:	<code>cbi VPORTA_OUT, n</code>	or	<code>sts PORTA_OUTCLR, r16</code>
Toggle output:	<code>sts PORTA_OUTTGL, r16</code> (atomic)		
Configure input:	(DIR default = 0 – no action required)		
Enable pull-up:	<code>sts PORTB_PIN2CTRL, r16</code>	(PORT_PULLUPEN_bm = 0x08)	
Read input:	<code>in r16, VPORTB_IN</code>	or	<code>sbic/sbis VPORTB_IN, n</code>

Debounce: after detecting edge, wait ~20 ms, re-check, wait for release, wait ~20 ms again before returning to poll loop.

9.13 Exercises

1. Reconfigure the example so the LED is on PC0 instead of PA3. Which registers change? Which instructions change? Is there any difference in the generated machine code size?
2. Write a routine `toggle_n_times` that accepts a count in r24 and toggles PA3 that many times with a 100 ms delay between each toggle. Use `PORTA_OUTTGL` and a subroutine-based delay.
3. The debounce loop calls `delay_20ms` twice: once after press, once after release. Under what button-press pattern could the LED be toggled without a real press if you removed the release debounce?
4. Rewrite the button poll using `SBIC` instead of `SBIS`. Draw the new control flow. Are the two approaches equivalent in terms of cycles per loop iteration?
5. How many cycles does the `delay_20ms` subroutine take from `RCALL` to the final `RET`? Show the sum: prologue, loop body × N, final loop exit, return.
6. Add a second button on PA5 (active-low, pull-up). When BTN1 (PB2) is pressed, turn the LED on; when BTN2 (PA5) is pressed, turn it off. Use only `SBIC/SBIS` — no `IN` instruction.

Next: Chapter 10 — Interrupts

Chapter 10

Interrupts

Polling burns CPU cycles checking whether something has happened. An **interrupt** lets hardware signal the CPU directly: the CPU pauses whatever it was doing, runs a short handler (ISR), and resumes exactly where it left off. This chapter covers the interrupt mechanism, the ATtiny3217 vector table, the SEI/CLI/RETI instructions, and the discipline required to write a correct ISR.

This is where the program stops being a single straight-line thread of control and starts reacting to the outside world asynchronously.

10.1 Why Interrupts?

The polling loop from Chapter 9:

```
.Lloop:
    sbis VPORTB_IN, BTN_BIT    ; check button every iteration
    rjmp .Lpressed
    rjmp .Lloop
```

This consumes 100% of the CPU checking one pin. In a real application, the CPU may need to: update a display, process serial data, manage a state machine — all while remaining responsive to button presses. Polling every possible event source and handling each one in priority order becomes unwieldy.

Interrupts invert the model: the CPU runs its main work and is **notified** when an event occurs. The ISR handles the event in a bounded time, then execution resumes with main code unmodified.

That last requirement is the hard part. An interrupt is only useful if the main code can continue as though nothing important was damaged by the pause.

10.2 The Interrupt Mechanism

10.2.1 What Happens When an Interrupt Fires

On ATtiny3217, think of interrupt entry as a short fixed hardware sequence: finish the current instruction, push the PC, then execute the vector entry. With the JMP-based vector table used by ATtiny3217, the first ISR instruction starts after **6 cycles minimum**: one cycle to finish the current instruction, two cycles to store the PC, and three cycles to execute the vector JMP.

That means interrupt latency is not arbitrary. It is mostly predictable, which is one reason small microcontrollers work well for real-time event handling.

1. CPU finishes the current instruction.
2. PC is pushed onto the stack (2 bytes on ATtiny3217).
3. CPUINT records the active interrupt level.
4. The interrupt vector entry executes (`JMP`` in this chapter's layout).
5. CPU begins the first instruction of the ISR body.

When the ISR executes RETI:

1. PC is popped from the stack (`SP += 2`).
2. CPUINT returns to the previous interrupt state.
3. CPU resumes the instruction after where it was interrupted.

The interrupted code is **unaware** an interrupt occurred — as long as the ISR leaves all registers and SREG flags exactly as it found them.

That “as long as” is the whole discipline of ISR writing.

10.2.2 The I Flag (Global Interrupt Enable)

SREG bit layout:

```

Bit:  7    6    5    4    3    2    1    0
      I    T    H    S    V    N    Z    C
      ↑
      Global Interrupt Enable
      1 = interrupts enabled
      0 = interrupts disabled (masked)

```

The I flag is a master switch. Even if a peripheral has its interrupt configured and pending, the CPU will not respond until `I = 1`. On ATtiny3217, Microchip specifies that this bit is **not** cleared by hardware when entering an ISR and is **not** set by RETI; control it explicitly with SEI, CLI, or a direct write to SREG.

So enabling a peripheral interrupt is never the whole story. The peripheral must be configured, and the global interrupt gate must also be open.

Microchip gives the initialization order explicitly:

1. Configure CPUINT only if the defaults are not enough.
2. Configure the peripheral interrupt condition and enable that source.
3. Set SREG.I globally with SEI.

For this chapter, the CPUINT defaults are enough: normal vector table, fixed priority by vector address, no round-robin scheduling, and no level-1 vector. Advanced CPUINT bits such as IVSEL, CVT, LVL0RR, and LVL1VEC are left unchanged.

10.2.3 Interrupt Priority

On ATtiny3217, lower vector address = higher priority. RESET (0x0000) has the highest priority. If two interrupts become pending simultaneously, the one with the lower vector address fires first.

CPUINT tracks the current interrupt level. By default, all maskable interrupt vectors are level 0, so another level 0 interrupt is kept pending until the current handler returns. One vector can optionally be assigned to level 1 (high priority), and a level 1 interrupt can interrupt a level 0 handler.

For most firmware, leave the default level-0 behavior alone unless there is a clear timing reason to introduce a high-priority interrupt.

CPUINT also has a Compact Vector Table mode (`CPUINT.CTRLA.CVT`). When enabled, all level-0 interrupts share one vector. This saves flash but requires a shared dispatcher that

checks interrupt flags. This book uses the normal table because one vector per source is easier to inspect in assembly.

10.3 SEI and CLI

SEI

Operation: $SREG.I \leftarrow 1$
 Flags: $I = 1$
 Cycles: 1
 Note: The instruction immediately after SEI executes before any pending interrupt fires.

CLI

Operation: $SREG.I \leftarrow 0$
 Flags: $I = 0$
 Cycles: 1
 Note: Guaranteed to prevent any interrupt from firing after CLI.

The one-instruction delay after SEI is important:

```
; Guaranteed: the RJMP executes before any pending interrupt fires.
sei
rjmp .Lidle
```

That behavior prevents awkward races where an interrupt could fire immediately between SEI and the next instruction.

10.3.1 Critical Sections

CLI/SEI create a **critical section** — a code region that cannot be interrupted:

```
cli                ; disable interrupts – begin critical section
; ... read or write shared variables that the ISR also touches ...
sei                ; re-enable interrupts – end critical section
```

A cleaner pattern that preserves the interrupt state rather than unconditionally re-enabling:

```
in  r24, SREG      ; save current SREG (including I flag)
cli                ; disable interrupts
; ... critical section ...
out SREG, r24      ; restore SREG exactly (may or may not re-enable)
```

This is safe if the caller had interrupts disabled — SEI would have incorrectly re-enabled them, but restoring SREG preserves their disabled state.

This is the pattern to prefer in reusable code, because it does not assume that interrupts were enabled before the critical section began.

10.4 RETI — Return from Interrupt

RETI

Operation: $PC \leftarrow \text{stack (pop 2 bytes)}, SP \leftarrow SP + 2$; CPUINT exits ISR state
 Flags: SREG is not restored automatically
 Cycles: 4

On ATtiny3217, RETI is still required even though it does not set the I flag. It tells CPUINT that the interrupt handler has completed and returns the interrupt controller to the previous interrupt level. Never use RET to exit an ISR:

Using RET at end of ISR:	Using RETI at end of ISR:
PC returns, but CPUINT still thinks an ISR is active.	PC returns and CPUINT exits the interrupt state cleanly.
Same/lower-priority requests may remain blocked.	Pending interrupts can be serviced after one instruction.

The gap between executing the OUT SREG, r24 (in the epilogue) and RETI is safe because CPUINT has not yet completed the current interrupt. Pending interrupts are not serviced until after RETI, and Microchip documents that pending interrupts are served only after one instruction of the resumed program has executed.

That is why restoring SREG before RETI is correct even when the saved I bit is 1.

10.5 The Interrupt Vector Table

10.5.1 Layout on ATtiny3217

The vector table occupies the first 124 bytes of flash (31 entries $\times$ 4 bytes each). ATtiny3217 has more than 8 KB of flash, so avr-libc defines `_VECTOR_SIZE` as 4 bytes and every entry must fill a 2-word slot. This book uses one JMP instruction per vector entry.

Byte addr	Vector num	Peripheral	Trigger condition
0x0000	—	RESET	Power-on, external reset, WDT
0x0004	1	CRCSCAN_NMI	CRC scan failure (non-maskable)
0x0008	2	BOD_VLM	Voltage level monitor
0x000C	3	PORTA_PORT	Any PORTA pin with ISC configured
0x0010	4	PORTB_PORT	Any PORTB pin with ISC configured
0x0014	5	PORTC_PORT	Any PORTC pin with ISC configured
0x0018	6	RTC_CNT	RTC counter/compare
0x001C	7	RTC_PIT	RTC periodic interrupt
0x0020	8	TCA0_LUNF/TCA0_OVF	Timer/Counter A low underflow/overflow
0x0024	9	TCA0_HUNF	TCA high-byte underflow (split)
0x0028	10	TCA0_LCMP0/CMP0	TCA compare channel 0
0x002C	11	TCA0_LCMP1/CMP1	TCA compare channel 1
0x0030	12	TCA0_CMP2/LCMP2	TCA compare channel 2
0x0034	13	TCB0_INT	Timer/Counter B0
0x0038	14	TCB1_INT	Timer/Counter B1
0x003C	15	TCD0_OVF	Timer/Counter D overflow
0x0040	16	TCD0_TRIG	TCD trigger
0x0044	17	AC0_AC	Analog Comparator 0
0x0048	18	AC1_AC	Analog Comparator 1
0x004C	19	AC2_AC	Analog Comparator 2
0x0050	20	ADC0_RESRDY	ADC result ready
0x0054	21	ADC0_WCOMP	ADC window compare
0x0058	22	ADC1_RESRDY	ADC1 result ready
0x005C	23	ADC1_WCOMP	ADC1 window compare
0x0060	24	TWI0_TWIS	TWI slave
0x0064	25	TWI0_TWIM	TWI master
0x0068	26	SPI0_INT	SPI
0x006C	27	USART0_RXC	USART RX complete
0x0070	28	USART0_DRE	USART data register empty
0x0074	29	USART0_TXC	USART TX complete
0x0078	30	NVMCTRL_EE	EEPROM ready

10.5.2 Defining the Vector Table in .S

```
.section .vectors, "ax", @progbits
.global __vectors
__vectors:
    jmp main                /* 0x0000: RESET – must be first */
    jmp isr_default         /* 0x0004: CRCSCAN_NMI */
    jmp isr_default         /* 0x0008: BOD_VLM */
    jmp isr_default         /* 0x000C: PORTA_PORT */
    jmp portb_isr           /* 0x0010: PORTB_PORT ← our handler */
    jmp isr_default         /* 0x0014: PORTC_PORT */
    /* ... continue through vector 30 ... */
```

Each consecutive JMP is exactly 4 bytes, so the *n*th entry lands at byte address *n*×4 automatically — no .org directives needed.

Unhandled vectors should point to a catch-all handler that just RETIs, not to main (re-running main on a spurious interrupt would re-initialize your whole program). A simple trap for debugging: rjmp . (infinite loop).

Which default you choose depends on your goal:

- reti if you want unused interrupts to be harmless
- rjmp . if you want unexpected interrupts to halt the system for debugging

```
isr_default:
    reti                    /* ignore unexpected interrupts silently */
```

10.6 ISR Prologue and Epilogue

An ISR may fire at any point during the main code: between any two instructions. The main code has no way to know whether an ISR ran during a particular period. Therefore:

Every register written by an ISR must be saved before use and restored before RETI.

SREG must always be saved. Any instruction that modifies flags (arithmetic, logic, shifts, compares, even SBIW) will corrupt the flags that the interrupted code was relying on.

These are the core ISR rules. Violating them may not fail immediately, which is why ISR bugs can be so frustrating to diagnose.

10.6.1 Minimal ISR (uses r16 only)

```
my_isr:
    push r16                ; save r16 (1 cycle, 1 byte stack)
    in r16, SREG            ; save SREG into r16
    push r16               ; save SREG on stack (1 cycle, 1 byte)

    ; ... ISR body, may freely use r16 ...

    pop r16                ; restore SREG value into r16
    out SREG, r16          ; restore SREG
    pop r16               ; restore r16
    reti
```

10.6.2 Why This Order?

Step	Stack contents (grows down)	SP points to
(on ISR entry)	[ret_hi][ret_lo]	top of ret addr
push r16	[ret_hi][ret_lo][r16]	saved r16
in r16, SREG	r16 = SREG value	
push r16	[ret_hi][ret_lo][r16][SREG_val]	← SP
... body ...		
pop r16	r16 = SREG_val, stack shrinks	
out SREG, r16	SREG restored	
pop r16	r16 = original r16 value	
reti	pop return addr, exit CPUINT ISR state, jump back	

On ATtiny3217, OUT SREG, r16 restores the I bit to whatever value it had when the ISR saved SREG. RETI does not repair SREG; it returns the CPUINT state to the previous interrupt level. The two operations handle different state, so both are required.

That separation is one of the main reasons RETI exists as a distinct instruction instead of reusing ordinary RET.

10.6.3 Saving Multiple Registers

```
my_isr:
    push    r16
    push    r17
    push    r18
    in      r16, SREG
    push    r16                ; SREG last – now on top of stack

    ; ... ISR body using r16, r17, r18 ...

    pop     r16                ; SREG value
    out     SREG, r16
    pop     r18                ; restore in reverse order
    pop     r17
    pop     r16
    reti
```

10.6.4 ISR Stack Budget

Every PUSH costs 1 cycle on AVRxt and 1 byte of stack. The two-register minimal ISR (r16 + SREG) uses: - 2 bytes for the return address (pushed automatically by CPU) - 2 bytes for the two PUSH instructions

Total: 4 bytes of stack per invocation. If main code uses 200 bytes of stack and ISRs nest, add the ISR frame size to your stack budget.

So ISR design is not just about speed. It is also about bounded stack usage.

10.7 Configuring Pin Interrupts on ATtiny3217

Classic ATmega devices have dedicated INT0/INT1 external interrupt pins. ATtiny3217 takes a different approach: **any pin** can generate an interrupt by configuring its PINnCTRL register.

This is more flexible than the old dedicated-pin model, but it also means you think in terms of per-pin configuration rather than special global interrupt pins.

10.7.1 PINnCTRL Layout

Bit:	7	6:4	3	2:0
	INVEN	reserved	PULLUPEN	ISC[2:0]

PULLUPEN = 1: enable internal pull-up (bit 3 = 0x08)

ISC[2:0]: Interrupt Sense Configuration

0x00	=	PORT_ISC_INTDISABLE_gc	-	interrupt disabled, input buffer on
0x01	=	PORT_ISC_BOTHEDGES_gc	-	interrupt on both edges (rise + fall)
0x02	=	PORT_ISC_RISING_gc	-	interrupt on rising edge (low → high)
0x03	=	PORT_ISC_FALLING_gc	-	interrupt on falling edge (high → low)
0x04	=	PORT_ISC_INPUT_DISABLE_gc	-	input buffer off (ADC pins, saves power)
0x05	=	PORT_ISC_LEVEL_gc	-	interrupt while pin is low (level triggered)

For a button wired to GND (active-low, fires on press = falling edge):

```
ldi r16, (PORT_PULLUPEN_bm | PORT_ISC_FALLING_gc) ; = 0x08 | 0x03 = 0x0B
sts PORTB_PIN2CTRL, r16 ; PB2: pull-up + falling edge interrupt
```

Microchip describes the PORT interrupt source as INTn in PORTx.INTFLAGS being raised according to ISC in PORTx.PINnCTRL. One PORTx_PORT vector covers the whole port, so the ISR must inspect INTFLAGS if more than one pin on that port can interrupt.

When changing interrupt sense settings, avoid doing it while an interrupt is being synchronized. The datasheet notes that changing ISC during synchronization may lose a request, and re-enabling an input after INPUT_DISABLE may request an interrupt even with different settings.

10.7.2 PORT INTFLAGS — Clearing Interrupt Flags

When an interrupt fires, the corresponding bit in PORT_INTFLAGS is set. The flag is **sticky**: it stays set until explicitly cleared. If you return from the ISR without clearing the flag, the interrupt fires again immediately.

That sticky-flag behavior is easy to forget the first time you write an ISR. If the flag is not cleared, the hardware still believes the event is pending.

Clearing: write 1 to the flag bit (write-1-to-clear, or W1C semantics).

Method A — STS with full PORT address (any bit):

```
ldi r16, (1 << BTN_BIT) ; bit 2 = PB2
sts PORTB_INTFLAGS, r16 ; write 1 to bit 2 → clears flag
```

Method B — SBI on VPORTB_INTFLAGS (single bit, faster):

```
sbi VPORTB_INTFLAGS, BTN_BIT ; VPORTB_INTFLAGS = I/O 0x07
; SBI writes 1 to bit 2 → clears flag
```

SBI writing a 1 to a W1C register clears the flag. This is counterintuitive (SBI normally “sets” a bit) but correct: the hardware interprets a write-1 as “clear this flag”.

This is one of those peripheral rules you simply have to memorize: W1C means “write one to clear,” even if the instruction name sounds like the opposite.

The flag for PORTB_PORT covers all PORTB pins. If multiple pins are configured with ISC ≠ 0, read PORTB_INTFLAGS first to determine which pin(s) fired before clearing.

10.8 ISR Latency and Timing

With the JMP vector style used here, the CPU reaches the first ISR instruction in **6 cycles minimum** from acknowledgement on ATtiny3217: one cycle to finish the ongoing instruction, two cycles to store the PC on the stack, and three cycles for the vector-table JMP. If the interrupt arrives during a multi-cycle instruction, that instruction must finish first, so total observed latency is higher.

Finish current instruction
Push 2-byte PC to stack
Execute vector-table JMP
Begin ISR body

At 3.3 MHz, one cycle is about 303 ns, so 6 cycles is about 1.8 us before the ISR body starts, not counting any extra wait for a multi-cycle interrupted instruction or synchronizer delay in the peripheral itself.

That is fast enough for many GPIO, timer, and serial-response tasks, but still slow enough that cycle counting matters in tight timing work.

The minimal ISR prologue shown here (PUSH r16, IN r16, SREG, PUSH r16) adds **3 cycles** before the first useful ISR work: - PUSH = 1 cycle - IN = 1 cycle - PUSH = 1 cycle

For time-critical ISRs, minimize the prologue: save only what you use.

In ISR code, every unnecessary instruction is paid on every event.

10.9 Worked Example: Pin Interrupt Toggles LED

Curiosity Nano LED0 on PA3, external button on PB2 (falling edge). The ISR toggles the LED; main loops forever doing nothing.

This is intentionally minimal so the interrupt path is easy to see in isolation.

Source file: src/interrupts.S

```
#include <avr/io.h>

.equ LED_BIT,    3
.equ BTN_BIT,    2
.equ BTN_CTRL,   0x0B      /* PORT_PULLUPEN_bm | PORT_ISC_FALLING_gc */

/* — Vector table — */
.section .vectors, "ax", @progbits
.global __vectors
__vectors:
    jmp main          /* 0x0000: RESET */
    jmp isr_default    /* 0x0004: CRCSCAN_NMI */
    jmp isr_default    /* 0x0008: BOD_VLM */
    jmp isr_default    /* 0x000C: PORTA_PORT */
    jmp portb_isr      /* 0x0010: PORTB_PORT */
    jmp isr_default    /* 0x0014: PORTC_PORT */
    /* ... full source lists the remaining vectors as isr_default ... */

/* — Catch-all ISR — */
.section .text

isr_default:
    reti
```

```

/* — PORTB ISR ————— */
/*
 * Prologue:    PUSH r16, IN r16←SREG, PUSH r16      (3 cycles, 2 stack)
 * Body:        LDI r16, STS PORTA_OUTTGL, SBI VPORTB_INTFLAGS
 * Epilogue:    POP r16, OUT SREG←r16, POP r16, RETI  (9 cycles, 0 stack)
 */
portb_isr:
    push    r16
    in      r16, SREG
    push    r16

    ldi     r16, (1 << LED_BIT)
    sts     PORTA_OUTTGL, r16      /* toggle PA3 / LED0          */
    sbi     VPORTB_INTFLAGS, BTN_BIT /* clear flag (W1C via SBI) */

    pop     r16
    out     SREG, r16
    pop     r16
    reti

/* — main ————— */
.global main

main:
    ldi     r16, hi8(RAMEND)
    out     SPH, r16
    ldi     r16, lo8(RAMEND)
    out     SPL, r16

    sbi     VPORTA_DIR, LED_BIT    /* PA3 = output          */
    sbi     VPORTA_OUT, LED_BIT    /* LED off (active-low)  */

    ldi     r16, BTN_CTRL
    sts     PORTB_PIN2CTRL, r16    /* PB2: pull-up + falling edge interrupt */

    sei                                /* enable global interrupts */

.Lidle:
    rjmp    .Lidle                /* CPU idles, ISR handles everything */

```

10.9.1 Execution Flow Walkthrough

Power-on:

RESET vector (0x0000) → JMP main → main executes
 Configures PA3 output, PB2 pull-up + ISC, SEI
 Enters .Lidle loop

User presses button:

PB2 falls from 1 → 0 (falling edge)
 PORTB_PIN2CTRL ISC detects edge → sets PORTB_INTFLAGS bit 2
 CPU detects pending interrupt (I=1, so enabled)
 Finishes current RJMP .Lidle instruction
 Pushes PC onto stack and records the active interrupt level in CPUINT
 Fetches PORTB_PORT vector (0x0010) → JMP portb_isr
 Executes portb_isr:
 push r16 / save SREG
 sts PORTA_OUTTGL → PA3 toggles (LED changes state)


```
sbi VPORTB_INTFLAGS, 2 → clears flag
restore SREG / pop r16
RETI → pop PC, exit CPUINT ISR state, jump back to .Lidle
```

User presses button again:

Same sequence, LED toggles back

10.9.2 Why Not Debounce in the ISR?

The Chapter 9 polling example used a 20 ms delay loop after each press. That loop busy-waits, which is unacceptable inside an ISR (other time-critical events would be missed, and ISR latency would explode to 20 ms).

Options for interrupt-based debouncing: 1. **Hardware debounce**: RC + Schmitt trigger on the button line. 2. **Timer debounce**: in the ISR, start a timer; the real handler fires when the timer expires (Chapter 11). Until the timer expires, disable the pin interrupt. 3. **Level-based re-arm**: change ISC to ISC_LEVEL or use a state machine in main that re-enables the interrupt after a delay.

For this example, the button is held long enough that multiple edges do not cause visible problems. Finished firmware should use method 2 or 3.

The example demonstrates the interrupt mechanism cleanly; a finished button interface should add a real debounce strategy.

10.9.3 Build and Inspect

```
avr-gcc -mmcu=attiny3217 -x assembler-with-cpp -c -o interrupts.o src/interrupts.S
avr-gcc -mmcu=attiny3217 -nostartfiles -o interrupts.elf interrupts.o
avr-objdump -d interrupts.elf
avr-objcopy -O ihex interrupts.elf interrupts.hex
```

Inspect the disassembly to confirm: - The vector table starts at address 0x0000 with a sequence of jmps. - portb_isr is at the address pointed to by the entry at 0x0010. - The sts PORTA_OUTTGL inside portb_isr is a 4-byte instruction. - sbi VPORTB_INTFLAGS, 2 is a 2-byte instruction.

10.10 Summary

SEI – set I flag (enable global interrupts). 1 cycle.
 CLI – clear I flag (disable global interrupts). 1 cycle.
 RETI – return from ISR: pop PC, exit CPUINT ISR state. 4 cycles. Never use RET here.

Interrupt response sequence (ATtiny3217, 6 cycles minimum with JMP vectors):
 finish current instruction → push PC → execute vector JMP → ISR

ISR discipline:

1. PUSH all registers used. (1 cy, 1 stack byte each)
2. IN r16, SREG → PUSH r16. (save SREG – always)
3. ISR body.
4. POP r16 → OUT SREG, r16. (restore SREG)
5. POP all registers in reverse order.
6. RETI. (not RET!)

Vector table: .section .vectors – sequential JMP instructions.
 4 bytes each. Entry n at byte address n×4. RESET at 0x0000.

Unhandled vectors → `isr_default: reti.`

ATtiny3217 pin interrupt config:

```
sts PORTx_PINnCTRL, r16
PORT_PULLUPEN_bm      = 0x08 (bit 3: pull-up)
PORT_ISC_FALLING_gc    = 0x03 (bits 2:0: falling edge)
PORT_ISC_RISING_gc     = 0x02
PORT_ISC_BOTHEDGES_gc = 0x01
```

Clearing INTFLAGS (W1C):

```
sts PORTx_INTFLAGS, r16      (STS to full PORT – any bit)
sbi VPORTx_INTFLAGS, n      (faster – SBI writes 1 → clears flag)
```

Must be done before RETI or interrupt fires again.

10.11 Exercises

1. The current ISR saves r16 and SREG. If the ISR body also needed to use r17 and r18, where exactly (before/after which existing instruction) should the PUSH r17 and PUSH r18 go? Write the complete modified prologue and epilogue.
 2. What happens if you replace `reti` with `ret` at the end of `portb_isr`? Describe the symptom you would observe and explain why it occurs in terms of CPUINT interrupt state.
 3. The idle loop is `rjmp .Lidle`. If you replaced it with the AVR sleep instruction (opcode 0x9588), the CPU would enter idle sleep mode and halt until the next interrupt. What changes would be required in main to enable sleep mode? (Hint: the SLPCTRL peripheral needs to be configured — look at Chapter 1's linker-script walkthrough or the ATtiny3217 datasheet section on SLPCTRL.)
 4. Configure PB3 for a second button that also fires on falling edge. In `portb_isr`, read `PORTB_INTFLAGS` to determine which pin fired, clear only that flag, and call one of two handlers: `led_on` or `led_off`. Write the complete modified ISR.
 5. Using `avr-objdump -d interrupts.elf`, count the exact number of bytes used by the vector table. Does it match $31 \text{ vectors} \times 4 \text{ bytes} = 124 \text{ bytes}$? What byte address does the first non-vector instruction appear at?
 6. The ISR uses `sts PORTA_OUTTGL, r16` (a 4-byte instruction, 2 cycles). Could you replace it with a 2-byte instruction using the VPORT registers? If so, write the alternative. If not, explain why.
-

Next: Chapter 11 — Timer/Counter

Chapter 11

Power and Clock Options

Clock and power control sit between GPIO/interrupts and the timing peripherals. Before you trust a timer period, a USART baud rate, or a delay loop, you need to know what clock the CPU and peripheral bus are using. Before you put a device to sleep, you need to know which clock domains stop and which interrupt sources can wake the part again.

This chapter covers the ATtiny3217 clock controller (CLKCTRL), the sleep controller (SLPCTRL), and the brown-out detector / voltage level monitor (BOD / VLM) from the point of view of assembly code.

11.1 The Default Clock

After reset, the ATtiny3217 uses the internal 16/20 MHz oscillator (OSC20M) as the main clock source. On the Curiosity Nano parts used in this book, the oscillator fuse selects 20 MHz and the main prescaler starts enabled with a divide-by-6 setting:

$$20 \text{ MHz} / 6 = 3.333333 \text{ MHz}$$

That default is why earlier examples use 3.333 MHz for rough delay loops and timer math. If your firmware changes the prescaler, every software delay, timer TOP value, PWM period, and asynchronous serial baud calculation must be checked again.

The important distinction:

CLK_MAIN	selected oscillator before the main prescaler
CLK_PER	prescaled clock distributed to CPU-visible peripherals
CLK_CPU	CPU/SRAM/NVM clock derived from the same main clock path

Many datasheet tables and register descriptions say CLK_PER because the peripheral bus sees the prescaled clock. In most examples in this book, that is also the effective CPU instruction clock.

11.2 Clock Sources

CLKCTRL_MCLKCTRLA.CLKSEL chooses the main clock source:

CLKSEL	Source
0x0	OSC20M: 16/20 MHz internal oscillator
0x1	OSCULP32K: 32.768 kHz internal ultra-low-power oscillator
0x2	XOSC32K: 32.768 kHz external crystal oscillator
0x3	EXTCLK: external clock input

The internal oscillator is the normal choice for examples and general firmware. It needs no external parts and starts quickly. The 32 kHz sources are useful for low-power timing, RTC/PIT work, or applications where a watch crystal gives better long-term timing than the internal ULP oscillator.

External clock options require board-level care. The datasheet warns that when switching to an external main clock, hardware waits for clock edges before the switch completes. If the selected external clock is not present, firmware may not be able to switch away again without a reset. Treat external clock selection as a board configuration decision, not as a casual run-time experiment.

11.3 Main Prescaler

CLKCTRL_MCLKCTRLB controls the main prescaler:

Bit:	7:5	4:1	0
	-	PDIV	PEN

PEN = 0 disables the prescaler and passes CLK_MAIN through undivided. PEN = 1 enables the prescaler, with PDIV selecting the division factor:

PDIV	Division
0x0	/2
0x1	/4
0x2	/8
0x3	/16
0x4	/32
0x5	/64
0x8	/6
0x9	/10
0xA	/12
0xB	/24
0xC	/48

The reset value of MCLKCTRLB is 0x11: PDIV = 0x8, PEN = 1, so the prescaler divides by 6.

Use documented field values together with header bit-position and bit-mask names when possible:

```
.equ MCLKCTRLB_DIV24, ((0x0B << CLKCTRL_PDIV_gp) | CLKCTRL_PEN_bm)

ldi r17, MCLKCTRLB_DIV24
rcall clock_write_mclkctrlb
```

That is clearer and safer than hard-coding 0x17.

11.4 Configuration Change Protection

Clock registers that can destabilize the system are protected by CCP (Configuration Change Protection). To write one of these registers, write the I/O protection key (0xD8) to CPU_CCP, then write the protected register immediately.

For example, this subroutine writes a new value to CLKCTRL_MCLKCTRLB. The new prescaler byte is passed in r17:

```
.equ CCP_IOREG_KEY, 0xD8

clock_write_mclkctrlb:
```

```
ldi    r16, CCP_IOREG_KEY
sts    CPU_CCP, r16
sts    CLKCTRL_MCLKCTRLB, r17
ret
```

Do not insert a call, branch, long calculation, or unrelated register write between CPU_CCP and the protected register write. The protected write must land inside the timed CCP window.

11.5 Waiting for Clock Switches

When firmware changes the main clock source, CLKCTRL_MCLKSTATUS.SOSC reports that the system oscillator is still changing. Status bits also report whether the selected oscillator is stable:

Bit	Meaning
EXTS	external clock has started
XOSC32KS	32.768 kHz crystal oscillator is stable
OSC32KS	internal 32.768 kHz oscillator is stable
OSC20MS	16/20 MHz oscillator is stable
SOSC	main clock source switch in progress

A conservative clock-source switch follows this shape:

```
; configure/enable the new oscillator if needed
; write CLKCTRL_MCLKCTRLA through CCP

.lwait_switch:
lds    r16, CLKCTRL_MCLKSTATUS
sbrc   r16, CLKCTRL_SOSC_bp
rjmp   .lwait_switch
```

Prescaler-only changes do not normally need this wait loop, but they still need the CCP unlock.

11.6 Sleep Modes

SLPCTRL_CTRLA selects the sleep mode and enables sleep entry:

SMODE	Mode
0x0	Idle
0x1	Standby
0x2	Power-down

SEN must be set before executing the SLEEP instruction.

```
.equ   SLPCTRL_MODE_PDOWN, (0x02 << SLPCTRL_SMODE_gp)

ldi    r16, (SLPCTRL_MODE_PDOWN | SLPCTRL_SEN_bm)
sts    SLPCTRL_CTRLA, r16
sleep
```

Sleep does not mean “resume later by magic.” An enabled interrupt source must wake the device. The CPU wakes, executes the ISR, returns with RETI, and then continues at the instruction after SLEEP.

11.6.1 Idle

Idle stops CPU execution but leaves peripherals running. All interrupt sources can wake the device. Use Idle when you want lower power without changing much about the rest of the system.

11.6.2 Standby

Standby stops more of the system. Peripherals that need to keep working must support and enable their own RUNSTDBY bit. Wake-up time may include oscillator start-up time unless the needed oscillator is already running.

11.6.3 Power-down

Power-down is the deepest sleep mode. The datasheet lists BOD, WDT, and the RTC PIT as active. Wake sources are limited: pin-change interrupt, PIT, VLM, TWI address match, and CCL. Only the PIT part of the RTC is available in Power-down.

That limitation is important. A TCA0 compare interrupt will not wake a Power-down system because TCA0 is not running there. Use PIT or a pin-change interrupt for deep-sleep wakeups.

11.6.4 RTC/PIT clock-source changes

The RTC and PIT use CLK_RTC, which is configured via RTC_CLKSEL. The datasheet notes that the clock source for CLK_RTC must only be changed when the RTC/PIT peripheral is disabled. In practice that means: disable RTC (RTC_RTCEN=0) and PIT (RTC_PITEN=0), wait for any synchronization-busy flags to clear, then write RTC_CLKSEL, then re-enable the function you need.

11.7 Brown-Out and VLM

The Brown-Out Detector monitors VDD against a fuse-selected threshold and can reset the device before voltage drops too low for reliable operation. The Voltage Level Monitor is an early-warning interrupt source set above the BOD threshold.

Key points for firmware:

- The BOD level and the Active/Idle BOD operating mode are loaded from fuses at reset and cannot be changed by normal software.
- The Standby/Power-down BOD mode is also loaded from fuses, but the SLEEP field in BOD_CTRLA can be changed at run time through CCP.
- VLM follows the BOD mode. If BOD is disabled, enabling the VLM interrupt alone is not enough.
- VLM can wake from Power-down and can give code a chance to save state before a brown-out reset threshold is reached.

For many small examples, the right policy is simple: leave BOD fuse policy alone and do not spend power-tuning effort until the application has a measured current budget.

11.8 Worked Example: Change the Prescaler

src/clock_prescale.S shows a minimal firmware image that changes the main prescaler through CCP and blinks LED0. It starts from the factory default 3.333 MHz, switches to 20 MHz, blinks quickly, then switches to 416.7 kHz (20 MHz / 48) and blinks slowly.

Build:

```
avr-gcc -mmcu=attiny3217 -x assembler-with-cpp -c -o clock_prescale.o src/clock_prescale.S
avr-gcc -mmcu=attiny3217 -nostartfiles -o clock_prescale.elf clock_prescale.o
avr-objdump -d clock_prescale.elf
```

The important routine is short:

```
clock_write_mclkctrlb:
    ldi    r16, CCP_IOREG_KEY
    sts    CPU_CCP, r16
    sts    CLKCTRL_MCLKCTRLB, r17
    ret
```

If the LED timing changes when this program runs, the CPU really did change speed. That is a crude test, but it is a useful first confirmation before you start recalculating timer and baud-rate values.

11.9 Worked Example: Power-Down Wake from PIT

src/sleep_pit.S configures the RTC PIT from the internal 1 kHz clock, enables a 1024-cycle PIT interrupt, selects Power-down sleep, and sleeps in the foreground loop. Each PIT interrupt wakes the device and toggles LED0.

The foreground loop is intentionally boring:

```
.Lsleep_forever:
    sleep
    rjmp  .Lsleep_forever
```

The useful work happens in the wake ISR. This is the usual low-power shape: configure a wake source, enable global interrupts, enter sleep, do brief work after wake, and go back to sleep.

11.10 Practical Rules

1. Decide the clock frequency before writing timer, USART, ADC, or delay-loop code.
2. Use header bit masks/positions such as CLKCTRL_PDIV_gp and CLKCTRL_PEN_bm instead of scattering raw register bytes through code.
3. Keep the CCP unlock and protected write adjacent.
4. Do not enter sleep until at least one valid wake interrupt is configured and global interrupts are enabled.
5. Use Idle for simple interrupt-driven waiting, Standby for peripheral-assisted waiting, and Power-down only when the wake source is known to work there.
6. Treat BOD fuses as part of the hardware power policy. Runtime code can adjust only the sleep-mode BOD behavior.

Microchip source notes:

- ATtiny3216/3217 Complete Datasheet DS40002205A, sections 10 (CLKCTRL), 11 (SLPCTRL), and 17 (BOD).
- Microchip online peripheral address map lists SLPCTRL at 0x0050, CLKCTRL at 0x0060, and BOD at 0x0080.
- The avr-libc device header iotn3217.h defines the register names, bit masks, and bit positions used by the assembly examples.

11.11 Exercises

1. Change `clock_prescale.S` to use `/24` instead of `/48`. What byte should `MCLKCTRLB` receive?
2. Recalculate the Chapter 11 `TCA0_CMP0` value for a 20 MHz clock with the timer prescaler still set to 1024.
3. Change `sleep_pit.S` from PIT period field value `0x09` (1024 cycles) to `0x0A` (2048 cycles). How often should the LED toggle from the 1 kHz PIT clock?
4. Replace the PIT wake source with a PORT pin-change wake source. Which sleep modes can still wake from that interrupt?
5. Enable VLM interrupt-on-crossing in a test program. Why should this be tested with a current-limited bench supply instead of a normal USB cable?

Chapter 12

Timer/Counter

Timers are the heartbeat of embedded firmware. They let you generate precise waveforms, schedule periodic work, measure elapsed time, and produce PWM signals — all without busy-waiting and without tying up the CPU. This chapter covers the TCA0 (Timer/Counter Type A) peripheral on the ATtiny3217, its main waveform-generation modes, and the interrupts each mode can generate. The worked example produces a 1 Hz timer tick that toggles LED0 from a CMP0 compare-match interrupt.

This is where you start handing time-sensitive repetition over to hardware instead of burning CPU cycles in delay loops.

12.1 Timer Fundamentals

12.1.1 What a Timer Actually Is

A timer is a hardware counter clocked by a divided version of the CPU clock. Each clock tick, the counter register (CNT) increments by 1. When CNT reaches a programmed limit (TOP), something happens — an interrupt fires, an output pin toggles, CNT resets — depending on the mode. The CPU is not involved; the peripheral manages everything in hardware.

That is the mental model to keep in view: a timer is just a counter plus rules about what to do when the counter reaches certain values.

The key equation:

$$f_{\text{interrupt}} = f_{\text{cpu}} / (\text{prescaler} \times (\text{TOP} + 1))$$

Rearranged for TOP given a desired frequency:

$$\text{TOP} = (f_{\text{cpu}} / (\text{prescaler} \times f_{\text{desired}})) - 1$$

Most timer configuration problems reduce to this equation plus careful attention to which register plays the role of TOP in the chosen mode.

12.1.2 Prescaler

The prescaler divides the CPU clock before it reaches the counter. Without it, a 16-bit counter at 20 MHz would overflow in 3.2 ms — useful for microsecond timing but not for a 1 Hz blink. Available prescaler values on TCA0:

Prescaler	Divides by	Max period @ 3.333 MHz (16-bit counter)
-----------	------------	---

DIV1	1	19.7 ms (65535 ticks)
------	---	-----------------------

DIV2	2	39.3 ms
DIV4	4	78.6 ms
DIV8	8	157.3 ms
DIV16	16	314.6 ms
DIV64	64	1.26 s
DIV256	256	5.03 s
DIV1024	1024	20.1 s

For a 1 Hz timer tick, DIV1024 gives 20 seconds of headroom with a 16-bit counter — plenty of room.

So the prescaler is what lets one timer peripheral cover both fast and slow timescales.

12.2 TCA0 on ATtiny3217

ATtiny3217 has one Timer/Counter Type A (TCA0) and two Timer/Counter Type B (TCB0, TCB1). TCA0 is the most versatile:

- 16-bit counter (CNT)
- Three compare channels (CMP0, CMP1, CMP2) each with an output pin option
- Three waveform generation modes
- Single mode (16-bit) or split mode (two independent 8-bit timers)

This chapter uses **single mode** (the default).

That is the easiest mode to reason about because the whole 16-bit counter stays together as one timer.

12.2.1 TCA0 Register Map

Register	Address	Width	Purpose
TCA0_SINGLE_CTRLA	0x0A00	8	Prescaler, enable
TCA0_SINGLE_CTRLB	0x0A01	8	Waveform generation mode, compare enable
TCA0_SINGLE_CTRLC	0x0A02	8	Compare output value override
TCA0_SINGLE_CTRLD	0x0A03	8	Split mode enable (leave 0 for single)
TCA0_SINGLE_CTRLCLR	0x0A04	8	Command/update bits clear side
TCA0_SINGLE_CTRLSET	0x0A05	8	Command/update bits set side
TCA0_SINGLE_CTRLFCLR	0x0A06	8	Buffer-valid flags clear side
TCA0_SINGLE_CTRLFSET	0x0A07	8	Buffer-valid flags set side
Reserved	0x0A08	8	Reserved
TCA0_SINGLE_EVCTRL	0x0A09	8	Event control
TCA0_SINGLE_INTCTRL	0x0A0A	8	Interrupt enable bits
TCA0_SINGLE_INTFLAGS	0x0A0B	8	Interrupt flags (W1C)
TCA0_SINGLE_DBGCTRL	0x0A0E	8	Debug run override
TCA0_SINGLE_TEMP	0x0A0F	8	Temporary bits for 16-bit access
TCA0_SINGLE_CNT	0x0A20	16	Current counter value (L at 0x20, H at 0x21)
TCA0_SINGLE_PER	0x0A26	16	Period (TOP in normal/PWM mode)
TCA0_SINGLE_CMP0	0x0A28	16	Compare 0 (also TOP in FRQ mode)
TCA0_SINGLE_CMP1	0x0A2A	16	Compare 1
TCA0_SINGLE_CMP2	0x0A2C	16	Compare 2
TCA0_SINGLE_PERBUF	0x0A36	16	Period buffer (double-buffered update)
TCA0_SINGLE_CMP0BUF	0x0A38	16	CMP0 buffer
TCA0_SINGLE_CMP1BUF	0x0A3A	16	CMP1 buffer
TCA0_SINGLE_CMP2BUF	0x0A3C	16	CMP2 buffer

All TCA0 registers are outside the 0x00-0x3F I/O range reached by IN/OUT, so they require LDS/STS. The 16-bit registers use two consecutive bytes: the L suffix register at the lower

address must be read or written first (low byte first for 16-bit accesses on AVR).

That means timer code is one of the places where the full memory-mapped peripheral model becomes very visible.

12.2.2 CTRLA — Control A

Bit:	7:4	3:1	0
	_____	_____	_____
	reserved	CLKSEL	ENABLE

CLKSEL[3:1]:

TCA_SINGLE_CLKSEL_DIV1_gc	= (0x00 << 1)	prescaler /1
TCA_SINGLE_CLKSEL_DIV2_gc	= (0x01 << 1)	prescaler /2
TCA_SINGLE_CLKSEL_DIV4_gc	= (0x02 << 1)	prescaler /4
TCA_SINGLE_CLKSEL_DIV8_gc	= (0x03 << 1)	prescaler /8
TCA_SINGLE_CLKSEL_DIV16_gc	= (0x04 << 1)	prescaler /16
TCA_SINGLE_CLKSEL_DIV64_gc	= (0x05 << 1)	prescaler /64
TCA_SINGLE_CLKSEL_DIV256_gc	= (0x06 << 1)	prescaler /256
TCA_SINGLE_CLKSEL_DIV1024_gc	= (0x07 << 1)	prescaler /1024

ENABLE (bit 0):

- 1 = timer running
- 0 = timer stopped (CNT holds its value)

Write ENABLE last after configuring the mode and TOP — avoids glitches where the timer runs briefly with incorrect settings.

This is a recurring embedded pattern: configure first, enable last.

12.2.3 CTRLB — Control B

Bit:	7	6	5	4	3	2:0
	—	CMP2EN	CMP1EN	CMP0EN	ALUPD	WGMODE[2:0]

CMP2EN/CMP1EN/CMP0EN: 1 = compare output drives the associated pin

ALUPD: auto-lock update of double-buffered registers on OVF/TOP

WGMODE[2:0]:

TCA_SINGLE_WGMODE_NORMAL_gc	= 0x00	Normal — count to PER, interrupt on OVF
TCA_SINGLE_WGMODE_FRQ_gc	= 0x01	Frequency — count to CMP0, toggle pin, interrupt on CMP0
0x02		Reserved
TCA_SINGLE_WGMODE_SINGLESLOPE_gc	= 0x03	PWM — count to PER, duty via CMPn
TCA_SINGLE_WGMODE_DSTOP_gc	= 0x05	Dual-slope PWM (top)
TCA_SINGLE_WGMODE_DSBOTH_gc	= 0x06	Dual-slope (top + bottom)
TCA_SINGLE_WGMODE_DSBOTTOM_gc	= 0x07	Dual-slope (bottom)

12.2.4 INTCTRL and INTFLAGS

Bit:	6	5	4	0
	CMP2	CMP1	CMP0	OVF

INTCTRL: writing 1 to a bit enables that interrupt.

INTFLAGS: bit set when interrupt fires. Clear by writing 1 (W1C).

TCA_SINGLE_OVF_bm	= 0x01	Overflow / Underflow
TCA_SINGLE_CMP0_bm	= 0x10	Compare 0 match
TCA_SINGLE_CMP1_bm	= 0x20	Compare 1 match

(pin goes low) (output low again) (pin goes high, OVF interrupt)

W00 output: 

↑ CMP0 ↑ PER (OVF)

CNT counts to PER. WO0 is high from 0 to CMP0, then low from CMP0 to PER. Duty cycle is approximately $CMP0 / (PER + 1)$; CMP0 = BOTTOM produces a static low output and CMP0 > TOP produces a static high output. OVF interrupt fires when CNT wraps. This is the standard PWM mode for motor control, LED dimming, etc.

PWM mode is where the timer stops being just a scheduler and starts being a waveform generator.

PWM frequency:

$$\begin{aligned} f_{\text{pwm}} &= f_{\text{cpu}} / (\text{prescaler} \times (\text{PER} + 1)) \\ \text{Duty} &\approx \text{CMP0} / (\text{PER} + 1) \end{aligned}$$

12.4 Computing Timer Values

For a 1 Hz interrupt from a 3.333 MHz CPU using FRQ mode and prescaler 1024:

$$\begin{aligned} \text{TOP} &= (\text{f_cpu} / (\text{prescaler} \times \text{f_desired})) - 1 \\ &= (3333333 / (1024 \times 1)) - 1 \\ &= 3255.2 - 1 \\ &= 3254.2 \rightarrow \text{round to } 3254 \end{aligned}$$

Actual frequency check:

$$\begin{aligned} f &= 3333333 / (1024 \times (3254 + 1)) \\ &= 3333333 / 3332096 \\ &= 1.000371 \text{ Hz} \quad (\text{error} < 0.04\%) \end{aligned}$$

The 16-bit value 3254 in hex is 0x0CB6:

```
3254 = 0x0CB6
  lo8(3254) = 0xB6
  hi8(3254) = 0x0C
```

Low byte must be written first. The ATtiny3217 uses a 16-bit write latch: writing CMP0L latches the value, and writing CMP0H completes the transfer to the 16-bit register atomically. Reversing the order produces a glitch where the intermediate (wrong) value is briefly seen by the timer.

This is not stylistic ceremony. The write order is part of the hardware interface contract.

```
ldi    r16, lo8(3254)
sts    TCA0_SINGLE_CMP0L, r16    /* writes 0xB6, latches */
ldi    r16, hi8(3254)
sts    TCA0_SINGLE_CMP0H, r16    /* writes 0x0C, transfers 0x0CB6 atomically */
```

12.5 TCA0 Interrupt Vectors

From the ATtiny3217 vector table (Chapter 10):

Byte addr	Vector	Peripheral	Trigger
0x0020	8	TCA0_LUNF/OVF	Low underflow / overflow
0x0024	9	TCA0_HUNF	Split mode high-byte underflow
0x0028	10	TCA0_LCMP0/CMP0	Compare 0 match
0x002C	11	TCA0_LCMP1/CMP1	Compare 1 match
0x0030	12	TCA0_CMP2/LCMP2	Compare 2 match

In FRQ mode, the **TCA0_CMP0** interrupt (vector 10, address 0x0028) fires on every compare match. FRQ mode also reaches TOP at CMP0, so the OVF flag can be set at TOP, but this example does not enable the OVF interrupt. It enables only CMP0 and clears only CMP0.

So the interrupt source depends on the chosen mode, not just on the peripheral name.

To place `tca0_cmp0_isr` at vector 10, the vector table needs `jmp tca0_cmp0_isr` as vector entry 10 (counting RESET as entry 0):

```
__vectors:
    jmp main          /* 0x0000: RESET      */
    jmp isr_default   /* 0x0004: CRCSCAN    */
    jmp isr_default   /* 0x0008: BOD        */
    jmp isr_default   /* 0x000C: PORTA      */
    jmp isr_default   /* 0x0010: PORTB      */
    jmp isr_default   /* 0x0014: PORTC      */
    jmp isr_default   /* 0x0018: RTC_CNT    */
    jmp isr_default   /* 0x001C: RTC_PIT    */
    jmp isr_default   /* 0x0020: TCA0_OVF   */
    jmp isr_default   /* 0x0024: TCA0_HUNF  */
    jmp tca0_cmp0_isr /* 0x0028: TCA0_CMP0 ← */
```

12.6 The ISR

The TCA0 CMP0 ISR follows the same prologue/epilogue discipline from Chapter 10. It adds one step: **clear the INTFLAGS bit before returning**, or the interrupt fires again immediately.

```
tca0_cmp0_isr:
    push    r16
    in      r16, SREG
    push    r16

    ldi     r16, (1 << LED_BIT)
    sts     PORTA_OUTTGL, r16      /* toggle LED0 on PA3 (active low) */

    ldi     r16, TCA_SINGLE_CMP0_bm /* = 0x10 */
    sts     TCA0_SINGLE_INTFLAGS, r16 /* clear CMP0 flag (W1C) */

    pop     r16
    out     SREG, r16
    pop     r16
    reti
```

PORTA_OUTTGL is a toggle register — writing a 1 to bit *n* toggles PORTAn without affecting other bits. It is at address 0x0407 (outside I/O space, requires STS). Writing $(1 \ll \text{LED_BIT}) = 0x08$ toggles only PA3.

TCA0_SINGLE_INTFLAGS uses W1C semantics: writing TCA_SINGLE_CMP0_bm (bit 4 = 1) clears only the CMP0 flag. If you needed to clear all flags simultaneously: `ldi r16, 0xFF; sts`

TCA0_SINGLE_INTFLAGS, r16.

As with every interrupt source on AVR, forgetting to clear the flag turns one interrupt into an endless interrupt loop.

12.7 Main: Timer Initialization Sequence

Order matters. The safe sequence is:

1. Set CTRLB (waveform mode) – timer is not running yet
2. Write CMP0L then CMP0H – latch and load TOP value
3. Set INTCTRL – enable the interrupt we want
4. Set CTRLA – prescaler + enable bit (starts the timer)
5. SEI – enable global interrupts

Starting the timer before configuring CMP0 would cause the timer to run briefly with CNT=0 and CMP0=0, firing immediately. Starting before INTCTRL is set means the first match is missed silently (no interrupt enabled yet).

Initialization order matters because hardware begins reacting as soon as the enable bit is set.

main:

```
ldi    r16, hi8(RAMEND)
out    SPH, r16
ldi    r16, lo8(RAMEND)
out    SPL, r16

sbi    VPORTA_DIR, LED_BIT    /* PA3 = output (LED0 on Curiosity Nano) */
sbi    VPORTA_OUT, LED_BIT    /* LED off initially (active low: HIGH = off) */

/* 1. Waveform mode */
ldi    r16, TCA_SINGLE_WGMODE_FRQ_gc
sts    TCA0_SINGLE_CTRLB, r16

/* 2. TOP value: CMP0 = 3254 */
ldi    r16, lo8(3254)
sts    TCA0_SINGLE_CMP0L, r16
ldi    r16, hi8(3254)
sts    TCA0_SINGLE_CMP0H, r16

/* 3. Enable CMP0 interrupt */
ldi    r16, TCA_SINGLE_CMP0_bm
sts    TCA0_SINGLE_INTCTRL, r16

/* 4. Start timer: prescaler 1024, enable */
ldi    r16, (TCA_SINGLE_CLKSEL_DIV1024_gc | TCA_SINGLE_ENABLE_bm)
sts    TCA0_SINGLE_CTRLA, r16

/* 5. Global interrupt enable */
sei
```

.Lidle:

```
rjmp .Lidle
```

12.8 Worked Example: 1 Hz Timer Tick

Full source: `src/timer_ctc.S`

The complete program combines the vector table, the CMP0 ISR, and main into one file. No shared state is needed between main and the ISR — main initializes hardware and parks in the idle loop; the ISR does all the work.

That makes this example a good demonstration of timer-driven design: main sets the policy, the timer peripheral generates the schedule, and the ISR performs the periodic action.

12.8.1 Register Usage

`r16`: `scratch` – used in prologue/epilogue and ISR body
(saved/restored each invocation)

No other registers needed. The simplest possible ISR.

LED0 on Curiosity Nano is active low (PA3 high = LED off, PA3 low = LED on). `PORTA_OUTTGL` toggles PA3 on every CMP0 match. With a 1 Hz interrupt, the LED changes state once per second, so a complete on-off cycle takes about two seconds.

12.8.2 Timing Analysis

Event	Cycles
TCA0 CMP0 match (hardware)	0
Interrupt request to CPU	1–2 (sync)
CPU finishes current RJMP	2
Push PC, enter CPUINT ISR state	2
JMP <code>tca0_cmp0_isr</code>	3
Total to ISR first instruction	~7
ISR prologue (PUSH + IN + PUSH)	3 cycles (1+1+1)
Toggle (LDI + STS)	3 cycles (1+2)
Clear flag (LDI + STS)	3 cycles (1+2)
ISR epilogue (POP + OUT + POP + RETI)	9 cycles
Total ISR duration	~18 cycles

At 3.333 MHz: $18 \text{ cycles} \times 300 \text{ ns} \approx 5.4 \text{ } \mu\text{s}$ ISR runtime per tick

The timer period is $3333333 / (1024 \times 3255) \approx 1.0 \text{ s}$. The ISR runtime ($\sim 5.4 \text{ } \mu\text{s}$) is 0.00054% of the period — nearly zero CPU overhead for the tick.

This is exactly why timers are so valuable. The CPU only pays for the tiny ISR runtime, not for the whole one-second wait.

12.8.3 Build Commands

```
avr-gcc -mmcu=attiny3217 -x assembler-with-cpp -c -o timer_ctc.o src/timer_ctc.S
avr-gcc -mmcu=attiny3217 -nostartfiles -o timer_ctc.elf timer_ctc.o
avr-objcopy -O ihex timer_ctc.elf timer_ctc.hex

# Inspect disassembly
avr-objdump -d timer_ctc.elf
```

12.8.4 Verifying with avr-objdump

Expected output (excerpts):


```

00000000 <__vectors>:
  0: XX XX XX XX  jmp  main
  4: ...
 28: XX XX XX XX  jmp  tca0_cmp0_isr    (TCA0_CMP0 vector at 0x0028)

<tca0_cmp0_isr>:
  push r16
  in   r16, 0x3f          ; SREG = 0x3f in I/O space
  push r16
  ldi  r16, 0x08          ; (1 << 3) = PA3
  sts  0x0407, r16        ; PORTA_OUTTGL
  ldi  r16, 0x10          ; TCA_SINGLE_CMP0_bm
  sts  0x0a0b, r16        ; TCA0_SINGLE_INTFLAGS
  pop  r16
  out  0x3f, r16          ; restore SREG
  pop  r16
  reti

```

Check that `tca0_cmp0_isr` is at the address that vector 0x0028 points to.

12.9 Extending: Normal Mode with OVF Interrupt

For cases where you want a fixed period and don't need compare-match output, use Normal mode with the OVF interrupt. Change two registers:

```

/* Normal mode – OVF fires when CNT reaches PER */
ldi  r16, TCA_SINGLE_WGMODE_NORMAL_gc    /* = 0x00 */
sts  TCA0_SINGLE_CTRLB, r16

ldi  r16, lo8(3254)
sts  TCA0_SINGLE_PERL, r16               /* PER (not CMP0) */
ldi  r16, hi8(3254)
sts  TCA0_SINGLE_PERH, r16

ldi  r16, TCA_SINGLE_OVF_bm              /* = 0x01: overflow interrupt */
sts  TCA0_SINGLE_INTCTRL, r16

```

The ISR changes slightly — clear OVF instead of CMP0:

```

tca0_ovf_isr:
  push r16
  in   r16, SREG
  push r16

  ldi  r16, (1 << LED_BIT)
  sts  PORTA_OUTTGL, r16

  ldi  r16, TCA_SINGLE_OVF_bm             /* = 0x01 */
  sts  TCA0_SINGLE_INTFLAGS, r16         /* clear OVF flag */

  pop  r16
  out  SREG, r16
  pop  r16
  reti

```

And the vector changes to entry 8 (0x0020 — TCA0_OVF):

```

jmp tca0_ovf_isr    /* 0x0020: TCA0_OVF */
jmp isr_default    /* 0x0024: TCA0_HUNF */
jmp isr_default    /* 0x0028: TCA0_CMP0 */

```

Both modes produce identical 1 Hz interrupt behavior. FRQ mode is marginally simpler (no PER register to configure separately).

That kind of tradeoff is common in timer work: several modes can solve the same problem, but one usually fits the mental model more directly.

12.10 Extending: Single-Slope PWM

To dim an LED with 50% duty cycle at ~16.3 kHz using PA0 (WO0 pin):

```

/* PER = 204 → f_pwm = 3333333 / (1 × 205) ≈ 16.26 kHz */
/* CMP0 = 102 → duty ≈ 102 / (204 + 1) ≈ 50% */

ldi r16, TCA_SINGLE_WGMODE_SINGLESLOPE_gc /* CTRLB: SS-PWM + CMP0EN */
ori r16, TCA_SINGLE_CMP0EN_bm
sts TCA0_SINGLE_CTRLB, r16

ldi r16, lo8(204)
sts TCA0_SINGLE_PERL, r16
ldi r16, hi8(204)
sts TCA0_SINGLE_PERH, r16

ldi r16, lo8(102)
sts TCA0_SINGLE_CMP0L, r16
ldi r16, hi8(102)
sts TCA0_SINGLE_CMP0H, r16

/* Start timer with prescaler 1 (no division) */
ldi r16, (TCA_SINGLE_CLKSEL_DIV1_gc | TCA_SINGLE_ENABLE_bm)
sts TCA0_SINGLE_CTRLA, r16

```

No interrupt needed for PWM — the hardware toggles WO0 automatically. The PA0 pin must be set as output (SBI VPORTA_DIR, 0) before enabling the timer output, or WO0 is disconnected from the pin.

This is an important milestone: once PWM is configured, the peripheral can keep producing the waveform even if the CPU goes off and does something unrelated.

To change duty cycle at runtime, write a new value to CMP0BUF (the buffer). The ALUPD bit in CTRLB causes the buffer to transfer to CMP0 on the next OVF, preventing glitches mid-cycle.

12.11 Summary

TCA0 modes:

- NORMAL (0x00): CNT → PER → reset; OVF interrupt.
- FRQ (0x01): CNT → CMP0 → reset; CMP0 interrupt; WO0 toggles (if enabled).
- SS-PWM (0x03): CNT → PER → reset; WO0 set at BOTTOM, cleared on CMP0; OVF int.

Frequency formula:

$$f = f_{\text{cpu}} / (\text{prescaler} \times (\text{TOP} + 1))$$

$$TOP = (f_{cpu} / (prescaler \times f)) - 1$$

16-bit register write order: LOW byte first, then HIGH byte (latch).

Initialization order:

1. CTRLB (mode) 2. CMP0/PER (TOP) 3. INTCTRL 4. CTRLA (start) 5. SEI

Interrupt flag clear: write TCA_SINGLE_CMP0_bm (or TCA_SINGLE_OVF_bm) to TCA0_SINGLE_INTFLAGS before RETI – or interrupt fires again immediately.

ISR discipline: identical to Chapter 10.

PUSH all used regs → IN/PUSH SREG → body → POP/OUT SREG → POP regs → RETI.

ATtiny3217 TCA0 vectors:

0x0020	TCA0_OVF	– overflow
0x0028	TCA0_CMP0	– compare 0 match (FRQ mode trigger)
0x002C	TCA0_CMP1	
0x0030	TCA0_CMP2	

12.12 Exercises

1. The example uses prescaler 1024 and CMP0 = 3254. Calculate the CMP0 value needed for a 2 Hz blink with the same prescaler. What is the exact frequency error as a percentage?
 2. Modify the example to blink the LED three times fast (150 ms on, 150 ms off) then pause for 700 ms. First reconfigure the timer for a 150 ms tick, then use a counter variable in SRAM to step through the pattern.
 3. Calculate PER and CMP0 values for a 1 kHz PWM signal on WO0 with 25% duty cycle, using prescaler DIV1 at 3.333 MHz. Write the complete initialization sequence in assembly.
 4. The ISR clears TCA0_SINGLE_INTFLAGS with an STS instruction. STS is a 4-byte, 2-cycle instruction. TCA0_SINGLE_INTFLAGS is at 0x0A0B — not in the VPORT I/O range (0x00–0x1F), so SBI cannot be used here. Confirm this by checking whether 0x0A0B mapped to any VPORT address. What is the smallest instruction count you can use to clear the flag?
 5. Add a second compare channel (CMP1) to the FRQ-mode example. CMP1 should fire at 2 Hz (while CMP0 still fires at 1 Hz). What constraint does FRQ mode place on CMP1's relationship to CMP0? (Hint: re-read the FRQ mode description — CNT resets at CMP0, not PER.)
 6. Using `avr-objdump -d timer_ctc.elf`, verify that the `sts TCA0_SINGLE_CMP0L` instruction is encoded as a 4-byte instruction. What are the two 16-bit words of its encoding? (Refer to the AVR instruction set manual for the STS 32-bit format.)
-

Next: Chapter 12 — USART (Serial Communication)

Chapter 13

USART (Serial Communication)

USART (Universal Synchronous/Asynchronous Receiver/Transmitter) is the simplest way to communicate between a microcontroller and a PC. One wire transmits (TX), one receives (RX), and data flows as a timed sequence of bits. No clock line is needed — both sides agree on a baud rate (bits per second) in advance. This chapter covers the ATtiny3217's USART0 peripheral: register setup, polling transmit, polling receive, ISR-driven receive, and a complete echo terminal example.

This is often the first peripheral that makes a microcontroller feel “interactive” because it gives you a direct text path to a terminal.

13.1 UART Framing

13.1.1 The Bit Stream

USART uses NRZ (Non-Return-to-Zero) encoding: the line is high (idle) when no data is being sent. Each byte is framed with:

Idle: _____
Start: _____ | ← start bit (line goes low, 1 bit period)
 | D0 – D1 – D2 – D3 – D4 – D5 – D6 – D7 |
Stop: _____ (line high)

Timeline at 9600 baud (1 bit = 104.2 μ s):

1 bit	1 bit	... 8 bits ...	1 stop	idle ...
start	D0 LSB		D7 MSB	

One character transmission = 1 start + 8 data + 1 stop = **10 bit periods**. At 9600 baud: $10 \times 104.2 \mu\text{s} = 1.042 \text{ ms}$ per byte → maximum ~960 bytes/second.

That simple 10-bit framing cost is worth remembering. Even before software overhead, UART throughput is lower than the raw baud number suggests.

13.1.2 8N1 Format

8N1 means: - **8** data bits (LSB first) - **N** — no parity bit - **1** stop bit

This is the most common UART configuration and the ATtiny3217 USART0 default after reset.

So for ordinary 8-bit serial text, the default framing is already the framing you want.

13.2 USART0 Registers on ATtiny3217

Register	Address	Purpose
USART0_RXDATA_L	0x0800	Receive data low byte (read = receive byte)
USART0_RXDATA_H	0x0801	Receive data high byte (error flags + 9th bit)
USART0_TXDATA_L	0x0802	Transmit data (write = queue byte to send)
USART0_TXDATA_H	0x0803	Transmit data high byte (9th bit when used)
USART0_STATUS	0x0804	Status flags (RXCIF, TXCIF, DREIF, etc.)
USART0_CTRLA	0x0805	Interrupt enables, RS-485 mode
USART0_CTRLB	0x0806	TX/RX enable, start-frame detect, RX mode
USART0_CTRLC	0x0807	Frame format: baud mode, char size, parity, stop bits
USART0_BAUD_L	0x0808	Baud rate register low byte
USART0_BAUD_H	0x0809	Baud rate register high byte
USART0_CTRLD	0x080A	Auto-baud width field
USART0_DBGCTRL	0x080B	Debug run override
USART0_EVCTRL	0x080C	IrDA event input enable
USART0_TXPLCTRL	0x080D	IrDA transmitter pulse length
USART0_RXPLCTRL	0x080E	IrDA receiver pulse length

All are outside the 0x00-0x3F I/O range reached by IN/OUT, so use LDS/STS.

That means USART code on ATtiny3217 naturally uses the “full peripheral register” access style rather than the short IN/OUT style.

13.2.1 STATUS — Status Register (0x0804)

Bit:	7	6	5	4	3	2	1	0
	RXCIF	TXCIF	DREIF	RXSIF	ISFIF	—	BDF	WFB

- RXCIF — RX Complete Interrupt Flag: 1 when unread byte in RXDATA_L.
Cleared automatically by reading RXDATA_L.
- TXCIF — TX Complete Interrupt Flag: 1 when last byte shifted out fully.
Cleared by writing 1 (W1C).
- DREIF — Data Register Empty Flag: 1 when TXDATA_L is ready for new byte.
Cleared automatically when TXDATA_L is written.
On reset, DREIF = 1 (TX is ready immediately).
- RXSIF — Receive Start Interrupt Flag, used with start-frame detection.
- ISFIF — Inconsistent Sync Field Interrupt Flag, used with auto-baud/LIN modes.
- BDF — Break Detected Flag, used with break/sync detection.
- WFB — Wait For Break control/status bit.

Received-byte error flags are not in STATUS. They are stored with the received byte in RXDATA_H: BUFOVF at bit 6, FERR at bit 2, PERR at bit 1, and DATA[8] at bit 0 for 9-bit frames. Read RXDATA_H before RXDATA_L when you need those error bits.

For polling TX: wait until DREIF = 1, then write TXDATA_L. For polling RX: wait until RXCIF = 1, then read RXDATA_L. For ISR-driven RX: enable RXCIE in CTRLA; ISR fires when RXCIF sets.

Those three sentences are the whole USART mental model in miniature: one flag for “ready to transmit,” one flag for “data received,” and one interrupt-enable bit if you want hardware to notify you instead of polling.

13.2.2 CTRLB — Control B (0x0806)

Bit:	7	6	5	4	3	2:1	0
	RXEN	TXEN	—	SFDEN	ODME	RXMODE	MPCM

RXEN — 1 = receive enabled (USART monitors RXD pin)

TXEN – 1 = transmit enabled (USART drives TXD pin)
 SFDEN – 1 = start-frame detection enabled
 ODME – 1 = open-drain mode enabled
 RXMODE – 0 = normal (16x oversampling), 1 = CLK2X (8x), 2 = LINAUTO, 3 = GENAUTO
 MPCM – 1 = multiprocessor communication mode

Bit constants (from avr/io.h):

USART_TXEN_bm = 0x40
 USART_RXEN_bm = 0x80
 USART_SFDEN_bm = 0x10
 USART_ODME_bm = 0x08
 USART_RXCIE_bm is in CTRLA (see below)

13.2.3 CTRLA — Control A / Interrupt Enables (0x0805)

Bit:	7	6	5	4	3	2	1:0
	RXCIE	TXCIE	DREIE	RXSIE	LBME	ABEIE	RS485

RXCIE – RX Complete Interrupt Enable: fires ISR when RXCIF = 1
 TXCIE – TX Complete Interrupt Enable
 DREIE – Data Register Empty Interrupt Enable
 RXSIE – Receive Start Interrupt Enable
 LBME – Loop-back Mode Enable
 ABEIE – Auto-baud Error Interrupt Enable

Bit constants:

USART_RXCIE_bm = 0x80
 USART_TXCIE_bm = 0x40
 USART_DREIE_bm = 0x20

13.2.4 CTRLC — Frame Format (0x0807)

Bit:	7:6	5:4	3	2:0
	CMODE	PMODE	SBMODE	CHSIZE

CMODE (bits 7:6): communication mode

0x00 = USART_CMODE_ASYNCHRONOUS_gc – async (normal UART)
 0x01 = USART_CMODE_SYNCHRONOUS_gc – sync
 0x02 = USART_CMODE_IRCOM_gc – IrDA
 0x03 = USART_CMODE_MSPI_gc – SPI master

SBMODE (bit 3): stop bits

0 = 1 stop bit
 1 = 2 stop bits

PMODE (bits 5:4): parity

field 0x0 = disabled
 field 0x2 = even
 field 0x3 = odd

CHSIZE (bits 2:0):

0x03 = 8-bit (USART_CHSIZE_8BIT_gc) – most common

Default after reset = 0x03 = async, 8-bit, no parity, 1 stop = 8N1.
 No write needed if 8N1 is desired.

13.3 Baud Rate Calculation

13.3.1 The Formula

USART0 on ATtiny3217 uses a 16-bit baud register (BAUDL + BAUDH) with a fractional format. For **asynchronous normal mode** (16× oversampling):

$$\begin{aligned}\text{BAUD_REG} &= (\text{f_cpu} \times 64) / (16 \times \text{baud_rate}) \\ &= (\text{f_cpu} \times 4) / \text{baud_rate}\end{aligned}$$

The × 64 and ÷ 16 come from the internal fractional divider format — the top 10 bits of BAUD_REG are the integer part, and the bottom 6 bits are the fractional part.

You do not need to memorize the internal fractional details to use the peripheral well, but they explain why the formula looks stranger than a simple clock-divider equation.

For 9600 baud at 3.333 MHz:

$$\begin{aligned}\text{BAUD_REG} &= (3333333 \times 64) / (16 \times 9600) \\ &= 21333312 / 153600 \\ &= 1388.88875 \\ &\rightarrow \text{round to } 1389 = 0x056D\end{aligned}$$

$$\begin{aligned}\text{Actual baud rate} &= (3333333 \times 64) / (16 \times 1389) \\ &= 21333312 / 22224 \\ &= 9599.23 \text{ bps} \quad (-0.008\% \text{ error} - \text{negligible})\end{aligned}$$

Common baud rates at 3.333 MHz:

Baud	BAUD_REG (decimal)	BAUD_REG (hex)	Error
9600	1389	0x056D	-0.01%
19200	694	0x02B6	0.06%
38400	347	0x015B	0.06%
57600	231	0x00E7	0.21%
115200	116	0x0074	-0.22%

Writing BAUD (low byte first):

```
ldi r16, lo8(1389) /* = 0x6D */
sts USART0_BAUDL, r16
ldi r16, hi8(1389) /* = 0x05 */
sts USART0_BAUDH, r16
```

As with other 16-bit peripheral values on AVR, the write order is part of the hardware interface, not just style.

13.4 TX/RX Pins on ATtiny3217

USART0 uses **PORTB** by default:

Function	Pin	Direction	Notes
TXD	PB2	Output	USART overrides the pin when TXEN=1
RXD	PB3	Input	Leave as input (default); no pull-up needed

Microchip documents that the transmitter overrides normal port operation for the TXD pin when TXEN=1. The example still sets PB2 as an output before enabling the transmitter so the idle level is well-defined during setup.

```
sbi VPORTB_DIR, 2 /* PB2 = output for TXD */
/* PB3 (RXD) = input by default – no configuration needed */
```

Alternative pin mapping (PORTMUX_USARTROUTEA register) allows moving USART0 to other pins, but this chapter uses the default mapping.

13.5 Polling Transmit

The simplest TX approach: spin until DREIF = 1, then write the byte.

Polling transmit is usually acceptable because the program already knows when it wants to send something.

AVR Instruction: LDS Rd, A – load from SRAM/peripheral
SBRS Rd, b – skip next if bit b of Rd is set

```
/*
 * usart_tx_byte – transmit r16 via polling
 * Clobbers: r17
 */
usart_tx_byte:
    push    r17
.Lwait_dreif:
    lds     r17, USART0_STATUS
    sbrs    r17, USART_DREIF_bp    /* bit 5 – skip if DREIF = 1          */
    rjmp    .Lwait_dreif           /* DREIF = 0: TX busy, loop        */
    sts     USART0_TXDATA, r16     /* write byte: clears DREIF, starts TX */
    pop     r17
    ret
```

USART_DREIF_bp is the bit position (5) of DREIF in the STATUS register. SBRS r17, 5 skips the RJMP when bit 5 is 1 (DREIF set = TX ready).

So the loop is doing exactly one question over and over: “is the transmit data register empty yet?”

13.5.0.1 Why Not SBIS?

SBIS and SBIC only work with I/O space registers (addresses 0x00–0x1F in I/O space, or 0x20–0x3F in data space). USART0_STATUS is at 0x0804 — far outside I/O space. Use LDS + SBRS for peripheral registers.

13.5.0.2 Cycle Cost of Polling

Each loop iteration: LDS (3 cy) + SBRS (1 cy, not skipping) + RJMP (2 cy) = 6 cycles. At 3.333 MHz, 6 cycles = 1.8 µs per iteration. At 9600 baud, one bit period is 104 µs — the polling loop runs ~58 times per bit period. Acceptable for simple programs; for power-sensitive applications, use DREIE interrupt instead.

That is the tradeoff in one paragraph: polling is simple, but it spends CPU time checking a condition that usually is not ready yet.

13.5.1 Transmitting a String from Flash

```
/*
 * usart_tx_string – send bytes from flash
 * ZL:ZH = pointer to first byte, r18 = length
 * Clobbers: r16, r17, r18, Z
 */
usart_tx_string:
    tst     r18
```



```

    breq  .Ltx_done
.Ltx_loop:
    lpm   r16, Z+           /* load byte from flash, post-increment Z */
    rcall usart_tx_byte
    dec  r18
    brne .Ltx_loop
.Ltx_done:
    ret

```

String data in flash:

```

.Lready_str:
    .byte 'R', 'e', 'a', 'd', 'y', '\r', '\n'
.equ  READY_LEN, 7
.balign 2

```

Calling it:

```

ldi  ZL, lo8(.Lready_str)
ldi  ZH, hi8(.Lready_str)
ldi  r18, READY_LEN
rcall usart_tx_string

```

The `.byte` directive places raw bytes in the `.text` section at the label address. `LPM Z+` loads each byte from flash and advances `Z`. `READY_LEN` is a constant counting the bytes (7 = 5 chars + CR + LF). The `.balign 2` keeps the next instruction word-aligned after the 7-byte string.

This is a nice place where Chapter 5's flash-reading model becomes immediately useful in application code.

13.6 Polling Receive

Wait until `RXCIF = 1`, then read `RXDATAL`. Reading clears the flag.

Polling receive is conceptually symmetric with polling transmit, but it is much more wasteful in practice because the CPU does not know when the next incoming byte will arrive.

```

/*
 * usart_rx_byte – receive one byte into r16 (polling)
 * Clobbers: r17
 */
usart_rx_byte:
    push r17
.Lwait_rxcif:
    lds  r17, USART0_STATUS
    sbrs r17, USART_RXCIF_bp /* bit 7 – skip if RXCIF = 1 (byte ready) */
    rjmp .Lwait_rxcif
    lds  r16, USART0_RXDATAL /* read byte – clears RXCIF */
    pop  r17
    ret

```

`USART_RXCIF_bp = 7` (bit position). `SBRs r17, 7` skips the loop branch when the RX byte is ready.

13.6.1 Error Checking

Before reading `RXDATAL`, the error flags in `RXDATH` should be checked:

```

.Lwait_rxcif:
    lds    r17, USART0_STATUS
    sbrs   r17, USART_RXCIF_bp
    rjmp   .Lwait_rxcif
    lds    r17, USART0_RXDATAH      /* read high byte first for error flags */
    andi   r17, (USART_FERR_bm | USART_BUFOVF_bm | USART_PERR_bm)
    brne   .Lrx_error              /* jump to error handler if any flag set */
    lds    r16, USART0_RXDATAL      /* no error – read data */

```

RXDATAH contains the error flags (FERR, BUFOVF, PERR) and optionally the 9th data bit. For 8N1, only the error bits matter. RXDATAH must be read before RXDATAL — reading RXDATAL consumes the byte and discards the error flags.

That read order is another hardware contract: get the status context first, then consume the byte.

For the simple echo example, error checking is omitted for clarity.

13.7 ISR-Driven Receive

Polling TX is fine for transmit — the CPU controls when it sends. But polling RX is wasteful: the CPU spins doing nothing until a byte arrives, unable to perform other work. The RXCIE interrupt solves this: the ISR fires only when a byte is ready, and main code runs uninterrupted between bytes.

This is one of the clearest examples of when interrupts are worth the added complexity.

13.7.1 USART0_RXC Vector

From the ATtiny3217 vector table (Chapter 10):

Byte addr	Vector	Peripheral	Trigger
0x006C	27	USART0_RXC	RXCIF = 1 (byte received)
0x0070	28	USART0_DRE	DREIF = 1 (TX data register empty)
0x0074	29	USART0_TXC	TXCIF = 1 (TX shift register empty)

The USART0_RXC vector is at byte address 0x006C. In the vector table, it is the 28th entry (counting RESET as entry 0):

```

__vectors:
    jmp    main          /* 0x0000: entry 0  RESET */
    ...                /* entries 1–26 */
    jmp    usart0_rxc_isr /* 0x006C: entry 27 USART0_RXC */
    jmp    isr_default   /* 0x0070: entry 28 USART0_DRE */
    jmp    isr_default   /* 0x0074: entry 29 USART0_TXC */
    jmp    isr_default   /* 0x0078: entry 30 NVMCTRL_EE */

```

13.7.2 RXC ISR

Reading RXDATAL inside the ISR **automatically clears RXCIF**. No explicit flag clear is needed (unlike TCA0 INTFLAGS which requires W1C).

That makes USART receive interrupts a little cleaner than many timer or GPIO interrupt sources.

```

usart0_rxc_isr:
    push   r16

```

```

push    r17
in      r16, SREG
push    r16

lds     r16, USART0_RXDATA    /* read byte – clears RXCIF automatically */
rcall   usart_tx_byte         /* echo it back (polling TX) */

pop     r16
out     SREG, r16
pop     r17
pop     r16
reti

```

Note r17 is pushed/popped because `usart_tx_byte` clobbers it. The ISR calls a subroutine — this is legal but requires that the callee’s clobbers are also saved in the prologue.

The important rule is local and mechanical: if an ISR calls a helper, the ISR must also preserve any registers the helper uses.

13.7.2.1 Calling Subroutines from ISRs

Calling `RCALL` inside an ISR pushes another 2 bytes onto the stack. The total stack depth for this ISR:

On ISR entry (CPU pushes return addr):	2 bytes
PUSH r16:	1 byte
PUSH r17:	1 byte
PUSH r16 (SREG):	1 byte
RCALL <code>usart_tx_byte</code> (CPU pushes addr):	2 bytes
PUSH r17 (in <code>usart_tx_byte</code>):	1 byte

Total stack depth: 8 bytes

On ATtiny3217, SRAM is 2 KB. Typical use of a few hundred bytes of stack is safe, but deep call trees inside ISRs require careful accounting.

So ISR design is not just about correctness and latency; it is also about bounded stack growth.

13.7.2.2 Blocking TX Inside an ISR

The ISR calls `usart_tx_byte`, which polls `DREIF`. While polling, the ISR is still active, so another `USART0_RXC` interrupt will not run until `RETI`. This is safe here because at 9600 baud the TX wait is at most one byte period (~1.042 ms) and no other time-critical ISR is competing. If more bytes arrive while the ISR is blocked, they wait in the receive path; if software does not drain them in time, `BUFOVF` is reported with the received frame. In higher-throughput designs, use a circular buffer: the `RXC` ISR puts the byte in SRAM, the `DREIE` interrupt sends from the buffer, and the `RXC` ISR returns immediately.

That distinction is important: a simple echo ISR is fine for a teaching example, but blocking inside an ISR does not scale well.

13.8 Initialization Sequence

```

main:
    /* Stack */
    ldi    r16, hi8(RAMEND)

```

```

out    SPH, r16
ldi    r16, lo8(RAMEND)
out    SPL, r16

/* Set TXD idle high before TXEN overrides the pin */
sbi    VPORTB_OUT, TXD_BIT      /* PB2 */
sbi    VPORTB_DIR, TXD_BIT      /* PB2 */

/* Baud rate: 9600 at 3.333 MHz → BAUD_REG = 1389 = 0x056D */
ldi    r16, BAUD_LO             /* 0x6D */
sts    USART0_BAUDL, r16
ldi    r16, BAUD_HI             /* 0x05 */
sts    USART0_BAUDH, r16

/* CTRLC defaults = 8N1 async – no write needed */

/* Enable TX + RX */
ldi    r16, (USART_TXEN_bm | USART_RXEN_bm)
sts    USART0_CTRLB, r16        /* TXEN + RXEN */
/* RXCIE lives in CTRLA – separate write */
ldi    r16, USART_RXCIE_bm
sts    USART0_CTRLA, r16

sei

/* Send startup message */
ldi    ZL, lo8(.Lready_str)
ldi    ZH, hi8(.Lready_str)
ldi    r18, READY_LEN
rcall  usart_tx_string

.Lidle:
rjmp  .Lidle

```

Note on CTRLB vs CTRLA: TXEN and RXEN are in CTRLB (0x0806). RXCIE is in CTRLA (0x0805). They are separate registers — two STS writes.

This is exactly the sort of small register-detail bug that can waste a lot of debug time if you assume all enable bits live together.

13.9 Worked Example: Echo Terminal

Full source: `src/usart_echo.S`

The program transmits `Ready\r\n` on startup, then echoes every received byte back to the sender. No main-loop processing — all receive logic is in the ISR.

That makes the data path very easy to reason about: receive a byte, trigger the ISR, send the same byte back.

13.9.1 Data Flow

```

PC terminal → USB-UART adapter → RXD (PB3) → USART0_RXDATA
RXCIF = 1 → USART0_RXC ISR fires
              reads RXDATA (clears RXCIF)
              calls usart_tx_byte(r16)
              polls DREIF → writes TXDATA

```

TXD (PB2) → USB-UART adapter → PC terminal (sees same char echoed back)

13.9.2 Build and Test

```
# Build
avr-gcc -mmcu=attiny3217 -x assembler-with-cpp -c -o usart_echo.o src/usart_echo.S
avr-gcc -mmcu=attiny3217 -nostartfiles -o usart_echo.elf usart_echo.o
avr-objcopy -O ihex usart_echo.elf usart_echo.hex

# Flash
avrdude -c serialupdi -p attiny3217 -P /dev/ttyUSB0 -U flash:w:usart_echo.hex

# Connect serial terminal at 9600 8N1
# Linux/macOS:
screen /dev/ttyUSB0 9600
# or:
minicom -b 9600 -D /dev/ttyUSB0
```

13.9.3 Disassembly Check

```
avr-objdump -d usart_echo.elf
```

Key things to verify: - Vector at 0x006C (jmp usart0_rxc_isr) — counts 27 entries from RESET. - sts 0x0808, r16 for BAUDL, sts 0x0809, r16 for BAUDH. - sts 0x0806, r16 for CTRLB, sts 0x0805, r16 for CTRLA. - lds r16, 0x0800 for RXDATAL inside the ISR. - lds r17, 0x0804 for STATUS in the TX polling loop.

13.10 Circular Buffer for High-Throughput RX

The simple ISR copies one byte directly. At 115200 baud with rapid bursts, the CPU may not finish the TX echo before the next byte arrives — BUFOVF sets and bytes are lost.

A circular buffer decouples RX from TX:

SRAM layout:

```
rx_buf:    .space 16      ; 16-byte circular buffer
rx_head:    .space 1      ; write index (ISR writes here)
rx_tail:    .space 1      ; read index (main reads here)
```

The ISR writes to rx_buf[rx_head] and increments rx_head (modulo 16). Main code checks rx_head != rx_tail to know a byte is ready, reads from rx_buf[rx_tail], and increments rx_tail.

This is the standard way to decouple “data arrived now” from “main code is ready to process it now.”

```
/* ISR stores byte in circular buffer */
usart0_rxc_isr:
    push    r16
    push    r17
    push    ZL
    push    ZH
    in      r16, SREG
    push    r16

    lds     r16, USART0_RXDATAL    /* read byte, clear RXCIF */
```

```

lds    r17, rx_head          /* load write index          */
ldi    ZL, lo8(rx_buf)
ldi    ZH, hi8(rx_buf)
add    ZL, r17               /* Z = &rx_buf[rx_head]    */
adc    ZH, __zero_reg__
st     Z, r16                /* store byte               */

inc    r17
andi   r17, 0x0F             /* modulo 16 (power-of-2 wrap) */
sts    rx_head, r17

pop    r16
out    SREG, r16
pop    ZH
pop    ZL
pop    r17
pop    r16
reti

```

Full circular buffer implementation is left as Exercise 5.

13.11 Summary

USART0 pins (default): TXD = PB2 (output), RXD = PB3 (input)

Baud rate register:

BAUD_REG = (f_cpu × 64) / (16 × baud) [async normal mode]

At 3.333 MHz, 9600 bps → BAUD_REG = 1389 = 0x056D

Write: BAUDL first (low byte), then BAUDH.

Registers to configure:

USART0_BAUDL / BAUDH – baud rate (write low byte first)
 USART0_CTRLC – frame format (default = 8N1, no write needed)
 USART0_CTRLB – TXEN_bm (0x40) | RXEN_bm (0x80)
 USART0_CTRLA – RXCIE_bm (0x80) for ISR-driven RX

STATUS flags:

DREIF (bit 5) – TX data register empty: 1 = ready for new byte

RXCIF (bit 7) – RX complete: 1 = byte ready in RXDATA1

Both checked with: LDS Rd, USART0_STATUS then SBRS/SBRC

Polling TX:

wait DREIF = 1 → STS USART0_TXDATA1, Rr → DREIF auto-clears

Polling RX:

wait RXCIF = 1 → LDS Rd, USART0_RXDATA1 → RXCIF auto-clears

ISR-driven RX (USART0_RXC, vector 27, address 0x006C):

Enable: USART_RXCIE_bm in USART0_CTRLA

ISR reads RXDATA1 → clears RXCIF automatically

No manual flag clear needed (unlike TCA0 INTFLAGS)

ISR rules: same as always.

Save r16+SREG minimum. Save any register clobbered by callees.

Stack depth = 2 (ret addr) + saved regs + callee frames.

13.12 Exercises

1. Calculate BAUD_REG for 115200 baud at 3.333 MHz. What is the percentage error? If the error exceeds 2%, the link is likely to be unreliable — is 115200 feasible at this clock speed?
2. The polling TX function checks DREIF each iteration. At 9600 baud and 3.333 MHz, how many poll iterations occur while waiting for a byte to finish transmitting? (Use the LDS+SBRS+RJMP cycle count.)
3. The ISR calls `usart_tx_byte`, which polls DREIF. During this polling, could a second USART0_RXC interrupt fire? Explain why or why not. What would happen to the second received byte?
4. Modify `usart_echo.S` to echo each byte with its hex representation. Receiving A (0x41) should transmit the string 41\r\n. Write the `byte_to_hex` subroutine that converts r16 (0x00–0xFF) into two ASCII hex digits in r17 (high nibble) and r16 (low nibble).
5. Implement the 16-byte circular buffer described in Section 12.10. Write:
 - the SRAM variable declarations (.bss section)
 - the full ISR using the buffer
 - a `main_loop` that calls a `process_byte` subroutine for each byte in the buffer (use `rx_head != rx_tail` as the ready check)
6. Using `avr-objdump -d usart_echo.elf`, locate the LDS r17, USART0_STATUS instruction in `usart_tx_byte`. What is its 32-bit encoding? What address (0x0804) appears in which bytes of the instruction?

Next: Chapter 13 — SPI & TWI (I2C) Overview

Chapter 14

SPI & TWI (I2C) Overview

Serial communication protocols let microcontrollers talk to sensors, memory chips, displays, and other devices using only a handful of wires. This chapter covers the two most common synchronous serial buses used in embedded systems: SPI (Serial Peripheral Interface) and TWI (Two-Wire Interface, Microchip's name for I²C). Both are present on the ATtiny3217 as hardware peripherals; both are accessed through memory-mapped registers using LDS/STS, exactly as with USART0 in Chapter 12.

This chapter is less about talking to a terminal and more about talking to other chips.

14.1 SPI Fundamentals

SPI is a full-duplex, synchronous serial bus. Four wires connect one master to one or more slave devices:

Master (ATtiny3217)		Slave (25LC010A EEPROM)
MOSI (PA1)	→	SI (Master Out Slave In)
MISO (PA2)	←	SO (Master In Slave Out)
SCK (PA3)	→	SCK (clock, generated by master)
CS (PA4)	→	C $\bar{S}$ (chip select, active-low)

Full duplex: every clock pulse simultaneously shifts one bit out on MOSI and samples one bit in on MISO. An 8-bit transfer is 8 clock pulses. If the slave has nothing meaningful to return (a pure command), the byte received by the master is discarded (or a dummy byte 0xFF is sent to clock the slave's response out).

That full-duplex behavior is one of the first things that makes SPI feel different from UART. You are always sending and receiving at the same time, even if one direction is only carrying dummy data.

Chip select: SPI has no built-in addressing. Instead, the master drives CS low to select a particular slave before the transfer begins, then drives CS high after the transfer ends. Each slave needs its own CS line.

So SPI trades protocol simplicity for pin count: the bus itself is simple, but every slave needs a dedicated select signal.

14.1.1 SPI Clock Polarity and Phase (Mode)

The idle state of SCK and the edge on which data is sampled define four modes:

Mode	CPOL	CPHA	SCK idle	Data captured on
------	------	------	----------	------------------


0	0	0	low	rising edge	← 25LC010A compatible
1	0	1	low	falling edge	
2	1	0	high	falling edge	
3	1	1	high	rising edge	← also 25LC010A compatible

The 25LC010A EEPROM used in this chapter works in Mode 0 or Mode 3. We use Mode 0 (CPOL=0, CPHA=0) — SCK idles low, data captured on the rising edge.

In practice, “which SPI mode?” is one of the first things you check in a datasheet when a device does not respond correctly.

On ATtiny3217, the default SPI0 pins are PA1=MOSI, PA2=MISO, PA3=SCK, and PA4=SS. The datasheet also lists alternate SPI0 positions on PORTC: PC2=MOSI, PC1=MISO, PC0=SCK, and PC3=SS, selected through PORTMUX. This chapter keeps the default PA pins and uses PA4 as a software-controlled chip select.

14.1.2 Timing Diagram (Mode 0)



Data is set up before the rising edge and sampled on the rising edge.

14.2 SPI0 Registers on the ATtiny3217

The ATtiny3217 SPI0 peripheral uses five memory-mapped registers:

Register	Address	Purpose
SPI0_CTRLA	0x0820	Master/slave, clock, data order, enable
SPI0_CTRLB	0x0821	Buffer mode, slave-select, SPI mode
SPI0_INTCTRL	0x0822	Interrupt enables
SPI0_INTFLAGS	0x0823	Transfer-complete flag (IF), write-collision (WRCOL)
SPI0_DATA	0x0824	Transmit / receive data register

14.2.1 SPI0_CTRLA — Control A

Bit:	7	6	5	4	3	2	1	0
	—	DORD	MASTER	CLK2X	—	PRESC[1:0]	ENABLE	

Field	Bits	Description
DORD	6	Data order: 0 = MSB first (normal), 1 = LSB first
MASTER	5	1 = master mode, 0 = slave mode
CLK2X	4	Double the clock speed (halves effective prescaler)
PRESC	2:1	Clock prescaler: 00=/4, 01=/16, 10=/64, 11=/128
ENABLE	0	1 = enable SPI peripheral

At 3.333 MHz with PRESC=DIV4 and CLK2X=0: SCK = 3.333 MHz / 4 $\approx$ **833 kHz**. The 25LC010A maximum SCK is 10 MHz, so the default chapter setting is well inside the EEPROM limit. On this CPU clock, SPI0 can run as fast as $f_{\text{CPU}}/2 \approx 1.667$ MHz by setting CLK2X.

14.2.2 SPI0_CTRLB — Control B

Bit: 7 6 5 4 3 2 1 0
 BUFEN BUFWR – – – SSD MODE[1:0]

Field	Bits	Description
BUFEN	7	Enable buffer mode (more complex; not used here)
BUFWR	6	Buffer write mode enable (buffer mode only)
SSD	2	Slave Select Disable — set to 1 when driving CS manually
MODE	1:0	SPI mode: 00=Mode0, 01=Mode1, 10=Mode2, 11=Mode3

SSD must be set when driving CS as a GPIO (as we do in this chapter). If SSD=0 and the hardware SS pin (PA4) is driven low by external logic, the ATtiny3217 will switch itself to slave mode — a common source of confusion.

This is one of those details that feels obscure until it breaks the whole bus.

14.2.3 SPI0_INTFLAGS — Interrupt / Status Flags (non-buffered mode)

Bit: 7 6 5 4 3 2 1 0
 IF WRCOL – – – – – –

Flag	Bit	Set when	Cleared by
IF	7	8-bit transfer complete	Read INTFLAGS, then access DATA
WRCOL6		Write collision (SPI0_DATA written mid-transfer)	Read INTFLAGS, then access DATA

IF at bit position SPI_IF_bp (= 7). Poll with SBR5.

14.2.4 SPI0_DATA — Transmit / Receive

Reading SPI0_DATA after polling IF returns the received byte and completes the flag-clear sequence, because the poll already read SPI0_INTFLAGS. Writing SPI0_DATA (when IF is set or before the first transfer) queues the byte to be shifted out. Both transmit and receive share the same register address — write to transmit, read to receive.

14.3 SPI Initialisation (Master Mode)

Before any transfer, configure the three output pins, leave MISO as input, set SSD, set MODE, then enable as master.

That order is deliberate: get the pins and mode right first, then turn the peripheral on.

```
/* Pin directions: MOSI (PA1), SCK (PA3), CS (PA4) → output */
ldi  r16, (1 << 1) | (1 << 3) | (1 << 4)
out  VPORTA_DIR, r16
```

```

/* CS high (deselect) */
sbi  VPORTA_OUT, 4

/* SPI0_CTRLB: SSD=1 (manual CS), MODE=0 (SPI mode 0) */
ldi  r16, SPI_SSD_bm
sts  SPI0_CTRLB, r16

/* SPI0_CTLRA: master, DIV4 prescaler, MSB first, enable */
ldi  r16, (SPI_MASTER_bm | SPI_ENABLE_bm)
sts  SPI0_CTLRA, r16

```

SPI_MASTER_bm = (1 << SPI_MASTER_bp) = 0x20. SPI_ENABLE_bm = (1 << SPI_ENABLE_bp) = 0x01. Both are defined in <avr/io.h> for the ATtiny3217.

14.4 Transferring a Byte

The hardware handles the 8-bit shift; the CPU's job is to: 1. Write the byte to SPI0_DATA (starts the transfer). 2. Poll SPI0_INTFLAGS until IF = 1 (transfer done). 3. Read SPI0_DATA (clears IF; returns received byte).

Once you see SPI transfers this way, most device protocols reduce to “send a command byte, maybe send an address byte, maybe send or receive data bytes.”

```

/*
 * spi_transfer – exchange one byte over SPI
 * In:  r16 = byte to send
 * Out: r16 = byte received
 * Clobbers: r17
 */
spi_transfer:
    sts  SPI0_DATA, r16          /* write tx byte → starts 8-bit transfer */
.lwait_if:
    lds  r17, SPI0_INTFLAGS
    sbrc r17, SPI_IF_bp          /* spin until IF (bit 7) = 1 */
    rjmp .lwait_if
    lds  r16, SPI0_DATA          /* read received byte – clears IF */
    ret

```

Cycle count at 3.333 MHz with SCK = f/4:

Item	Duration
8 SPI bits	8 × 4 CPU cycles = 32 cycles on SCK
Missed poll	6 cycles (LDS 3 + SBRS 1 + R JMP 2)
Final poll	5 cycles (LDS 3 + taken SBRS skip 2)
DATA read	3 cycles (LDS r16, SPI0_DATA)

The exact number of poll iterations depends on where the CPU is inside the loop when IF becomes set. With this code and f/4 SCK, DATA is typically read about 40 CPU cycles after the transfer starts.

14.5 The 25LC010A SPI EEPROM

The 25LC010A is a 1-Kbit (128×8 -bit) SPI EEPROM. It communicates using single-byte commands followed by optional address and data bytes.

14.5.1 Key Characteristics

Parameter	Value
Capacity	128 bytes (1 Kbit)
Address width	7 bits (0x00 – 0x7F)
Page size	16 bytes (for page write)
SPI modes	Mode 0 and Mode 3
Max SCK	10 MHz
Write cycle time	5 ms maximum
Supply voltage	1.8 V – 5.5 V

14.5.2 Command Set

Opcode	Hex	Bytes after opcode	Description
READ	0x03	addr, [data...]	Read one or more bytes
WRITE	0x02	addr, data[...]	Write one byte (or page)
WREN	0x06	–	Set Write Enable Latch (WEL)
WRDI	0x04	–	Clear Write Enable Latch
RDSR	0x05	[dummy]	Read Status Register
WRSR	0x01	data	Write Status Register

Write Enable Latch (WEL): the EEPROM resets WEL after every write. You must send WREN (its own CS-framed transaction) before every WRITE command.

That makes accidental writes less likely, but it also means your software has to treat “enable writing” as part of every write sequence, not as a one-time setup step.

14.5.3 Status Register (returned by RDSR)

Bit:	7	6	5	4	3	2	1	0
	–	–	–	–	BP1	BP0	WEL	WIP

Flag	Bit	Meaning
WIP	0	Write In Progress — 1 while write cycle is running
WEL	1	Write Enable Latch — 1 after WREN, 0 after write completes
BP	3:2	Block-protect bits (control write-protect regions)

Poll WIP = 0 before issuing another WRITE command.

In other words, the WRITE command only starts the EEPROM’s internal work. The chip still needs time afterwards to commit the byte into nonvolatile storage.

14.5.4 CS Framing Rules

Each command must be bracketed by CS low/high transitions. Raising CS before a complete transfer aborts the current command. The typical patterns:

This is the other half of SPI correctness besides mode. Many slave devices care not just about the bytes, but about the exact CS framing around them.

```

Write Enable:          CS↓ → WREN(0x06) → CS↑
Write byte:           CS↓ → WRITE(0x02) → ADDR → DATA → CS↑
Read status (poll WIP): CS↓ → RDSR(0x05) → DUMMY → CS↑
                      read status byte from DUMMY transfer

```

14.6 Example: Write Byte to SPI EEPROM

The full annotated source is at `ch13_spi_twi/src/spi_eeprom.S`. Here we walk through the key subroutines.

14.6.1 eeprom_wren — Send Write Enable

```

eeprom_wren:
    rcall cs_assert          /* CS low */
    ldi    r16, EEPROM_WREN /* 0x06 */
    rcall spi_transfer       /* shift out opcode (received byte discarded) */
    rcall cs_deassert        /* CS high — WEL is now set */
    ret

```

14.6.2 eeprom_write_byte — Write One Byte

```

/*
 * r24 = 8-bit address, r22 = data byte
 */
eeprom_write_byte:
    rcall eeprom_wren        /* WEL must be set before each write */

    rcall cs_assert
    ldi    r16, EEPROM_WRITE /* 0x02 */
    rcall spi_transfer       /* opcode */
    mov    r16, r24          /* address */
    rcall spi_transfer
    mov    r16, r22          /* data */
    rcall spi_transfer       /* CS↑ latches the write; 5 ms cycle starts */
    rcall cs_deassert
    ret

```

14.6.3 eeprom_wait_wip — Poll WIP Until Write Done

```

eeprom_wait_wip:
.Lwip_poll:
    rcall cs_assert
    ldi    r16, EEPROM_RDSR /* 0x05 */
    rcall spi_transfer       /* send RDSR opcode */
    ldi    r16, 0xFF         /* dummy: clocks out the status register */
    rcall spi_transfer       /* r16 = status byte after return */
    rcall cs_deassert
    sbrc   r16, WIP_BIT      /* skip rjmp if WIP = 0 (write done) */
    rjmp   .Lwip_poll
    ret

```

After `spi_transfer` with the dummy byte, `r16` holds the status register. `SBRC r16, 0` skips the `RJMP` when bit 0 (WIP) is clear — loop exits.

This is a common SPI pattern: send a command, then send a dummy byte only to clock the response back from the slave.

14.6.4 main — Putting It Together

```
main:
    /* SP, pins, SPI0 init (see full source) */

    ldi    r24, 0x00          /* EEPROM address          */
    ldi    r22, 0xA5          /* data byte               */
    rcall  eeprom_write_byte /* WREN + WRITE command sequence */
    rcall  eeprom_wait_wip   /* wait up to 5 ms for write cycle */

.Lhalt:
    rjmp   .Lhalt
```

14.6.5 Build and Inspect

```
# Assemble and link with the GCC driver so <avr/io.h> is preprocessed
avr-gcc -mmcu=attiny3217 -x assembler-with-cpp -c -o spi_eeprom.o src/spi_eeprom.S
avr-gcc -mmcu=attiny3217 -nostartfiles -o spi_eeprom.elf spi_eeprom.o

# Disassemble – verify instruction encoding
avr-objdump -d spi_eeprom.elf

# Generate flashable HEX
avr-objcopy -O ihex spi_eeprom.elf spi_eeprom.hex
```

Expected objdump output for spi_transfer:

```
0000007e <spi_transfer>:
 7e: 00 93 24 08    sts    0x0824, r16    ; SPI0_DATA ← tx byte
 82: 10 91 23 08    lds    r17, 0x0823    ; read SPI0_INTFLAGS
 86: 17 ff          sbrs   r17, 7         ; wait for IF (bit 7)
 88: fc cf          rjmp   0x82          ; loop
 8a: 00 91 24 08    lds    r16, 0x0824    ; read received byte, clears IF
 8e: 08 95          ret
```

(Exact symbol labels and addresses vary; with `-nostartfiles`, `objdump` may also show `__ctors_end` at the same address as the first text symbol. The important part is the LDS/STS encoding: 24 08 encodes address 0x0824 in little-endian.)

That is another good reminder that on AVR, peripheral access cost depends on where the register lives.


14.7 TWI (I²C) Fundamentals

TWI uses two wires: SDA (data) and SCL (clock). Both lines are open-drain with pull-up resistors — any device can pull a line low, but no device drives it high (the pull-up does that). This allows multi-master arbitration and clock stretching by slaves.

On ATtiny3217, the default TWI0 pins are PB1=SDA and PB0=SCL. Alternate TWI0 signals are available on PA1=SDA and PA2=SCL through PORTMUX. The bus still needs pull-up resistors in either location.

This is a very different bus philosophy from SPI. TWI uses fewer wires and built-in addressing, but the protocol is more stateful and the electrical behavior is less direct.

14.7.1 Bus Topology




Every device has a 7-bit address (0x00–0x7F). The master sends the address at the start of each transaction; only the matching slave responds.

So TWI solves the “which slave am I talking to?” problem in-band, with an address byte, instead of with separate CS wires.

14.7.2 Transaction Structure

A master write transaction (master sends data to slave) looks like:



- **START (S)**: SDA falls while SCL is high.
- **Address + R/W**: 7-bit address followed by a Read(1) or Write(0) bit.
- **ACK (A)**: slave pulls SDA low during the 9th SCK pulse to acknowledge.
- **STOP (P)**: SDA rises while SCL is high.

14.7.3 I²C Speeds

Mode	Max SCL	Common use
Standard	100 kHz	Most sensors, EEPROMs
Fast	400 kHz	High-speed peripherals
Fast-Plus	1 MHz	ATtiny3217 supports this

14.8 TWI0 Master Mode Registers on ATtiny3217

The ATtiny3217’s TWI0 peripheral has separate register sets for master and slave operation. We use master-mode registers only.

Register	Address	Purpose
TWI0_CTRLA	0x0810	SDA setup/hold time, fast-mode-plus enable
TWI0_DBGCTRL	0x0812	Debug run control
TWI0_MCTRLA	0x0813	Master enable, interrupts, smart mode
TWI0_MCTRLB	0x0814	Flush, ACK action, bus command (STOP etc.)
TWI0_MSTATUS	0x0815	WIF, RIF, CLKHOLD, RXACK, ARBLOST, BUSERR, BUSSTATE
TWI0_MBAUD	0x0816	SCL baud rate register
TWI0_MADDR	0x0817	Send START + address (write to initiate)

TWI0_MDATA 0x0818 Master data (read = receive, write = transmit)

14.8.1 TWI0_MCTRLA — Master Control A

Bit: 7 6 5 4 3 2 1 0
 RIEN WIEN — QCEN TIMEOUT[1:0] SMEN ENABLE

Field	Description
RIEN	Read Interrupt Enable (fires on RIF)
WIEN	Write Interrupt Enable (fires on WIF)
QCEN	Quick Command Enable (ACK sent automatically)
TIMEOUT	Bus timeout (00 = disabled)
SMEN	Smart Mode Enable (ACK sent automatically on read)
ENABLE	Enable TWI master

For polling mode (no interrupts): set only TWI_ENABLE_bm.

14.8.2 TWI0_MSTATUS — Master Status

Bit: 7 6 5 4 3 2 1 0
 RIF WIF CLKHOLD RXACK ARBLOST BUSERR BUSSTATE[1:0]

Flag	Description
RIF	Read Interrupt Flag — byte received and ready in MDATA
WIF	Write Interrupt Flag — address/data write accepted, ACK received
CLKHOLD	Clock hold — SCL is being stretched; write MDATA/MCTRLB to release
RXACK	Received ACK: 0 = ACK (slave acknowledged), 1 = NACK
ARBLOST	Arbitration lost (multi-master: another master won the bus)
BUSERR	Bus error (illegal START/STOP condition detected)
BUSSTATE	00=unknown, 01=idle, 10=owner, 11=busy

WIF is the flag you poll most: it sets after the master successfully sends an address byte or a data byte. Check RXACK immediately after — if RXACK=1 the slave sent NACK (not present or busy).

That RXACK polarity is easy to misread because success is encoded as zero.

14.8.3 TWI0_MBAUD — Baud Rate Register

$$f_{\text{SCL}} = f_{\text{CPU}} / (10 + 2 \times \text{BAUD} + f_{\text{CPU}} \times T_{\text{rise}})$$

Ignoring T_{rise} (assume short wires, small capacitance):

$$\text{BAUD} = (f_{\text{CPU}} / f_{\text{SCL}} - 10) / 2$$

At 3.333 MHz for 100 kHz:

$$\text{BAUD} = (3333333/100000 - 10) / 2 = 11.67$$

Use **12** if you want to stay just below 100 kHz when ignoring rise time: $3333333 / (10 + 2 \times 12) \approx 98.0 \text{ kHz}$.

```
.equ TWI_BAUD_100K, 12      /* 3.333 MHz → ~98 kHz SCL, ignoring rise time */
```


14.8.4 TWI0_MCTRLB — Master Control B (Bus Commands)

Write to TWI0_MCTRLB to issue a bus command after each transaction phase. The MCMD field (bits 1:0) controls the state machine, while ACKACT (bit 2) selects whether the next acknowledge action sends ACK or NACK.

Bit:	7	6	5	4	3	2	1	0
	—	—	—	—	FLUSH	ACKACT	MCMD[1:0]	

ACKACT = 0 sends ACK; ACKACT = 1 sends NACK. ACKACT and MCMD can be written in the same STS. The ACK action matters for host reads; during a host write, the peripheral waits for the next data byte or command.

MCMD value	Symbolic name	Effect
0b00	TWI_MCMD_NOACT_gc	No action
0b01	TWI_MCMD_REPSTART_gc	ACK/NACK action, then repeated START
0b10	TWI_MCMD_RECVTRANS_gc	ACK/NACK action, then receive/transmit
0b11	TWI_MCMD_STOP_gc	ACK/NACK action, then STOP

To send STOP after a write: sts TWI0_MCTRLB, TWI_MCMD_STOP_gc

14.9 TWI Initialisation

```

/* TWI0 pins (default mux): SDA = PB1, SCL = PB0
 * Both open-drain — no DDR configuration needed.
 * The TWI peripheral takes control of the pins when ENABLE = 1.
 * External 4.7 kΩ pull-ups to VCC required. */

/* Force bus state to IDLE before we start */
ldi r16, TWI_BUSSTATE_IDLE_gc /* 0x01 */
sts TWI0_MSTATUS, r16

/* Set baud rate: 100 kHz at 3.333 MHz */
ldi r16, TWI_BAUD_100K /* 12 */
sts TWI0_MBAUD, r16

/* Enable master */
ldi r16, TWI_ENABLE_bm /* 0x01 */
sts TWI0_MCTRLA, r16

```

Writing TWI_BUSSTATE_IDLE_gc (= 0x01) to TWI0_MSTATUS forces the bus state machine to IDLE. Without this, after reset the state is UNKNOWN and the peripheral will not initiate transactions.

This is one of those TWI-specific startup details that feels odd until you remember the peripheral is tracking bus state, not just shifting bytes.

14.10 Master Write Transaction (Polling)

A complete write transaction — send one data byte to a slave device:

Step	Action
------	--------

1. Write (slave_addr << 1 | 0) to TWI0_MADDR → generates START + address
2. Wait for WIF set

3. Check RXACK == 0 (slave ACK); if RXACK == 1 → NACK, abort
4. Write data byte to TWI0_MDATA
5. Wait for WIF set
6. Check RXACK == 0
7. Write TWI_MCMD_STOP_gc to TWI0_MCTRLB → generates STOP

That list is longer than the SPI case because TWI has more protocol structure: START, address, ACK/NACK, data, ACK/NACK, STOP.

In assembly (using r24 = slave address, r22 = data byte):

```
/*
 * twi_write_byte – write one byte to TWI slave
 * r24 = 7-bit slave address (unshifted, 0x00–0x7F)
 * r22 = data byte
 * Returns: r24 = 0 on success, 1 on NACK
 * Clobbers: r16, r17
 */
twi_write_byte:
    /* Send START + address + W (R/W bit = 0) */
    lsl    r24                /* shift left: addr occupies bits[7:1], bit0=0 */
    sts    TWI0_MADDR, r24    /* write triggers START + 9-bit address phase */

    /* Wait for WIF */
.Lwif_addr:
    lds    r17, TWI0_MSTATUS
    sbrs   r17, TWI_WIF_bp    /* WIF = bit 6 */
    rjmp   .Lwif_addr

    /* Check RXACK */
    sbrc   r17, TWI_RXACK_bp /* RXACK = bit 4; skip ret if ACK (bit = 0) */
    rjmp   .Ltwi_nack

    /* Send data byte */
    sts    TWI0_MDATA, r22

    /* Wait for WIF again */
.Lwif_data:
    lds    r17, TWI0_MSTATUS
    sbrs   r17, TWI_WIF_bp
    rjmp   .Lwif_data

    /* Check RXACK for data */
    sbrc   r17, TWI_RXACK_bp
    rjmp   .Ltwi_nack

    /* Send STOP */
    ldi    r16, TWI_MCMD_STOP_gc /* 0x03 */
    sts    TWI0_MCTRLB, r16

    ldi    r24, 0              /* return 0 = success */
    ret

.Ltwi_nack:
    ldi    r16, TWI_MCMD_STOP_gc
    sts    TWI0_MCTRLB, r16 /* release bus even on failure */
    ldi    r24, 1              /* return 1 = NACK error */
    ret
```

14.10.1 SBRC vs SBRS for RXACK

RXACK = 0 means ACK received (success). The condition for *success* is bit = 0, so we use SBRC (Skip if Bit in Register is Cleared) to skip the NACK handler: if the bit is cleared (ACK), skip over `rjmp .Ltwi_nack`.

This is the opposite of USART's SBRS on DREIF — read the bit polarity carefully for each flag.

That is a general embedded rule: never assume status bits use the same success polarity across peripherals.

14.11 Example: Write to PCF8574 I²C I/O Expander

The PCF8574 is a popular 8-bit I²C I/O expander. Its 7-bit address is `0b0100xxx` where `xxx` is set by hardware pins A2:A0. With A2:A0 = `0b000`, the address is `0x20`.

To write a byte to the PCF8574's output latch, a single data byte is sent after the address — no register address needed.

That simplicity is why expanders like the PCF8574 are common first TWI devices: the bus protocol is still there, but the device-side command format is minimal.

```
START | 0x20 W | ACK | 0b10101010 | ACK | STOP
      (slave)                (slave)
```

This drives outputs P7..P0 = 1,0,1,0,1,0,1,0 (alternating).

```
.equ PCF8574_ADDR, 0x20

main:
/* TWI0 init (as shown in 13.9) */
ldi  r16, TWI_BUSSTATE_IDLE_gc
sts  TWI0_MSTATUS, r16
ldi  r16, TWI_BAUD_100K
sts  TWI0_MBAUD, r16
ldi  r16, TWI_ENABLE_bm
sts  TWI0_MCTRLA, r16

/* Write 0xAA to PCF8574 */
ldi  r24, PCF8574_ADDR /* 0x20 */
ldi  r22, 0b10101010 /* output data */
rcall twi_write_byte

.Lhalt:
rjmp .Lhalt
```

14.11.1 Build

```
avr-gcc -mmcu=attiny3217 -nostartfiles -o pcf8574.elf pcf8574.S
avr-objcopy -O ihex pcf8574.elf pcf8574.hex
avrdude -c serialupdi -p t3217 -P /dev/ttyUSB0 -U flash:w:pcf8574.hex
```

14.12 SPI vs TWI Comparison

Feature	SPI	TWI (I ² C)
---------	-----	------------------------

Wires	4 (MOSI, MISO, SCK, CS×n)	2 (SDA, SCL) + pull-ups
Multi-slave	One CS pin per slave	Addressed (7-bit)
Duplex	Full duplex	Half duplex
Max speed here	833 kHz default, 1.667 MHz max	100 kHz / 400 kHz / 1 MHz
Protocol overhead	None (raw shift register)	START/STOP/ACK overhead
Hardware	Shift register + IF flag	State machine (WIF/RIF)
Complexity	Low	Medium
Typical devices	EEPROM, flash, ADC, display	Sensor ICs, I/O expanders

The practical split is usually this: use SPI when you want speed and simple byte streaming, and use TWI when you want several slower peripherals to share the same two wires.

14.13 Summary

SPI0 setup (master, mode 0, f/4, manual CS):

```
SPI0_CTRLB ← SPI_SSD_bm           ; disable hardware SS
SPI0_CTRLA ← SPI_MASTER_bm | SPI_ENABLE_bm
```

SPI transfer:

```
STS SPI0_DATA, r16      ; start transfer
poll SPI_IF_bp in SPI0_INTFLAGS
LDS r16, SPI0_DATA      ; get received byte, clears IF
```

25LC010A write sequence:

1. CS↓, send WREN (0x06), CS↑
2. CS↓, send WRITE (0x02), addr, data, CS↑
3. CS↓, send RDSR (0x05), dummy, CS↑ – check WIP (bit 0)
4. Repeat step 3 until WIP = 0 (max 5 ms)

TWI0 setup (master, 100 kHz):

```
TWI0_MSTATUS ← TWI_BUSSTATE_IDLE_gc ; force idle
TWI0_MBAUD    ← 12                   ; ~98 kHz at 3.333 MHz
TWI0_MCTRLA   ← TWI_ENABLE_bm
```

TWI master write:

```
STS TWI0_MADDR, (addr << 1)          ; START + address + W
wait WIF, check RXACK == 0
STS TWI0_MDATA, data_byte
wait WIF, check RXACK == 0
STS TWI0_MCTRLB, TWI_MCMD_STOP_gc    ; STOP
```

RXACK polarity: 0 = ACK (slave present), 1 = NACK (slave absent/busy)

14.14 Exercises

1. The `spi_transfer` function polls `SPI_IF_bp` in `SPI0_INTFLAGS`. At 3.333 MHz with `SCK = f/4`, how many CPU cycles elapse from writing `SPI0_DATA` until `IF` becomes set? How many poll iterations does the CPU execute during that time? (Use the `LDS+SBRs+RJMP` cycle count.)
2. The 25LC010A has a 16-byte page buffer. A page write starts with address 0x00 and fills up to 16 bytes before `CS↑`. Modify `eeeprom_write_byte` into `eeeprom_write_page` that accepts a pointer (`Z`) and a byte count (`r20`, max 16). Send `WREN`, then one `WRITE`

transaction with all bytes inside the same CS-framed transfer. How does the CS framing differ from the single-byte version?

3. Change the SPI clock to the fastest value this 3.333 MHz CPU clock can generate while remaining below the 25LC010A's 10 MHz limit. Which combination of PRESC and CLK2X gives that SCK? Write the new SPI0_CTRLA value.
4. In `twi_write_byte`, the `BUSSTATE` field of `TWI0_MSTATUS` is not checked before sending. What value does it hold between transactions (after a STOP)? Under what conditions might starting a transaction with `BUSSTATE = BUSY` cause a problem? How would you guard against it?
5. Write `twi_read_byte`: read one byte from a TWI slave after sending its address with `R/W = 1`. After receiving the byte, send NACK + STOP (to tell the slave you don't want more bytes). Which `ACKACT` bit value selects NACK, which `MCMD` value issues STOP, and which `MSTATUS` flag indicates the byte is ready?
6. Using `avr-objdump -d spi_eeprom.elf`, find the STS `SPI0_DATA`, `r16` instruction. What is its 32-bit encoding? Identify which bytes encode the destination address (0x0824) and which bytes encode the source register (`r16 = R16 = register 16`). Recall that STS with a 16-bit address uses the encoding `1001 001d dddd 0000 | kkkk kkkk kkkk kkkk`.

Next: Chapter 14

Chapter 15

CRC and Data Integrity in Assembly

CRC routines are a compact way to detect corrupted messages, tables, and code images. This chapter stays in assembly: a software CRC-8 routine for small data blocks, then the ATtiny3217 CRCSCAN peripheral for checking program Flash.

The two jobs are different:

Software CRC routine:

- Runs over bytes you choose.

- Useful for packets, EEPROM records, protocol frames, and small tables.

CRCSCAN peripheral:

- Runs over Flash sections selected by hardware configuration.

- Useful for startup or runtime program-memory integrity checks.

Microchip documents CRCSCAN on ATtiny3217 as a CRC-16-CCITT memory scanner that can check the entire Flash, the boot + application code section, or the boot section. It can also be configured through FUSE.SYSCFG0 to run during reset initialization before normal code execution starts.

15.1 CRC Basics

A cyclic redundancy check treats a byte stream as a polynomial over bits. Each incoming bit shifts the current remainder. When a one is shifted out, the remainder is XORed with a generator polynomial.

For the byte-oriented routines in this chapter:

CRC-8/SMBus:

- Width: 8 bits

- Polynomial: $0x07 \ (x^8 + x^2 + x + 1)$

- Initial CRC: $0x00$

- Bit order: MSB first

CRC-8 is small enough to write directly as a bit loop. It is not as strong as CRC-16, but it is useful for short records and teaching the mechanics.

For an 8-bit CRC accumulator:

for each input byte:

- `crc = crc XOR byte`

- repeat 8 times:

```

    shift crc left
    if the old bit 7 was one:
        crc = crc XOR polynomial

```

On AVR, LSL shifts bit 7 into SREG.C. That means the branch condition is already available:

```

lsl   r24      ; old bit 7 moves into C
brcc  no_xor   ; if old bit 7 was 0, skip feedback
eor   r24, r25 ; apply polynomial

```

15.2 Software CRC-8 Register Plan

The chapter example computes a CRC over bytes stored in Flash.

Input:

```

Z      byte address of first Flash byte
r20    byte count

```

Output:

```

r24    CRC result

```

Scratch:

```

r18    current input byte
r19    bit counter
r25    polynomial constant

```

Because this is a standalone assembly routine, the interface is simply the register contract above. There is no hidden runtime setup and no external language call boundary.

15.3 Flash Test Data

The example uses the byte sequence 01 02 03 04. For CRC-8/SMBus with polynomial 0x07 and initial value 0x00, the expected CRC is 0xE3.

```

.section .text
.Lmsg:
    .byte 0x01, 0x02, 0x03, 0x04
.balign 2

```

The `.balign 2` matters because the next instruction must start on an even byte address. Raw byte data in `.text` can leave the location counter odd.

15.4 CRC-8 Implementation

```

.equ CRC8_POLY, 0x07

.global crc8_flash
crc8_flash:
    ldi   r24, 0x00      /* CRC accumulator */
    ldi   r25, CRC8_POLY /* polynomial */

    tst   r20
    breq  .Lcrc_done

```

```

.Lcrc_byte:
    lpm    r18, Z+           /* read next Flash byte */
    eor    r24, r18         /* mix byte into CRC */
    ldi    r19, 8           /* 8 bits per byte */

.Lcrc_bit:
    lsl    r24              /* old bit 7 -> C */
    brcc   .Lcrc_no_xor
    eor    r24, r25

.Lcrc_no_xor:
    dec    r19
    brne   .Lcrc_bit

    dec    r20
    brne   .Lcrc_byte

.Lcrc_done:
    ret

```

Important details:

- LPM r18, Z+ reads program memory and advances the Flash byte pointer.
- EOR is used both to mix the input byte and to apply the polynomial.
- LSL sets carry from the old MSB, so no separate mask instruction is needed.
- r24 is reused as the accumulator and final result.

For four bytes, the bit loop runs 32 times. Each bit takes:

No feedback path: LSL + BRCC taken + DEC + BRNE

Feedback path: LSL + BRCC not taken + EOR + DEC + BRNE

The exact cycle count depends on the data because the feedback path includes the extra EOR and different branch timing.

15.5 Standalone Demo

The full source is [crc8.S](#). It computes the CRC and stores the result in SRAM for debugger inspection.

```

.section .bss
.global crc_result
crc_result:
    .skip 1

.global main
main:
    ldi    r16, hi8(RAMEND)
    out    SPH, r16
    ldi    r16, lo8(RAMEND)
    out    SPL, r16

    ldi    ZL, lo8(.Lmsg)
    ldi    ZH, hi8(.Lmsg)
    ldi    r20, 4
    rcall  crc8_flash

    sts    crc_result, r24

```



```
.Lhalt:
    rjmp .Lhalt
```

Build and inspect:

```
avr-gcc -mmcu=attiny3217 -x assembler-with-cpp -c -o crc8.o src/crc8.S
avr-gcc -mmcu=attiny3217 -nostartfiles -o crc8.elf crc8.o
avr-objdump -d crc8.elf
```

Expected signs in the disassembly:

```
lpm    r18, Z+
eor    r24, r18
lsl    r24
brcc   ...
eor    r24, r25
sts    <crc_result>, r24
```

The expected stored value is 0xE3.

15.6 Table-Driven CRC

The bit loop is small, but it costs eight iterations per byte. A faster design uses a 256-byte lookup table in Flash:

```
index = crc XOR input_byte
crc = table[index]
```

On AVR, a table lookup uses Z and LPM. If the table starts on a 256-byte boundary, the index can be placed directly in ZL while ZH holds the page. If the table is not aligned, add the index to the full 16-bit table pointer.

Tradeoff:

Bit loop:

Small code and no table.
Cost grows with 8 bit-steps per input byte.

Table lookup:

Larger Flash use: 256 bytes for CRC-8.
Much faster per byte.

For short packets or rare checks, the bit loop is usually fine. For long streams, the table can be worth the Flash.

15.7 ATtiny3217 CRCSCAN Peripheral

The ATtiny3217 has a dedicated CRCSCAN peripheral at 0x0120.

Register	Address	Purpose

CRCSCAN_CTRLA	0x0120	Enable, NMI enable, reset
CRCSCAN_CTRLB	0x0121	Source selection and Flash access mode
CRCSCAN_STATUS	0x0122	Busy and OK flags

CTRLA bits:

```

Bit 7  RESET   CRCSCAN_RESET_bm   = 0x80
Bit 1  NMIEN   CRCSCAN_NMIEN_bm   = 0x02
Bit 0  ENABLE  CRCSCAN_ENABLE_bm  = 0x01

```

CTRLB fields:

```

SRC[1:0]
0x00  CRCSCAN_SRC_FLASH_gc        entire Flash
0x01  CRCSCAN_SRC_APPLICATION_gc  boot + application code section
0x02  CRCSCAN_SRC_BOOT_gc         boot section

```

```

MODE[5:4]
0x00  CRCSCAN_MODE_PRIORITY_gc    priority access to Flash

```

STATUS bits:

```

Bit 1  OK      CRCSCAN_OK_bm      = 0x02
Bit 0  BUSY    CRCSCAN_BUSY_bm    = 0x01

```

Microchip documents the hardware algorithm as CRC-16-CCITT:

Polynomial: $x^{16} + x^{12} + x^5 + 1$
 Initial checksum register value: 0xFFFF
 Input order: byte by byte, starting at byte 0, MSB first

The last two bytes of the selected Flash section must contain the precomputed 16-bit checksum. For full Flash, Microchip places CHECKSUM[15:8] at FLASHEND-1 and CHECKSUM[7:0] at FLASHEND. Boot and boot+application checks use the corresponding BOOTEND or APPEND fuse boundary. If the calculated value matches the placed checksum, CRCSCAN_STATUS.OK is set.

15.8 Starting a CRCSCAN in Assembly

The software sequence is short:

1. Write CRCSCAN_CTRLB to select the Flash source.
2. Write CRCSCAN_CTRLA with ENABLE.
3. Poll STATUS.BUSY until it clears.
4. Test STATUS.OK.

Example: scan entire Flash without NMI.

```

crcscan_flash:
    ldi    r16, CRCSCAN_SRC_FLASH_gc
    sts    CRCSCAN_CTRLB, r16

    ldi    r16, CRCSCAN_ENABLE_bm
    sts    CRCSCAN_CTRLA, r16

.Lcrcscan_busy:
    lds    r16, CRCSCAN_STATUS
    sbrc   r16, CRCSCAN_BUSY_bp
    rjmp   .Lcrcscan_busy

    sbrs   r16, CRCSCAN_OK_bp
    rjmp   .Lcrcscan_failed
    ret

.Lcrcscan_failed:
    rjmp   .Lcrcscan_failed

```

Microchip notes that a software-started scan begins after three cycles. In priority mode, CRCSCAN has priority access to Flash and stalls the CPU until the scan is complete. That makes runtime scans easy to reason about but not free; schedule them where a pause is acceptable.

15.9 NMI on CRC Failure

CRCSCAN can trigger a Non-Maskable Interrupt when a scan fails:

```
ldi    r16, (CRCSCAN_ENABLE_bm | CRCSCAN_NMIEN_bm)
sts    CRCSCAN_CTRLA, r16
```

The CRCSCAN NMI vector is vector 1. On ATtiny3217, vector entries are 4-byte JMP entries, so vector 1 starts at byte address 0x0004:

```
__vectors:
    jmp    main                /* 0x0000 RESET */
    jmp    crcscan_nmi         /* 0x0004 CRCSCAN_NMI */
```

A CRC failure means program memory may be unreliable, so the NMI handler should be extremely conservative. A typical response is to put outputs in a safe state and stop execution or force a reset path. Do not treat a failed program-memory CRC like an ordinary recoverable event. Once CRCSCAN triggers an NMI, the NMI request remains active until a system reset and cannot be disabled in software.

15.10 Startup CRC Through Fuses

FUSE_SYSCFG0 contains the CRCSRC field at bits 7:6. Microchip documents that this field can make CRCSCAN run during reset initialization. If the startup CRC check fails, normal code execution is not allowed to start. The failure can still be observed through the debug interface.

The ATtiny3217 fuse encodings are:

```
CRCSRC[1:0]
0x00  CRCSRC_FLASH_gc    entire Flash
0x01  CRCSRC_BOOT_gc     boot section
0x02  CRCSRC_BOOTAPP_gc  boot + application code section
0x03  CRCSRC_NOCRC_gc    startup CRC disabled
```

That mode is different from a software-started runtime scan:

Runtime scan:

Code starts first, then software enables CRCSCAN.

Startup scan:

Reset sequence checks Flash before normal code execution.

Use startup CRC for stronger boot-time assurance. Use runtime CRC when the program needs to periodically verify a Flash region while running.

15.11 Summary

Software CRC-8:

Polynomial 0x07

Initial value 0x00
 LSL moves old bit 7 into SREG.C
 BRCC decides whether to XOR the polynomial

ATtiny3217 CRCSCAN:

Registers: CTRLA=0x0120, CTRLB=0x0121, STATUS=0x0122
 Hardware algorithm: CRC-16-CCITT
 Polynomial: $x^{16} + x^{12} + x^5 + 1$
 Initial value: 0xFFFF
 Sources: entire Flash, boot + application code, or boot section
 STATUS.BUSY bit 0, STATUS.OK bit 1
 Optional NMI on failure
 Optional startup scan through FUSE_SYSCFG0.CRCSRC

15.12 Exercises

1. Step through `crc8_flash` for the first byte 0x01. After the EOR, what is the accumulator? Which bit iterations take the feedback path?
 2. Modify `crc8_flash` to compute over SRAM instead of Flash. Which instruction changes, and how does the pointer setup differ?
 3. Add a second test message in Flash and store two result bytes in SRAM.
 4. Rewrite the CRC-8 routine so the polynomial is passed in r21 instead of assembled as `CRC8_POLY`.
 5. Write a `crcscan_boot` routine that selects `CRCSRC_SCAN_BOOT_gc`, starts a scan, waits for `BUSY=0`, and returns `r24=0` on OK or `r24=1` on failure.
 6. Estimate the runtime cost of a CRCSCAN priority scan over 32 KB Flash using Microchip's note that the peripheral fetches one 16-bit word every third main clock cycle in priority mode.
-

Next: Chapter 15

Chapter 16

Optimisation Techniques

Assembly optimisation is useful only when it is measured. The ATtiny3217 is predictable enough that many improvements can be proven from the instruction listing alone: count the instructions on the hot path, apply the cycle counts from Microchip's AVR Instruction Set Manual, then compare the result with the flash and SRAM cost.

This chapter keeps the examples assembly-only. The benchmark harness in `src/bench_harness.S` calls the routines directly with registers, gives the linker real symbols to place, and leaves the result visible to `avr-objdump`, `avr-nm`, GDB, or MDB.

The techniques:

1. **Cycle counting** - build a baseline from the disassembly.
2. **Register residency** - keep hot values in registers, not SRAM.
3. **Loop unrolling** - trade flash for fewer branch and counter cycles.
4. **Flash lookup tables** - trade flash for much faster fixed-domain results.
5. **Size checks** - verify that speed did not quietly consume too much flash.

Measure first, optimise second.

16.1 Build and Inspect the Chapter Examples

Build from the chapter directory:

```
cd ch15_optimisation

avr-gcc -mmcu=attiny3217 -x assembler-with-cpp -g3 -Wa,--gdwarf-2 \
  -nostartfiles -o bench.elf \
  src/bench_harness.S src/memcpy_naive.S src/memcpy_fast.S src/lut_sin.S
```

Inspect symbols and sizes:

```
avr-nm --size-sort --print-size bench.elf
```

Inspect instructions:

```
avr-objdump -d -S bench.elf
```

Useful focused views:

```
avr-objdump -d bench.elf | grep -A 20 '<memcpy_naive>'
avr-objdump -d bench.elf | grep -A 40 '<memcpy_fast>'
avr-objdump -d bench.elf | grep -A 12 '<sin8>'
avr-objdump -s -j .text bench.elf
```

The examples are linked into one ELF so the addresses and code sizes are real. No generated PDF rebuild is required for this chapter edit.

16.2 Cycle Reference for This Chapter

The ATtiny3217 uses the AVRxt timing column in Microchip’s AVR Instruction Set Manual. Always check the manual for instructions not listed here, especially when carrying code to older AVR, AVRxm, or reduced-core AVRrc devices.

Common timing used in this chapter:

Instruction	ATtiny3217 cycles	Notes
NOP	1	
MOV, MOVW	1	
LDI	1	r16-r31 only
CLR	1	encoded as EOR Rd,Rd
ADD, ADC	1	
AND, ANDI, EOR	1	
INC, DEC	1	
LSR	1	
TST	1	encoded as AND Rd,Rd
BREQ/BRNE not taken	1	
BREQ/BRNE taken	2	
RJMP	2	
RCALL	2	
RET	4	larger PC-width AVRr can differ
LD Rd, Z+	2	
ST X+, Rr	1	
LPM Rd, Z	3	
LPM Rd, Z+	3	
STS k, Rr	2	two-word instruction

Two details cause most bad cycle counts:

- A conditional branch costs 2 cycles when taken and 1 cycle when it falls through.
- Count the last loop iteration separately, because the loop branch usually falls through only once.

For a loop that executes N times:

```
branch cycles = (N - 1) * 2 + 1
```

16.3 Baseline: One-Byte Copy Loop

The file `src/memcpy_naive.S` copies one byte per loop iteration.

Register interface:

R25:R24 destination pointer
R23:R22 source pointer
R20 byte count

Inner loop:

```
.lnaive_loop:
    ld    r18, Z+          /* 2 cycles */
    st    X+, r18          /* 1 cycle  */
```

```

dec    r20                /* 1 cycle */
brne   .Lnaive_loop       /* 2 taken, 1 final fall-through */

```

Taken-path loop cost:

$2 + 1 + 1 + 2 = 6$ cycles per byte

For $\text{len} = 16$, the branch is taken 15 times and falls through once:

```

15 taken iterations: 15 * 6 = 90
1 final iteration:   2 + 1 + 1 + 1 = 5
inner-loop total:    95 cycles

```

Full subroutine cost for $\text{len} = 16$ on ATtiny3217:

```

setup:      movw + movw + tst + breq(not taken) = 4
copy loop:  95
ret:        4
subroutine: 103 cycles
with RCALL: 105 cycles from the harness call site

```

The important baseline is the loop body: **95 cycles for 16 bytes**, or about 5.94 cycles per byte including the final branch fall-through.

16.4 Keep Hot Values in Registers

SRAM is much slower than a register. Direct SRAM access also costs flash because LDS and STS are two-word instructions on this device family.

Bad pattern for hot loop state:

```

.Lloop:
    lds    r18, counter
    inc    r18
    sts    counter, r18
    dec    r20
    brne   .Lloop

```

That loop repeatedly pays for direct SRAM access. If the value is only needed by this routine, keep it in a register:

```

    clr    r18
.Lloop:
    inc    r18
    dec    r20
    brne   .Lloop
    sts    counter, r18    /* store once if an observable result is needed */

```

The same rule applies to pointers and accumulators. Plan register use before writing the loop:

```

Routine: block_sum
Z        source pointer
r20      byte count
r18      loaded byte
r24      accumulator/result
r1       zero source for ADC carry propagation when needed

```

Example:

```

.global block_sum
.type    block_sum, @function

block_sum:
    movw    ZL, r22
    clr     r24
    tst     r20
    breq    .Lsum_done

.Lsum_loop:
    ld      r18, Z+
    add     r24, r18
    dec     r20
    brne    .Lsum_loop

.Lsum_done:
    ret
.size block_sum, .-block_sum

```

No SRAM is touched inside the loop. The accumulator stays live in r24 until the routine returns.

When registers run short, do not spill in the inner loop by reflex. First try:

- reusing argument registers once their original value is no longer needed
- shortening the lifetime of temporary values
- splitting a large loop into smaller named phases
- saving a register once at entry and restoring it once at exit

One PUSH/POP pair outside a hot loop is often much cheaper than LDS/STS inside every iteration.

16.5 Loop Unrolling

Each iteration of the naive copy loop spends 3 cycles on loop overhead:

```

DEC    = 1
BRNE   = 2 when taken

```

For a tiny loop body, that overhead is a large fraction of the total. Unrolling processes multiple bytes per branch.

The file `src/memcpy_fast.S` processes four bytes per block:

```

.Lfast_block:
    ld      r18, Z+           /* 2 */
    ld      r19, Z+           /* 2 */
    ld      r22, Z+           /* 2 */
    ld      r23, Z+           /* 2 */
    st      X+, r18           /* 1 */
    st      X+, r19           /* 1 */
    st      X+, r22           /* 1 */
    st      X+, r23           /* 1 */
    dec     r20               /* 1 */
    brne    .Lfast_block     /* 2 taken, 1 final */

```

Taken-path block cost:


```

4 loads + 4 stores + DEC + BRNE
= 8 + 4 + 1 + 2
= 15 cycles per 4 bytes
= 3.75 cycles per byte

```

For len = 16, there are four blocks:

```

3 taken blocks:      3 * 15 = 45
1 final block:       8 + 4 + 1 + 1 = 14
block-loop total:    59 cycles

```

The fast routine also has setup and tail handling:

```

setup before block loop: 9 cycles
block loop:              59 cycles
tail test, no tail:      3 cycles
ret:                    4 cycles
subroutine:              75 cycles
with RCALL:              77 cycles from the harness call site

```

Compare the hot copy work:

```

naive inner loop, len=16: 95 cycles
unrolled block loop:      59 cycles
speedup in copy loop:     95 / 59 = 1.61x

```

Compare the full call from the harness:

```

naive with RCALL:         105 cycles
fast with RCALL:          77 cycles
full-call speedup:        105 / 77 = 1.36x

```

Both numbers matter. The first tells you what happens on long buffers where setup fades away. The second tells you what this exact len = 16 benchmark costs.

16.6 Tail Handling

Unrolling by four creates a remainder problem. Lengths such as 13 or 17 do not divide evenly by four.

The fast routine computes:

```

mov    r21, r20
andi   r21, 0x03      /* tail = len & 3 */
lsr    r20
lsr    r20             /* blocks = len >> 2 */
breq   .Lfast_tail    /* len < 4 */

```

Then the block loop copies groups of four and the tail loop copies the last 0-3 bytes:

```

.Lfast_tail:
    tst    r21
    breq   .Lfast_done
    ld     r18, Z+
    st     X+, r18
    dec    r21
    brne   .Lfast_tail

```

This is the price of a general-purpose unrolled routine. The fast path is cheaper per byte, but the control flow is larger than the naive loop.

Unrolling is usually worth it when:

- the loop runs many times
- the useful work per iteration is small
- flash budget is available
- the tail path is rare or short

It is often not worth it when:

- the loop normally runs only a few times
- the loop body already contains expensive work
- the routine is cold startup code
- flash space is the limiting resource

16.7 Flash Lookup Tables

Some computations have a small input domain. If every possible input fits in one byte, a 256-entry table can replace a long calculation.

The file `src/lut_sin.S` contains a 256-entry unsigned sine table:

```
.section .progmem.data,"a",@progbits
.global sin_table
sin_table:
    .byte 128,131,134,137,140,143,146,149
    /* full 256-entry table in src/lut_sin.S */
.size sin_table, .-sin_table
```

The table value is:

```
128 + 127 * sin(n * 2 * pi / 256)
```

Rounded to bytes, angle 0 gives 128, angle 64 gives 255, angle 128 gives 128, and angle 192 gives 0.

The lookup routine:

```
.global sin8
.type sin8, @function

sin8:
    ldi ZL, lo8(sin_table)
    ldi ZH, hi8(sin_table)
    add ZL, r24
    adc ZH, r1
    lpm r24, Z
    ret
.size sin8, .-sin8
```

Cycle count on ATtiny3217:

```
LDI + LDI + ADD + ADC + LPM + RET
= 1 + 1 + 1 + 1 + 3 + 4
= 11 cycles
```

With RCALL from the harness: 13 cycles

The table costs 256 flash bytes. The routine costs 12 flash bytes. For a value that is used repeatedly, that is often a good trade.

16.7.1 Why ADC ZH, r1 Works

The `ADD ZL, r24` instruction may carry from the low byte of the table address into the high byte. `ADC ZH, r1` adds that carry to ZH.

This requires r1 to be zero. The chapter harness clears r1 before calling the routines:

```
clr    r1
```

If a standalone assembly program uses r1 as a zero source, it must preserve that convention itself. There is no startup code in these examples to do it for you.

16.7.2 Flash Address Limits

LPM addresses the low 64 KB of program memory. The ATtiny3217 has 32 KB of flash, so ordinary LPM is sufficient. Larger AVR devices may require ELPM and RAMPZ for tables above the low 64 KB region.

16.7.3 Verifying Table Placement

The default AVR linker script places program-memory input sections into the final flash image. In the linked ELF, inspect the final `.text` output section:

```
avr-nm -n bench.elf | grep sin_table
avr-objdump -s -j .text bench.elf
```

Do not assume that an input section name such as `.progmem.data` remains a separate output section after linking.

16.8 Size Results from the Current Sources

With the current chapter files, `avr-nm --size-sort --print-size bench.elf` reports the relevant symbols as:

```
00000202 0000000c T sin8
00803820 00000010 B dst_fast
00803810 00000010 B dst_naive
00803800 00000010 B src_data
000001bc 00000012 T memcpy_naive
000001ce 00000034 T memcpy_fast
0000017e 0000003e T main
0000007c 00000100 T sin_table
```

Meaning:

Symbol	Size	Meaning
<code>memcpy_naive</code>	18 bytes	Smallest copy routine.
<code>memcpy_fast</code>	52 bytes	Faster for larger lengths, 34 bytes larger.
<code>sin8</code>	12 bytes	Lookup routine only.
<code>sin_table</code>	256 bytes	Flash table used by <code>sin8</code> .
<code>src_data</code>	16 bytes	SRAM source buffer for the harness.
<code>dst_naive</code>	16 bytes	SRAM destination for naive copy.
<code>dst_fast</code>	16 bytes	SRAM destination for unrolled copy.

<code>sin_sample</code>	1 byte	SRAM location storing one sine lookup result.
-------------------------	--------	---

Speed is not free. In this example, the unrolled copy routine saves cycles by spending 34 extra flash bytes. The sine lookup spends 256 flash bytes to replace a much longer calculation with one table load.

16.9 Reading `avr-objdump` Correctly

Example disassembly:

```
000001c4 <.lnaive_loop>:
1c4: 21 91      ld      r18, Z+
1c6: 2d 93      st      X+, r18
1c8: 4a 95      dec     r20
1ca: e1 f7      brne    .-8
```

Read it as:

```
1c4      byte address in flash
21 91    encoded instruction bytes, little-endian in the listing
ld ...   decoded instruction
```

AVR instructions are stored as 16-bit words, but the ELF view uses byte addresses. Most instructions advance by 2 bytes. Two-word instructions such as STS, LDS, CALL, and JMP advance by 4 bytes.

Use these commands often:

```
avr-objdump -d -S bench.elf
avr-objdump -h bench.elf
avr-nm -n bench.elf
avr-nm --size-sort --print-size bench.elf
```

If a symbol reports an unexpected size, check that the source has matching `.type` and `.size` directives.

16.10 Common Errors in Hand Optimisation

16.10.1 Optimising Cold Code

Speed work belongs on hot paths. A routine that runs once during setup is rarely worth making larger or harder to inspect. Use a timer, GPIO toggle, logic analyzer, simulator trace, or debugger stop points to identify where time is actually spent.

16.10.2 Ignoring Setup Cost

An unrolled loop can be faster in steady state but slower for tiny lengths. Do both calculations:

```
steady-state loop cost
full routine cost including setup, tail, and return
```

For `len = 16`, the unrolled copy is still faster. For very short lengths, the naive loop may win.

16.10.3 Forgetting the Final Branch

The last conditional branch usually costs one cycle, not two. Count taken and not-taken paths separately.

16.10.4 Using a Zero Register Without Clearing It

The lookup example uses `r1` as a zero source. In a fully standalone assembly program, clear it yourself before relying on it:

```
clr r1
```

If a routine temporarily uses `r1` for some other purpose, restore it to zero before returning to code that expects the zero-register convention.

16.10.5 Putting Hot State in SRAM

Repeated LDS/STS inside a tight loop is often the first thing to remove. Keep counters, pointers, accumulators, masks, and temporary bytes in registers for as long as their live ranges allow.

16.10.6 Unrolling Until the Source Becomes Fragile

Unrolling `x2` or `x4` is often manageable. Unrolling `x8` or `x16` can become hard to audit, especially when there is a tail path, status flags, or overlapping source and destination ranges. Code that cannot be reviewed reliably is not a professional optimisation.

16.11 Optimisation Checklist

Use this sequence before changing a routine:

1. Mark the exact hot loop or hot path.
2. Build the ELF.
3. Disassemble with `avr-objdump -d -S`.
4. Count cycles for the current path.
5. Record flash size for the current routine.
6. Apply one change.
7. Rebuild and recount.
8. Check flash size again.
9. Keep the change only if the gain is worth the size and complexity.

For AVR assembly, this discipline works well because the machine is small and the instruction timings are documented.

16.12 Source Notes

Facts in this chapter are based on:

- Microchip AVR Instruction Set Manual DS40002198C, especially the AVRxt cycle timing column used by ATtiny3217-class devices.
 - Microchip ATtiny3217 documentation for the 32 KB flash size that makes plain LPM sufficient for the chapter's lookup table.
 - The local linked ELF produced from `src/bench_harness.S`, `src/memcpy_naive.S`, `src/memcpy_fast.S`, and `src/lut_sin.S`.
-

16.13 Exercises

1. Count the exact full-call cycle cost of `memcpy_naive` for `len = 8`, including `RCALL`, `setup`, `loop`, and `RET`.
2. Count the exact full-call cycle cost of `memcpy_fast` for `len = 8`. Does the unrolled routine still win?
3. Modify `memcpy_fast` to unroll by eight instead of four. Measure the new symbol size and calculate the new loop cost for `len = 64`.
4. Write `block_sum`, an assembly routine that sums `len` bytes from `Z` into `r24`. Compare the cycle count before and after unrolling by two.
5. Change the sine table to a 64-entry quarter-wave table and reconstruct the quadrant in code. How many flash bytes are saved, and how many cycles are added?
6. Use GDB or MDB to stop at `bench_done` and inspect `dst_naive`, `dst_fast`, and `sin_sample`. Explain what each byte proves about the benchmark.

Next: Chapter 16 - Bootloaders

Chapter 17

Self-Programming Flash: NVMCTRL on tinyAVR 1-Series

Classic AVR devices (ATmega328P and friends) use SPM from a classic boot section to write flash. The **ATtiny3217**, like other modern tinyAVR 1-series parts, takes a different approach: self-programming uses normal stores to the mapped Flash page buffer plus NVMCTRL commands, not the classic SPM flow, and there is no BOOTRST fuse. Instead, a memory-mapped peripheral called **NVMCTRL** (Non-Volatile Memory Controller) handles self-programming, while the BOOTEND and APPEND fuses divide flash into BOOT, APPCODE, and APPDATA sections with directional write protection.

This chapter covers everything needed to write assembly code that programs its own flash on the ATtiny3217 Curiosity Nano board:

1. Why AVRXMEGA3 replaced SPM with NVMCTRL.
2. How NVMCTRL registers and commands work.
3. The page buffer and the exact write sequence.
4. Configuration Change Protection (CCP) — the timing-critical unlock.
5. A complete 128-byte page-write routine in assembly.
6. A UART-driven self-update loop that receives Intel HEX and writes it to flash live.
7. The BOOTEND and APPEND fuses: defining protected flash sections.
8. Building, uploading, and common pitfalls.

This is one of the more hardware-specific chapters in the book. The goal is to turn “the chip can rewrite its own flash” from a mysterious feature into a sequence you can reason about step by step.

17.1 tinyAVR 1-Series vs Classic AVR: No SPM, NVMCTRL Instead

17.1.1 Architecture Family: AVRXMEGA3

The ATtiny3217 belongs to the **AVRXMEGA3** sub-architecture. The three key differences that affect self-programming are:

Feature	ATmega328P (classic)	ATtiny3217 (AVRXMEGA3)
Self-program flow	SPM-driven	ST to mapped Flash + NVMCTRL
Boot section	fixed classic boot area	configurable BOOT section
BOOTRST fuse	yes	none
Flash write ctrl.	SPMCSR register	NVMCTRL peripheral

Flash address space	separate word-address	byte-address, unified
Page size	128 bytes (64 words)	128 bytes (64 words)
Write privileges	boot section only	BOOT can write APPCODE/APPDATA; APPCODE can write APPDATA
UPDI programming	no	yes (PA0, replaces ISP)

The AVRXMEGA3 memory map is **unified**: flash, SRAM, and peripherals all share the same byte-addressed data space. Flash byte address 0 appears at data address 0x8000 when accessed by ordinary LD/ST instructions. The CPU's program counter still addresses executable flash from 0x0000. You will see both conventions in this chapter.

17.2 Microchip App Notes Used

The main Microchip reference for this chapter is **AN2634, Bootloader for tinyAVR 0- and 1-series, and megaAVR 0-series**. The relevant points from that application note are:

- tinyAVR 0/1-series and megaAVR 0-series devices program flash one page at a time.
- The page buffer has the same size as the flash page: 64 or 128 bytes depending on the device. The ATtiny3217 Flash page size is 128 bytes.
- The target page must be erased before the page buffer is written to flash. Programming an unerased page corrupts its contents.
- PAGEERASEWRITE starts the erase and write as one NVMCTRL command.
- The page buffer is automatically cleared after NVMCTRL commands execute.
- For tinyAVR 0/1-series, flash is mapped at data address 0x8000 for normal load/store access. Microchip's headers call this MAPPED_PROGMEM_START.
- BOOTEND and APPEND define BOOT, APPCODE, and APPDATA sections in 256-byte units.
- The documented write privileges are directional: BOOT can write APPCODE and APPDATA; APPCODE can write APPDATA; APPDATA cannot write flash or EEPROM.
- AN2634 warns against issuing CHIPERASE during self-programming because it erases the flash array.

The ATtiny3216/3217 data sheet adds two important details:

- The page buffer write uses the least significant address bits as the offset in the page buffer, and new page-buffer data is combined with previous buffer contents by a binary AND.
- After the self-programming CCP signature is written to CPU_CCP, the protected NVMCTRL command must be executed within the next four instructions.

17.2.1 What Replaces SPM?

On tinyAVR 1-series, you fill the flash **page buffer** by writing normally to flash addresses using ST (Store Indirect). There is no special opcode. The NVMCTRL peripheral watches these writes and accumulates them in a 128-byte buffer. When you issue a command to NVMCTRL_CTRLA, it commits the buffer to the actual flash array.

Classic AVR SPM flow:

```
load R1:R0 ← data word
write SPMCSR ← mode
SPM ← fills buffer
(repeat for each word)
write SPMCSR ← PGERS
SPM ← erases page
write SPMCSR ← PGWRT
SPM ← writes page
```

tinyAVR NVMCTRL flow:

```
(nothing special)
wait FBUSY = 0
ST Z+, rn ← fills buffer (just write!)
(repeat for each word)
CCP unlock (0x9D → 0x0034)
STS NVMCTRL_CTRLA ← 0x03
wait FBUSY = 0
```


The NVMCTRL approach is simpler to understand: write your data to the target flash addresses, then send the “erase+write” command. The hardware does the rest.

The important mental shift is that self-programming no longer looks like a special instruction stream. It looks like ordinary memory-mapped peripheral control wrapped around a careful write sequence.

17.2.2 Programming Interface: UPDI

The ATtiny3217 has no ISP pins. All external programming is done over **UPDI** (Unified Program and Debug Interface), a single-wire protocol on **PA0**. The Curiosity Nano board’s on-board debugger bridges USB to UPDI automatically. Use `pymcuprog` for the on-board debugger:

```
pymcuprog write -d attiny3217 -t uart -f firmware.hex
```

Once the self-update firmware is running, the device reprograms itself — no external programmer needed for firmware updates.

17.3 NVMCTRL Overview

17.3.1 Flash Parameters

ATtiny3217 flash layout

Total flash:	32 KB (32768 bytes)
Page size:	128 bytes
Number of pages:	256
Flash data space:	0x8000–0xFFFF (byte addresses when read as data)
LPM address range:	0x0000–0x7FFF (byte addresses for program memory reads)
SRAM:	0x3800–0x3FFF (2 KB)

The 128-byte page size is critical: every NVMCTRL flash write operation commits one page. If your update stream ends with a short final page, pad the unwritten bytes in the SRAM page image with 0xFF before issuing `PAGEERASEWRITE`.

That last point matters because it explains why page alignment and buffering show up everywhere in self-programming code. Flash is not rewritten one byte at a time, even if your firmware update logic receives bytes one at a time.

17.3.2 NVMCTRL Register Map

All NVMCTRL registers are outside the I/O space and require LDS/STS:

Register	Address	Width	Purpose
NVMCTRL_CTRLA	0x1000	1 byte	Command register – write command here
NVMCTRL_CTRLB	0x1001	1 byte	Config: BOOTLOCK, APCWP write-protect
NVMCTRL_STATUS	0x1002	1 byte	Status flags: FBUSY, EEBUSY, WRERROR
NVMCTRL_INTCTRL	0x1003	1 byte	Interrupt enable: EEREADY (bit 0)
NVMCTRL_INTFLAGS	0x1004	1 byte	Interrupt flags (write 1 to clear)
NVMCTRL_DATA	0x1006	2 bytes	Data register (used for EEPROM word)
NVMCTRL_ADDR	0x1008	4 bytes	Address register (used for EEPROM)

For flash programming, you only need `NVMCTRL_CTRLA` and `NVMCTRL_STATUS`.

17.3.3 NVMCTRL_STATUS — Status Register (0x1002)

Bit:	7	6	5	4	3	2	1	0
	—	—	—	—	—	WRERROR	EEBUSY	FBUSY

FBUSY — Flash Busy: 1 while a flash erase/write operation is in progress.
Poll this bit; do not issue a new command while FBUSY = 1.

EEBUSY — EEPROM Busy: 1 while an EEPROM operation is in progress.

WRERROR — Write Error: set if a write was attempted to a protected region.
Write 1 to clear (W1C).

17.3.4 NVMCTRL Commands (written to NVMCTRL_CTRLA)

Value	Name	Description
0x00	NOCMD	No operation
0x01	PAGEWRITE	Write page buffer to flash (no erase first)
0x02	PAGEERASE	Erase one flash page (Z must point into page)
0x03	PAGEERASEWRITE	Erase + write in one atomic step ← use this
0x04	PAGEBUFFERASE	Clear the page buffer (all bytes → 0xFF)
0x05	CHIPERASE	Erase flash array; do not use in this updater
0x06	EEERASE	Erase EEPROM
0x07	FUSEWRITE	Write fuse bytes (restricted, rarely needed)

PAGEERASEWRITE (0x03) is the command you will use for firmware updates. It atomically erases the target page and writes the page buffer in one operation — safer than doing them separately, because a power failure between separate erase and write steps cannot leave a page half-erased.

In practice, that combined command is what makes the sequence robust enough to use in the field. Fewer separate steps means fewer partially completed states.

17.4 The Page Buffer and Write Sequence

17.4.1 How the Page Buffer Works

The NVMCTRL page buffer is a 128-byte staging area internal to the NVM controller. It is **not** directly addressable. You populate it by writing to the corresponding flash addresses in the data address space:

```
mapped_flash_address = 0x8000 + flash_byte_address
```

When the Z pointer points into flash address space and you execute ST Z+, Rn, the byte goes into the page buffer at the offset corresponding to Z's address. The hardware handles this automatically whenever the NVM controller recognises the write target as flash.

Page buffer write – step by step:

Flash page at byte address 0x0000 appears in data space at 0x8000.

```
Z = 0x8000          ; points to byte 0 of flash page 0
st Z+, r16          ; → page buffer[0] = r16
st Z+, r17          ; → page buffer[1] = r17
...                 ; fill all 128 bytes
st Z+, r30          ; → page buffer[127] = r30
                    ; Z = 0x8080 (past end of page)
```

The page buffer is erased to all ones after reset, after page write/erase commands, after a Clear Page Buffer command, and after wake-up from sleep. The least significant address

bits select the offset inside the page buffer. If you write the same page-buffer byte more than once, the resulting value is the binary AND of the old buffer value and the new value.

That AND behavior matters. A partial page update is not automatically merged with the existing flash page for you. For a robust partial update, either fill unwritten bytes with 0xFF when programming a fresh page, or read the old page into SRAM, modify the intended bytes, erase the flash page, and write the full merged 128-byte image.

17.4.2 Complete Flash Page Write Sequence

Step 1: Wait for any previous NVM operation to finish.
Poll NVMCTRL_STATUS.FBUSY until it reads 0.

Step 2: Clear any stale write error.
STS NVMCTRL_STATUS, 0x04 ; WRERROR is write-one-to-clear

Step 3: Fill the page buffer.
Set Z = 0x8000 + (target_page × 128) ; data-space address
ST Z+, r0 ; ST Z+, r1 ; ... ; 128 sequential writes

Step 4: Issue PAGEERASEWRITE (CCP-protected).
LDI r16, 0x03 ; PAGEERASEWRITE command, prepared first
LDI r17, 0x9D ; CCP_SPM_gc key, prepared first
STS 0x0034, r17 ; write to CCP register
STS NVMCTRL_CTRLA, r16 ; next instruction: protected command

Step 5: Wait for FBUSY to clear again.
Poll NVMCTRL_STATUS.FBUSY until 0.

17.4.3 The Z Pointer Convention

On ATtiny3217 the Z pointer (R31:R30) is a 16-bit register, and 0x8000 is well within that range. For page-buffer writes with ST Z+, Rn, set Z to the mapped program-memory address:

```
/* To write page at flash byte address 0x0200 (page 4): */
ldi ZL, lo8(0x8200)
ldi ZH, hi8(0x8200)
/* Now ST Z+, Rn writes to page buffer slots starting at offset 0 */
```

Equivalently:

```
ldi ZL, lo8(0x0200)
ldi ZH, hi8(0x0200)
subi ZH, -0x80 /* add mapped flash base 0x8000 */
```

The examples in Sections 16.5 and 16.6 use this mapped-address convention.

That convention keeps the assembly easier to follow because the same flash byte address can be used consistently through buffering, writing, and later jumps.

17.5 CCP — Configuration Change Protection

17.5.1 Why CCP Exists

Some NVMCTRL commands (and other sensitive operations like changing clock settings) can corrupt firmware if executed accidentally — for example, if a runaway pointer writes to NVMCTRL_CTRLA, or if a flash write executes during a brownout. **CCP** (Configuration

Change Protection) prevents accidental execution by requiring a timed unlock sequence before each protected write.

So CCP is not there to make your life difficult. It is there because an accidental flash-control write is serious enough that the hardware demands a deliberate two-step action.

17.5.2 The CCP Register (0x0034)

Address: 0x0034 (in I/O space: use OUT or STS)

CCP keys:

- 0x9D – CCP_SPM_gc Unlocks NVM Self-Programming commands
- 0xD8 – CCP_IOREG_gc Unlocks protected I/O register writes (e.g. CLKCTRL)

After writing a key, the protected write must occur within the next four CPU instructions. Interrupt requests are held pending during this short configuration-change window.

17.5.3 The Four-Instruction Rule

The CCP window opens the moment you write the key. Use the strictest practical rule in assembly: make the protected NVMCTRL_CTRLA write the very next instruction after the CPU_CCP write:

```
sts    CPU_CCP, r17
sts    NVMCTRL_CTRLA, r16
```

Assembly idiom:

```
ldi    r16, 0x03          /* PAGEERASEWRITE command */
ldi    r17, 0x9D          /* CCP_SPM_gc key */
sts     0x0034, r17        /* open the CCP window */
sts     NVMCTRL_CTRLA, r16 /* next instruction: protected write */
```

Critical: do not put any instruction between STS CPU_CCP and STS NVMCTRL_CTRLA — not even a NOP.

This is one of the rare cases where instruction spacing is not just an optimization detail. It is part of the functional correctness of the code.

17.6 Flash Write Routine in Assembly

The routine `nvm_write_page` writes one 128-byte page. It expects:

- **Z** (R31:R30): mapped flash data address of the target page: `0x8000 + flash_byte_address` (must be 128-byte aligned).
- **X** (R27:R26): data address of the 128-byte source buffer in SRAM.

It clobbers R16, R17, R18 (scratch) and leaves Z past the end of the written page. The source buffer pointer X is incremented and left past the last byte.

Leaving both pointers advanced is a small design choice, but it makes the routine easier to chain when writing consecutive pages from a stream.

17.6.1 Register Allocation

Z (R31:R30)	mapped flash destination – page base address, then advanced
X (R27:R26)	SRAM source pointer – 128-byte buffer of new content
R16	scratch: NVMCTRL command, status poll byte
R17	loop counter (128 iterations)
R18	data byte being copied into page buffer

17.6.2 Full Listing: src/nvmctrl_write.S

See src/nvmctrl_write.S for the complete annotated source. Key routines:

```
/*
 * nvm_wait – spin until NVMCTRL_STATUS.FBUSY = 0
 *
 * Clobbers: R16
 * Poll loop while busy: LDS (3) + SBRC (1) + RJMP (2) = 6 cycles per poll
 */
nvm_wait:
    lds    r16, NVMCTRL_STATUS          /* 3 cy: load status          */
    sbrc   r16, NVMCTRL_FBUSY_bp        /* 1 cy: skip next if FBUSY=0 */
    rjmp   nvm_wait                    /* 2 cy: still busy, loop     */
    ret                                     /* 4 cy                      */
```

SBRC Rd, b — Skip if Bit in Register is Cleared. Takes 1 cycle when not skipping, 2 when skipping (pipeline flush). NVMCTRL_FBUSY_bp = 0 (bit position 0 of NVMCTRL_STATUS).

```
/*
 * nvm_write_page – write 128-byte SRAM buffer to one flash page
 *
 * Entry:  Z = mapped flash data address (0x8000 + flash byte address)
 *         and 128-byte aligned
 *         X = SRAM pointer to 128-byte source buffer
 * Exit:   Z advanced by 128, X advanced by 128
 * Clobbers: R16, R17, R18
 */
nvm_write_page:
    /* Step 1: wait for any previous NVM operation */
    rcall nvm_wait                /* wait for FBUSY to clear */

    /* Step 2: clear previous write error (WRERROR is W1C) */
    ldi    r16, NVMCTRL_WRError_bm /* r16 = 0x04 */
    sts    NVMCTRL_STATUS, r16

    /* Step 3: fill page buffer from SRAM source */
    ldi    r17, 128                /* loop counter = page size */

.Lfill_loop:
    ld     r18, X+                 /* 2 cy: load byte from SRAM */
    st     Z+, r18                 /* write byte to flash page buffer */
    dec    r17                    /* 1 cy: decrement counter */
    brne   .Lfill_loop            /* 2 cy taken / 1 cy fall-through */
    /* Mapped-Flash page-buffer writes are NVM accesses; do not use SRAM ST
       timing for exact wall-clock estimates. */

    /* Step 4: rewind Z to page start for PAGEERASEWRITE */
    subi   ZL, lo8(128)           /* Z -= 128 */
    sbci   ZH, hi8(128)

    /* Step 5: CCP-protected PAGEERASEWRITE command */
    cli                                     /* 1 cy: disable interrupts */
    ldi    r16, NVMCTRL_CMD_PAGEERASEWRITE_gc /* r16 = 0x03 */
    ldi    r17, CCP_SPM_gc        /* r17 = 0x9D */
    sts    CPU_CCP, r17           /* 2 cy: unlock CCP window (cycles 1-2) */
    sts    NVMCTRL_CTRLA, r16    /* 2 cy: issue command (cycles 3-4) */
    sei                                     /* 1 cy: re-enable interrupts */

    /* Step 6: wait for completion */
    rcall nvm_wait
```



```

├─ init: stack, USART0 (PB2 TX, PB3 RX, 9600 baud at 3.333 MHz)
├─ LED0 (PA3) on → indicates update mode active
├─ send banner "SU\r\n"
└─ record_loop:
    ├─ wait for ':' (discards noise before start code)
    ├─ recv_hex_byte × 4: byte_count, addr_hi, addr_lo, type
    ├─ if type = 0x01 (EOF): send "OK\r\n", wait FBUSY, IJMP to 0x0000
    ├─ if type = 0x00 (data):
    │   ├─ set Z = record address (flash byte address)
    │   ├─ receive LL data bytes → SRAM page buffer
    │   ├─ if buffer >= 128 bytes: call nvm_write_page
    │   └─ verify checksum → send '.' (ACK) or '!' (NACK)
    └─ repeat

```

17.7.3 Register Allocation for selfupdate.S

Z (R31:R30)	flash write address (maintained across records)
X (R27:R26)	SRAM page buffer pointer (buf_base .. buf_base+127)
R16	scratch / UART byte / NVMCTRL command
R17	checksum accumulator (zeroed at each record start)
R18	byte count from record header
R19	record type (00 or 01)
R24	address low byte from record header
R25	address high byte from record header
R20, R21	scratch in helpers
R15	saved high nibble in recv_hex_byte

17.7.4 USART0 Initialisation at 3.333 MHz

The ATtiny3217's default clock is 3.333 MHz (20 MHz internal oscillator with the default /6 prescaler in CLKCTRL). The USART0 fractional baud register for 9600 baud:

```

BAUD_REG = (f_cpu × 64) / (16 × baud)
          = (3333333 × 64) / (16 × 9600)
          = 213333312 / 153600
          = 1389 (0x056D)

```

```

.equ BAUD_REG, 1389                                /* 9600 baud @ 3.333 MHz, error < 0.01% */

usart_init:
    sbi    VPORTB_DIR, 2                            /* PB2 = TXD output */
    ldi    r16, lo8(BAUD_REG)
    sts    USART0_BAUDL, r16
    ldi    r16, hi8(BAUD_REG)
    sts    USART0_BAUDH, r16
    ldi    r16, USART_TXEN_bm | USART_RXEN_bm
    sts    USART0_CTRLB, r16                        /* enable TX and RX */
    ret
/* CTRLC defaults to 8N1 – no write needed */

```

Note: VPORTB_DIR is at I/O address 0x04 (within sbi/cbi range) — we use sbi rather than sts to set the single bit without disturbing other PORTB direction bits.

17.7.5 Full Listing: src/selfupdate.S

The complete source is at src/selfupdate.S. The core hex-receive loop and page-write call are shown below for reference:

```

.Ldata_loop:
    rcall recv_hex_byte      /* receive next data byte from HEX record */
    st      X+, r16          /* store in SRAM page buffer */
    add     r17, r16         /* accumulate checksum */
    dec     r18              /* decrement byte count */

    /* check if page buffer is full (128 bytes written) */
    mov     r20, XL
    andi    r20, 0x7F        /* XL & 127 – non-zero means not at boundary */
    brne    .Lcheck_done     /* not at 128-byte boundary: continue */

    /* page buffer is full: save X, write page, restore */
    movw    r20, XL          /* save X */
    rcall   nvm_write_page   /* writes page, advances Z by 128 */
    movw    XL, r20          /* restore X */

.Lcheck_done:
    tst     r18
    brne    .Ldata_loop     /* more bytes in this record */

```

17.7.6 Jumping to the Updated Application

After the EOF record is received and the last page is written:

```

.Leof:
    rcall   recv_hex_byte    /* consume EOF checksum byte (not verified) */
    rcall   nvm_wait         /* ensure final write is complete */
    /* turn off LED0 (PA3) to signal update complete */
    sbi     VPORTA_OUT, 3    /* PA3 high = LED off (active-low) */
    /* send "OK\r\n" */
    ldi     r16, 'O'
    rcall   usart_tx
    ldi     r16, 'K'
    rcall   usart_tx
    ldi     r16, '\r'
    rcall   usart_tx
    ldi     r16, '\n'
    rcall   usart_tx
    /* jump to new application */
    clr     ZL
    clr     ZH
    ijmp                                /* PC ← Z = 0x0000 (new firmware entry) */

```

IJMP (Indirect Jump): sets the program counter to Z (2 cycles). With Z = 0, execution continues from the reset vector of the freshly written firmware.

Conceptually, that final jump is just a handoff. The updater finishes its work, then transfers control to the program it just installed.

17.7.7 Uploading Firmware

Use the existing `src/send_hex.py` helper, or any serial terminal. The selfupdate firmware sends `SU\r\n` instead of `BL\r\n`:

The helper rewrites ordinary Intel HEX input into 128-byte page records before sending. That matches the small updater's assumptions: each data record starts on a page boundary, and a short final page is padded with `0xFF`.


```
# Flash the self-update firmware via the Curiosity Nano debugger (one-time):
pymcuprog write -d attiny3217 -t uart -f selfupdate.hex

# Reset the board – LED0 turns on (active-low, so PA3 = 0 = LED on)
# Then immediately send application firmware:
python3 src/send_hex.py /dev/ttyUSB0 app.hex
```

Expected output:

Waiting for self-updater on /dev/ttyUSB0...

Self-updater ready

record 1/18 OK

record 2/18 OK

...

record 18/18 OK

Upload complete. Application starting.

LED0 turns off when the update completes, and the new firmware starts.

17.8 BOOTEND Fuse: Protecting a Bootloader Region

17.8.1 The Problem

The self-update firmware in Section 16.6 writes to any flash address, including address 0x0000 — which is where the self-update firmware itself lives. If the update process is interrupted after erasing page 0 but before writing the new content, the device is bricked (no valid firmware at the reset vector). It also means buggy application firmware could accidentally overwrite the updater.

17.8.2 BOOTEND: Defining a Protected Region

The ATtiny3217 provides the **BOOTEND** and **APPEND** fuses. These divide flash into BOOT, APPCODE, and APPDATA sections in 256-byte units. For the simple updater layout, set BOOTEND to reserve space for the updater and leave APPEND = 0 so the rest of flash is APPCODE:

FUSES.BOOTEND value, with APPEND = 0:

0x01 – 256 bytes BOOT (pages 0–1)

0x02 – 512 bytes BOOT (pages 0–3)

0x04 – 1024 bytes BOOT (pages 0–7)

...

Boot section: 0x0000 .. (BOOTEND × 256 - 1)


APPCODE section: (BOOTEND × 256) .. 0x7FFF

With BOOTEND = 0x04 (1 KB boot section), the self-update firmware occupies pages 0–7 (addresses 0x0000–0x03FF). The application firmware lives at 0x0400 onward. The NVMCTRL enforces section privileges. BOOT code can write APPCODE and APPDATA. APPCODE can write APPDATA. APPDATA cannot write flash or EEPROM. The CPU cannot write the BOOT section.

This is the safety boundary that turns a self-programming demo into something you can trust more. Without it, the updater can erase the very code that is needed to recover from a failed update.

17.8.3 Flash Layout With BOOTEND = 0x04

ATtiny3217 flash (32 KB) with BOOTEND = 0x04



17.8.4 Application Entry Point

When `BOOTEND = 0x04`, the application vector table starts at **0x0400**. The hardware Reset vector still starts execution at `0x0000`; there is no `BOOTRST` fuse on this device. After a successful update, the self-update firmware must jump to the application's entry at `0x0400` rather than back to its own reset entry at `0x0000`:

```
/* With BOOTEND = 0x04: application starts at 0x0400 */
.equ APP_START, 0x0400

.Leof:
    rcall recv_hex_byte
    rcall nvm_wait
    ldi   ZL, lo8(APP_START)
    ldi   ZH, hi8(APP_START)
    jmp  /* jump to application section start */
```

The application must be built with its `.text` section starting at `0x0400`:

```
avr-gcc -mmcu=attiny3217 -nostartfiles \
    -Wl,--section-start=.text=0x0400 \
    -o app.elf app.S
```

If the bootloader uses interrupts while running from the `BOOT` section, it must set `CPUINT.CTRLA.IVSEL` to move the interrupt vector table to the start of Flash, then restore the application vector location before handing control to the application. The reset vector itself remains at `0x0000`.

17.8.5 Run-Time Locks: BOOTLOCK and APCWP

`NVMCTRL_CTRLB` has two run-time lock bits that remain active until reset:

- `APCWP` prevents further writes to `APPCODE`.
- `BOOTLOCK` prevents reads and execution from `BOOT`.

Setting `APCWP` after the update completes provides defense-in-depth:

```
/* After successful update, lock application section against further writes */
ldi   r16, 0x01 /* APCWP bit */
sts   NVMCTRL_CTRLB, r16
```

This is cleared on every reset, so it does not permanently lock the device.

Think of `APCWP` as a temporary guardrail for the current boot session, not as a replacement for fuse-based protection.

17.9 Building and Flashing

17.9.1 Building nvmctrl_write.S (standalone test)

```
# Assemble the NVM write routine alone (produces an ELF for inspection)
avr-gcc -mmcu=attiny3217 -nostartfiles -nodefaultlibs \
    -o nvmctrl_write.elf src/nvmctrl_write.S

# Check sections and disassemble to verify instruction encoding
avr-objdump -h nvmctrl_write.elf
avr-objdump -d nvmctrl_write.elf
```

17.9.2 Building selfupdate.S

```
# Assemble the self-update firmware
# Without BOOTEND protection: starts at 0x0000
avr-gcc -mmcu=attiny3217 -nostartfiles -nodefaultlibs \
    -o selfupdate.elf src/selfupdate.S

# With BOOTEND=0x04: starts at 0x0000, app section starts at 0x0400
# The selfupdate firmware itself is the same; linker placement is 0x0000
avr-gcc -mmcu=attiny3217 -nostartfiles -nodefaultlibs \
    -o selfupdate.elf src/selfupdate.S

# Size check – .text must fit within chosen boot section (1 KB = BOOTEND 0x04)
avr-objdump -h selfupdate.elf
```

Expected section output excerpt:

Idx	Name	Size	VMA
1	.text	000001c6	00000000
2	.bss	00000080	00803800

0x1c6 is 454 bytes of flash code, which fits comfortably within the 1 KB BOOT section. The 128-byte .bss allocation is the SRAM page buffer.

17.9.3 Converting to HEX

```
avr-objcopy -O ihex selfupdate.elf selfupdate.hex
avr-objcopy -O ihex app.elf app.hex
```

17.9.4 Flashing with pymcuprog

```
# Program selfupdate firmware via UPDI (requires Curiosity Nano USB connection)
pymcuprog write -d attiny3217 -t uart -f selfupdate.hex

# Set BOOTEND fuse to 0x04 (1 KB boot section)
# WARNING: verify the exact pymcuprog fuse syntax for your installed version.
# BOOTEND is fuse byte 8; write value 0x04.

# Read back fuse to verify
pymcuprog read -d attiny3217 -t uart
```

Warning: BOOTEND is in fuse byte 8 on ATtiny3217. An incorrect value can make the device appear to run no code. Always read back fuses after writing. UPDI can always recover the device regardless of fuse settings (unlike ISP on ATmega parts where the BOOTRST fuse can complicate recovery).

That recovery path is important operationally: experimenting with self-update code is much less risky when you know UPDI can still get you back in.

17.9.5 Verifying the Disassembly

After building, verify the CCP sequence is encoded correctly:

```
avr-objdump -d selfupdate.elf | grep -A5 "CPU_CCP\|0034"
```

Look for two back-to-back sts instructions with addresses 0x0034 (CCP) and 0x1000 (NVMCTRL_CTRLA). No instruction should appear between them in the listing.

17.9.6 Building an Application for the Protected Layout

```
# Application must be linked at 0x0400 when BOOTEND = 0x04
avr-gcc -mmcu=attiny3217 -nostartfiles -nodefaultlibs \
    -Wl,--section-start=.text=0x0400 \
    -o app.elf app.S
avr-objcopy -O ihex app.elf app.hex
```

```
# Send via self-update firmware:
python3 src/send_hex.py /dev/ttyUSB0 app.hex
```

17.10 Common Pitfalls

17.10.1 Pitfall 1 — CCP Window Violated

Symptom: NVMCTRL_CTRLA write appears to succeed (no error) but no flash write occurs; WRError is set in STATUS.

Cause: An instruction was inserted between STS CPU_CCP, r16 and STS NVMCTRL_CTRLA, r16, or an interrupt fired inside the window and consumed cycles.

Fix: Use CLI before the CCP sequence and SEI immediately after the protected write. Ensure no instruction sits between the two STS calls.

```
/* WRONG – ldi between STS instructions wastes a cycle inside the window */
sts CPU_CCP, r16
ldi r16, 0x03 /* ← this is inside the window, shrinks margin */
sts NVMCTRL_CTRLA, r16

/* CORRECT – load command value BEFORE writing CCP key */
ldi r16, 0x03 /* load command first (outside window) */
...
ldi r17, CCP_SPM_gc
sts CPU_CCP, r17 /* open window */
sts NVMCTRL_CTRLA, r16 /* issue command (within window) */
```

Note the subtlety: if r16 holds the CCP key when you write CPU_CCP, you must reload r16 with the command before the protected write, or use two registers. The corrected pattern above uses two registers to keep the command value ready before opening the CCP window.

That is a good example of why “almost correct” is not good enough here. A sequence can look visually right and still fail because one register was reused at the wrong moment.

17.10.2 Pitfall 2 — Page Buffer Not Aligned

Symptom: The first few bytes of a written page are corrupted or contain old flash content.

Cause: The mapped flash address passed to `nvm_write_page` is not 128-byte aligned. The page buffer fill loop always fills from byte 0 to byte 127 of the chosen page. If `Z` points into the middle of a page, the first `ST Z+, Rn` writes to a mid-page buffer slot. The earlier slots contain the current page-buffer state, not an automatic copy of the old flash page.

Fix: Always pass a 128-byte aligned address:

```
/* Align mapped Z to 128-byte page boundary: clear bits 6:0 */
andi ZL, 0x80 /* mask: 1000 0000 – clear bits 6:0 */
```

Page alignment bugs are especially deceptive because most of the write may look correct. Only the offsets near the page boundary reveal the mistake.

17.10.3 Pitfall 3 — Writing to Flash During Interrupt

Symptom: Application ISR fires while `nvm_write_page` is between the page buffer fill and the `PAGEERASEWRITE` command. The ISR reads stale flash content via `LPM` from the page being written.

Cause: The page buffer is not yet committed. If an ISR attempts to read instruction or data from flash in the middle of a `PAGEERASEWRITE`, it reads the erased (0xFF) state during the erase phase.

Fix: Disable global interrupts before the entire write sequence and re-enable after:

```
cli
rcall nvm_write_page
sei
```

The `nvm_write_page` routine already disables interrupts around the CCP write, but a complete `CLI/SEI` around the entire page write is safer if you have ISRs active.

17.10.4 Pitfall 4 — Forgetting to Wait for FBUSY Before Each Command

Symptom: Random corruption of flash pages, or `WRERROR` flag set frequently.

Cause: A second `PAGEERASEWRITE` command was issued before the first completed. `NVMCTRL_STATUS.FBUSY` must be 0 before issuing any command.

Fix: The `nvm_wait` routine at the start of `nvm_write_page` ensures this. If you issue `NVMCTRL` commands from multiple places in your code, always call `nvm_wait` first.

17.10.5 Pitfall 5 — Writing the Boot Section From Application Code

Symptom: Write to pages 0–(`BOOTEND`-based boundary) triggers `WRERROR` but no actual write.

Cause: With `BOOTEND` set, `NVMCTRL` enforces section privileges. The CPU cannot write the `BOOT` section through `NVMCTRL`.

Fix: This is intentional behaviour — the boot section protection is working correctly. If you need to update the self-update firmware, use `UPDI` with `pymcuprog`.

In other words, `WRERROR` here is a feature, not a malfunction.

17.10.6 Pitfall 6 — UPDI Port Floating (PA0)

Symptom: `avrdude` cannot connect; “timeout communicating with programmer.”

Cause: On the Curiosity Nano, PA0 is used for UPDI. If you configure PA0 as a GPIO output in your firmware, UPDI communication breaks. The on-board on-board debugger pulls PA0 to its own UPDI driver.

Fix: Never configure PA0 as GPIO output. Leave PA0 at its power-on default (UPDI function). If your application needs an extra GPIO and PA0 was used, remap to a different pin.

17.10.7 Pitfall 7 — Baud Rate at Non-Default Clock

Symptom: Garbage characters received; UART communication fails after changing CLKCTRL prescaler.

Cause: The BAUD_REG constant was calculated for 3.333 MHz (default clock). If your firmware changes CLKCTRL_MCLKCTRLB to select a different frequency before calling `usart_init`, the baud rate will be wrong.

Fix: Recalculate BAUD_REG for the actual `f_cpu`, or initialise USART before changing the clock. The self-update firmware deliberately runs at the default clock to avoid this dependency.

17.11 Summary

ATtiny3217 self-programming (AVRMEGA3):

Flash geometry:

32 KB total, 128-byte pages, 256 pages
Data-space flash base address: 0x8000
Page buffer fill: ST Z+, Rn (plain store instruction)

NVMCTRL registers:

NVMCTRL_CTRLA 0x1000 command register
NVMCTRL_STATUS 0x1002 FBUSY (bit 0), WRERROR (bit 2)

Write sequence for one 128-byte page:

1. Poll NVMCTRL_STATUS until FBUSY = 0
2. STS NVMCTRL_STATUS, 0x04 (clear WRERROR; write-one-to-clear)
3. 128 × ST Z+, Rn (fill page buffer at mapped flash address)
4. SUBI/SBCI Z, 128 (rewind Z to page base)
5. CLI
 - STS CPU_CCP, 0x9D (unlock CCP)
 - STS NVMCTRL_CTRLA, 0x03 (PAGEERASEWRITE as next instruction)
 - SEI
6. Poll NVMCTRL_STATUS until FBUSY = 0

CCP (Configuration Change Protection):

Register: 0x0034 (CPU_CCP)
Key for NVMCTRL: 0x9D (CCP_SPM_gc)
Window: next four CPU instructions
Rule: STS CPU_CCP must be immediately followed by STS NVMCTRL_CTRLA

BOOTEND fuse (fuse byte 8):

0x04 = 1 KB boot section (selfupdate.S fits here)
App section starts at BOOTEND × 256

Key instructions used:

SUBI/SBCI – 16-bit pointer adjustment for +/-128

IJMP - indirect jump through Z (2 cy)

Build:

```
avr-gcc -mmcu=attiny3217 -nostartfiles -nodefaultlibs \
-o selfupdate.elf src/selfupdate.S
avr-objcopy -O ihex selfupdate.elf selfupdate.hex
pymcuprog write -d attiny3217 -t uart -f selfupdate.hex
```

17.12 Exercises

1. The `nvm_wait` polling loop costs 6 cycles per iteration while `FBUSY` is set at 3.333 MHz. A `PAGEERASEWRITE` operation on ATtiny3217 takes approximately 2.5 ms. How many loop iterations does `nvm_wait` execute while waiting? At 3.333 MHz, what is the duration of each iteration in microseconds? Is a timer-based timeout (to detect a stuck NVM controller) practical within the same loop?
2. The CCP window is four CPU instructions. Why does the source place `STS NVMCTRL_CTRLA`, `r16` immediately after `STS CPU_CCP`, `r17` even though the hardware allows a small window?
3. The `nvm_write_page` routine fills the page buffer with 128 iterations of `LD r18, X+ + ST Z+, r18`. Why is the exact cycle count for the `ST` portion not the same as a normal SRAM store? Estimate the loop overhead, then compare it with the millisecond-scale NVM erase/write time.
4. The self-update firmware in `selfupdate.S` does not verify the EOF record's checksum. Add checksum verification: if the EOF checksum is incorrect, send '!', turn LED0 on (it was turned off on successful EOF), and loop back to `record_loop` without jumping to the application. What register holds the checksum accumulator at the point of the EOF handler, and what condition do you test?
5. The current `selfupdate.S` jumps to address `0x0000` on EOF (the reset vector). If `B00TEND = 0x04`, the application section begins at `0x0400`. Modify the EOF handler to: (a) check that at least one page in the application section was written (track the highest written address in a register), (b) jump to `0x0400` instead of `0x0000`. Show the assembly changes.
6. Implement a **write-verify** step: after writing each page with `nvm_write_page`, read back the 128 bytes using `LPM Z+, r16` and compare against the SRAM buffer. If any byte mismatches, retransmit '!' (NACK) and re-request the record. Where in the `selfupdate.S` main loop would you add this? What additional register or memory is needed?
7. Add a page-address guard that rejects any incoming data record below `B00TEND * 256`. On rejection, send '!' and discard the record data. Which two header bytes do you compare, and where should the check occur?
8. Using `avr-objdump -d selfupdate.elf`, locate the two `STS` instructions that form the CCP sequence. Record their addresses. What is the byte offset between them? Confirm that no other instruction appears at an address between them. Now deliberately insert a `NOP` between them in the source and rebuild — does `avr-objdump` show the `NOP` between the two `STS` instructions? Describe the hardware consequence of this change.

Next: Chapter 17 — Bit Manipulation and Fixed-Point Arithmetic

Chapter 18

Bit Manipulation & Fixed-Point Math

AVR has no hardware divider and no floating-point unit. The ATtiny3217 does include the AVR hardware multiplier, but shift-and-add multiplication remains useful for understanding the machine model and for AVR parts that omit MUL. When computation must be fast and deterministic — sensor scaling, display formatting, signal processing — the answers live in bit manipulation and fixed-point arithmetic.

This chapter covers:

1. **Shift-and-add multiply** — multiplication using only shifts and adds; useful as a portable fallback and as a teaching algorithm.
2. **Fixed-point Q8.8 arithmetic** — representing and operating on fractional numbers in 16-bit register pairs.
3. **BCD arithmetic** — binary-coded decimal for human-readable display.
4. **Practical example** — convert a 16-bit binary sensor reading to BCD for a four-digit 7-segment display, using all three techniques.

The common thread is that all of these techniques trade mathematical elegance for predictability. On a small MCU, that is usually a good trade.

18.1 Shift Operations

AVR provides five shift/rotate instructions. All take a single register operand and execute in **1 cycle**.

Instruction	Operation	Flags
LSL Rd	$C \leftarrow Rd7$; $Rd \leftarrow Rd \ll 1$; $Rd0 = 0$	C, Z, N, V, H
LSR Rd	$Rd7 = 0$; $Rd \leftarrow Rd \gg 1$; $C \leftarrow Rd0$	C, Z, N, V
ASR Rd	$Rd7$ unchanged; $Rd \leftarrow Rd \gg 1$; $C \leftarrow Rd0$	C, Z, N, V
ROL Rd	$C \leftarrow Rd7$; $Rd \leftarrow Rd \ll 1$; $Rd0 = \text{old } C$	C, Z, N, V, H
ROR Rd	$Rd7 = \text{old } C$; $Rd \leftarrow Rd \gg 1$; $C \leftarrow Rd0$	C, Z, N, V

LSL/LSR: logical shift, zero fill. LSL Rd multiplies by 2; LSR Rd divides by 2 (unsigned).

ASR: arithmetic shift right — sign-extends during right shift, preserving the sign of a signed value. ASR Rd divides a signed byte by 2 (rounds toward negative infinity).

ROL/ROR: rotate through carry. The carry flag acts as a 9th bit. This is how multi-byte shifts are chained:

That is the key idea behind most “wide” arithmetic on AVR. You do not get a 16-bit or 32-bit shift instruction; you build one byte at a time and let carry connect the pieces.

```
/* Logical left shift a 16-bit value in R25:R24 by 1 bit */
lsl  r24      /* shift low byte left; bit 7 → C */
rol  r25      /* shift high byte left; C → bit 0 */
```

Multi-byte right shift (logical):

```
/* Logical right shift R25:R24 by 1 bit */
lsr  r25      /* shift high byte right; bit 0 → C */
ror  r24      /* shift low byte right; C → bit 7 */
```

18.1.1 Register Pair Diagram — 16-bit LSL

Before: R25[b7..b0] R24[b7..b0] └ (bit 0 = 0 after LSL)

After:

R24: b6 b5 b4 b3 b2 b1 b0 0 ← LSL r24 (b7 of R24 → C)
 R25: b6 b5 b4 b3 b2 b1 b0 C ← ROL r25 (C → bit 0 of R25)
 C = old R25[b7]

18.2 Shift-and-Add Multiplication

The MUL instruction ($8 \times 8 \rightarrow 16$ unsigned) exists on ATtiny3217 and many other AVR devices. Some smaller or older AVR parts omit hardware multiply, so the shift-and-add algorithm is still worth knowing.

The shift-and-add algorithm works on any AVR: for each bit in the multiplier, if the bit is set, add the multiplicand (shifted appropriately) to the accumulator.

18.2.1 Algorithm — Unsigned $8 \times 8 \rightarrow 16$

```
product = 0
for bit = 0 to 7:
    if multiplier[bit] == 1:
        product += multiplicand << bit
    (equivalently: shift multiplicand left each iteration,
     add if LSB of multiplier is 1, then shift multiplier right)
```

18.2.2 Implementation

Instruction	Operands	Cycles	Flags
LSL Rd	Rd	1	C,Z,N,V,H
LSR Rd	Rd	1	C,Z,N,V
ROL Rd	Rd	1	C,Z,N,V,H
ROR Rd	Rd	1	C,Z,N,V
ADD Rd,Rr	Rd, Rr	1	C,Z,N,V,H
ADC Rd,Rr	Rd, Rr	1	C,Z,N,V,H
BRCC k	k (−64..63)	1/2	—
DEC Rd	Rd	1	Z,N,V

```

/*
 * mul8u – unsigned 8×8 → 16 multiply (shift-and-add)
 *
 * Entry:  R16 = multiplicand
 *         R17 = multiplier
 * Exit:   R25:R24 = R16 × R17 (unsigned 16-bit product)
 * Clobbers: R16, R17, R18, R20
 * Cycles: 71 + popcount(multiplier), including setup and RET on AVRxt
 */
mul8u:
    clr    r24                /* product = 0 */
    clr    r25
    clr    r20                /* multiplicand high byte = 0 */
    ldi    r18, 8             /* 8 bits to process */
.Lmul8u_loop:
    lsr    r17                /* multiplier >>= 1; old bit 0 → C */
    brcc   .Lmul8u_skip
    add    r24, r16           /* product_lo += multiplicand_lo */
    adc    r25, r20           /* product_hi += multiplicand_hi + carry */
.Lmul8u_skip:
    lsl    r16                /* multiplicand_lo <<= 1; old bit 7 → C */
    rol    r20                /* multiplicand_hi <<= 1; C → bit 0 */
    dec    r18
    brne   .Lmul8u_loop
    ret

```

The key detail is that the multiplicand is tracked as a 16-bit value (R20:R16), so left shifts preserve the carry-out from the low byte.

That is what makes the routine correct for products larger than 255. If you shifted only the low byte, the upper half of the multiplicand would vanish.

Trace — 3×5 (i.e. $0b00000011 \times 0b00000101 = 15$):

r16=3, r17=5, r20=0, r24=0, r25=0

Iter 1: lsr r17 → r17=2, C=1 (bit 0 of 5)
 add r24, r16 → r24=3; adc r25, r20 → r25=0
 lsl r16 → r16=6, C=0; rol r20 → r20=0

Iter 2: lsr r17 → r17=1, C=0 (bit 0 of 2)
 skip
 lsl r16 → r16=12; rol r20 → r20=0

Iter 3: lsr r17 → r17=0, C=1 (bit 0 of 1)
 add r24, r16 → r24=3+12=15; adc r25, r20 → r25=0
 lsl r16 → r16=24; rol r20 → r20=0

Iters 4-8: r17=0, C=0 each time → all skip

Result: r25:r24 = 0x000F = 15 OK

18.2.3 Signed Multiply — mul8s (8×8 → 16 signed)

Signed multiply extends naturally: run the unsigned loop for 7 bits, then treat the MSB of the multiplier as negative (two's complement weight -128):

```

/*
 * mul8s – signed 8×8 → 16 multiply
 *
 * Entry:  R16 = multiplicand (signed)

```

```

*      R17 = multiplier    (signed)
* Exit:  R25:R24 = R16 × R17 (signed 16-bit product)
* Clobbers: R16, R17, R18, R20
*/
mul8s:
    clr    r24
    clr    r25
    /* sign-extend R16 into R20 */
    mov    r20, r16
    lsl    r20          /* old bit 7 → C */
    sbc    r20, r20      /* r20 = 0xFF if negative, 0x00 if positive */
    ldi    r18, 7        /* process 7 unsigned bits */
.Lmul8s_loop:
    lsr    r17
    brcc   .Lmul8s_skip
    add    r24, r16
    adc    r25, r20
.Lmul8s_skip:
    lsl    r16
    rol    r20
    dec    r18
    brne   .Lmul8s_loop
    /* bit 7 of multiplier: subtract instead of add (two's complement) */
    lsr    r17          /* bit 7 → C */
    brcc   .Lmul8s_done
    sub    r24, r16      /* product -= (multiplicand << 7) */
    sbc    r25, r20
.Lmul8s_done:
    ret

```

SBC Rd, Rr subtracts Rr and the carry flag from Rd. Here, after `lsr r17`, C holds the sign bit of the original multiplier. If set, we subtract the current (shifted) multiplicand — the two’s complement correction for the negative weight of bit 7.

This is one of those places where two’s complement stops being abstract. The top bit is not just “another positive weight”; it contributes a negative term, so the final step has to subtract instead of add.


18.3 Fixed-Point Q8.8 Arithmetic

Fixed-point represents a real number as an integer scaled by a power of two. **Q8.8** uses a 16-bit register pair where the high byte is the integer part and the low byte is the fractional part:

The advantage is that the CPU still performs ordinary integer arithmetic. The “fractional” meaning comes from how you interpret the bits, not from special hardware support.

Q8.8 value = register_pair / 256

Bit layout (R25:R24 example):



Range: -128.0 to +127.99609375 (signed)
 0.0 to 255.99609375 (unsigned)

Resolution: $1/256 \approx 0.0039$

18.3.1 Converting to Q8.8

```
/* Integer 42 → Q8.8: shift left 8 bits (multiply by 256) */
ldi r24, 0 /* fractional part = 0 */
ldi r25, 42 /* integer part = 42 */
/* R25:R24 = 0x2A00 = 10752 raw = 42.0 in Q8.8 */

/* Fractional constant 0.5 → Q8.8 = 128 = 0x0080 */
ldi r24, 128
clr r25
/* R25:R24 = 0x0080 = 0.5 in Q8.8 */

/* Convert from Q8.8 back to integer (truncate) */
mov r16, r25 /* integer part = high byte */
```

18.3.2 Q8.8 Addition

Identical to 16-bit integer addition — no adjustment needed:

That simplicity is one of the main reasons fixed-point is attractive. Once two values share the same scaling, addition and subtraction are just normal integer operations.

```
/* R25:R24 += R23:R22 (Q8.8 + Q8.8) */
add r24, r22
adc r25, r23
```

18.3.3 Q8.8 Multiplication (Q8.8 × Q8.8 → Q8.8)

Two Q8.8 values multiplied give a Q16.16 result in a 32-bit space. To get a Q8.8 result, extract bits [23:8] of the 32-bit product — effectively shift right 8 bits:

```
A = a15..a0 (Q8.8)   B = b15..b0 (Q8.8)
A × B = 32-bit Q16.16
Q8.8 result = (A × B) >> 8 = bits [23:8] of the 32-bit product
```

In practice, build the 32-bit product as four 8×8 partial products:

```
A = AH:AL (high byte : low byte)
B = BH:BL
```

```
A × B = AH×BH (weight 216) + AH×BL (weight 28)
        + AL×BH (weight 28) + AL×BL (weight 20)
```

```
Q8.8 result = shift right 8 → take bits [23:8]:
bits [15:8] of (AH×BL + AL×BH) + AH×BH (lower byte) + carry from AL×BL
```

The ATtiny3217 is an AVRxt device with the hardware multiply instructions available; Microchip's instruction set manual lists MUL, MULS, and MULSU as two-cycle operations that write the 16-bit product to R1:R0. The routine below deliberately uses mul8u from Section 17.2 so the byte decomposition is visible and so the same listing still works on smaller AVR devices that omit hardware multiply.

The bookkeeping looks tedious, but the idea is straightforward: break the 16-bit multiply into byte-sized pieces, then keep only the middle 16 bits that correspond to the Q8.8 result.

```
/*
 * mulq88 — unsigned Q8.8 × Q8.8 → Q8.8 using mul8u
```

```

*
* Entry:  R25:R24 = A (Q8.8),  R23:R22 = B (Q8.8)
* Exit:   R25:R24 = A × B (Q8.8, truncated)
*/
mulq88:
    push    r12
    push    r13
    push    r14
    push    r15

    mov     r12, r24    /* AL */
    mov     r13, r25    /* AH */
    mov     r14, r22    /* BL */
    mov     r15, r23    /* BH */

    /* AL × BL */
    mov     r16, r12
    mov     r17, r14
    rcall   mul8u
    mov     r21, r25    /* result low byte starts with high(AL×BL) */
    clr     r22          /* result high byte */

    /* AL × BH */
    mov     r16, r12
    mov     r17, r15
    rcall   mul8u
    add     r21, r24
    adc     r22, r25

    /* AH × BL */
    mov     r16, r13
    mov     r17, r14
    rcall   mul8u
    add     r21, r24
    adc     r22, r25

    /* AH × BH */
    mov     r16, r13
    mov     r17, r15
    rcall   mul8u
    add     r22, r24    /* low byte of AH×BH contributes at result high byte */

    mov     r24, r21
    mov     r25, r22
    pop     r15
    pop     r14
    pop     r13
    pop     r12
    ret

```

Full listing in src/fixedpoint.S.

18.4 BCD Arithmetic

Binary-Coded Decimal stores each decimal digit as a 4-bit nibble. Two formats:

Unpacked BCD: one digit per byte (low nibble used, high nibble = 0)
 Decimal 47 → byte 0x04, byte 0x07

Packed BCD: two digits per byte
 Decimal 47 → single byte 0x47

18.4.1 Decimal Adjust on AVR

Unlike some 8-bit CPUs, AVR does **not** provide a general DAA instruction for packed-BCD correction. After a binary ADD/ADC, packed BCD must be corrected manually using the half-carry and carry flags.

18.4.2 H Flag and BCD Correction

The H (half-carry) flag is set when a carry propagates from bit 3 to bit 4 during addition — i.e., when the low nibble overflows. This is exactly the condition that requires BCD correction:

After binary ADD of two packed BCD bytes:

If low nibble > 9 OR H flag set: add 0x06 to correct low nibble

If high nibble > 9 OR C flag set: add 0x60 to correct high nibble

```

/*
 * bcd_add – add two single packed BCD bytes
 *
 * Entry:  R16 = BCD byte A (e.g. 0x47 for decimal 47)
 *         R17 = BCD byte B
 * Exit:   R16 = BCD sum (low byte), C = 1 if result > 99
 * Clobbers: R18, R19, R20
 */
bcd_add:
    add    r16, r17        /* binary add */
    clr    r18              /* correction byte: 0x00, 0x06, 0x60, or 0x66 */
    clr    r19              /* decimal carry-out marker */

    /* Capture H and C from the original ADD before later tests change SREG. */
    brhc   .Lbcd_no_h
    ori    r18, 0x06

.Lbcd_no_h:
    brcc   .Lbcd_no_c
    ori    r18, 0x60
    ldi    r19, 1

.Lbcd_no_c:

    /* Low nibble > 9 also needs +0x06. */
    mov    r20, r16
    andi   r20, 0x0F
    cpi    r20, 10
    brlo   .Lbcd_check_high
    ori    r18, 0x06

.Lbcd_check_high:
    /* High nibble > 9 also needs +0x60 and a decimal carry out. */
    mov    r20, r16
    andi   r20, 0xF0
    cpi    r20, 0xA0
    brlo   .Lbcd_done
    ori    r18, 0x60
    ldi    r19, 1

```

```

.Lbcd_done:
    add    r16, r18
    tst    r19
    breq   .Lbcd_no_decimal_carry
    sec
    ret
.Lbcd_no_decimal_carry:
    clc
    ret

```

BRHS (Branch if Half-carry Set) and BRHC (Branch if Half-carry Clear) branch directly on the H flag.

That matters because packed BCD correction is really nibble arithmetic. The H flag tells you whether the low decimal digit overflowed even when the full byte did not.

Instruction	Branch condition	Cycles
BRHS k	H = 1	1/2
BRHC k	H = 0	1/2

18.5 Binary to BCD Conversion

Converting a multi-digit binary number to BCD for display is a recurring task. The **double-dabble** (shift-and-add-3) algorithm converts an N-bit binary number to packed BCD in N iterations, requiring no division:

Algorithm:

```

Place binary value in a shift register alongside N/4 BCD digits (all zero).
Repeat N times:
    For each BCD digit: if digit > 4, add 3 to that digit.
    Shift the entire combined register left by 1 bit.
After N iterations, BCD digits contain the decimal representation.

```

The “add 3” step ensures that a BCD digit that is about to overflow (≥ 5) will produce the correct carry into the next BCD digit after the shift (since shifting left doubles, and $5 \times 2 = 10$ wraps in BCD; pre-adding 3 makes $(5+3) \times 2 = 16$, which carries correctly as 6 in the upper digit and 0 in the lower).

If double-dabble feels magical at first, focus on that one fact: digits 5..9 need help before the shift so that a binary left shift turns into a valid decimal carry.

18.5.1 16-bit Binary to 5-Digit BCD

A 16-bit value (0–65535) needs 5 BCD digits (0–9 each), packed into 3 bytes (two packed + one partial):

Packed BCD storage: [tens-thousands][thousands|hundreds][tens|ones]
 Example: 65535 → 0x06, 0x55, 0x35

```

/*
 * bin16_to_bcd – convert 16-bit unsigned binary to 5 packed BCD digits
 *
 * Entry:  R25:R24 = 16-bit value (0–65535)
 * Exit:   R21:R20:R19 = packed BCD (5 digits)
 *         R21[3:0] = ten-thousands digit
 *         R20[7:4] = thousands digit
 *         R20[3:0] = hundreds digit

```

```

*          R19[7:4] = tens digit
*          R19[3:0] = ones digit
* Clobbers: R16, R17, R22
*/
bin16_to_bcd:
    clr    r19          /* BCD digits [tens|ones] = 0 */
    clr    r20          /* BCD digits [thousands|hundreds] = 0 */
    clr    r21          /* BCD digit [ten-thousands] = 0 */
    ldi    r16, 16      /* loop counter: 16 bits to process */

.Lbcd_loop:
    /* For each BCD nibble: if > 4, add 3 */
    rcall  bcd_adjust   /* adjust R21:R20:R19 */

    /* Shift binary value and BCD digits left by 1 */
    lsl    r24          /* binary low byte: bit 7 → C */
    rol    r25          /* binary high byte: C → bit 0; bit 7 → C */
    rol    r19          /* BCD byte 0: C in, bit 7 → C */
    rol    r20          /* BCD byte 1: C in, bit 7 → C */
    rol    r21          /* BCD byte 2: C in (no further shift needed) */

    dec    r16
    brne   .Lbcd_loop
    ret

/*
* bcd_adjust – add 3 to any BCD nibble > 4 in R21:R20:R19
*/
bcd_adjust:
    /* Check and adjust each nibble */
    /* R19 low nibble (ones) */
    mov    r22, r19
    andi   r22, 0x0F
    cpi    r22, 5
    brlo   .Lbcd_adj_r19h
    subi   r19, -3      /* add 3 to low nibble */

.Lbcd_adj_r19h:
    /* R19 high nibble (tens) */
    mov    r22, r19
    andi   r22, 0xF0
    cpi    r22, 0x50
    brlo   .Lbcd_adj_r20l
    subi   r19, -0x30

.Lbcd_adj_r20l:
    /* R20 low nibble (hundreds) */
    mov    r22, r20
    andi   r22, 0x0F
    cpi    r22, 5
    brlo   .Lbcd_adj_r20h
    subi   r20, -3

.Lbcd_adj_r20h:
    /* R20 high nibble (thousands) */
    mov    r22, r20
    andi   r22, 0xF0
    cpi    r22, 0x50
    brlo   .Lbcd_adj_r21
    subi   r20, -0x30

```



```
.Lbcd_adj_r21:
    /* R21 low nibble (ten-thousands) – max digit is 6, never > 4 until late */
    mov    r22, r21
    andi   r22, 0x0F
    cpi    r22, 5
    brlo   .Lbcd_adj_done
    subi   r21, -3
.Lbcd_adj_done:
    ret
```

18.5.2 Build and Test

```
avr-gcc -mmcu=attiny3217 -nostartfiles \
        -o bin2bcd.elf src/bin16_to_bcd.S
avr-objdump -d bin2bcd.elf
```

Simulate with simavr:

```
simavr -m attiny3217 bin2bcd.elf
```

18.6 Complete Example: Sensor Reading to 7-Segment Display

A 10-bit ADC reading (0–1023) from a temperature sensor is scaled to a °C value (0–99.9°C, one decimal place), then displayed on a four-digit 7-segment display interface connected over SPI.

18.6.1 Scaling: ADC → Temperature

Assume linear sensor: $0 \rightarrow 0.0^{\circ}\text{C}$, $1023 \rightarrow 99.9^{\circ}\text{C}$.

```
temp10 = (adc_reading × 63998 + 32768) >> 16
```

Where $63998 = \text{round}(999 / 1023 \times 65536)$. The result is an integer in units of 0.1°C , so 365 means 36.5°C . The added 32768 rounds before the final right shift.

Or, equivalently with Q8.8:

```
scale_factor = 99.9 × 256 / 1023 ≈ 24.999... ≈ Q8.8 value 0x1900
temp_Q88 = adc_reading_Q88 × scale_factor_Q88 (Q8.8 multiply)
```

We use the simpler fixed shift approach here to avoid an extra Q8.8 multiply.

The full example is in `src/display.S`. Below are the key routines.

This example is intentionally end-to-end. It shows why bit tricks and numeric formats matter: eventually the value has to become something a human can read.

18.6.2 Register Map for display.S

R25:R24	ADC reading (10-bit, 0–1023)
R23:R22	scaled temperature × 10 (e.g. 365 = 36.5°C)
R21:R20:R19	BCD digits (five digits packed)
R18	digit index (0–3)
R17	segment data byte to SPI
R16	scratch
Z (R31:R30)	pointer to 7-segment LUT in flash

18.6.3 7-Segment LUT

```
.section .rodata
seg7_lut:
/* segments: gfedcba (bit 6..0), active high */
.byte 0x3F /* 0: abcdef = 0b0111111 */
.byte 0x06 /* 1: bc      = 0b0000110 */
.byte 0x5B /* 2: abdeg  = 0b1011011 */
.byte 0x4F /* 3: abcdg  = 0b1001111 */
.byte 0x66 /* 4: bcfg    = 0b1100110 */
.byte 0x6D /* 5: acdfg   = 0b1101101 */
.byte 0x7D /* 6: acdefg  = 0b1111101 */
.byte 0x07 /* 7: abc      = 0b0000111 */
.byte 0x7F /* 8: all     = 0b1111111 */
.byte 0x6F /* 9: abcdfg  = 0b1101111
```

18.6.4 Scale ADC to Temperature

The implementation in `src/display.S` uses a 32-bit intermediate for:

```
temp10 = (adc × 63998 + 32768) >> 16
```

That keeps the result accurate across the full 0-1023 ADC range and avoids the overflow problems that appear in a 16-bit-only sketch.

That is a practical fixed-point lesson by itself: the final answer may fit in 16 bits, but an intermediate step often needs more headroom.

18.6.5 Extract BCD Digits and Drive SPI

```
/*
 * display_temp – convert R23:R22 (temp×10) to BCD and send to SPI display
 *
 * Digit layout: [blank][tens][ones].[tenths]
 * Entry: R23:R22 = temperature × 10 (0–999)
 */
display_temp:
    movw r24, r22
    rcall bin16_to_bcd /* R21:R20:R19 = BCD digits */

    /* temp×10 = abc in decimal corresponds to display "ab.c":
       hundreds digit = tens of degrees
       tens digit     = ones of degrees
       ones digit     = tenths of degrees. */

    /* Send digit 3: always blank for the 0.0..99.9 range */
    ldi r17, 0x00 /* blank segment */
    rcall spi_send

    /* Send digit 2: tens of degrees */
    mov r16, r20
    andi r16, 0x0F
    rcall send_digit

    /* Send digit 1: ones of degrees, with decimal point */
    mov r16, r19
    swap r16
    andi r16, 0x0F
```

```

rcall send_digit_dp

/* Send digit 0: tenths of degrees */
mov  r16, r19
andi r16, 0x0F
rcall send_digit
ret

/*
 * send_digit – look up 7-seg pattern for digit in R16, send via SPI
 */
send_digit:
    ldi  ZL, lo8(seg7_lut)
    ldi  ZH, hi8(seg7_lut)
    add  ZL, r16          /* index into LUT */
    adc  ZH, r1
    lpm  r17, Z           /* r17 = segment pattern */
    rjmp spi_send

/*
 * spi_send – send R17 via SPI (polling)
 */
spi_send:
    sts  SPDR, r17        /* load SPI data register */
.Lspi_wait:
    lds  r16, SPSR
    sbrs r16, SPIF        /* wait for SPIF (transfer complete) bit */
    rjmp .Lspi_wait
ret

```

18.6.6 Build

```

# display.S in this directory targets ATmega328P SPI hardware
avr-gcc -mmcu=atmega328p -nostartfiles \
    -o display.elf src/display.S src/fixedpoint.S src/mul8u.S src/bin16_to_bcd.S
avr-objcopy -O ihex display.elf display.hex
avr-size display.elf

```

18.7 Instruction Reference for This Chapter

Instruction	Syntax	Cycles	Flags	Notes
LSL	lsl Rd	1	C,Z,N,V,H	Rd ← Rd«1; C←Rd7
LSR	lsr Rd	1	C,Z,N,V	Rd ← Rd»1; C←Rd0
ASR	asr Rd	1	C,Z,N,V	Signed »1, sign extends
ROL	rol Rd	1	C,Z,N,V,H	Rotate left through C

Instruction	Syntax	Cycles	Flags	Notes
ROR	ror Rd	1	C,Z,N,V	Rotate right through C
SWAP	swap Rd	1	—	Exchange high/low nibbles
MUL	mul Rd,Rr	2	Z,C	R1:R0 ← Rd×Rr unsigned
MULS	muls Rd,Rr	2	Z,C	Signed × signed
MULSU	mulsu Rd,Rr	2	Z,C	Signed × unsigned
BRHS	brhs k	1/2	—	Branch if H set
BRHC	brhc k	1/2	—	Branch if H clear
ADIW	adiw Rd,K	2	Z,C,N,V,S	Add immediate (0-63) to register pair

SWAP Rd is extremely useful for packed BCD: it moves the high nibble to the low position for digit extraction without shifting:

That is much cleaner than doing four single-bit shifts just to expose one decimal digit.

```
/* Extract high nibble of R16 as a value 0-15 */
swap r16          /* high nibble → low nibble */
andi r16, 0x0F    /* mask off high bits */
```

18.8 Common Pitfalls

18.8.1 Pitfall 1 — Using R1 Without Zeroing After MUL

MUL, MULS, and MULSU write the result to R1:R0. Many AVR assembly projects reserve R1 as a zero register because it makes multi-byte arithmetic and pointer math simpler. If your project follows that convention, copy the product out and restore R1 before code that expects it to be zero:

```
mul    r16, r17
mov    r24, r0      /* use result */
mov    r25, r1
clr    r1           /* restore project zero-register convention */
```

18.8.2 Pitfall 2 — ASR vs LSR on Signed Values

LSR fills the MSB with 0, which is wrong for negative values:

-8 in 8-bit two's complement: 0b11111000 = 0xF8

LSR → 0b01111100 = 0x7C = +124 WRONG

ASR → 0b11111100 = 0xFC = -4 CORRECT (-8 / 2)

Always use ASR for signed right shifts.

If you pick the wrong shift instruction on signed data, the result is not “slightly off.” It is usually catastrophically wrong.

18.8.3 Pitfall 3 — Q8.8 Overflow

Q8.8 integer range is ± 127 (signed) or 0–255 (unsigned). Multiplication of two large Q8.8 values easily overflows:

$200.0 \times 200.0 = 40000.0$ – exceeds unsigned Q8.8 max (255.99)

Check that the mathematical range fits before choosing Q8.8. For larger ranges, use Q4.12 (4-bit integer, 12-bit fraction) or Q12.4.

Choosing a fixed-point format is mostly a range-versus-resolution trade. More fraction bits give you finer steps; more integer bits give you more headroom.

18.8.4 Pitfall 4 — BCD Correction Order

The BCD correction must check low nibble **before** high nibble. Correcting the high nibble first can leave the low nibble invalid after the correction carry propagates:

```
/* WRONG: high nibble corrected first */
subi r16, -0x60 /* might corrupt low nibble */
subi r16, -0x06

/* CORRECT: low nibble first */
subi r16, -0x06
subi r16, -0x60
```

18.8.5 Pitfall 5 — Shift Count for Multi-Byte Values

AVR shifts operate on single registers only. There is no `lsl r25:r24` opcode. Multi-byte shifts require one instruction per byte:

```
/* 24-bit left shift (R22:R21:R20) */
lsl r20
rol r21
rol r22
```

Forget the middle `rol` and bits disappear silently.

That is a recurring AVR pattern: multi-byte arithmetic usually fails by losing a carry in the middle, not by producing an obvious syntax error.

18.9 Summary

Shift recap:

- LSL / LSR – logical shift (unsigned), zero fill
- ASR – arithmetic right shift (signed), sign extend
- ROL / ROR – rotate through C; chain with LSL/LSR for multi-byte shifts
- SWAP – swap nibbles in 1 cycle; essential for packed BCD

Shift-and-add multiply (portable fallback when hardware MUL is unavailable):

8 iterations: `lsr multiplier; brcc skip; add accumulator; shift 16-bit multiplicand`

71..79 cycles for 8×8 vs 2 cycles for MUL on AVRxt devices such as ATtiny3217

Fixed-point Q8.8:

```

value = raw_int / 256
add: same as 16-bit integer add
multiply: 4 partial products, extract bits [23:8] of 32-bit result

```

BCD conversion (double-dabble):

```

16 iterations of: add-3-if-nibble>4, then shift-left
No division, no LUT, bounded data-dependent cycle count

```

Binary-to-BCD use cases:

```

ADC reading → decimal display
Timer counter → MM:SS display
Any sensor value needing human-readable output

```

Build:

```

avr-gcc -mmcu=attiny3217 -nostartfiles -nodefaultlibs -o ch17.elf src/mul8u.S src/fixedpoint.
avr-objdump -d ch17.elf      # verify cycle counts
avr-size ch17.elf           # check flash budget

```

Microchip source notes:

```

AVR Instruction Set Manual, device-core appendix: ATtiny3217 is AVRxt and
its missing-instruction list does not include MUL/MULS/MULSU
AVR Instruction Set Manual, MUL/MULS/MULSU descriptions: 8x8 products write
R1:R0 and execute in 2 cycles on AVRxt

```

18.10 Exercises

1. The `mul8u` routine takes 71..79 cycles, depending on multiplier bits. For a specific use case — multiplying an 8-bit value by the constant 17 — write a dedicated routine using only shifts, adds, and moves that computes the result much faster than the general routine. (Hint: $17 = 16 + 1$.)
2. Q4.12 format stores 4 integer bits and 12 fractional bits in 16 bits. Write the Q4.12 multiply: two Q4.12 values → Q4.12 result. Where do you extract bits from the 32-bit intermediate product? What is the new maximum integer value?
3. The `bin16_to_bcd` routine calls `bcd_adjust` 16 times — once per shift. Each `bcd_adjust` call checks all 5 nibbles. How many of those 5-nibble checks do useful work on iteration 1 (when all BCD bytes are 0)? Optimise `bcd_adjust` to skip nibbles that cannot yet be > 4 based on loop iteration count. How many instruction words does this save?
4. Implement `bcd16_add` — add two 4-digit packed BCD values (two bytes each, representing 0000-9999) with carry out. Use the H flag and manual correction. Test with $9999 + 0001$ (result: 0000, carry = 1).
5. The `seg7_lut` uses active-high segment encoding. Many 7-segment driver ICs use active-low (a 0 bit turns the segment on). Rewrite `send_digit` to invert the LUT byte before sending. Use a single `COM Rd` instruction (one's complement, 1 cycle) — do not rewrite the LUT itself.
6. Profile `display_temp` using `avr-objdump -d`. Count the cycles for the path where no digit is blank. Then count cycles for the path where the hundreds digit is zero (blanked). Which path is longer and by how much?
7. The double-dabble loop shifts the binary input value leftward until it is consumed. After 16 iterations, $R25:R24 = 0$. Verify this is always true regardless of the input value, and explain why it does not cause a problem that $R25:R24$ is destroyed during the conversion.

Next: Chapter 18 — Real-Time Patterns

Chapter 19

Real-Time Patterns

Real-time embedded systems must respond to events within bounded, predictable time. An AVR running bare-metal assembly has no OS scheduler, no kernel, and no hidden overhead — every cycle is accounted for. That is both its strength and its responsibility.

This chapter covers:

1. **Cooperative task scheduling** — a minimal round-robin scheduler in assembly; no dynamic allocation, no context switch overhead beyond register saves.
2. **Deterministic ISR latency analysis** — counting cycles from interrupt trigger to first ISR instruction; identifying and eliminating latency sources.
3. **Example: two-task scheduler with TCB0** — a TCB0 periodic interrupt fires every 1 ms; cooperative task switching drives an LED blinker and a USART heartbeat sender on the ATtiny3217 Curiosity Nano.

The main goal is not to build a tiny preemptive kernel. It is to show how timing, interrupts, and task structure fit together when you want behavior that stays bounded and explainable.

19.1 Real-Time Fundamentals

A real-time system is not necessarily fast — it is **predictable**. Two key properties:

Deadline: a task must complete before a specific time. Missing a deadline is a failure, regardless of whether the computation eventually finishes.

Determinism: given the same inputs, execution time is bounded and known in advance. The bound must hold under all interrupt patterns and all data values.

19.1.1 Two Scheduling Models

Preemptive: a higher-priority task can interrupt a lower-priority task at any instruction boundary. The interrupted task's full register state must be saved and restored, usually on a per-task stack.

Cooperative: tasks run to completion (or to an explicit yield point) before the scheduler runs the next task. No mid-task preemption. Context switch is cheaper; latency is bounded only by the longest task segment between yields.

That last clause is the defining tradeoff. Cooperative systems are simple because nothing interrupts a task halfway through, but that also means one task can delay everything behind it if it runs too long.

This chapter implements cooperative scheduling. It is appropriate when:

- All tasks have known, short run times.
- The system has only a few tasks (two to eight is typical on an 8-bit AVR).
- Worst-case response time = sum of all task segments + scheduler overhead, and that total meets all deadlines.

If you can write that sum down and it fits comfortably inside your timing budget, cooperative scheduling is often the cleanest answer on an AVR.

19.2 Cooperative Task Scheduler Design

19.2.1 Task Representation

Each task is a subroutine that runs, does some work, and returns. The scheduler calls each task in rotation. A task that needs to wait (e.g., for a UART byte) must either poll with a non-blocking check and return immediately, or use a flag set by an ISR.

Scheduler main loop:

```
loop forever:
    call task_0
    call task_1
    call task_2
    ...
```

This is the simplest possible cooperative scheduler. Its properties:

- **No stack switch:** each task runs on the same stack, called directly.
- **No saved context:** tasks return normally; all register state is the task's own responsibility (save what you use).
- **Tick-driven variant:** tasks are only called when a timer tick has elapsed, providing a fixed time base.

Notice what is missing: there is no saved "task state" in CPU registers and no separate per-task stack. These tasks are really just well-behaved subroutines that run on a schedule.

19.2.2 Tick-Driven Scheduler

A TCB0 periodic interrupt fires every 1 ms (the *tick*). Each tick sets a flag. The main loop spins until the flag is set, clears it, then runs all tasks:

```
tcb0_isr:
    ldi    r16, 1
    sts    tick_flag, r16
    reti

main_loop:
.lwait_tick:
    lds    r16, tick_flag
    tst    r16
    breq   .lwait_tick
    clr    r16
    sts    tick_flag, r16
    rcall  task_led
    rcall  task_uart
    rjmp   main_loop
```

In assembly, `tick_flag` is a byte in SRAM (`.bss`). The ISR sets it; the main loop polls and clears it.

That one-byte handoff is a classic embedded pattern: ISR records that an event happened, main-loop code does the slower work later.

19.2.3 Task Counter Pattern

Each task maintains its own counter and acts only when its counter reaches its period:

```
task_led:
    increment led_counter
    compare led_counter with LED_PERIOD
    branch to task_led_done when below period
    clear led_counter
    toggle LED
task_led_done:
    ret

task_uart:
    increment uart_counter
    compare uart_counter with UART_PERIOD
    branch to task_uart_done when below period
    clear uart_counter
    send heartbeat byte
task_uart_done:
    ret
```

With a 1 ms tick: - LED_PERIOD = 500 → LED toggles every 500 ms (1 Hz blink) - UART_PERIOD = 1000 → UART byte sent every 1 s

This pattern scales well because each task keeps its own notion of time without needing a complex global scheduler.

19.3 TCB0 Configuration — 1 ms Tick

On the ATtiny3217 Curiosity Nano the CPU runs at **3.333 MHz** (internal 20 MHz oscillator divided by 6, set by CLKCTRL_MCLKCTRLB).

TCB0 in **Periodic Interrupt mode** (CNTMODE = 0) counts up from 0, reloads from CCMP when it matches, and fires an interrupt on each match. With no prescaler (CLK_PER):

```
CCMP = (F_CPU / F_tick) - 1
      = (3 333 333 / 1000) - 1
      ≈ 3333.333 - 1
      = 3332 (0x0D04)
```

CCMP = 3332 gives a 3333-cycle period, about 999.9 μs at 3.333333 MHz. The remaining fractional cycle is far below the timing error of the default oscillator.

19.3.1 TCB0 Register Map

Address	Name	Description
0x0A40	TCB0_CTRLA	Control A – clock select, enable
0x0A41	TCB0_CTRLB	Control B – mode select
0x0A44	TCB0_EVCTRL	Event control
0x0A45	TCB0_INTCTRL	Interrupt enable (bit 0: CAPT)
0x0A46	TCB0_INTFLAGS	Interrupt flags (bit 0: CAPT; write 1 to clear)
0x0A4A	TCB0_CNTL	Counter low byte
0x0A4B	TCB0_CNTH	Counter high byte

0x0A4C TCB0_CCMP_L Compare/capture match low byte
 0x0A4D TCB0_CCMP_H Compare/capture match high byte

TCB0_CTRLA bit layout:

Bit	Name	Description
0	ENABLE	1 = timer runs
3:1	CLKSEL	000 = CLK_PER (no prescaler) 001 = CLK_PER/2 110 = TCA0 CLK (divide-down from TCA0 prescaler output)

TCB0_CTRLB bit layout:

Bit	Name	Description
2:0	CNTMODE	000 = Periodic Interrupt (INT mode) 001 = Timeout Check 010 = Input Capture on Event 011 = Input Capture Frequency Measurement 100 = Input Capture Pulse-Width Measurement 101 = Input Capture Frequency and Pulse-Width Measurement 110 = Single Shot 111 = 8-bit PWM

19.3.2 Initialization Code

```
/* tcb0_init – configure TCB0 for 1 ms periodic interrupt
 *
 * F_CPU = 3 333 333 Hz, CCMP = 3332 = 0x0D04
 * Clobbers: R16, R17
 */
tcb0_init:
    /* Set CCMP = 3332 (low byte first, high byte second – TCB locks on CCMPH write) */
    ldi r16, lo8(3332)      /* 0x04 */
    sts TCB0_CCMP_L, r16
    ldi r16, hi8(3332)      /* 0x0D */
    sts TCB0_CCMP_H, r16

    /* Mode: periodic interrupt (CNTMODE = 0) */
    ldi r16, 0
    sts TCB0_CTRLB, r16

    /* Enable CAPT interrupt */
    ldi r16, TCB_CAPT_bm     /* bit 0 */
    sts TCB0_INTCTRL, r16

    /* Enable timer, clock = CLK_PER (no prescaler) */
    ldi r16, TCB_ENABLE_bm  /* bit 0 */
    sts TCB0_CTRLA, r16

    ret
```

Instruction	Cycles
LDS Rd, k	3
STS k, Rr	2

LDS/STS access SRAM-mapped peripherals (addresses > 0x001F). On AVRxt, LDS costs 3

cycles and STS costs 2 cycles.

That is an easy place to lose cycles in timer setup and ISR code on newer AVR parts: many peripheral registers are no longer reachable with the short I/O instructions.

19.4 ATtiny3217 Interrupt Vector Table

Vectors are 4 bytes each (one JMP instruction). The table starts at address 0x0000 in flash.

Byte addr	Vector #	Source	Description
0x0000	0	RESET	Power-on, external, WDT, BOD reset
0x0004	1	CRCSCAN_NMI	CRC scan non-maskable interrupt
0x0008	2	BOD_VLM	Brown-out detector voltage level monitor
0x000C	3	PORTA_PORT	Port A pin-change
0x0010	4	PORTB_PORT	Port B pin-change
0x0014	5	PORTC_PORT	Port C pin-change
0x0018	6	RTC_CNT	RTC counter overflow/compare
0x001C	7	RTC_PIT	RTC periodic interrupt timer
0x0020	8	TCA0_LUNF	TCA0 low byte underflow (split mode)
0x0024	9	TCA0_OVF/HUNF	TCA0 overflow / high byte underflow
0x0028	10	TCA0_CMP0	TCA0 compare match 0
0x002C	11	TCA0_CMP1	TCA0 compare match 1
0x0030	12	TCA0_CMP2	TCA0 compare match 2
0x0034	13	TCB0_INT	TCB0 capture/compare (our tick ISR)
0x0038	14	TCB1_INT	TCB1 capture/compare
0x003C	15	TCD0_OVF	TCD0 overflow
0x0040	16	TCD0_TRIG	TCD0 trigger
0x0044	17	AC0_AC	Analog comparator 0
0x0048	18	AC1_AC	Analog comparator 1
0x004C	19	AC2_AC	Analog comparator 2
0x0050	20	ADC0_RESRDY	ADC0 result ready
0x0054	21	ADC0_WCOMP	ADC0 window comparator match
0x0058	22	ADC1_RESRDY	ADC1 result ready
0x005C	23	ADC1_WCOMP	ADC1 window comparator match
0x0060	24	TWI0_TWIS	TWI0 slave
0x0064	25	TWI0_TWIM	TWI0 master
0x0068	26	SPI0_INT	SPI0
0x006C	27	USART0_RXC	USART0 receive complete
0x0070	28	USART0_DRE	USART0 data register empty
0x0074	29	USART0_TXC	USART0 transmit complete
0x0078	30	NVMCTRL_EE	NVMCTRL EEPROM ready

In assembly, the vector table uses jmp for each slot:

```
.section .vectors, "ax", @progbits
    jmp reset_handler      /* 0x0000 RESET */
    jmp default_isr        /* 0x0004 CRCSCAN_NMI */
    jmp default_isr        /* 0x0008 BOD_VLM */
    jmp default_isr        /* 0x000C PORTA_PORT */
    jmp default_isr        /* 0x0010 PORTB_PORT */
    jmp default_isr        /* 0x0014 PORTC_PORT */
    jmp default_isr        /* 0x0018 RTC_CNT */
    jmp default_isr        /* 0x001C RTC_PIT */
    jmp default_isr        /* 0x0020 TCA0_LUNF */
    jmp default_isr        /* 0x0024 TCA0_OVF */
    jmp default_isr        /* 0x0028 TCA0_CMP0 */
```

```

jmp    default_isr    /* 0x002C TCA0_CMP1 */
jmp    default_isr    /* 0x0030 TCA0_CMP2 */
jmp    tcb0_isr       /* 0x0034 TCB0_INT ← our tick */
jmp    default_isr    /* 0x0038 TCB1_INT */
jmp    default_isr    /* 0x003C TCD0_OVF */
jmp    default_isr    /* 0x0040 TCD0_TRIG */
jmp    default_isr    /* 0x0044 AC0_AC */
jmp    default_isr    /* 0x0048 AC1_AC */
jmp    default_isr    /* 0x004C AC2_AC */
jmp    default_isr    /* 0x0050 ADC0_RESRDY */
jmp    default_isr    /* 0x0054 ADC0_WCOMP */
jmp    default_isr    /* 0x0058 ADC1_RESRDY */
jmp    default_isr    /* 0x005C ADC1_WCOMP */
jmp    default_isr    /* 0x0060 TWI0_TWIS */
jmp    default_isr    /* 0x0064 TWI0_TWIM */
jmp    default_isr    /* 0x0068 SPI0_INT */
jmp    default_isr    /* 0x006C USART0_RXC */
jmp    default_isr    /* 0x0070 USART0_DRE */
jmp    default_isr    /* 0x0074 USART0_TXC */
jmp    default_isr    /* 0x0078 NVMCTRL_EE */

```

The explicit-table style is a little longer, but it keeps real code and data out of vector-table slots and makes interrupt ownership obvious when you inspect the binary later.

19.5 ISR Prologue and Epilogue

Every ISR on AVR must:

1. **Save SREG** — instructions inside the ISR affect flags; they must be restored before RETI.
2. **Save any registers it uses** — if the ISR uses R16–R29 or any other register that the interrupted code relies on.
3. **Acknowledge the interrupt source correctly** — for TCB0, write 1 to TCB0_INTFLAGS bit 0 (TCB_CAPT_bm). Other peripherals may clear their flags by different means, such as a data-register read.
4. **Restore saved registers and SREG.**
5. **RETI** — return from interrupt and let CPUINT leave the active interrupt state.

That save/restore discipline is the price you pay for asynchronous execution. If the ISR leaves flags or registers disturbed, the interrupted code resumes in an impossible state.

19.5.1 Minimal Tick ISR

The tick ISR only sets a 1-byte flag — tick_flag in SRAM:

```

/*
 * tcb0_isr – TCB0 periodic interrupt (1 ms tick)
 *
 * Sets tick_flag = 1.
 * Saves/restores: SREG, R16.
 * Latency from interrupt trigger to first useful instruction: see §18.6.
 */
tcb0_isr:
    push    r16          /* 1 cycle: save R16 to stack on AVRxt */
    in      r16, _SFR_IO_ADDR(SREG) /* 1 cycle: save SREG */
    push    r16          /* 1 cycle */

```

```

/* clear interrupt flag (write 1 to TCB_CAPT_bm = bit 0) */
ldi   r16, TCB_CAPT_bm
sts   TCB0_INTFLAGS, r16      /* 2 cycles */

/* set tick flag */
ldi   r16, 1
sts   tick_flag, r16         /* 2 cycles */

pop   r16                    /* 1 cycle: restore SREG value */
out   _SFR_IO_ADDR(SREG), r16 /* 1 cycle: restore SREG */
pop   r16                    /* 1 cycle: restore R16 */
reti  /* 4 cycles on ATtiny3217 */

```

Total ISR body: 1+1+1+1+2+1+2+1+1+1+4 = **16 cycles** from first instruction to end of RETI.

On AVRxt devices such as the ATtiny3217, RETI costs 4 cycles with the 16-bit program counter. Microchip's instruction set manual notes that CPUINT tracks interrupt state on AVRxm/AVRxt devices; RETI returns through that interrupt-controller state rather than being just a plain RET. Keep it only at the true end of the handler, after every register and flag has been restored.

19.6 ISR Latency Analysis

ISR latency = cycles from interrupt trigger condition to the first instruction of the ISR executing.

19.6.1 Sources of Latency

On ATtiny3217 (avrxtmega3), the CPU-side interrupt response after the request has been recognized is:

Source	Cycles
Peripheral/request synchronization	0–2
Finish current instruction	1–4
Push PC to stack	2
Jump from vector to ISR (jmp)	3
Total (best case)	6
Total (worst case with 2 sync cycles and a 4-cycle instruction)	11

Microchip documents a minimum CPU interrupt response of 6 cycles on devices with more than 8 KB of flash: 1 cycle to finish a single-cycle instruction, 2 cycles to push the PC, and 3 cycles for the vector jmp. Peripheral-level synchronization can add a small device-specific delay before the CPU begins that sequence. In practice:

Interrupt asserted during instruction N

- request is synchronized into the CPU domain (0–2 cycles)
- instruction N completes (1–4 cycles total, depending on instruction length)
- CPU pushes PC (2 cycles)
- CPU executes JMP at vector (3 cycles)
- first ISR instruction executes (cycle 6–11 from assertion)

This is the right way to think about interrupt latency on AVR: not as one magic number, but as a short chain of small, countable delays.

19.6.2 Push/Pop Overhead

The prologue `push r16 / in r16, SREG / push r16` costs **3 cycles** before any useful work. For a time-critical ISR (e.g., one that must sample a pin within N cycles), count:

```
latency_to_useful_work = ISR_entry_latency + prologue_cycles
                        = 6 (best) + 3 = 9 cycles minimum
                        = 11 (worst) + 3 = 14 cycles worst case
```

At 3.333 MHz, 14 cycles = **4.2 µs** worst-case latency to the first useful instruction.

19.6.3 Eliminating Sources of Latency

Technique	Cycles saved
Keep ISRs short — move work to main loop	Reduces re-entry latency
Use VPORT for output (1-cycle <code>sbi/cbi</code>)	0 — saves body cycles, not entry
Avoid long instructions (<code>LPM</code> , <code>CALL</code>) near critical ISR entry points	0–4
Reduce number of registers pushed in prologue	2 per saved/restored register
Reserve a known scratch register for a dedicated ISR and document ownership	2 per omitted push/pop pair

For the tick ISR, latency is unimportant — the tick fires every 1 ms and a few microseconds of jitter is negligible. Latency matters for ISRs that must capture hardware events (e.g., `TCB0` in input-capture mode, or a pin-change ISR measuring pulse width).

So always ask “latency relative to what?” before optimizing an ISR. A scheduler tick and an edge-timestamp ISR care about very different things.

19.6.4 Measuring Latency

Toggle a debug pin at the start and end of the ISR, then use an oscilloscope or logic analyser to measure ISR entry and execution width relative to another observable system event:

```
tcb0_isr:
    sbi    VPORTB_OUT, 3    /* debug pin PB3 high — mark ISR entry    */
    push   r16
    ...
    cbi    VPORTB_OUT, 3    /* debug pin PB3 low  — mark ISR exit    */
    reti
```

`SBI VPORTB_OUT, 3` executes in **1 cycle** and does not affect `SREG` — safe before saving `SREG`, as long as the debug pin is not used by the application.

19.7 Complete Two-Task Scheduler

19.7.1 Hardware Setup — Curiosity Nano

LED0 = PA3, active low (HIGH = off, LOW = on)
 USART0 TX = PB2 (to on-board debugger virtual COM port)
 TCB0 tick = every 1 ms (CCMP = 3332, CLK_PER, F_CPU = 3.333 MHz)

19.7.2 Register Allocation

R16 – scratch in ISR, scheduler, and peripheral setup
 R24:R25 – 16-bit task counters while updating SRAM state
 R1 – zero register convention for carry propagation

SRAM variables (.bss):

```
tick_flag    1 byte – set by ISR, cleared by main loop
led_cnt      2 bytes – 16-bit counter for LED task
uart_cnt     2 bytes – 16-bit counter for UART task
```

Using SRAM variables (accessed via LDS/STS) ensures the ISR and main loop share data correctly — SRAM is in the data address space accessible to both.

It also keeps the design honest: anything shared between asynchronous code paths should live in memory, not only in a caller's temporary register plan.

19.7.3 Source Code

```
/* scheduler.S – two-task cooperative scheduler, ATtiny3217 Curiosity Nano
 *
 * Tasks:
 *   task_led – toggle LED0 (PA3) every 500 ms
 *   task_uart – send 'H' over USART0 every 1000 ms
 *
 * Tick: TCB0 periodic interrupt, 1 ms
 * F_CPU: 3 333 333 Hz (20 MHz / 6, default Curiosity Nano clock)
 */

#include <avr/io.h>

#define F_CPU      333333UL
#define LED_PIN    3          /* PA3, Curiosity Nano LED0 */
#define LED_PERIOD 500        /* ticks (= 500 ms) */
#define UART_PERIOD 1000      /* ticks (= 1 s) */

/* USART0 baud: 9600 at 3.333 MHz
 * BAUD register = 64 * F_CPU / (16 * baud) = 64 * 333333 / (16 * 9600)
 *                = 21333333 / 153600 ≈ 1389 = 0x056D
 */
#define BAUD_9600 1389

/* — Vector table — */
.section .vectors, "ax", @progbits
jmp reset_handler      /* 0x0000 RESET */
jmp default_isr         /* 0x0004 CRCSCAN_NMI */
jmp default_isr         /* 0x0008 BOD_VLM */
jmp default_isr         /* 0x000C PORTA_PORT */
jmp default_isr         /* 0x0010 PORTB_PORT */
jmp default_isr         /* 0x0014 PORTC_PORT */
jmp default_isr         /* 0x0018 RTC_CNT */
jmp default_isr         /* 0x001C RTC_PIT */
jmp default_isr         /* 0x0020 TCA0_LUNF */
jmp default_isr         /* 0x0024 TCA0_OVF */
jmp default_isr         /* 0x0028 TCA0_CMP0 */
jmp default_isr         /* 0x002C TCA0_CMP1 */
jmp default_isr         /* 0x0030 TCA0_CMP2 */
jmp tcb0_isr           /* 0x0034 TCB0_INT */
jmp default_isr         /* 0x0038 TCB1_INT */
```

```

    jmp default_isr      /* 0x003C TCD0_OVF */
    jmp default_isr      /* 0x0040 TCD0_TRIG */
    jmp default_isr      /* 0x0044 AC0_AC */
    jmp default_isr      /* 0x0048 AC1_AC */
    jmp default_isr      /* 0x004C AC2_AC */
    jmp default_isr      /* 0x0050 ADC0_RESRDY */
    jmp default_isr      /* 0x0054 ADC0_WCOMP */
    jmp default_isr      /* 0x0058 ADC1_RESRDY */
    jmp default_isr      /* 0x005C ADC1_WCOMP */
    jmp default_isr      /* 0x0060 TWI0_TWIS */
    jmp default_isr      /* 0x0064 TWI0_TWIM */
    jmp default_isr      /* 0x0068 SPI0_INT */
    jmp default_isr      /* 0x006C USART0_RXC */
    jmp default_isr      /* 0x0070 USART0_DRE */
    jmp default_isr      /* 0x0074 USART0_TXC */
    jmp default_isr      /* 0x0078 NVMCTRL_EE */

/* — SRAM variables ————— */
.section .bss
tick_flag: .byte 0      /* 1 = tick has fired; 0 = not yet */
led_cnt:   .word 0      /* LED task counter (16-bit) */
uart_cnt:  .word 0      /* UART task counter (16-bit)

/* — TCB0 ISR ————— */
.section .text
tcb0_isr:
    push r16
    in r16, _SFR_IO_ADDR(SREG)
    push r16

    ldi r16, TCB_CAPT_bm
    sts TCB0_INTFLAGS, r16 /* clear interrupt flag */

    ldi r16, 1
    sts tick_flag, r16     /* signal tick to main loop */

    pop r16
    out _SFR_IO_ADDR(SREG), r16
    pop r16
    reti

default_isr:
    reti                  /* ignore unexpected interrupts */

/* — Initialization ————— */
reset_handler:
    /* Stack pointer: top of SRAM = 0x3FFF */
    ldi r16, hi8(RAMEND)
    out _SFR_IO_ADDR(SPH), r16
    ldi r16, lo8(RAMEND)
    out _SFR_IO_ADDR(SPL), r16
    clr r1                /* project zero register for ADC-with-carry use */

    /* LED0: PA3 output, initial state HIGH (LED off – active low) */
    sbi VPORTA_DIR, LED_PIN
    sbi VPORTA_OUT, LED_PIN

```



```

/* USART0: set TX pin PB2 idle-high, then output */
sbi  VPORTB_OUT, 2
sbi  VPORTB_DIR, 2

rcall tcb0_init
rcall usart0_init

sei                                     /* enable global interrupts */

rjmp  main_loop

/* — TCB0 setup — */
tcb0_init:
ldi   r16, lo8(3332)
sts   TCB0_CCMP_L, r16
ldi   r16, hi8(3332)
sts   TCB0_CCMP_H, r16
ldi   r16, 0
sts   TCB0_CTRLB, r16      /* CNTMODE = 0: periodic interrupt */
ldi   r16, TCB_CAPT_bm
sts   TCB0_INTCTRL, r16    /* enable CAPT interrupt */
ldi   r16, TCB_ENABLE_bm
sts   TCB0_CTRLA, r16      /* CLK_PER, enable */
ret

/* — USART0 setup — */
usart0_init:
/* BAUD = 1389 = 0x056D */
ldi   r16, lo8(BAUD_9600)
sts   USART0_BAUD_L, r16
ldi   r16, hi8(BAUD_9600)
sts   USART0_BAUD_H, r16
/* CTRLB: TX enable (TXEN bit 6) */
ldi   r16, USART_TXEN_bm
sts   USART0_CTRLB, r16
/* CTRLC: async USART, 8N1 (default reset value = 0x03 = 8-bit, 1 stop) */
/* No explicit write needed – reset value is correct */
ret

/* — Main loop — */
main_loop:
/* wait for tick */
.Lwait_tick:
lds   r16, tick_flag
tst   r16
breq  .Lwait_tick

/* clear flag */
clr   r16
sts   tick_flag, r16

/* run tasks */
rcall task_led
rcall task_uart

rjmp  main_loop

```

```

/* — task_led ————— */
task_led:
    /* increment 16-bit counter */
    lds    r24, led_cnt
    lds    r25, led_cnt+1
    adiw   r24, 1
    sts    led_cnt, r24
    sts    led_cnt+1, r25

    /* compare to LED_PERIOD (500) */
    cpi    r24, lo8(LED_PERIOD)
    ldi    r16, hi8(LED_PERIOD)
    cpc    r25, r16
    brlo   .Lled_done          /* counter < period: nothing to do */

    /* reset counter */
    clr    r24
    clr    r25
    sts    led_cnt, r24
    sts    led_cnt+1, r25

    /* toggle PA3 via PORTA_OUTTGL (1-byte write toggles the pin) */
    ldi    r16, (1 << LED_PIN)
    sts    PORTA_OUTTGL, r16

.Lled_done:
    ret

/* — task_uart ————— */
task_uart:
    /* increment 16-bit counter */
    lds    r24, uart_cnt
    lds    r25, uart_cnt+1
    adiw   r24, 1
    sts    uart_cnt, r24
    sts    uart_cnt+1, r25

    /* compare to UART_PERIOD (1000) */
    cpi    r24, lo8(UART_PERIOD)
    ldi    r16, hi8(UART_PERIOD)
    cpc    r25, r16
    brlo   .Luart_done

    /* reset counter */
    clr    r24
    clr    r25
    sts    uart_cnt, r24
    sts    uart_cnt+1, r25

    /* send 'H' (0x48): wait for DREIF (Data Register Empty Flag) */
.Luart_wait_dreif:
    lds    r16, USART0_STATUS
    sbrs   r16, USART_DREIF_bp    /* skip next if DREIF=1 (ready) */
    rjmp   .Luart_wait_dreif

    ldi    r16, 'H'
    sts    USART0_TXDATA1, r16

```

```
.Luart_done:
    ret
```

19.7.4 Build

```
MCU=attiny3217
avr-gcc -mmcu=${MCU} -nostartfiles -nodefaultlibs -o scheduler.elf src/scheduler.S
avr-objdump -h scheduler.elf
avr-objcopy -O ihex scheduler.elf scheduler.hex
avrdude -p t3217 -c serialupdi -P /dev/ttyUSB0 -U flash:w:scheduler.hex:i
```

19.8 Timing Analysis — Worst-Case Response

With two tasks per tick, the cooperative overhead before the main loop re-checks `tick_flag` is:

```
scheduler_overhead = wait_tick_poll (1 iteration) + clear_flag
                   = (3 + 1 + 1) + (1 + 2) = ~8 cycles
```

```
task_led_max = lds×2 + adiw + sts×4 + cpi + ldi + cpc + brlo + ... toggle path
              ≈ 2×3 + 2 + 2×2 + 1 + 1 + 1 + 1 + clr×2 + 2×2 + ldi + sts + ret
              ≈ 29 cycles (toggle path, excluding RCALL)
              ≈ 21 cycles (no-toggle path, excluding RCALL)
```

```
task_uart_max (no-send path) ≈ 21 cycles (same structure as task_led)
task_uart_send_fast_path ≈ 34 cycles (USART immediately ready, excluding RCALL)
```

```
total_cycles_per_tick (when USART is immediately ready) ≈ 8 + 2 + 29 + 2 + 34 + 2 = ~77 cycles max
time_at_3.333MHz = 77 / 3 333 333 ≈ 23 μs
```

The tick period is 1 ms = 3333 cycles. On the fast path, the scheduler consumes ≈ 77 of those cycles per tick — about 2.3% CPU utilisation.

Worst-case tick latency (time between ISR firing and main loop responding):

```
worst_case_tick_latency_fast_path ≈ task_led_max + task_uart_send_fast_path
                                   ≈ 29 + 34 + RCALL overheads = ~67 cycles ≈ 20 μs
```

This occurs when both tasks do their full work in the tick that just preceded the one being measured. The scheduler handles one tick at a time; if both tasks fire simultaneously on the same tick, `tick_flag` is set for the next tick while the current work finishes.

Important caveat: the `scheduler.S` UART task is still **blocking**. If the code reaches `.Luart_wait_dreif` just after the USART has become busy, the loop can stall for almost a full character time. At 9600 baud with 8N1, one frame is about $10/9600 \text{ s} \approx \mathbf{1.04 \text{ ms}}$, which is longer than the 1 ms tick. So the fast-path cycle count above is useful for instruction budgeting, but it is **not** a strict worst-case bound for the blocking UART design.

That is the most important design lesson in the chapter: a system is not truly bounded just because its arithmetic paths are short. One blocking peripheral wait can dominate the whole schedule.

19.9 Non-Blocking Task Extension

The fix is to decouple task execution from peripheral transmit timing. The non-blocking variant in `src/scheduler_nb.S` uses:

- a small SRAM ring buffer,
- a short `uart_enqueue` routine used by the task, and
- the `USART0_DRE` interrupt to feed bytes into `TXDATA1` whenever the UART is ready for another byte.

Conceptually:

```
task_uart:
    once per second:
        enqueue "Hbeat\r\n"
        enable DRE interrupt

USART0_dre_isr:
    compare tx_tail with tx_head
    if equal, disable DRE interrupt and reti
    load tx_buf[tx_tail]
    write byte to TXDATA1
    advance tx_tail
    reti
```

This keeps `task_uart` short and bounded: it copies bytes into the queue and returns. The USART hardware timing is handled asynchronously by the DRE ISR. In the example implementation, bytes are dropped if the ring buffer is full; that policy keeps the task non-blocking and deterministic.

This is a typical real-time trade: it is often better to lose optional output than to let one slow device destroy the timing guarantees of the whole system.

Full implementation in `src/scheduler_nb.S`.

Two details matter in that source file:

- The heartbeat string is placed in `.text`, not `.rodata`, so LPM reads from a flash byte address. It is aligned after the 7-byte string so the following AVR instructions still start on a word boundary.
- `uart_enqueue` saves `R18` because the caller uses it as the heartbeat string loop counter. A non-blocking helper is only deterministic if its register ownership is documented and obeyed.

19.10 Scaling to More Tasks

The round-robin scheduler scales by adding more `rcall task_N` instructions in the main loop. Each task must be **short** — its worst-case execution time must fit within the tick budget:

```
sum(task_worst_case_cycles) < tick_period_cycles
sum ≤ 3333 cycles (1 ms at 3.333 MHz)
```

For a task that exceeds this (e.g., a DFT computation), split it across multiple ticks using a state machine:

```
dft_task:
    /* process one input sample per tick – spread N-point DFT over N ticks */
    load dft_state
    branch to the matching state label
SAMPLE_0:
```

```

    compute partial sum for k=0
    store SAMPLE_1 into dft_state
    ret
SAMPLE_1:
    compute partial sum for k=1
    store SAMPLE_2 into dft_state
    ret

```

Breaking work into per-tick slices is one of the most useful habits in small embedded systems. It keeps long algorithms compatible with a simple scheduler.

19.10.1 When to Move to a Preemptive Kernel

A cooperative scheduler works until either: 1. Any task's worst-case time can exceed the tick period, or 2. Tasks have different priorities that require preemption.

For those cases, a preemptive kernel provides priority scheduling and proper stack switching. The assembly patterns in this chapter (ISR prologue/epilogue, timer setup, ISR-to-task signalling) are the same low-level building blocks.

So even if you eventually move to a preemptive kernel, the reasoning here still matters. The kernel is built out of the same timer and interrupt mechanics.

19.11 Summary

Topic	Key Point
Cooperative scheduling	Tasks run to completion; scheduler calls them in rotation
Tick source	TCB0 periodic interrupt, CCMP=3332 for 1 ms at 3.333 MHz
ISR entry latency	6-11 cycles (hardware) + 3 cycles (prologue)
SREG save	Mandatory in any ISR that uses arithmetic or branch instructions
RETI	Returns through CPUINT interrupt state; costs 4 cycles on ATtiny3217
VPORT	1-cycle SBI/CBI for fast output; does not affect SREG
Tick budget	3333 cycles/ms; cooperative core uses < 80 cycles on the fast path
Scaling	Add <code>rcall task_N</code> ; split long tasks into per-tick state machines

Microchip source notes:

- ATtiny3216/3217 data sheet, TCB Periodic Interrupt mode: TCB counts to the capture value, restarts from BOTTOM, and generates a CAPT interrupt/event.
- ATtiny3216/3217 data sheet, CPUINT Interrupt Response Time: devices with more than 8 KB Flash use a 3-cycle `jmp` vector, after 1 cycle to finish a single-cycle instruction and 2 cycles to store the PC.
- AVR Instruction Set Manual, RETI: AVRxt devices use CPUINT interrupt state on return; RETI is 4 cycles with a 16-bit program counter.
- Microchip USART guidance: asynchronous normal mode uses $BAUD = 64 * f_CLK_PER / (16 * f_BAUD)$.

Next: Chapter 19 — CORDIC on AVR

Chapter 20

CORDIC on AVR

CORDIC (COordinate Rotation DIgital Computer) is one of the classic tricks for small processors. It computes trigonometric, inverse-trigonometric, and vector rotation results using only:

- shifts,
- adds/subtracts,
- comparisons, and
- a small table of constants.

That makes it a natural fit for AVR. The ATtiny3217 includes the AVR hardware multiplier, but Microchip's instruction set still has no general divide instruction and the device has no floating-point execution unit. When you need sine, cosine, vector angle, or vector magnitude with predictable execution time, CORDIC is often a better fit than general software floating-point math.

This chapter covers:

1. **Why CORDIC is useful on AVR** — where it beats general math routines, and where a lookup table is still better.
2. **Rotation mode** — computing sine and cosine from an angle.
3. **Vectoring mode** — reducing a vector to magnitude and angle.
4. **Fixed-point formats** — a practical AVR-friendly choice: Q1.15 for vector components and BAM16 for angles.
5. **Examples** — AVR assembly stage examples, register plans, and practical integration patterns.

The important theme is the same one that runs through the rest of the book: replace expensive, hard-to-bound math with small deterministic steps.

20.1 Why Use CORDIC on AVR?

CORDIC is attractive on AVR for three practical reasons:

1. **No floating-point execution unit.** General software floating-point trigonometry is much larger and harder to budget than a fixed shift-add kernel.
2. **Fixed iteration count.** A 12-stage CORDIC always does 12 stages. That is useful when you care about real-time bounds.
3. **One core, many functions.** The same shift-add structure can be used for sine/cosine, vector rotation, atan2 , and magnitude approximation.

Typical AVR uses:

- waveform generation in a DDS/NCO,

- phase tracking,
- I/Q demodulation,
- motor-control angle estimation,
- tilt/heading math from sensors, and
- display-oriented trig without a large general-purpose math package.

20.1.1 When Not to Use It

CORDIC is not magic. Sometimes a lookup table is better:

- If you only need sine values at 8-bit resolution, a 256-byte flash table is usually faster.
- If you need one-off trig rarely, a larger general-purpose math routine may be acceptable.
- If your MCU has enough flash and your timing budget is loose, a polynomial or LUT approach may be easier to maintain.

So the real question is not “is CORDIC clever?” It is “does repeated deterministic shift-add math fit this problem better than tables or general software arithmetic?”

20.2 The Core Idea

CORDIC works by applying a sequence of tiny rotations whose tangents are powers of two:

$\text{atan}(2^0)$, $\text{atan}(2^{-1})$, $\text{atan}(2^{-2})$, $\text{atan}(2^{-3})$, ...

Because $\tan(\theta_i) = 2^{-i}$, each stage can be written without a multiply:

```
x' = x plus_or_minus (y >> i)
y' = y minus_or_plus (x >> i)
z' = z minus_or_plus atan(2^-i)
```

The shifts replace a true multiply by $\tan(\theta_i)$.

There are two standard modes:

20.2.1 Rotation Mode

Start with a unit vector on the x-axis and rotate it toward the requested angle. After the final stage:

- x approximates $\cos(\theta)$
- y approximates $\sin(\theta)$

20.2.2 Vectoring Mode

Start with an input vector (x, y) and rotate it toward the x-axis. After the final stage:

- z approximates the vector angle (atan2)
- x approximates the vector magnitude, scaled by the CORDIC gain
- y approaches zero

This is why one algorithm can serve both “give me sine/cosine” and “tell me the angle/magnitude of this vector.”

20.3 Fixed-Point Choices for AVR

CORDIC is almost always done in fixed-point on AVR.

20.3.1 Vector Format: Q1.15

Use signed 16-bit values for x and y:

Q1.15: signed fraction in the range about -1.0 to +0.99997

```
0x4000 = +0.5
0x7FFF = +0.99997
0x0000 = 0.0
0xC000 = -0.5
0x8000 = -1.0
```

This is a good fit for sine and cosine because those values naturally live in the interval [-1, +1].

20.3.2 Angle Format: BAM16

For angles, a very AVR-friendly choice is **BAM16** (Binary Angular Measure):

```
Full turn = 65536 units
90°       = 16384 = 0x4000
180°      = 32768 = 0x8000
360°      = wraps to 0x0000
```

Why BAM16 works well:

- addition and wraparound are natural in 16-bit unsigned arithmetic,
- phase accumulators use the same representation,
- quadrants are easy to identify from the top two bits.

20.3.3 CORDIC Gain

The repeated pseudo-rotations scale the vector by a constant gain:

```
1 / K ≈ 1.646760258
K      ≈ 0.607252935
```

After 12 stages the gain is already close to the infinite-stage value:

```
K12 = 0.607252959
1/K12 = 1.646760193
```

For a 12-stage sine/cosine CORDIC, preload:

```
K_Q15 = round(0.607252959 × 32768) = 19898 = 0x4DBA
```

That means rotation mode typically starts with:

```
x = 0x4DBA ; +0.60725 in Q1.15
y = 0x0000
```

and the final result lands near properly scaled sine/cosine outputs.

20.4 The Angle Table

For 12 stages in BAM16:

i	atan(2 ⁻ⁱ)	BAM16 value
0	45.000°	0x2000
1	26.565°	0x12E4
2	14.036°	0x09FB

i	$\text{atan}(2^{-i})$	BAM16 value
3	7.125°	0x0511
4	3.576°	0x028B
5	1.790°	0x0146
6	0.895°	0x00A3
7	0.448°	0x0051
8	0.224°	0x0029
9	0.112°	0x0014
10	0.056°	0x000A
11	0.028°	0x0005

This table is tiny. Even stored as 16-bit words in flash, 12 stages only cost 24 bytes.

20.4.1 Convergence Range

Basic rotation-mode CORDIC converges only for angles near $\pm 99.88^\circ$, not over the full 360° range. On AVR, the usual fix is **quadrant reduction**:

1. Determine the quadrant from bits [15:14] of the BAM16 angle.
2. Reflect the angle into the first quadrant [0° , 90°].
3. Run the CORDIC core there.
4. Apply the final sign correction to sin and cos.

That is why BAM16 is convenient: the top two bits already tell you the quadrant.

One fixed-point detail matters at the quadrant boundaries: signed Q1.15 can represent -1.0 exactly (0x8000), but not +1.0. A production sine/cosine wrapper should special-case the exact axis angles (0° , 90° , 180° , 270°) or saturate positive full-scale outputs to 0x7FFF.

20.5 Rotation Mode for Sine and Cosine

Register-transfer view using Q1.15 vectors and BAM16 angles:

```
x = K_Q15
y = 0
z = reduced_angle      ; first quadrant, 0..90 degrees

repeat for stages i = 0..11:
    temp_x = arithmetic_shift_right(x, i)
    temp_y = arithmetic_shift_right(y, i)
    test sign bit of z
    when z is non-negative:
        x = x - temp_y
        y = y + temp_x
        z = z - atan_table[i]
    when z is negative:
        x = x + temp_y
        y = y - temp_x
        z = z + atan_table[i]
```

After the loop:

- x is the first-quadrant cosine
- y is the first-quadrant sine

Then apply quadrant signs. For exact axis angles, use the special-case or saturation rule above before applying the signs:

Quadrant	Reduced angle	cos sign	sin sign
0	angle	+	+
1	180°-angle	-	+
2	angle-180°	-	-
3	360°-angle	+	-

20.5.1 Example Values

With 12 stages and $K_{Q15} = 0x4DBA$, away from the exact axis-angle cases:

- 30° (0x1555) gives about $\cos \approx 0.8660$, $\sin \approx 0.5001$
- 45° (0x2000) gives about $\cos \approx 0.7070$, $\sin \approx 0.7072$
- 300° (0xD555) gives about $\cos \approx 0.5001$, $\sin \approx -0.8660$

That is easily good enough for control loops, DDS, UI graphics, and many sensor applications on AVR.

20.6 AVR Assembly Structure

An AVR-assembly CORDIC routine should be organized around three fixed pieces:

- quadrant reduction on the BAM16 input angle,
- a 12-entry atan table in flash, and
- either a counted loop or a fully unrolled set of stages.

20.6.1 Table Layout

One straightforward flash table is:

```
cordic_atan_bam16:
.word 0x2000, 0x12E4, 0x09FB, 0x0511
.word 0x028B, 0x0146, 0x00A3, 0x0051
.word 0x0029, 0x0014, 0x000A, 0x0005
```

Put this table in .text or another flash-resident section and align it to a word boundary. When reading it with LPM, remember that LPM fetches bytes; read the low byte and high byte separately, then advance Z to the next word.

For rotation mode, preload the vector with $K_{Q15} = 0x4DBA$:

```
ldi r24, lo8(0x4DBA) ; x low
ldi r25, hi8(0x4DBA) ; x high
clr r22 ; y low
clr r23 ; y high
```

20.6.2 Core Loop Skeleton

At a high level, the loop does this each stage:

```
1. form temp_x = x >> i      using signed shifts
2. form temp_y = y >> i      using signed shifts
3. if z >= 0:
    x = x - temp_y
    y = y + temp_x
    z = z - atan[i]
else:
    x = x + temp_y
```

```

    y = y - temp_x
    z = z + atan[i]
4. advance i and table pointer

```

That description is intentionally close to the register transfers you will actually code. On AVR, the main implementation choice is whether *i* is variable inside a loop or fixed in an unrolled stage sequence.

20.7 How This Maps to AVR Assembly

At the algorithm level, CORDIC is simple. At the assembly level, the main cost is the **variable arithmetic shift** by *i*.

For a general loop, $x \gg i$ and $y \gg i$ mean:

1. copy the 16-bit value to a temporary pair,
2. shift it right *i* times using ASR on the high byte and ROR on the low byte,
3. use the shifted temporary in the add/subtract step.

That is why many hand-tuned AVR CORDIC routines are either:

- partially unrolled, or
- fully unrolled.

When *i* is a compile-time constant, the shifts become fixed instruction sequences and the loop-control overhead disappears.

The file `src/cordic_helpers.S` contains a buildable support routine for the variable signed 16-bit shift:

```

/* Entry: R25:R24 = signed value, R17 = count. Exit: R25:R24 shifted. */
cordic_asr16:
    tst    r17
    breq   .Lasr16_done
.Lasr16_loop:
    asr    r25
    ror    r24
    dec    r17
    brne   .Lasr16_loop
.Lasr16_done:
    ret

```

20.7.1 Example: One Unrolled Stage (*i* = 3)

Suppose:

- R25:R24 = *x*
- R23:R22 = *y*
- R21:R20 = *z*

One positive-rotation stage for *i* = 3 looks like:

```

/* temp_y = y >> 3 (signed shift) */
movw   r19, r22          /* r19:r18 = y copy */
asr     r19
ror     r18
asr     r19
ror     r18
asr     r19
ror     r18

```

```

/* x = x - temp_y */
sub    r24, r18
sbc    r25, r19

/* temp_x = old x >> 3
 * In a real implementation, preserve old x before modifying it. */

```

The mechanics are exactly the same as any other multi-byte signed shift on AVR: ASR the high byte, ROR the low byte, repeat as many times as needed.

20.7.2 Practical Register Plan

For a hand-written AVR version, a reasonable plan is:

```

R25:R24    x
R23:R22    y
R21:R20    z
R19:R18    temp shifted copy
R17        stage counter i
R16        scratch / table byte
Z          pointer into atan table

```

This is workable, but you can see why register pressure gets real quickly. CORDIC is a good example of a routine where AVR assembly quality depends as much on register planning as on the arithmetic itself.

20.8 Vectoring Mode: Magnitude and Angle

Vectoring mode starts with an input vector (x, y) and rotates it toward the x-axis:

```

z = 0
repeat for stages i = 0..11:
    temp_x = arithmetic_shift_right(x, i)
    temp_y = arithmetic_shift_right(y, i)
    test sign bit of y
    when y is non-negative:
        x = x + temp_x
        y = y - temp_y
        z = z + atan_table[i]
    when y is negative:
        x = x - temp_x
        y = y + temp_y
        z = z - atan_table[i]

```

After the last stage:

- z approximates the four-quadrant vector angle
- x approximates magnitude / K
- y is near zero

The magnitude is scaled upward by the CORDIC gain $1/K$, so you either:

- multiply the result by K afterwards,
- pre-scale the input vector by K, or
- calibrate the rest of the application around the scaled value.

This gain matters for overflow. A Q1.15 input vector with magnitude near 1.0 cannot safely be fed directly into a 16-bit vectoring core, because the final x wants to grow to

about 1.64676. Either pre-scale the vector by K, use a wider internal format such as 24-bit or 32-bit fixed point, or accept a smaller input range.

20.8.1 Where AVR Uses This

Vectoring mode is especially useful for:

- phase estimation from I/Q samples,
- converting Cartesian sensor axes to angle,
- simple magnitude estimation without a square root.

This is where CORDIC often beats ad-hoc math on AVR: one fixed-iteration loop replaces both angle extraction and square-root style magnitude calculation.

20.9 Example 1: DDS / NCO Sine Generator

A common AVR use is a numerically controlled oscillator:

1. keep a 16-bit phase accumulator,
2. add a phase step each sample period,
3. feed the phase into CORDIC rotation mode,
4. send the sine result to a PWM or DAC.

```
phase += phase_step;
call cordic_sincos_bam16
use high byte of sine as PWM sample
```

Why CORDIC can make sense here:

- no large sine table in flash,
- easy frequency control through phase_step,
- constant run time every sample.

If you need the very highest sample rate, a LUT is still usually faster. But if flash is tight and you want both sine and cosine from the same phase input, CORDIC is an attractive compromise.

20.10 Example 2: Tilt Angle from Two Axes

Suppose an AVR reads two already-filtered accelerometer channels:

```
x_axis = forward/back tilt component
y_axis = vertical tilt component
```

You want an angle estimate:

```
tilt = four-quadrant angle from x_axis and y_axis
```

Vectoring CORDIC gives the AVR a fixed-time integer path:

1. load (x, y) into Q1.15,
2. run vectoring mode,
3. interpret z in BAM16,
4. convert BAM16 to degrees only if the UI needs it.

That avoids general software floating-point math and keeps the control-loop representation in a format that is cheap to update.

20.11 Accuracy, Flash, and Cycle Tradeoffs

20.11.1 Iteration Count

More stages mean more accuracy and more cycles.

Typical choices on AVR:

- 8 stages: quick, coarse
- 12 stages: good practical middle ground
- 16 stages: near the limit of 16-bit usefulness

With Q1.15 outputs, 12 stages is usually enough because the result resolution is already limited by the 16-bit format.

20.11.2 LUT vs CORDIC

Rough comparison:

Method	Flash	Cycles	Strength
256-byte sine LUT	moderate	very low	fastest for sine only
CORDIC (12 stages)	small table + code	moderate, fixed	sine/cosine/angle/magnitude from one core
Polynomial fixed-point	small to moderate	data-dependent unless carefully bounded	good for one narrow function

The main advantage of CORDIC is not “fewest cycles.” It is that one deterministic integer core can cover several mathematically different jobs.

20.12 Common Pitfalls

20.12.1 Pitfall 1 — Forgetting Quadrant Reduction

Basic rotation mode does not converge over the full 360° range.

If you feed 210° directly into the core without reducing it first, the result will be wrong even if the stage math itself is perfect.

20.12.2 Pitfall 2 — Using Logical Right Shift on Signed Values

The shifts inside CORDIC are signed for x and y.

On AVR assembly:

```
asr    r25
ror    r24
```

is correct for a signed 16-bit right shift.

Using:

```
lsr    r25
ror    r24
```

would break negative vector components.

20.12.3 Pitfall 3 — Ignoring the Gain

CORDIC results are scaled. If you forget to preload with K in rotation mode, your sine/cosine amplitude will be too large. If you forget the gain in vectoring mode, your magnitude estimate will be off by a constant factor.

20.12.4 Pitfall 4 — Overflow in Narrow Formats

Q1.15 is ideal for sine/cosine outputs, but it is tight for vectoring-mode inputs.

If the input magnitude can approach 1.0, pre-scale by K or move to a wider fixed-point format before running vectoring CORDIC. Otherwise the internal $x = \text{magnitude} / K$ result can overflow even though both original components were inside $[-1, +1]$.

The same representation issue shows up at rotation-mode axis angles: $+1.0$ does not fit in signed Q1.15, so exact $\cos(0^\circ)$ or $\sin(90^\circ)$ must be saturated to $0x7FFF$ or handled as a special case.

20.12.5 Pitfall 5 — Hand-Writing the Loop Before Proving Correctness

CORDIC is simple enough to look obvious and subtle enough to get wrong.

On AVR, prove the math with a small set of known angles and compare the assembly against those expected values stage by stage. That is faster than debugging a broken fully-unrolled assembly version from scratch.

20.13 Summary

CORDIC on AVR:

- uses shifts + adds/subtracts + compare + small atan table
- no floating-point execution unit required
- fixed iteration count → deterministic timing

Useful formats:

- x, y = Q1.15 signed fractions
- angle = BAM16 (65536 units per turn)
- K_{Q15} = $0x4DBA$ for 12-stage rotation mode

Rotation mode:

- input = angle
- output = $\cos(\text{theta})$, $\sin(\text{theta})$

Vectoring mode:

- input = x, y vector
- output = angle and scaled magnitude

Main AVR tradeoff:

- slower than a small LUT for sine only
- more flexible than a LUT when you need sine/cosine/angle/magnitude

Build:

```
avr-gcc -mmcu=attiny3217 -nostartfiles -ndefaultlibs -o ch19.elf src/cordic_helpers.S
avr-objdump -h -d ch19.elf
```

This builds an inspection ELF for the helper routines. It is not a complete firmware image and should not be flashed by itself because it does not provide a reset vector or startup code.

Microchip source notes:

- ATtiny3217 product documentation lists an 8-bit AVR CPU with hardware multiplier, up to 20 MHz operation, 32 KB Flash, 2 KB SRAM, and 256 bytes EEPROM.
 - The AVR Instruction Set Manual device-core appendix lists ATtiny3217 as AVRxt. Its missing-instruction list does not include MUL, but the AVR instruction set has no general divide instruction.
 - The AVR Instruction Set Manual defines ASR as an arithmetic right shift and ROR as rotate right through carry; together they implement signed multi-byte right shifts for CORDIC stage terms.
-

20.14 Exercises

1. Implement an 8-stage CORDIC sine/cosine routine and compare its maximum error against the 12-stage version. How many cycles do you save?
 2. Store the BAM16 atan table in flash and read it with LPM through Z. Compare the flash cost and cycle cost against hard-coding the constants in a fully unrolled routine.
 3. Write an assembly helper that performs a signed 16-bit arithmetic right shift by a variable count in R17. Use ASR/ROR in a loop. How many cycles does it take for counts 0, 1, 4, and 8?
 4. Write a sine-only wrapper around the rotation-mode assembly core. Does a dedicated calling convention save any useful register moves, or are sine and cosine effectively “free” once the same 12 stages have already run?
 5. Implement vectoring mode for inputs pre-scaled by K and test it on a vector proportional to (0.6, 0.8). What BAM16 angle do you get, and how close is the corrected magnitude to 1.0?
 6. Compare three ways to generate a PWM sine on AVR:
 - a. 256-entry flash LUT
 - b. 12-stage CORDIC
 - c. fixed-point polynomial approximationMeasure flash usage with `avr-objdump -h` and inspect the call paths with `avr-objdump -d`.
-

Next: Appendices

Chapter 21

ADC and Analog Input

Digital inputs answer one question: is the pin low or high? Analog inputs answer a different question: how far is this voltage between ground and a reference voltage? The ADC converts that voltage into a number.

The ATtiny3217 has two 10-bit ADC peripherals. This chapter uses **ADC0** and a single external input channel. The goal is practical bring-up: choose a pin, choose a reference, choose an ADC clock, start a conversion, wait for RESRDY, and read the result correctly.

21.1 The ADC Mental Model

An ADC conversion has four parts:

1. Select an input channel with `ADC0_MUXPOS`.
2. Select a voltage reference and ADC clock with `ADC0_CTRLA`.
3. Start a conversion by writing `ADC_STCONV_bm` to `ADC0_COMMAND`.
4. Wait for `ADC_RESRDY_bm`, then read `ADC0_RESL` and `ADC0_RESN`.

For an ideal 10-bit single-ended conversion:

```
result = 1023 * Vin / Vref
Vin     = result * Vref / 1023
```

So with a 1.1 V reference:

```
0x000 = about 0 V
0x200 = about 0.55 V
0x3FF = about 1.1 V or higher
```

Input voltages above the selected reference saturate at the maximum result. The ADC does not tell you “too high”; it just returns the largest code.

21.2 ADC0 Register Path

The full register reference is in Appendix A. For a first polling conversion, these are the important registers:

Register	Purpose
VREF_CTRLA	Select internal reference voltage for ADC0
ADC0_CTRLA	Enable ADC, resolution, free-running mode

ADC0_CTRL	Reference source and ADC clock prescaler
ADC0_CTRLD	Initial delay and sample delay
ADC0_SAMPCTRL	Extra sample length for high-impedance sources
ADC0_MUXPOS	Positive input channel selection
ADC0_COMMAND	Start conversion
ADC0_INTFLAGS	Result-ready flag
ADC0_RES	Result low byte
ADC0_RES	Result high byte

All ADC0 registers are memory-mapped above low I/O space, so use LDS and STS, not IN and OUT.

21.3 Choosing an Analog Pin

The ADC input mux names channels as AIN0, AIN1, and so on. On the ATtiny3217:

AIN0..AIN7	PA0..PA7
AIN8	PB5
AIN9	PB4
AIN10	PB1
AIN11	PB0

For simple experiments, avoid:

- PA0, because it is also UPDI.
- PA3, because the Curiosity Nano LED is on PA3.
- pins already used by USART, SPI, or TWI examples in your project.

This chapter uses AIN4 on PA4.

Before using a pin as an analog input, disable its digital input buffer:

```
ldi    r16, 0x04
sts    PORTA_PIN4CTRL, r16    ; ISC = input buffer disabled
```

That disconnects the digital input path from the analog signal. It reduces power and avoids the digital input buffer reacting to slow voltages between logic low and logic high.

21.4 Choosing the Reference

The reference voltage defines the top of the conversion range. ADC0 can use:

- the internal reference from VREF
- VDD
- an external reference on VREFA

For repeatable examples, this chapter uses the internal 1.1 V reference:

```
ldi    r16, (0x01 << 4)    ; ADC0REFSEL = 1.1 V, bits 6:4
sts    VREF_CTRLA, r16

ldi    r16, VREF_ADC0REFEN_bm ; enable ADC0 reference output
sts    VREF_CTRLB, r16
```

Then select the internal reference source in ADC0_CTRL by leaving REFSEL[1:0] = 0.

The practical rule is simple:

- Use an internal reference when you want the result to mean volts.

- Use VDD when you want the result to be ratiometric with the supply.
- Use an external reference when the board provides a precise analog reference.

Ratiometric measurement is useful for sensors powered by the same supply as the MCU. If both the sensor output and ADC reference move with VDD, the code can stay stable even when the exact supply voltage changes.

21.5 Choosing the ADC Clock

For 10-bit resolution with a 1.1 V or higher reference, the datasheet specifies an ADC clock from 200 kHz to 1.5 MHz. The ADC clock comes from CLK_PER through the ADC prescaler:

PRES	Division
0	/2
1	/4
2	/8
3	/16
4	/32
5	/64
6	/128
7	/256

At the factory-default Curiosity Nano clock, CLK_PER is usually 20 MHz / 6, about 3.333 MHz. Good choices are:

DIV8	416.7 kHz
DIV16	208.3 kHz

This chapter uses DIV16:

```
ldi    r16, 0x03                ; REFSEL=internal, PRESC=DIV16
sts    ADC0_CTRL, r16
```

A normal conversion takes 13 ADC clocks, plus any configured sampling delay and sample length. At 208.3 kHz, 13 clocks is about 62 us before software overhead.

When using or changing an internal reference, give the reference and input mux time to settle before the first sample. ADC0_CTRLD provides INITDLY for that purpose. This chapter uses 32 ADC clocks:

```
ldi    r16, 0x40                ; INITDLY = 32 CLK_ADC cycles
sts    ADC0_CTRLD, r16
```

21.6 Sampling Time

The ADC does not measure the pin forever. It briefly connects the input to an internal sample-and-hold capacitor, then converts the held voltage.

Low-impedance sources charge that capacitor quickly. High-impedance sources need more time. The datasheet calls out about 10 kOhm or less as the easy case. For larger source impedances, increase ADC0_SAMPCTRL.

The sample-time formula is:

```
SampleTime = (2 + SAMPDLY + SAMPLEN) / CLK_ADC
```

This chapter uses a small extra sample length:

```
ldi    r16, 8
sts    ADC0_SAMPCTRL, r16
```

That is conservative for breadboard experiments, potentiometers, and simple sensor outputs. It costs conversion time, but it makes first measurements less fragile.

21.7 Polling Conversion Routine

The simplest ADC routine starts one conversion and waits until the result is ready:

```
adc0_read_ain4:
    ldi    r16, ADC_STCONV_bm
    sts    ADC0_COMMAND, r16

.Lwait_adc:
    lds    r16, ADC0_INTFLAGS
    sbrs   r16, ADC_RESRDY_bp
    rjmp   .Lwait_adc

    lds    r24, ADC0_RESL
    lds    r25, ADC0_RESR

    ldi    r16, ADC_RESRDY_bm
    sts    ADC0_INTFLAGS, r16
    ret
```

r25:r24 now holds the right-adjusted 10-bit result. The useful bits are:

```
r25:r24 = 000000xx xxxxxxxx
           high    low
```

Read low byte first. That matches the normal AVR order for 16-bit peripheral results.

The Microchip ADC examples clear RESRDY by writing a 1 to the flag bit after the conversion is complete. The routine above does the same with `ADC0_INTFLAGS = ADC_RESRDY_bm`. This keeps the polling pattern explicit and matches the write-one-to-clear style used by other AVR peripherals.

21.8 Complete Example: Read AIN4 and Drive an LED

The example reads PA4/AIN4. If the 10-bit result is at least half scale (≥ 512), it turns the Curiosity Nano LED off. If the result is below half scale, it turns the LED on.

The LED on PA3 is active-low:

```
PA3 = 0  LED on
PA3 = 1  LED off
```

Source file: `src/adc_poll.S`

```
#include <avr/io.h>

.equ LED_BIT,      3           ; PA3, active-low LED0
.equ ADC_PINCTRL,  0x04       ; PORT ISC input buffer disabled
.equ ADC_REF_1V1,  0x10       ; VREF ADC0REFSEL = 1.1 V, bits 6:4
.equ ADC_REFEN,    0x02       ; VREF ADC0REFEN
.equ ADC_PRESC_DIV16, 0x03    ; ADC0_CTRL: internal ref, prescaler /16
```

```

.equ ADC_INITDLY32, 0x40      ; ADC0_CTRL0D: 32 CLK_ADC startup delay
.equ ADC_AIN4, 0x04          ; PA4 / AIN4
.equ ADC_SAMPLEN8, 8

.section .vectors, "ax", @progbits
.global __vectors
__vectors:
    jmp    main
    .rept 30
    jmp    default_isr
    .endr

.section .text
.global main
main:
    ldi    r16, hi8(RAMEND)
    out    SPH, r16
    ldi    r16, lo8(RAMEND)
    out    SPL, r16

    sbi    VPORTA_DIR, LED_BIT
    sbi    VPORTA_OUT, LED_BIT      ; LED off

    ldi    r16, ADC_PINCTRL
    sts    PORTA_PIN4CTRL, r16      ; disable digital input on PA4

    ldi    r16, ADC_REF_1V1
    sts    VREF_CTRLA, r16          ; internal 1.1 V reference for ADC0

    ldi    r16, ADC_REFEN
    sts    VREF_CTRLB, r16          ; enable ADC0 reference output

    ldi    r16, ADC_PRESC_DIV16
    sts    ADC0_CTRL0C, r16

    ldi    r16, ADC_INITDLY32
    sts    ADC0_CTRL0D, r16

    ldi    r16, ADC_SAMPLEN8
    sts    ADC0_SAMPCTRL, r16

    ldi    r16, ADC_AIN4
    sts    ADC0_MUXPOS, r16

    ldi    r16, ADC_ENABLE_bm
    sts    ADC0_CTRLA, r16

.Lmain_loop:
    rcall  adc0_read_ain4

    sbrs   r25, 1                    ; bit 9 set means result >= 512
    rjmp   .Lbelow_half

    sbi    VPORTA_OUT, LED_BIT      ; LED off
    rjmp   .Lmain_loop

.Lbelow_half:

```

```

    cbi    VPORTA_OUT, LED_BIT        ; LED on
    rjmp   .Lmain_loop

adc0_read_ain4:
    ldi    r16, ADC_STCONV_bm
    sts    ADC0_COMMAND, r16

.Lwait_adc:
    lds    r16, ADC0_INTFLAGS
    sbrs   r16, ADC_RESRDY_bp
    rjmp   .Lwait_adc

    lds    r24, ADC0_RESL
    lds    r25, ADC0_RESR

    ldi    r16, ADC_RESRDY_bm
    sts    ADC0_INTFLAGS, r16
    ret

default_isr:
    reti

```

Build:

```

avr-gcc -mmcu=attiny3217 -nostartfiles -nodefaultlibs \
-o adc_poll.elf src/adc_poll.S
avr-objdump -d adc_poll.elf
avr-objcopy -O ihex adc_poll.elf adc_poll.hex

```

Hardware test:

- Connect a potentiometer between 1.1 V or a safe reference voltage and GND, with the wiper to PA4.
- Do not drive PA4 above VDD.
- If using the internal 1.1 V ADC reference, values above about 1.1 V saturate.

For a first bench test, a divided voltage below 1.1 V is easier to reason about than a full 0-5 V potentiometer.

21.9 Polling vs Interrupts

Polling is good when:

- the sample rate is low
- the code has nothing better to do
- the conversion is part of a simple foreground loop

ADC interrupts are better when:

- conversions happen periodically
- another peripheral or timer triggers sampling
- the main loop should keep doing other work

The interrupt path uses the same hardware flag. The difference is that you set `ADC_RESRDY_bm` in `ADC0_INTCTRL`, provide an `ADC0_RESRDY` ISR, and enable global interrupts with `SEI`.

The ISR must read the result and store it somewhere stable for the main loop. Do not do long filtering, formatting, serial printing, or display work inside the ADC ISR. Keep the ISR short and move processing into foreground code.

21.10 Accumulation and Simple Filtering

ADC0 can accumulate multiple samples in hardware. That reduces interrupt or polling overhead when you want an averaged reading, but it changes the scale of the result.

For example, if 16 samples are accumulated:

```
raw accumulated range = 0 .. 16 * 1023 = 0 .. 16368
average                = accumulated >> 4
```

Hardware accumulation is useful for noise reduction, but it is not magic. It helps with random noise. It does not fix:

- a bad reference
- a floating input
- a source impedance too high for the sample time
- digital switching noise coupled into the analog wiring
- sampling at exactly the wrong point in a periodic waveform

For slow sensors, a simple foreground low-pass filter is often enough:

```
filtered = filtered + ((new_sample - filtered) >> n)
```

With $n = 3$, each new sample contributes one eighth of the correction. That kind of filter belongs in the fixed-point/control material later in the book.

21.11 Free-Running Mode

In single-conversion mode, software writes `ADC_STCONV_bm` each time it wants a sample. In free-running mode, software starts the first conversion and hardware immediately starts the next one after each result.

The control change is small:

```
ldi    r16, (ADC_ENABLE_bm | ADC_FREERUN_bm)
sts    ADC0_CTRLA, r16

ldi    r16, ADC_STCONV_bm
sts    ADC0_COMMAND, r16      ; starts the continuous stream
```

The read side still waits for `RESRDY`, reads the result, and clears the flag:

```
adc0_read_latest:
.Lwait_adc:
    lds    r16, ADC0_INTFLAGS
    sbrs   r16, ADC_RESRDY_bp
    rjmp   .Lwait_adc

    lds    r24, ADC0_RESL
    lds    r25, ADC0_RESH

    ldi    r16, ADC_RESRDY_bm
    sts    ADC0_INTFLAGS, r16
    ret
```

Use free-running mode when you want a steady stream from one channel. It is less convenient when changing channels, because after a channel change the next result may still

reflect the previous channel selection. Discard one result after switching channels unless you have verified the timing carefully.

Source file: `src/adc_freerun.S`

21.12 Hardware Sample Accumulation

Hardware accumulation performs several conversions and adds the results into one result register value. Microchip's TB3209 examples use this pattern to average multiple samples with less software overhead.

For a 16-sample accumulation:

```
ldi    r16, 0x04                ; SAMPNUM = ACC16
sts    ADC0_CTRLB, r16
```

The raw result is now a sum:

```
raw range = 0 .. 16 * 1023
average   = raw >> 4
```

In assembly, right-shift the 16-bit result four times:

```
; r25:r24 = accumulated result
lsr    r25
ror    r24
lsr    r25
ror    r24
lsr    r25
ror    r24
lsr    r25
ror    r24
```

Hardware accumulation helps when noise is random and centered. It does not repair a floating input, poor grounding, a bad reference, or a source that cannot charge the sample-and-hold capacitor quickly enough.

Source file: `src/adc_accum16.S`

21.13 Window Comparator

The ADC window comparator lets hardware compare the result against a threshold or range. It can set WCMP and optionally request an interrupt when the result is below, above, inside, or outside the configured window.

For a simple low-threshold detector:

```
; WINLT = 0x0200 (half scale)
ldi    r16, 0x00
sts    ADC0_WINLTL, r16
ldi    r16, 0x02
sts    ADC0_WINLTH, r16

ldi    r16, 0x01                ; WINCM = BELOW
sts    ADC0_CTRLB, r16
```

After each conversion, check `ADC_WCMP_bm` in `ADC0_INTFLAGS`:


```
lds    r16, ADC0_INTFLAGS
sbrc   r16, ADC_WCMP_bp      ; skip next if WCMP = 0
rcall  handle_below_threshold

ldi    r16, (ADC_RESRDY_bm | ADC_WCMP_bm)
sts    ADC0_INTFLAGS, r16
```

The window comparator is useful when the CPU only cares about threshold crossings. It is also useful with accumulation: the comparison happens after the final sample in the accumulation burst, so thresholds apply to the accumulated value unless you scale them to match an averaged result.

Source file: `src/adc_window.S`

21.14 Event-Triggered Conversion

Software-triggered conversions are easy, but they happen whenever the foreground loop reaches the start instruction. Event-triggered conversions move that start signal into hardware. A timer, RTC/PIT, pin edge, or other event generator can start the ADC at a repeatable time.

The ADC side only needs one bit:

```
ldi    r16, ADC_STARTEI_bm
sts    ADC0_EVCTRL, r16      ; ADC starts on incoming event
```

Then EVSYS connects a generator to the ADC0 user. On ATtiny3217, asynchronous user 1 is ADC0. The user value 0x06 selects ASYNCCH3:

```
ldi    r16, 0x0D              ; ASYNCCH3 generator = PIT_DIV1024
sts    EVSYS_ASYNCCH3, r16

ldi    r16, 0x06              ; ADC0 listens to ASYNCCH3
sts    EVSYS_ASYNCUSER1, r16
```

Finally enable the RTC PIT source:

```
ldi    r16, ((9 << 3) | RTC_PITEN_bm) ; PIT CYC1024 period select + enable
sts    RTC_PITCTRLA, r16
```

After that, the foreground code can simply wait for RESRDY. No write to ADC0_COMMAND is needed for each sample; the PIT event starts conversions.

Event-triggered ADC is the clean path for periodic sampling because the sample time no longer depends on foreground loop timing.

Source file: `src/adc_event_pit.S`

21.15 Common Pitfalls

- Forgetting to disable the digital input buffer on the analog pin.
- Using a reference lower than the input signal and wondering why the result is always near 0x3FF.
- Reading RESH first and getting inconsistent 16-bit results.
- Sampling a high-impedance source with the default sample length.
- Switching channels in free-running mode and assuming the very next result is already from the new channel.

- Using IN/OUT for ADC registers. ADC0 is memory-mapped; use LDS/STS.
 - Forgetting that PA0 is UPDI and choosing it as a casual analog input.
-

21.16 Summary

ADC bring-up is mostly sequencing:

1. Configure the analog pin.
2. Select the reference.
3. Select a valid ADC clock.
4. Configure startup/sample timing.
5. Select the channel.
6. Enable ADC0.
7. Start a conversion.
8. Wait for RESRDY.
9. Read RESL, then RESH.
10. Clear RESRDY by writing 1 to the flag bit.

Once that path works, the more advanced features are extensions of the same model: interrupts replace polling, events replace software starts, accumulation replaces manual averaging, and the window comparator replaces software threshold checks.

21.17 Exercises

1. Change the example to use AIN5 on PA5. Which PORTA_PINnCTRL register and MUXPOS value change?
2. Use VDD as the ADC reference instead of the internal reference. What does the result mean if the sensor is also powered from VDD?
3. Increase ADC0_SAMPCTRL from 8 to 16. How does that change conversion time?
4. Convert the polling example into an interrupt-driven example that stores the latest sample in two SRAM bytes.
5. Add a 16-sample software average in the foreground loop. Which registers do you need for the accumulator?
6. Read the GND internal channel. What result should an ideal ADC produce? What might a real device produce?
7. Read the INTREF internal channel using a VDD reference. How can this be used to estimate supply voltage?

Chapter 22

ATtiny3217 Register Reference

Quick lookup for registers used throughout this book. The ATtiny3217 is an AVRXMEGA3 device. All peripheral registers are memory-mapped and require LDS/STS unless noted otherwise.

VPORT exception: Virtual PORT registers at 0x0000–0x001F are within the I/O space and can be accessed with IN/OUT, SBI/CBI, SBIS/SBIC in a single cycle. All other peripheral registers require LDS/STS.


CPU registers: SREG (0x003F), SPH (0x003E), SPL (0x003D) are accessible with IN/OUT using their I/O addresses directly (no +0x20 offset needed on AVRXMEGA3 — these addresses are already the memory-mapped addresses, and the I/O address equals the memory address for registers at 0x0000–0x003F).

22.1 CPU Registers

Register	Mem Addr	Access	Description
SREG	0x003F	IN/OUT	Status Register (I T H S V N Z C)
SPH	0x003E	IN/OUT	Stack Pointer High
SPL	0x003D	IN/OUT	Stack Pointer Low
CCP	0x0034	LDS/STS	Configuration Change Protection

SREG, SPH, and SPL sit within the range 0x0000–0x003F and are accessible via IN/OUT using the address directly. The CCP register requires LDS/STS.

22.1.1 SREG Bit Layout



22.1.2 Configuration Change Protection (CCP)

Certain registers are protected against accidental writes. Before writing to a protected register, write the appropriate key to CCP. The protected register must then be written within 4 CPU cycles.

Key	Value	Protects
-----	-------	----------

IOREG	0xD8	I/O registers (e.g. CLKCTRL, BOD, OSC settings)
SPM	0x9D	NVM (flash/EEPROM) write commands

Example — disable clock prescaler to run at 20 MHz:

```
ldi    r16, 0xD8          ; CCP key for I/O register changes
sts    0x0034, r16        ; write to CCP
sts    0x0061, r1         ; write 0 to CLKCTRL_MCLKCTRLB (clear PEN)
                        ; must complete within 4 cycles of CCP write
```

22.2 CLKCTRL — Clock Controller

Base address: 0x0060. All registers require LDS/STS and CCP protection where noted.

Register	Mem Addr	CCP?	Description
CLKCTRL_MCLKCTRLA	0x0060	Yes	Main clock source select
CLKCTRL_MCLKCTRLB	0x0061	Yes	Main clock prescaler
CLKCTRL_MCLKLOCK	0x0062	Yes	Lock register (prevents further changes)
CLKCTRL_MCLKSTATUS	0x0063	No	Status (read-only)

22.2.1 MCLKCTRLA — Main Clock Source Select

Bit:	7	6:2	1:0
	CLKOUT	—	CLKSEL

CLKSEL	Source
--------	--------

0x0	20 MHz internal oscillator ← factory default
0x1	32 KHz internal ULP oscillator
0x2	32.768 KHz external crystal (XOSC32K)
0x3	External clock on CLKI pin

Setting CLKOUT (bit 7) routes the main clock to the CLKOUT pin.

22.2.2 MCLKCTRLB — Main Clock Prescaler

Bit:	7:5	4:1	0
	—	PDIV	PEN

PEN Prescaler enable (1 = prescaler active, 0 = disabled)

PDIV Division factor (when PEN=1)

0x0	/2
0x1	/4
0x2	/8
0x3	/16
0x4	/32
0x5	/64
0x8	/6 ← factory default (20 MHz / 6 = 3.333 MHz)
0x9	/10

0xA /12
 0xB /24
 0xC /48

Factory default: PDIV=0x8 (divide by 6), PEN=1 → **3.333 MHz** system clock.

To run at full **20 MHz**: write 0x00 to MCLKCTRLB (clears PEN) — requires CCP key 0xD8 first. To run at **10 MHz**: write PEN=1, PDIV=0x0 (divide by 2).

22.2.3 MCLKSTATUS — Status (read-only)

Bit:	7	6	5	4	3	2	1	0
	EXTS	XOSC32KS	SOSC	OSC32KS	OSC20MS	—	—	SOSC

Bit name	Description
----------	-------------

OSC20MS	1 = 20 MHz oscillator is stable and running
OSC32KS	1 = 32 KHz oscillator is stable
SOSC	1 = main clock source switch is in progress
EXTS	1 = External clock source is active

22.3 SLPCTRL — Sleep Controller

Base address: 0x0050.

Register	Mem Addr	Description
SLPCTRL_CTRLA	0x0050	Sleep mode select and enable

22.3.1 CTRLA — Sleep Control

Bit:	7:3	2:1	0
	—	SMODE	SEN

SMODE	Sleep mode
-------	------------

0x0	Idle
0x1	Standby
0x2	Power-down

SEN	Sleep enable. Must be 1 before executing SLEEP.
-----	---

22.4 BOD — Brown-Out Detector / Voltage Level Monitor

Base address: 0x0080. Some BOD_CTRLA fields are loaded from fuses. The SLEEP field is CCP-protected when changed at run time.

Register	Mem Addr	Description
BOD_CTRLA	0x0080	BOD sample frequency and active/sleep modes
BOD_CTRLB	0x0081	BOD threshold level (fuse-loaded)
BOD_VLMCTRLA	0x0088	VLM threshold above BOD level
BOD_INTCTRL	0x0089	VLM interrupt enable/configuration
BOD_INTFLAGS	0x008A	VLM interrupt flag (W1C)

BOD_STATUS 0x008B VLM status

22.4.1 BOD_CTRLA

Bit:	7:5	4	3:2	1:0
	—	SAMPFREQ	ACTIVE	SLEEP

SAMPFREQ 0 = 1 kHz sampled mode, 1 = 125 Hz sampled mode

ACTIVE Active/Idle BOD mode (fuse-loaded, not changed by software)

0x0	Disabled
0x1	Continuous
0x2	Sampled
0x3	Continuous, wake-up halted until BOD is ready

SLEEP Standby/Power-down BOD mode (runtime writable through CCP)

0x0	Disabled
0x1	Continuous
0x2	Sampled

22.4.2 VLM Registers

BOD_VLMCTRLA.VLMLVL

0x0	VLM threshold 5% above BOD level
0x1	VLM threshold 15% above BOD level
0x2	VLM threshold 25% above BOD level

BOD_INTCTRL

bit 0	VLMIE: enable VLM interrupt
bits 2:1	VLMCFG: interrupt below, above, or crossing VLM threshold

BOD_INTFLAGS

bit 0	VLMIF: write 1 to clear
-------	-------------------------

22.5 GPIO Registers

The AVRXMEGA3 GPIO model is completely different from classic AVR. Each port has a full set of direction/output/input/control registers at a fixed base address. Atomic set/clear/toggle registers eliminate the need for read-modify-write sequences.

22.5.1 Port Base Addresses

Port	Base Addr
PORTA	0x0400
PORTB	0x0420
PORTC	0x0440

22.5.2 Per-Port Register Map (offset from base)

Offset	Register	Description
+0x00	PORTx_DIR	Data direction register (1=output, 0=input)
+0x01	PORTx_DIRSET	Write 1 to set direction bits (atomic)
+0x02	PORTx_DIRCLR	Write 1 to clear direction bits (atomic)

+0x03	PORTx_DIRTGL	Write 1 to toggle direction bits (atomic)
+0x04	PORTx_OUT	Output value register
+0x05	PORTx_OUTSET	Write 1 to set output bits (atomic)
+0x06	PORTx_OUTCLR	Write 1 to clear output bits (atomic)
+0x07	PORTx_OUTTGL	Write 1 to toggle output bits (atomic)
+0x08	PORTx_IN	Input value (reads current pin state)
+0x09	PORTx_INTFLAGS	Pin-change interrupt flags (W1C)
+0x0A	PORTx_PORTCTRL	Port control register
+0x10	PORTx_PIN0CTRL	Pin 0 config (pull-up, invert, ISC)
+0x11	PORTx_PIN1CTRL	Pin 1 config
+0x12	PORTx_PIN2CTRL	Pin 2 config
+0x13	PORTx_PIN3CTRL	Pin 3 config
+0x14	PORTx_PIN4CTRL	Pin 4 config
+0x15	PORTx_PIN5CTRL	Pin 5 config
+0x16	PORTx_PIN6CTRL	Pin 6 config
+0x17	PORTx_PIN7CTRL	Pin 7 config

22.5.3 PORTA Absolute Addresses

Register	Mem Addr	Description
PORTA_DIR	0x0400	Direction
PORTA_DIRSET	0x0401	Set direction (atomic)
PORTA_DIRCLR	0x0402	Clear direction (atomic)
PORTA_DIRTGL	0x0403	Toggle direction (atomic)
PORTA_OUT	0x0404	Output value
PORTA_OUTSET	0x0405	Set output (atomic)
PORTA_OUTCLR	0x0406	Clear output (atomic)
PORTA_OUTTGL	0x0407	Toggle output (atomic)
PORTA_IN	0x0408	Input value
PORTA_INTFLAGS	0x0409	Interrupt flags
PORTA_PORTCTRL	0x040A	Port control
PORTA_PIN0CTRL	0x0410	PA0 config (UPDI pin – do not reconfigure)
PORTA_PIN1CTRL	0x0411	PA1 config (SPI MOSI default)
PORTA_PIN2CTRL	0x0412	PA2 config (SPI MISO default)
PORTA_PIN3CTRL	0x0413	PA3 config (LED0 – Curiosity Nano)
PORTA_PIN4CTRL	0x0414	PA4 config (SPI SS default)
PORTA_PIN5CTRL	0x0415	PA5 config
PORTA_PIN6CTRL	0x0416	PA6 config
PORTA_PIN7CTRL	0x0417	PA7 config

22.5.4 PORTB Absolute Addresses

Register	Mem Addr	Description
PORTB_DIR	0x0420	Direction
PORTB_DIRSET	0x0421	Set direction (atomic)
PORTB_DIRCLR	0x0422	Clear direction (atomic)
PORTB_DIRTGL	0x0423	Toggle direction (atomic)
PORTB_OUT	0x0424	Output value
PORTB_OUTSET	0x0425	Set output (atomic)
PORTB_OUTCLR	0x0426	Clear output (atomic)
PORTB_OUTTGL	0x0427	Toggle output (atomic)
PORTB_IN	0x0428	Input value
PORTB_INTFLAGS	0x0429	Interrupt flags
PORTB_PORTCTRL	0x042A	Port control
PORTB_PIN0CTRL	0x0430	PB0 config (TWI SDA default)
PORTB_PIN1CTRL	0x0431	PB1 config (TWI SCL default)


PORTB_PIN2CTRL	0x0432	PB2 config (USART0 TX – Curiosity Nano)
PORTB_PIN3CTRL	0x0433	PB3 config (USART0 RX – Curiosity Nano)
PORTB_PIN4CTRL	0x0434	PB4 config
PORTB_PIN5CTRL	0x0435	PB5 config
PORTB_PIN6CTRL	0x0436	PB6 config
PORTB_PIN7CTRL	0x0437	PB7 config

22.5.5 PORTC Absolute Addresses

Register	Mem Addr	Description
PORTC_DIR	0x0440	Direction
PORTC_DIRSET	0x0441	Set direction (atomic)
PORTC_DIRCLR	0x0442	Clear direction (atomic)
PORTC_DIRTGL	0x0443	Toggle direction (atomic)
PORTC_OUT	0x0444	Output value
PORTC_OUTSET	0x0445	Set output (atomic)
PORTC_OUTCLR	0x0446	Clear output (atomic)
PORTC_OUTTGL	0x0447	Toggle output (atomic)
PORTC_IN	0x0448	Input value
PORTC_INTFLAGS	0x0449	Interrupt flags
PORTC_PORTCTRL	0x044A	Port control
PORTC_PIN0CTRL	0x0450	PC0 config
PORTC_PIN1CTRL	0x0451	PC1 config
PORTC_PIN2CTRL	0x0452	PC2 config
PORTC_PIN3CTRL	0x0453	PC3 config
PORTC_PIN4CTRL	0x0454	PC4 config
PORTC_PIN5CTRL	0x0455	PC5 config
PORTC_PIN6CTRL	0x0456	PC6 config (not present on 20-pin package)
PORTC_PIN7CTRL	0x0457	PC7 config (not present on 20-pin package)

22.5.6 PINnCTRL Bit Layout

Each PORTx_PINnCTRL register configures one pin independently.



ISC[2:0] Input Sense Configuration

0x0	INTDISABLE	No interrupt; digital input buffer enabled
0x1	BOTHEDGES	Interrupt on both rising and falling edges
0x2	RISING	Interrupt on rising edge only
0x3	FALLING	Interrupt on falling edge only
0x4	INPUT_DISABLE	Digital input buffer disabled (reduces power)
0x5	LEVEL	Interrupt on low level

22.5.7 Virtual PORT Registers (VPORT) – Fast I/O

VPORTs mirror the key PORT registers into I/O space (0x0000–0x001F). They can be accessed with IN/OUT (1 cycle), SBI/CBI, and SBIS/SBIC. Writing to VPORT_IN toggles the corresponding output bit (same as OUTTGL).

Register	Addr	Mirrors
----------	------	---------

VPORTA_DIR	0x0000	PORTA_DIR	(1=output)
VPORTA_OUT	0x0001	PORTA_OUT	
VPORTA_IN	0x0002	PORTA_IN	(write 1 to toggle output)
VPORTA_INTFLAGS	0x0003	PORTA_INTFLAGS	
VPORTB_DIR	0x0004	PORTB_DIR	
VPORTB_OUT	0x0005	PORTB_OUT	
VPORTB_IN	0x0006	PORTB_IN	
VPORTB_INTFLAGS	0x0007	PORTB_INTFLAGS	
VPORTC_DIR	0x0008	PORTC_DIR	
VPORTC_OUT	0x0009	PORTC_OUT	
VPORTC_IN	0x000A	PORTC_IN	
VPORTC_INTFLAGS	0x000B	PORTC_INTFLAGS	

Use VPORT registers with IN/OUT for single-cycle access. Use the full PORT registers with LDS/STS when you need atomic set/clear/toggle or pin config.

22.6 Curiosity Nano Pin Assignments

Signal	Port Pin	Mem Addr (PINnCTRL)	Notes
LED0	PA3	0x0413	Active low – drive LOW to light
USART0 TX	PB2	0x0432	To on-board debugger virtual COM port
USART0 RX	PB3	0x0433	From on-board debugger virtual COM port
UPDI	PA0	0x0410	Programming/debug – do not reconfigure

22.6.1 LED and Button Quick Reference

LED0 on (light): STS PORTA_OUTCLR, (1<<3) ; clear PA3 (active low)
 LED0 off (dark): STS PORTA_OUTSET, (1<<3) ; set PA3
 LED0 toggle: STS PORTA_OUTTGL, (1<<3) ; toggle PA3

The ATtiny3217 Curiosity Nano does not provide a dedicated user pushbutton. For button examples, wire an external active-low button to a free GPIO and enable that pin's internal pull-up.

22.7 TCA0 — Timer/Counter Type A (16-bit)

Base address: 0x0A00 (single-slope mode, the default). TCA0 is a general-purpose 16-bit timer supporting normal, CTC, and PWM modes.

Register	Mem Addr	Description
TCA0_SINGLE_CTRLA	0x0A00	Prescaler select (CLKSEL), enable (ENABLE)
TCA0_SINGLE_CTRLB	0x0A01	Waveform mode (WGMODE), compare enables
TCA0_SINGLE_CTRLC	0x0A02	Force compare outputs
TCA0_SINGLE_CTRLD	0x0A03	Split mode enable
TCA0_SINGLE_CTRLCLR	0x0A04	Control E clear
TCA0_SINGLE_CTRLSET	0x0A05	Control E set
TCA0_SINGLE_INTCTRL	0x0A06	Interrupt enable (OVF, CMP0, CMP1, CMP2)
TCA0_SINGLE_INTFLAGS	0x0A07	Interrupt flags (W1C)
TCA0_SINGLE_EVCTRL	0x0A09	Event control
TCA0_SINGLE_DBGCTRL	0x0A0E	Debug control

TCA0_SINGLE_TEMP	0x0A0F	Temp register for 16-bit read/write
TCA0_SINGLE_CNT	0x0A20	Counter value (16-bit, write low byte first)
TCA0_SINGLE_PER	0x0A26	Period / TOP value (16-bit)
TCA0_SINGLE_CMP0	0x0A28	Compare/capture 0 (16-bit)
TCA0_SINGLE_CMP1	0x0A2A	Compare/capture 1 (16-bit)
TCA0_SINGLE_CMP2	0x0A2C	Compare/capture 2 (16-bit)
TCA0_SINGLE_PERBUF	0x0A36	Period buffer (double-buffered)
TCA0_SINGLE_CMP0BUF	0x0A38	Compare 0 buffer
TCA0_SINGLE_CMP1BUF	0x0A3A	Compare 1 buffer
TCA0_SINGLE_CMP2BUF	0x0A3C	Compare 2 buffer

22.7.1 CTRLA — Prescaler and Enable

Bit: 7:4 3:1 0
 — CLKSEL ENABLE

CLKSEL	Prescaler
0x0	DIV1 (no prescaling)
0x1	DIV2
0x2	DIV4
0x3	DIV8
0x4	DIV16
0x5	DIV64
0x6	DIV256
0x7	DIV1024

Example — 1 Hz overflow at 3.333 MHz with DIV64 (prescaled clock = 52,083 Hz, PER = 52082):

```
ldi r16, 0x0B           ; CLKSEL=DIV64 (bits 3:1 = 101), ENABLE=1
sts 0x0A00, r16         ; TCA0_SINGLE_CTRLA
ldi r16, lo8(52082)
sts 0x0A26, r16         ; TCA0_SINGLE_PER low byte
ldi r16, hi8(52082)
sts 0x0A27, r16         ; TCA0_SINGLE_PER high byte
ldi r16, 0x01           ; OVF interrupt enable
sts 0x0A06, r16         ; TCA0_SINGLE_INTCTRL
```

22.7.2 CTRLB — Waveform Mode

WGMODE[2:0] (bits 2:0)	Mode
0x0	Normal (count up, overflow at PER)
0x1	Frequency (toggle W0, reset at CMP0)
0x2	Single-slope PWM
0x5	Dual-slope PWM (overflow at BOTTOM and TOP)
0x6	Dual-slope PWM (overflow at TOP only)
0x7	Dual-slope PWM (overflow at BOTTOM only)

22.8 TCB0 / TCB1 — Timer/Counter Type B (16-bit)

Type B timers are 16-bit, single-channel timers suited for periodic interrupts, input capture, and PWM. TCB0 and TCB1 share the same register layout.

TCB0 base: 0x0A40
 TCB1 base: 0x0A50

Offset	Register	Description
+0x00	TCBn_CTRLA	Clock select, enable
+0x01	TCBn_CTRLB	Compare/capture mode, output enable
+0x04	TCBn_EVCTRL	Event control
+0x05	TCBn_INTCTRL	CAPT interrupt enable (bit 0)
+0x06	TCBn_INTFLAGS	CAPT interrupt flag (bit 0, W1C)
+0x07	TCBn_STATUS	RUN flag (bit 0, read-only)
+0x08	TCBn_DBGCTRL	Debug control
+0x09	TCBn_TEMP	Temp register for 16-bit read/write
+0x0A	TCBn_CNTL	Counter low byte
+0x0B	TCBn_CNTH	Counter high byte
+0x0C	TCBn_CCMP_L	Compare/capture low byte
+0x0D	TCBn_CCMP_H	Compare/capture high byte

22.8.1 CTRLA — Clock Select and Enable

Bit: 7:3 2:1 0
 — CLKSEL ENABLE

CLKSEL Clock source

0x0	CLKPER	(same as system clock)
0x1	CLKPER/2	
0x2	TCA0	(TCB clocked from TCA0 overflow)

22.8.2 CTRLB — Mode (CMODE)

CMODE[2:0] (bits 2:0) Mode

0x0	Periodic Interrupt ← default, most common
0x1	Timeout Check
0x2	Input Capture on Event
0x3	Input Capture Frequency Measurement
0x4	Input Capture Pulse-Width Measurement
0x5	Input Capture Freq+PW Measurement
0x6	Single Shot
0x7	8-bit PWM

In Periodic Interrupt mode (CMODE=0), the counter counts up to CCMP, fires the CAPT interrupt, then resets to zero. CCMP sets the period (TOP).

Example — 1 kHz periodic interrupt at 3.333 MHz:

```
ldi r16, lo8(3332)      ; CCMP = (f_cpu / f_irq) - 1 = 3332
sts 0x0A4C, r16         ; TCB0_CCMP_L
ldi r16, hi8(3332)
sts 0x0A4D, r16         ; TCB0_CCMP_H
ldi r16, 0x01           ; CAPT interrupt enable
sts 0x0A45, r16         ; TCB0_INTCTRL
ldi r16, 0x01           ; CLKSEL=CLKPER, ENABLE=1
sts 0x0A40, r16         ; TCB0_CTRLA
```

22.9 USART0

Base address: 0x0800. On the Curiosity Nano board, USART0 is connected to the on-board debugger (PB2=TX, PB3=RX), providing a virtual COM port over USB.

Register	Mem Addr	Description
USART0_RXDATA	0x0800	Receive data (low byte, 8-bit data here)
USART0_RXDATAH	0x0801	Receive data high + error flags
USART0_TXDATA	0x0802	Transmit data (write byte here to send)
USART0_TXDATAH	0x0803	Transmit data high (9-bit mode only)
USART0_STATUS	0x0804	Status flags
USART0_CTRLA	0x0805	Interrupt enables
USART0_CTRLB	0x0806	Receiver/transmitter enable
USART0_CTRLC	0x0807	Frame format (default = 8N1)
USART0_BAUDL	0x0808	Baud rate register low byte
USART0_BAUDH	0x0809	Baud rate register high byte
USART0_CTRLD	0x080A	Auto-baud, IRCOM
USART0_DBGCTRL	0x080B	Debug control
USART0_EVCTRL	0x080C	Event control
USART0_TXPLCTRL	0x080D	TX pin LINAUTO
USART0_RXPLCTRL	0x080E	RX pin LINAUTO

22.9.1 STATUS Bits

Bit: 7 6 5 4 3 2 1:0
 RXCIF TXCIF DREIF FERR BUFOVF PERR –

RXCIF Receive Complete Interrupt Flag (1 = unread data in RXDATA)
 TXCIF Transmit Complete Interrupt Flag (1 = TX shift register empty)
 DREIF Data Register Empty Flag (1 = OK to write next byte to TXDATA)
 FERR Frame Error
 BUFOVF Buffer Overflow
 PERR Parity Error

22.9.2 CTRLA — Interrupt Enables

Bit: 7 6 5 4:2 1 0
 RXCIE TXCIE DREIE – LBME ABEIE

RXCIE Receive Complete IE
 TXCIE Transmit Complete IE
 DREIE Data Register Empty IE (triggers when ready for next transmit byte)

22.9.3 CTRLB — Enable Bits

Bit: 7 6 5:4 3 2 1:0
 RXEN TXEN MPCM RXMODE ODME –

RXEN 1 = enable receiver
 TXEN 1 = enable transmitter
 RXMODE 0x0=NORMAL, 0x1=CLK2X, 0x2=GENAUTO, 0x3=LINAUTO

22.9.4 CTRLC — Frame Format

Default value: 0x03 = 8N1 (8 data bits, no parity, 1 stop bit)

Bit: 7:6 5:4 3 2:0
 CMODE PMODE SBMODE CHSIZE

CMODE 0x0=ASYNCHRONOUS, 0x1=SYNCHRONOUS, 0x2=IRCOM, 0x3=MSPI
 PMODE 0x0=DISABLED, 0x2=EVEN, 0x3=ODD

SBMODE 0=1 stop bit, 1=2 stop bits

CHSIZE 0x0=5-bit, 0x1=6-bit, 0x2=7-bit, 0x3=8-bit, 0x4-0x6=reserved, 0x7=9-bit

22.9.5 Baud Rate Formula and Common Values

Normal asynchronous mode:

$$\begin{aligned}\text{BAUD_REG} &= (\text{f_cpu} \times 64) / (16 \times \text{baud_rate}) \\ &= (\text{f_cpu} \times 4) / \text{baud_rate}\end{aligned}$$

Baud Rate	f_cpu = 3.333 MHz BAUD_REG (hex)	f_cpu = 20 MHz BAUD_REG (hex)
9600	0x056D (1389)	0x2098 (8344)
19200	0x02B7 (694)	0x104C (4172)
38400	0x015B (347)	0x0826 (2086)
57600	0x00E2 (226)	0x0571 (1393)
115200	0x0074 (116)	0x02B9 (697)

22.10 SPI0

Base address: 0x08C0. Default pin assignments: MOSI=PA1, MISO=PA2, SCK=PA3, SS=PA4.

Register	Mem Addr	Description
SPI0_CTRLA	0x08C0	Control A: DORD, MASTER, CLK2X, PRESC[1:0], ENABLE
SPI0_CTRLB	0x08C1	Control B: BUFEN, SSD, MODE[1:0]
SPI0_INTCTRL	0x08C2	Interrupt enables: IE, SSIE, TXCIE, RXCIE
SPI0_INTFLAGS	0x08C3	Interrupt flags: IF(7), WRCOL(0)
SPI0_DATA	0x08C4	TX/RX data register

22.10.1 CTRLA Bits

Bit:	7	6	5	4:3	2	1	0
	DORD	MASTER	CLK2X	PRESC	—	—	ENABLE

DORD 0=MSB first, 1=LSB first

MASTER 0=slave, 1=master

CLK2X 1=double clock speed

PRESC 0x0=DIV4, 0x1=DIV16, 0x2=DIV64, 0x3=DIV128

ENABLE 1=enable SPI

22.10.2 SPI Clock Rates

PRESC	CLK2X=0	CLK2X=1
0x0	f_cpu / 4	f_cpu / 2
0x1	f_cpu / 16	f_cpu / 8
0x2	f_cpu / 64	f_cpu / 32
0x3	f_cpu / 128	f_cpu / 64

22.11 TWI0 — Two-Wire Interface (I2C)

Base address: 0x08A0. Default pins: SDA=PB0, SCL=PB1. Requires external 4.7 kΩ pull-up resistors on SDA and SCL lines.

Register	Mem Addr	Description
TWI0_CTRLA	0x08A0	SDA hold time, fast-mode plus enable
TWI0_DUALCTRL	0x08A1	Dual-mode control
TWI0_MCTRLA	0x08A2	Master: RIEN, WIEN, QCEN, TIMEOUT, SMEN, ENABLE
TWI0_MCTRLB	0x08A3	Master: FLUSH, ACKACT, MCMD[1:0]
TWI0_MSTATUS	0x08A4	Master status flags
TWI0_MBAUD	0x08A5	Master baud rate
TWI0_MADDR	0x08A6	Master address (write triggers START condition)
TWI0_MDATA	0x08A7	Master data register (read/write)
TWI0_SCTRLA	0x08A8	Slave: DIEN, APIEN, PIEN, PMEN, SMEN, ENABLE
TWI0_SCTRLB	0x08A9	Slave: ACKACT, SCMD[1:0]
TWI0_SSTATUS	0x08AA	Slave status flags
TWI0_SADDR	0x08AB	Slave address
TWI0_SDATA	0x08AC	Slave data
TWI0_SADDRMASK	0x08AD	Slave address mask

22.11.1 MSTATUS Bits

Bit:	7	6	5	4	3	2	1:0
	RIF	WIF	CLKHOLD	RXACK	ARBLIST	BUSERR	BUSSTATE

RIF	Read Interrupt Flag
WIF	Write Interrupt Flag
CLKHOLD	Clock hold (SCL stretched by master)
RXACK	Received ACK (0=ACK received, 1=NACK)
ARBLIST	Arbitration lost
BUSERR	Bus error
BUSSTATE	0x0=UNKNOWN, 0x1=IDLE, 0x2=OWNER, 0x3=BUSY

22.11.2 MCTRLB — Master Command

MCMD[1:0]	Command
0x0	NOACT No action
0x1	REPSTART Issue repeated START
0x2	RECVTRANS Continue (ACK) / send byte
0x3	STOP Issue STOP condition

22.11.3 Baud Rate Formula

$$\text{BAUD} = (\text{f_cpu} / \text{f_scl} - 10) / 2 \quad (\text{for f_cpu in Hz, f_scl in Hz})$$

100 kHz at 3.333 MHz:	BAUD = (3333333 / 100000 - 10) / 2 ≈ 12
400 kHz at 3.333 MHz:	BAUD = (3333333 / 400000 - 10) / 2 ≈ 2 (use fast-mode+)
100 kHz at 20 MHz:	BAUD = (20000000 / 100000 - 10) / 2 = 95
400 kHz at 20 MHz:	BAUD = (20000000 / 400000 - 10) / 2 = 20

22.12 ADC0

Base address: 0x0600. The ATtiny3217 ADC is 10-bit (or 8-bit selectable).

Register	Mem Addr	Description
ADC0_CTRLA	0x0600	Enable, resolution, free-run, left-adjust
ADC0_CTRLB	0x0601	Sample accumulation (SAMPNUM)
ADC0_CTRLC	0x0602	Prescaler, reference select, sample capacitor
ADC0_CTRLD	0x0603	Init delay, auto-sample delay
ADC0_CTRLF	0x0604	Window comparator mode
ADC0_SAMPCTRL	0x0605	Sample length extension
ADC0_MUXPOS	0x0606	Input mux (selects channel)
ADC0_COMMAND	0x0608	Write bit 0 (STCONV) to start conversion
ADC0_EVCTRL	0x0609	Event trigger enable
ADC0_INTCTRL	0x060A	RESRDY(0) and WCMP(1) interrupt enables
ADC0_INTFLAGS	0x060B	Interrupt flags (W1C)
ADC0_DBGCTRL	0x060C	Debug run control
ADC0_TEMP	0x060D	Temporary register
ADC0_RESL	0x0610	Result low byte
ADC0_RESB	0x0611	Result high byte
ADC0_WINLT	0x0612	Window comparator low threshold
ADC0_WINHT	0x0614	Window comparator high threshold
ADC0_CALIB	0x0616	Calibration

22.12.1 CTRLA Bits

Bit:	7	6:5	4	3	2	1	0
	—	—	LEFTADJ	RESSEL	FREERUN	RUNSTDBY	ENABLE

LEFTADJ 1=result left-adjusted in 16-bit register (8-bit in RESB)
RESSEL 0=10-bit result, 1=8-bit result
FREERUN 1=continuous conversions
ENABLE 1=enable ADC

22.12.2 CTRLC — Prescaler and Reference

Bit:	7:6	5:4	3	2:0
	—	REFSEL	SAMPCAP	PRESC

REFSEL 0x0=INTREF (internal), 0x1=VDDI02/10, 0x2=reserved, 0x3=VREFA (external)
SAMPCAP 1=reduced sample capacitance (for high-impedance sources)
PRESC ADC clock prescaler (see table below)

PRESC Division

0x0	DIV2
0x1	DIV4
0x2	DIV8
0x3	DIV16
0x4	DIV32
0x5	DIV64
0x6	DIV128
0x7	DIV256

ADC requires 50–1500 kHz clock. At 3.333 MHz use DIV8 (417 kHz) or DIV16 (208 kHz).
At 20 MHz use DIV32 (625 kHz) or DIV64 (312 kHz).

22.12.3 MUXPOS — Input Mux

MUXPOS Channel

0x00	AIN0 (PA0 – also UPDI, avoid)
------	-------------------------------

0x01	AIN1 (PA1)
0x02	AIN2 (PA2)
0x03	AIN3 (PA3)
0x04	AIN4 (PA4)
0x05	AIN5 (PA5)
0x06	AIN6 (PA6)
0x07	AIN7 (PA7)
0x08	AIN8 (PB5)
0x09	AIN9 (PB4)
0x0A	AIN10 (PB1)
0x0B	AIN11 (PB0)
0x1B	Reserved for ADC0 / PTC
0x1C	DAC0
0x1D	INTREF (internal reference voltage)
0x1E	TEMPSENSE (internal temperature sensor)
0x1F	GND

22.12.4 Internal Reference Voltage

The internal reference is configured via VREF peripheral (base 0x00A0):

Register	Mem Addr	Description
VREF_CTRLA	0x00A0	ADC0 and DAC0 reference select
VREF_CTRLB	0x00A1	ADC0 reference force enable

VREF_CTRLA bits 6:4 (ADC0REFSEL):

0x0	0.55V
0x1	1.1V
0x2	2.5V
0x3	4.3V
0x4	1.5V

Default internal reference: 1.1V.

22.13 NVMCTRL — Non-Volatile Memory Controller

Base address: 0x1000. Writing flash or EEPROM requires writing the CCP_SPM key (0x9D) to CCP before the command, then completing the write within 4 cycles.

Register	Mem Addr	Description
NVMCTRL_CTRLA	0x1000	Command register (write CCP key first)
NVMCTRL_CTRLB	0x1001	Boot lock, application code write protect
NVMCTRL_STATUS	0x1002	Status flags (read-only)
NVMCTRL_INTCTRL	0x1003	EEREADY interrupt enable
NVMCTRL_INTFLAGS	0x1004	EEREADY interrupt flag (W1C)
NVMCTRL_DATA	0x1006	Data register (16-bit)
NVMCTRL_ADDR	0x1008	Address register (16-bit)

22.13.1 CTRLA — NVM Commands

CMD[2:0] (bits 2:0)	Command
0x00	NOCMD
0x01	PAGEWRITE
0x02	PAGEERASE

0x03	PAGEERASEWRITE	Erase and write in one operation
0x04	PAGEBUFCLR	Clear page write buffer
0x05	CHIPERASE	Erase entire chip (flash + EEPROM)
0x06	EEERASE	Erase EEPROM
0x07	FUSEWRITE	Write fuse bytes

22.13.2 STATUS Bits

Bit:	2	1	0
	WRERROR	EEBUSY	FBUSY

FBUSY 1 = flash write/erase in progress
 EEBUSY 1 = EEPROM write/erase in progress
 WRERROR 1 = write error occurred

22.13.3 Memory Layout

Region	Start	End	Size	Notes
Flash	0x0000	0x7FFF	32 KB	Program memory (word-addressed by PC)
SRAM	0x3800	0x3FFF	2 KB	Data memory (avr-ld origin 0x803800)
EEPROM	0x1400	0x14FF	256 B	(memory-mapped for read; use NVMCTRL to write)
Fuses	0x1280	0x128A		Written via UPDI/avrdude

Flash page size: 64 bytes. EEPROM page size: 32 bytes.

22.14 Interrupt Vector Table

All vectors are 4-byte (2-word) entries. Use JMP (4-byte instruction) for targets anywhere in the 32 KB flash, or RJMP + NOP (2+2 bytes) to keep the same slot size.

Vector	Byte	Addr	Name	Description
0	0x0000		RESET	Power-on / external / WDT reset
1	0x0004		CRCSCAN_NMI	CRC scan non-maskable interrupt
2	0x0008		BOD_VLM	Brown-out detection voltage level monitor
3	0x000C		PORTA_PORT	Port A pin-change interrupt
4	0x0010		PORTB_PORT	Port B pin-change interrupt
5	0x0014		PORTC_PORT	Port C pin-change interrupt
6	0x0018		RTC_CNT	RTC overflow / compare match
7	0x001C		RTC_PIT	RTC periodic interrupt timer
8	0x0020		TCA0_OVF / TCA0_LUNF	TCA0 overflow (single) / low underflow (split)
9	0x0024		TCA0_HUNF	TCA0 high byte underflow (split mode)
10	0x0028		TCA0_CMP0 / TCA0_LCMP0	TCA0 compare 0 match
11	0x002C		TCA0_CMP1 / TCA0_LCMP1	TCA0 compare 1 match
12	0x0030		TCA0_CMP2 / TCA0_LCMP2	TCA0 compare 2 match
13	0x0034		TCB0_INT	TCB0 capture interrupt
14	0x0038		TCB1_INT	TCB1 capture interrupt
15	0x003C		TWI0_TWIS	TWI0 slave interrupt
16	0x0040		TWI0_TWIM	TWI0 master interrupt
17	0x0044		SPI0_INT	SPI0 transfer complete
18	0x0048		USART0_RXC	USART0 receive complete
19	0x004C		USART0_DRE	USART0 data register empty
20	0x0050		USART0_TXC	USART0 transmit complete
21	0x0054		AC0_AC	Analog comparator interrupt
22	0x0058		ADC0_RESRDY	ADC0 result ready

23	0x005C	ADC0_WCMP	ADC0 window comparator match
24	0x0060	CCL_CCL	Configurable custom logic interrupt
25	0x0064	NVMCTRL_EE	EEPROM ready

Minimal startup with vector table in assembly:

```
.section .vectors,"ax",@progbits
jmp _start          ; 0x0000 RESET – must be first
jmp _unhandled      ; 0x0004 CRCSCAN_NMI
; ... one jmp per vector ...

.section .text
_unhandled:
    rjmp _unhandled    ; spin on unhandled interrupts

_start:
    ; ... init stack, SRAM, then jump to main
```

22.15 Fuses

Fuses are one-time-programmable configuration bytes written via UPDI using avrdude. Factory defaults are shown. Unlike ATmega, there is no separate CKDIV, BOOTRST, or BOOTSZ fuse.

Fuse	Mem Addr	Default	Description
FUSE0	0x1280	0x00	APPEND[7:0] – application code end page
FUSE1	0x1281	0x00	BOOTEND[7:0] – boot section end page (0=none)
FUSE2	0x1282	0x02	OSCCFG – oscillator config
FUSE4	0x1284	0xF6	SYSCFG0 – CRC, reset pin, UPDI pin config
FUSE5	0x1285	0x07	SYSCFG1 – startup time SUT[2:0]
FUSE6	0x1286	0x00	(same as FUSE0 on this device)
FUSE7	0x1287	0x00	(same as FUSE1)
FUSE8	0x1288	0xAA	LOCKBIT

22.15.1 FUSE2 – OSCCFG

Bit: 1 0
FREQSEL FREQSEL

FREQSEL 0=16 MHz internal oscillator base
1=20 MHz internal oscillator base ← factory default (0x02)

The clock prescaler in CLKCTRL_MCLKCTRLB further divides this frequency. Factory default: 20 MHz / 6 = 3.333 MHz.

22.15.2 FUSE4 – SYSCFG0

Bit: 7:6 5:4 3:2 1:0
CRCSRC RSTPINCFG UPDIPINCFG –

RSTPINCFG 0x0=GPIO, 0x1=UPDI, 0x2=RESET (0xF6 default = RESET=GPIO, UPDI active)
CRCSRC 0x3=no CRC check (default)

22.15.3 Writing Fuses with avrdude

```
avrdude -c serialupdi -p t3217 -P /dev/ttyUSB0 -b 115200 \
```

```
-U fuse2:w:0x02:m
```

```
avrdude -c serialupdi -p t3217 -P /dev/ttyUSB0 -b 115200 \  
-U fuse0:w:0x00:m -U fuse1:w:0x00:m -U fuse2:w:0x02:m \  
-U fuse4:w:0xF6:m -U fuse5:w:0x07:m
```

Warning: Setting RSTPINCFG incorrectly can disable the UPDI programming interface, effectively bricking the device. Never modify FUSE4 unless certain.

Chapter 23

AVR Instruction Set Summary

Complete instruction set for AVR (classic core, enhanced core with MUL). Columns: **Op-code** — mnemonic and operands; **Operation** — what it does; **Cycles** — clock cycles (branch cycles shown as taken/not-taken); **Flags** — SREG bits modified (I T H S V N Z C).

Rd = destination register (R0-R31 unless noted). Rr = source register. K = 8-bit constant (0-255). k = relative branch offset (−64..63 for 7-bit, −2048..2047 for 12-bit). b = bit index (0-7). A = I/O address (0-63). q = displacement (0-63) for Y/Z indirect with displacement.

23.1 How to Read the AVR Instruction Set Manual

This appendix is a quick reference. The Microchip instruction set manual is the authoritative source. If you only use the summary table, you can write code. If you learn to read the manual pages directly, you can answer harder questions:

- Why does this instruction only work on r16–r31?
- Why is this branch 1 cycle sometimes and 2 cycles other times?
- Why does SBIS skip differently from BRNE?
- Why does OUT work here but STS is required there?
- Why does an instruction update Z but not C?

The manual is not hard once you know what each field is trying to tell you. The problem is that it reads like a hardware document, not a tutorial.

23.1.1 What One Instruction Page Contains

Each instruction in the manual is usually presented in the same structure:

- **Description / operation:** what value changes logically.
- **Syntax:** the assembly form, such as ADD Rd, Rr.
- **Operands:** legal ranges, such as 0 <= d <= 31.
- **Program Counter:** how PC changes after execution.
- **Opcode:** the encoded bit pattern.
- **Status Register (SREG) and Boolean Formula:** which flags change and why.
- **Words / Cycles:** code size and execution time, often by AVR core family.

When reading a page, do not read it top to bottom like prose. Read it in this order instead:

1. **Syntax:** what form the instruction has.
2. **Operands:** which registers or constants are legal.
3. **Operation:** what value is computed.
4. **SREG effects:** which later branches or compares it can support.

5. **Words / Cycles**: whether it is a good fit for tight loops.
6. **Opcode**: only when you want to understand encoding limits or disassembly.

That order matches how you actually use the information while programming.

23.1.2 Start with Syntax, but Trust the Operand Limits

Beginners often read `LDI Rd, K` as “load any register with any 8-bit constant”. The syntax alone does not tell the whole story. The **Operands** field does:

```
LDI Rd, K
Operands: 16 <= d <= 31, 0 <= K <= 255
```

That means `ldi r0, 1` is impossible, not merely discouraged. The reason is in the opcode encoding: `LDI` only has four bits for the register number, so the encoded register is really `Rd - 16`.

This pattern appears all over AVR:

- `LDI`, `ANDI`, `ORI`, `SUBI`, `SBCI`, `CPI` only work on `r16–r31`.
- `ADIW` and `SBIW` only work on specific register pairs.
- `MOVW` only uses even-numbered register pairs.
- `SBI`, `CBI`, `SBIC`, `SBIS` only work on low I/O addresses.

If the syntax looks general but the operand range is narrow, believe the range. The encoding is the real reason.

23.1.3 Read “Words” Before “Cycles”

AVR documentation uses **words**, not bytes, as the natural unit of program storage. One word is 16 bits, so:

- **1 word** = 2 bytes
- **2 words** = 4 bytes

This matters because AVR instructions are either one word or two words. That affects all of these at once:

- code size in flash
- instruction fetch behaviour
- skip-instruction timing
- absolute vs relative addressing choices

Examples:

- `OUT` is one word; `STS` is two words.
- `RJMP` is one word; `JMP` is two words.
- `RCALL` is one word; `CALL` is two words.

So when the manual says an instruction is “Words: 2”, translate that immediately into “this costs 4 bytes and may also affect skip timing”.

23.1.4 Relative Addresses Are Not Absolute Addresses

AVR manuals use several similar-looking operand letters:

- `K` often means an immediate constant.
- `k` often means a relative code offset.
- `q` means a small displacement from `Y` or `Z`.
- `A` means a low I/O address.

Case matters. `K` and `k` are not interchangeable.

For branches and relative calls:

- RJMP k means jump relative to the current PC
- RCALL k means call relative to the current PC
- BRNE k means branch a short relative distance if $Z = 0$

The manual usually writes this as:

```
PC <- PC + k + 1
```

That + 1 exists because the PC normally advances to the next instruction word before the relative offset is applied. The important practical result is:

- RJMP and RCALL do not encode a full absolute address.
- BRxx instructions have much shorter range than RJMP.
- Disassemblers may show targets as labels, but the hardware stores offsets.

23.1.5 Read the SREG Section as “What Can I Test Next?”

The **Status Register (SREG)** is not just bookkeeping. It determines what branches are meaningful after an instruction.

The bits are:

- I: global interrupt enable
- T: transfer bit for BST / BLD
- H: half-carry, mainly for nibble/BCD-style carries
- S: sign bit, defined as $N \text{ xor } V$
- V: signed overflow
- N: negative
- Z: zero
- C: carry / borrow

When reading an instruction page, ask:

- Does this instruction update flags at all?
- Which flags does it update?
- Are those flags enough for unsigned comparison, signed comparison, or both?

Examples:

- ADD updates Z, C, N, V, S, H, so it supports later arithmetic decisions.
- INC does **not** update C, so it is not a drop-in replacement for ADD Rd, 1 if carry matters.
- TST is really a flag-producing logical test, useful before BREQ or BRMI.
- CP changes flags without changing registers, which is why branches usually follow it directly.

The fastest way to understand a flag table is to map it to branch usage:

- For **unsigned** comparisons, care mostly about C and Z.
- For **signed** comparisons, care mostly about S, V, N, and Z.
- For equality, care about Z.

That is why AVR has both unsigned-style branches (BRLO, BRSH) and signed-style branches (BRLT, BRGE).

23.1.6 Do Not Panic About the Boolean Formulas

The manual often gives flag equations in Boolean form. Those formulas are precise, but you usually do not need to derive them by hand while writing ordinary code.

Use them in three situations:

1. When a flag result surprises you.
2. When writing multi-byte arithmetic where carry/borrow chains matter.
3. When reviewing whether two instructions are truly interchangeable.

For most programming, this simpler interpretation is enough:

- C: carry out of bit 7 for addition, borrow behaviour for subtraction
- Z: result was zero
- N: result bit 7 is 1
- V: signed overflow happened
- S: signed comparison helper, $N \text{ xor } V$
- H: carry from bit 3 to bit 4

If you are not doing decimal-style arithmetic or bit-exact arithmetic tricks, H is usually less important than Z, C, and S.

23.1.7 Cycles Are Family-Specific

The instruction set manual often shows cycle counts across multiple AVR core families. This book targets the ATtiny3217, which is an **AVRxt** device, so the AVRxt column is the one that matters most here.

Do not assume:

- a timing number from an older ATmega necessarily matches AVRxt
- all loads and stores cost the same on every AVR family
- every manual example was written with your exact core in mind

When timing matters, read the correct family column first, then read any note attached to it. Notes are where the manual hides the details that break naive assumptions.

23.1.8 Skip Instructions Need Extra Attention

CPSE, SBRC, SBRS, SBIC, and SBIS do not branch. They **skip the next instruction** if the condition is true.

That sounds simple, but the manual’s word count matters here:

- If the next instruction is 1 word, the skip jumps over 2 bytes.
- If the next instruction is 2 words, the skip jumps over 4 bytes.
- That is why skip instructions have multiple cycle counts.

So this is safe:

```
sbrc r16, 0
rjmp bit_was_set
```

But this behaves differently from a true branch because the skip is based on the size of the *next* instruction, not on a stored relative destination.

When reading the manual, any cycle count like 1/2/3 or 2/3/4 should make you ask: “is this a skip instruction, and is the next instruction one or two words long?”

23.1.9 Learn the Address Spaces Separately

AVR documentation is much easier once you stop treating “memory” as one thing. The manual switches among several address spaces:

- **register file:** r0–r31
- **I/O space:** low peripheral addresses used by IN / OUT
- **data space:** SRAM and memory-mapped locations used by LD / ST / LDS / STS
- **program memory:** flash, read with LPM

Many confusing instruction choices become obvious once you ask which space is being accessed.

Example:

- OUT PORTA_DIR, r16 only works if that register is in low I/O space.
- STS PORTA_DIR, r16 works with a data-space address, but costs more.

- LPM is for flash, not SRAM.

If two instructions seem to do the same job, check whether they are really talking to different address spaces.

23.1.10 Aliases and Pseudo-Instructions Can Hide the Real Opcode

Assembly source often uses friendly names that are not separate hardware instructions.

Examples:

- CLR Rd is assembled as EOR Rd, Rd.
- SER Rd is assembled as LDI Rd, 0xFF.
- LSL Rd is effectively ADD Rd, Rd.
- named branches like BREQ and BRNE are readable aliases for specific BRBS / BRBC conditions.

That matters because the manual may document one spelling while the assembler accepts another. If something looks inconsistent, check whether you are dealing with:

- a real opcode
- an alias mnemonic
- a pseudo-instruction expanded by the assembler

23.1.11 A Good Manual-Reading Workflow

When you meet an unfamiliar instruction, use this checklist:

1. Read the syntax and operand limits.
2. Identify the address space involved: register, I/O, SRAM, or flash.
3. Check whether it is 1 word or 2 words.
4. Check which SREG flags it updates.
5. Check whether the cycle count depends on branch taken/not taken, skip size, or AVR core family.
6. Check whether there is a shorter equivalent instruction for your case.

That method is enough to decode most of the manual quickly.

23.1.12 What This Means for the Rest of the Book

Use Chapter 4 to learn the common instruction patterns, then come back to this appendix when you need the full table. Use the Microchip manual when a question depends on one of these details:

- exact flag behaviour
- exact operand limits
- instruction encoding constraints
- family-specific cycle counts
- special notes about program memory, skip timing, or core differences

Once you start reading the manual this way, it stops looking like a wall of tables and starts looking like a set of answers to very specific questions.

23.2 Arithmetic and Logic

Opcode	Operation	Cycles	Flags
ADD Rd, Rr	$Rd \leftarrow Rd + Rr$	1	H S V N Z C
ADC Rd, Rr	$Rd \leftarrow Rd + Rr + C$	1	H S V N Z C

ADIW Rd, K	$Rd+1:Rd \leftarrow Rd+1:Rd + K$ (0–63)	2	S V N Z C
SUB Rd, Rr	$Rd \leftarrow Rd - Rr$	1	H S V N Z C
SUBI Rd, K	$Rd \leftarrow Rd - K$	1	H S V N Z C
SBC Rd, Rr	$Rd \leftarrow Rd - Rr - C$	1	H S V N Z C
SBCI Rd, K	$Rd \leftarrow Rd - K - C$	1	H S V N Z C
SBIW Rd, K	$Rd+1:Rd \leftarrow Rd+1:Rd - K$ (0–63)	2	S V N Z C
AND Rd, Rr	$Rd \leftarrow Rd \text{ AND } Rr$	1	S V N Z
ANDI Rd, K	$Rd \leftarrow Rd \text{ AND } K$ (R16–R31)	1	S V N Z
OR Rd, Rr	$Rd \leftarrow Rd \text{ OR } Rr$	1	S V N Z
ORI Rd, K	$Rd \leftarrow Rd \text{ OR } K$ (R16–R31)	1	S V N Z
EOR Rd, Rr	$Rd \leftarrow Rd \text{ XOR } Rr$	1	S V N Z
COM Rd	$Rd \leftarrow 0xFF - Rd$ (bitwise NOT)	1	S V N Z C
NEG Rd	$Rd \leftarrow 0x00 - Rd$ (two's comp)	1	H S V N Z C
INC Rd	$Rd \leftarrow Rd + 1$	1	S V N Z
DEC Rd	$Rd \leftarrow Rd - 1$	1	S V N Z
MUL Rd, Rr	$R1:R0 \leftarrow Rd \times Rr$ (unsigned)	2	Z C
MULS Rd, Rr	$R1:R0 \leftarrow Rd \times Rr$ (signed)	2	Z C
	Rd, Rr: R16–R31 only		
MULSU Rd, Rr	$R1:R0 \leftarrow Rd \times Rr$ (sgn $\times$ uns)	2	Z C
	Rd, Rr: R16–R23 only		
FMUL Rd, Rr	$R1:R0 \leftarrow (Rd \times Rr) \ll 1$ (uns)	2	Z C
FMULS Rd, Rr	$R1:R0 \leftarrow (Rd \times Rr) \ll 1$ (sgn)	2	Z C
FMULSU Rd, Rr	$R1:R0 \leftarrow (Rd \times Rr) \ll 1$ (s $\times$ u)	2	Z C

23.3 Shift and Rotate

Opcode	Operation	Cycles	Flags
LSL Rd	$Rd \leftarrow Rd \ll 1$; $Rd0=0$; $C \leftarrow Rd7$	1	H S V N Z C
LSR Rd	$Rd \leftarrow Rd \gg 1$; $Rd7=0$; $C \leftarrow Rd0$	1	S V N Z C
ASR Rd	$Rd \leftarrow Rd \gg 1$; $Rd7$ unchanged	1	S V N Z C
ROL Rd	$Rd \leftarrow Rd \ll 1$; $Rd0=C$; $C \leftarrow Rd7$	1	H S V N Z C
ROR Rd	$Rd \leftarrow Rd \gg 1$; $Rd7=C$; $C \leftarrow Rd0$	1	S V N Z C
SWAP Rd	$Rd[7:4] \leftrightarrow Rd[3:0]$	1	—

23.4 Bit Operations

Opcode	Operation	Cycles	Flags
SBI A, b	$I/O[A][b] \leftarrow 1$ (A: 0–31)	2	—
CBI A, b	$I/O[A][b] \leftarrow 0$ (A: 0–31)	2	—
BST Rd, b	$T \leftarrow Rd[b]$	1	T
BLD Rd, b	$Rd[b] \leftarrow T$	1	—
SEC	$C \leftarrow 1$	1	C
CLC	$C \leftarrow 0$	1	C
SEN	$N \leftarrow 1$	1	N
CLN	$N \leftarrow 0$	1	N
SEZ	$Z \leftarrow 1$	1	Z
CLZ	$Z \leftarrow 0$	1	Z
SEI	$I \leftarrow 1$ (global interrupt on)	1	I
CLI	$I \leftarrow 0$ (global interrupt off)	1	I
SES	$S \leftarrow 1$	1	S
CLS	$S \leftarrow 0$	1	S

SEV	$V \leftarrow 1$	1	V
CLV	$V \leftarrow 0$	1	V
SET	$T \leftarrow 1$	1	T
CLT	$T \leftarrow 0$	1	T
SEH	$H \leftarrow 1$	1	H
CLH	$H \leftarrow 0$	1	H
BSET b	$SREG[b] \leftarrow 1$	1	(b)
BCLR b	$SREG[b] \leftarrow 0$	1	(b)

23.5 Compare and Test

Opcode	Operation	Cycles	Flags
CP Rd, Rr	$Rd - Rr$ (flags only)	1	H S V N Z C
CPC Rd, Rr	$Rd - Rr - C$ (flags only)	1	H S V N Z C
CPI Rd, K	$Rd - K$ (flags only, R16–R31)	1	H S V N Z C
TST Rd	Rd AND Rd (flags only)	1	S V N Z

23.6 Branch Instructions

Opcode	Condition	Cycles	Notes
RJMP k	—	2	$PC \leftarrow PC + k + 1$; k: -2048..2047
JMP k	—	3	$PC \leftarrow k$; 32-bit instruction
IJMP	—	2	$PC \leftarrow Z$ (16-bit)
EIJMP	—	2	$PC \leftarrow EIND:Z$ (>128KB flash)
RCALL k	—	3	Push PC, $PC \leftarrow PC + k + 1$
CALL k	—	4	Push PC, $PC \leftarrow k$; 32-bit
ICALL	—	3	Push PC, $PC \leftarrow Z$
EICALL	—	3	Push PC, $PC \leftarrow EIND:Z$
RET	—	4	Pop PC
RETI	—	4	Pop PC, $I \leftarrow 1$
BRBS b, k	$SREG[b] = 1$	1/2	Branch if bit set; 1=not taken, 2=taken
BRBC b, k	$SREG[b] = 0$	1/2	Branch if bit clear
/* Named branch mnemonics (aliases for BRBS/BRBC) */			
BREQ k	$Z = 1$	1/2	Branch if equal
BRNE k	$Z = 0$	1/2	Branch if not equal
BRCS k	$C = 1$	1/2	Branch if carry set
BRCC k	$C = 0$	1/2	Branch if carry clear
BRSH k	$C = 0$	1/2	Branch if same or higher (unsigned $\geq$)
BRL0 k	$C = 1$	1/2	Branch if lower (unsigned $<$)
BRMI k	$N = 1$	1/2	Branch if minus
BRPL k	$N = 0$	1/2	Branch if plus
BRGE k	$S = 0$	1/2	Branch if greater or equal (signed $\geq$)
BRLT k	$S = 1$	1/2	Branch if less than (signed $<$)
BRHS k	$H = 1$	1/2	Branch if half-carry set
BRHC k	$H = 0$	1/2	Branch if half-carry clear
BRTS k	$T = 1$	1/2	Branch if T set
BRTC k	$T = 0$	1/2	Branch if T clear
BRVS k	$V = 1$	1/2	Branch if overflow set

BRVC k	V = 0	1/2	Branch if overflow clear
BRIE k	I = 1	1/2	Branch if interrupt enabled
BRID k	I = 0	1/2	Branch if interrupt disabled

23.6.1 Skip Instructions

Opcode	Condition	Cycles	Notes
CPSE Rd, Rr	Rd = Rr → skip next	1/2/3	skip 1 or 2-word instruction
SBRC Rd, b	Rd[b] = 0 → skip	1/2/3	
SBRs Rd, b	Rd[b] = 1 → skip	1/2/3	
SBIC A, b	I/O[A][b] = 0 → skip	2/3/4	A: 0–31 only
SBIS A, b	I/O[A][b] = 1 → skip	2/3/4	A: 0–31 only

Skip cycle counts: 1 = no skip; 2 = skip 1-word instruction; 3 = skip 2-word instruction (LDS, STS, CALL, JMP). SBIC/SBIS add 1 due to I/O read.

23.7 Load and Store

Opcode	Operation	Cycles	Notes
MOV Rd, Rr	Rd ← Rr	1	
MOVW Rd, Rr	Rd+1:Rd ← Rr+1:Rr	1	Rd, Rr: even registers
LDI Rd, K	Rd ← K (R16–R31 only)	1	
LDS Rd, k	Rd ← SRAM[k]	2	k: 16-bit address
STS k, Rr	SRAM[k] ← Rr	2	k: 16-bit address
/* X register (R27:R26) indirect */			
LD Rd, X	Rd ← SRAM[X]	2	
LD Rd, X+	Rd ← SRAM[X]; X ← X+1	2	
LD Rd, -X	X ← X-1; Rd ← SRAM[X]	2	
ST X, Rr	SRAM[X] ← Rr	2	
ST X+, Rr	SRAM[X] ← Rr; X ← X+1	2	
ST -X, Rr	X ← X-1; SRAM[X] ← Rr	2	
/* Y register (R29:R28) indirect */			
LD Rd, Y	Rd ← SRAM[Y]	2	
LD Rd, Y+	Rd ← SRAM[Y]; Y ← Y+1	2	
LD Rd, -Y	Y ← Y-1; Rd ← SRAM[Y]	2	
LDD Rd, Y+q	Rd ← SRAM[Y+q] q: 0–63	2	
ST Y, Rr	SRAM[Y] ← Rr	2	
ST Y+, Rr	SRAM[Y] ← Rr; Y ← Y+1	2	
ST -Y, Rr	Y ← Y-1; SRAM[Y] ← Rr	2	
STD Y+q, Rr	SRAM[Y+q] ← Rr q: 0–63	2	
/* Z register (R31:R30) indirect */			
LD Rd, Z	Rd ← SRAM[Z]	2	
LD Rd, Z+	Rd ← SRAM[Z]; Z ← Z+1	2	
LD Rd, -Z	Z ← Z-1; Rd ← SRAM[Z]	2	
LDD Rd, Z+q	Rd ← SRAM[Z+q] q: 0–63	2	
ST Z, Rr	SRAM[Z] ← Rr	2	
ST Z+, Rr	SRAM[Z] ← Rr; Z ← Z+1	2	
ST -Z, Rr	Z ← Z-1; SRAM[Z] ← Rr	2	
STD Z+q, Rr	SRAM[Z+q] ← Rr q: 0–63	2	

23.7.1 Load from Program Memory (Flash)

Opcode	Operation	Cycles	Notes
LPM	$R0 \leftarrow \text{FLASH}[Z]$	3	Z = byte address
LPM Rd, Z	$Rd \leftarrow \text{FLASH}[Z]$	3	
LPM Rd, Z+	$Rd \leftarrow \text{FLASH}[Z]; Z \leftarrow Z+1$	3	
ELPM	$R0 \leftarrow \text{FLASH}[\text{RAMPZ}:Z]$	3	>64KB flash devices
ELPM Rd, Z	$Rd \leftarrow \text{FLASH}[\text{RAMPZ}:Z]$	3	
ELPM Rd, Z+	$Rd \leftarrow \text{FLASH}[\text{RAMPZ}:Z]; Z++$	3	
SPM	$\text{FLASH}[Z] \leftarrow R1:R0$	varies	Boot section only

23.8 I/O Instructions

Opcode	Operation	Cycles	Notes
IN Rd, A	$Rd \leftarrow I/O[A]$	1	A: 0–63
OUT A, Rr	$I/O[A] \leftarrow Rr$	1	A: 0–63
PUSH Rr	$\text{SRAM}[SP] \leftarrow Rr; SP \leftarrow SP-1$	2	
POP Rd	$SP \leftarrow SP+1; Rd \leftarrow \text{SRAM}[SP]$	2	

23.9 Miscellaneous

Opcode	Operation	Cycles	Notes
NOP	No operation	1	
SLEEP	Enter sleep mode	1	MCU-defined behaviour
WDR	Reset watchdog timer	1	
BREAK	Stop for on-chip debug system	1	

23.10 Instruction Encoding Summary

Most AVR instructions are 16-bit (one word). 32-bit instructions:

32-bit instructions (2 words):

- LDS Rd, k – word 1: opcode+Rd; word 2: 16-bit address k
- STS k, Rr – word 1: opcode+Rr; word 2: 16-bit address k
- CALL k – word 1: $0x940E+kH$; word 2: kL (22-bit address)
- JMP k – word 1: $0x940C+kH$; word 2: kL (22-bit address)

All other instructions are 16-bit. The assembler handles encoding automatically.

23.11 GCC ABI Register Conventions

Registers	Role	Save by
R0	Temporary, MUL result	Caller (assume destroyed)
R1	Always zero (GCC)	Callee must restore if modified
R2–R17	Call-saved	Callee (push/pop if used)

R18–R27	Call-clobbered	Caller (assume destroyed)
R28–R29 (Y)	Frame pointer	Callee
R30–R31 (Z)	Call-clobbered	Caller

Return values:

8-bit: R24

16-bit: R25:R24

32-bit: R25:R24:R23:R22 (little-endian: R22=lowest byte)

Arguments (first to last):

arg1: R25:R24 (or R24 for 8-bit)

arg2: R23:R22

arg3: R21:R20

arg4: R19:R18

More: pushed on stack (rightmost first)

After MUL/MULS/MULSU: always restore R1 = 0 before returning or calling any C function. R0 is implicitly scratch.

Chapter 24

GAS Directive Reference

GNU Assembler (avr-as) directives used in AVR assembly. Directives begin with `.` and are not instructions — they control how the assembler lays out output.

24.1 Section Directives

Directive	Effect
<code>.section name [,"flags"[@type]]</code>	Switch to named section. Common flags: a = allocatable w = writable x = executable Types: @progbits (data), @nobits (BSS, zero-init)
<code>.text</code>	Switch to .text (code, flash)
<code>.data</code>	Switch to .data (initialised SRAM variables)
<code>.bss</code>	Switch to .bss (zero-initialised SRAM)
<code>.rodata</code>	Read-only data (constants in flash on AVR)

Common section usage on AVR:

```
.section .text          /* executable code (flash) */
.section .data          /* initialised variables (SRAM; copied from flash at startup) */
.section .bss           /* zero-initialised variables (SRAM) */
.section .rodata        /* constants (flash; access via LPM) */
.section .vectors,"ax",@progbits /* interrupt vector table */
.section .noinit,"aw",@nobits  /* SRAM variables NOT initialised at reset */
```

24.2 Symbol Directives

Directive	Effect
<code>.global symbol</code>	Export symbol (visible to linker / other files)
<code>.local symbol</code>	Mark symbol local (file scope only)
<code>.weak symbol</code>	Weak symbol: overridden if another file defines it
<code>.type symbol, @function</code>	Mark symbol as function (for debugging/ELF info)
<code>.type symbol, @object</code>	Mark symbol as data object
<code>.size symbol, size</code>	Set symbol size (bytes)

```
.set    symbol, expr      Define symbol = expr (reassignable)
.equ    symbol, expr      Define symbol = expr (not reassignable; synonym: .equiv)
```

Example:

```
.global uart_init
.type   uart_init, @function
uart_init:
    /* ... */
    ret
.size   uart_init, . - uart_init /* size = current address minus symbol start */
```

24.3 Data Directives

Directive	Size per item	Example
.byte expr[,...]	1 byte	.byte 0x3F, 42, 'A'
.2byte expr[,...]	2 bytes LE	.2byte 0x1234 /* = 0x34, 0x12 in memory */
.word expr[,...]	2 bytes LE	.word 1000 /* AVR word = 2 bytes */
.long expr[,...]	4 bytes LE	.long 0xDEADBEEF
.quad expr[,...]	8 bytes LE	.quad 0
.ascii "str"	n bytes	.ascii "hello" /* no null terminator */
.asciz "str"	n+1 bytes	.asciz "hello" /* null terminated */
.string "str"	n+1 bytes	.string "hello" /* alias for .asciz */
.space n [,fill]	n bytes	.space 16, 0xFF /* fill with 0xFF */
.zero n	n zero bytes	.zero 64 /* alias for .space n, 0 */
.skip n [,fill]	n bytes	.skip 4 /* alias for .space */
.fill reps,size,val		.fill 8, 1, 0xAA /* 8 bytes of 0xAA */

24.4 Alignment Directives

Directive	Effect
.align n	Align to 2 ⁿ byte boundary (with NOP padding in .text)
.balign n [,fill]	Align to n byte boundary
.p2align n	Align to 2 ⁿ boundary (same as .align on most targets)
.org expr	Set location counter to expr (absolute address)

On AVR, .align 1 aligns to 2-byte boundary (one AVR word). .align 2 aligns to 4 bytes. Useful before jump tables and 16-bit data.

24.5 Conditional Assembly

Directive	Effect
.ifdef sym	Include if sym is defined
.ifndef sym	Include if sym is NOT defined
.if expr	Include if expr != 0
.else	Else branch
.elif expr	Else-if
.endif	End conditional block
.error "msg"	Abort assembly with message

`.warning "msg"` Emit warning, continue

Example — select multiply implementation at assemble time:

```
#ifdef NOMUL
    rcall mul8u_shiftadd
#else
    mul    r16, r17
    movw   r24, r0
    clr    r1
#endif
```

When using `avr-gcc` as the front-end, `#ifdef/#define` (C preprocessor) work in `.S` files. For plain `avr-as`, use `.ifdef/.equ`.

24.6 Macro Directives

Directive	Effect
<code>.macro name [params]</code>	Begin macro definition
<code>.endm</code>	End macro definition
<code>.exitm</code>	Exit macro early
<code>.rept n</code>	Repeat enclosed block n times
<code>.endr</code>	End <code>.rept</code> block
<code>.irp sym, values</code>	Iterate: sym takes each value in turn
<code>.endr</code>	
<code>.irpc sym, chars</code>	Iterate over characters of a string
<code>.endr</code>	

Example — macro for a 16-bit load:

```
.macro ldw reg_h, reg_l, value
    ldi    \reg_l, lo8(\value)
    ldi    \reg_h, hi8(\value)
.endm

/* Usage */
ldw    r25, r24, 1999    /* expands to: ldi r24, lo8(1999); ldi r25, hi8(1999) */
```

Example — `.rept` for vector table padding:

```
__vectors:
    rjmp   reset_handler
    .rept  25
    rjmp   dummy_isr      /* fill remaining 25 vectors */
    .endr
```

24.7 Include Directives

Directive	Effect
<code>.include "file"</code>	Textually include file at this point
<code>#include <file></code>	C preprocessor include (only when assembled via <code>avr-gcc</code>)

For AVR with `avr-gcc -x assembler-with-cpp` (the default for `.S` files):


```
#include <avr/io.h>      /* defines I/O register names, bit names, RAMEND, etc. */
```

<avr/io.h> selects the correct device header based on -mmcu= flag.

24.8 Listing and Debug Directives

Directive	Effect
.file "name"	Set source file name (for debug info)
.line n	Set source line number
.loc file line [col]	DWARF location (used by avr-gcc; rarely hand-written)
.stabs "str",n,n,n,n	STABS debug info (legacy)
.ident "string"	Embed comment string in .comment ELF section

24.9 Expression Operators

GAS expressions can use these operators in directives and immediate operands:

Operator	Example	Meaning
+	.byte 1+2	Addition
-	.byte 5-3	Subtraction
*	.word 16*1000	Multiplication
/	.byte 255/2	Integer division
%	.byte 17%10	Modulo
<<	.word 1<<8	Left shift
>>	.word 256>>1	Right shift
&	.byte 0xFF & val	Bitwise AND
	.byte flags mask	Bitwise OR
^	.byte a ^ b	Bitwise XOR
~	.byte ~mask	Bitwise NOT
lo8(x)	ldi r16, lo8(val)	Low byte of 16-bit value
hi8(x)	ldi r16, hi8(val)	High byte of 16-bit value
hlo8(x)		Byte 2 (bits 23:16) of 32-bit value
hhi8(x)		Byte 3 (bits 31:24) of 32-bit value
pm(x)	.word pm(label)	Program memory word address (byte addr >> 1)
.	(dot)	Current location counter (address)

24.10 Common Patterns

24.10.1 Define a constant

```
.equ F_CPU,    16000000
.equ BAUD,     9600
.equ UBRR_V,   (F_CPU / (16 * BAUD)) - 1    /* = 103 */
```

24.10.2 Reserve SRAM variable

```
.section .bss
my_var: .byte 0      /* 1-byte variable */
my_word: .2byte 0    /* 2-byte variable, aligned */
my_buf: .space 64    /* 64-byte buffer */
```

24.10.3 Initialised variable (copied to SRAM at startup by `avr-libc` `_init`)

```
.section .data
init_val: .byte 42
```

24.10.4 Constant in flash (accessed via LPM)

```
.section .rodata
my_table:
.byte 0, 1, 1, 2, 3, 5, 8, 13 /* Fibonacci */
```

24.10.5 Local labels

```
func:
    ldi    r16, 10
1:      /* local label – referenced as 1f (forward) or 1b (back) */
    dec    r16
    brne   1b /* branch back to label 1 */
    ret
```

Local labels (numeric) can be reused across the file; 1f means “next 1: going forward”, 1b means “previous 1: going backward”. They do not pollute the symbol table.

Chapter 25

Linker Script Walkthrough

The linker script tells `avr-ld` where to place each section in the output ELF file and ultimately in the device's address space. For most projects `avr-gcc` selects a built-in script automatically via `-mmcu=`. Understanding linker scripts is necessary when:

- Writing a bootloader (section placement at a specific flash address).
 - Placing variables in a specific SRAM region.
 - Using `.noinit` for variables that survive a software reset.
 - Adding custom sections (e.g., a parameter page at a known flash address).
-

25.1 Anatomy of a Linker Script

A linker script has three main constructs:

`MEMORY { }` – declares memory regions (name, origin, length)
`SECTIONS { }` – maps input sections to output sections and memory regions
`PROVIDE() / ENTRY()` – define special symbols and the entry point

25.1.1 Minimal AVR Linker Script

```
/* minimal_avr.ld – bare-bones script for ATmega328P */

MEMORY
{
    flash    (rx)  : ORIGIN = 0x000000, LENGTH = 32K
    sram     (rwx) : ORIGIN = 0x800100, LENGTH = 2K
}

SECTIONS
{
    /* Interrupt vector table – must be at flash origin */
    .vectors :
    {
        KEEP(*(.vectors))
    } > flash

    /* Executable code */
    .text :
    {
        *(.text)
    }
}
```

```

    *(.text.*)
} > flash

/* Read-only data (constants accessed via LPM) */
.rodata :
{
    *(.rodata)
    *(.rodata.*)
} > flash

/* Initialised data: stored in flash, copied to SRAM at startup */
.data :
{
    _data_start = .;
    *(.data)
    *(.data.*)
    _data_end = .;
} > sram AT > flash          /* VMA in sram, LMA in flash */

/* Zero-initialised data: only reservation in SRAM */
.bss :
{
    _bss_start = .;
    *(.bss)
    *(.bss.*)
    *(COMMON)
    _bss_end = .;
} > sram

/* Variables preserved across software reset (not zeroed at startup) */
.noinit (NOLOAD) :
{
    *(.noinit)
} > sram

/* Stack: grows downward from RAMEND */
_stack_top = ORIGIN(sram) + LENGTH(sram) - 1;
}

```

25.2 MEMORY Region Flags

r – readable
 w – writable
 x – executable
 ! – negate following flags

Common combinations:

(rx) – read-execute (flash)
 (rw!x) – read-write, not executable (SRAM)
 (rx!w) – read-execute, not writable (flash, explicit)

25.3 Address Spaces on AVR

AVR uses a **Harvard architecture** with separate address spaces:

Flash (program memory):

avr-ld uses byte addresses with origin 0x000000

ATmega328P: 0x000000–0x007FFF (32 KB)

SRAM (data memory):

avr-ld uses byte addresses with origin 0x800000 + hardware_offset

ATmega328P hardware SRAM starts at 0x0100

avr-ld SRAM origin: 0x800100

Registers: 0x0000–0x001F (R0–R31) – not in linker script

I/O: 0x0020–0x005F – not in linker script

Ext I/O: 0x0060–0x00FF – not in linker script

SRAM: 0x0100–0x08FF (2 KB on ATmega328P)

The 0x800000 offset is an avr-ld convention to distinguish data addresses from code addresses in the ELF file. The hardware only sees the lower 16 bits.

25.4 VMA vs LMA

VMA (Virtual Memory Address): where the section lives at runtime (SRAM for .data).

LMA (Load Memory Address): where the section is stored in the ELF output (flash for .data, since flash is what gets programmed).

```
.data :
{
  _data_start = .;          /* VMA: SRAM address where data lives at runtime */
  *(.data)
  _data_end = .;
} > sram AT > flash        /* LMA: in flash, immediately after .rodata */
```

The startup code (__init in avr-libc) copies .data from its LMA (flash) to its VMA (SRAM) before main() runs:

```
/* avr-libc __init pseudo-code */
ldi ZL, lo8(_data_lma_start) /* source: flash LMA */
ldi ZH, hi8(_data_lma_start)
ldi YL, lo8(_data_start)    /* dest: SRAM VMA */
ldi YH, hi8(_data_start)
1: lpm r0, Z+
st Y+, r0
cpi ZL, lo8(_data_lma_end)
brne 1b
```

If you use -nostartfiles, **you must perform this copy yourself** if any .data variables are used. The simplest approach: avoid initialised variables entirely, or use .equ constants in flash and load them explicitly.

25.5 KEEP()

The garbage-collector (--gc-sections) removes unreferenced sections. KEEP() prevents this:

```
KEEP(*(.vectors))    /* never discard the vector table */
```

When building with `avr-gcc -Wl,--gc-sections` (common for size optimisation), sections not reachable from the entry point are removed. `KEEP()` forces retention regardless of reachability.

25.6 Bootloader Placement

Place the bootloader at the boot section start address by overriding `.text` origin:

```
avr-gcc -Wl,--section-start=.text=0x7C00 ...
```

Or in a dedicated linker script:

```
/* bootloader.ld */
MEMORY
{
    boot_flash (rx) : ORIGIN = 0x007C00, LENGTH = 1K
    sram        (rw!x): ORIGIN = 0x800100, LENGTH = 2K
}

SECTIONS
{
    .text : { *(.vectors) *(.text) *(.text.*) } > boot_flash
    .bss  : { *(.bss) } > sram
}
```

```
avr-gcc -mmcu=atmega328p -nostartfiles \
-T bootloader.ld \
-o bootloader.elf bootloader.S
```

25.7 Parameter Page at Fixed Flash Address

Store calibration data at a known, fixed flash address so the application can find it regardless of code size changes:

```
SECTIONS
{
    /* ... normal sections ... */

    .params :
    {
        *(.params)
    } > flash
    . = 0x7000; /* force next section to start at 0x7000 */
    .calibration :
    {
        KEEP(*(.calibration))
    } > flash
}
```

In source:

```
.section .calibration
cal_offset: .word 0x0032 /* factory calibration: ADC offset */
cal_gain:   .word 0x0100 /* Q8.8 gain factor */
```

Access at runtime:

```
ldi    ZL, lo8(cal_offset)
ldi    ZH, hi8(cal_offset)
lpm    r24, Z+
lpm    r25, Z      /* r25:r24 = cal_offset value */
```

25.8 .noinit Section

Variables in `.noinit` are **not zeroed** at startup — they retain whatever value was in SRAM before reset. Useful for reset reason tracking:

```
.section .noinit,"aw",@nobits
reset_reason: .byte 0      /* written before software reset; survives it */
```

Application writes a code before resetting:

```
ldi    r16, 0xAA
sts    reset_reason, r16
/* ... trigger watchdog reset ... */
```

After reset, check `reset_reason` before the BSS clear to determine why the device restarted. BSS clear (zeroing `.bss`) happens in `__init` — `.noinit` is placed in a separate section that `__init` never touches.

25.9 Useful Linker Symbols

These symbols are defined by the default `avr-gcc` linker script and available in assembly via `lo8()/hi8()`:

Symbol	Value
<code>__data_start</code>	VMA start of <code>.data</code> in SRAM
<code>__data_end</code>	VMA end of <code>.data</code> in SRAM
<code>__data_load_start</code>	LMA of <code>.data</code> in flash (copy source for <code>__init</code>)
<code>__bss_start</code>	Start of <code>.bss</code> in SRAM
<code>__bss_end</code>	End of <code>.bss</code> in SRAM
<code>__heap_start</code>	Start of heap (after <code>.bss</code>)
<code>__stack</code>	Initial stack pointer value (= <code>RAMEND</code>)
<code>__TEXT_REGION_ORIGIN__</code>	Start of <code>.text</code> (usually 0)
<code>__DATA_REGION_ORIGIN__</code>	Start of SRAM (usually <code>0x800100</code>)

Access in assembly:

```
ldi    r26, lo8(__bss_start)
ldi    r27, hi8(__bss_start)
ldi    r24, lo8(__bss_end - __bss_start) /* BSS length */
```

25.10 Inspecting Linker Output

```
# Show section layout (addresses, sizes)
avr-objdump -h firmware.elf

# Show all symbols with addresses
avr-nm --numeric-sort firmware.elf

# Show memory map (verbose linker output)
avr-gcc -Wl,-Map=firmware.map ... && cat firmware.map

# Show disassembly with source interleaved
avr-objdump -dS firmware.elf

# Verify flash usage
avr-size --mcu=atmega328p --format=avr firmware.elf
```

Sample avr-size output with AVR format:

AVR Memory Usage

Device: atmega328p

Program: 486 bytes (1.5% Full)
(.text + .data + .bootloader)

Data: 3 bytes (0.1% Full)
(.data + .bss + .noinit)

25.11 Common Linker Errors

Error	Cause / Fix
undefined reference to `symbol'	Symbol not defined in any linked file; add the .S or .o file that defines it.
relocation truncated to fit: R_AVR_13_PCREL against `symbol'	Branch offset exceeds range (e.g. RJMP only reaches $\pm 2\text{KB}$); use JMP or reorganise code.
section `.text' will not fit in region `flash'	Total code exceeds flash; reduce size or use a larger device.
cannot find entry symbol reset_handler	Define reset_handler or use -e reset_handler linker flag.
multiple definition of `__vector_N'	Two .S files both define the same ISR vector entry; remove the duplicate.

Chapter 26

avrdude Flashing Cheat-Sheet

avrdude is the standard tool for programming AVR microcontrollers. It communicates with the device via a programmer (USB, serial, SPI ISP) and can read/write flash, EEPROM, and fuse bytes.

26.1 Basic Syntax

```
avrdude -p PARTNO -c PROGRAMMER [-P PORT] [-b BAUD] -U MEMTYPE:OP:FILE[:FORMAT]
```

-p PARTNO Device (e.g. m328p, t3217, m2560). See Section E.4 for list.
-c PROGRAMMER Programmer type. See Section E.3.
-P PORT Port (e.g. /dev/ttyUSB0, /dev/ttyACM0, usb).
-b BAUD Baud rate (for serial programmers; default depends on programmer).
-U Memory operation (can be repeated for multiple operations).
-v Verbose output. -vv for more detail.
-n Dry run: do not write anything.
-e Erase flash before programming.
-F Override signature check (use with caution).

26.1.1 -U Flag Format

MEMTYPE:OP:FILE[:FORMAT]

MEMTYPE:

flash – program memory
eeprom – EEPROM
lfuse – low fuse byte
hfuse – high fuse byte
efuse – extended fuse byte
lock – lock bits
signature – device signature (read only)
calibration – oscillator calibration (read only)

OP:

r – read from device, write to FILE
w – write FILE to device
v – verify device matches FILE

FORMAT (optional, auto-detected if omitted):

i – Intel HEX

s – Motorola S-record
r – raw binary
e – ELF
h – hex string
b – byte value (for single-byte ops like fuses)
m – immediate value (e.g. 0xDE)

26.2 Most-Used Commands

26.2.1 Flash a HEX file

```
# Arduino Uno (ATmega328P, USB-serial bootloader on /dev/ttyUSB0)
avrdude -p m328p -c arduino -P /dev/ttyUSB0 -b 115200 \
        -U flash:w:firmware.hex:i

# USBasp programmer (no port needed)
avrdude -p m328p -c usbasp \
        -U flash:w:firmware.hex:i

# AVRISP mkII (USB)
avrdude -p m328p -c avrispmkII \
        -U flash:w:firmware.hex:i
```

26.2.2 Read flash to file

```
avrdude -p m328p -c usbasp \
        -U flash:r:backup.hex:i
```

26.2.3 Verify only (no write)

```
avrdude -p m328p -c usbasp \
        -U flash:v:firmware.hex:i
```

26.2.4 Erase then write

```
avrdude -p m328p -c usbasp -e \
        -U flash:w:firmware.hex:i
```

26.2.5 Write EEPROM

```
avrdude -p m328p -c usbasp \
        -U eeprom:w:data.hex:i
```

26.2.6 Read fuse bytes

```
avrdude -p m328p -c usbasp \
        -U lfuse:r:-:h -U hfuse:r:-:h -U efuse:r:-:h
# Output: e.g. 0xff 0xd9 0x07 (lfuse hfuse efuse)
# '-' = stdout, 'h' = hex string format
```

26.2.7 Write fuse bytes

```
# WARNING: incorrect fuse values can brick the device.
# Always calculate fuses with an online calculator and verify before writing.

# Set hfuse = 0xDE (BOOTRST=0, BOOTSZ=10 for 1KB bootloader section)
avrdude -p m328p -c usbasp \
        -U hfuse:w:0xDE:m

# Restore ATmega328P factory fuses (lfuse=0xFF, hfuse=0xD9, efuse=0x07)
avrdude -p m328p -c usbasp \
        -U lfuse:w:0xFF:m \
        -U hfuse:w:0xD9:m \
        -U efuse:w:0x07:m
```

Warning: Fuses control clock source, brown-out detection, bootloader region, and JTAG. Setting CKSEL bits for an external crystal when no crystal is connected will make the device unresponsive — recovery requires an ISP programmer that provides a clock signal (High-Voltage Serial Programming).

26.2.8 Multiple operations in one command

```
avrdude -p m328p -c usbasp \
        -U flash:w:firmware.hex:i \
        -U eeprom:w:eeprom_data.hex:i
```

26.3 Common Programmers

Programmer ID	Description
arduino	Arduino bootloader via serial (auto-reset on DTR)
arduino-ft232r	Arduino via FTDI FT232R cable
avrisp	Atmel AVRISP v1 (serial)
avrisp2	Atmel AVRISP mkII (USB HID); also: avrispmkII
usbasp	USBasp (cheap USB ISP clone, widely available)
usbtiny	USBTiny / Adafruit USBTinyISP
dragon_isp	AVR Dragon in ISP mode
dragon_jtag	AVR Dragon in JTAG mode
jtag2isp	JTAGICE mkII in ISP mode
jtag2	JTAGICE mkII
atmelice_isp	Atmel-ICE in ISP mode
atmelice_pdi	Atmel-ICE in PDI mode (XMEGA)
atmelice_updi	Atmel-ICE in UPDI mode (tinyAVR/megaAVR 0/1/2-series)
pymcuprog	Microchip MPLAB Snap / Curiosity Nano via pymcuprog bridge
serialupdi	UPDI via USB-serial adapter (cheap, 1-wire)

26.3.1 UPDI Devices (ATtiny3217 and tinyAVR 1-series)

The ATtiny3217 uses UPDI (Unified Program and Debug Interface), a single-wire protocol. Use serialupdi with a USB-serial adapter connected to the UPDI pin via a 4.7kΩ resistor:

```
# ATtiny3217 via serialupdi on /dev/ttyUSB0
avrdude -p t3217 -c serialupdi -P /dev/ttyUSB0 \
        -U flash:w:firmware.hex:i
```

Read fuses

```
avrdude -p t3217 -c serialupdi -P /dev/ttyUSB0 \
-U fuse0:r:-:h -U fuse1:r:-:h -U fuse2:r:-:h \
-U fuse4:r:-:h -U fuse5:r:-:h -U fuse6:r:-:h -U fuse7:r:-:h
```

UPDI wiring (1-wire):

```

USB-serial TX ┌───┐
                │   ┌── 4.7kΩ ─ UPDI pin (ATtiny3217 pin 16)
USB-serial RX ─┘──┘
USB-serial GND ─── GND

```

Alternatively, use the Atmel-ICE or MPLAB PICKit 4 for UPDI.

26.4 Device Part Numbers

Part Number	Device
m328p	ATmega328P (Arduino Uno, Nano)
m328pb	ATmega328PB
m2560	ATmega2560 (Arduino Mega)
m32u4	ATmega32U4 (Arduino Leonardo, Pro Micro)
m168p	ATmega168P
m88p	ATmega88P
m48p	ATmega48P
t3217	ATtiny3217 (book target for tinyAVR 1-series)
t1616	ATtiny1616
t816	ATtiny816
t416	ATtiny416
t85	ATtiny85
t84	ATtiny84
t45	ATtiny45
t13	ATtiny13
x128a4u	ATxmegal28A4U

Full list: `avrdude -p ?`

26.5 Fuse Byte Reference (ATmega328P)

26.5.1 Low Fuse (LFUSE)

Bit	Name	Description
7	CKDIV8	Divide clock by 8 (0 = enabled; factory default)
6	CKOUT	Clock output on PB0 (0 = enabled)
5:4	SUT1:0	Start-up time (10 = default 14CK + 65ms)
3:0	CKSEL3:0	Clock source (0010 = internal 8MHz RC; 1111 = external crystal)

Factory default (internal RC, /8): LFUSE = 0x62

No divide (internal 8MHz full speed): LFUSE = 0xE2

External 16MHz crystal (Arduino): LFUSE = 0xFF

26.5.2 High Fuse (HFUSE)

Bit	Name	Description
-----	------	-------------

7	RSTDISBL	Disable reset pin (0 = PC6 becomes I/O; DANGEROUS)
6	DWEN	debugWIRE enable (0 = enabled; disables ISP when set)
5	SPIEN	SPI programming enable (0 = enabled; keep this 0)
4	WDTON	Watchdog always on (0 = always on)
3	EESAVE	Preserve EEPROM during chip erase (0 = preserved)
2:1	BOOTSZ1:0	Boot section size (11=256B, 10=512B, 01=1KB, 00=2KB)
0	BOOTRST	Boot reset vector (0 = reset jumps to boot section)

Factory default: HFUSE = 0xD9

(SPI enabled, WDT off, EEPROM erased, 2KB boot, reset to 0x0000)

Bootloader (1KB boot, reset to boot section): HFUSE = 0xDE

26.5.3 Extended Fuse (EFUSE)

Bit	Name	Description
-----	------	-------------

2:0	BODLEVEL2:0	Brown-out detection level
		111 = BOD disabled (factory)
		110 = 1.8V
		101 = 2.7V
		100 = 4.3V

Factory default: EFUSE = 0x07 (BOD disabled; only bits 2:0 used)

26.5.4 Online Fuse Calculator

The Engbedded AVR Fuse Calculator provides a GUI for calculating fuse values. Search “AVR fuse calculator” — cross-check with the device datasheet before programming.

26.6 Generating Intel HEX from ELF

```
avr-objcopy -O ihex firmware.elf firmware.hex

# EEPROM section separately
avr-objcopy -O ihex -j .eeprom \
  --set-section-flags=.eeprom=alloc,load \
  --change-section-lma .eeprom=0 \
  firmware.elf eeprom.hex
```

26.7 Complete Build + Flash Script

```
#!/bin/bash
# build_flash.sh — assemble, link, and flash ATmega328P
set -e

MCU=atmega328p
PORT=/dev/ttyUSB0
BAUD=115200
SRC=firmware.S
```

```

avr-gcc -mmcu=${MCU} -nostartfiles -o firmware.elf ${SRC}
avr-size firmware.elf
avr-objcopy -O ihex firmware.elf firmware.hex
avrdude -p ${MCU} -c arduino -P ${PORT} -b ${BAUD} \
        -U flash:w:firmware.hex:i
echo "Done."

# For ATtiny3217 via serialupdi
MCU=attiny3217
PORT=/dev/ttyUSB0

avr-gcc -mmcu=${MCU} -nostartfiles -o firmware.elf firmware.S
avr-objcopy -O ihex firmware.elf firmware.hex
avrdude -p ${MCU} -c serialupdi -P ${PORT} \
        -U flash:w:firmware.hex:i

```

26.8 Troubleshooting

Problem	Likely cause / fix
avrdude: stk500_recv(): timeout	Wrong port, wrong baud, board not reset. Try -P /dev/ttyACM0, check USB cable.
avrdude: Expected signature for ATmega328P is 1E 95 0F	Wrong -p part number; use avrdude -p ? to list all; read signature with -U signature:r:-:h.
avrdude: verification error	Flash write failed; bad connection, insufficient power, or write-protected region.
avrdude: error: usbasp, Error: Could not find USB device	USBasp firmware too old for target; update USBasp firmware, or add -B 10 flag (slow SCK).
avrdude: initialization failed rc=-1; check connections and try again, or use -F	Target not responding; check power, RESET pin, JTAG enable fuse (JTAGEN=0 disables ISP on some devices).
Can't open device /dev/ttyUSB0 Permission denied	Add user to dialout group: sudo usermod -a -G dialout \$USER Then log out and back in.

26.8.1 USBasp Slow Clock

If ISP communication fails with a freshly-soldered board or a device that has been running at low speed (CKDIV8 fuse set):

```

# Slow down SCK to 125 kHz (-B = bitclock period in μs; -B 8 = 125kHz SCK)
avrdude -p m328p -c usbasp -B 8 \
        -U lfuse:w:0xFF:m # then reprogram fuses for 16MHz crystal

```

26.8.2 Arduino Auto-Reset Circuit

Arduino boards use a 100 nF capacitor on DTR/RST. avrdude -c arduino toggles DTR to reset the MCU before uploading. If the reset does not trigger:

```
# Manual: press reset button 2 seconds before avrdude starts, then release
avrdude -p m328p -c arduino -P /dev/ttyUSB0 -b 115200 \
-U flash:w:firmware.hex:i
```

Or disable the auto-reset capacitor (cut the RESET EN trace on the Arduino PCB) when running a custom bootloader that does not expect the timing.